

Advanced Computer Architectures - Notes -
v1.1.1

260236

February 2026

Preface

Every theory section in these notes has been taken from two sources:

- Computer Architecture: A Quantitative Approach. [2]
- Pipelining slides. [3]
- Course slides. [4]

About:

 [GitHub repository](#)



These notes are an unofficial resource and shouldn't replace the course material or any other book on advanced computer architectures. It is not made for commercial purposes. I've made the following notes to help me improve my knowledge and maybe it can be helpful for everyone.

As I have highlighted, a student should choose the teacher's material or a book on the topic. These notes can only be a helpful material.

For the exam, I created a simple formulary that follows the professor's instructions. You can access it via the link below or by scanning the QR code.



Contents

1 Pipelining	8
1.1 Basic Concepts	8
1.2 RISC-V Pipelining	13
1.2.1 Pipelined execution of instructions	16
1.2.2 Pipeline Implementation	18
1.3 Problem of Pipeline Hazards	20
1.3.1 RISC-V Optimized Pipeline	22
1.3.2 Solutions to RAW Hazards	25
1.4 Performance evaluation	27
2 Control Hazards and Branch Prediction	30
2.1 Conditional Branch Instructions	30
2.2 Control Hazards	32
2.3 Naïve Solutions to Control Hazards	34
2.4 Intro to Branch Prediction	37
2.5 Static Branch Prediction	38
2.5.1 Branch Always Not Taken	39
2.5.2 Branch Always Taken	41
2.5.3 Backward Taken Forward Not Taken (BTFNT)	44
2.5.4 Profile-Driven Prediction	45
2.5.5 Delayed Branch	46
2.5.5.1 From Before	48
2.5.5.2 From Target	49
2.5.5.3 From Fall-Through	51
2.5.5.4 From After	53
2.6 Dynamic Branch Prediction	54
2.6.1 1-bit Branch History Table	56
2.6.2 2-bit Branch History Table	59
2.6.3 Branch Target Buffer	61
2.6.4 Correlating Branch Predictors	65
2.6.4.1 (1,1) Correlating Predictors	67
2.6.4.2 (2,2) Correlating Predictors	69
2.6.5 Two-Level Adaptive Branch Predictors	71
3 Instruction Level Parallelism	74
3.1 The problem of dependencies	74
3.1.1 Data Dependencies	75
3.1.2 Name Dependencies	76
3.1.3 Control Dependencies	79
3.2 Multi-Cycle Pipelining	80
3.2.1 Multi-Cycle In-Order Pipeline	81
3.2.2 Multi-Cycle Out-of-Order Pipeline	84
3.3 Dynamic Scheduling	87
3.4 Multiple-Issue Processors	89
3.4.1 Introduction to Multiple-Issue Pipelines	89
3.4.2 Evolution Towards Superscalar Execution	92
3.4.3 Superscalar Processors	95
3.4.4 Static vs Dynamic Scheduling	98

3.5	ILP Limitations & Alternatives	101
3.6	Scoreboard: Dynamic Scheduling Algorithm	103
3.6.1	Assumptions and Architecture	103
3.6.2	Pipeline Stage Refinement	105
3.6.3	Hazard Management (RAW, WAR, WAW)	107
3.6.4	Control Logic and Stages	109
3.6.5	Summary	111
3.6.6	Scoreboard Data Structures	113
3.6.7	In-Depth Execution Example	115
3.7	Tomasulo's Algorithm	140
3.7.1	Introduction	140
3.7.2	Register Renaming: Static vs. Implicit	142
3.7.3	Basic Concepts of Tomasulo's Algorithm	144
3.7.4	Architecture	146
3.7.5	Stages	151
3.7.5.1	Stage 1: Issue	151
3.7.5.2	Stage 2: Start Execution	154
3.7.5.3	Stage 3: Write Result	157
3.7.6	In-Depth Execution Example	159
3.7.7	Tomasulo vs. Scoreboarding	183
3.7.8	Register Renaming	186
3.7.8.1	Introduction	186
3.7.8.2	Loop Unrolling and Code Scheduling	188
3.7.8.3	How Tomasulo Overlaps Iterations of Loops	193
3.7.8.4	Tomasulo Loop Execution	194
3.7.9	Reorder Buffer (ROB)	222
3.7.9.1	Hardware-based Speculation	222
3.7.9.2	Why ROB is really needed	223
3.7.9.3	ROB as a Data Communication Mechanism	226
3.7.9.4	Architecture	228
3.7.9.5	Speculative Tomasulo Algorithm with ROB	238
4	Performance Evaluation	246
4.1	Basic Concepts and Performance Metrics	246
4.2	Amdahl's Law	251
4.3	Pipelined Processors	253
4.4	Memory Hierarchy	257
5	VLIW (Very Long Instruction Word)	266
5.1	Introduction	266
5.2	Data dependencies	270
5.3	Statically Scheduled Processors	273
5.4	Code Scheduling	275
5.4.1	Scheduling Basics	276
5.4.2	Dependence Graph and Critical Path	277
5.4.3	ASAP Scheduling Algorithm (As Soon As Possible)	279
5.4.4	List-Based Scheduling Algorithm	281
5.4.5	Local vs Global Scheduling	283
5.4.6	Local Scheduling Techniques	285
5.4.6.1	Loop Unrolling	285

5.4.6.2	Software Pipelining	289
5.4.7	Global Scheduling	292
5.4.7.1	Trace Scheduling	292
5.4.7.2	Superblock Scheduling	294
6	Advanced Memory	297
6.1	Introduction	297
6.2	Principle of Locality	300
6.3	So, what is a Cache?	301
6.4	Cache Performance Metrics	303
6.5	Cache Architecture	304
6.5.1	Block Placement: <i>Where can a block be placed?</i>	305
6.5.2	Block Identification: <i>How is a block found?</i>	308
6.5.3	Replacement Strategy: <i>Which block should be replaced?</i>	309
6.5.4	Write Strategy: <i>What happens on a write?</i>	311
6.6	Miss Penalty Reduction	315
6.7	Design Space of Cache	317
6.8	Cache Miss Classification	319
6.9	Improving Cache Performance	321
6.9.1	Reducing Miss Rate Techniques	323
6.9.1.1	Increasing Cache Size	323
6.9.1.2	Increasing Block Size	324
6.9.1.3	Increasing Associativity	326
6.9.1.4	Victim Cache	328
6.9.1.5	Pseudo-Associativity & Way Prediction	330
6.9.1.6	Hardware Prefetching (Instructions & Data)	335
6.9.1.7	Software Prefetching (Compiler-Controlled)	338
6.9.1.8	Compiler Optimizations (for Cache Performance)	340
6.9.2	Reducing Miss Penalty Techniques	342
6.9.2.1	Read Priority over Write on Miss	342
6.9.2.2	Sub-block Placement	345
6.9.2.3	Early Restart & Critical Word First	346
6.9.2.4	Non-Blocking Caches	348
6.9.2.5	Second-Level and Multi-Level Caches	350
6.9.2.6	Merging Write Buffers	353
6.9.3	Reducing Hit Time Techniques	356
6.9.3.1	Small and Simple L1 Caches	356
6.9.3.2	Avoid address translation (VIPT caches)	357
6.9.3.3	Pipelined Writes	366
6.9.3.4	Small Sub-Blocks for Write-Through Caches	367
6.9.4	Summary	368
7	Multithreading (TLP)	371
7.1	Introduction	371
7.2	What is a Thread?	373
7.3	Multithreading Basics	376
7.4	Hardware Support for Multithreading	378
7.5	Typology of Hardware Multithreading	380
7.6	SMT on Wide Superscalars	383
7.7	Multicore + SMT	385

8 SIMD & Vector Architectures	387
8.1 Why DLP now? Limits of ILP	387
8.2 Flynn's taxonomy	388
8.3 SIMD Architecture	390
8.4 Vector Architectures vs Generic SIMD	394
8.5 Anatomy of a Classic Vector Machine	396
8.6 VMIPS as Didactic Reference Machine	399
8.7 Case Study: DAXPY	402
8.8 Chaining, Convoys, and Chimes	403
8.9 Multiple Lanes	405
8.10 Handling Vector Lengths	407
8.11 Control Divergence Inside Loops	409
8.12 Memory System for Vectors	411
8.13 SIMD ISA Extensions vs True Vectors	414
9 GPGPU Computing	416
10 Multiprocessors	417
10.1 What is a multiprocessor?	417
10.2 Key Design Questions	418
10.3 From single bus to interconnection networks	420
10.4 Network performance metrics	421
10.5 Classic Multiprocessor Topologies	422
10.6 Address-Space Models	425
10.7 Physical Memory Organization	427
10.8 Small-scale multicores	432
10.9 Distributed Shared Memory	434
10.10 Distributed-Memory Clusters	438
10.11 Communication/Programming Models	440
10.12 Why caches matter	444
10.13 Why Cache Coherence?	447
10.14 Snooping (bus-based) Protocols	449
10.14.1 Snoopy bus idea, broadcast, and scalability limits	449
10.14.2 Practicalities: duplicate tag to avoid stalling the core	450
10.14.3 Write-invalidate vs write-update (bandwidth vs latency)	452
10.14.4 MSI Protocol: Finite-State Machine (FSM)	453
10.14.5 MESI Protocol	458
10.15 Directory-Based Protocols	462
10.15.1 Basics	462
10.15.2 Message Types	465
10.15.3 Walkthroughs on a 4-node system	468
10.15.4 Directory FSM and Invariants	476
11 Exams	479
11.1 2026	479
11.1.1 January 8	479
11.2 2025	506
11.2.1 Midterm - May 5	506
11.2.2 February 11	542

Index

584

1 Pipelining

1.1 Basic Concepts

Pipelining is a fundamental technique in computer architecture aimed at improving instruction throughput by overlapping the execution of multiple instructions. The main idea behind pipelining is to divide the execution of an instruction into distinct stages and process different instructions simultaneously in these stages. This approach significantly increases the efficiency of instruction execution in modern processors.

❖ Understanding the RISC-V instruction set

Before delving into pipelining, it is essential to understand the basic instruction set of the RISC-V architecture. The instruction set consists of three major categories:

- **ALU Instructions (Arithmetic and Logic Operations)**

- Performs addition between registers:

```
1 add rd, rs1, rs2
```

Performs the addition between the values in registers `rs1` and `rs2` and stores the result in register `rd`.

$$rd \leftarrow rs1 + rs2$$

- Performs an addition between a constant and a register:

```
1 addi rd, rs1, 4
```

Performs the addition between the value in register `rs1` and the value `4` and stores the result in register `rd`.

$$rd \leftarrow rs1 + 4$$

- **Load/Store Instructions (Memory Operations)**

- Loads data from memory:

```
1 ld rd, offset(rs1)
```

Load data into register `rd` from an address formed by adding `rs1` to a signed `offset`.

$$rd \leftarrow M[rs1 + offset]$$

- Stores data in memory:

```
1 sd rs2, offset(rs1)
```

Store data from register `rs2` to an address formed by adding `rs1` to a signed `offset`.

$$M[rs1 + offset] \leftarrow rs2$$

- **Branching Instructions (Control Flow Management)**

- **Conditional Branches**

- * Branch on equal:

```
1 breq rs1, rs2, L1
```

Branch to the label L1 if the value in register **rs1** is equal to the value in register **rs2**.

$$\text{rs1} = \text{rs1} \xrightarrow{\text{go to}} \text{L1}$$

- * Branch on not equal:

```
1 bne rs1, rs2, L1
```

Branch to the label L1 if the value in register **rs1** is not equal to the value in register **rs2**.

$$\text{rs1} \neq \text{rs2} \xrightarrow{\text{go to}} \text{L1}$$

- **Unconditional Jumps**

- * Jump to the label (jump):

```
1 j L1
```

Jump directly to the L1 label.

- * Jump to the address stored in a register (jump register):

```
1 jr ra
```

Take the value in register **ra** and use it as the address to jump to. So it is assumed that **ra** contains an address.

These basic instructions will be used throughout the course.

☰ Execution phases in RISC-V

1. **IF (Instruction Fetch)**: The instruction is **fetched** from memory.
2. **ID (Instruction Decode)**: The instruction is **decoded**, and the **required registers** are read.
3. **EX (Execution)**: The instruction is **executed**, typically involving ALU operations.
4. **ME (Memory Access)**: For *load/store* instructions, this stage **reads from** or **writes to** memory.
5. **WB (Write Back)**: The **result is written back** to the destination register.

These five stages form the basis of the RISC-V pipeline.

Implementation of the RISC-V Data Path

The **RISC-V Data Path** is a fundamental component of the processor's architecture, responsible for **executing instructions efficiently by coordinating various hardware units**. It defines how instructions flow through different stages of execution, interacting with memory, registers, and the Arithmetic Logic Unit (ALU).

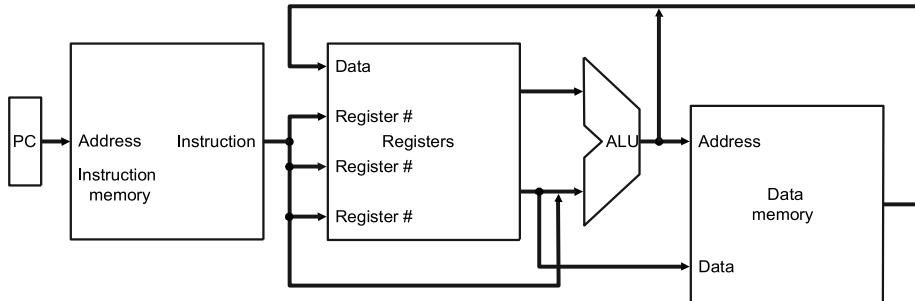


Figure 1: Generic implementation of the RISC-V Data Path.

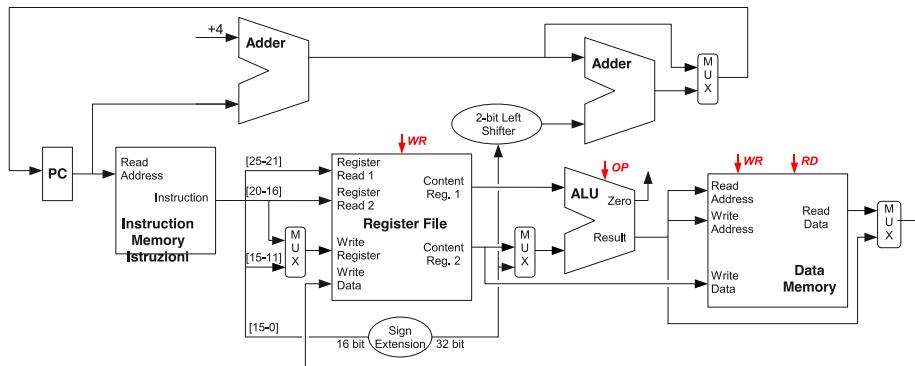


Figure 2: Specific implementation of the RISC-V Data Path.

Its fundamental components include:

- **Instruction Memory and Data Memory Separation.** RISC-V adopts a Harvard Architecture style, where the **Instruction Memory (IM)** and **Data Memory (DM)** are separate. This **prevents structural hazards** where instruction fetch and memory access could conflict in a single-memory design (this topic will be addressed later).
- **General-Purpose Register File (RF).** It consists of **32 registers**, each **32-bit wide**. The register file has **two read ports** and **one write port** to support simultaneous read and write operations. This setup allows faster register access, which is crucial for pipelined execution.
- **Program Counter (PC).** It holds the address of the next instruction to be fetched. Automatically increments during execution, typically by 4 bytes (for 32-bit instructions).

- **Arithmetic Logic Unit (ALU)**. Performs arithmetic and logical operations required by instructions. Inputs to the ALU come from registers or immediate values decoded from the instruction.

Other components that we can see in the general implementation of the RISC-V data path are:

- **Register File**. Stores temporary values used by instructions. Contains read ports (two registers can be read simultaneously for ALU operations) and write port (one register can be updated per clock cycle). The register file ensures high-speed execution of operations by reducing memory accesses.
- **Instruction Fetch (IF)**. The PC (Program Counter) retrieves the next instruction from Instruction Memory. The PC is incremented using an adder ($PC + 4$), ensuring sequential instruction flow.
- **Instruction Decode (ID)**. Extracts opcode (determines the instruction type), source and destination registers, immediate values (if present). It reads values from the Register File based on instruction requirements.
- **Execution (EX)**. The ALU performs arithmetic and logical operations. A multiplexer (MUX) selects the second operand: a register value (for R-type instructions) or an immediate value (for I-type instructions like addi). The ALU result is forwarded to the next stage.
- **Memory Access (ME)**. Load (ld) and Store (sd) instructions interact with data memory. Data is either loaded from memory into a register or stored from a register into memory.
- **Write Back (WB)**. The result from ALU or memory is written back to the Register File.

Example 1: Data Path Execution Example

Let's consider a simple RISC-V **load instruction** (`ld x10, 40(x1)`) passing through the data path:

1. **IF Stage:** Instruction Fetch
 - PC → Instruction Memory → `ld x10, 40(x1)` fetched
 - PC updated to PC + 4
2. **ID Stage:** Instruction Decode
 - Registers read: `x1` (base register for memory access)
 - Immediate value extracted: 40
3. **EX Stage:** Execution
 - ALU calculates memory address: `x1 + 40`
4. **ME Stage:** Memory Access
 - Data is loaded from $M[x_1 + 40]$
5. **WB Stage:** Write Back
 - Data stored in `x10`

1.2 RISC-V Pipelining

Pipelining is analogous to an assembly line in a factory. Instead of waiting for one instruction to complete before starting the next, **different instructions are executed simultaneously in different stages**.

If we consider a **non-pipelined execution**:

- Each instruction completes all five stages sequentially before the next instruction starts.
- If each instruction stage (IF stage, ID stage, etc.) takes, say, 2 nanoseconds, executing all stages of an instruction (IF, ID, EX, MEM, WB) takes 5 times 2 nanoseconds, then 10 nanoseconds. If we also want to execute 5 instructions, we need 10 nanoseconds times 5, then 50 nanoseconds!

Now, we consider a **pipelined execution**:

- Once the first instruction moves to the second stage, the next instruction starts in the first stage.
- The **pipeline becomes fully utilized** after the first few cycles, significantly **improving throughput**.

In an **ideal scenario**, a 5-stage pipeline should provide a speedup of $5 \times$ reducing execution time to:

$$(5 + 4) \times 2 \text{ ns} = 18 \text{ ns}$$

Where 5 are the steps of the first instruction, 5 are the steps of the last instruction, minus 1 because one step is already counted in the first instruction, so 4. Therefore, 9 is multiplied by 2 nanoseconds, the time taken by each stage. The result, 18 nanoseconds, is the time it takes the pipeline to execute 5 instructions in an ideal scenario.

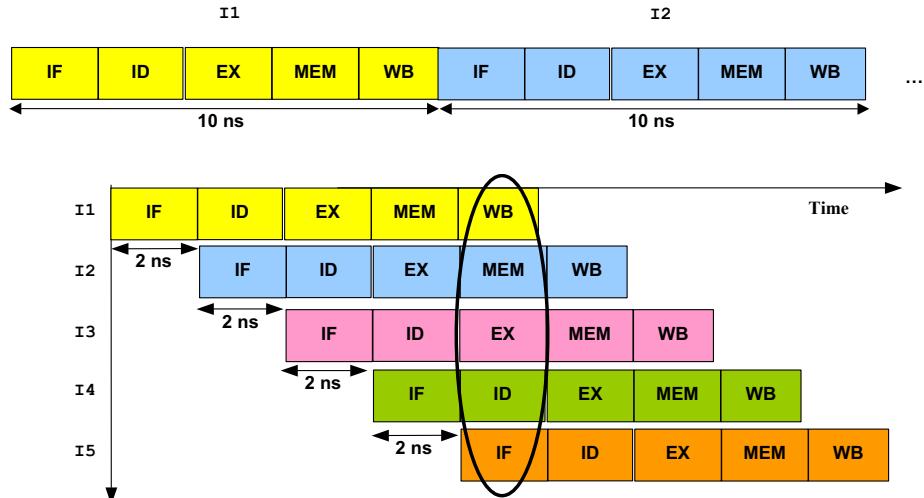


Figure 3: Sequential vs. Pipelining execution.

⌚ Pipeline Performance and Speedup

The ideal performance improvement from pipelining is derived from the fact that **once the pipeline is filled, a new instruction completes every cycle**. The key performance metrics include:

- **Latency (Execution Time)**: The total time to complete a single instruction does not change (sequential or pipeline).
- **Throughput (Instructions per Unit Time)**: The number of completed instruction per unit time significantly increases.
- **Speedup Calculation**
 - A non-pipelined CPU with 5 execution cycles of 2 ns would take 10 ns per instruction.
 - A pipelined CPU with 5 stages of 2 ns results in 1 instruction completing every 2 ns.
 - This gives a theoretical speedup of $5 \times$ (ideal case).

Unfortunately, real-world implementations are subject to **pipeline hazards** that reduce efficiency.

Understanding Pipelining Performance

Pipelining **improves instruction throughput** by allowing multiple instructions to be processed simultaneously in different stages. The **execution of an instruction is divided into 5 pipeline stages**:

1. IF (Instruction Fetch)
2. ID (Instruction Decode)
3. EX (Execution)
4. MEM (Memory Access)
5. WB (Write Back)

Each stage takes 2 ns (a *pipeline cycle*), meaning that an **instruction moves from one stage to the next every 2 ns**. Now, let's analyze the timeline of instruction execution:

Clock Cycle	IF	ID	EX	MEM	WB
1st (0-2 ns)	I1				
2nd (2-4 ns)	I2	I1			
3rd (4-6 ns)	I3	I2	I1		
4th (6-8 ns)	I4	I3	I2	I1	
5th (8-10 ns)	I5	I4	I3	I2	I1
6th (10-12 ns)	I6	I5	I4	I3	I2
7th (12-14 ns)	I7	I6	I5	I4	I3
8th (14-16 ns)	I8	I7	I6	I5	I4
9th (16-18 ns)	I9	I8	I7	I6	I5

Table 1: Pipelining timeline execution in an ideal case.

- The first instruction I1 takes 5 (clock) cycles to complete, i.e., 10 ns.
- However, starting from cycle 5, a new instruction finishes every cycle (every 2 ns).
- In a non-pipelined system, each instruction would take 10 ns (5 stages × 2 ns each).
- In a pipelined system, once the pipeline is full, an instruction completes every cycle (every 2 ns), achieving a 5× speedup compared to the non-pipelined execution.

Thus, after an initial “fill” time (1st, 2nd, 3rd, 4th), **a new instruction completes every 2 ns** (from 5th to 6th, I1 is finished; from 6th to 7th, I2 is finished; from 7th to 8th, I3 is finished), which is the duration of a single pipeline stage.

1.2.1 Pipelined execution of instructions

Each RISC-V instruction follows the five pipeline stages, but their interactions with the pipeline vary depending on the instruction type.

- **ALU Instructions** (e.g., `op $x, $y, $z`)

These are register-based operations that do not require memory access. Since there is no memory operation, the instruction **bypasses the ME stage**.

Stage	Description
IF	Fetch instruction from memory
ID	Decode instruction, read registers $\$y$ and $\$z$
EX	Perform ALU operation ($\$x = \$y + \$z$)
ME	No memory access (skipped)
WB	Write the ALU result to $\$x$

- **Load Instructions** (e.g., `lw $x, offset($y)`)

These instructions retrieve data from memory and store it in a register. The **memory access stage (ME) is crucial** here since the instruction must fetch data from memory.

Stage	Description
IF	Fetch instruction from memory
ID	Decode instruction, read base register $\$y$
EX	Compute memory address ($\$y + \text{offset}$)
ME	Read data from memory
WB	Write data into destination register $\$x$

- **Store Instructions** (e.g., `sw $x, offset($y)`)

These instructions write data from a register into memory. Unlike `lw`, **store instructions do not require the WB stage**, as data is written directly into memory.

Stage	Description
IF	Fetch instruction from memory
ID	Decode instruction, read base register $\$y$ and source register $\$x$
EX	Compute memory address ($\$y + \text{offset}$)
ME	Write $\$x$ into memory at the computed address
WB	No write-back stage (skipped)

- **Conditional Branches** (e.g., `beq $x, $y, offset`)

Branching introduces control hazards, as the pipeline needs to determine whether the branch is taken or not. Branches can introduce **stalls** due to dependencies on comparison results. This issue is typically mitigated using branch prediction.

Stage	Description
IF	Fetch instruction from memory
ID	Decode instruction, read registers $\$x$ and $\$y$
EX	Compare $\$x$ and $\$y$, compute target address
ME	No memory access (skipped)
WB	Update PC if branch is taken

This section breaks down how **different types of instructions behave in the pipeline**:

- ALU Instructions complete in the EX stage and do not use memory.
- Load Instructions require a memory access in the ME stage.
- Store Instructions write to memory instead of registers.
- Branch Instructions introduce control hazards because they may change the PC.

This means that **not all instructions behave the same** in the pipeline. Some instructions **skip certain stages** (e.g., stores do not have WB, ALU instructions skip ME), and some instructions **introduce potential problems** (e.g., branches can cause delays).

In conclusion, this section sets the stage for understanding pipeline stalls, forwarding, and hazard resolution techniques that are essential for designing high-performance processors.

1.2.2 Pipeline Implementation

The **RISC-V pipeline implementation** is designed to efficiently execute multiple instructions simultaneously, following the classical five-stage pipeline model:

1. IF (Instruction Fetch)
2. ID (Instruction Decode)
3. EX (Execution)
4. MEM (Memory Access)
5. WB (Write Back)

Each clock cycle, a new instruction enters the pipeline while previous instructions move to the next stage, allowing **five different instructions to be in execution at the same time**.

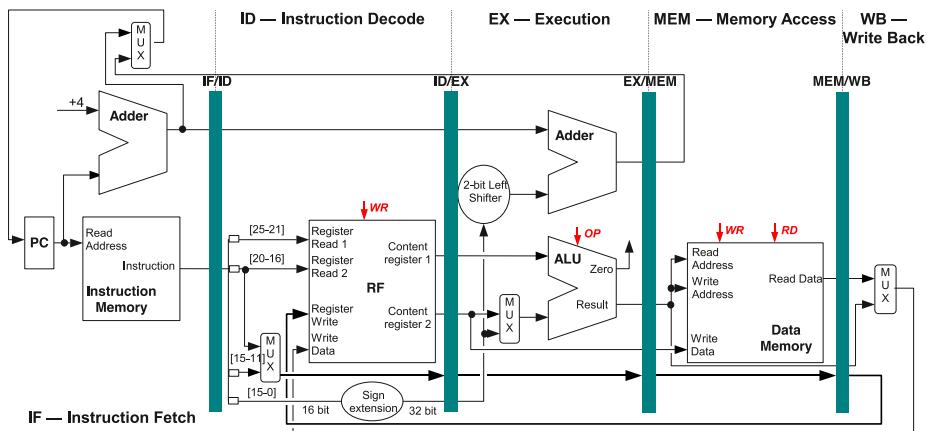


Figure 4: Structure of RISC-V pipeline.

❖ Execution Stages and Pipeline Modules

Each stage of the pipeline corresponds to a specific hardware module in the CPU. The RISC-V pipeline is composed of five primary hardware modules:

- **Instruction Fetch (IF) Module:** Fetches instructions from instruction memory and updates the PC.
- **Instruction Decode (ID) Module:** Decodes the fetched instruction and reads register values.
- **Execution (EX) Module:** Performs arithmetic/logical operations in the ALU or computes memory addresses.
- **Memory Access (MEM) Module:** Reads from or writes data to memory.

- **Write Back (WB) Module:** Writes the computed result back into the register file.

Each module is responsible for a specific **stage of execution**, and together they allow overlapping execution of multiple instructions.

Pipeline Registers

To maintain separation between stages, **pipeline registers** are used (see Figure 4, page 18). These registers store intermediate results and ensure proper communication between stages:

- **IF/ID Register:** Holds fetched instruction and updated PC.
- **ID/EX Register:** Stores decoded instruction, read register values, and control signals.
- **EX/MEM Register:** Holds ALU results, destination register, and memory access information.
- **MEM/WB Register:** Stores memory data or ALU result to be written back to registers.

These pipeline registers eliminate the need for re-fetching or re-decoding instructions at each cycle, thus maintaining pipeline efficiency.

1.3 Problem of Pipeline Hazards

⚠ Assumptions Made

Until now, our discussion on the RISC-V pipeline implementation has relied on several key assumptions to simplify the analysis and focus on fundamental concepts. These **assumptions help in understanding the ideal case of pipelining** before introducing complexities like hazards and optimizations.

1. All instructions are independent, so there are no dependencies between them.
2. No branches or jumps that change execution flow.

This is a theoretical idealization, because in real-world scenarios, **hazards** (structural, data, and control) **interfere with smooth execution**. Also, our second assumption ignores **branch instructions** (`beq`, `bne`, `j`, `jr`), which **cause control hazards** that require branch prediction or pipeline flushing.

❓ What is a Pipeline Hazard?

Now that we have understood the ideal execution of a RISC-V pipeline, we must discuss pipeline hazards, which are obstacles that prevent the pipeline from operating at maximum efficiency.

A **Hazard** (or conflict) is a phenomenon that occurs when the **overlapping execution of instructions in the pipeline changes the expected order of instruction execution**. This can lead to incorrect results or the **need to insert stalls** (*pipeline bubbles*), reducing performance.

In other words, **hazards cause the next instruction in the pipeline to be delayed, which reduces the ideal throughput of 1 instruction per cycle**. Thus, hazards disrupt the smooth flow of instructions and require techniques to resolve them.

☰ Classes of Pipeline Hazards

- **Structural Hazards:** Attempt to use the same resource from different instructions simultaneously.
 ❓ **Example:** Single memory for both instruction and data access.
- **Data Hazards:** Attempt to use a result before it is ready.
 ❓ **Example:** Instruction depending on a result of a previous instruction still in the pipeline.
- **Control Hazards:** Try to make a decision about the next statement to execute before the condition is evaluated.
 ❓ **Example:** Conditional branch execution.

✓ Structural Hazards

A **structural hazard** occurs when multiple pipeline stages need to use the same hardware resource at the same time.

✓ **Structural Hazard cannot be applied to RISC-V.** This is a great thing, because thanks to the Harvard Architecture, **RISC-V uses separate instruction and data memory**, and this adoption avoids structural hazards.

⌚ Control Hazards

A **control hazard** (section 2, page 30) occurs when the **pipeline does not know which instruction to fetch next, usually due to a branch or jump instruction**. It is discussed in the following sections.

⌚ Data Hazards

A **data hazard** (section 3.1, page 74) occurs when an **instruction depends on the result of a previous instruction that is still in the pipeline**.

There are several types of data hazards:

- **RAW (Read After Write).** An instruction tries to read a register before a previous instruction writes to it.

⌚ Example:

```
1 lw x2, 0(x1)
2 add x3, x2, x4
```

The **add** instruction needs **x2**, but **x2** is still being fetched from memory in the **MEM** stage. Without hazard resolution, the processor would get the wrong value for **x2**.

- **WAR (Write After Read).** A later instruction writes to a register before an earlier instruction reads it (rare in RISC).
- **WAW (Write After Write).** Two instructions try to write to the same register in the wrong order.

1.3.1 RISC-V Optimized Pipeline

The **RISC-V optimized pipeline** introduces refinements that **reduce stalls, improve data access, and enhance instruction throughput**. The key optimizations include:

- ✓ **Efficient Register File Access.** In the standard RISC-V pipeline, register accesses **happen in two stages**:
 - ID (Instruction Decode) → Reads register values.
 - WB (Write Back) → Writes computed values back to registers.

⌚ Optimization: Read and Write in the Same Cycle

- In the optimized pipeline:
 - Register **writing** happens in the **first half** of the **clock cycle**;
 - While register **reading** happens in the **second half** of the **clock cycle**.
- This means an instruction can write its result to a register in WB, and the **next instruction can immediately read** that value in ID during the **same cycle**.

This optimization **removes unnecessary stalls** when an instruction immediately depends on a result written in the previous cycle.

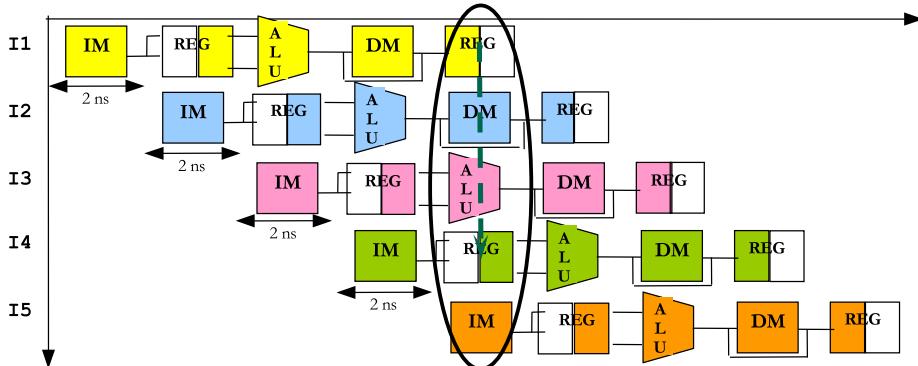


Figure 5: Visual **example** of an optimized pipeline; here the result (WB stage) of I1 is written in the first half of the clock cycle and the read (ID stage) of I4 is done in the second half of the clock cycle. So there is no hazards!

- ✓ **Forwarding (Bypassing) to Reduce Stalls.** **Forwarding** (also called **bypassing**) is a hardware technique that **eliminates stalls by providing ALU results directly to dependent instructions without waiting for the WB stage**. It is a possible solution for Data Hazards.

- ⌚ **Forwarding Paths:** To support forwarding, the **pipeline includes extra paths** that allow instructions to fetch values from intermediate pipeline registers instead of waiting for WB.

- **EX/EX Path.** Allows ALU results to be forwarded from **EX stage output** to the next **EX stage input**. Used when an **instruction depends on an arithmetic result of the previous instruction**.

Example 2: EX/EX Forwarding

```
1 sub x2, x1, x3    # Compute x2 = x1 - x3
2 and x12, x2, x5  # Use x2 immediately
```

Cycle	sub x2, x1, x3	and x12, x2, x5
1	IF	
2	ID	IF
3	EX	ID
4	MEM	<i>Stall</i>
5	WB	<i>Stall</i>
6		EX
7		MEM
8		WB

The **and** instruction **must wait until WB writes x2 to the register file**. Two stall cycles are introduced and this wastes execution time.

Instead of waiting for WB, we forward the ALU result from the EX stage of **sub** directly to the EX stage of **and**.

Cycle	sub x2, x1, x3	and x12, x2, x5
1	IF	
2	ID	IF
3	EX	ID
4	MEM	EX (forwarded x2 from EX)
5	WB	MEM
6		WB

In cycle 4, **and x12, x2, x5** gets the forwarded **x2** from the **EX stage of sub, removing stalls**.

This is **EX/EX forwarding**, taking ALU results from one EX stage directly into the next EX stage.

- **MEM/EX Path.** Forwards the ALU result from **MEM** stage to **EX** stage. Used when an **instruction depends on an ALU operation two cycles before**.

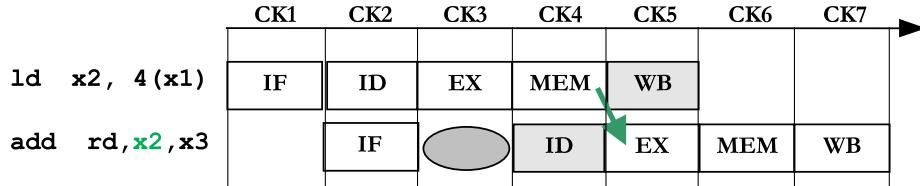


Figure 6: Example of MEM/EX path.

- **MEM/MEM Path.** Forwarding directly between two memory operations in the **MEM** stage. It removes stalls in Load/Store dependencies.

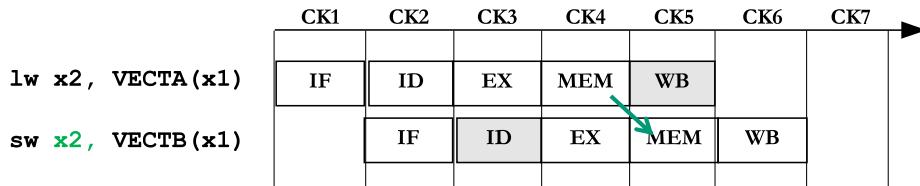


Figure 7: Example of MEM/MEM path.

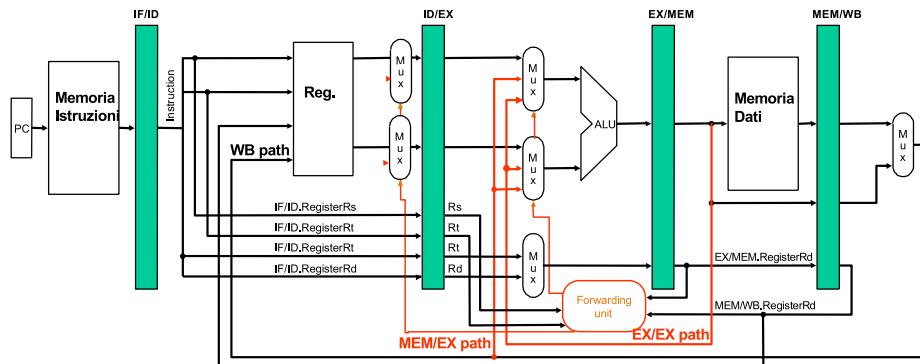


Figure 8: Implementation of RISC-V with Forwarding Unit.

1.3.2 Solutions to RAW Hazards

To handle RAW hazards, we can **use both static (compile-time) and dynamic (hardware-based) techniques**. These include:

- **Static (compile-time):**

- ✓ **nop insertion:** compiler adds empty instructions to delay execution.
- ✓ **Instruction Scheduling:** compiler reorders instructions to avoid conflicts.

- **Dynamic (hardware-based):**

- ✓ **Pipeline Stalling (bubbles):** inserts delay cycles when necessary.
- ✓ **Forwarding (bypassing):** uses intermediate values from the pipeline instead of waiting.

✓ **Static (compile-time) solution: Inserting nops (naïve)**

One simple way to handle RAW hazards is to **insert nop instructions manually between dependent instructions**. This gives the pipeline time to complete the write-back of the needed value.

Key takeaway of inserting nops:

- ✗ **Simple, but inefficient because it wastes clock cycles.** It should be the very last solution considered.
- ✗ Instead of using useful instructions, the **processor waits**, reducing performance.

Example 3: nop insertion

```

1 sub x2, x1, x3
2 nop          # Delay slot (bubble)
3 and x12, x2, x5 # Now x2 is ready

```

✓ **Static (compile-time) solution: Instruction Scheduling**

A more efficient technique is **instruction reordering**, also known as **compiler scheduling**. The **compiler reorders instructions to avoid data hazards without inserting nops**.

Key takeaway of instruction scheduling:

- Instruction reordering is a **compiler optimization**.
- ✓ It works well if **independent instructions are available**.
- ✗ In some cases, no independent instructions exist, so **stalling or forwarding is needed**.

Example 4: Instruction Scheduling

```

1 sub x2, x1, x3
2 # Independent instruction
3 # (can execute while sub is completing)
4 add x4, x10, x11
5 and x12, x2, x5    # Now x2 is ready

```

Instead of a nop, we insert add x4, x10, x11, which does not depend on x2. This keeps the pipeline utilized while avoiding RAW hazards.

✓ Dynamic (hardware-based): Pipeline Stalling (Bubble Insertion)

When no independent instructions can be scheduled, the **hardware must stall the pipeline** by inserting a **bubble (stall cycle)**.

Key takeaway of pipeline stalling:

- ✗ Stalling is simple but **reduces performance** (pipeline sits idle).
- ✓ We prefer **forwarding** (next solution) instead of stalling.

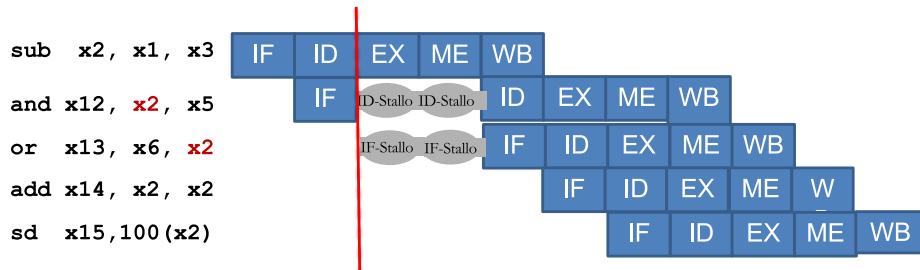


Figure 9: Example of inserting stalls.

✓ Dynamic (hardware-based): Forwarding (Bypassing)

Forwarding is an optimized hardware technique that avoids pipeline stalls by **directly passing results between pipeline registers**. The entire implementation has already been explained on page 22.

Key takeaway of forwarding:

- ✓ **Forwarding is the best solution** because it eliminates stalls and maximizes performance.
- It requires extra hardware (MUX and control logic), but it significantly improves throughput.

1.4 Performance evaluation

Evaluating the performance of a pipelined processor is essential to understanding the impact of stalls, hazards, and instruction throughput. In an **ideal scenario**, a **pipeline achieves one instruction per cycle** ($CPI = 1$), but real-world execution includes pipeline stalls, which degrade performance.

Several **key metrics** are used to evaluate the efficiency of pipelining:

- **Instruction Count (IC)**. Represents the **total number of instructions executed**. Used as a basis for performance calculations.
- **Clocks Per Instruction (CPI)**. CPI measures the **average number of clock cycles required to execute one instruction**. Ideal CPI for a pipelined processor is **1**, but hazards and stalls increase CPI.

$$CPI = \frac{\text{Total Clock Cycles}}{\text{Instruction Count (IC)}} \quad (1)$$

Where the total clock cycles is:

$$\text{Total Clock Cycles} = IC + 4 + \text{Stall Cycles} \quad (2)$$

Where $+4$ is the **fill time of the first instruction**. The $+4$ represents the initial pipeline fill time required before the pipeline reaches full execution throughput.

Example 5: Why is the Pipeline Startup Overhead $+4$?

A 5-stage pipeline (IF, ID, EX, MEM, WB) requires 4 extra cycles before the first instruction completes. Consider the following scenario:

Clock Cycle	IF	ID	EX	MEM	WB
1	I1				
2	I2	I1			
3	I3	I2	I1		
4	I4	I3	I2	I1	
5	I5	I4	I3	I2	I1
6	I6	I5	I4	I3	I2
7	I7	I6	I5	I4	I3

The first instruction (I1) requires 5 cycles to complete. The next instruction (I2) completes in cycle 6, and so on. After the first 5 cycles, the **pipeline reaches steady state**, completing 1 instruction per cycle (ideal scenario, no hazards).

- **Instruction Per Clock (IPC)**. IPC is the inverse of CPI:

$$IPC = \frac{1}{CPI} \quad (3)$$

Measures **how many instructions complete per clock cycle**.

- **Millions of Instructions Per Second (MIPS)**. Evaluates processor speed in terms of **millions of instructions executed per second**:

$$\text{MIPS} = \frac{f_{\text{clock}}}{\text{CPI} \times 10^6} \quad (4)$$

Higher **clock frequency** (f_{clock}) and lower CPI result in better MIPS.

Example 6: Performance Calculation

Given:

- Instruction Count (IC) = 5
- Stall Cycles = 2
- Clock Frequency = 500MHz

Metrics:

- Total Clock Cycles:

$$\text{Clock Cycles} = \text{IC} + \text{Stall Cycles} + 4 = 5 + 2 + 4 = 11$$

- CPI Calculation:

$$\text{CPI} = \frac{11}{5} = 2.2$$

- MIPS Calculation:

$$\text{MIPS} = \frac{500 \text{ MHz}}{2.2 \times 10^6} = 227$$

Without stalls, CPI would be 1 (ideal pipeline). But stalls increase CPI, reducing MIPS and overall efficiency.

⌚ Performance in Loops and Asymptotic Analysis

When evaluating **loops** or **long-running programs**, we use asymptotic performance metrics.

For a loop with:

- **m instructions** per iteration.
- **k stall cycles** per iteration.
- **n iterations**.

We have:

- **Clock Cycles per Iteration:**

$$\text{Clock Cycles per Iteration} = m + k + 4 \quad (5)$$

- **CPI per Iteration:**

$$\text{CPI}_{\text{iter}} = \frac{(m + k + 4)}{m} \quad (6)$$

- **MIPS per Iteration:**

$$\text{MIPS}_{\text{iter}} = \frac{f_{\text{clock}}}{\text{CPI}_{\text{iter}} \times 10^6} \quad (7)$$

For **large n** , the impact of pipeline startup delay (+4 cycles) is reduced:

- **CPI per Iteration:**

$$\begin{aligned} \text{CPI}_{\text{AS}} &= \lim_{n \rightarrow \infty} \frac{(\text{IC}_{\text{AS}} + \text{Stall Cycles}_{\text{AS}} + 4)}{\text{IC}_{\text{AS}}} \\ &= \lim_{n \rightarrow \infty} \frac{(m \times n + k \times n + 4)}{(m \times n)} \\ &= \frac{(m + k)}{m} \end{aligned} \quad (8)$$

- **Millions of Instructions Per Second (MIPS):**

$$\text{MIPS}_{\text{AS}} = \frac{f_{\text{clock}}}{\text{CPI}_{\text{AS}} \times 10^6} \quad (9)$$

For **large programs**, startup stalls become negligible, and **performance depends mainly on stall cycles per iteration**. **Minimizing k (stalls per iteration)** is crucial to achieving high efficiency.

⌚ Why CPI is Greater than 1 in Real Pipelines

Even with an **optimized pipeline**, **real-world execution is affected by hazards**. Thus, actual CPI is always greater than 1, even in well-optimized designs.

2 Control Hazards and Branch Prediction

2.1 Conditional Branch Instructions

In pipelined processor architectures, control flow is not always linear, and decisions about the next instruction to execute are often dependent on certain conditions. This introduces the necessity for **conditional branch instructions**, particularly relevant in RISC-V architectures, where typical instructions include:

- **beq** (branch if equal): `beq rs1, rs2, L1`

Transfers execution to the label L1 if the contents of registers `rs1` and `rs2` are equal.

- **bne** (branch if not equal): `bne rs1, rs2, L1`

Transfers control to L1 if `rs1` and `rs2` hold different values.

These branch instructions are essential in implementing control structures such as loops, conditionals (`if/else`), and function returns.

At the hardware level, the **Branch Target Address (BTA)** plays a central role. This address represents where the processor should continue execution if the branch is taken (i.e., if the condition specified by the branch instruction evaluates as true). When the condition is satisfied, the processor updates the **Program Counter (PC)** with the BTA, thus redirecting the flow of instruction fetch.

Conversely, if the condition is not satisfied, the branch is not taken, and the processor continues sequential execution. In RISC-V, since instructions are generally 32 bits (4 bytes) long, the next instruction is fetched from `PC + 4`.

Understanding whether a branch is taken or not is crucial for instruction fetch in pipelined architectures. **Mispredicting** this can introduce **performance penalties**, which are addressed in detail through the study of control hazards and branch prediction techniques in later sections.

Execution of Branches in Pipelined Architectures

When executing a branch instruction, such as `beq rx, ry, L1`, the processor must compare two registers (`rx, ry`) to determine whether the branch should be taken (i.e., jump to label L1) or not taken (continue sequentially). In RISC-V, the **Branch Outcome** (Taken/Not Taken) and **Branch Target Address (BTA)** are calculated during the **EX stage** (when the ALU performs arithmetic and logical operations):

- EX Stage:
 - Compare `rx` and `ry` using the ALU.
 - Compute `PC + offset` to obtain the BTA.

- ME Stage:
 - Based on the comparison, update the PC to either $PC + 4$ (if not taken) or $PC + offset$ (if taken).

MIPS follows a similar structure but emphasizes that **branch decisions are finalized at the end of the EX stage**, with the **PC update happening at the ME stage**. This introduces a **delay in resolving the branch**, which becomes critical for understanding control hazards.

⚠ Implication for IF

Since new instructions are fetched every clock cycle, the **processor needs to decide early which instruction to fetch next**. However, with branches, this decision is **not immediately clear because the branch condition hasn't yet been evaluated**. The Program Counter (PC) **cannot be updated correctly until the branch outcome is known**, leading to potential pipeline stalls or incorrect instruction fetches.

2.2 Control Hazards

In a pipelined architecture, one of the primary challenges in achieving high performance is dealing with **Control Hazards**, also known as **branch hazards**. These arise due to the presence of conditional branch instructions, where the processor must decide the next instruction to fetch before knowing whether the branch will be taken.

② What really causes Control Hazards?

To sustain the pipeline and **avoid idle stages**, a processor needs to **fetch one instruction per clock cycle**. However, with branch instructions, this becomes problematic because the **branch decision** (branch outcome) and the **branch target address (BTA)** are **not immediately available**. Specifically, **during the IF stage**, when the next instruction is fetched, the **processor still does not know whether the branch will be taken or not**, because this information typically becomes available later in the pipeline.

This leads to uncertainty:

- ③ *Which instruction should be fetched after a branch?*
- ③ *Should it be sequential instruction ($PC + 4$) or the instruction at the BTA?*

If the processor fetches the wrong instruction, it might **need to discard** or “**flush**” it later, **wasting valuable cycles**. Alternatively, the **processor may stall**, delaying the fetching of any instruction until the branch decision is known, which also hurts performance.

The key issue is this: in a pipeline, the **instruction stream needs to continue**, but the **correct path is unclear** until the branch condition is evaluated. Thus:

1. Either **wrong-path instructions are fetched** (requiring flushing later)
2. Or **pipeline stalls** are introduced (causing delay and loss of ideal speedup)

Key Takeaways: Control Hazards

- **Definition:** Pipeline hazard due to **uncertainty in branch outcome** during instruction fetch.
- **Cause:** The **branch condition is unresolved** when the next instruction must be fetched (IF stage).
- **Instructions Involved:** Conditional branches (`beq`, `bne`) and jumps, all **instructions modifying the PC**.
- **Pipeline Timing Conflict:** BO and BTA known only in EX or later, but instruction fetch **must occur every cycle**.

- **Main Problem:** Cannot decide whether to fetch **next sequential instruction** ($PC + 4$) or **BTA**.
- **Possible Outcomes:**
 - **Stall:** Delay fetch until branch resolved.
 - **Fetch wrong instruction** → flush.
- **Performance Impact:** Loss of **ideal pipelining speedup**; reduced throughput due to stalls or wasted fetches.
- **Goal of Solutions:** Mitigate stalls and improve fetch accuracy through early evaluation or prediction.

2.3 Naïve Solutions to Control Hazards

To manage control hazards in a simple and reliable way, one of the **earliest approaches** developed was the **conservative solution** of introducing **branch stalls**. The idea is straightforward: when a branch instruction enters the pipeline, the **processor stalls the pipeline until the branch decision is known and the correct next instruction can be safely fetched**.

⌚ How does the conservative solution work?

In the typical 5-stage pipeline, the Branch Outcome (i.e., whether the branch is taken or not) and the Branch Target Address (BTA) are usually resolved **at the end of the EX stage**. However, the **Program Counter (PC)** is actually **updated at the end of the ME stage**. This introduces a **delay of multiple cycles** between the branch instruction entering the pipeline and the point when the next instruction can be fetched with certainty.

To avoid fetching an incorrect instruction, the **processor simply pauses instruction fetch** for **3 clock cycles** after the branch instruction enters the pipeline. These are called **stalls**, essentially empty cycles where no new instructions enter the pipeline. Once the PC is correctly updated based on the branch outcome, instruction fetch resumes.

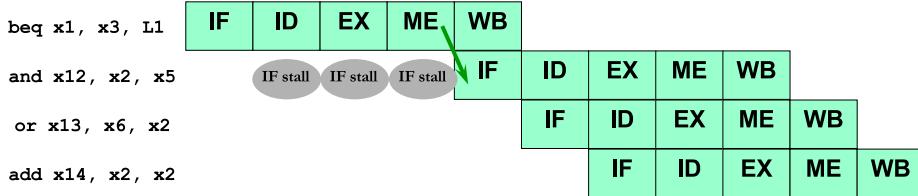


Figure 10: Example of stalls inserted in the pipeline to read the correct value after a branch condition.

⚠️ Performance Impact

It's pretty obvious that if **each branch introduces a penalty of 3 cycles**, it will **significantly degrade performance**, especially in programs with frequent branching. Since pipelining aims to maximize instruction throughput, **this solution sacrifices speed for correctness**. In fact, it is **called conservative because it does not attempt to guess or speculate about the branch outcome**. Instead, it **waits for certainty, favoring reliability over efficiency**.

⌚ Can optimized evaluation at the ID stage improve performance?

Although the conservative solution degrades performance, it can be relatively optimized thanks to hardware optimization. Processor designers have introduced **hardware optimizations** that allow the **branch outcome (BO)** and the **branch target address (BTA)** to be **computed earlier** in the pipeline. Specifically, these computations can be **moved from the EX stage to the**

ID stage (during the decoding phase). This optimization is often referred to as **Early Branch Evaluation**.

❖ How Early Branch Evaluation Works

To achieve this, the **Instruction Decode (ID)** stage must be **enhanced with additional hardware logic** that allows it to:

1. **Compare register values** (rx and ry) to determine the branch condition.
2. **Compute the BTA** using the **sign-extended offset** from the instruction and the current PC value.
3. **Update the PC** as soon as BO and BTA are known.

By doing this, **both BO and BTA are available at the end of the ID stage**, allowing the processor to update the PC immediately and fetch the correct instruction in the following cycle.

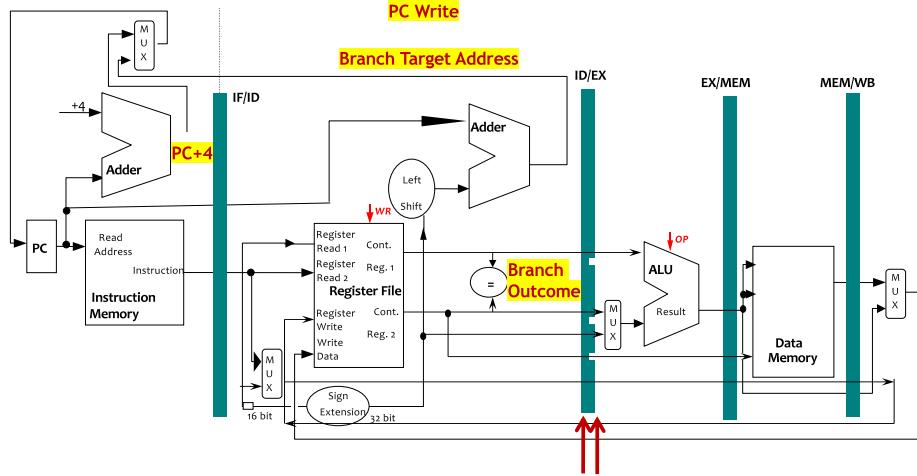


Figure 11: Hardware features modifications to allow early branch evaluation.

✖ Hardware Overhead: Complexity vs. Performance

This optimization requires **more complex hardware**, as the ID stage now includes:

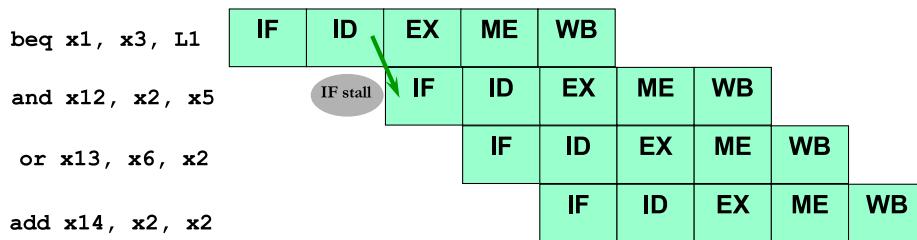
- **ALU logic** for comparison and addition.
- Additional **multiplexer and control signals** to direct the PC update.
- Expanded **data paths** for handling the offset and register values.

✓ Effect on Pipeline Execution

Let's consider an example. In a pipeline using early evaluation:

- The processor **only stalls for 1 cycle** after a branch instruction, as opposed to the 3 cycles required by the conservative approach.
- This **one-cycle stall** allows the processor to fetch the correct instruction **immediately after the branch** is resolved.

The following diagram illustrates that the instruction fetch after the branch is delayed by only one stall cycle, resulting in a **smaller performance hit**.



⌚ Conclusion

In summary, by anticipating branch evaluation at the ID stage, the processor reduces the branch penalty to 1 cycle per branch. This is a significant improvement over the 3-cycle stall of conservative stalling. While it introduces **hardware overhead**, it offers a **better balance between performance and correctness** and serves as a stepping stone toward even more advanced techniques, such as branch prediction.

2.4 Intro to Branch Prediction

In modern computer architectures, achieving high performance requires efficient instruction-level parallelism (ILP)¹. However, one of the **major obstacles to ILP is the occurrence of branch hazards**, which happen when the processor encounters a branch instruction (e.g., `if`, `for`, `while`) and cannot immediately determine which instruction to execute next. To mitigate the performance loss caused by these hazards, branch prediction is employed.

Branch Prediction is essentially a **speculative execution technique** where the processor *guesses the outcome of a branch instruction*, whether the branch will be taken (control jumps to the branch target) or not taken (execution continues sequentially), **before the actual result is known**. Instead of stalling the pipeline and waiting for the branch condition to be resolved, the processor **proceeds based on the predicted outcome**. If the prediction:

- ✓ Is **correct**, performance is preserved.
- ✗ Is **wrong** (a **misprediction**), the incorrectly fetched **instructions are flushed**, and **execution restarts at the correct address**, causing a **performance penalty**.

Branch Prediction categories

Branch prediction techniques are generally classified into two main categories:

- **Static Branch Prediction Techniques.** In this method, the **branch direction** (taken/untaken) is **decided at compile time** and **remains fixed during the program's execution**. Static prediction often relies on compiler heuristics or profiling data to guess likely outcomes. Since the behavior doesn't adapt to runtime changes, this **technique works best when branch outcomes are highly predictable and consistent** across executions.
- **Dynamic Branch Prediction Techniques.** Unlike static methods, dynamic prediction uses **hardware mechanisms to observe past branch behavior at runtime** and make predictions accordingly. The **prediction adapts to actual program execution**, making it more effective for applications with **complex or data-dependent control flow**. This method can dynamically switch its guess depending on the *branch history*.

It's important to note that in both static and dynamic techniques, the **processor must avoid updating its internal state** (registers, memory, etc.) **until the branch outcome is known with certainty**. This ensures speculative execution doesn't cause side effects in case of misprediction.

Additionally, **hybrid approaches** are possible, where static and dynamic predictions are combined to optimize performance further.

¹**Instruction-Level Parallelism (ILP):** A measure of **how many of the operations in a computer program can be performed simultaneously**. High ILP enables multiple instructions to be executed in parallel within a single processor cycle, exploiting the parallelism inherent in sequential instruction streams through techniques like pipelining, superscalar execution, and out-of-order execution.

2.5 Static Branch Prediction

Static branch prediction represents one of the **simplest approaches** to handling branch hazards. In this technique, the **prediction** regarding whether a branch will be taken or not is **made at compile time** and **remains unchanged throughout program execution**. This method **relies heavily on heuristics or compiler-generated hints**, which estimate the likely behavior of each branch without any consideration of the program's actual runtime behavior.

✓ When does static branch prediction work well?

This approach is particularly effective in scenarios where the **branch behavior is stable and highly predictable**, such as in embedded or domain-specific applications. In such cases, the overhead and complexity of dynamic prediction mechanisms may not justify the potential benefits, making static prediction a practical alternative.

⚠ RISC-V assumption

A key architectural note here is the assumption that we are working with a **RISC-V processor**, which is **optimized for early branch evaluation during the Instruction Decode (ID) stage** (see more in section 2.3, page 35). This means that in RISC-V, the decision to predict a branch direction occurs early in the pipeline, minimizing the potential for instruction fetch delays if the prediction is accurate.

2.5.1 Branch Always Not Taken

The **Branch Always Not Taken** strategy is the simplest form of static branch prediction. It operates under the **assumption that the branch condition will never be satisfied**, i.e., the control flow of the program will continue sequentially as if the branch is not taken. As a result, instructions immediately following the branch in program order are fetched and executed without any need to determine or access the **Branch Target Address (BTA)**.

➊ When is Branch Always Not Taken effective?

This approach is especially effective for **certain control flow patterns**, such as **if-then-else** structures where the **then** clause is more likely to be executed than the **else** clause. For example:

```

1 z = x + y;
2 if (z > 0)
3     w = x;
4 else
5     w = y;
```

Assuming z is typically positive, the branch is not taken because execution proceeds sequentially to $w = x$. This makes predict-not-taken a suitable and effective default strategy for such cases.

✖ Implementation Details

The prediction is made at the end of the **Instruction Fetch (IF) stage**, without calculating or knowing the **BTA** (since the branch is always not taken and the next instruction to execute is the $PC + 4$, as always). This makes the approach **lightweight and efficient**.

⚠ Misprediction Case

If the actual **branch outcome** (BO) evaluated during the Instruction Decode (ID) stage is **not taken**, then the **prediction is correct**, and **no penalty cycles** are incurred. The pipeline proceeds as planned.

Otherwise, if the actual **branch outcome** (BO) turns out to be **taken**, then the **prediction was incorrect**, leading to a **misprediction penalty**. The processor must:

1. **Flush** the fetched instruction(s) after the branch (turned into NOPs).
2. **Fetch the instruction** at the Branch Target Address (**BTA**) and **restart execution from there**.

This results in a **one-cycle branch penalty**, which is minimal but still affects performance.

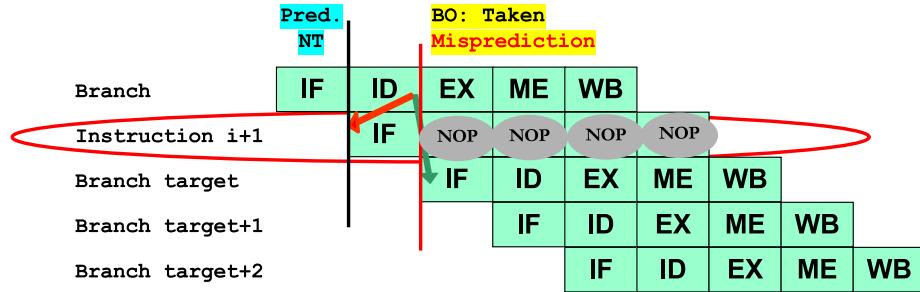


Figure 12: Branch Always Not Taken techniques failed and the processor must flush instruction $i+1$ and restart execution from the BTA.

2.5.2 Branch Always Taken

This approach represents the dual case of the previous technique (page 39): it assumes that **every branch will be taken**, meaning the **control flow will jump to the branch target address** rather than continue sequentially. This method is especially **useful for backward branches**, which occur in **loops** such as **for**, **while**, and **do-while**, since these branches are typically **taken repeatedly during loop iterations**.

❓ Implementation Challenge

Unlike the not-taken strategy, where the processor simply continues to $PC + 4$, the **taken strategy requires knowledge of the Branch Target Address (BTA)** at the Instruction Fetch (IF) stage. This is **non-trivial** because:

- ❓ The **BTA depends on the branch instruction's target**, which typically requires decoding.
- ✓ To solve this, we introduce a **Branch Target Buffer (BTB)**, a special hardware structure.

❓ What is BTB and why do we need it?

The **Branch Target Buffer (BTB)** is a **specialized cache** in the processor designed to predict the target address of a taken branch instruction **before the branch condition is actually resolved**.

In Branch Always Taken, we assume that the program will jump to a new address (the Branch Target Address, BTA). However, this **BTA is not immediately known during Instruction Fetch (IF)** because it typically requires decoding the branch instruction (Instruction Decode, ID). **To avoid delays**, the **BTB remembers past branch target addresses**, allowing the processor to quickly predict where to jump when encountering a branch.

❖ How does BTB work? Quickest explanation

- **BTB Structure**, it is a kind of lookup table or cache where:
 - Key: **address of the branch instruction** (the PC value where the branch resides)
 - Value: Predicted Target Address (PTA), i.e., where to jump if the branch is taken
- **BTB Lookup**: when fetching a branch instruction, the processor simultaneously queries the BTB via the branch PC.
 - ✓ If a **match is found** (*cache hit*), the **BTB immediately provides the Predicted Target Address (PTA)**, and the **processor starts fetching from that address**, before knowing if the branch is actually taken.

- ✖ If a **no match** (*cache miss*), the processor might **default to sequential execution** (PC + 4) or **wait for the BTA** to be calculated, which causes delay.

Example 1: BTB and Branch Always Taken technique

Let's say:

- A loop branch at address 0x100 typically jumps to 0x80.
- The BTB stores: 0x100 → 0x80 (key → value).

When the branch at 0x100 is fetched again:

- The BTB predicts the next instruction will be at 0x80 (taken).
- The processor **starts fetching from 0x80, without waiting to evaluate the branch condition.**

If it turns out the **branch was not taken**, the processor **flushes the incorrect fetch** from 0x80 and resumes at 0x104.

✓ Correct Prediction Path

If the **branch** is indeed **taken**, and the **BTB correctly supplies the BTA**, the processor proceeds **without penalty**. **Execution continues from the target address** just as expected.

✖ Misprediction Case

If the **actual outcome** is **not taken**, the **prediction is incorrect**:

1. The Instruction Fetched (IF) from the **target address** is **flushed (NOP)**.
2. The **processor must fetch the sequential instruction** at PC + 4.
3. **One-cycle penalty** incurred, similar to the not-taken misprediction case.

❓ When is this technique effective?

This method is **well-suited for loop constructs**, where **branches typically go backward** and are **taken with high probability**. For example:

- In a **do-while** loop, the branch is taken almost every time except the last iteration.
- Conversely, in forward branches like **if-then-else**, the branch is less likely to be taken, making this technique less effective.

This underscores that **branch direction** (forward or backward) **can influence the effectiveness of prediction strategies**.

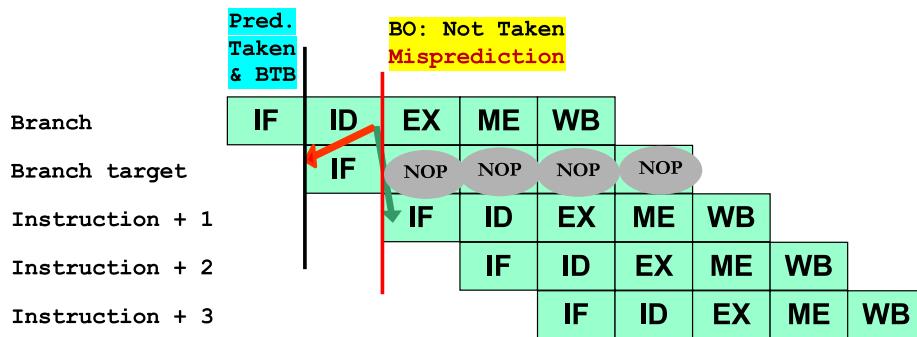


Figure 13: Branch Always Taken techniques failed and the processor must flush instruction $i+1$ and restart execution from the BTA.

2.5.3 Backward Taken Forward Not Taken (BTFNT)

The **Backward Taken Forward Not Taken (BTFNT)** strategy represents a refinement of static prediction that uses a simple yet effective heuristic: the direction of the branch, whether it jumps backward or forward in memory, can be used to predict its outcome.

Prediction Rule

- **Backward-going branches** (i.e., branches where the target address is lower than the current PC) are predicted as **taken**.
 - These branches **often occur in loops**, where execution loops back to an earlier instruction (e.g., in `for`, `while`, or `do-while` constructs).
- **Forward-going branches** (i.e., target address is greater than the current PC) are predicted as **not taken**.
 - These branches typically correspond to **if-then-else constructs**, where the `else` path is **less probable** and control usually **proceeds sequentially**.

Why does this work?

The rationale behind BTFNT lies in **empirical observations**:

- **Loops** tend to execute **multiple times**, hence **backward branches are mostly taken**.
- **Conditional statements** often have **rarely taken else paths**, hence **forward branches are mostly not taken**.

Pros and Cons

- ✓ Simple to implement because BTFNT requires just a comparison of the **target address vs the current PC**:
 - $\text{target address} < \text{PC} \Rightarrow \text{predict taken}$
 - $\text{target address} > \text{PC} \Rightarrow \text{predict not taken}$

Also, better accuracy than uniform always-taken or always-not-taken, especially for mixed codebases.

- ✗ Not adaptive; fails for atypical control flows where direction doesn't align with expected behavior.

2.5.4 Profile-Driven Prediction

Profile-Driven Prediction is a static prediction technique that **uses empirical data from previous program executions** to guide the prediction of branch outcomes. Rather than relying solely on heuristics or branch direction, this method **leverages profiling to derive probabilistic insights** about how branches behave under typical conditions.

❖ How does it work?

1. The **target application** is **executed multiple times beforehand**, using **diverse data sets** to simulate realistic execution scenarios.
2. During these early runs, the **behavior of each branch instruction is recorded**. Specifically, how often it was taken or not taken.
3. This **profiling produces statistics** for each branch, e.g., a pattern like:

T T T T T T T T NT NT NT

“Taken” is most probable.

4. Once the profiling is complete, the **compiler encodes a hint** directly into each branch instruction (e.g., in a dedicated **hint bit** in the instruction format):
 - 1: if the branch is **usually taken**.
 - 0: if the branch is **usually not taken**.

This enables the **processor** to **consult the hint during execution** and predict accordingly, without requiring runtime monitoring.

✓ Advantages

- ✓ Offers **higher accuracy than heuristics alone**, especially for **applications with stable branch behavior**.
- ✓ **No runtime hardware cost**, since prediction decisions are guided by **static hints**.

✗ Limitations

- ✗ **Static nature**: Predictions don't adapt to runtime variability; changes in data patterns may invalidate the profile.
- ✗ Requires **extra compilation effort**: profiling and hint encoding add complexity to the build process.
- ✗ **Less effective** for programs with **input-dependent control flow**.

2.5.5 Delayed Branch

The **Delayed Branch** technique is a static scheduling approach where the compiler plays a central role in mitigating branch penalties. Unlike traditional branch prediction, which involves guessing the outcome of a branch, delayed branching **reorders instructions so that useful work is done regardless of the branch direction**.

Core Concept

When a **branch instruction is executed**, it typically introduces a **delay before the processor can determine where to fetch the next instruction**. During this delay (known as the **branch delay slot**), rather than letting the pipeline sit idle or fetch incorrect instructions, the **compiler schedules an independent instruction to execute in that slot**.

- The instruction in the **branch delay slot** is always executed, regardless of whether the branch is taken or not.
- This allows useful work to be completed during what would otherwise be a stall or pipeline bubble.

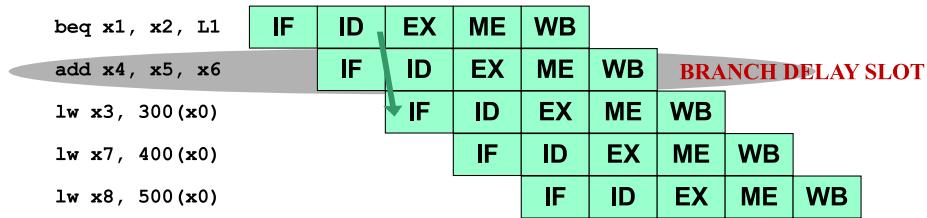


Figure 14: Example Scenario. In this case, the `add` instruction is scheduled after the branch, always executed, and does not affect the branch condition or outcomes.

Compiler Responsibility

A **critical task for the compiler** is to find a valid and useful instruction to place in the branch delay slot. The instruction must be:

- **Independent** from the branch decision.
- **Safe to execute** whether the branch is taken or not.

To guide this, the compiler can choose an instruction:

1. Section 2.5.5.1, page 48 - From **before** the branch.
2. Section 2.5.5.2, page 49 - From the **target** of the branch.
3. Section 2.5.5.3, page 51 - From the **fall-through path** (i.e., the sequential next instruction).
4. Section 2.5.5.4, page 53 - From **after** the branch.

We'll explore each of these four scheduling strategies step-by-step next.

Technique	Prediction Source	Complexity
Always Not Taken	Assume PC + 4	Very Low
Always Taken	Assume jump to BTA	Moderate (BTB)
BTFTNT	Direction-based	Low
Profile-Driven	Prior run data	High (compile time)
Delayed Branch	Compiler scheduled	High (compiler)
Technique	Risk	Best for
Always Not Taken	Mispredict backward branches	if-then-else
Always Taken	Mispredict forward branches	Loops (backward branches)
BTFTNT	Errors in irregular control flow	Mixed code (simple logic)
Profile-Driven	Profile mismatch	Stable behavior, performance-tuned
Delayed Branch	Wasted work if not efficient	RISC pipelines, e.g., MIPS processors

Table 2: Quick Comparison Table.

2.5.5.1 From Before

In the “**From Before**” strategy of delayed branch scheduling, the **compiler** selects an instruction that appears *before* the branch in program order and moves it into the branch delay slot. The selected **instruction must be independent of the branch decision** and safe to execute regardless of whether the branch is taken.

Key Characteristics

- The **instruction in the branch delay slot is always executed**.
- This instruction **will never be flushed**, since it is **guaranteed to execute** irrespective of the branch’s outcome.
- After the delay slot, execution continues normally, either to the branch target or the fall-through instruction, depending on whether the branch is taken.

Example 2: From Before

Original code:

```
1 add x1, x2, x3
2 beq x2, x4, L1
3 [delay slot, stall]
```

After scheduling:

```
1 beq x1, x2, x3
2 add x1, x2, x3      # delay slot filled
```

Here, the **add** is originally before the **beq** and has **no dependency** on the branch condition. It is safely moved into the delay slot.

Pipeline Behavior

Branch Not Taken

- The instruction in the delay slot is executed.
- Execution continues sequentially with the next instruction after the branch.

Branch Taken

- The delay slot instruction still executes.
- Execution jumps to the branch target after the delay slot.

The **instruction moved to the delay slot is always executed**.

Advantages

- ✓ **No need for instruction flushing**: the delay slot instruction is always valid.
- ✓ **Efficiency**: reuses existing instructions from earlier in the program to **hide the branch delay**.

2.5.5.2 From Target

In the “**From Target**” strategy, the compiler schedules an instruction from the *branch target* into the delay slot. This technique is useful when the branch is likely to be taken, as the delay slot instruction corresponds to what would naturally execute next in that control path.

Key Characteristics

- The delay slot contains an instruction from the branch target path.
- This strategy is typically used when the branch is taken with high probability, such as in loops.
- **Challenge:** If the branch is not taken, this instruction may be invalid and might have to be flushed (if mispredicted), or it must be safe to execute even if not needed.

Example 3: From Target

Original code:

```

1 sub x4, x5, x6      # target instruction
2 add x1, x2, x3      # branch instruction
3 if x1 == 0 then      # branch condition
4 [delay slot, stall]

```

After scheduling:

```

1 add x1, x2, x3      # if branch taken, go here!
2 if x1 == 0 then      # branch condition
3 sub x4, x5, x6      # delay slot filled

```

Here, the `sub` instruction from the branch target is moved into the delay slot. If the branch is taken, execution proceeds smoothly. If not, we either flush `sub` or ensure it causes no side effects.

Pipeline Behavior

- **Branch Taken**
 - The **delay slot** instruction is part of the intended control flow.
 - Execution continues with the **next instruction in the target path**.
- **Branch Not Taken**
 - Delay slot **may need to be flushed** (as it's not part of the sequential path), or must be **safe to execute anyway** (**no side effects or wasted computation**).

⚠ Instruction Duplication

When we move an instruction from the branch target into the delay slot, we still need to keep it at the target location because other parts of the code might also jump there.

Let's take an example to illustrate the problem. Let's say:

- We have **two branches that can jump** to label L1.
- **Instruction X** is the first instruction at L1.
- We move **Instruction X** into the delay slot of one branch, **but the other branch still needs to find instruction X at L1.**

Here's what happens:

```

1 Branch A → L1
2 Branch B → L1
3
4 L1:
5     Instruction X
6     Instruction Y

```

If Branch A decides to move Instruction X inside its delay slot, Branch B cannot see Instruction X anymore! Therefore, we have to keep Instruction X in two places:

1. In Branch A's delay slot
2. At Label L1 for Branch B

Okay, and that should be a problem? For three reasons:

- We've **duplicated Instruction X.**
- **More code = more memory = larger executable.**
- **Harder to maintain:** if we change Instruction X in one place, we might **forget to update the duplicate.**

✓ Best Use Case

Loops, particularly **do-while** constructs, where **backward branches are taken most of the time.**

2.5.5.3 From Fall-Through

In the “**From Fall-Through**” strategy, the **compiler selects an instruction that comes after the branch in program order** (i.e., from the fall-through path) and **moves it into the branch delay slot**. This method is **suitable when the branch is unlikely to be taken**, as execution will naturally continue sequentially.

Key Characteristics

- The **fall-through path** is **taken when the branch is not taken**.
- The **delay slot instruction** comes from this path, meaning it is **executed anyway if the branch is not taken**.
- If the **branch is taken**, the delay slot instruction is either:
 - **Flushed** (discarded), or
 - Must be **safe to execute** (no side effects), even though it becomes useless work.

Example 4: From Fall-Through

Original code:

```

1 add x1, x2, x3
2 if x1 == 0 then      # branch condition
3 [delay slot, stall]
4 or x7, x8, x9       # execute if branch is not taken
5 sub x4, x5, x6       # execute if branch is taken

```

After scheduling:

```

1 add x1, x2, x3
2 if x1 == 0 then      # branch condition
3 or x7, x8, x9       # delay slot filled
4 sub x4, x5, x6       # execute if branch is taken

```

Here, **or x7, x8, x9** is **moved into the delay slot** from the instruction that would **normally execute next** if the branch is **not taken**.

2 Pipeline Behavior

- **Branch Not Taken** (Mist Likely)
 - **Delay slot instruction** is correctly **executed**.
 - Execution proceeds sequentially.
- **Branch Taken**
 - **Delay slot instruction** is **not needed**.
 - It must be **flushed**, or **safe to execute** even though its result is discarded.

 **When is this strategy used?**

- ✓ When the **branch is not likely to be taken**.
- ✓ **Common in forward branches**, such as `if-then-else`, where `else` is rare.

Strategy	Delay Slot Instruction	Executed when branch
From Fall-Through	Instruction at PC + 4 (next in sequence)	Not Taken (common case)
From Target	Instruction at BTA (label target)	Taken (common case)

Table 3: Comparison between “From Target” and “From Fall-Through”.

2.5.5.4 From After

The “From After” scheduling technique is **rarely used because it is too complex to be practical**. However, in the “**From After**” strategy, the **instruction** scheduled in the **branch delay slot** is **taken** from a later point the code, specifically, from **after the fall-through instruction**.

Let’s number the instructions to make it easy:

```

1 Instr A      # Before branch
2 Branch       # Branch condition
3 Instr B      # PC + 4 (fall-through)
4 Instr C      # After fall-through ← from after
5 Instr D

```

In “from after”, the compiler **moves Instr C into the delay slot**, even though Instr B ($PC + 4$) should come right after the branch in normal execution (if not taken).

⚠ Why is this hard?

To safely move Instr C up into the delay slot, the **compiler must guarantee**:

1. **No Data Dependency Conflicts**

- Instr C must not use or modify data that depends on Instr A, B, or the branch outcome. For example, if Instr C uses a value computed in Instr B, moving it before Instr B causes incorrect results.

2. **Safe if Executed Early**

- Even if the branch is taken and Instr C should never execute, now it always executes in the delay slot.
- So Instr C must be safe to execute even when it’s not needed.
- We call this a *speculatively safe instruction*.

3. **No Control Flow Violation**

- If Instr C should only run after a condition is met, moving it earlier might break program logic.

2.6 Dynamic Branch Prediction

While static branch prediction relies on fixed rules or compile-time knowledge, **dynamic branch prediction** aims to **learn and adapt during program execution**. It uses hardware mechanisms to observe past branch behavior and predict future outcomes at runtime.

Core Idea

Dynamic prediction is **based on a key assumption**: if a branch behaved a certain way in the past, it's likely to behave the same way again. Therefore, instead of guessing statically, the processor monitors each branch at runtime and uses past outcomes to inform future predictions.

Hardware Components

Dynamic prediction relies on **two tightly-coupled hardware blocks**, both situated in the **Instruction Fetch (IF)** stage:

1. **Branch Outcome Predictor (BOP)**:

- Predicts branch direction: Taken (T) or Not Taken (NT).
- Based on runtime history (past outcomes of this or other branches).

2. **Branch Target Buffer (BTB)**:

- Predicts the target address to jump to if the branch is predicted taken.
- Returns the **Predicted Target Address (PTA)**².
- Useful only when BOP predicts Taken; irrelevant if Not Taken.

How it works

During instruction fetch:

1. BOP predicts T/NT.
2. If Taken (T), BTB provides the PTA.
3. The processor fetches the next instruction accordingly.

²**Predicted Target Address (PTA)**: The memory address that the processor predicts as the destination for a taken branch. If the branch is predicted taken, the Branch Target Buffer (BTB) provides the PTA so that instruction fetch can continue from this address without waiting for the branch condition to be resolved.

② Execution Scenarios

- **Prediction: Not Taken (PC + 4)**

- If Branch Outcome = Not Taken $\Rightarrow \checkmark$ Correct prediction \Rightarrow No penalty.
- If Branch Outcome = Taken $\Rightarrow \times$ Misprediction:
 1. Flush next instruction (NOP)
 2. Fetch from BTA (to understand where to jump)
 3. One-cycle penalty

- **Prediction: Taken (BTB used)**

- If Branch Outcome = Taken $\Rightarrow \checkmark$ Correct prediction \Rightarrow No penalty.
- If Branch Outcome = Not Taken $\Rightarrow \times$ Misprediction:
 1. Flush fetched target instruction (NOP) provided by BTB
 2. Fetch PC + 4 (next instruction sequentially)
 3. One-cycle penalty

Unlike static prediction, **dynamic prediction is adaptive**. If a branch changes its behavior at runtime, future predictions adjust accordingly.

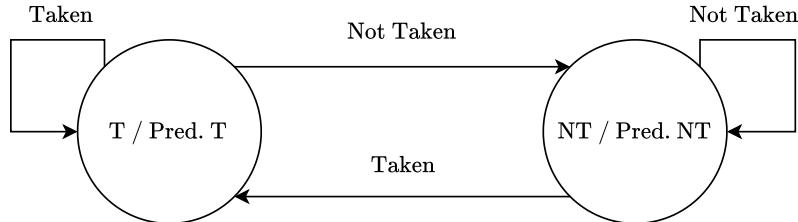
2.6.1 1-bit Branch History Table

In general, the **Branch History Table (BHT)**, or **Branch Prediction Buffer**, is a dynamic hardware structure that predicts branch outcomes based on recent execution behavior. Specifically, it is a Program Counter (PC)-indexed table that associates branch instructions with a prediction state, and dynamically predicts whether a branch will be *taken* or *not taken* using small finite-state machines (FSMs), such as 1-bit or 2-bit saturating counters.

The BHT can be seen as a **cache-like structure** that maps branch addresses to prediction states. Rather than storing an explicit sequence of past outcomes, each entry **implicitly encodes the recent behavior of the branch** through its FSM state, which evolves according to previous branch resolutions. This implicit history allows the processor to adapt predictions dynamically and to tolerate occasional anomalies, such as loop exits.

The simplest version uses 1 bit per branch to remember whether the branch was recently **taken or not taken**. For this reason, it is called a **1-bit Branch History Table (1-bit BHT)**. It operates at runtime and uses a **Final State Machine (FSM)** with 1-bit history:

- If last outcome was Taken \Rightarrow predict Taken.
- If last outcome was Not Taken \Rightarrow predict Not Taken.



❖ How it works

- Each branch instruction's address (or part of it) **indexes a table entry**.
- That entry holds a single bit (T/NT) representing the **last observed outcome**.
- On next encounter:
 - Use the **bit to predict the outcome**.
 - After actual branch resolution:
 - * **Correct prediction** \Rightarrow keep bit unchanged
 - * **Incorrect prediction** \Rightarrow flip bit

Indexing the Table

- Use k lower bits (right side) of the branch's address as the index.
- ⚠ **No tags:** any branch with the same low-order bits shares the entry (can cause interference).
- 2^k entries total.

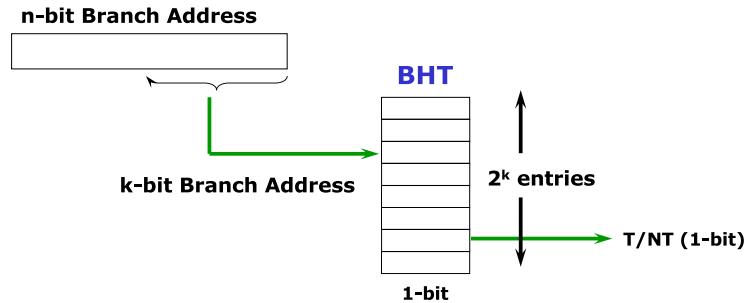


Figure 15: Visual representation of the 1-bit Branch History Table.

Example 5: Accuracy Issue

Consider a loop that executes 10 iterations. The expected behavior of the branch is:

T T T T T T T T T T NT

Where the last is Not Taken because the code must exit from the loop and must continue (and not jump).

There are **two mispredictions**:

- **At the end:** Since iteration 9 is marked as Taken, the 10th iteration is predicted to be Taken, since the BHT contains Taken. This throws a misprediction because the branch result is Not Taken.
- **On the next time loop starts:** since the last iteration is stored in the BHT as Not Taken, the BHT has to flip the bit on the next time loop starts, and a misprediction occurs.

As a result, with 100% of 10 iterations, the BHT only catches 8 out of 10 iterations, the accuracy is 80%.

✖ Shortcomings

- ✖ **Flipping prediction after 1 misprediction causes instability**, especially in loops.
- ✖ **Conflict problem:** two different branches with same index overwrite each other's bit.

 **Partial Solutions**

To reduce interference:

- Increase table size (more k bits).
- Use hashing to mix address bits better.

2.6.2 2-bit Branch History Table

2-bit Branch History Table (BHT) is an improvement over 1-bit BHT designed to increase prediction stability, especially for loops, and reduce mispredictions.

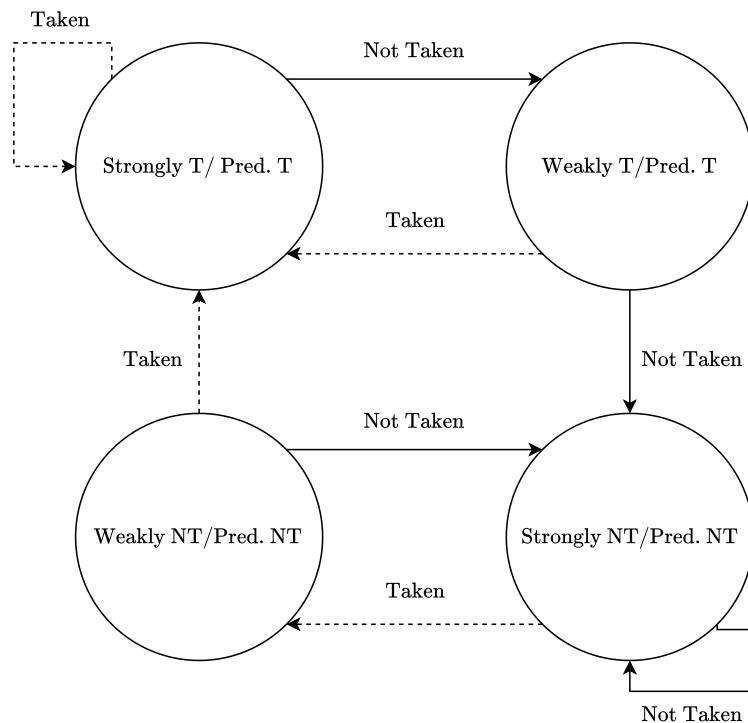
⌚ Why 2-bit? The Problem with 1-bit BHT

In loops, 1-bit BHT suffers from **flip-flopping**: 1 misprediction is enough to change the prediction. This causes two misprediction per loop:

- Exiting the loop (NT mispredicted as T).
- Re-entering the loop (T mispredicted as NT).

The 2-bit BHT introduces a **4-state FSM** using 2 bits per entry. It requires **2 consecutive mispredictions to change the predicted outcome**, thus adding stability. The FSM states are:

1. **Strongly Taken (ST)** → Predict Taken
2. **Weakly Taken (WT)** → Predict Taken
3. **Weakly Not Taken (WNT)** → Predict Not Taken
4. **Strongly Not Taken (SNT)** → Predict Not Taken



✓ Effect on Loops

Assume a loop with:

T T T T T T T T T NT

- Exit NT causes 1 misprediction, but FSM moves from Strongly Taken to Weakly Taken. So the prediction remains Taken.
- Re-enter on the loop causes a Branch Outcome Taken, and the 2-bit BHT predicts correctly because it is on the WT state.

Only 1 misprediction per loop, improving accuracy to 90% (from 80%).

✓ Benefits

- Improved accuracy in loops and repetitive patterns.
- Reduces misprediction penalty in typical branch-heavy code.
- Balances prediction stability and adaptability.

2.6.3 Branch Target Buffer

The **Branch Target Buffer (BTB)** is a **specialized cache used to store target address of taken branches**. The stored Predicted Target Address (PTA) allows the processor to fetch instructions from the target without delay when a branch is predicted taken. The PTA is typically stored in PC-relative format (offset from current PC).

Core Idea

While the Branch History Table (BHT) predicts *whether* a branch will be taken, the Branch Target Buffer (BTB) **predicts *where* the program should go if the branch is taken**. The **BTB stores Predicted Target Addresses (PTAs)** for previously encountered branches and **enables fast redirection of control flow**.

How Is the BTB Structured?

- The BTB is designed as a **direct-mapped cache**:
 - The **address of the branch instruction is used to index the BTB**.
 - **Tags** are used for associative lookup to **confirm correctness** (i.e., ensure the indexed entry really belongs to the current branch)
- Components per entry:
 - **Tags**: Identifies the branch instruction.
 - **PTA**: The Predicted Target Address.
 - Often combined with **T/NT bits** from a **Branch History Table** (1-bit or 2-bit) for branch outcome prediction.

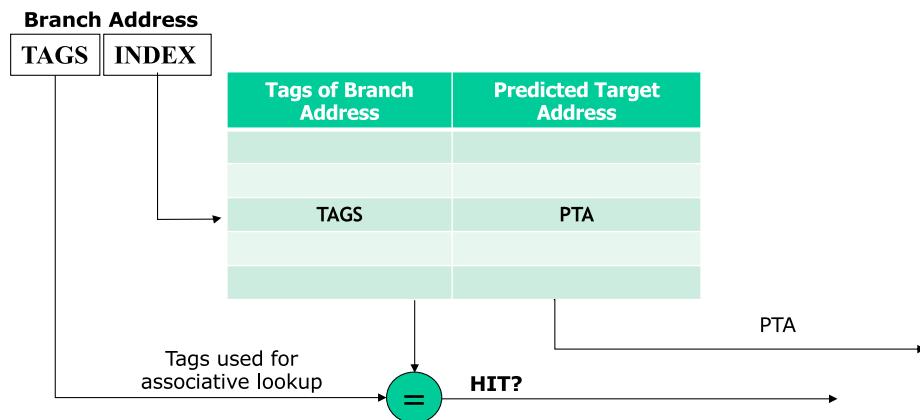


Figure 16: Branch Target Buffer without Branch Outcome Predictor.

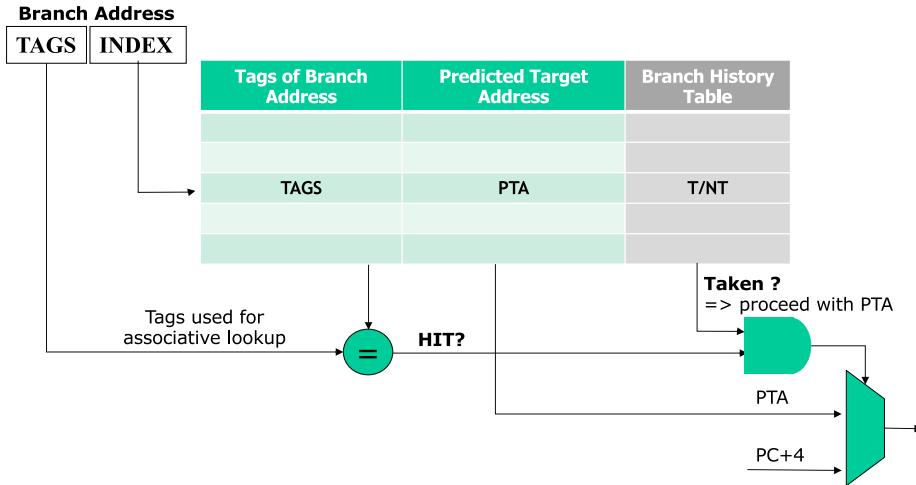


Figure 17: Branch Target Buffer with Branch Outcome Predictor.

BTB in the Pipeline

It is placed in the Instruction Fetch (IF) phase. During fetch:

- The BTB is queried using the current PC (branch address).
- If **hit** and **BHT predicts Taken**, **fetch from PTA**.
- If **miss** or **predict Not Taken**, **continue at PC + 4**.

Prediction	BTB use	Action
Predict Not Taken	BTB not used	Fetch from PC + 4
Predict Taken (BTB hit)	BTB used	Fetch from PTA stored in BTB
Predict Taken (BTB miss)	BTB miss	Stall or default to PC + 4, then calculate BTA

Table 4: Summary of Behavior.

Advantages

- ✓ **Eliminates delay** from calculating the **branch target address** manually.
- ✓ Enables speculative instruction fetch from correct target, **improving pipeline efficiency**.

Deepening: Branch Target Buffer

✖ What is the problem without a BTB?

Consider a pipelined processor. When we fetch an instruction at address PC:

- If it's **not a branch** instruction → fetch next instruction at PC + 4.
- If it's a **branch** instruction, we must know:
 1. *Will it be taken or not?* (predicted by BHT)
 2. *If taken, where do we jump to?* (target address)

Now the key point is that **the branch target address is usually computed later in the pipeline** (e.g., in ID or EX stage using PC and an offset or a register value). So even if we can *predict* that the branch will be taken, **we still don't know where to fetch from**.

⚠ Assume we have a Branch History Table (BHT) that predicts the branch is **taken**, but the target address is not yet known. Without a BTB, the processor must **stall instruction fetch, or fetch sequentially and later flush**. Either way, this loses cycles and speculation becomes weak and inefficient. So, this is the **control hazard BTB is designed to solve**.

❓ What a BTB actually does?

A **Branch Target Buffer (BTB)** is a small, cache-like table that remembers **the target addresses of previously executed branches**, indexed by the branch PC. Each BTB entry typically contains:

- **Tag**: A branch instruction address, PC. It identifies the branch.
- **Predicted Target Address (PTA)**: Target address. The address to jump to if the branch is taken.
- (often) **Prediction State**: A bit or counter indicating taken/not taken (sometimes combined with BHT).

✖ How BTB works in practice?

During the **instruction fetch (IF) stage** (i.e., when fetching the branch instruction):

1. The PC is used to:
 - Index the **BHT** → predict *taken* or *not taken*.
 - Index the **BTB** → get the *Predicted Target Address* if the branch is taken.

2. If the BHT predicts **taken** and there is a **BTB hit** (i.e., the branch PC matches the BTB tag), the **next PC is immediately set to the target address**, in the **same cycle**. This allows no stalling, no waiting, and fetching continues smoothly.

✗ What BTB does not do? The BTB does **not** decide *whether* the branch is taken or not (that's the BHT's job), **nor** does it compute the target address (that's done later in the pipeline), **nor** does it replace BHT. It simply **remembers previously computed target addresses** for quick access during fetch, like a cache for branch targets.

⚠ What happens on BTB miss?

In the previous explanation, we always assumed that there is a **BTB hit** when the BHT predicts taken. However, like any cache, the BTB can also **miss**. If the BHT predicts taken but there is a **BTB miss** (i.e., no entry for that branch PC):

- The processor must wait until the target is computed later in the pipeline (e.g., in ID or EX stage).
- This causes a **stall or delayed fetch**.
- Once the target is computed, the BTB can be **updated** with the new target address for future predictions.

In other words, **it's as if there's no BTB for that branch at all**, which leads to the same control hazard problem that we wanted to avoid. Clearly, the best option is to have a BTB with a certain rate of misses rather than no BTB at all.

2.6.4 Correlating Branch Predictors

❷ What is the problem?

With standard BHT, we **predict each branch individually**, based only on its **own past behavior**. But real programs often have **branches that influence each other**. For example, let's look at the following code:

```

1 if (x > 0) // Branch A
2 ...
3 if (x > 0) // Branch B
4 ...

```

- Branch B often behaves like Branch A, because they depend on the same condition ($x > 0$).
- A normal predictor doesn't know this. It treats A and B independently.

■ Key Idea

Use **global branch history** (outcomes of previous branches) to **improve prediction** for the current branch. This approach exploits correlation between different branches. This technique is called **Correlating Predictors** or **2-level Predictors**.

■ General Case: (m, n)

In a **(m, n) correlating predictor**, the **past outcomes of the last m branches** are used to select among 2^m prediction tables, each of which uses n -bit prediction entries.

- m : The number of **global history bits**.
- n : The number of **bits per prediction entry in the BHTs** (e.g., 1 or 2).

It works like this:

1. **Track the Last m Branches:** Store the outcomes (T/NT) of the last m branches in a **Global History Register (GHR)**. This forms an m -bit global history pattern.
2. **Use GHR to Select Table:** The m -bit GHR selects 1 out of 2^m Branch History Tables (BHTs). Each BHT contains n -bit entries.
3. **Index the Selected Table:** Use low-order bits of the branch instruction address (e.g., PC bits) to **index an entry in the selected table**.
4. **Predict Using n -bit Entry:** Use the n -bit entry to predict:
 - 1-bit BHT: predict Taken or Not Taken.
 - 2-bit BHT: use 4-state FSM (Strong and Weak Taken and Not Taken)

So what we have in the **memory** is:

- **Total tables:** 2^m BHTs.
- **Each BHT has:** 2^k entries (k is the number of PC bits used).
- **Total entries:** $2^m \times 2^k$.

✓ Advantages

- Captures patterns across multiple branches.
- Helps in complex control flow where a branch's outcome depends on prior branches.
- More accurate than per-branch-only prediction.

2.6.4.1 (1,1) Correlating Predictors

Use the **result of the last executed branch** (global history, $m = 1$ bit) to **choose between two prediction tables**, each of which has 1-bit entries.

- **1-bit Global History:** Stores last branch outcome (Taken = 1, Not Taken = 0).
- **2 BHTs (T1 & T2):** Each is a 1-bit predictor table, selected based on global history. We use a 1-bit Branch History Table technique.
- **Indexing:** Use PC low-order bits to index into the selected table.

Consider a pseudo code:

```

1 if (x > 0) // Branch A
2 ...
3 if (x > 0) // Branch B
4 ...

```

Let's say if A is true, B is usually true. The execution walkthrough:

- Cycle 1: **First Execution of Branch A**
 1. Global History: unknown or NT (0); because it doesn't track anything yet. We assume Not Taken (0).
 2. Use Table T2 (since GH = 0).
 3. Index into T2 with **Branch A**'s PC bits.
 - Predict: Not Taken!
 - Unfortunately, the Branch Outcome (BO) says Taken $\Rightarrow \times$ **Mispredict** \Rightarrow update T2 entry to T.
 4. Update Global History = 1 (Taken).
- Cycle 2: **Now Executing Branch Branch B**
 1. Global History = 1 (Taken) from **Branch A**.
 2. Use Table T1 (since GH = 1).
 3. Index with **Branch B**'s PC bits.
 - T1 says: "Try to predict as Taken".
 - The Branch Outcome (BO) says Taken $\Rightarrow \checkmark$ **Correct prediction**.
 4. No update needed. The Global History doesn't need an update either, because it is already 1 (Taken).

Since **Branch A** was Taken, **Branch B** is likely to be Taken. This is a smart technique because normal predictors treat **Branch B** alone. Instead, the Correlating Branch Predictor uses context, and the last branch helps predict this one.

Branch	Global History	Table Used	Prediction	Outcome	Update
Branch A	0	T2	NT	T	T2 entry to T
Branch B	1	T1	T	T	No change

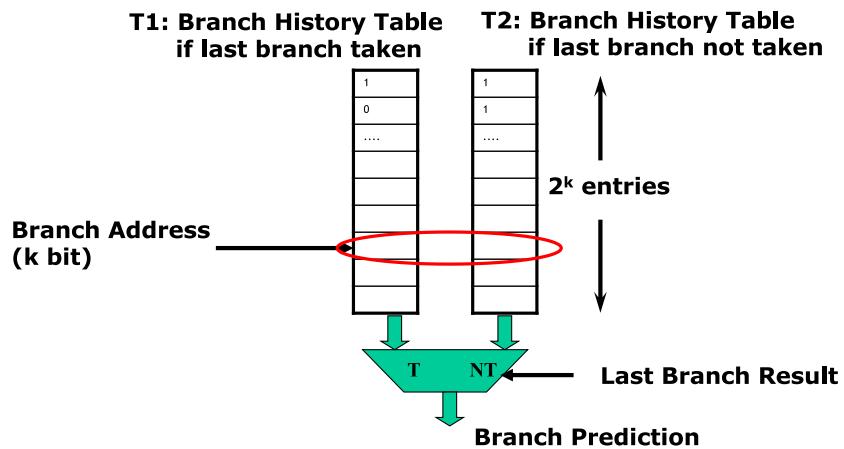


Figure 18: Visual representation of the (1,1) Correlating Predictor.

Aspect	Description
$m = 1$	Use 1-bit GHR (last branch result).
$n = 1$	1-bit prediction, (T or NT).
Tables	2 BHTs (for GHR = 0 and GHR = 1).
Selection Logic	If last branch T, use Table T1; else T2.

Table 5: (1,1) Correlating Predictor.

2.6.4.2 (2,2) Correlating Predictors

The correlating predictors with $m = 2$ and $n = 2$ have the following components:

- **2-bit Global History:** Stores outcomes of the last 2 branches. Forms 4 patterns:
 - 00: both Not Taken.
 - 01: last branch Taken, penultimate branch Not Taken.
 - 10: last branch Not Taken, penultimate branch Taken.
 - 11: both Taken.
- **4 Prediction Tables:** One for each global history pattern ($2^2 = 4$ BHTs).
- **2-bit entries per BHT:** Each BHT uses 2-bit saturating counters for stable predictions.
- **Indexing:** Use PC low bits + global history to access an entry in a BHT.

Consider a pseudo code:

```

1 if (A)           // Branch 1
2 ...
3 if (B)           // Branch 2
4 ...
5 if (C)           // Branch 3
6 ...

```

Let's simulate Branch 3's prediction, influenced by Branches 1 & 2.

0. Initial State

- Global History Register (GHR) = 00 (no branches taken yet).
- BHT for history 00 selected.
- Predicts Branch 3 using its 2-bit counter in BHT[00].

1. Cycle 1: Branch 1 = Taken

- GHR: 00 → 01 (shift in T = 1, Taken).
- Update BHT[00] (for Branch A), since we used GHR = 00 before A.

2. Cycle 2: Branch 2 = Not Taken

- GHR: 01 → 10 (T, NT).
- Update BHT[01] (for Branch B), since GHR = 01 before B.

3. Cycle 3: Predict Branch 3

- GHR = 10, so select the BHT that contains 10 for prediction.
- Use Branch 3's PC low bits + GHR = 10 to index BHT[10].

💡 Check the 2-bit FSM in this entry. Assume the FSM state is Weakly Taken, so the **prediction is Taken**.

💡 The outcome of Branch 3 is Taken, so the **prediction is correct** and we update the FSM of entry BHT[10].

Cycle	Branch	Outcome	GHR Before	GHR After	BHT Used for Update
1	A	T	00	01	BHT[00]
2	B	NT	01	10	BHT[01]
3	C	T	10	01	BHT[10]

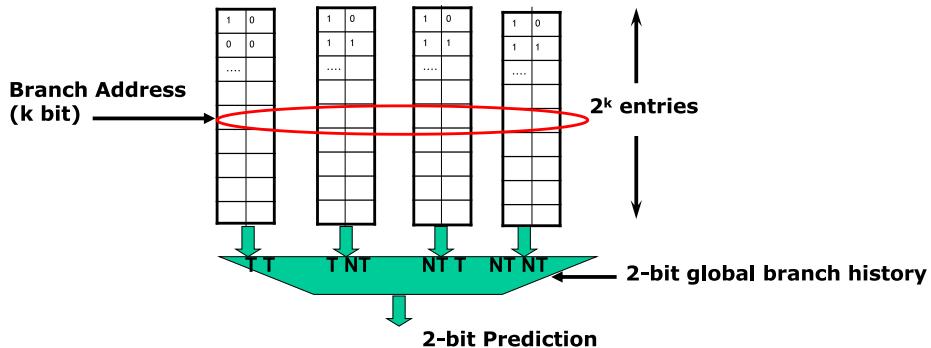


Figure 19: Visual representation of the (2,2) Correlating Predictor.

Aspect	Description
$m = 2$	Track outcomes of last 2 branches (GHR = 2 bits)
$n = 2$	2-bit prediction entries per BHT
Tables	4 BHTs (for GHR = 00, 01, 10, 11).
Indexing	4-bit PC + 2-bit GHR \rightarrow 6-bit index for accessing a table
Prediction Stability	More robust due to 2-bit FSM per entry.

Table 6: (2,2) Correlating Predictor.

2.6.5 Two-Level Adaptive Branch Predictors

Two-Level Adaptive Branch Predictors are advanced techniques that aim to provide highly accurate and adaptive predictions by **combining history tracking with pattern-based decision-making**. Unlike simpler predictors that use only the outcome of the last branch or last few outcomes of a single branch, these predictors **consider patterns over time and across different branches** to improve prediction accuracy.

Core Concept

The two-level approach consists of:

1. A **history-tracking component**: to record the outcomes of recent branches.
2. A **pattern-based prediction component**: to use that history to make accurate predictions.

This **design allows the processor to learn and adapt to recurring patterns in branch behavior**, which is particularly useful for complex control flows and loops.

Structure

1. Branch History Register (BHR)

- A **k -bit shift register** that records the outcomes of the k most recent branches (e.g., T, NT, NT, T).
- The BHR can be either:
 - **Global**: one register for all branches.
 - **Local**: separate register for each branch.

2. Pattern History Table (PHT)

- A table of **2-bit saturating counters** (like in 2-bit BHT).
- **Indexed** using the **content of the BHR**.
- Each **entry** provides a **prediction** (Taken/Not Taken) and adapts over time.

3. Prediction Process

- (a) Use the BHR value to index the PHT.
- (b) Read the 2-bit counter at that entry.
- (c) Predict Taken if in a Taken state, otherwise Not Taken.
- (d) After the actual branch outcome:
 - i. Update the 2-bit counter accordingly.
 - ii. Shift the actual outcome into the BHR.

Global Adaptive Predictor (GA)

The **Global Adaptive Predictor (GA)** is a specific form of the two-level predictor where **global history (BHR)** is used to index the PHT.

- The **BHR** is shared across all branches, and thus captures the global correlation among different branches.
- The **PHT** is local in the sense that it provides per-entry adaptation via 2-bit counters.

The main **advantage** is that by correlating the current branch with the behavior of previous branches (stored in the BHR), the predictor can detect global patterns and make more informed predictions.

GShare Predictor

GShare is a variation of the Global Adaptive Predictor, designed to improve the indexing of the PHT and reduce aliasing (i.e., different branches mapping to the same PHT entry).

Instead of directly using BHR to index the PHT, **GShare performs an XOR** between:

- The **BHR** (global history of recent outcomes).
- The **low-order bits of the program counter (PC)** of the current branch.

The XOR operation **mixes the global history with branch-specific information**, making it more likely that **different branches will access different entries** in the PHT, thus **reducing prediction interference** (aliasing). This allows GShare to **reduce aliasing and have a global and local view**.

Predictor	History used	Indexing to PHT	Benefit
Global Adaptive (GA)	Global BHR	BHR value directly indexes PHT	Simple, effective for globally correlated branches
GShare	Global BHR + PC	BHR XOR PC bits index PHT	Reduces aliasing, captures global + local context

Table 7: Summary of Global Adaptive and GShare.

Deepening: What is Aliasing?

In general, **Aliasing** refers to a situation where **two or more distinct entities become indistinguishable to the hardware or software**, leading to **incorrect behavior or loss of information**.

It appears in **multiple layers of the system**, but the underlying idea is always the same: **different things look the same to a limited**

representation.

More formally, **Aliasing** occurs when different signals, addresses, or data objects map to the same representation, so the system cannot tell them apart. This limitation may come from: finite resolution, limited bits, reduced indexing, or partial information.

- **Memory Aliasing** (also called **Address Aliasing**): two different memory locations appear to be the same to some hardware structure (e.g., cache, TLB).
 - ?
 - Causes:** Typical causes include virtual-to-physical address translation and cache indexing with partial address bits.
 - ⚠ **Consequences:** can lead to false cache hits, silent data corruption, and coherence issues.

Example 6: Cache Aliasing

Consider a cache that uses only the lower bits of the address for indexing. Two different physical addresses:

```
0x0000_1234
0x1000_1234
```

Map to the same cache line. They **alias** in the cache, leading to potential false hits and data corruption.

- **Register Aliasing:** occurs when different architectural registers map to the same physical register in the hardware, leading to false dependencies. We will explore this in future sections like Tomasulo's algorithm (page 140). However, the core idea is that **distinct logical registers appear identical at the physical level**, causing stalls and reduced parallelism. The solution is **register renaming** (page 77) where hardware dynamically maps architectural registers to distinct physical registers.
- **Branch Predictor Aliasing:** occurs when different branches map to the same entry in a branch predictor table (e.g., BHT, PHT). This can lead to **interference**, where the prediction for one branch is affected by the behavior of another, reducing accuracy. Techniques like GShare (page 72) help mitigate this by mixing branch-specific information (PC) with global history to create more unique indices.

Aliasing is a fundamental challenge in computer architecture, as it arises from the need to **balance limited hardware resources with the complexity of modern programs**.

3 Instruction Level Parallelism

3.1 The problem of dependencies

Instruction-Level Parallelism (ILP) is a foundational concept in modern processor design that aims to improve performance by executing multiple instructions simultaneously within a single processor core. The fundamental premise of ILP is that many instructions within a program can be executed independently, and thus, can be overlapped in time. This section explores the principles of pipelining as a means to exploit ILP, emphasizing its benefits, ideal performance metrics, and the inherent limitations.

However, **instruction dependencies** represent critical constraints on this parallelism. Understanding these dependencies is essential for analyzing the potential parallelism in a program and for designing hardware or compilers that can exploit ILP safely and efficiently.

Instruction dependencies determine which instructions can be executed simultaneously and which ones must respect a specific order of execution. These dependencies are classified into three broad categories: **data dependencies**, **name dependencies**, and **control dependencies**.

✓ Correct Program Behavior

For correct program behavior, two program properties must always be preserved during instruction scheduling:

1. **Data Flow**. The correct values must be produced and consumed in the proper order.
2. **Exception Behavior**. Reordering must not alter the way exceptions are raised and handled in the program.

While dependencies are intrinsic to the program semantics, hazards are an architectural artifact of the pipeline implementation.

3.1.1 Data Dependencies

Data Dependencies, also called **True Data Dependencies**, are the most fundamental type of instruction dependencies in a program. They **express the real flow of data from one instruction to another** and are dictated by the *semantics* of the program. These dependencies **must be strictly preserved** during any reordering or parallel execution of instructions, otherwise the correctness of the program is compromised.

Formally, we say there is a data dependency from instruction I_i to instruction I_j (where I_j follows I_i in program order), if I_j **reads a value that is produced by I_i** . In other words, I_j needs the output of I_i as its input. This is known as a **Read After Write (RAW)** hazard in pipeline terminology.

❓ Why is it called “true”?

This type of dependence is “*true*” because the **second instruction cannot proceed correctly until the first one completes its write operation**. It reflects an **actual requirement for program correctness**.

Example 1: RAW Hazard

```
1 I1:   r3 ← r1 + r2    # produces a value in r3
2 I2:   r4 ← r3 + r5    # consumes the value from r3
```

Here, I_2 is data-dependent on I_1 because it reads from register r_3 , which is written by I_1 . The instructions must execute in order:

- I_1 must execute and complete its write to r_3 before.
- I_2 reads r_3 to perform its own computation.

If this order is violated, e.g., I_2 executes before I_1 finishes, then I_2 will read an incorrect or undefined value.

⚠ Why Data dependencies Matter for ILP

Data dependencies define which **instructions must not be executed in parallel**, because doing so would result in violating program semantics.

- In a **pipelined processor**, data dependencies may cause **pipeline stalls**.
- In **out-of-order processors**, special mechanisms (like reservation stations and the reorder buffer) track and resolve data dependencies to allow other independent instructions to proceed while dependent ones wait for operands.

3.1.2 Name Dependencies

Unlike true data dependencies, **Name Dependencies** arise when **two instructions use the same register or memory location**, but **there is no actual flow of data between them**. These dependencies are called **False Dependencies** or **Pseudo-Dependencies** because they are **not required** for program correctness from a data perspective, **but still impose constraints** on instruction scheduling.

These constraints are due to the reuse of names (i.e., identifiers like register names), not due to real dependencies in the data values. They may still cause hazards in a pipeline and need to be addressed, especially when trying to execute instructions in parallel or out of order.

☰ Types of Name Dependencies

There are two main types:

- **Anti-Dependence (Write After Read - WAR)**. An anti-dependence occurs when:

- A first instruction **reads** from a location (register/memory);
- A second instruction **writes** to that same location, after it.

This introduces a WAR hazard: if the **second instruction is executed too early** (before the first instruction finishes reading), **it might overwrite the value before the first instruction uses it**.

Example 2: WAR Hazard

```
1 I1:   r3 ← r1 + r2    # reads r1 and r2
2 I2:   r1 ← r4 + r5    # writes to r1
```

In this case:

- I1 reads from r1
- I2 writes to r1

There is no data flow between the two (i.e., I1 doesn't use the result of I2, and vice versa), but **if I2 executes before I1 finishes, the read in I1 may get a corrupted value**.

- **Output Dependence (Write After Write - WAW)**. An output dependence occurs when **two instructions write to the same location** (register or memory). This results in a WAW hazard: executing the **second instruction first may overwrite the location**, changing the final value from what the program originally intended.

Example 3: WAW Hazard

```

1 I1:   r3 ← r1 + r2    # writes to r3
2 I2:   r3 ← r4 + r5    # writes to r3

```

There's no direct data flow between the two, but the **ordering matters**. If I2 is supposed to overwrite r3 after I1, reversing the order would result in I1's result being incorrectly seen as the final value in r3.

✓ Resolving Name Dependencies: Register Renaming

The key idea in dealing with name dependencies is to **recognize that they are not real** and can be **eliminated if we avoid the reuse of names**. The technique used to **eliminate these artificial constraints** is called **Register Renaming**. The idea is simple:

- If two instructions refer to the same register but don't actually share data, assign them **different physical registers**.

This is only possible when the underlying hardware (or compiler) provides more physical registers than the number of logical registers visible in the ISA.

Example 4: Resolving WAR and WAW

Original code (WAR):

```

1 I1:   r3 ← r1 + r2
2 I2:   r1 ← r4 + r5

```

Renamed:

```

1 I1:   r3 ← r1 + r2
2 I2:   r9 ← r4 + r5    # write to r9 instead of r1

```

Now, there is no conflict, I2 can proceed independently of I1.

Original code (WAW):

```

1 I1:   r3 ← r1 + r2
2 I2:   r3 ← r4 + r5

```

Renamed:

```

1 I1:   r3 ← r1 + r2
2 I2:   r9 ← r4 + r5    # write to a new register

```

🔗 Hardware vs. Software Register Renaming

- **Hardware (Dynamic Renaming).** Performed at runtime by structures such as the Register Alias Table (RAT), typically in out-of-order superscalar processors.
 - ✓ Flexible
 - ✗ Adds hardware complexity
- **Software (Static Renaming).** Performed at compile time by the compiler, particularly for VLIW or statically scheduled processors.
 - ✓ Simpler in hardware
 - ✗ Puts more pressure on compiler technology

3.1.3 Control Dependencies

While data and name dependencies arise from how instructions read and write operands, **Control Dependencies** stem from **the flow of control in the program**, that is, the **presence of branches and conditional execution**.

Control dependencies are fundamentally about **deciding whether an instruction should execute at all**, based on the **result of a preceding branch or conditional instruction**.

Formally, an instruction I_j is **control-dependent** on a branch instruction I_b if:

- I_j must only execute **if a particular outcome** of I_b is taken.
- But the **decision** made by I_b (e.g., whether to branch or not) is **not yet known** when I_j enters the pipeline.

This introduces uncertainty: *should I_j be fetched and executed, or not?*

Example 5: Control Dependencies

```
1 if (x > 0)
2     A; // Instruction A is control dependent on the condition
      (x > 0)
```

In assembly:

```
1 I1: bgtz r1, LABEL    # branch if r1 > 0
2 I2: ...                 # instruction before LABEL
3 I3: LABEL: A           # instruction A
```

- A should only execute **if the branch is taken**.
- But we don't know whether the branch is taken until the condition is resolved, which happens later in the pipeline.

❷ Why Control Dependencies matter for ILP

Control dependencies **limit instruction parallelism**:

- We cannot freely reorder or speculate on the instructions following a branch.
- Waiting for the branch outcome introduces **stalls** in the pipeline.

Thus, exploiting ILP requires **breaking or relaxing control dependencies**, without violating program semantics.

✓ Control Dependencies Solution

We have dedicated an entire section to this topic, see 2, page 30.

3.2 Multi-Cycle Pipelining

As processor microarchitectures evolved to support more complex instructions and higher performance demands, the basic model of a uniform single-cycle pipeline became insufficient. In practice, **many instructions**, especially those involving floating-point operations, memory access, or division, **require more than one clock cycle** to complete their execution or memory stages.

This leads to the development of **multi-cycle pipelines**, where **individual stages** (particularly EX and MEM) may **last for multiple cycles**, depending on the instruction type and runtime events. In such architectures, the ability to manage instruction progress intelligently becomes central to maintaining high throughput and correctness.

Motivation and Assumptions

In a classical 5-stage pipeline (IF, ID, EX, MEM, WB), all **stages are assumed to complete in one clock cycle**. However, this assumption doesn't hold in realistic systems:

- **Integer instructions** may complete in 1-2 cycles.
- **Floating-point operations**, like multiplication or division, can take 3 to 10+ cycles.
- **Memory access** times are unpredictable due to cache hits and misses, which can add variable delays.
- **Instruction fetch** may also stall due to instruction cache misses or branch resolution delays.

To accommodate these characteristics, **processors adopt multi-cycle pipelines**, where:

- Execution latency varies by operation type.
- Memory access can take multiple cycles.
- Functional units are not necessarily pipelined, particularly for floating-point operations.

Basic assumptions in this model:

1. The processor is **single-issue** (one instruction issued per cycle).
2. **Instructions** are typically **issued in-order** (fetched and passed into the pipeline in the order that they appear in the program).
3. Execution and memory stages may involve **multiple functional units with variable latencies**.
4. **Write-back** is often **delayed or synchronized** to ensure consistent state updates and avoid hazards.

3.2.1 Multi-Cycle In-Order Pipeline

In a **Multi-Cycle In-Order Pipeline**, instructions are:

- **Issued in program order (in-order issue).**
- **Executed on dedicated functional units**, each potentially requiring multiple cycles.
- **Committed in order**, i.e., write-back to the architectural state occurs in the order instructions were issued (**in-order commit**).

This model retains a **strict discipline**:

- Even if a later instruction finishes earlier (because it uses a faster unit), it **cannot write back until its turn arrives**.
- This **avoids WAR** (Write After Read) **and WAW** (Write After Write) hazards by ensuring in-order commit.

✓ Advantages

- ✓ Simpler control logic.
- ✓ Preserves the precise exception model.
- ✓ **No need for register renaming or reorder buffers.**

✗ Disadvantages

- ✗ **Poor ILP exploitation:** independent instructions may stall behind slow ones.
- ✗ All instructions are serialized through the same issue logic, even if no true dependence exists.

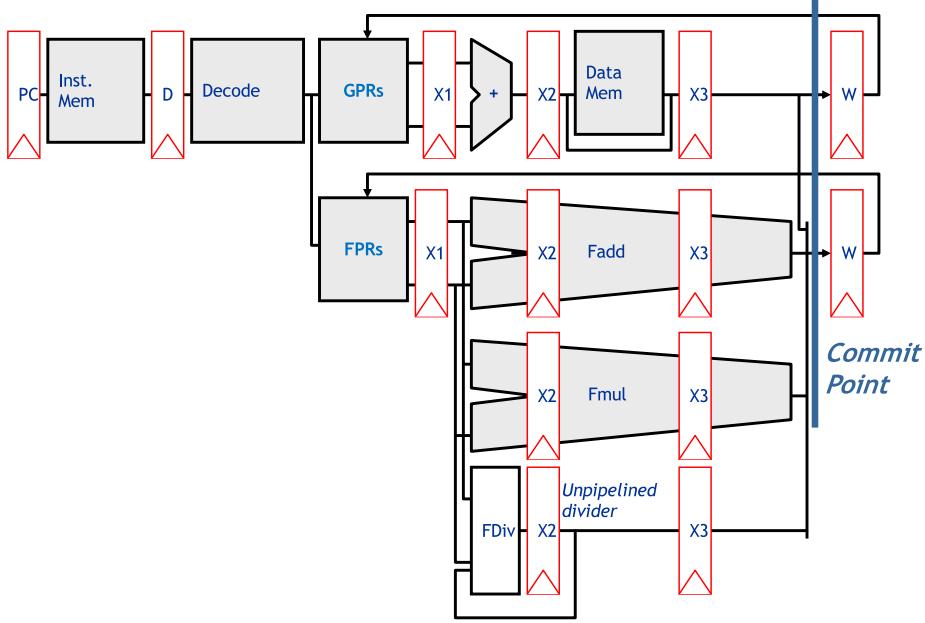


Figure 20: Multi-Cycle In-Order Pipeline architecture. This processor includes **different execution units**, each optimized for a specific operation type:

- X1-X2: 2-stage execution, used for basic integer ALU.
- Fadd: 3-stage execution, used for floating-point addition.
- Fmul: 3-stage execution, used for floating-point multiplication.
- FDiv: not pipelined, used for floating-point division (long)

So integer operations may take 2 cycles, add/mul take 3, and divide takes many cycles and cannot overlap because it's not pipelined. The instruction flow moves through:

- IF (Instruction Fetch): from PC and instruction memory.
- D (Decode): identifies operand registers, selects execution unit.
- X1, X2, ...: execution pipeline stages, depending on instruction type.
- Data Mem: if needed (e.g., for load/store).
- W (commit point) + GPRs/FPRs: Write-Back (WB) stage to either General-Purpose Registers or Floating-Point Registers.

Feature	Description
Issue	In order
Execution	In order
Completion (write-back)	In order (even if execution latency differs)
Architectural State Updates	In order; results are committed exactly in program order
ILP	Limited, stalls propagate even to independent instructions
Complexity	Moderate, simpler control logic, no renaming or reorder buffer needed
Exceptions	Always precise, easy to track and recover since instructions complete in order

Table 8: Summary of Multi-Cycle In-Order Pipeline.

3.2.2 Multi-Cycle Out-of-Order Pipeline

To overcome the limitations of in-order execution, processors adopt **Multi-Cycle Out-of-Order (OoO) Pipelines**. It is a more sophisticated architecture that aims to maximize ILP by **executing independent instructions as early as possible**, regardless of program order, and **allowing instructions to complete out of order**.

In this model:

- **Instructions** are still fetched and decoded **in-order** (**in-order issue**).
- After decoding, **instructions are placed into issue queues or reservation stations** (**out-of-order execution**).
- As soon as operands are available and a suitable functional unit is free, instructions execute, regardless of original order.
- **Write-back and commit may also occur out-of-order**, although the final architectural state is updated in program order to preserve correctness (**out-of-order commit**).

✓ Advantages

- ✓ **Maximizes ILP** by letting independent instructions execute as soon as possible.
- ✓ **Improves throughput** by keeping all functional units busy.
- ✓ **Hides long latencies**, like FP division or memory misses.

⚠ Challenges

Out-of-Order execution introduces serious architecture challenges:

- **WAR and WAW Hazards**. If later instructions write back before earlier ones:
 - ✗ They might overwrite data that's still needed.
 - ✓ Hardware **must detect and prevent** these scenarios.

This is typically handled using: register renaming and scoreboard / reservation stations.

- **Imprecise Exceptions**. If an exception occurs (e.g., divide-by-zero), but the processor has already executed and committed later instructions, the architectural state is **no longer consistent** with the point of the fault. To fix this, high-performance CPUs use:

- ✓ **Reorder Buffers (ROB)** to store results until it's safe to commit them in program order.

- ✓ **Checkpointing and rollback mechanisms** to recover precise state.

Formally, an **exception is imprecise** occurs if the processor state when an exception is thrown does not look exactly as if the instructions were executed in order.

Feature	Description
Issue	In order
Execution	Out of order
Completion (write-back)	Out of order
Architectural State Updates	May occur out of order (but real designs use re-order buffers to enforce in-order commit)
ILP	High
Complexity	High - needs hazard detection, renaming, ROBs
Exceptions	Risk of imprecision without commit logic

Table 9: Summary of Multi-Cycle Out-of-Order Pipeline.

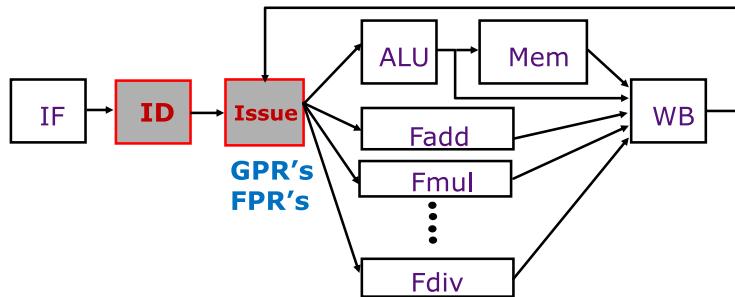


Figure 21: High-Level Multi-Cycle Out-of-Order Pipeline architecture.

- IF (Instruction Fetch): fetches the next instruction in program order from instruction memory using the program counter (PC).
- ID (Instruction Decode). This stage is now split into two sub-stages:
 1. ID: decoding the instruction format and operation type.
 2. Issue: reading registers and checking availability of operands.

This split is key to **preparing instructions early**, even if they're not ready to execute immediately.

- Functional Units. The processor has **multiple independent execution units**, each possibly multi-cycle and with different latencies:
 - ALU: used for integer arithmetic, logic, and takes 1-2 cycles.
 - Mem: used for load/store, with cache hits/misses, and takes a variable number of cycles.
 - Fadd: used for floating-point addition, and takes ≈ 3 cycles.
 - Fmul: used for floating-point multiplication, and takes ≈ 3 cycles.
 - Fdiv: used for floating-point division, and takes multiple cycles and is not always pipelined.

Each unit operates independently, and several can be active at once.

- GPRs and FPRs (General/Floating-Point Registers): architectural registers where results are ultimately written. But in this pipeline, results may first go to **temporary storage** until committed (through not shown here explicitly, concepts like **reorder buffer (ROB)** are implied).

Unlike the in-order pipeline, there is **no single commit point** shown. This means out-of-order commit, and introduces the risk of: WAR and WAW hazards, and imprecise exceptions.

3.3 Dynamic Scheduling

As we've seen, static in-order pipelines are limited in their ability to exploit ILP because they stall the entire pipeline when a single instruction is blocked. **Dynamic Scheduling** solves this by allowing **instructions to be issued, executed, and even completed out of order**, as long as doing so does not violate program correctness.

⌚ The Need for Dynamic Scheduling

Consider the following instruction sequence:

```

1 I1: F0 ← F2 / F4    # long-latency divide
2 I2: F10 ← F0 + F8   # depends on F0
3 I3: F12 ← F8 - F14  # independent of I1 and I2

```

In a naive in-order pipeline:

- I2 stalls waiting for F0, and
- I3 stalls behind I2, even though it's independent.

This results in lost parallelism. In contrast, **dynamic scheduling** would:

- Allow I1 to begin and proceed through the divider.
- Stall I2 because it depends on F0.
- Allow I3 to **proceed and complete** immediately, despite the stall.

This is possible because the processor can **track operand availability** and issue instructions **as soon as dependencies are satisfied**, not based solely on their program order.

🛠 How Dynamic Scheduling Works

Instructions are issued in order but may **execute and complete out of order**, depending on operand readiness and unit availability. The processor uses dedicated **hardware structures** to manage this:

- Reservation Stations (or Issue Queues)
- Reorder Buffer (ROB)
- Register Renaming Tables
- Common Data Bus (CDB) for broadcasting results

In a dynamically scheduled pipeline, stages are typically:

- **Fetch (IF)**: Get instruction from memory.
- **Decode (ID)**: Determine opcode, operands, destination.
- **Issue**: Place instruction into reservation station if operands aren't ready.
- **Execute (EX)**: Start when all operands are available.

- **Write Result:** Write result to a temporary buffer or broadcast to waiting instructions.
- **Commit:** Update architectural registers in order.

✓ Benefits

- ✓ **Higher ILP:** Instructions don't wait unnecessarily.
- ✓ **Resource Utilization:** Keeps functional units busy.
- ✓ **Latency Hiding:** Tolerates cache misses, long FP ops.
- ✓ **Exploits Independence:** Independent ops no longer block one another.

⚠ Challenges Introduced

While powerful, dynamic scheduling is **complex**:

- ✗ **WAR and WAW Hazards.** With out-of-order execution, later instructions might write before earlier ones. Solution: use register renaming to remove name dependencies.
- ✗ **Imprecise Exceptions.** If a later instruction causes an exception, but earlier ones have already modified state, it becomes hard to roll back. Solution: use a Reorder Buffer (ROB) that holds results temporarily and commits them in program order.
- ✗ **Hardware Cost and Complexity.** Additional logic is needed for:
 - Dependency tracking
 - Wakeup and select logic
 - Common data bus broadcasting
 - Instruction window buffering

3.4 Multiple-Issue Processors

3.4.1 Introduction to Multiple-Issue Pipelines

In previous sections, we explored how pipelining (section 3.2) and dynamic scheduling (section 3.3) help improve instruction throughput by exploiting instruction-level parallelism (ILP). However, traditional scalar **pipelines are fundamentally limited**: they can **issue only one instruction per clock cycle**. To overcome this limitation and achieve even higher performance, computer architects developed Multiple-Issue Processors.

Definition 1: Multiple-Issue Processors

Multiple-Issue Processors are processors designed to fetch, decode, issue, and execute **more than one instruction per clock cycle**, with the goal of increasing instruction throughput and exploiting instruction-level parallelism (ILP).

They achieve higher performance than scalar processors by issuing multiple independent instructions in parallel, using either hardware-based dynamic scheduling (as in superscalar architectures) or compiler-driven static scheduling (as in VLIW architectures).

This section introduces the key principles of multiple-issue pipelines and lays the foundation for understanding both superscalar and VLIW architectures.

⚠ The Limits of Scalar Pipelines

In a **scalar pipeline**, **only one instruction is issued and completed per clock cycle**, even if other instructions are independent and could be executed in parallel.

Let's consider a classic 5-stage pipeline:

$$\text{IF} \rightarrow \text{ID} \rightarrow \text{EX} \rightarrow \text{MEM} \rightarrow \text{WB}$$

In an ideal case, the pipeline achieves an IPC (Instructions Per Cycle) of 1. That is:

- 1 instruction finishes per cycle.
- Corresponding CPI (Cycles Per Instruction) is also 1:

$$\text{CPI}_{\text{ideal}} = 1, \quad \text{IPC}_{\text{ideal}} = 1$$

But in reality, hazards (data, control, structural) can cause stalls and the IPC can fall below 1. As we have already discussed in the previous sections.

⚠ Key Limitation: Even if the program contains many **independent instructions**, the scalar **pipeline processes them sequentially**, one at a time.

⌚ Raising Performance: Introducing Multiple Issue

To extract more parallelism and achieve better throughput, multiple-issue processors aim to:

- Fetch **multiple instructions** per cycle.
- Issue and execute **multiple instructions** in parallel.
- Increase **IPC above 1 ↑** and reduce **CPI below 1 ↓**

This means:

- The processor is no longer limited by the sequential issue constraint.
- ILP is exploited across multiple instructions simultaneously.

A simple example is the **dual-issue pipeline**, where up to two instructions can be issued and completed per clock cycle. It allows two independent instructions to proceed through the pipeline in parallel, potentially doubling the instruction throughput compared to a scalar pipeline:

$$\text{IPC}_{\text{ideal}} = 2 \quad \text{CPI}_{\text{ideal}} = \frac{1}{\text{IPC}_{\text{ideal}}} = \frac{1}{2} = 0.5$$

Definition 2: Dual-Issue Pipeline

A **Dual-Issue Pipeline** is a **type** of multiple-issue processor pipeline that can fetch, decode, and issue **up to two instructions per clock cycle**.

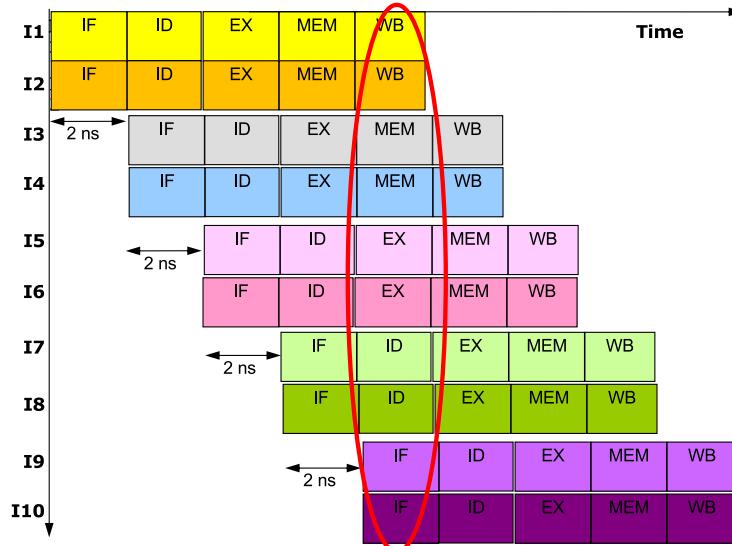


Figure 22: Dual-Issue Pipeline timeline.

■ Architectural Requirements

It's pretty obvious that multi-issues processors require more hardware resources to support parallelism:

- **Wider Instruction Fetch (IF) units:** able to fetch 2 instructions per cycle from instruction memory.
- **Parallel Instruction Decoders (IDs):** 2 independent decode units to process both instructions in parallel.
- **Multi-Ported Register File (RF):**
 - 4 Read Ports: to read up 2 source operands per instruction.
 - 2 Write Ports: to write results from both instructions simultaneously.
- **Duplicated Functional Units:** at least 2 independent units (e.g., 1 ALU or branch, 1 load/store) to allow parallel execution.

These **additions increase complexity, area, and power consumption**, but allow significant performance gains.

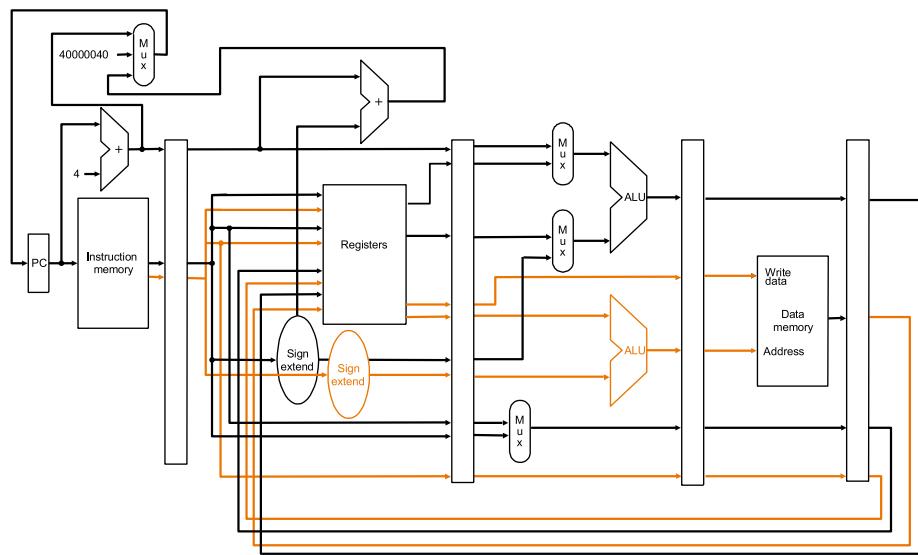


Figure 23: Dual-Issue Pipeline architecture.

3.4.2 Evolution Towards Superscalar Execution

The transition from simple scalar pipelines to high-performance superscalar processors is not abrupt. It is the result of a **progressive refinement** of microarchitectural techniques aimed at exposing and exploiting more Instruction-Level Parallelism (ILP). This section traces the key evolutionary steps that bridge the gap between single-issue scalar designs and fully dynamic multiple-issue architectures.

☒ Step 1: Single-Issue, In-Order Execution

This is the traditional scalar baseline and was explained in the earlier sections (section 3.2.1, page 81).

- Only **one instruction** is **issued** and **committed** per clock cycle.
- All instructions are fetched, decoded, executed, and written back in **strict program order**.
- Hazards (data, structural, control) cause pipeline stalls that affect all subsequent instructions.
- CPI ≥ 1 , IPC ≤ 1 .

This model is **simple and ensures precise state at all times**, but it is **severely limited in ILP exploitation**.

☒ Step 2: Single-Issue, Out-of-Order Execution

To overcome unnecessary stalls caused by instruction dependencies, processors began executing instructions **out of order**, while still **issuing only one instruction per cycle** (section 3.2.2, page 84).

- Instructions are fetched and issued **in program order**.
- But **independent instructions** are allowed to **execute and complete out of order**, as soon as their operands are ready and a functional unit is available.
- Techniques like **dynamic scheduling** (e.g., Tomasulo's algorithm) are used to manage dependencies and operand availability.
- A **commit stage ensures in-order architectural updates**, preserving program correctness and exception handling.

This **significantly improves ILP**, but **throughput is still constrained by the single-issue limit**.

☒ Step 3: Multiple-Issue, In-Order Execution (Dual-Issue Pipeline)

This step involves fetching, decoding, and executing **more than one instruction per cycle**, but **in program order**.

- Typical example: **dual-issue pipelines** (e.g., MIPS dual-issue architecture, section 3.4.1, page 89).
- The **hardware allows the issue of up to two instructions per clock**, provided they are independent and compatible (e.g., one ALU + one Load/Store).
- Requires hardware additions such as:
 - Multiple functional units,
 - Multi-ported register file,
 - Hazard detection logic across simultaneously issued instructions.

This model **increases IPC (ideal IPC = 2)**, but still suffers from limitations:

- ✖ **Dependent instructions must wait**, even if others could proceed.
- ✖ **Static scheduling** (compiler) or simple hardware interlocks determine issue feasibility.

☒ Step 4: Multiple-Issue, Out-of-Order Execution

This is the **most flexible and powerful configuration**, forming the **basis of superscalar processors**.

- Fetches and decodes **multiple instructions per cycle**.
- Uses **dynamic scheduling logic** to decide which subset can be issued and executed out of order.
- Independent **instructions proceed as soon as their operands are ready**, regardless of program order.
- Results are **committed in order** to maintain a precise architectural state.
- This model requires **complex hardware**:
 - Reservation stations
 - Register renaming
 - Reorder buffer (ROB)
 - Instruction window/wakeup-select logic

This architecture provides the **highest ILP**, as it combines the breadth of multiple issue with the flexibility of out-of-order execution.

The transition toward superscalar execution is a gradual process, where each step builds upon the previous to **mitigate limitations**, **maximize resource utilization**, and **exploit greater ILP**. Superscalar architectures represent the culmination of this evolution, dynamically scheduling and executing multiple instructions in parallel, out of order, while maintaining program correctness and exception safety.

Step	Issue	Execution	Commit	Example
1	In-order	In-order	In-order	Scalar pipeline
2	In-order	Out-of-order	In-order	Dynamic single-issue
3	In-order	In-order	In-order	Dual-issue pipeline
4	In-order	Out-of-order	In-order	Superscalar processor

Table 10: Evolution towards superscalar execution.

3.4.3 Superscalar Processors

The culmination of the architectural evolution toward higher ILP is the **superscalar processor**, a class of processors capable of **issuing, executing, and committing multiple instructions per clock cycle, dynamically and out of order**. Superscalar architectures aim to exploit maximum parallelism hidden within sequential instruction streams, while preserving the illusion of sequential execution.

Definition 3: Superscalar Processor

A **Superscalar Processor** is a dynamically scheduled, multiple-issue architecture capable of issuing and executing **several instructions per clock cycle**, using complex hardware mechanisms to **detect and exploit instruction-level parallelism at runtime**.

Unlike static VLIW processors³, where instruction parallelism is exposed at compile time, **superscalar designs rely on hardware to discover and schedule parallel instructions on the fly**.

Key Characteristics

1. **Multiple-Issue Width**, also called issue width, is the **maximum number of instructions** a processor can fetch, decode, issue, and begin executing in a **single clock cycle**. It defines the **instruction throughput potential** of a multiple-issue processor.

- Therefore, this value is given by the superscalar processor, but obviously it is the maximum and it cannot be always reached (hazards, etc.).
- The maximum number of instructions per single clock cycle **affects the theoretical IPC limit**:

$$\text{CPI}_{\text{ideal}} = \frac{1}{\text{issue width}}$$

2. **Dynamic Scheduling**

- **Hardware analyzes instruction dependencies in real time**.
- **Independent instructions** are allowed to proceed **out of order**.

3. **Out-of-Order Execution**

- Instructions **execute as soon as operands and resources are ready**.
- Improves pipeline utilization and hides latencies.

³**VLIW (Very Long Instruction Word) Processor:** A type of multiple-issue processor where the compiler statically schedules multiple operations into a single wide instruction word. Each operation in the bundle is executed in parallel, assuming they are independent. VLIW architectures rely on simple hardware and place the burden of dependency checking and scheduling on the compiler rather than the processor.

4. In-Order Commit

- Despite out-of-order execution, the processor updates architectural state in program order.
- Ensures **precise exceptions** and consistent state.

Core Hardware Components

To support superscalar execution, the **architecture must include**:

- **Multiple Functional Units:** ALUs, FPUs, load/store units, enough to support parallel execution.
- **Register Renaming:** Eliminates WAR and WAW hazards by mapping architectural registers to a larger set of physical registers.
- **Reservation Stations / Issue Queues:** Buffer instructions waiting for operands or functional units.
- **Reorder Buffer (ROB):** Holds results of completed instructions until they are safe to commit. Ensures correct program order and precise exceptions.
- **Common Data Bus (CDB):** Broadcasts results to dependent instructions.
- **Instruction Window:** Sliding window of in-flight instructions from which the scheduler selects those ready to issue.

Benefits

- ✓ **High ILP:** Multiple independent instructions can be executed in parallel.
- ✓ **Dynamic Parallelism:** No need for compiler to expose ILP, hardware finds it at runtime.
- ✓ **Latency Hiding:** Long operations (e.g., cache misses, FP div) can be overlapped with other instructions.
- ✓ **General Purpose:** Works well with a wide range of programs, even without manual tuning.

Challenges

- ✗ **Complex Hardware:** Scheduling, renaming, hazard detection logic increases **area and power**.
- ✗ **Issue Logic Scalability:** The complexity of deciding which N instructions to issue out of M in-flight grows quickly.

- ✖ **Branch and Memory Dependencies:** Control and memory dependencies still limit achievable ILP.
- ✖ **Diminishing Returns:** After 3-4 issue width, real programs rarely expose enough parallelism to fully utilize the issue bandwidth.

Superscalar processors represent a powerful solution to the ILP problem, combining **multiple-issue capability** with **hardware-driven dynamic scheduling**. They dominate the high-performance general-purpose processor space (e.g., x86 and ARM cores), though their complexity and scalability limitations have motivated complementary techniques such as multithreading, vectorization, and heterogeneous architectures.

3.4.4 Static vs Dynamic Scheduling

Instruction-level parallelism (ILP) can be **exploited** either **at compile time** by the compiler or **at runtime** by the hardware. These two approaches give rise to two distinct classes of architectures:

- **Statically scheduled processors**, such as VLIW (Very Long Instruction Word) architectures.
- **Dynamically scheduled processors**, such as superscalar architectures.

This section compares these two philosophies, focusing on their mechanisms, advantages, drawbacks, and the contexts in which each is most appropriate.

Scheduling Type	Performed By	When?
Static	Compiler	At compile time
Dynamic	Hardware	At runtime

Table 11: Key difference: who does the scheduling?

⌚ What is VLIW (Static Scheduling)?

VLIW (Very Long Instruction Word) is a **statically scheduled, multiple issue processor architecture** in which the **compiler selects and packs multiple independent operations into a single, wide instruction word** that is **executed in parallel** by the processor.

The **instruction word** in a VLIW architecture consists of **multiple operations** (e.g., an ALU op, a memory op, and a floating-point op) that are **intended to be executed simultaneously**. The **compiler is responsible** for:

- Detecting ILP in the program.
- Scheduling independent instructions to avoid hazards.
- Filling empty slots with NOPs when no instruction fits.

In contrast to superscalar processors (which discover ILP dynamically at runtime), **VLIW processors rely entirely on compile-time scheduling**.

- ✓ **Simple hardware**: no need for out-of-order logic, dependency checks at runtime, or renaming hardware.
- ✓ **Predictable performance**: useful in embedded or real-time systems.
- ✓ **Energy efficient**: avoids complex runtime scheduling logic.
- ✗ **Compiler must do all the work**: requires sophisticated analysis and scheduling. So the performance depends on the quality of the compiler and the amount of visible ILP.

- ✖ **Binary compatibility issues (limited portability)**: compiled code often tied to a specific machine configuration (e.g., number of functional units). Cannot adapt to unpredictable latencies at runtime (e.g., cache misses, branch misprediction).
- ✖ **Wasted instruction slots**: when insufficient ILP is found, unused slots become NOPs, reducing efficiency.

② Superscalar: Dynamic Scheduling

In dynamically scheduled architectures:

- The compiler generates sequential code (as usual, no additional effort).
- The processor's **hardware detects ILP at runtime**, using structures like reservation stations, reorder buffers, and register renaming.
- The processor decides **which instructions to issue and execute** based on operand availability and resource status.
- ✓ Automatically adapts to **unpredictable latencies** and instruction dependencies.
- ✓ **Improved performance portability**: no need to recompile code for each variant.
- ✓ Better at **exploiting ILP in general-purpose programs**.
- ✖ **Higher hardware complexity**, area, and power consumption.
- ✖ **Scheduling logic becomes a bottleneck** at wider issue widths.
- ✖ Greater difficulty in verifying and validating timing behavior.

Feature	Static (VLIW)	Dynamic (Superscalar)
Scheduling responsibility	Compiler	Hardware
Instruction issue	Fixed and pre-planned	Determined at runtime
Flexibility at runtime	Low	High
Hardware complexity	Low	High
Compiler complexity	High	Moderate
Portability of compiled code	Low (machine-dependent)	High
Latency tolerance	Poor (fixed schedule)	Good (adaptive execution)
ILP exploitation	Only what compiler exposes	Also includes dynamic/hidden parallelism

Table 12: Static vs Dynamic Scheduling.

Static and dynamic scheduling represent two fundamentally different approaches to exploiting ILP.

- **VLIW** is ideal for **predictable workloads**, embedded systems, or domain specific processors, where simplicity and determinism matter.
- **Superscalar processors** excel in **general-purpose computing**, where **dynamic behavior and runtime variability** make hardware-managed scheduling more effective.

Ultimately, both models aim to improve throughput, but the trade-off between hardware complexity and compiler sophistication defines their respective domains of success.

3.5 ILP Limitations & Alternatives

Instruction-Level Parallelism (ILP) has been the cornerstone of high performance processor design for decades. Techniques such as pipelining, multiple-issue architectures, dynamic scheduling, and register renaming have pushed ILP to impressive levels. However, **ILP alone has fundamental limits**, both theoretical and practical, which restrict its scalability and efficiency in modern workloads. This section explores **why ILP hits a wall**, and how architects are moving toward **complementary and alternative forms of parallelism**, including multithreading, SIMD, and heterogeneous computing, to sustain performance growth.

⚠ Limitations of ILP

1. Limited Parallelism in Programs

- Many programs are inherently sequential in logic (e.g., control-intensive code, algorithms with tight dependencies).
- Available ILP is often limited to short instruction windows. An **instruction window** is the set of instructions that the processor can see and analyze at a give time to find parallelism. Due to hardware constraints (area, power, timing), the instruction window usually holds a few dozen to a few hundred instructions.
It means that even if the entire program contains parallelism, the processor can only exploit what it sees in its current instruction window. This is a practical limit.
- **Amdahl's Law bounds the speedup achievable by parallel execution of a sequential program.**⁴ Even if we have unlimited issue width, perfect branch prediction, and ideal memory, this law says that we can't speed up the parts of the program that are inherently sequential (e.g., control logic, data dependencies). For example, if 10% of our program is serial ($f = 0.1$), then the best speedup we can ever get is: $\frac{1}{0.1} = 10$.

⁴ **Amdahl's Law** states that the maximum theoretical speedup of a program is limited by the fraction of the program that must be executed sequentially, even if the rest can be infinitely accelerated or parallelized. Let:

- S : speedup
- f : fraction of the program that is serial (cannot be parallelized)
- $(1 - f)$: fraction that is parallelizable
- p : speedup of the parallel portion (e.g., number of processors)

Then:

$$\text{Speedup}(p) = \frac{1}{f + \frac{(1-f)}{p}}$$

And the maximum speed (when $p \rightarrow \infty$) is:

$$\text{Speedup}_{\max} = \frac{1}{f}$$

Even with infinite hardware resource, the speedup is limited by the serial part of the code.

2. Dependency Constraints

- True data dependences (RAW) cannot be bypassed or parallelized because the value doesn't exist yet. No matter how many cores, execution units, or parallel tricks we have, if one instruction computes a value that another must use, the second must wait for the first to finish.
- Some instructions must wait for preceding results, creating **bubbles** in the issue pipeline.

3. Control Dependencies and Branches

- Branches disrupt instruction flow.
- Even with speculative execution and prediction, **misppeculation causes flushes and wasted cycles**.

4. Memory Latency and Aliasing

- Cache misses introduce long, unpredictable delays.
- **Memory dependencies** (e.g., between loads and stores) are difficult to resolve safely at runtime, limiting aggressive scheduling.

5. Hardware Complexity and Power

- The **logic** needed for dependency checking, wakeup-select, register renaming, and instruction window scaling **grows rapidly**.
- Superscalar processors beyond 4-6 issue width become infeasible to scale due to **power, area, and control path complexity**.

⌚ Alternative Forms of Parallelism

1. **Thread-Level Parallelism (TLP) - Multithreading.** Execute **multiple threads** in parallel on the same core or across multiple cores. TLP hides long-latency events (e.g., cache misses) by switching to ready threads.
2. **Data-Level Parallelism (DLP) - SIMD and Vectorization.** Exploits uniform operations over data arrays (e.g., matrix ops, DSP, graphics). Single Instruction, Multiple Data (SIMD), one instruction operates on multiple data elements.
3. **Heterogeneous Computing - Specialized Accelerators.** Use of domain-specific architectures (DSAs) optimized for specific tasks: GPU, TPUs, NPUs, FPGAs. Offload compute-intensive or parallel workloads from the CPU to accelerators.

3.6 Scoreboard: Dynamic Scheduling Algorithm

3.6.1 Assumptions and Architecture

The **Scoreboard** is a **dynamic scheduling** mechanism introduced in the [CDC 6600](#) that enables **out-of-order execution** while maintaining program correctness. It coordinates the flow of instructions in a way that allows independent instructions to execute in parallel, despite pipeline stalls caused by data or structural hazards. This approach was fundamental to enhancing instruction-level parallelism (ILP) without relying on complex compiler-level optimizations.

Assumptions of the Scoreboard Model

To analyze the behavior of the scoreboard, it's crucial to understand the initial architectural assumptions:

- **Single-Issue Processor:** only one instruction can be fetched and issued per cycle, enforcing a serialized dispatching model despite the internal parallelism.
- **In-Order Issue:** instructions are issued in the program order (page 80). However, once issued, they are allowed to execute and complete out-of-order depending on operand availability.
- **No Forwarding Mechanism:** unlike Tomasulo's algorithm, which allows results to be forwarded from functional units directly to waiting instructions, the Scoreboard lacks this feature. Operands are only considered available once written back to the Register File (RF).
- **Multiple Pipelined Functional Units (FUs):** the architecture assumes the presence of multiple pipelined FUs, e.g. floating-point add, multiply, divide, and integer units; each with potentially variable latency.
- **Latency-Aware Execution:** both the **Execution Stage (EX)** and **Memory Access Stage (ME)** are allowed to span multiple cycles depending on the operation type and cache behavior.
- **Out-of-Order Execution and Commit:** execution and result write-back (or commit) can happen out-of-order, introducing hazards such as:
 - Write After Write (WAW)
 - Write After Read (WAR)

These are especially critical since there's no register renaming mechanism (page 77) to avoid false dependencies.

This configuration **allows** the scoreboard to **bypass pipeline stalls** by executing independent instructions out-of-order, while relying on a **centralized control logic to track hazards and resource usage**.

■ Architectural Scheme

The Scoreboard orchestrates execution by separating three phases:

1. **Instruction Issue (in-order)**
2. **Instruction Execution (out-of-order)**
3. **Instruction Completion (out-of-order)**

This setup breaks the rigid in-order pipeline flow and increases functional unit utilization. Furthermore:

- Multiple instructions can be in execution simultaneously.
- Precise exceptions are **not** guaranteed due to the possibility of earlier instructions committing after later ones, a model referred to as **imprecise interrupts**.

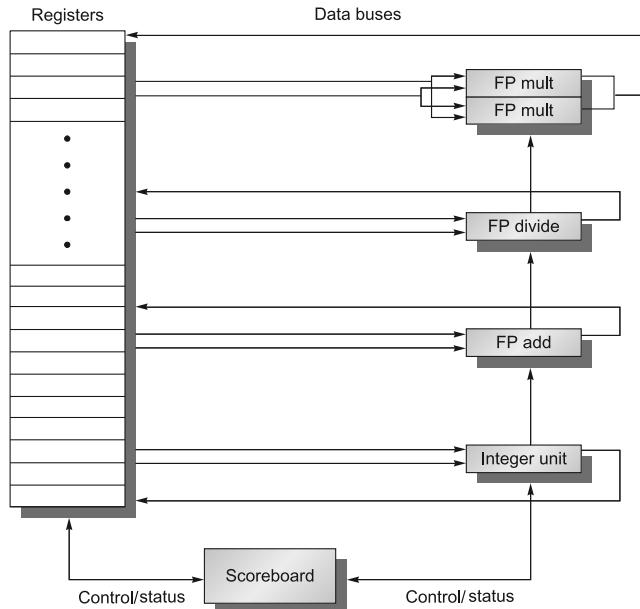


Figure 24: The basic structure of a RISC V processor with a scoreboard. [2]

- A shared Register File feeds data into multiple data buses.
- Each Functional Unit (FU) is independently pipelined and connected to the scoreboard. Units include two FP multipliers, FP adder, FP divider, and integer unit.
- A centralized Scoreboard logic block maintains: Control/Status signals, Dependency tracking, Issue constraints.
- There's a separate Memory Unit, handled similarly to functional units, responsible for data memory operations.

3.6.2 Pipeline Stage Refinement

In a **traditional pipeline**:

- The **Instruction Decode (ID)** stage performs both decoding and operand reading at once.
- But this **assumes operands are always ready**, which is **not true** in a dynamically scheduled out-of-order pipeline.

So the **scoreboard splits Instruction Decode (ID) stage into**:

1. Issue Stage

- Responsible for **decoding the instruction**.
- Checks for **structural hazards** (page 21), particularly whether the appropriate functional unit (FU) is available and whether the scoreboard's bookkeeping allows the instruction to proceed (this will become clearer later).
- Enforces **in-order issue**, instructions are considered strictly in the sequence fetched from memory.

2. Read Operands Stage (RR)

- Waits until **operands are available and not blocked** by earlier instructions.
- Specifically, **avoids RAW** (Read After Write) hazards (page 21) by deferring operand reads until the register is no longer "reserved" by an active instruction writing to it. In other words, delays reading operands **until they're truly ready**, that is, the producer instruction has completed writing them.
- **Operands** are then **read from the register file** (since forwarding is not available).

This separation increases the scoreboard's ability to exploit Instruction-Level Parallelism (ILP) while maintaining control over dependency tracking and avoiding illegal hazards.

✓ Flexible Execution Behavior

After the two front-end stages:

- **Out-of-Order Execution**: Once operands are read, instructions may enter the execution stage as soon as the corresponding FU is available, regardless of program order.
- **Variable Latency Handling**: Functional units (FUs) may have different latencies (e.g., FP divide vs. add), so instructions finish execution at different times and write back their results out-of-order.
- **Out-of-Order Commit**: Because instructions complete independently and there is **no reorder buffer**, the commit (or write-back) stage is also **out-of-order**, unlike more modern precise pipelines.

Stage	Behavior	Order Enforcement
Issue	Decode, FU check	In-Order
Read Operands	Wait for availability	Out-of-Order
Execute	Run in FU	Out-of-Order
Write Result	Commit to Reg. File	Out-of-Order

Table 13: Key features of the scoreboard.

⌚ What is the scoreboard constantly tracking?

We can think of the scoreboard as a “Control Office” inside the processor. The main job of the scoreboard is to **keep track of every instruction** that’s in the pipeline at the same time, and **make sure they don’t mess each other up**.

To do this, it monitors four key things:

1. **Availability of Source Operands.** Every instruction needs to read its inputs (like F2, F4, etc.). The scoreboard checks: **are those registers ready, or is another instruction still going write them (busy)?** If they’re **not ready** yet, the **instruction waits** in the Read Operands (RR) stage.
⌚ Why? To avoid **RAW** (Read After Write) hazards, reading too early before the data is correct.
2. **Status of each Functional Unit (FU).** It knows which **units** (like the adder, multiplier) are **busy** or **free**. It won’t assign two instructions to the same FU at the same time, that would **cause structural hazards**.
⌚ Why? So it knows **which instruction can be issued** and which needs to **wait for hardware**.
3. **Pending writes and register conflicts.** If two instructions plan to write to the same register, it keeps track of this. This **helps prevent**:
 - ✓ **WAW** (Write After Write): two instructions writing to the same register in the wrong order.
 - ✓ **WAR** (Write After Read): an instruction overwriting a value that another one still needs to read.**⌚ Why?** This avoids **wrong results**, even if instructions execute out-of-order.
4. **Which instructions have completed.** It **tracks** when each instruction finishes **execution** and when it’s **allowed to write back** its result.
 Remember: there’s **no reorder buffer**, so the scoreboard must **carefully manage write-backs to prevent conflicts**.
⌚ Why? So it knows when to **release resources** and **update registers safely**.

3.6.3 Hazard Management (RAW, WAR, WAW)

A hazard occurs when the pipeline execution of instructions might lead to incorrect results. Hazards arise from:

- Resource conflicts
- Data dependencies
- Instruction ordering mismatches

The scoreboard handles three types of data hazards dynamically:

RAW (Read After Write) - True Dependency

Occurs when:

- Instruction B tries to **read** a register before instruction A has **written** its result.

Scoreboard handling:

- The scoreboard stalls **Instruction B** in the **Read Operands (RR)**⁵ stage until **Instruction A** completes its write.

Handled in the RR (Read Operands) stage

WAR (Write After Read) - Anti-Dependency

Occurs when:

- **Instruction A** needs to **read** a register before **Instruction B** **overwrites** it.

Scoreboard handling:

- The scoreboard stalls **Instruction B** in the **Write Result** stage until **Instruction A** has read the register.

Handled in the WR (Write Result) stage

WAW (Write After Write) - Output Dependency

Occurs when:

- Two instructions write to the same register in the wrong order.

Scoreboard handling:

- The scoreboard stalls **the second instruction** in the **Issue** stage until the first instruction has written its result.

Handled in the Issue stage (or sometimes delayed to Write stage)

⁵Note that the Read Operands (RR) stage is the second stage in the Instruction Decode stage. Its purpose is to wait until all source operands are available and not blocked by an active instruction. Only then can the instruction safely read its operands from the register file.

So Scoreboard solves WAR/WAW **explicitly via stalls**, instead of using *register renaming*. This makes scoreboard simpler, but **limits how much parallelism it can safely exploit** compared to more modern approaches.

Hazard	Cause	Scoreboard Action	Handled In
RAW	Read before prior write	Stall reader until value ready	Read Operands
WAR	Write before earlier read	Stall writer until read completes	Write Result
WAW	Write before earlier write	Stall issuer	Issue

Table 14: Hazards managed by the scoreboard.

3.6.4 Control Logic and Stages

The Scoreboard architecture divides instruction **execution into four dynamic control stages**, each governed by centralized logic. These stages (**Issue**, **Read Operands**, **Execution**, and **Write Result**), replace the traditional ID, EX, and WB stages of a standard RISC pipeline.

The key idea is that the scoreboard **monitors dependencies and structural hazards in hardware** and makes real-time decisions on when each instruction can safely advance to the next stage.

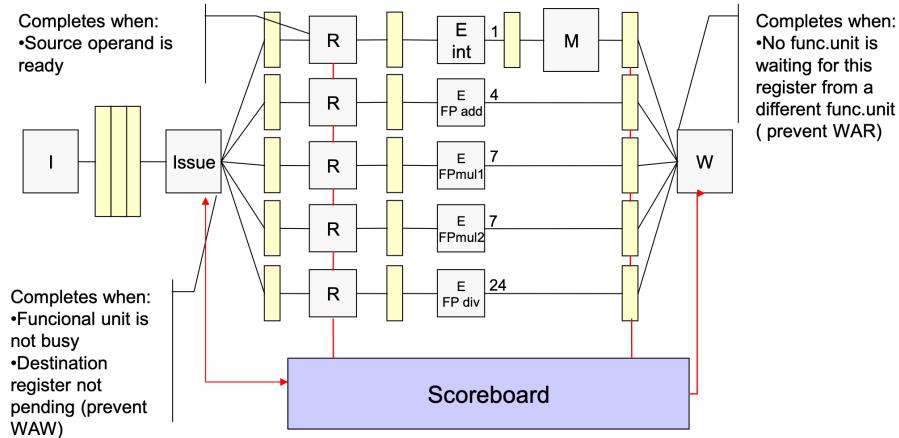


Figure 25: The Scoreboard architecture. Pipeline Flow is: Instruction Fetch (I), Read Operands Blocks (R), Execution Units (E), Memory Stage (M) and Write Result Stage (W).

1. **Issue Stage (In-Order)**. This is the first stage after Instruction Fetch (IF).
 - (a) The instruction is **decoded**.
 - (b) The scoreboard **checks**:
 - i. **Structural Hazards**: *is the required Function Unit (FU) available?*
 - ii. **WAW (Write After Write) Hazards**: *is another instruction already writing to the destination register?*

If **neither hazard exists**, the **instruction is issued** and marked in the scoreboard's internal tables. **Otherwise**, the **instruction stalls**.

Performance Optimizations: **WAW hazards** are typically **detected** here, but optimizations may postpone this to the **write-back stage**.

2. **Read Operands Stage (Out-of-Order)**. The scoreboard waits for both source operands to become available.

- (a) **RAW** (Read After Write) hazards are **checked dynamically**.
- (b) If any **operand** is still **pending** (i.e. will be written by another instruction), the **scoreboard stalls this stage**.
- (c) Once **ready**, operands are **read from the Register File (RF)** (no forwarding!).
- (d) The instruction is then sent to the Functional Unit (FU) to begin execution.

This is the **stage that enables out-of-order execution**, as independent instructions may pass each other based on operand readiness.

3. Execution Stage

- (a) The instruction **executes** in the assigned Functional Unit.
- (b) Functional Units (FUs) may have **variable latency**, depending on the operation (e.g., divides vs. add).
- (c) Upon completion, the **unit signals the scoreboard**.

This phase also includes **additional memory access latency for load-store instructions** affected by cache hit/miss.

4. Write Result Stage (Out-of-Order)

- (a) Before writing the result to the destination register, the **scoreboard checks** for:
 - **WAR** (Write After Read) hazards: *is any previous instruction still waiting to read this register?*
 - **Structural hazards**: *are the Register File (RF) write ports available?*

If **clear**, the **instruction writes back** the result. If a **WAR hazard exists**, the **scoreboard stalls this instruction** until the reading instruction completes.

Stage	Hazard Checked	Order	Description
Issue	Structural, WAW	In-Order	FU availability + dest reg conflict
Read Operands	RAW	Out-of-Order	Waits for source operands
Execute	-	Out-of-Order	Runs in FU (latency varies)
Write Result	WAR, Structural (RF)	Out-of-Order	Writes result if safe

Table 15: Hazards managed by the scoreboard.

3.6.5 Summary

In this section, we present a summary of the Scoreboard's dynamic scheduling algorithm.

The Scoreboard implements a **classic dynamic scheduling mechanism** where instruction progress through the pipeline is dictated not just by structural availability, but also by **true data readiness**.

Pipeline Stage	Order	Hazard Checked	Notes
Issue	In-order	Structural, WAW	Instruction decoded, FU reserved
Read Operands	Out-of-order	RAW, structural (RF ports)	Wait for all inputs to be ready
Execution	Out-of-order	-	Variable latency depending on FU
Write Result	Out-of-order	WAR, structural (write port)	Write to reg file if safe

The scoreboard enforces **precise tracking** at each stage to dynamically resolve hazards without register renaming or forwarding.

Execution Properties

- **In-Order Issue**

- Simplifies the hardware: instructions are always issued in program order.
- Helps detect **WAW hazards** early.

- **Out-of-Order Read Operands**

- Once issued, instructions **wait until all operands are available**, then read them from the Register File (RF).
- No data forwarding! **Operands are read only from the Register File (RF)**.
- Allows **independent instructions** to leapfrog stalled ones.

- **Out-of-Order Execution**

- Instructions execute as soon as their operands are ready and the FU is free.
- Multiple instructions can execute **simultaneously** in parallel FUs or pipelined units.
- Leads to **higher FU utilization and throughput**.

- **Out-of-Order Completion**

- Results are written back when ready, **unless a WAR hazard** is detected.
- This breaks precise exception semantics, i.e., exceptions can be **imprecise**.

☒ No Forwarding, No Renaming

- **No data forwarding:** causes **extra stalls** at operand read stage.
- **No register renaming:** makes the scoreboard **vulnerable to WAR and WAW hazards**, which it **handles by stalls and centralized checks**.

☒ Control Logic Centralization

All control decisions (hazard detection, operand availability, resource usage) are made by a **central scoreboard table**. This avoids complex distributed hardware (as in Tomasulo), but limits the potential for speculation or aggressive scheduling.

3.6.6 Scoreboard Data Structures

At the heart of the scoreboard's centralized control logic are **three hardware data structures** that track the status of instructions, functional units, and register dependencies. These structures allow the scoreboard to make safe, real-time decisions about instruction scheduling, execution, and result writing, all while avoiding hazards.

1. **Instruction Status Table.** Tracks the **lifecycle of each instruction** through the pipeline. For each instruction, the scoreboard record whether it has:

- Been **issued**
- **Read operands**
- **Completed execution**
- **Written back** the result

We can think of this as a per-instruction timeline: it tracks which stage the instruction is currently in.

2. **Functional Unit Status Table.** Tracks the **current state** of each **Functional Unit (FU)**. Each FU entry includes:

- **Busy:** whether the FU is currently in use.
- **Op:** operation being performed (e.g., ADD, MULT).
- **Fi:** destination register of the operation.
- **Fj, Fk:** source register operands.
- **Qj, Qk:** functional units producing **Fj** and **Fk**.
- **Rj, Rk:** boolean flags indicating if **Fj, Fk** are ready.

These fields help the scoreboard:

- (a) Decide when operands are ready (for RAW)
- (b) Prevent WAW and WAR
- (c) Handle operand read scheduling

3. **Register Result Status Table.** Tracks **which FU will produce each register value**. For each register (e.g., F0, F2, ..., F30), it stores:

- The **name of the FU** that will write to it.
- Or blank (–, don't care) if no instruction is scheduled to write it.

This structure is essential to:

- Detect WAW hazards at **issue** stage.
- Detect WAR hazards at **write-back** stage.
- Ensure only the latest producing instruction claims the register.

Example 6

Let's say MULT f0, f2, f4 is issued to Mult1 (functional unit). The scoreboard will:

1. Mark Mult1 as Busy
2. Set
 - Op = MULT
 - Fi = F0
 - Fj = F2
 - Fk = F4
3. Fill Qj and Qk if other FUs are writing F2 or F4
4. In the **Register Result Status**, assign F0 = Mult1

This coordination ensures:

- Other instructions know F0 will be produced by Mult1
- F2 and F4 are only read when available
- Subsequent instructions that depend on F0 will wait

3.6.7 In-Depth Execution Example

The goal of this section is to observe how the scoreboard manages dependencies, tracks resource usage, and handles all threats over time using an example.

★ Initial Setup: Instruction List and Dependencies

The instructions are:

```

1 LD      F6, 34(R2)
2 LD      F2, 45(R3)
3 MULTD  F0, F2, F4      # RAW on F2
4 SUBD  F8, F6, F2      # RAW on F6, F2
5 DIVD  F10, F0, F6     # RAW on F0, F6
6 ADDD  F6, F8, F2      # WAW & WAR on F6, RAW on F8 & F2

```

We have a mix of:

- RAW hazards: F2, F6, F0, F8
- WAW/WAR: around register F6

During the example, we will show the status of three main hardware data structures introduced in the section 3.6.6, page 113:

- **Instruction Status Table:** tracks the lifecycle state of each instruction in the pipeline.
- **Functional Unit Status Table:** tracks the usage and readiness of each functional unit (FU), and the dependency state of the operands.
- **Register Result Status Table:** tracks which FU is **scheduled to write** to each floating-point register.

1. Cycle 1

- LD F6, 34(R2) is issued and begins execution.
- Integer unit is now **busy**.
- All other instructions wait.

No **hazard** yet. This sets up the first data dependency (F6 will be written soon)

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1			
LD F2, 45(R3)				
MULTD F0, F2, F4				
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	Yes	Load	F6		R2				Yes
	Mult1	No								
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
1				Integer					

Register result status.

2. Cycle 2

⚠ Cannot issue LD F2, 45(R3) yet, due to **structural hazard** on the Integer unit.

- Execution of first LD continues.

Stall due to structural hazard, despite in-order issue.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1		2	
LD F2, 45(R3)				
MULTD F0, F2, F4				
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	Yes	Load	F6		R2				Yes
	Mult1	No								
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
2				Integer					

Register result status.

3. Cycle 3

- First load finishes execution.
- Still can't issue second load.

Memory latency is ideal (1 cycle), but scoreboard doesn't allow overcommit of the integer unit.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2		3
LD F2, 45(R3)				
MULTD F0, F2, F4				
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	Yes	Load	F6		R2				Yes
	Mult1	No								
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
3				Integer					

Register result status.

4. Cycle 4

- F6 is written to the register file (RF)
- Integer unit is now **free**.

Register F6 becomes available, other instructions depending on it can now move (when their turn comes).

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)				
MULTD F0, F2, F4				
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	No								
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
4				Integer					

Register result status.

5. Cycle 5

- LD F2, 45(R3) is issued and starts execution.
- Instruction 2 enters pipeline, finally.

No data hazards here, but we're about to enter the **RAW jungle** starting next cycle.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)		5		
MULTD F0, F2, F4				
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	Yes	Load	F2		F3				Yes
	Mult1	No								
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
5		Integer							

Register result status.

6. Cycle 6

- ✓ MULTD F0, F2, F4 is issued to Mult1.
- ✗ **RAW hazard** on F2: operand not yet available (second load not completed).
- MULTD waits in the Read Operands stage.
- LD F2 is executing (started in Cycle 5).

Highlights:

- ✓ MULTD is issued because: the Mult1 functional unit is free (no structural hazard), and no other instruction is writing to F0 (no WAW hazard).
- ✗ But execution (of MULTD) is blocked because:
 - F2 (a source operand) is still being loaded by LD F2, and this is a RAW (Read After Write) hazard.
 - MULTD must wait until LD F2 writes its result into register F2.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6		
MULTD F0, F2, F4	6			
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	Yes	Load	F2		F3				Yes
	Mult1	Yes	Mult	F0	F2	F4	Integer		No	Yes
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
6	Mult	Integer							

Register result status.

7. Cycle 7

- ✓ LD F2 completes (data cache hit), result will be written to F2 next cycle.
- ✓ SUBD F8, F6, F2 is issued to Add unit (free).
- ✗ But SUBD cannot start execution yet:
 - It needs F6 and F2 as operands.
 - F2 just finished loading and is **not yet written back**, so still blocked (**RAW hazard**).
 - F6 was written earlier by LD F6 and is now available.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	
MULTD F0, F2, F4	6			
SUBD F8, F6, F2	7			
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	Yes	Load	F2		F3			Yes	
	Mult1	Yes	Mult	F0	F2	F4	Integer		No	Yes
	Mult2	No								
	Add	Yes	Sub	F8	F6	F2		Integer	Yes	No
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
7	Mult	Integer			Add				

Register result status.

8. Cycle 8

- ✓ DIVD F10, F0, F62 is issued to Divide unit (free).
- ✓ LD F2 writes back to F2, now MULTD and SUBD can read F2 (cycle 9).
- ✗ But execution is blocked because:
 - It needs F0 (being written by MULTD), so **RAW hazard** on F0.
 - It also uses F6, which is already ready.
 - F2 just finished loading and is **not yet written back** (this cycle), so still blocked (**RAW hazard**).

This stage shows **cascading dependency chains** forming: DIVD waits for MULTD, which waits for LD F2.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6			
SUBD F8, F6, F2	7			
DIVD F10, F0, F6	8			
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
	Add	Yes	Sub	F8	F6	F2	Integer	Yes	Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
8	Mult			Add	Divide				

Register result status.

9. Cycle 9

- ✓ MULTD and SUBD **read operands** (in parallel): scoreboard uses a **multi-port register file** (e.g., 4 read ports) to allow simultaneous operand reads.
- MULTD begins execution on Mult unit (10-cycle latency, table “Functional unit status”).
- SUBD begins execution on Add unit (2-cycle latency).
- ✗ ADDD cannot be issued because the Add unit is **already in use** by SUBD (structural hazard).

Scoreboard enables **parallel out-of-order read** and execution.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9		
DIVD F10, F0, F6	8			
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
10	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
2	Add	Yes	Sub	F8	F6	F2		Integer	Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
9	Mult				Add	Divide			

Register result status.

10. Cycle 10

- MULTD and SUBD are executing.
- DIVD is still waiting for F0 (from MULTD).
- ADDD remains stalled.

This block shows classic **out-of-order read and execution** behavior.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9		
DIVD F10, F0, F6	8			
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
9	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
1	Add	Yes	Sub	F8	F6	F2	Integer	Yes	Yes	
	Divide	Yes	Div	F10	F0	F6	Mult	No	Yes	

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
10	Mult				Add	Divide			

Register result status.

11. Cycle 11

- ✓ SUBD finishes execution (2-cycle latency complete).
- ✓ SUBD is now ready to write back to F8.
- ✗ But SUBD **cannot write yet**, it will occur in the next cycle.
 - MULTD is still executing (10-cycle latency).
 - DIVD continues to **wait for F0**, which is still in production by MULTD.
 - ADDD remains **stalled**, since the Add unit is busy with SUBD.

The key idea here is that the scoreboard **only allows one instruction to be written per cycle**; even completed executions must wait their turn.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9		11
DIVD F10, F0, F6	8			
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
8	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
0	Add	Yes	Sub	F8	F6	F2		Integer	Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
11	Mult				Add	Divide			

Register result status.

12. Cycle 12

- ✓ SUBD writes result to F8
- ✓ The Add unit becomes available again.
 - MULTD is still executing.
 - DIVD still waiting for F0.
 - ADDD can now potentially be issued, since the Add unit is no longer busy.

Hazard check: the **WAW hazard on F6** (from ADDD) is now **clear** because no one is writing to F6 at this moment.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2				

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
7	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
	Add	No								
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
12	Mult					Divide			

Register result status.

13. Cycle 13

- ✓ ADDD is issued to the Add unit: no structural hazard and WAW on F6 is clear.
- ✓ It waits in the Read Operands (RR) stage.
- ✗ Still blocked by a **WAR hazard on F6**:
 - DIVD is supposed to read F6, but ADDD wants to write it.
 - ADDD must wait until DIVD reads F6 to avoid overwriting it too early.
- MULTD still executing.

Classic WAR hazard: writer (ADDD) must not overwrite a register until readers (DIVD) finish reading it.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13			

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
6	Integer	No								
	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
13	Mult			Add		Divide			

Register result status.

14. Cycle 14

- ✓ ADDD reads its operands F8 and F2:
 - F8 just written by SUBD (cycle 12).
 - F2 is available since cycle 8.

✗ DIVD is still waiting on F0 (from MULTD), so hasn't read F6 yet.

- ⌚ Event though DIVD was issued before ADDD, here ADDD performs operand read first. This is out-of-order read!

Scoreboard allows out-of-order read operands when no hazards exist.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14		

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
5	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
2	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
14	Mult			Add		Divide			

Register result status.

15. Cycle 15

✓ ADDD starts execution on Add unit.

- MULTD and DIVD still wait:

- MULTD still processing.
- DIVD still cannot read F0, so it also hasn't read F6; keeping the **WAR hazard active** on F6.

Execution parallelism: ADDD and MULTD are both executing, in separate units.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14		

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
4	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
1	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
15	Mult			Add		Divide			

Register result status.

16. Cycle 16

✓ ADDD finishes execution.

✗ Cannot write back yet due to **WAR** hazard:

- ADDD wants to write to F6, but DIVD hasn't yet read it.
- DIVD is **still stalled**, waiting for F0 (result of MULTD).
- MULTD continues execution (almost done!).

Scoreboard protection: even though ADDD finished early, it must **wait** to avoid corrupting data DIVD still needs.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14	16	

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
3	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
0	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
16	Mult			Add		Divide			

Register result status.

17. Cycle 17

- Same situation:
 - ADDD is waiting to write to F6.
 - DIVD is waiting to read both F0 and F6.
- ⚠ WAR hazard on F6 still active.
- MULTD is almost done (last cycle of execution).

Note: WAR hazards delay write-back, not execution.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14		16

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
2	Integer	No								
	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
17	Mult			Add		Divide			

Register result status.

18. Cycle 18

- Again, same situation:
 - ADDD is **still waiting** to write F6 (WAR hazard).
 - DIVD is still blocked from reading F0 (pending from MULTD).
- MULTD completes execution (latency ends here)

Now the scoreboard is ready for MULTD to write its result.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9		
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14	16	

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
1	Integer	No								
	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
18	Mult			Add		Divide			

Register result status.

19. Cycle 19

- ✓ MULTD writes to F0.
- ✓ This finally allows DIVD to read operands in the **next cycle** (20).
- ⚠ ADDD still waits on WAR hazard, DIVD still hasn't read F6.

This is a key cycle: **RAW** on F0 is **resolved**, allowing DIVD to make progress at last.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9	19	
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14	16	

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
0	Integer	No								
	Mult1	Yes	Mult	F0	F2	F4			Yes	Yes
	Mult2	No								
	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6	Mult		No	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
19	Mult			Add		Divide			

Register result status.

20. Cycle 20

- ✓ DIVD reads both F0 and F6:
 - F0 from MULTD, now available
 - F6 from earlier LD.
- ✓ With both operands read, **WAR hazard on F6 is gone.**
- ✓ This **unlocks** ADDD, which can now write F6.

Scoreboard logic synchronizes dependent events: DIVD completes operand read, then ADDD can write, then hazard avoided.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9	19	20
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8			
ADDD F6, F8, F2	13	14	16	

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	No								
	Mult2	No								
	Add	Yes	Add	F6	F8	F2			Yes	Yes
	Divide	Yes	Div	F10	F0	F6			Yes	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
20				Add		Divide			

Register result status.

21. Cycle 21

- ✓ DIVD finally reads its operands:
 - F0 written by MULTD in cycle 19.
 - F6 read now, which **clears the WAR hazard** blocking ADDD.
- ✓ ADDD can now **safely write back to F6** in the **next cycle**.

Key transitions: DIVD finishes operand read (last one of the program); ADDD gets the green light to write since no instructions are waiting to read F6.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9	19	20
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8	21		
ADDD F6, F8, F2	13	14	16	

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	No								
	Mult2	No								
	Add	Yes	Add	F6	F8	F2			Yes	Yes
40	Divide	Yes	Div	F10	F0	F6			Yes	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
21				Add		Divide			

Register result status.

22. Cycle 22

- ✓ ADDD writes result to F6: this completes the final **WAW dependency** involving F6.
- DIVD is now **executing** (started after operand read).
- No structural or data hazards remain, all previous dependencies are solved.

Everything is now in-flight or completed:

- All operands have been read.
- All issued instructions are either executing or have completed.
- The scoreboard is now **idling** except for the ongoing DIVD.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9	19	20
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8	21		
ADDD F6, F8, F2	13	14	16	22

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	No								
	Mult2	No								
	Add	No								
39	Divide	Yes	Div	F10	F0	F6			Yes	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
22						Divide			

Register result status.

61. Cycle 61

✓ DIVD completes its execution.

- Recall: DIVD was issued in cycle 8, and with a long latency (40 cycles), it finally ends execution here.
- The scoreboard marks the Divide functional unit as **ready to write**.

This shows how the scoreboard tracks long-latency FUs **without blocking the pipeline**. Other instructions have long since finished.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9	19	20
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8	21	61	
ADDD F6, F8, F2	13	14	16	22

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	No								
	Mult2	No								
	Add	No								
0	Divide	Yes	Div	F10	F0	F6			Yes	Yes

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
61						Divide			

Register result status.

62. Cycle 62

✓ DIVD writes result to F10.

- All instructions have now:
 - Been **issued**
 - **Executed**
 - **Written back**

All functional units are **idle**, the pipeline is now **completely drained**.

Instruction	Issue	Read Op.	Exec Comp	Write Res
LD F6, 34(R2)	1	2	3	4
LD F2, 45(R3)	5	6	7	8
MULTD F0, F2, F4	6	9	19	20
SUBD F8, F6, F2	7	9	11	12
DIVD F10, F0, F6	8	21	61	62
ADDD F6, F8, F2	13	14	16	22

Instruction status.

Time	Name	Busy	Op	Fi	Fj	Fk	Qj	Qk	Rj	Rk
	Integer	No								
	Mult1	No								
	Mult2	No								
	Add	No								
	Divide	No								

Functional unit status.

Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
62									

Register result status.

The 62-cycle example is a demonstration of: efficient hazard management, robust scheduling logic, clean and scalable hardware coordination. The scoreboard shines in showing how **simple rules, carefully enforced, can deliver powerful out-of-order behavior** without complexity of modern speculative or superscalar techniques.

3.7 Tomasulo's Algorithm

3.7.1 Introduction

Tomasulo's Algorithm represents a pivotal innovation in the domain of dynamic scheduling and out-of-order execution within high-performance computing. Developed in 1967 at IBM for the [IBM System/360 Model 91](#), it was introduced as a means to exploit **Instruction-Level Parallelism (ILP)** in the absence of compiler support or source-level reordering. The essential goal was to **overcome pipeline stalls** due to data hazards, particularly **Write After Write (WAW)** and **Write After Read (WAR)** hazards, both of which are challenging to resolve through simple pipeline control mechanisms.

Core Idea

Tomasulo's algorithm enables **instructions to execute out of program order**, yet maintains **data correctness** via hardware-level mechanisms. Central to this approach is the concept of **implicit register renaming**, which dynamically assigns storage locations (reservation stations) to values rather than using architectural register names directly. This mechanism ensures that no two instructions mistakenly read or overwrite the same register unless there is a true data dependency (RAW, Read After Write).

New features introduced: a kind of Dynamic Scheduling 2.0

- **Implicit Register Renaming:** avoids WAR and WAW hazards by assigning intermediate results to reservation stations rather than architecture registers.
- **Dynamic Scheduling:** unlike static instruction scheduling (done at compile time), Tomasulo's algorithm uses runtime analysis to decide the order of instruction execution (yes, like the Scoreboard algorithm, but smarter).
- **Out-of-Order Execution:** instructions can be issued and begin execution as soon as operands are ready, independently of program order, provided that dependencies are resolved.
- **Common Data Bus (CDB):** results are broadcast to all units waiting for them, further enabling parallelism.

Historical Significance

Tomasulo's work appeared just three years after the CDC 6600 Scoreboarding mechanism (Seymour Cray's design), which was the first dynamic scheduling mechanism. Unlike the scoreboard, Tomasulo's algorithm **distributes control** among the **functional units** via **reservation stations**, offering more scalable and parallel data communication through the CDB.

It served as the architectural blueprint for later processors such as:

- Alpha 21264



- Intel Pentium II



- HP PA-8000



- PowerPC 604



- MIPS R10000



Tomasulo vs. Scoreboarding: What's the difference?

Both **Tomasulo's Algorithm** and the **Scoreboard Algorithm** are techniques for **dynamic instruction scheduling**, meaning they allow instructions to be executed **out of program order** while still preserving correctness. But Tomasulo goes a step further in terms of efficiency and cleverness:

Feature	Scoreboard (CDC 6600)	Tomasulo (IBM 360/91)
Register Renaming	✗ No	✓ Yes (implicit via reservation stations)
WAR/WAW Hazards Handling	✗ Needs to stall	✓ Avoided by renaming
Data Communication	✓ Writes to Register File	✓ Used CDB to forward directly
Control	Centralized scoreboard	Distributed (each FU has its own RS)
Execution Start Condition	Wait for operand <i>registers</i>	Wait for operand <i>values</i> or <i>tags</i>

Table 16: Tomasulo vs. Scoreboarding.

3.7.2 Register Renaming: Static vs. Implicit

❷ First of all, why Register Renaming?

One of the main challenges in pipelined and out-of-order execution is handling **false dependencies**. **False Dependencies** (page 76) occur when instructions appear to depend on each other because they use the same register name, but there is no true data dependency between them. These are name-related hazards, not value-related. There are two main types:

- **WAR (Write After Dependencies)**: a later instruction writes to a register that a previous instruction needs to read.
- **WAW (Write After Write)**: two instructions write to the same register, but the second one is issued before the first finishes.

These are not *true* data dependencies (like Read After Write, RAW), but **name conflicts**, where different values want to use the same register name.

❸ **Register renaming** solves this by dynamically or statically mapping registers to different physical storage locations.

☒ Static Register Renaming (Compiler-Based)

In **Static Renaming**, the compiler performs the renaming at compile time, allocating **temporary** (non-architectural) registers to break WAR and WAW dependencies.

For example:

```

1 DIV.D F0, F2, F4
2 ADD.D F6, F0, F8      ; RAW on F0
3 S.D   F6, 0(R1)       ; RAW on F6
4 MUL.D F6, F10, F8    ; WAW & WAR on F6

```

There is a **WAW hazard** (both ADD.D and MUL.D write to F6), and a **WAR hazard** (S.D reads F6 while MUL.D wants to write it). With static register renaming, the compiler assigns a new register (e.g., S) to avoid the conflict:

```

1 DIV.D F0, F2, F4
2 ADD.D S, F0, F8      ; RAW still valid
3 S.D   S, 0(R1)       ; Now reads S
4 MUL.D F6, F10, F8   ; Safe to use F6

```

Now the hazards are gone: ADD.D writes to S, which is consumed by the store; MUL.D writes to F6 independently.

⚠ **Static Register Renaming - Limitation.** Static Renaming requires **predicting all hazards** in advance and knowing the **full execution path** (hard with branches, loops, dynamic inputs, etc.). Also, it requires **many more architectural registers** to be encoded in the ISA, not always feasible.

✓ Implicit Register Renaming (Hardware-Based - Tomasulo's way)

Tomasulo's algorithm takes a smarter **dynamic** approach. Instead of using physical register names, it uses **Reservation Stations (RS)** as **Temporary Names (tags)** for operands. This renaming is done **implicitly** by the hardware at runtime.

Using the previous example:

```

1 DIV.D F0, F2, F4
2 ADD.D F6, F0, F8      ; RAW on F0
3 S.D   F6, 0(R1)       ; RAW on F6
4 MUL.D F6, F10, F8    ; WAW & WAR on F6

```

Tomasulo rewrites it with **RS identifiers**:

```

1 DIV.D F0, F2, F4
2 ADD.D RS1, F0, F8      ; ADD result goes to RS1
3 S.D   RS1, 0(R1)       ; Store reads from RS1
4 MUL.D F6, F10, F8     ; Now safe, F6 is available

```

- ADD.D **doesn't write to F6**, but to a *reservation station* named RS1.
- S.D **reads from RS1**, not from F6.
- This avoids **WAW** (two writes to F6) and **WAR** (store reads F6 before MUL.D writes).

Tomasulo's algorithm **automatically tracks which RS (Reservation Stations) produces what and only writes to the actual register file once the instruction retires**. Until then, everything is handled with RS tags.

This is the **core strength** of Tomasulo over simpler approaches like Scoreboarding or static renaming. It allows **aggressive out-of-order execution** without risking data hazards, making it fundamental in modern CPU design.

Here we have only presented the *secret sauce* of Tomasulo's power, in future sections we will gradually reveal how the tag-based mechanism works in practice.

3.7.3 Basic Concepts of Tomasulo's Algorithm

★ Goals of Tomasulo's Design

Tomasulo's algorithm was designed to solve a major performance bottleneck in pipelined processors: **pipeline stalls caused by operand unavailability** due to data hazards. The solution? Introduce a distributed, smart scheduling mechanism that:

- Avoids **WAR** and **WAW** hazards (false dependencies)
- Allows **out-of-order execution**
- Enables **register renaming implicitly**
- Uses **Reservation Stations (RSs)** and a **Common Data Bus (CDB)**

⌚ Reservation Stations (RSs): Tomasulo's Brain

Rather than having a central scoreboard (as in CDC 6600), Tomasulo distributes the **control logic and buffering close to the Functional Units (FUs)** using **Reservation Stations**.

Each functional unit (like a floating-point adder or multiplier) **has its own RSs** in front of it. These are **small buffers** that:

- **Hold instruction operands** (or *tags* pointing to where the operand will come from).
- **Wait until operands are ready.**
- **Dispatch instructions into the FU** as soon as everything is available.

This local storage of operands **removes the need to stall** the entire pipeline, **each unit become self-scheduling**.

▣ Implicit Register Renaming with RS Tags

Instead of keeping track of operand names (e.g., F2, F4, F6, etc.), **Tomasulo tracks**:

- Either the **value** of the operand (if available)
- Or the **tag** of the RS that will produce that value (if not yet ready)

This is very powerful because:

- ✓ Registers are replaced by **RS names or actual values**
- ✓ WAR and WAW hazards are **completely avoided**
- ✓ Instruction scheduling becomes **data-driven**

We no longer wait for registers, we **wait for values**, and when they're ready, we go.

Tags, RSs, and the CDB

A critical component that ties everything together is the **Common Data Bus (CDB)**:

1. When a **functional unit finishes** execution, the result is **broadcast on the CDB**.
2. Any **RS waiting for that result's tag** will grab the value and store it in its local buffer.
3. Also, the result is written to the **Register File**, but only if no newer instruction is overwriting that register.

This broadcasting mechanism allows Tomasulo to perform a kind of **hardware-level forwarding**, operands are handed off *before* they hit the register file.

Feature	Scoreboarding	Tomasulo
Operand Waiting	Wait for register	Wait for value or tag
Operand Tracking	Centralized	Distributed in RSs
WAR/WAW hazards	Cause stalls	Avoided via renaming
Communication	Implicit write-back	Broadcast over CDB
Renaming	✗ None	✓ Implicit via RS

Table 17: Compared to Scoreboarding.

Tomasulo replaces rigid, centralized scheduling with a **fluid, decentralized approach**. Reservation Stations **track availability**, **rename registers**, and **drive execution**. The Common Data Bus **broadcasts results** to all who need them.

The processor becomes **dataflow-like**: instructions execute *when their operands are ready*, not when some global scheduler says so.

3.7.4 Architecture

The classic architecture includes:

- An **Instruction Queue** that feeds instructions in program order.
- Several Reservation Stations (RSs) sitting in front of Functional Units (FUs) (like ADD, MUL, DIV units).
- A **Register File** that doesn't just store values, but also tracks **where the next value will come from**.
- A **Common Data Bus (CDB)** that broadcasts results to all RSs and the Register File.

❖ Components of a Reservation Station (RS)

Each **Reservation Station (RS)** is a mini-instruction buffer attached to a specific type of Functional Unit. It holds **1 instruction** and includes the following fields:

Field	Description
Tag (name)	Unique ID of the RS (e.g., RS1, RS2, ...)
Busy	A boolean flag: <i>is the RS holding an instruction?</i>
OP	Operation type (e.g., ADD.D, MUL.D)
V _j , V _k	Value of operands (if available)
Q _j , Q _k	Tag of RS producing V _j , V _k (if not available)

- If V_j/V_k are valid → operands are ready.
- If Q_j/Q_k are nonzero → wait for the result tagged with that RS name to arrive on the CDB.

Only **one of (V, Q)** is valid for each operand at any given time. Finally, note that for memory instructions (like Load/Store), V_j often holds the **address**, not a register value.

❖ Register File and Store Buffers

The Register File in Tomasulo's architecture does **more than just store values**. Each register contains:

- V_i , current value (if available)
- Q_i , tag of the RS that will produce this value (if pending)

All depends on the Q_i value:

- If $Q_i = 0 \rightarrow$ the value is already available \rightarrow use V_i .
- If $Q_i \neq 0 \rightarrow$ the register is waiting for an instruction to produce that value.

This tagging makes the register file an elegant interface for:

- **Implicit register renaming**
- **Hazard tracking**
- **Register freeing**

Likewise, **Store Buffers** behave like a special type of RS:

- They **wait** for both **address** and **data** to be ready **before sending it to memory**
- Both elements may depend on other RS outputs, so Store Buffers also monitor the CDB.

❖ Load/Store Buffers

Memory operations need special handling due to:

- Address calculation
- Unknown data dependencies
- Potential memory hazards

So we use:

- **Load Buffers**, consisting of two fields: **Busy**, **Address**. Calculated in two steps:
 1. Address calculation using **base register + offset**.
 2. Wait for memory unit availability to load the data.
- **Store Buffers**, wait for the **data to be stored** and the **address**. Once both are ready, send to memory.

Store Buffers act like RSs, they hold partial information and wait on operand values (tags), same as ADD or MUL units.

💡 Why this architectural design works?

- **Decouple instruction issue from operand readiness**, then out-of-order execution becomes natural.
- **Multiple RSs per FU**, then enables multiple instructions to be in-flight for the same unit type.
- **Register file doesn't need to "remember everything"**, then RSs temporarily hold data and manage dependencies.
- **Broadcast via CDB**, then all waiting units instantly get what they need, without central control.

In other words, the Tomasulo's architecture is a **beautifully decentralized, tag-driven pipeline**:

- Each **Reservation Station** acts as a micro-scheduler for its FU.
- The **Register File** manages *what is available* vs. *what is still pending*.
- **Load/Store Buffers** elegantly handle memory-side hazards and operand waits.
- The **CDB** ties it all together, broadcasting results as soon as they're ready.

This structure sets the stage for the dynamic instruction lifecycle, which we'll explore in detail in the next section.

❖ Example of Tomasulo structure

The following figure shows an example of RISC-V structure using Tomasulo. It is a very good example to understand how the architecture works.

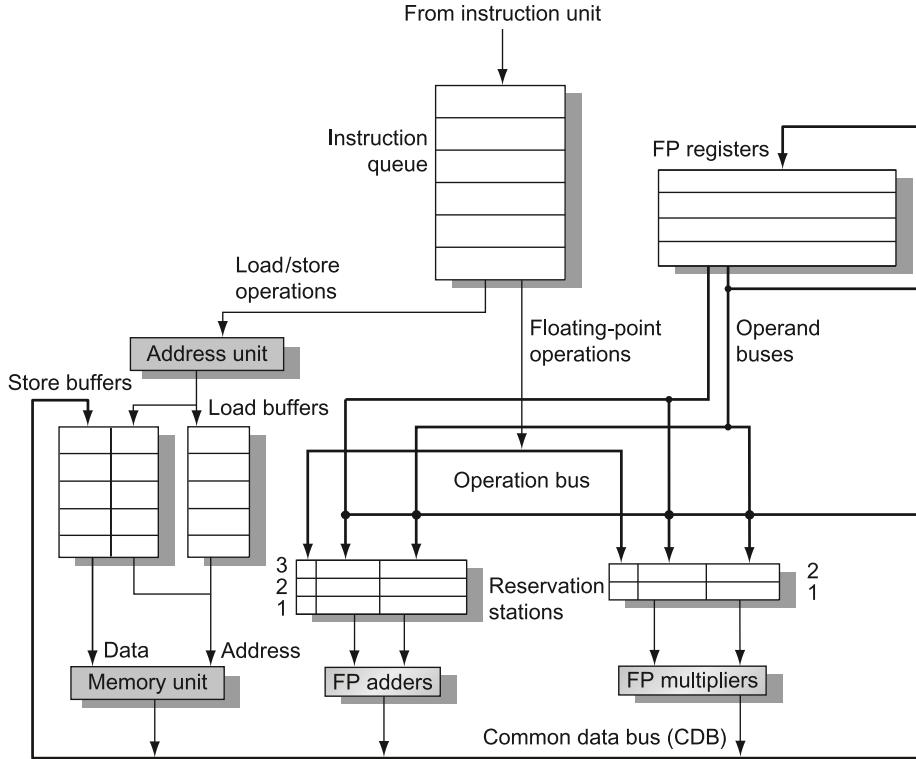


Figure 26: The Basic Structure of a Tomasulo-Style FPU. [2] It illustrates the **hardware datapath and control structure** used in a **RISC-V Floating-Point Unit (FPU)** that supports **dynamic scheduling** via **Tomasulo's algorithm**. This architecture shows **how a processor executes FP instructions out of order**, tracks dependencies using **tags**, and delivers results using the **Common Data Bus (CDB)**.

- **Instruction Queue.** Holds instructions waiting to be issued. The instructions are issued in-order (First In, First Out) to **Reservation Stations**.
- **Reservation Stations (RSs).** Act as **temporary instruction buffers** for each type of FU (e.g., FP ADD, MULT, DIV). Each RS holds:
 - Operation Type (e.g., ADD.D, MUL.D)
 - Operand values (V_j/V_k) or tags (Q_j/Q_k)
 - Status bits (busy, ready, etc.)

Waits until **both operands are ready** and the FU is free, then issues to the FU.

- **Function Units (FUs)**. Includes separate **pipelined units** for ADD, MULT, and DIV. Each can execute one instruction at a time. Upon completion, sends result to the **CDB**.
- **Register File (RF)**. Stores FP register values (F0-F31). Each register has: **Value field (Vi)**, and **Tag field (Qi)** that is the ID of the RS that will produce its value. When reading a register:
 - If $Qi = 0$, value is ready → use Vi .
 - If $Qi \neq 0$, value is pending → use tag to track it.
- **Common Data Bus (CDB)**. Broadcasts **completed results** (value and tag) to: **all waiting RSs** (those with Qj/Qk matching the tag), **Register File** (if that tag matches a register Qi). Enables **hardware forwarding**: results don't need to wait for write-back to the register file.
- **Load/Store Queue** (optional). In a full Tomasulo implementation, memory operations are handled by **Load and Store buffers**. This diagram is focused on **FP register-register instructions**, so memory is abstracted.

❖ Execution Flow

1. Issue (Instruction Queue → RS)

- Fetch next **Floating Point (FP)** instruction.
- If a Reservation Station (RS) for the operation type is available:
 - Copy instruction into RS
 - Fetch operands from Register File (RF)
 - If ready → store value in Vj/Vk
 - If not → store producer RS tag in Qj/Qk
 - Update the destination register's Qi with this RS tag (register renaming).

2. Execute (RS → FU)

- Once operands are ready ($Qj = Qk = 0$) and FU is free:
 - RS issues to its Functional Unit (FU)
 - FU computes the result (may take multiple cycles)
 - Execution is **out of order**, instructions are scheduled based on readiness, not program order.

3. Write Result (FU → CDB)

- When the FU finishes:
 - Sends the result with **its tag** on the **CDB**.
 - All RSs waiting on that tag update their operand fields.
 - Register file updates any registers whose Qi matches the tag.
 - The RS becomes free again.

3.7.5 Stages

Tomasulo's execution model is based on three main pipeline stages:

1. **Issue (in-order)**: decode instruction, send to RS.
2. **Execute (out-of-order)**: wait for operands, perform operation.
3. **Write Result (out-of-order)**: broadcast result via CDB to RF and RSs.

Each instruction flows through these three stages, but thanks to the algorithm's design, they don't have to do so in program order after issuing.

3.7.5.1 Stage 1: Issue

The **Issue Stage** is the first step in Tomasulo's 3-phase pipeline. In this stage, instructions are **decoded and dispatched** from the **Instruction Queue (IQ)** into **Reservation Stations (RSs)** (see Figure 26, page 149), where they wait until operands and functional units are ready.

❷ Main Responsibilities of the Issue Stage

1. **Instruction decoding**. Identify the opcode and register operands.
2. **Reservation Station assignment**. Check for an available RS of the appropriate type (e.g., an ADD.D goes to an RS attached to the FP ADD unit).
3. **Operand availability check**. For each source register:
 - ✓ If **ready**, read its **value** → store in V_j/V_k .
 - ✗ If **not ready**, read its **tag** → store in Q_j/Q_k (wait for it!).
4. **Destination register renaming**. Update the **register file's tag field** (Q_i). It now points to the RS handling this instruction. This is **implicit register renaming**.

❸ Why is this stage in-order?

Tomasulo always issues instructions **in program order**, one at a time. Why? Three reasons:

- Ensures **precise exceptions**, it helps the machine know which instruction caused a fault.
- Keeps memory operations (like stores) in **correct order**.
- **Prevents complexity** from instruction reordering too early.

Out-of-Order behavior starts **after issue**, in the **Execute** stage.

⚠ What can cause a stall during Issue?

1. **Structural Hazard:** no free Reservation Station.
2. **Load/Store queue full:** for memory operations.
3. (in more advanced versions) **Pending branch:** speculative instructions must wait.

✓ Why this Stage is powerful

- Begins **implicit renaming**, eliminating **WAR** and **WAW** hazards.
- Allows instructions to wait for operands without blocking the pipeline.
- Decouples **instruction arrival** from **execution readiness**.

Example 7: Issue stage

Say the instruction is:

¹ ADD.D F6, F2, F4

1. Tomasulo checks: “Is there a free ADD reservation station?”

✗ No, insert a **stall**: avoid **structural hazard**.

- ✓ If there are available reservation stations, then checks:

- (a) Is F2 ready?
 - ✓ Yes $\rightarrow V_j = \text{value of } F2$
 - ✗ No $\rightarrow Q_j = \text{tag of RS producing } F2$
- (b) Is F4 ready?
 - ✓ Yes $\rightarrow V_k = \text{value of } F4$
 - ✗ No $\rightarrow Q_k = \text{tag of RS producing } F4$

2. Finally, it updates: $F6.Q_i = RS1$ (renaming F6 to RS1). This means: “F6 is not ready yet, it will be produced by RS1”.

From this point on, **any instruction that needs F6 will wait for RS1**, not for F6 directly.

For example, a second instruction appears:

¹ MUL.D F8, F6, F10

Tomasulo checks: “what’s the status of F6?”. It finds:

$$F6.Q_i = RS1$$

F6 does not contain a value yet, but its value will come from RS1. So now, in the RS for the MUL instruction (say RS2), we set:

- $Q_j = RS1$ (because F6 is not ready, and its result will come from RS1).
- $V_k = \text{value of } F10$ (assuming F10 is ready).

This is the magic: RS2 knows that it must **listen for RS1’s result** on the Common Data Bus (CDB).

⌚ Quick Recap

1. Decode instruction
2. Allocate Reservation Station (RS)
3. Read operand values or tags
4. Rename destination register via tag (Q_i)
5. If no RS → stall

This stage is like **preparing an instruction for its flight**, loading its bags, assigning it a gate, and scheduling it to depart; but waiting for the runway (operands and Functional Unit) to be clear.

3.7.5.2 Stage 2: Start Execution

This is the phase where the instruction **actually performs the computation** in the appropriate **Functional Unit (FU)**, like ADD, MULT, or DIV. But, and this is key, it can only happen **when all operand values are ready and the Functional Unit (FU) is available**.

So, execution isn't just about doing the math, it's also about **waiting for the right moment**.

⌚ Preconditions: When can an Instruction start executing?

To start execution, three things must happen:

1. **Both operands are available.** In the Reservation Station (RS), this means:

$$Q_j = Q_k = 0$$

No tags, just valid values in V_j and V_k .

2. **The Functional Unit (FU) is free**, e.g., the FP ADD unit is idle.
3. (optional) **Instruction dependencies are resolved**, e.g., no pending branches or memory issues ahead.

Only **when all conditions are met**, the instruction *leaves* the Reservation Stations (RS) and enters the **Functional Unit (FU)**.

❖ What Happens in the Execute Stage?

Once **instructions are in Reservation Stations (RSs)** and waiting for operands, the **Execute Stage begins**. This is where:

- The instruction is **dispatched to the Functional Unit (FU)**.
- It **executes for its latency**.
- The result is **held**, ready to be broadcast in Stage 3.

But unlike a classic pipeline, **Tomasulo waits for readiness**, not for a fixed cycle count.

Let's break it into key sub-steps:

1. **Operand Readiness Check.** Each Reservation Station (RS) monitors:

- Q_j
- Q_k

If both are zero, it means that the **values V_j and V_k are ready**, and no other Reservation Station needs to broadcast them.

- **In other words, when both operands are ready, the instruction can start executing.**

2. **Check FU Availability.** Tomasulo now checks if the **relevant Functional Unit** (e.g., FP ADD, FP MUL, FP DIV) is **free**.

- ✓ If **yes** → dispatch instruction from RS to FU.
- ✗ If **no** → **stall inside the RS** until FU becomes free.

Multiple Reservation Stations might be ready, but **only one can use a given Functional Unit at a time**. A scheduler chooses which RS gets the FU (usually oldest-first).

3. **Instruction Execution (Latency Counts).** Once inside the Functional Unit (FU):

- The instruction performs its **actual arithmetic operation**.
- This can take multiple cycles depending on the instruction type.

Instruction	Latency (example)
ADD.D	2-3 cycles
MUL.D	4-7 cycles
DIV.D	10+ cycles

This time is **internal to the FU**, the RS has already “let go” of the instruction. The RS is now **free** and can be reused after the instruction completes Stage 3 (write-back).

4. **Memory Instructions (LOAD/STORE) - Special Case.** Unlike register-based instructions (e.g. ADD.D F6, F2, F4), memory instructions **read or write actual memory**, and memory access has some *dangerous side effects*:

- It's **global**: memory affects all parts of the program.
- It must be **precise**: no reordering can change program behavior.
- It can **raise exceptions** (e.g., accessing invalid addresses).

So Tomasulo handles **LOAD and STORE more carefully** than other instructions.

- **LOAD Instruction L.D F4, 0(R1).** Load the value from memory address R1 + 0 into register F4. What happens:
 - (a) **Wait for base register (R1) to be ready.**
 - If R1 is not ready, the Load Buffer stores a **tag**, just like in RSs.
 - Once ready, calculate **effective address** (e.g., 1000).
 - (b) **Wait for access to memory.** If there's no memory conflict, proceed to load.
 - (c) **Stage 3 (write).** The value is broadcast on the **CDB** to: waiting RSs (if any), register F4 (if still tagged).

LOADs are usually allowed to execute out-of-order (relative to other LOADs), *as long as we're sure no STORE before them is writing to the same address.*

- **STORE Instruction S.D F4, 0(R1).** Store the value in F4 to memory address R1 + 0. This is trickier, because:

- STORE instructions do **not write results to a register**.
- They write **into memory**, which is a **shared global state**.
- If we mess up the ordering, we can cause **wrong program results**.

What happens:

- (a) **Wait for both:**

- i. The **data to be stored** (F4) → may need to track a tag.
- ii. The **address to store to** (R1 + 0) → also may depend on a tag.

- (b) **Wait for memory access permission:**

- Tomasulo ensures **store ordering** is preserved.
- This means **STOREs happen in program order**.
- We cannot let a later STORE “jump ahead” of an earlier one.

- (c) **When both data and address are ready** → write to memory (not via CDB).

STORE instructions do not write anything on the CDB, because they don't produce a result for future instructions, they only update memory.

⌚ Classical Pipeline vs. Tomasulo

Classical Pipeline	Tomasulo
Fixed issue → execute	Waits dynamically for operands
Register-based data flow	Tag-based tracking via RSs
RAW hazards stall issue	RAW stalls only inside RSs
WAR/WAW hazards possible	Eliminated via register renaming

Table 18: Classical Pipeline vs. Tomasulo.

This is **dataflow execution**: instructions run **as soon as data is ready**, not when the program says so.

⌚ Quick Recap

1. Check if Q_j and Q_k are 0 (ready).
2. If yes and FU is free → dispatch.
3. Execute in FU (multi-cycle).
4. Prepare to broadcast result on CDB.

3.7.5.3 Stage 3: Write Result

This is the **moment the instruction finishes its computation** in the Functional Unit (FU), and the result is made **globally visible** by broadcasting it over the **Common Data Bus (CDB)**. It's the stage where:

- The **value computed by the FU** is sent to:
 - All **Reservation Stations (RSs)** waiting for it;
 - The **Register File (RF)** (if still tagged with RS).
- The instruction is now officially “**done**”.

❖ What happens step-by-step

Let's say we have this instruction in Reservation Station 3 (RS3):

1 MUL.D F6, F2, F4

RS3 starts execution in the MUL unit, takes several cycles, and finally finishes. Now we begin:

1. **Broadcast on CDB.** The Functional Unit (FU) places the result on the **Common Data Bus (CDB) along with its tag (RS3)**:

$$\text{CDB} \leftarrow \langle \text{tag} = \text{RS3}, \text{value} = \text{result} \rangle$$

This is the Tomasulo version of a “public announcement”: “*RS3 has finished computing a result. Anyone waiting for this, come get it!*”.

2. **Reservation Stations listen for Tags.** Every Reservation Station (RS) in the system checks: “*does my Qj or Qk match the tag on the CDB?*”. If yes:

- The Reservation Station (RS) grabs the **value**.
- Stores it in Vj or Vk.
- Clears the tag:

$$Qj \leftarrow 0 \quad \vee \quad Qk \leftarrow 0$$

Meaning: operand is now ready.

This process allows many instructions to **simultaneously wake up** when the result they needed finally becomes available.

3. **Register File Update.** The **Register File** checks: “*is there any register whose Qi = RS3?*”. If yes:

- It writes the result into that register.
- Clears the tag field ($Qi = 0$), this means “F6 is now valid and available”.

If the register has already been renamed again (e.g., $Qi = RS4$), we **don't write** to it, this avoid **WAW hazards**.

4. **Free the RS and FU.** The instruction in RS3 is now complete, RS3 is marked free, and the Functional Unit (FU) becomes available again for other instructions.

Example 8: why don't we write to a register when it has been renamed?

Suppose we have two instructions writing to the **same register**:

```
1 ADD.D F6, F2, F4      ; instruction A
2 MUL.D F6, F6, F8      ; instruction B
```

Instruction 1 is issued first, but instruction 2 also **writes to F6** (again). This is a **Write After Write (WAW)** situation.

1. Instruction 1 (ADD.D) is issued → assigned to RS1
 - Tomasulo sets: $F6.Qi \leftarrow RS1$ ("F6 will come from RS1").
2. Instruction 2 (MUL.D) is issued **immediately after** → assigned to RS2
 - Tomasulo updates again: $F6.Qi \leftarrow RS2$ ("F6 will now come from RS2").

Now we have:

- Two instructions writing to F6.
- But **only the result of RS2 should actually end up in F6**.
- RS1's result is now **obsolete**, because it was overwritten by a newer instruction.

When RS1 finishes (WRITE STAGE), RS1 broadcasts its result on the CDB:

$\langle Tag = RS1, Value = \dots \rangle$

Tomasulo checks the Register File: "does any register have $Qi = RS1$ ":

- ✓ If **yes**, update that register.
- ✗ If **no** (this case), $F6.Qi = RS2$, not RS1 anymore.

So we skip the write. Because if we allowed RS1 to write to F6 now, it would **overwrite the result of RS2**, which hasn't been computed yet, and then we'd break correctness.

The is how **Tomasulo avoids WAW hazards**: only write to a register if it's still waiting for our result. If someone newer came along and renamed the register again, we're out. Our value is no longer needed in the register file.

- CDB allows **many instructions** to grab the result **at the same time**.
- Values are **forwarded before being stored** (no need to wait for register write-back).
- **False dependencies (WAR/WAW)** are avoided automatically.
- **RAW dependencies** are resolved **dynamically**, as instructions "listen" for the values they need.

3.7.6 In-Depth Execution Example

Example 9: Dependence and Hazard Analysis

We're given a sequence of floating-point instructions to be executed by Tomasulo's algorithm:

```

1 L.D F6, 34(R2)      ; F6 ← Mem[R2 + 34]
2 L.D F2, 45(R3)      ; F2 ← Mem[R3 + 45]
3 MUL.D F0, F2, F4    ; F0 ← F2 × F4
4 SUB.D F8, F6, F2    ; F8 ← F6 - F2
5 DIV.D F10, F0, F6   ; F10 ← F0 ÷ F6
6 ADD.D F6, F8, F2    ; F6 ← F8 + F2

```

We begin by identifying **data dependencies** (RAW, WAR, WAW) between these instructions. Let's go instruction by instruction and analyze dependencies:

- Instruction 1: L.D F6, 34(R2) → No dependencies, loads into F6.
- Instruction 2: L.D F2, 45(R3) → No dependencies, loads into F2.
- Instruction 3: MUL.D F0, F2, F4
 - **RAW** dependency on F2, produced by Instr. 2
 - F4 is assumed ready (not in this instruction stream)
- Instruction 4: SUB.D F8, F6, F2
 - **RAW** dependency on F6, from Instr. 1
 - **RAW** dependency on F2, from Instr. 2
- Instruction 5: DIV.D F10, F0, F6
 - **RAW** dependency on F0, from Instr. 3
 - **RAW** dependency on F6, from Instr. 1 (but will be overwritten later)
- Instruction 6: ADD.D F6, F8, F2
 - **RAW** dependency on F8, from Instr. 4
 - **RAW** dependency on F2, from Instr. 2
 - **WAW** hazard: Instr. 1 wrote F6, this one **overwrites** it again.
 - **WAR** hazard: Instr. 5 reads F6, but this instruction **writes** it again.

Some key observations:

- Instr. 1 and 2 are **independent loads**, they will issue immediately.
- Instr. 3 (MUL.D) has to **wait for** F2, but F4 is assumed ready.

- Instr. 4 and 5 have **multiple dependencies**, they will wait in Reservation Station (RS) until results are available.
- Instr. 6 creates **WAW and WAR** hazards on F6, this is where **Tomasulo's register renaming** shines:
 - The second write to F6 gets a **new tag**.
 - The WAR/WAW conflicts are eliminated by tracking **RS tags instead of F6 name**.

Example 10: Register Renaming

The goal of this section is to demonstrate how **register renaming eliminates false dependencies** (WAR and WAW) and allows **out-of-order execution** by tagging instructions with Reservation Station IDs instead of register names.

Let's revisit the original instruction sequence:

```

1 L.D F6, 34(R2)
2 L.D F2, 45(R3)
3 MUL.D F0, F2, F4
4 SUB.D F8, F6, F2
5 DIV.D F10, F0, F6
6 ADD.D F6, F8, F2

```

Instead of using F6, F2, etc. directly, Tomasulo assigns a Reservation Station (RS) to track **who will produce** each value. This Reservation Station (RS) **renames** the register in the hardware temporarily. This helps **avoid**:

- ✓ **WAW** hazard → only the most recent instruction writes to the final destination.
- ✓ **WAR** hazard → consumers of earlier values don't get overwritten by later writers.

Here's a sample of how registers and reservation stations look during execution:

Instruction	RS	Dest Reg.	Qi Field in RF	Renaming Effect
L.D → F6	LB1	F6	F6.Qi ← LB1	F6 will be produced by LB1
L.D → F2	LB2	F2	F2.Qi ← LB2	F2 will be produced by LB2
MUL.D → F0	MUL1	F0	F0.Qi ← MUL1	F0 will be produced by MUL1
SUB.D → F8	SUB1	F8	F8.Qi ← SUB1	F8 will be produced by SUB1
DIV.D → F10	DIV1	F10	F10.Qi ← DIV1	F10 will be produced by DIV1
ADD.D → F6	ADD1	F6	F6.Qi ← ADD1	F6 now renamed again by ADD1

Let's follow F6 across the timeline:

1. Instruction 1: $F6 \leftarrow L.D \Rightarrow RF: F6.Qi = LB1$

2. Instruction 5: DIV.D uses F6 \Rightarrow waits on LB1's result

3. Instruction 6: F6 \leftarrow ADD.D \Rightarrow RF: F6.Qi = ADD1

So by the time ADD.D issues:

- **Register File (RF) no longer tracks LB1**, even though that result hasn't been written yet.
- DIV.D is still listening for LB1's result.
- **WAR/WAW hazards are avoided** because consumers and writers track **RS tags**, not the physical register name F6.

Now consider propagating the result with tags. Let's say MUL1 finishes first:

1. It broadcasts $\langle \text{Tag} = \text{MUL1}, \text{Value} = (\text{result}) \rangle$

2. The **Register File** updates:

- If Qi = MUL1 \rightarrow write result into the register.
- If Qi \neq MUL1 \rightarrow skip write.

3. RSs waiting on Qj = MUL1 or Qk = MUL1, grab the value.

Same logic applies for LB1, ADD1, etc. *Why Register Renaming Solves WAR and WAW?*

Hazard	Traditional Pipeline	Tomasulo's Solution
WAR	Later write may overwrite a value before it's read.	Each reader remembers which tag to wait for, not the register name.
WAW	Multiple instructions write to same register \rightarrow wrong order.	Only the last tag (Qi) is honored; old ones are discarded.

Register renaming in Tomasulo means:

- Every instruction output gets a **temporary, unique name**, called the *Reservation Station tag*.
- The **Register File tracks tags**, not values, while results are still pending.
- When a **result is ready**:
 1. It's broadcast with its tag
 2. Listeners receive it
 3. Register File updated only if **still needed**.

This make execution:

- **Safe** (no false hazards)
- **Flexible** (out-of-order)
- **Clean** (decouples values from register names)

Now we bring Tomasulo to life cycle-by-cycle. We'll track each instruction through the **three stages**:

- Issue
- Start Execution (Execute Start)
- Write Result (Write)

The instructions are:

```

1 L.D F6, 34(R2)
2 L.D F2, 45(R3)
3 MUL.D F0, F2, F4
4 SUB.D F8, F6, F2
5 DIV.D F10, F0, F6
6 ADD.D F6, F8, F2

```

❷ Why are SUBD and DIVD merged into ADDD and MULD units?

ADDD and SUBD are both **floating-point addition operations**, differing only in the operation:

- ADDD F6, F2, F4 → F6 = F2 + F4
- SUBD F6, F2, F4 → F6 = F2 - F4 = F2 + (-F4)

They both use the **same operand types**, have the **same input/output behavior** and can be performed by the **same kind of floating-point adder/-subtractor unit**. Therefore, a single ADD unit and its reservation stations can be reused for both operations. That's why in the RS table, SUBD uses an ADD-type RS (e.g., ADD1, ADD2).

Same for MULD and DIVD. These are both **multiplicative FP operations**. Architecturally, many processors use a **shared pipeline or RS pool** for:

- MULD (fastest latency)
- DIVD (slower latency but same operand structure)

So MULD and DIVD share **MUL reservation stations** and a **shared execution unit or Functional Unit slot**. That's why we see:

- MULD F0, F2, F4 issued to MUL1
- DIVD F10, F0, F6 is waiting in MUL2 for F0

They both use the **same RS pool**, because they're functionally similar in structure, and sharing resources saves hardware area. Also note that in floating-point arithmetic we have:

$$A \div B = A \times \left(\frac{1}{B}\right)$$

Instead of implementing a full-blown divider, many systems compute the **reciprocal of B** ($\frac{1}{B}$) and then multiply:

```

1 DIV.D    F10, F0, F6      ; F10 ← F0 ÷ F6
2 ≈
3 RECIP    F12, F6          ; F12 ← 1 ÷ F6
4 MULD     F10, F0, F12    ; F10 ← F0 × F12

```

This transformation is:

- **Valid numerically**, though it may lose precision.
- **Common in vector processors, GPUs**, and some scalar floating-point pipelines (e.g., ARM, GPUs).
- Used when latency and area are more important than perfect accuracy.

This allows the processor to **reuse the multiplier** (which is fast and compact) instead of having a separate, slow divider.

1. Cycle 1

- LD F6, 34(R2) is issued.
 - LoadBuffer1 allocated.
 - F6.Qi = Load1
- Base address R2 assumed ready → can start address calculation in next cycle.

Instruction	Issue	Start	Execute	Write	Result
LD F6, 34(R2)	1				
LD F2, 45(R3)					
MULTD F0, F2, F4					
SUBD F8, F6, F2					
DIVD F10, F0, F6					
ADDD F6, F8, F2					

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1	34		v(R2)	
Load2				
EXLoad				

EXLoad (or EX_LD). Tracks the **status of the Load Buffer(LB)** and **Load execution units**.

Name	Vj	Qj	Vk	Qk
ADD1				
ADD2				
ExADD				

ExADD. Tracks the use of the **addition** arithmetic **Functional Unit** (FU).

Name	Vj	Qj	Vk	Qk
MUL1				
MUL2				
ExMUL				

ExMUL. Tracks the use of the **multiplication** arithmetic **Functional Unit** (FU).

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi													Load1

Register Result Status. Shows the state of **register renaming** during execution. In other words, it shows the state of each floating-point register (F0-F31). It is a **snapshot of the Qi field** for each register in the floating-point register file.

2. Cycle 2

- LD F2, 45(R3) is issued
 - LoadBuffer2 allocated
 - F2.Qi = Load2
- LD F6 (Instr. 1) starts execution:
 - Address calculated (R2 + 34)
 - Memory load begins

Instruction	Issue	Start	Execute	Write	Result
LD F6, 34(R2)	1		2		
LD F2, 45(R3)		2			
MULTD F0, F2, F4					
SUBD F8, F6, F2					
DIVD F10, F0, F6					
ADDD F6, F8, F2					

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1	34		v(R2)	
Load2	45		v(R3)	
EXLoad	34		v(R2)	

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1				
ADD2				
ExADD				

ExADD.

Name	Vj	Qj	Vk	Qk
MUL1				
MUL2				
ExMUL				

ExMUL.

RF 0 1 2 3 4 5 6 7 8 9 10 ... 31	
Qi	Load2

Register Result Status.

3. Cycle 3

- MULTD F0, F2, F4 is issued
 - Reservation Station = MUL1
 - Needs:
 - * F2 → still pending from Load2 → Qj = Load2
 - * F4 → assumed ready → Vk = value = v(F4)
 - F0.Qi = MUL1
- LD F2 (Instr. 2) starts execution: Address calculated (R3 + 45); Memory load begins.
- Load (LD F6) has a delay and is still not finished. This is important because it causes **dependent instructions to wait**. However, this is the brilliance of Tomasulo's dataflow: these **instructions don't block the pipeline**, they just **wait for the tag (Load1) to appear in the CDB**.

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1		2	
LD F2, 45(R3)		2		
MULTD F0, F2, F4		3		
SUBD F8, F6, F2				
DIVD F10, F0, F6				
ADDD F6, F8, F2				

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1	34		v(R2)	
Load2	45		v(R3)	
EXLoad	34		v(R2)	

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1				
ADD2				
ExADD				

ExADD.

Name	Vj	Qj	Vk	Qk
MUL1		Load2	v(F4)	
MUL2				
ExMUL				

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		Load2			Load1							

Register Result Status.

4. Cycle 4

- SUBD F8, F6, F2 is issued
 - Reservation Station = SUB1
 - Needs:
 - * F6 → pending from Load1 → Qj = Load1
 - But the RS Load listens to the CDB for the Load1 tag. It sees the result of the CDB broadcast and overwrites its state:

$$\Rightarrow Qj = 0 \quad Vj = v(F6)$$

- * F2 → pending from Load2 → Qk = Load2
- F8.Qi = SUB1

- No execution start yet for MUL or SUB, still waiting on operands. BUT, the result of loading the first instruction (value from memory) is sent on the Common Data Bus (CDB). And the SUB instruction has grabbed it and is now ready to start executing.

Also, in the register result status table, we insert the value of F6 ($v(F6)$) instead of leaving the Load1 tag. This is because the CDB sends the result to all listeners.

Instruction	Issue	Start Execute	Write Result
LD F6, 34(R2)	1	2	4
LD F2, 45(R3)	2		
MULTD F0, F2, F4	3		
SUBD F8, F6, F2	4		
DIVD F10, F0, F6			
ADDD F6, F8, F2			

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1	34		$v(R2)$	
Load2	45		$v(R3)$	
EXLoad	34		$v(R2)$	

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1	v(F6)	Load1 ⁰		Load2
ADD2				
ExADD				

ExADD.

Name	Vj	Qj	Vk	Qk
MUL1		Load2	v(F4)	
MUL2				
ExMUL				

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		Load2			v(F6)		ADD1					

Register Result Status.

5. Cycle 5

- DIVD F10, F0, F6 is issued
 - Reservation Station = DIV1
 - Needs:
 - * F0 → pending from MUL1 → Qj = MUL1
 - * F6 is available thanks to the CDB
 - F10.Qi = DIV1
- LD F2 (Instr. 2) starts execution:
 - Addressed calculated (R3 + 45)
 - Memory load begins

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1		2	4
LD F2, 45(R3)	2		5	
MULTD F0, F2, F4	3			
SUBD F8, F6, F2	4			
DIVD F10, F0, F6	5			
ADDD F6, F8, F2				

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2	45		v(R3)	
EXLoad	45		v(R3)	

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk	Name	Vj	Qj	Vk	Qk
ADD1	v(F6)				MUL1			Load2	v(F4)
ADD2					MUL2		MUL1	v(F6)	
ExADD					ExMUL				

ExADD.

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		Load2				v(F6)		ADD1		MUL2		

Register Result Status.

6. Cycle 6

- ADDD F6, F8, F2 is issued
 - Reservation Station = ADD2
 - Needs:
 - * F8 → pending from ADD1 (SUB1) → Qj = ADD1
 - * F2 → pending from Load2 → Qk = Load2
 - F6.Qi = $\text{Load1} \xrightarrow{\text{ADD2}}$
- WAR on F6 has been eliminated:
 - Tomasulo renames the destination register of the ADDD instruction: F6 will now be written by ADD2, not by Load1 (cycle 1).
 - But crucially, DIVD was issued earlier (previous cycle); it already read F6's value from the Reservation Station of the Load1. It stored that value in V_k inside its RS (tag MUL2).
 - So even though ADDD will overwrite F6 later, it doesn't matter: DIVD already took what it needed and moved on.

The same rule is applied to SUBD.

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1	2		4
LD F2, 45(R3)	2	5		
MULTD F0, F2, F4	3			
SUBD F8, F6, F2	4			
DIVD F10, F0, F6	5			
ADDD F6, F8, F2	6			

Instruction status.

Name	V _j	Q _j	V _k	Q _k
Load1				
Load2	45		v(R3)	
EXLoad	45		v(R3)	

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1	v(F6)			Load2
ADD2		ADD1		Load2
ExADD				

ExADD.

Name	Vj	Qj	Vk	Qk
MUL1		Load2	v(F4)	
MUL2		MUL1	v(F6)	
ExMUL				

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		Load2			ADD2		ADD1		MUL2			

Register Result Status.

7. Cycle 7

- LD F2 completes execution. It broadcasts its result (value of F2) via the Common Data Bus (CDB). Forwarding takes place to both:
 - Register File: if $Q_i = \text{Load2}$, then $F_2 = \text{value} \Rightarrow Q_i = 0$
 - Reservation Stations:
 - * MULD (MUL1): $Q_j = \text{Load2} \rightarrow \text{now } V_j = v(F_2) \text{ and } Q_j = 0$
 - * SUBD (ADD1): $Q_k = \text{Load2} \rightarrow \text{now } V_k = v(F_2) \text{ and } Q_k = 0$
 - * ADDD (ADD2): $Q_k = \text{Load2} \rightarrow \text{now } V_k = v(F_2) \text{ and } Q_k = 0$
- ✖ DIVD F10, F0, F6 still waiting. Needs F0, which is being computed by MULD (Tag = MUL1). Can't execute yet.
- ✖ ADDD F6, F8, F2 still waiting. Needs F0, but still waiting on F8, tag = ADD1 (result of SUBD).

Instruction	Issue	Start Execute	Write Result
LD F6, 34(R2)	1	2	4
LD F2, 45(R3)	2	5	7
MULTD F0, F2, F4	3		
SUBD F8, F6, F2	4		
DIVD F10, F0, F6	5		
ADDD F6, F8, F2	6		

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2	45		v(R3)	
EXLoad	45		v(R3)	

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1	v(F6)		v(F2)	
ADD2		ADD1	v(F2)	
ExADD				

ExADD.

Name	Vj	Qj	Vk	Qk
MUL1	v(F2)		v(F4)	
MUL2		MUL1	v(F6)	
ExMUL				

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		v(F2)				ADD2		ADD1		MUL2		

Register Result Status.

8. Cycle 8

- **MULTD F0, F2, F4 starts execution in MUL1:**
 - It was waiting on F2 (now received via CDB) and F4 was already ready
 - Now: $Q_j = 0, Q_k = 0 \rightarrow \checkmark$ ready to execute
 - Executing in MUL1 Functional Unit
- **SUBD F8, F6, F2 starts execution in ADD1:**
 - F6 was already ready from Cycle 4 (CDB)
 - F2 now arrives, both operands are ready
 - Executing in ADD1 Functional Unit

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1	2		4
LD F2, 45(R3)	2	5		7
MULTD F0, F2, F4	3		8	
SUBD F8, F6, F2	4		8	
DIVD F10, F0, F6	5			
ADDD F6, F8, F2	6			

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk	Name	Vj	Qj	Vk	Qk
ADD1	v(F6)		v(F2)		MUL1	v(F2)		v(F4)	
ADD2			ADD1	v(F2)	MUL2			MUL1	v(F6)
ExADD	v(F6)		v(F2)		ExMUL	v(F2)		v(F4)	

ExADD.

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		v(F2)				ADD2		ADD1		MUL2		

Register Result Status.

10. Cycle 10

- SUBD F8, F6, F2 finishes execution and writes result:
 - The result of $F8 = F6 - F2$ is now available
 - **Broadcasts on the CDB:**
 - * Tag = ADD1 (RS that held SUBD)
 - * Value = result of $F6 - F2$
- Forwarding from CDB occurs:
 - **Register File:**
 - * $F8.Qi == ADD1$, then write result in F8 and set $F8.Qi = 0$
 - **Reservation Stations:**
 - * ADDD F6, F8, F2 is waiting for F8;
 - * If $Qj == ADD1$, then $Vj = v(F8)$ and $Qj = 0$
- The current execution MULD is still running (latency is 10 cycles) and will finish around cycle 17.

Instruction	Issue	Start Execute	Write Result
LD F6, 34(R2)	1	2	4
LD F2, 45(R3)	2	5	7
MULTD F0, F2, F4	3	8	
SUBD F8, F6, F2	4	8	10
DIVD F10, F0, F6	5		
ADDD F6, F8, F2	6		

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1	v(F6)		v(F2)	
ADD2	v(F8)		v(F2)	
ExADD	v(F6)		v(F2)	

Name	Vj	Qj	Vk	Qk
MUL1	v(F2)		v(F4)	
MUL2			MUL1	v(F6)
ExMUL	v(F2)		v(F4)	

ExADD.

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1	v(F2)				ADD2		v(F8)		MUL2			

Register Result Status.

11. Cycle 11

- ADDD F6, F8, F2 starts execution.
- MULTD F0, F2, F4 continues execution.
- DIVD still waiting.

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1	2		4
LD F2, 45(R3)	2	5		7
MULTD F0, F2, F4	3	8		
SUBD F8, F6, F2	4	8		10
DIVD F10, F0, F6	5			
ADDD F6, F8, F2	6		11	

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk	Name	Vj	Qj	Vk	Qk
ADD1					MUL1	v(F2)		v(F4)	
ADD2	v(F8)		v(F2)		MUL2		MUL1	v(F6)	
ExADD	v(F8)		v(F2)		ExMUL	v(F2)		v(F4)	

ExADD.

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		v(F2)		ADD2		v(F8)		MUL2				

Register Result Status.

13. Cycle 13

- ADDD F6, F8, F2 completes execution and write result.
- WAW hazard on F6 is avoided:
 - LD F6, 34(R2) originally wrote to F6 → F6.Qi = Load1
 - ADDD F6, F8, F2 overwrites F6 → F6.Qi = ADD2
 - But DIVD, which uses F6, was issued in cycle 5:
 - * It read the **value of F6** at that time from Load 1
 - * It stored it into its reservation station
 - * So even if F6 gets overwritten later by ADDD, it doesn't matter
 - * DIVD already has the value it needed

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1	2		4
LD F2, 45(R3)	2	5		7
MULTD F0, F2, F4	3	8		
SUBD F8, F6, F2	4	8		10
DIVD F10, F0, F6	5			
ADDD F6, F8, F2	6	11		13

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk	Name	Vj	Qj	Vk	Qk
ADD1					MUL1	v(F2)		v(F4)	
ADD2	v(F8)		v(F2)		MUL2		MUL1	v(F6)	
ExADD	v(F8)		v(F2)		ExMUL	v(F2)		v(F4)	

ExADD.

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	MUL1		v(F2)		v(F6)		v(F8)		MUL2				

Register Result Status.

18. Cycle 18

- MULTD F0, F2, F4 **completes execution**. Executed in MUL1 from cycle 8 to 18 (10 cycles total latency) and now broadcasts result on the Common Data Bus (CDB).
- DIVD is waiting for F0, but it is **now fully ready to start execution in the next cycle**.

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1	2		4
LD F2, 45(R3)	2	5		7
MULTD F0, F2, F4	3	8		18
SUBD F8, F6, F2	4	8		10
DIVD F10, F0, F6	5			
ADDD F6, F8, F2	6	11		13

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1				
ADD2				
ExADD				

Name	Vj	Qj	Vk	Qk
MUL1	v(F2)		v(F4)	
MUL2	v(F0)		v(F6)	
ExMUL	v(F2)		v(F4)	

ExMUL.

ExADD.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	v(F0)		v(F2)			v(F6)		v(F8)		MUL2			

Register Result Status.

19. Cycle 19

- DIVD F10, F0, F6 begins execution in MUL2.
- Division latency is not shown, but based on typical values in FP pipelines (and possibly earlier context), it may be **up to 20 cycles**. Therefore, DIVD will likely finish in a later cycle (e.g., 38+ depending on design).

Instruction	Issue	Start	Execute	Write Result
LD F6, 34(R2)	1	2		4
LD F2, 45(R3)	2	5		7
MULTD F0, F2, F4	3	8		18
SUBD F8, F6, F2	4	8		10
DIVD F10, F0, F6	5		19	
ADDD F6, F8, F2	6	11		13

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1				
ADD2				
ExADD				

Name	Vj	Qj	Vk	Qk
MUL1				
MUL2	v(F0)		v(F6)	
ExMUL	v(F0)		v(F6)	

ExMUL.

ExADD.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	v(F0)		v(F2)			v(F6)		v(F8)		MUL2			

Register Result Status.

59. Cycle 59

- DIVD F10, F0, F6 completes execution and writes result after a long 40-cycle latency.
- All instructions have issued, executed and written their results. No hazards occurred, and Tomasulo's algorithm correctly handled:
 - ✓ RAW (by waiting with tags).
 - ✓ WAR and WAW (by renaming registers with Qi, Qj, Qk).
 - ✓ Long execution delays (like the 40-cycle DIVD) without blocking unrelated instructions.

Instruction	Issue	Start Execute	Write Result
LD F6, 34(R2)	1	2	4
LD F2, 45(R3)	2	5	7
MULTD F0, F2, F4	3	8	18
SUBD F8, F6, F2	4	8	10
DIVD F10, F0, F6	5	19	59
ADDD F6, F8, F2	6	11	13

Instruction status.

Name	Vj	Qj	Vk	Qk
Load1				
Load2				
EXLoad				

EXLoad (or EX_LD).

Name	Vj	Qj	Vk	Qk
ADD1				
ADD2				
ExADD				

Name	Vj	Qj	Vk	Qk
MUL1				
MUL2	v(F0)		v(F6)	
ExMUL	v(F0)		v(F6)	

ExADD.

ExMUL.

RF	0	1	2	3	4	5	6	7	8	9	10	...	31
Qi	v(F0)		v(F2)				v(F6)		v(F8)		v(F10)		

Register Result Status.

3.7.7 Tomasulo vs. Scoreboarding

Both Tomasulo's algorithm and Scoreboarding (introduced in the CDC 6600) are **dynamic instruction scheduling** techniques designed to:

- Improve **Instruction-Level Parallelism (ILP)**
- Allow **out-of-order execution**
- Handle **data hazards** without compiler involvement

But they differ in:

- *How they track operands*
- *Where renaming happens*
- *How results are forwarded*
- *How decentralized the logic is*

⌚ How Tomasulo improves over Scoreboarding

1. **Eliminates WAR and WAW hazards.** Scoreboarding can only resolve RAW hazards, WAR and WAW cause stalls in Scoreboarding. Tomasulo's register renaming via **RS tags** removes these conflicts entirely.
2. **Allows hardware-level dataflow execution.** Scoreboarding relies on centralized logic to check operand availability. Tomasulo allows **instructions to execute as soon as operands are available**, independently. This creates a **data-driven execution model**, more similar to modern superscalar and out-of-order CPUs.
3. **Forwarding via the Common Data Bus (CDB)**
 - In Scoreboarding: instructions must wait for register file updates
 - In Tomasulo: values are **broadcast to all RSs and RF** as soon as they're computed

4. **Reservation Stations = Smarter Instruction Buffers.**

Hold operands, operation type, and source RS tags. Track readiness locally and act like **mini control units**, freeing up the need for centralized logic (Scoreboard).

However, despite its limitations, Scoreboarding:

- **Simpler to implement**
- Historically foundational (used in the CDC 6600)
- More **hardware-economical** for small-scale or embedded designs

But for high-performance out-of-order execution, Tomasulo's model is the clear winner.

Scoreboarding		Tomasulo's Algorithm
Operand Tracking		
Uses register availability		Uses tag-based tracking via RS
Register Renaming		
✗ None		✓ Implicit via RSs ($Q_i/Q_j/Q_k$)
WAR/WAW Hazards		
✗ Must be stalled explicitly		✓ Eliminated by renaming
Operand Read Timing		
At execution start		At issue (or when available) via CDB
Operand Forwarding		
✗ No, must wait for register write-back		✓ Yes, uses Common Data Bus (CDB)
Control Model		
Centralized Scoreboard		Decentralized logic at each RS + FU
Functional Units		
Check Scoreboard for availability		RS check tags and readiness locally
Instruction Readiness		
RAW only (no renaming)		RAW + renamed operands (data-driven)
Common Bus		
✗ No shared result forwarding		✓ Shared broadcast on CDB
Complexity		
Simpler renaming logic		More complex, but higher performance

Table 19: Tomasulo vs. Scoreboarding.

⚠ Drawbacks of Tomasulo's Algorithm

While Tomasulo's algorithm was a **revolutionary advancement** in dynamic scheduling, especially for floating-point pipelines, it is not without **limitations**. As computer architecture evolved into superscalar and multithreaded designs, certain drawbacks of Tomasulo's original formulation became more evident.

1. **Complex Hardware Implementation.** Tomasulo's distributed, tag-driven control system requires:

- **Reservation Stations (RS)** with operand tracking logic
- **Register File with tag fields (Q_i)**
- **A Common Data Bus (CDB)** that broadcasts to **many destinations**
- Logic to **match tags and grab values** dynamically

While this provides high performance, it increases **hardware area, control complexity** and **power consumption**.

In particular, **scaling** the CDB becomes harder with many Functional Units and many RSs, since broadcasting to all of them becomes a bottleneck.

2. **Centralized Common Data Bus (CDB) Bottleneck.** Tomasulo relies on a **single CDB** to broadcast results to all waiting RSs and update the register file. But:

- Only **one instruction can write** on the CDB per cycle
- This **limits write throughput**
- It becomes a **bottleneck** in wide-issue (superscalar) processors

Modern designs mitigate this with multiple CDBs, local bypassing networks, and reorder buffers (ROB) with centralized commit stages (advanced topics).

3. **Lack of Support for Precise Exceptions.** Tomasulo as originally designed does **not support precise exceptions**.⁶ This problem is due to the structure of the algorithm:

- Tomasulo allows **out-of-order completion** and **in-place register writes**
- Once a value is written to the register file via CDB, it's hard to "undo" it.

Scoreboarding can handle precise exceptions more gracefully. However, modern CPUs integrate Tomasulo-style scheduling with ROB (Reorder Buffer) to restore precise exceptions.

4. **Limited Scalability for Superscalar Designs.** Original Tomasulo was built for **one instruction per cycle** pipelines (like IBM 360/91). But in **superscalar** processors:

- Multiple instructions issue and retire each cycle;
- The CDB and tag-matching logic must scale accordingly;
- Reservation Stations structures become more complex with many FUs and wider pipelines.

Modern out-of-order cores use advanced techniques such as Register Alias Tables (RAT), physical register files, and ROB-based commit logic. These ideas **evolved from Tomasulo**, but allow better scaling and modularity.

⁶In precise exception handling, all instructions **before the faulting one** must complete, and none **after it** should affect the architectural state.

3.7.8 Register Renaming

The Register Renaming techniques have already been discussed in the previous sections. In section 3.1.2 (page 77) we gave a very short presentation dedicated to understand how to solve name dependencies, and in section 3.7.2 (page 142) we have seen what is the difference between static register renaming and implicit register renaming. In this section we will go deeper and understand how to apply this technique to the tomasulo algorithm.

3.7.8.1 Introduction

In modern out-of-order execution processors, one of the most significant **bottlenecks** to exploiting Instruction-Level Parallelism (ILP) lies in **false data dependencies**. Among these, **WAR (Write After Read)** and **WAW (Write After Write)** hazards can severely restrict the freedom to reorder instructions during execution. These are not **true dependencies** (like RAW, Read After Write), but are constraints imposed by the **limited number of registers** visible in the Instruction Set Architecture (ISA).

To tackle these issues, **Register Renaming** is employed, a mechanism that dynamically or statically replaces **architectural registers with a larger set of physical registers**, breaking these false dependencies and enabling more parallel execution.

☰ Types of Register Renaming

There are **two primary approaches** to register renaming:

- **Static Register Renaming**, handled at **compile-time**, often via techniques like **loop unrolling** and aggressive register allocation. This relies on **ISA-visible registers** and requires the compiler to explicitly assign new register names to avoid reuse conflicts.
- **Implicit (Dynamic) Register Renaming**, which is employed in hardware at **runtime**, often as part of advanced scheduling mechanisms such as **Tomasulo's algorithm**. This method uses structures like **reservation stations** and **reorder buffers** to rename registers implicitly and track dependencies through tags instead of register names.

⌚ Why Register Renaming is important?

Without register renaming, WAR and WAW hazards can stall instruction issue unnecessarily:

- A WAR hazard arises when an instruction wants to write to a register that a previous instruction still needs to read.
- A WAW hazard occurs when two instructions write to the same register, and their order must be preserved to maintain correctness.

By decoupling the programmer-visible registers from the physical storage used in the microarchitecture, renaming allows:

- ✓ **Multiple instructions** writing to the “same” architectural register to execute out of order.
- ✓ **Concurrent execution** of independent iterations of loops.
- ✓ **Elimination of false dependencies**, improving pipeline utilization and throughput.

Register renaming is a fundamental enabler of modern ILP exploitation. Whether implemented statically via compiler transformations or dynamically through hardware mechanisms like Tomasulo's algorithm, it allows us to **overlap instruction execution** aggressively and avoid performance penalties due to naming limitations of the architectural register file.

3.7.8.2 Loop Unrolling and Code Scheduling

Tomasulo's algorithm is a classic dynamic scheduling method that **naturally incorporates implicit register renaming**. It is designed to **overlap instruction execution** even when instructions involve true or false dependencies, by dynamically resolving hazards at runtime.

In this section, we focus specifically on **how Tomasulo's algorithm deals with a single loop**, showing the issues of data hazards and motivating why **renaming**, even if implicit, is critical. Furthermore, we show three techniques, included register renaming, to reach the best optimization. We also show three techniques, including register renaming, to achieve the best optimization.

Note that these topics are already covered in the Tomasulo section (3.7.2, page 143).

❷ Why Renaming?

Consider the following floating-point loop:

```

1 Loop: LD    F0, 0(R1)      ; Load word into F0 from address 0+R1
2      MULTD F4, F0, F2     ; Multiply F0 and F2, result into F4
3      SD    F4, 0(R1)      ; Store word from F4 to address 0+R1
4      SUBI R1, R1, #8       ; Decrement R1 by 8
5      BNEZ R1, Loop        ; Branch to Loop if R1 ≠ 0

```

Each loop iteration consists of five instructions. We assume branch prediction correctly predicts the loop branch as taken. *But are there dependencies?* Yes, **several hazards arise**:

1. RAW (Read After Write) on F0 and F4
2. WAW (Write After Write) on F0 and F4 across iterations
3. WAR (Write After Read) on R1 due to address updates and branching

These **hazards limit how aggressively we can overlap loop iterations**. Let's break down the hazard types across consecutive iterations:

- **Within a single iteration:**
 - MULTD depends on the result of LD (needs F0).
 - SD depends on the result of MULTD (needs F4).
 - BNEZ depends on SUBI to determine the loop condition.
- **Across iterations:**
 - LD of next iteration **writes** F0, creating a **WAW hazard** with the previous LD.
 - MULTD of next iteration **writes** F4, also leading to **WAW hazard**.
 - SUBI and BNEZ reuse R1, creating **WAR hazards**.

Without renaming, these hazards **would serialize iterations**, nullifying any gain from speculative execution.

Tomasulo's Contribution: Implicit Register Renaming

Tomasulo's algorithm automatically resolves these hazards by **renaming registers dynamically**:

- It uses **reservation stations** that hold **tags** instead of register names.
- When an instruction needs a value, it can either:
 - Use the value directly if it's ready;
 - Or wait for a tag associated with the producing reservation station.
- WAW and WAR hazards are naturally eliminated because **register names are no longer the synchronization point**, tags are!

Thus, **each iteration of the loop can proceed independently** as soon as its operands are available, even if they involve the same architectural register names.

Loop Unrolling combined with Register Renaming

While Tomasulo's algorithm handles implicit renaming at runtime, an alternative strategy to reduce the impact of dependencies, particularly across loop iterations, is to act **at compile time**, through a technique called **Loop Unrolling** combined with **Register Renaming**.

Loop Unrolling is a **compiler optimization technique** (or sometimes a manual optimization) that **expands the loop body** by replicating its operations multiple times in sequence, reducing the number of iterations and minimizing the overhead of branch instructions. Instead of executing a small amount of work many times, we “**stretch out**” multiple loop iterations into a single, larger block of straight-line code.

In our context, Loop Unrolling aims to **expose more parallelism statically** and to **reduce the loop overhead** by executing multiple iterations at once. Register Renaming in this context is crucial: **it eliminates artificial WAW and WAR dependencies** that would otherwise serialize execution.

Let's revisit the same loop we discussed before:

```

1 Loop: LD      F0, 0(R1)
2      MULTD F4, F0, F2
3      SD      F4, 0(R1)
4      SUBI   R1, R1, #8
5      BNEZ   R1, Loop

```

If we simply replicate this loop multiple times without renaming registers, WAW and WAR hazards would prevent full exploitation of the increased instruction stream. **To eliminate such hazards, we must rename the registers properly during unrolling.**

Here's the unrolled version of the loop, unrolled four times, with **explicit register renaming** to avoid hazards:

```

1 Loop: LD      F0, 0(R1)
2          MULTD F4, F0, F2
3          SD      F4, 0(R1)
4
5          LD      F6, -8(R1)
6          MULTD F8, F6, F2
7          SD      F8, -8(R1)
8
9          LD      F10, -16(R1)
10         MULTD F12, F10, F2
11         SD      F12, -16(R1)
12
13         LD      F14, -24(R1)
14         MULTD F16, F14, F2
15         SD      F16, -24(R1)
16
17         SUBI   R1, R1, #32
18         BNEZ   R1, Loop

```

- Observations

- New registers are introduced: F6, F8, F10, F12, F14, F16.
- **No two iterations reuse the same destination register.**
- **WAW hazards** between iterations are eliminated.
- **More operations are ready to issue independently.**
- Original loop: 5 instructions per iteration. Unrolled loop: 14 instructions for 4 iterations, so 3.5 instructions per iteration on average. Thus, unrolling not only **exposes parallelism** but also **increases code density efficiency**.

✓ Benefits

- ✓ **Higher ILP:** Each iteration's operations are now largely independent.
- ✓ **Reduced Loop Overhead:** The branch and counter updates are performed once every 4 iterations, decreasing control dependencies.
- ✓ **Efficient Use of Hardware:** Multiple functional units can work simultaneously, better exploiting available execution bandwidth.
- ✓ **Preparation for Further Optimizations:** The code can now be *rescheduled* to minimize stalls (we'll see this in future sections).

Finally, Tomasulo's algorithm **does not perform loop unrolling** by itself. Loop Unrolling is a **compiler optimization**. However, thanks to **dynamic register renaming and scheduling**, Tomasulo can **overlap multiple loop iterations even without unrolling**. This behavior mimics some of benefits of **loop unrolling at runtime**, but it doesn't physically unroll the code or replicate instructions.

⚡ Even more performance: Code Scheduling with Register Renaming

Once we have **unrolled a loop** and **renamed registers** to avoid false dependencies, the next logical optimization is **code scheduling**. The goal is now to **rearrange the unrolled instructions to minimize stalls**, especially those caused by **true RAW (Read After Write)** dependencies or **long execution latencies** (e.g., memory loads or floating-point operations). Thus, **even with register renaming, the order in which instructions are issued still matters to achieve maximum parallelism**.

Code Scheduling is the process of **reordering the instructions** in a program (especially after loop unrolling and register renaming) to **maximize parallelism** during execution. It is a **compiler-level** or sometimes **hardware-level** optimization. The **goal is to feed the CPU continuously with ready-to-execution instructions**.

Let's look at the previous code again, and by unrolling the loop by a factor of 4 and renaming the registers correctly, we had this sequence:

```

1 LD      F0, 0(R1)
2 MULTD  F4, F0, F2
3 SD      F4, 0(R1)
4 LD      F6, -8(R1)
5 MULTD  F8, F6, F2
6 SD      F8, -8(R1)
7 LD      F10, -16(R1)
8 MULTD  F12, F10, F2
9 SD      F12, -16(R1)
10 LD     F14, -24(R1)
11 MULTD  F16, F14, F2
12 SD      F16, -24(R1)
13 SUBI   R1, R1, #32
14 BNEZ   R1, Loop

```

⚠ Problem: if we execute this exactly in order, some instructions will have to wait unnecessarily for their operands because of true dependencies (particularly between LD and MULTD). So the **key idea** of code scheduling is:

- **Group independent instructions together** as much as possible.
- **Delay dependent instructions** just enough to satisfy their true data dependencies.
- **Hide long latencies** (e.g., waiting for a memory load) by **doing useful work in parallel**.

After careful rescheduling, the reordered code looks like:

```

1 LD      F0, 0(R1)
2 LD      F6, -8(R1)
3 LD      F10, -16(R1)
4 LD      F14, -24(R1)
5 MULTD  F4, F0, F2
6 MULTD  F8, F6, F2
7 MULTD  F12, F10, F2
8 MULTD  F16, F14, F2
9 SD      F4, 0(R1)

```

```

10 SD      F8, -8(R1)
11 SD      F12, -16(R1)
12 SUBI   R1, R1, #32
13 BNEZ   R1, Loop
14 SD      F16, -24(R1) ; In branch delay slot

```

The changes are:

- All **loads** are moved upfront, memory latency can be hidden while computing.
- **Multiplications** (dependent on loads) are grouped next.
- **Stores** happen **only after the corresponding multiply is finished**, respecting the RAW dependency.
- SUBI and BNEZ are moved right before the branch to reduce branching penalties.
- The **last store** (SD F16) is placed in the **branch delay slot**, making clever use of a wasted cycle.

And the benefits are:

- ✓ **Load latency hidden.** Loads start early while compute units are idle.
- ✓ **Maximal Floating Point unit usage.** Multiple MULTDs can be issued in parallel.
- ✓ **Minimal pipeline stalls.** Operands ready when needed.
- ✓ **Reduced branch penalties.** Good use of branch delay slot.

Technique	How Tomasulo handles it
Register Renaming	✓ Yes, dynamic, using RSs and tags.
Code Scheduling	✓ Yes, dynamic, based on operand availability.
Loop Unrolling	✗ No, but it can dynamically overlap iterations thanks to renaming + scheduling.

Table 20: Does Tomasulo's algorithm adopt Loop Unrolling, Register Renaming, and Code Scheduling?

3.7.8.3 How Tomasulo Overlaps Iterations of Loops

Tomasulo's algorithm can **dynamically overlap multiple iterations of a loop even without unrolling the code**, thanks to three main tricks:

- **Implicit Register Renaming via Reservation Stations.** Tomasulo does **not use the architectural register names** (like F0, F4) to track data anymore. Instead, it replaces every operand with a **dynamic tag** that points to a **Reservation Station (RS)**. This means: two iterations that both write to F0 are treated separately, because each LD or MULTD producing F0 will have its own RS tag.
✓ WAR and WAW hazards disappear.
- **Each instruction gets its own Reservation Station.** In Tomasulo, each new instruction (even from a new loop iteration) **gets its own Reservation Station**. So multiple instances of LD, MULTD, etc., from different loop iterations **can be active at the same time**. Even if the instructions all logically use F0, F4, etc., **hardware sees them as completely different** because they are tied to **different RS entries**.
✓ Dynamic loop unrolling effect: many loop iterations proceed in parallel without modifying the loop code.
- **Multiple Functional Units and Branch Prediction.** To exploit this properly, we must have:
 1. **Multiple Functional Units** (e.g., several FP multipliers and memory load units).
 2. **Branch Prediction** that allows **speculative execution** beyond the loop control instruction (like BNEZ).

Thus, Tomasulo **speculatively fetches, issues, and schedules instructions from many iterations ahead**, even **before knowing if the loop will finish**.

- ✓ Instruction issuing moves faster than control flow.

Tomasulo overlaps loop iterations by **making the Reservation Stations act like an extended virtual register file**, where each operation, even from different iterations, can proceed independently and dynamically, breaking naming limitations. This is sometimes called: **Dynamic Loop Unrolling**, **Dynamic Speculative Execution**, **Dynamic Out-of-Order Parallelism**. All enabled automatically by Tomasulo's mechanism.

3.7.8.4 Tomasulo Loop Execution

We now analyze a **concrete example** of how Tomasulo's algorithm **dynamically overlaps multiple loops iterations**, using **register renaming** and **out-of-order execution**.

The loop we consider is:

```

1 Loop:
2   LD    F0, 0(R1)
3   MULTD F4, F0, F2
4   SD    F4, 0(R1)
5   SUBI R1, R1, #8
6   BNEZ R1, Loop

```

❖ Setup Assumptions

- Cache miss on the **first LD** → 8 clock cycles.
- Cache hit on the **second LD** → 1 clock cycles.
- FP MULTD operation **latency** → 4 clock cycles.
- **Branch Prediction** → branch assumed **taken** (so iterations continue).
- Only SUBI and BNEZ clock times are shown explicitly for clarity.
- 5 instructions per loop iteration.

Thus, **memory accesses are slow**, and **floating-point operations are moderately slow** compared to integer operations.

⚠ Hazards

- ✗ **RAW hazards** on F0, F4, R1.
- ✓ (**eliminated dynamically**, thanks to Tomasulo's implicit register renaming) **WAW hazards** mainly on R1.
- ✓ (**eliminated dynamically**, thanks to Tomasulo's implicit register renaming) **WAW hazards** on F0 and F4 across iterations.

█ Structures in Tomasulo's algorithm

- **Instruction Status Table** (or Instruction Queue, or Tracking Table), tracks the status of each issued instruction through its three major stages: issue, execute and write result. Allow us to monitor the instruction flow cycle-by-cycle.
- **Reservation Stations Table**, temporarily hold instructions waiting to execute, **along with operand values or operand tags** if values are not yet available. Performs dynamic register renaming and resolve dependencies internally before issuing instructions to functional units.

- **Register Result Status Table**, tells which **Reservation Station (RS)** is responsible for generating the current value of each architectural register. Allows forwarding and prevents WAR/WAW hazards, by making reads wait for the right instruction to finish.
- **Load and Store Buffers (Memory Buffers)**, manage memory operations separately, since memory addresses take time to compute and loads/-stores have different latencies. Supports out-of-order memory access while ensuring correct memory ordering when necessary.

1. Cycle 1

- The **first instruction** from iteration 1, LD F0, 0(R1), is issued. It is assigned to Load Buffer 1 (**Load1**), marked as busy. The memory address to load from is calculated: R1 = 80, so the load address is 80.
- F0 is now associated with Load1 in the register status table. This means:
 - Any instruction that wants to read F0 must wait for Load1 to complete and broadcast its result.
 - Tomasulo does not write to F0 immediately, tag Load1 is what matters now.
- The other instructions (MULTD, SD, etc.) are **not issued yet** in this cycle. They are waiting because their source operands depend on F0 (the load result), which is **not ready** yet. Only the first load is active at this point.

At cycle 1, we **prepare the first data** to flow through Tomasulo's architecture: we **rename F0 dynamically**, **assign** it to a load buffer, and **start the memory access** (which will take 8 cycles because of cache miss). All following instructions **must track this tag** (Load1) to know when F0 will be ready.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1			
1	MULTD F4, F0, F2				
1	SD F4, 0(R1)				
2	LD F0, 0, R1				
2	MULTD F4, F0, F2				
2	SD F4, 0(R1)				

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	No		\
Load3	No		\
Store1	No		
Store2	No		
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	No					
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
1	80	Load1								

Register result status.

2. Cycle 2

- Instruction MULTD F4, F0, F2 (from iteration 1) is **issued**. It is assigned to Mult1 (a multiplication Reservation Station). Operands:
 - F2 is ready, its value is available $\rightarrow V_j = R(F2)$
 - F0 is *not* ready, **waiting on** Load1 to produce it $\rightarrow Q_j = \text{Load1}$.
- F0 is still pending, being loaded by Load1. F4 will eventually be produced by the MULTD operation (currently waiting for F0 to complete).

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1			
1	MULTD F4, F0, F2		2		
1	SD F4, 0(R1)				
2	LD F0, 0, R1				
2	MULTD F4, F0, F2				
2	SD F4, 0(R1)				

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	No		\
Load3	No		\
Store1	No		
Store2	No		
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	Yes	MULTD		R(F2)	Load1	
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
2	80	Load1		Mult1						

Register result status.

3. Cycle 3

- Instruction SD F4, 0(R1) (store) from iteration 1 is **issued**. It is assigned to **Store1** (store buffer). The store buffer must wait for F4 to be computed by **Mult1**. We record the tag of the functional unit that will produce the needed value F4.

Iter.	Instruction	Issue	Exec Comp	Write Res
1	LD F0, 0, R1	1		
1	MULTD F4, F0, F2	2		
1	SD F4, 0(R1)	3		
2	LD F0, 0, R1			
2	MULTD F4, F0, F2			
2	SD F4, 0(R1)			

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	No		\
Load3	No		\
Store1	Yes	80	Mult1
Store2	No		
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	Yes	MULTD		R(F2)	Load1	
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
3	80	Load1		Mult1						

Register result status.

6. Cycle 6

- We have skipped cycle 4 and 5 because we are only waiting for LD, MULTD or SD updates. In cycle 5, register R1 becomes 72 due to the operation `SUBI R1, R1, #8`.
- Instruction LD F0, 0(R1) from iteration 2 is now issued into Load2 buffer. This new load has a cache hit (remember: only the first load had a cache miss).
- F0 is again associated with Load2 (the second load). The previous F0 result (from Load1) is **overwritten** by the new F0 pending load. Even though the logical register F0 is overwritten, the data dependency graph remains correct, because **Tomasulo tracks data flow by tags**, not register names. Thus:
 - MULTD F4, F0, F2 (iteration 1): waits for Load1 to complete (old F0)
 - MULTD F4, F0, F2 (iteration 2): will wait for Load2 to complete (new F0)
- WAW hazard avoided** by dynamic renaming. New instructions get the new F0 and old instructions already locked the old value.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1			
1	MULTD F4, F0, F2	2			
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6			
2	MULTD F4, F0, F2				
2	SD F4, 0(R1)				

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	Yes	72	\
Load3	No		\
Store1	Yes	80	Mult1
Store2	No		
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	Yes	MULTD		R(F2)	Load1	
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
6	72	Load2		Mult1						

Register result status.

7. Cycle 7

- Even though the second load is ready (cache hit, 1 cycle latency, and issued in cycle 6), it must wait for the first load to complete because they share the same architectural address and could have aliasing or store/load order dependencies. So unless advanced speculation is used, loads and stores to the same address must be done in program order.

This is called **Memory Disambiguation**. In out-of-order processors, memory disambiguation is the logic that decides whether two loads/stores/loads can safely reorder. Tomasulo in its original form does not perform aggressive disambiguation.

- MULTD F4, F0, F2 (iteration 2) is issued.

Iter.	Instruction	Issue	Exec	Comp	Write	Res
1	LD F0, 0, R1	1				
1	MULTD F4, F0, F2	2				
1	SD F4, 0(R1)	3				
2	LD F0, 0, R1	6				
2	MULTD F4, F0, F2	7				
2	SD F4, 0(R1)					

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	Yes	72	\
Load3	No		\
Store1	Yes	80	Mult1
Store2	No		
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	Yes	MULTD		R(F2)	Load1	
	Mult2	Yes	MULTD		R(F2)	Load2	

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
7	72	Load2		Mult2						

Register result status.

8. Cycle 8

- SD F4, 0(R1) (iteration 2) is issued. It is assigned to **Store2** (store buffer). The store buffer must wait for F4 to be computed by Mult2. We record the tag of the functional unit that will produce the needed value F4.
- **WAW (Write After Write) hazards on registers F0 and F4 have been resolved**, meaning the first and second iterations are now fully overlapped.

Iter.	Instruction	Issue	Exec	Comp	Write	Res
1	LD F0, 0, R1	1				
1	MULTD F4, F0, F2	2				
1	SD F4, 0(R1)	3				
2	LD F0, 0, R1	6				
2	MULTD F4, F0, F2	7				
2	SD F4, 0(R1)	8				

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	Yes	72	\
Load3	No		\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	Yes	MULTD		R(F2)	Load1	
	Mult2	Yes	MULTD		R(F2)	Load2	

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
8	72	Load2		Mult2						

Register result status.

9. Cycle 9

- LD F0, 0(R1) (iteration 1, address 80, cache miss) finally completes and writes back in the next cycle (1 cycle latency). Now the LD F0, 0(R1) from iteration 2 can start its execution.
- Even though we don't track the SUBI instruction, it is issued in this cycle (iteration 2).

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1		9	
1	MULTD F4, F0, F2	2			
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6			
2	MULTD F4, F0, F2	7			
2	SD F4, 0(R1)	8			

Instruction status.

	Busy	Addr	FU
Load1	Yes	80	\
Load2	Yes	72	\
Load3	No		\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	Yes	MULTD		R(F2)	Load1	
	Mult2	Yes	MULTD		R(F2)	Load2	

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
9	72	Load2		Mult2						

Register result status.

10. Cycle 10

- LD F0, 0(R1) (iteration 1) now writes back (1 cycle latency). It also sends the result through the CDB. The value written can be represented as M[80], indicating the location of memory at address 80.

Mult1 was listening in the CDB for the Load1 value; now it receives the update, so it sets the Qj entry to zero and updates its Vj value with the value in memory at address 80 (M[80]).

- The second load (Load2) completes its execution in 1 cycle thanks to the cache hit. In the next cycle, the result is written back. This instruction was blocked by the previous load to avoid too aggressive memory disambiguation.
- Even though we don't track the BNEZ instruction, it is issued in this cycle (iteration 2). This updates the value in register R1 from 72 to 64 (minus 8).

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1		9	10
1	MULTD F4, F0, F2		2		
1	SD F4, 0(R1)		3		
2	LD F0, 0, R1		6		10
2	MULTD F4, F0, F2		7		
2	SD F4, 0(R1)		8		

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	Yes	72	\
Load3	No		\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
4	Mult1	Yes	MULTD	M[80]	R(F2)		
	Mult2	Yes	MULTD		R(F2)	Load2	

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
10	64	Load2		Mult2						

Register result status.

11. Cycle 11

- LD F0, 0(R1) (iteration 2) now writes back (1 cycle latency). It also sends the result through the CDB. The value written can be represented as M[72], indicating the location of memory at address 72.
- Mult2 was listening in the CDB for the Load2 value; now it receives the update, so it sets the Qj entry to zero and updates its Vj value with the value in memory at address 72 (M[72]).
- The third load (Load3) starts its execution because the cycle continues to execute. So we are now at iteration number 3, at the first load command: LD F0, 0(R1). So we add the entry Load3 to the table *Load and Store Buffers*.
- Mult1 continues its execution, it has 3 cycles left to complete.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1		9	10
1	MULTD F4, F0, F2	2			
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6		10	11
2	MULTD F4, F0, F2	7			
2	SD F4, 0(R1)	8			

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	Yes	64	\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
3	Mult1	Yes	MULTD	M[80]	R(F2)		
4	Mult2	Yes	MULTD	M[72]	R(F2)		

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
11	64	Load3		Mult2						

Register result status.

12. Cycle 12

- MULTD F4, F0, F2 (iteration 3) cannot be issued, because in Tomasulo's algorithm, to issue a new instruction there must be a free Reservation Station available for that operation type (here, floating-point multiply). If no RS is free, the instruction cannot be issued and must wait.

Since both Mult1 and Mult2 are occupied executing previous multiplies, the third multiply (from iteration 3) cannot issue yet.

- Mult1 continues its execution, it has 2 cycles left to complete.
- Mult2 continues its execution, it has 3 cycles left to complete.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1		9	10
1	MULTD F4, F0, F2	2			
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6		10	11
2	MULTD F4, F0, F2	7			
2	SD F4, 0(R1)	8			

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	Yes	64	\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
2	Mult1	Yes	MULTD	M[80]	R(F2)		
3	Mult2	Yes	MULTD	M[72]	R(F2)		

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
12	64	Load3		Mult2						

Register result status.

13. Cycle 13

- LD F0, 0(R1) completes its execution and writes the result back to the CDB ($M[64]$).
- Mult1 continues its execution, it has 1 cycle left to complete.
- Mult2 continues its execution, it has 2 cycles left to complete.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1		9	10
1	MULTD F4, F0, F2	2			
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6		10	11
2	MULTD F4, F0, F2	7			
2	SD F4, 0(R1)	8			

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	No		\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
1	Mult1	Yes	MULTD	M[80]	R(F2)		
2	Mult2	Yes	MULTD	M[72]	R(F2)		

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
13	64	M[64]		Mult2						

Register result status.

14. Cycle 14

- Mult1 finishes its execution at the end of this cycle! This will allow the Store1 to write the result in memory at the address 80, and also to the MULTD instruction in the iteration 3 to be issued.
- Mult2 continues its execution, it has 1 cycle left to complete.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1		9	10
1	MULTD F4, F0, F2	2		14	
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6		10	11
2	MULTD F4, F0, F2	7			
2	SD F4, 0(R1)	8			

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	No		\
Store1	Yes	80	Mult1
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
0	Mult1	Yes	MULTD	M[80]	R(F2)		
1	Mult2	Yes	MULTD	M[72]	R(F2)		

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
14	64	M[64]		Mult2						

Register result status.

15. Cycle 15

- MULTD (iteration 1) writes the result to the CDB; this triggers the **Store1** listener, which waits for the result to be stored in memory (the result of the load, M[80], multiplied by the value of register F2), and the MULTD of iteration 3, which waits for a free reservation station to be issued.
- Mult2 finishes its execution at the end of this cycle! This will allow the **Store2** to write the result in memory at the address 72.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1	9		10
1	MULTD F4, F0, F2	2	14		15
1	SD F4, 0(R1)	3			
2	LD F0, 0, R1	6	10		11
2	MULTD F4, F0, F2	7		15	
2	SD F4, 0(R1)	8			

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	No		\
Store1	Yes	80	M[80] × F2
Store2	Yes	72	Mult2
Store3	No		

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
4	Mult1	Yes	MULTD	M[64]	R(F2)		
0	Mult2	Yes	MULTD	M[72]	R(F2)		

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
15	64	M[64]		Mult1						

Register result status.

16. Cycle 16

- MULTD (iteration 2) writes the result to the CDB; this triggers the Store2 listener, which waits for the result to be stored in memory (the result of the load, M[72], multiplied by the value of register F2).

Iter.	Instruction	Issue	Exec Comp	Write Res
1	LD F0, 0, R1	1	9	10
1	MULTD F4, F0, F2	2	14	15
1	SD F4, 0(R1)	3	16	
2	LD F0, 0, R1	6	10	11
2	MULTD F4, F0, F2	7	15	16
2	SD F4, 0(R1)	8		

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	No		\
Store1	Yes	80	M[80] × F2
Store2	Yes	72	M[72] × F2
Store3	Yes	64	Mult1

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
3	Mult1	Yes	MULTD	M[64]	R(F2)		
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
16	64	M[64]		Mult1						

Register result status.

17. Cycle 17

- SD (iteration 1) writes the result of the expression $M[80] \times F2$ to memory at address 80.
- Since the SD instruction (iteration 1) has finished its execution, the second store (iteration 2) can start its execution and send the result to memory between M[72] and F2.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1	9		10
1	MULTD F4, F0, F2	2	14		15
1	SD F4, 0(R1)	3	16		17
2	LD F0, 0, R1	6	10		11
2	MULTD F4, F0, F2	7	15		16
2	SD F4, 0(R1)	8		17	

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	No		\
Store1	No		
Store2	Yes	72	$M[72] \times F2$
Store3	Yes	64	Mult1

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
2	Mult1	Yes	MULTD	$M[64]$	R(F2)		
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
17	64	M[64]		Mult1						

Register result status.

18. Cycle 18

- SD (iteration 2) writes the result of the expression $M[72] \times F2$ to memory at address 72.

Iter.	Instruction	Issue	Exec	Comp	Write Res
1	LD F0, 0, R1	1	9		10
1	MULTD F4, F0, F2	2	14		15
1	SD F4, 0(R1)	3	16		17
2	LD F0, 0, R1	6	10		11
2	MULTD F4, F0, F2	7	15		16
2	SD F4, 0(R1)	8	17		18

Instruction status.

	Busy	Addr	FU
Load1	No		\
Load2	No		\
Load3	No		\
Store1	No		
Store2	No		
Store3	Yes	64	Mult1

Load and Store Buffers.

Time	Name	Busy	Op	Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
1	Mult1	Yes	MULTD	M[64]	R(F2)		
	Mult2	No					

Reservation Stations.

Clock	R1	F0	F2	F4	F6	F8	F10	F12	...	F30
18	64	M[64]		Mult1						

Register result status.

3.7.9 Reorder Buffer (ROB)

3.7.9.1 Hardware-based Speculation

Speculation in hardware is a foundational technique in high-performance processors, used to execute instructions **before it is known whether they are needed**. This is especially relevant when instructions are **control-dependent** on unresolved branches.

In particular, hardware-based speculation enables:

- **Speculative instruction execution** before knowing branch outcomes.
- **Rollback** in case a mispredicted path was taken.

This approach increases Instruction-Level Parallelism (ILP), allowing more instructions to be in-flight, even if some turn out to be unnecessary.

❷ Why speculation needs the ROB

Executing instructions speculatively introduces a challenge: “*how do we prevent side effects (e.g., register or memory updates) from mispredicted instructions?*”. The **solution**: **defer the architectural state update** until it’s safe to commit.

This is exactly what the **Reorder Buffer (ROB)** is for:

- ✓ It holds results of instructions that have **finished execution** but are **not yet committed**.
- ✓ It allows **instruction results to be written in-order**, preserving program semantics.
- ✓ In case of a misprediction, the ROB enables the processor to **flush invalid speculative results** quickly and precisely.

But the ROB is **critical** for:

- ⚠ Decoupling **execution completion** from **state update (commit)**.
- ⚠ Supporting **precise exceptions** (so the CPU can cleanly stop at the last correct instruction).
- ⚠ Managing **speculative and non-speculative instruction tracking**.

3.7.9.2 Why ROB is really needed

Modern high-performance processors use **out-of-order execution (OoO)** to improve instruction-level parallelism. However, they must still ensure **in-order commit** to preserve correct program behavior. This is where the **Reorder Buffer (ROB)** comes into play.

❷ Why Out-of-Order Execution and In-Order Commit?

Out-of-Order execution because:

- ✗ **Dependencies** (e.g., RAW hazards) and **delays** (e.g., memory access) prevent some instructions from being executed immediately.
- ✓ OoO execution **allows the processor to bypass stalled instructions** and execute independent ones earlier.

❸ This helps keep functional units busy and **improves throughput**.

For example, if an instruction is waiting for a memory load, another instruction, for example an ALU op, can execute immediately.

Unfortunately, Out-of-Order execution **introduces risk**:

- ⚠ Executing instructions out of order can **break the precise exception model**.
- ⚠ Also, **speculative** instructions (e.g., those after a branch) could be **incorrect**.

To preserve **program correctness, instructions must commit (retire) in order**, exactly as they appear in the original program.

✓ ROB's role in bridging OoO and In-Order Commit

The **Reorder Buffer (ROB)** is a hardware structure used in modern out-of-order processors to: **support out-of-order execution while enforcing in-order commitment of instructions**, ensuring correct architectural state and precise exception handling.

The ROB is the key mechanism enabling this balance:

- During **issue**: **allocate an entry** in the ROB, marking the instruction's place in program order.
- During **execution**: **store the result** in the ROB as soon as the instruction finishes.
- During **commit**: **only commit** the instruction (i.e., update architectural registers or memory) **if**
 1. It is **at the head** of the ROB
 2. Its **execution has completed**

3. It is not speculative

This design ensures:

- ✓ Architectural state (register file, memory) is **updated in program order only**.
- ✓ **Speculative results** are kept in the ROB until validated.
- ✓ It is possible to **undo speculated instructions** if needed (e.g., branch misprediction).

The ROB allows a **processor to execute instructions as soon as their operands are ready**, regardless of program order, while **ensuring they commit their results in-order**.

⌚ What about Precise Exception support?

An exception is *precise* if:

1. The processor can **stop at a well-defined point**, corresponding to a specific instruction.
2. All **previous instructions are fully committed**, and
3. **No subsequent instruction** has modified the architectural state.

This allows the OS or exception handler to reliably identify and handle the error. In **out-of-order execution**, instructions complete in a different order than they appear in the program. **Without careful control** instructions *after* the faulting one could have modified registers or memory and this would **leave the system in an inconsistent state**.

The **Reorder Buffer (ROB)** solves this by:

- ✓ **Storing all results temporarily**: Instructions write their output to the ROB instead of the register file or memory.
- ✓ **Committing results in-order**: Only when an instruction is at the **head** of the ROB and marked **completed**, its result is written to the architectural state.
- ✓ **Rejecting speculative results**: If an exception or misprediction occurs:
 - (a) The ROB **flushes all speculative entries after the faulting instructions**.
 - (b) Execution restarts from the faulting instruction or the exception handler.

This rollback is clean because the **architectural state remains untouched** beyond the last committed instruction.

Example 11: Precise Exception support

Imagine a sequence:

```

1 Instr 1 → OK
2 Instr 2 → OK
3 Instr 3 → Exception (e.g., divide by zero)
4 Instr 4 → Executed early (0o0), wrote to reg R5

```

✗ Without the ROB:

- Instruction 4 might modify R5 *before* the exception at *Instr 3* is recognized.
- This makes the exception imprecise, corrupting the state.

✓ With the ROB:

- Instruction 4's result is held in the ROB.
- Since *Instr 3* caused an exception, no later instruction commits.
- R5 is unaffected, precise state is preserved.

❖ Functional Roles of the ROB

1. **Result Buffering.** Holds **results of instructions** that:
 - Have **completed execution**.
 - But have **not yet committed** to the architectural state (e.g., register file or memory).
2. **Speculative Result Propagation.** ROB acts as a **buffer to pass results among instructions that have started speculatively after a branch**. This allows speculative instructions (e.g., those after a predicted branch) to forward their results internally via the ROB without prematurely updating architectural registers.
3. **Precise Interrupt Support.** Originally introduced in 1988, the ROB was created to:
 - Preserve the **precise interrupt model**.
 - Guarantee that **only committed instructions affect** the architectural state.
 - Enable **rollback on branch misprediction or exceptions** by flushing speculative ROB entries.

The ROB is not merely a commit buffer, it is also a **speculation-aware result forwarding structure**, enabling safe communication among instructions that may never actually commit.

3.7.9.3 ROB as a Data Communication Mechanism

In Tomasulo's original algorithm, **Reservation Stations (RSs)** handled register renaming and operand forwarding. However, with the introduction of the ROB, it **replaces the renaming and forwarding function of RSs**. Before we explain how, we need to understand some *basic* concepts of ROB.

💡 What are ROB numbers?

A **ROB number** is simply an **index** (or tag, in Tomasulo's algorithm) **identifying** a specific **entry in the Reorder Buffer**. Each instruction that is issued receives a unique ROB number that corresponds to its place in the ROB, its **slot identifier** (slot ID).

The ROB number is then used in two key ways:

1. **As a destination tag (Register Renaming).** When an **instruction** is issued and it **will write to a register**:
 - The destination register is **not renamed to another physical register** (like in a pure register renaming scheme).
 - Instead, the **architectural register is mapped to the instruction's ROB number**.

It is the *identical* logic of tag in Tomasulo algorithm.

2. **For Operand Dependency Resolution.** Later **instructions that depend on the result** of the ROB number do not need to stall:
 - They **record the ROB number** as the operand source.
 - Once the ROB number becomes “ready” (i.e., the value is written to the ROB), dependent instructions can **read the value from the ROB number** and proceed to execution.

In tomasulo algorithm this feature is provided by the CDB, where each instruction listens to that because it waits for the result of the tag.

As happens in tomasulo algorithm, this mechanism:

- ✓ Avoids **WAR and WAW hazards**.
- ✓ Enables **data forwarding** even before instructions commit.

Tomasulo	ROB-based Tomasulo
Tags = RS entry IDs	Tags = ROB entry numbers
Values broadcast via CDB	Same, but identified by ROB number
RS buffers result locally	ROB stores result globally
Commit on write-back	Commit delayed and via ROB head

Table 21: Quick comparison between Tomasulo with and without ROB.

✖ Updated Role of Reservation Stations

The ROB not only buffers instruction results but also **propagates those results to dependent instructions** as soon as they are ready, even if not yet committed. So, with the ROB managing renaming and data forwarding:

- ✖ **Reservation Stations** are **no longer responsible for naming/tagging**.
- ✓ Their role is now focused on:
 - **Buffering decoded instructions** before they're issued to the Functional Units (FUs);
 - **Holding operand values** (or ROB tags) temporarily;
 - Helping reduce **structural hazards**.

Reservation Stations now act like **staging areas**, not tracking results or resolving dependencies directly.

Function	Tomasulo	ROB-based Tomasulo
Register Renaming	RS	ROB
Data forwarding	RS	ROB
Instruction buffering	RS	RS (same)
Operand availability tracking	RS	RS (via ROB tags)

Table 22: Summary of architectural changes.

In modern speculative Tomasulo architectures, the ROB becomes the central structure for both result tracking and inter-instruction communication. RSs are demoted to lightweight instruction and operand buffers.

3.7.9.4 Architecture

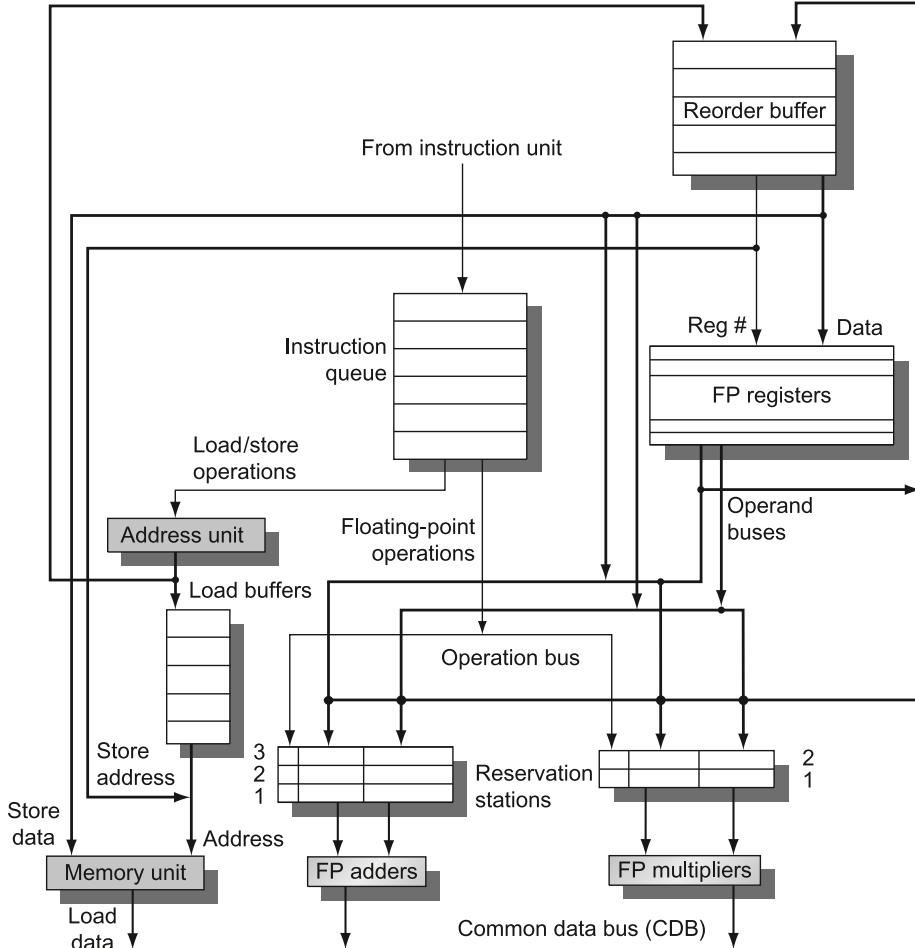


Figure 27: The basic structure of a FP unit using Tomasulo's algorithm and extended to handle speculation (ROB). [2]

This architecture shows a **speculative version of Tomasulo's algorithm**, integrating a Reorder Buffer (ROB) and eliminating the classic **store buffer**. It's designed for **floating-point**, but the same design applies to general OoO pipelines.

- **Instruction Queue.** Holds fetched instructions awaiting issue. Supplies instructions to **reservation stations** and the **ROB** in parallel.
- **Reservation Stations.** Acts as **temporary buffers** between issue and execution. Each FP unit (adder, multiplier) has dedicated RSs. Stores instructions with operand tags or values. **Unlike the classic Tomasulo, renaming is handled via the ROB.**
- **Floating-Point Units.** Includes **FP Adders** and **FP Multipliers**. Execution units receive instructions from RSs when operands are ready.

- **Common Data Bus (CDB)**. Used to broadcast result values and their corresponding **ROB tag**. All waiting instructions and ROB entries **listen** to the CDB.
- **Reorder Buffer (ROB)**. Central to this architecture, replaces:
 1. Register renaming logic
 2. Store buffers

Each ROB entry holds critical metadata about one in-flight instruction:

- **Busy field**. Indicates whether the ROB entry is currently active (busy) or free (available).
- **Instruction type field**. Specifies the instruction category:
 - * Branch (no destination result),
 - * Store (destination is memory address),
 - * Load/ALU (destination is a register).
- **Destination field**
 - * For load and ALU instructions: **target register** number.
 - * For store instructions: **memory address** where the value must be written.
- **Value field**. Holds the **result** of the instruction after execution, kept until commit.
- **Ready Bit**. Set when execution has completed and the result is valid.
- **Speculative Flag**. Shows whether the instruction is executed speculatively (e.g., after a predicted branch) or not.

- **Load Buffers & Address Unit**.

Load Buffers is a **queue or small table** that holds load instructions waiting to access **memory**. It **temporarily buffers** loads *after issue* and *before* they actually perform a memory access. Its purpose is shown in the example on page 194.

Address Unit is a specialized hardware block **dedicated to calculating effective addresses** for loads and stores. Given a base register value and an offset, it computes an effective address.

⌚ How Load Buffers and Address Unit Work Together?

1. **Issue Stage**: the instruction is issued. Allocates an entry in the ROB and Load Buffer.
2. **Address Calculation**: Address Unit computes the effective address.
3. **Memory Access Decision**: if there is no preceding store to the same address (or speculation allows), the load can access memory early. If not, the load must wait until memory disambiguation clears it.
4. **Load Execution**: read data from memory. Write the result into the ROB entry. When ready and safe, commit to the architectural register file.

- **Register File (FP Registers)**: not updated directly after execution. Instead, it is **updated only when instructions commit via the ROB**, preserving the **precise state model**.

❷ Key Innovations in this design

- **ROB replaces store buffers**: the store buffer doesn't exist anymore, now memory writes are delayed until commit.
- **Register renaming is moved from RSs to ROB**: simplifies dependency tracking.
- **Precise exceptions and speculative execution are both supported** via ROB tracking and flushing.
- **CDB remains critical** for result forwarding and readiness detection.

❸ Ready Bit

The **Ready Bit** is a control bit that indicates whether the instruction associated with this ROB entry has **completed execution** and its **result is available**. It is **set to true** when:

1. The instruction finishes execution.
2. The result is written into the ROB.

Its purpose is to **inform dependent instructions that they can now read this value** (via the CDB or directly from the ROB). Also, it enables commit: an instruction can **only commit if its ready bit is set**.

In other words, the ready bit **tracks the availability of the computed result**.

❹ Speculative Flag

The **Speculative Flag** is a control bit that indicates whether the **instruction was issued after an unresolved branch prediction** (or other speculative control decision). It is:

- ❶ Set to true only when the instruction is **issued speculatively, before** the control flow is confirmed.
- ❷ Cleared when the speculation is resolved (e.g., branch prediction validated).

Its purpose is to ensure that **speculative instructions do not prematurely update** the architectural state. Allows all speculative instructions to be **flushed** upon misprediction or exception.

In other words, the speculative flag tracks whether the instruction is **tentative** or **safe**.

✖ ROB Structure and Operation: Circular Buffer

The Reorder Buffer (ROB) is implemented as a **circular buffer** (circular FIFO, First-In-First-Out, queue), managed with **two pointers**:

- **Head pointer.** Points to the **oldest instruction** that is **next to commit** (i.e., retire).
The head **advances** when instructions commit.
- **Tail pointer.** Points to the **next free entry**, where a new issued instruction will be inserted.
The tails **advances** when new instructions are issued.

Example 12: ROB Circular Buffer

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes	In		No	F0	—	No
ROB1	No	—		—	—	—	—
ROB2	No	—		—	—	—	—
ROB3	No	—		—	—	—	—

- **Head** points at ROB0 (first to commit once ready).
- **Tail** points at ROB1 (next free entry for a new instruction).

⌚ Why Two Pointers?

The head and tail pointers allow the ROB to **efficiently manage** dynamic instruction issue and commit while **preserving program order**, and they **minimize hardware complexity** by avoiding costly entry movement.

1. **Sequential Insertion at the Tail (Issue Stage).** When a new instruction issues, it needs a free ROB entry. The tail pointer indicates where to insert the new instruction. After insertion, the tail advances to the next free slot. This preserves the program order at issue time.

New instructions always go at the tail.

2. **Sequential Retirement from the Head (Commit Stage).** When an instruction completes and satisfies all commit conditions (ready, not speculative, prior instructions retired), it retires. The head pointer shows which instruction should commit next. After commit, the head advances to the next oldest instruction.

Commit always happens starting from the head, ensuring in-order commit.

3. **Circular Buffer Efficiency.** Memory and hardware are limited: we don't want an infinitely growing ROB. Using a circular buffer:

- When the tail reaches the end of the buffer, it wraps around to position 0.
- Same for the head.

This reuses space efficiently without needing to shift entries manually.

Circular structure avoids expensive data movement and saves silicon area.

4. **Detecting Full and Empty Conditions.** With just head and tail:

- If head = tail and entry busy → ROB is full (cannot issue more instructions).
- If head = tail and entry not busy → ROB is empty (no instructions to retire).

Simple hardware checks based on two pointers.

Example 13: ROB and Rename Table

The **Rename Table** is a small **hardware table** that records where the *most recent* value of each register will come from, while the instruction is still speculative or not yet committed. It connects logical (architecture) registers (like F0, F2, etc.) to ROB entries that are producing their updated value.

It works like this:

- When an instruction is issued that **writes** a register:
 - The Rename Table **maps** that register to its **new ROB number**.
- When an instruction needs to **read** a register:
 - It **checks the Rename Table**:
 - * If an ROB entry is mapped: the instruction reads the pending value from the ROB.
 - * If no mapping: the instruction reads from the architectural register file.

Now we present an example using the ROB and the Rename Table. At the beginning of this example, we have a ROB almost empty and a Rename Table partially filled.

1. Cycle 1, first step.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes	In		No	F0	—	No
ROB1	Yes	In+1		No	F2	—	No
ROB2	No	—		—	—	—	—
ROB3	No	—		—	—	—	—

Head is at ROB0 and Tail is at ROB2.

Two instructions were issued:

- In writes to F0, mapped to ROB0.
- In+1 writes to F2, mapped to ROB1.

Both are non-speculative (SPEC = No)

Register	F0	F2	F4	F6	F8
Pointer	ROB0	ROB1			

F0 and F2 are now renamed: if anyone uses F0/F2, they will look at the respective ROB entry, not the old register.

2. Cycle 2, second step.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes	In		No	F0	—	No
ROB1	Yes	In+1		No	F2	—	No
ROB2	Yes	In+2		No	F4	—	No
ROB3	No	—		—	—	—	—

Head is still at ROB0 and Tail is now at ROB3.

Another instruction issued:

- In+2 writes to F4, mapped to ROB2.

Still non-speculative.

Register	F0	F2	F4	F6	F8
Pointer	ROB0	ROB1	ROB2		

3. Cycle 3, speculative instruction issued.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes	In		No	F0	—	No
ROB1	Yes	In+1		No	F2	—	No
ROB2	Yes	In+2		No	F4	—	No
ROB3	Yes	In+3		No	F6	—	Yes

Head is still at ROB0. As for Tail, since all four entries (ROB0 through ROB3) are occupied and there are no free entries, Tail cannot move forward until there is a free ROB entry. Tail must logically “point” to the next free entry, but since there is none available, it “hovers” behind or at ROB0, waiting for a free entry after a commit. So Tail is *stalled* at ROB0 because the ROB is full.

A new instruction (In+3) was issued and allocated into ROB3. In+3 writes to register F6. The speculative bit is set to “Yes”, so this instruction depends on an unresolved branch (it’s speculative).

Register	F0	F2	F4	F6	F8
Pointer	ROB0	ROB1	ROB2	ROB3	

4. Cycle 4, first execution result ready.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes	In		Yes	F0	10	No
ROB1	Yes	In+1		No	F2	—	No
ROB2	Yes	In+2		No	F4	—	No
ROB3	Yes	In+3		No	F6	—	Yes

Head is still at ROB0. Tail is still waiting.

The instruction at ROB0 (In) **completed execution**: it is ready (Yes) and the result is 10. **But it has not yet committed** because commits happen separately.

The value 10 is still in the ROB, not yet in the register file. F0 is still mapped to ROB0 (waiting commit).

Register	F0	F2	F4	F6	F8
Pointer	ROB0	ROB1	ROB2	ROB3	

5. Cycle 5, first commit happens.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	No		—	—	—	—	—
ROB1	Yes		In+1	No	F2	—	No
ROB2	Yes		In+2	No	F4	—	No
ROB3	Yes		In+3	No	F6	—	Yes

Head is now at ROB1. Tail finally finds an empty entry and points to it (ROB0).

Instruction In (ROB0) committed:

- ✓ It was ready (Ready = Yes);
- ✓ Not speculative (Spec = No);
- ✓ Head pointer advanced from ROB0 → ROB1.

Register File updated:

- F0 now officially = 10 (not pending anymore).

Rename Table updated:

- F0 no longer points to a ROB entry, but to the **final committed value** (10).

Register	F0	F2	F4	F6	F8
Pointer	10	ROB1	ROB2	ROB3	

6. Cycle 6, new instructions, partial execution.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes		In+4	No	F8	—	Yes
ROB1	Yes		In+1	Yes	F2	20	No
ROB2	Yes		In+2	No	F4	—	No
ROB3	Yes		In+3	No	F6	—	Yes

Head is still at ROB1. Tail stalled because ROB is full.

A new instruction (In+4) was issued and allocated to ROB0:

- Destination: F8
- Marked speculative (SPEC = Yes)

Meanwhile the instruction In+1 (in ROB1) has **completed execution** (READY = Yes) with the value 20.

Register	F0	F2	F4	F6	F8
Pointer	10	ROB1	ROB2	ROB3	ROB0

7. Cycle 7, commit of a non-speculative instruction.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes		In+4	No	F8	—	Yes
ROB1	No		—	—	—	—	—
ROB2	Yes		In+2	No	F4	—	No
ROB3	Yes		In+3	No	F6	—	Yes

Head is now at ROB2. Tail is now at ROB1.

Instruction In+1 in ROB1 has committed, and its destination register F2 now holds the value 20. ROB1 is now free (BUSY = No) and the head moves forward to ROB2.

Register	F0	F2	F4	F6	F8
Pointer	10	20	ROB2	ROB3	ROB0

8. Cycle 8, misprediction.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	Yes		In+4	No	F8	—	Yes
ROB1	No		—	—	—	—	—
ROB2	Yes		In+2	No	F4	—	No
ROB3	Yes		In+3	No	F6	—	Yes

Head is still at ROB2. Tail is still at ROB1.

ROB0 and ROB3 still contain **speculative instructions** waiting for execution. ROB2 contains non-speculative instruction but not ready yet.

At some point, **misprediction detected!** The branch speculation fails and all speculative instructions must be **flushed** (discarded).

Register	F0	F2	F4	F6	F8
Pointer	10	20	ROB2	ROB3	ROB0

8. Cycle 8, after misprediction flush.

ROB#	BUSY	INSTR.	TYPE	READY	DEST	VALUE	SPEC
ROB0	No		—	—	—	—	—
ROB1	No		—	—	—	—	—
ROB2	Yes		In+2	No	F4	—	No
ROB3	No		—	—	—	—	—

Head is still at ROB2. Tail is moved to ROB3 because after a misprediction, the **tail must move to the next free ROB entry after the last non-speculative instruction**, maintaining strict program order.

ROB0 (In+4) and ROB3 (In+3) were speculative → flushed immediately. Only ROB2 remains:

- It is non-speculative (safe);
- Still waiting to execute or ready.

Register	F0	F2	F4	F6	F8
Pointer	10	20	ROB2	ROB3 (<i>obsolete</i>)	ROB0 (<i>obsolete</i>)

3.7.9.5 Speculative Tomasulo Algorithm with ROB

The speculative Tomasulo architecture extends the original Tomasulo's algorithm by integrating a Reorder Buffer (ROB). This allows:

- ✓ Dynamic register renaming using ROB entries;
- ✓ In-order commit;
- ✓ Precise exception handling;
- ✓ Efficient speculative execution with safe rollback.

➊ Main Innovations Introduced by the ROB

1. **Pointers towards ROB Entries.** Operands are no longer tracked via Reservation Stations. Instead, **operands are identified by ROB numbers**. This ensures centralized tracking of instruction dependencies.
2. **Implicit Register Renaming.** When a destination register is renamed:
 - It is mapped to a **ROB entry**;
 - This removes **WAR** and **WAW hazards** automatically;
 - Enables **dynamic loop unrolling** without conflicts.
3. **Delayed Update of Architectural State.** Registers and memory are **not updated at execution time**. They are updated **only when the instruction is at the head of the ROB and is ready to commit**. This guarantees correct program state even in speculative execution.
4. **Easy Speculation Management.** By holding all results inside the ROB until commit:
 - **Mispredictions** can be handled by **flushing speculative entries**;
 - **Precise exceptions** are maintained, ensuring a clean rollback mechanism.

➋ Operating phases of the Speculative Tomasulo Algorithm

The execution flow is divided into four main stages:

1. **Issue: Fetch and Prepare the Instruction**

The **Issue phase** is the step where an instruction is fetched from the **Instruction Queue** and prepared for execution. In this phase:

- The processor checks if there are **resources available** to allow the instruction to proceed.
- It allocates both a **Reservation Station (RS) slot** and a **ROB entry**.

Only if both are available, the instruction can continue.

- (a) **Check for Available Slots.** The issue logic checks:

- Is there a free Reservation Stations?
- Is there a free ROB entry?

If yes, the instruction is issued into the pipeline. If no, the **instruction stalls** in the Instruction Queue and waits.

- (b) **Fetch Operands.** Once issued, the instruction's **source operands are fetched**. If an operand is:

- **Ready in the Register File** → send it immediately to the Reservation Station.
- **Pending in the ROB** → send the ROB number (tag) to the Reservation Station instead (to wait for the value).

Thus, the RS either: receives the actual operand value (if available), or a **tag** (ROB number) pointing to where the value will appear later.

- (c) **Allocate ROB Entry for the Result.** A **ROB entry** is assigned to the instruction. The **ROB number** is sent to the Reservation Station. This number will later be used to **tag the result on the Common Data Bus (CDB)** after execution. This **ROB number** replaces the register name during speculative execution.

⚠ What Happens If No Space? If either no RS slot is free, or no ROB entry is free, then the instruction **cannot issue** and **must stall**. This ensures the system **never overflows** the ROB or Reservation Stations, maintaining control over execution resources.

2. Execution Started

After the instruction is successfully issued and resides in a **Reservation Station (RS)**, the next step is **starting execution**. Execution can start only when all required **operands are available**.

- (a) **Check Operand Availability.** The Reservation Station **monitors the availability** of operands. If **both operands are available** (i.e., no more pending data hazards):

- Instruction **immediately starts execution** on its operands.
- This stage is called **EX** (Execute).

⚠ RAW hazards (Read After Write) must already be solved to start execution.

- (b) **Monitoring for Pending Operands.** If one or more operands are **not yet ready** (e.g., still waiting for a previous instruction to produce them):

- The RS **listens** to the **Common Data Bus (CDB)**.
- As soon as the missing operand appears on the CDB, it captures it and becomes ready.

⚠ Dynamic readiness: the RS waits without blocking the pipeline.

(c) **Special Case: Store Instructions.** For **store** instructions:

- Only the **base register** (address computation) must be available to start execution.
- At this point, the **effective memory address** is computed.

⌚ **The store value itself can still arrive later, closer to the commit phase.**

Instructions don't wait statically, they dynamically listen and proceed as soon as they can, maximizing execution overlap.

3. Execution Completed & Write Result

Once the **execution of the instruction finishes**, the result must be **broadcast** and **stored** so that dependent instructions can proceed and the instruction can eventually commit.

(a) **Broadcast the Result on the Common Data Bus (CDB).** The completed instruction **places its result** on the **Common Data Bus**. This **broadcasts** the result to:

- All **Reservation Stations** that are waiting for it;
- The corresponding **ROB entry** (updating its **Value** field).

⌚ **Any instruction that was stalled waiting for this operand can now proceed.**

(b) **Update the ROB Entry.** The **Value** field of the instruction's **ROB entry** is updated with the computed result. The instruction's **Ready bit** inside the ROB is set to **true**, meaning the instruction is now completed and awaiting commit.

⌚ **The ROB now safely holds the completed result, ready for in-order retirement.**

(c) **Release the Reservation Station.** The Reservation Station that held this instruction is now **marked as free**. It becomes available for **issuing new instructions**.

⌚ **This keeps the hardware resource usage efficient and avoids blocking incoming instructions.**

(d) **Special Case: Store Instructions.** For **store** operations:

- The **data to be stored** (not just the address) is written into the ROB's **Value** field.
- If the value wasn't ready at the start of the store's execution, the RS monitors the CDB until the data is captured.

⌚ **This ensures the store operation has both address and data ready by the time it needs to commit.**

Results are not immediately written to registers; they are buffered safely into the ROB and forwarded to dependent instructions dynamically.

4. Commit: Updating the Architectural State

The **Commit** phase is when the result of an instruction is **safely written into the Register File or memory**, and the instruction is **retired from the ROB**. **Commit happens strictly in program order**:

- ✓ Only the instruction at the **head of the ROB** can be considered for commit.
- ✓ The instruction must have **finished execution** and have its **result ready**.

(a) **Normal Commit**. If the instruction at the head is **completed** (Ready bit = true):

- Its **result is written into the Register File** (if it's a computational instruction);
- The **ROB entry is freed**.

⌚ **The architectural state is updated in strict program order.**

(b) **Store Commit**. If the instruction at the head is a **store**:

- Instead of updating the Register File, it **writes the data into memory**;
- Afterward, the instruction is removed from the ROB.

⌚ **Memory operations happen precisely at the correct point, preserving correct memory ordering.**

(c) **Branch Commit and Misprediction Handling**. If the instruction at the head is a **branch**:

- ✓ If the branch was **predicted correctly**, it is committed normally, no special action.
- ✗ If the branch was **mispredicted**:
 - The **ROB is flushed**, all speculative instructions after the branch are discarded.
 - Execution **restarts** at the correct successor (true path).

⌚ **Mispredictions are corrected cleanly, and speculative damage is avoided.**

This flushing process is sometimes called **graduation** in technical literature⁷.

Only when the instruction is at the ROB head, completed, and non-speculative, it can safely update the real architectural state (register file or memory). This phase ensures that the processor maintains a **precise state** even under heavy speculation and out-of-order execution.

⁷**Graduation** is the final commit of an instruction from the ROB into the architectural state (registers or memory), happening in program order after execution is completed and speculation has been validated. In literature like *Computer Architecture: A Quantitative Approach* [2], graduation and commit are used interchangeably.

In other words, it is another technical term for instruction commit in speculative out-of-order processors.

Example 14: Speculative Tomasulo Algorithm with ROB - Execution Overlap

We are executing this loop:

```

1 Loop:
2     LD    F0, 0(R1)
3     MULTD F4, F0, F2  # RAW hazard on F0
4     SD    F4, 0(R1)   # RAW hazard on F4
5     SUBI  R1, R1, #8
6     BNEZ  R1, Loop   # RAW and WAR hazards on R1

```

- LD loads from memory address pointed by R1 into F0.
- MULTD depends on F0 being ready → must wait for the LD.
- SD depends on F4 → must wait for MULTD.
- SUBI modifies R1 (address computation for next iteration).
- BNEZ depends on the result of SUBI to decide whether to branch.
- **First iteration issued**

ROB#	BUSY	INSTR. TYPE	READY	DEST.	VALUE	SPEC.
ROB0	Yes	LD F0, 0(R1)	No (exec. start)	F0	—	No
ROB1	Yes	MULTD F4, F0, F2	No (issued)	F4	—	No
ROB2	Yes	SD F4, 0(R1)	No (issued)	M[0+[R1]]	—	No
ROB3	Yes	SUBI R1, R1, #8	No (issued)	R1	—	No
ROB4	Yes	BNEZ R1, Loop	No (issued)	—	—	No
ROB5	No					
ROB6	No					
ROB7	No					

The head is at ROB0 and the tail is at ROB5.

Register #	F0	F2	F4	F6	F8	F10
Pointer	ROB0		ROB1			

The actions are:

- All these 5 instructions are issued one after another.
- Each gets a **ROB entry**.
- The **Rename Table** is updated:
 - * F0 is renamed to ROB0;
 - * F4 to ROB1;
 - * Memory address for SD points to ROB2;
 - * R1 to ROB3 (because SUBI modifies R1).

ROB entries are allocated from ROB0 to ROB4.

• **After Some Cycles (ROB Full)**

ROB#	BUSY	INSTR. TYPE	READY	DEST.	VALUE	SPEC.
ROB0	Yes	LD F0, 0(R1)	Ready to Commit	F0	M[0+[R1]]	No
ROB1	Yes	MULTD F4, F0, F2	No (issued)	F4	—	No
ROB2	Yes	SD F4, 0(R1)	No (issued)	M[0+[R1]]	—	No
ROB3	Yes	SUBI R1, R1, #8	Exec. Started	R1	—	No
ROB4	Yes	BNEZ R1, Loop	No (issued)	—	—	No
ROB5	Yes	LD F0, 0(R1)	No (issued)	F0	—	Yes
ROB6	Yes	MULTD F4, F0, F2	No (issued)	F4	—	Yes
ROB7	Yes	SD F4, 0(R1)	No (issued)	M[0+[R1]]	—	Yes

The head is still at ROB0 and the tail is stall.

Register #	F0	F2	F4	F6	F8	F10
Pointer		ROB5		ROB6		

Now, after some execution, the situation is:

- **ROB is full:** we used all 8 entries (ROB0 to ROB7).
- Meanwhile:
 - * LD (ROB0) has completed (Ready = Yes).
 - * MULTD (ROB1) is still waiting to execute (waiting on F0).
 - * SUBI (ROB3) has started executing (updating R1).
- Second iteration of the loop started speculative issuing:
 - * A new LD into F0 (ROB5);
 - * A new MULTD into F4 (ROB6);
 - * A new SD (ROB7).

Speculative entries being appearing (ROB5-ROB7 marked speculative). New renamings happen:

- F0 is now renamed to ROB5 (overriding previous ROB0).
- F4 is renamed to ROB6.

The processor allows speculative execution past the BNEZ, but only under the assumption that the branch prediction was correct.

- **First Commit and Progress**

ROB#	BUSY	INSTR. TYPE	READY	DEST.	VALUE	SPEC.
ROB0	No					
ROB1	Yes	MULTD F4, F0, F2	Exec. Started	F4	*	No
ROB2	Yes	SD F4, O(R1)	No (issued)	M[0+[R1]]	—	No
ROB3	Yes	SUBI R1, R1, #8	Exec. Started	R1	—	No
ROB4	Yes	BNEZ R1, Loop	No (issued)	—	—	No
ROB5	Yes	LD F0, O(R1)	No (issued)	F0	—	Yes
ROB6	Yes	MULTD F4, F0, F2	No (issued)	F4	—	Yes
ROB7	Yes	SD F4, O(R1)	No (issued)	M[0+[R1]]	—	Yes

*: the value $M[0+[R1]] * F2$ is under computation. The head is at ROB1 and the tail is at ROB0.

Next cycle:

- ROB0 commits:
 - * LD completed \rightarrow value is written to the Register File (F0 updated).
 - * ROB0 is freed (Busy = No).
- MULTD at ROB1 finally starts execution (*): now that F0 is available, it can compute $F4 = F0 \times F2$.
- SUBI at ROB3 also continues execution.

Instructions are allowed to move forward even if speculative ones are pending behind.

- **ROB Remains Full (and Critical Situation)**

ROB#	BUSY	INSTR. TYPE	READY	DEST.	VALUE	SPEC.
ROB0	Yes	SUBI R1, R1, #8	No (issued)	R1	—	Yes
ROB1	Yes	MULTD F4, F0, F2	Exec. Started	F4	*	No
ROB2	Yes	SD F4, O(R1)	No (issued)	M[0+[R1]]	—	No
ROB3	Yes	SUBI R1, R1, #8	Exec. Started	R1	—	No
ROB4	Yes	BNEZ R1, Loop	No (issued)	—	—	No
ROB5	Yes	LD F0, O(R1)	No (issued)	F0	—	Yes
ROB6	Yes	MULTD F4, F0, F2	No (issued)	F4	—	Yes
ROB7	Yes	SD F4, O(R1)	No (issued)	M[0+[R1]]	—	Yes

*: the value $M[0+[R1]] * F2$ is still under computation. The head is still at ROB1 and the tail is stall.

The situation is:

- ROB remains fully occupied because new speculative instructions filled it (ROB5 to ROB7).

- The ongoing speculative execution **relies heavily** on correct branch prediction (**BNEZ**).
- **Problem:** If the branch was **mispredicted**, all speculative instructions (**ROB5**, **ROB6**, **ROB7**) must be **flushed**.

This example shows how the ROB manages **execution overlap**, **hazard resolution**, **speculation**, and **in-order commit** even under complex dependencies.

4 Performance Evaluation

All formulas and metrics can be found in the Index section, page 584, under the word F (Formula).

4.1 Basic Concepts and Performance Metrics

Performance can be considered from two key viewpoints:

- **Purchasing Perspective.** Given multiple machines, the goal is to determine:
 - *Which machine has the best performance.*
 - *Which machine has the lowest cost.*
 - *Which machine offers the best performance-to-cost ratio.*
- **Design Perspective.** When facing different architectural or technological options, the aim is to identify:
 - *Which design achieves the greatest performance improvement.*
 - *Which design minimizes cost.*
 - *Which design optimize the performance-to-cost balance.*

Both perspective require a **basic for comparison** and **clear performance metrics**. The main goal is to understand how architectural choices impact both **performance** and **cost**.

📘 Two Fundamental Notions of Performance

1. **Execution Time (Latency, Response Time):** refers to the **time needed to complete a task**.
Lower execution time = better.
2. **Throughput (Bandwidth):** measures the **number of tasks completed per unit of time** (e.g., jobs per hour).
Higher throughput = better.

Response time and throughput often **oppose** each other. Optimizing for one can degrade the other.

Example 1: Planes Comparison, Concorde vs Boeing 747

Concorde is a supersonic airliner.

- **Speed:** Concorde is faster (1350 mph vs 610 mph).
- **Execution Time:** Concorde takes less time to fly Washington DC to Paris (3h vs 6.5h).
- **Throughput:** Boeing transports more passengers per mile (higher “passenger-miles per hour”).

Basic Definitions and Calculations

- **Performance Formula**

$$\text{Performance} = \frac{1}{\text{Execution Time}} \quad (10)$$

- **Relative Performance.** If X is $n\%$ faster than Y , then:

$$\text{Performance}(X) = \left(1 + \frac{n}{100}\right) \times \text{Performance}(Y) \quad (11)$$

Or:

$$\text{Execution Time}(Y) = \left(1 + \frac{n}{100}\right) \times \text{Execution Time}(X) \quad (12)$$

Example 2

If machine A executes a program in 10 sec and machine B executes same program in 15 sec: A is 50% faster than B or A is 33% faster than B? Using equation 12, if A is 50% faster than B, then it should be true that:

$$\begin{aligned} \text{Execution Time}(B) &= \left(1 + \frac{n}{100}\right) \times \text{Execution Time}(A) \\ 15 &= \left(1 + \frac{50\%}{100}\right) \times 10 \\ 15 &= \frac{3}{2} \times 10 \\ 15 &= \frac{30}{2} \checkmark \end{aligned}$$

Clock Cycles and Clock Frequency

Modern processors operate by executing instructions synchronously according to a **clock signal**:

- **Clock Cycle Time T_{CLK}** : the **duration** of a single clock cycle, measured in seconds (s).
- **Clock Frequency f_{CLK}** : the **number of clock cycles per second**, measured in Hertz (Hz). Defined as:

$$f_{\text{CLK}} = \frac{1}{T_{\text{CLK}}} \quad (13)$$

Some examples:

$$\begin{aligned} - f_{\text{CLK}} = 500 \text{ MHz} &\Leftrightarrow T_{\text{CLK}} = \frac{1}{500 \times 10^6} = \frac{1}{500'000'000} = 2 \text{ ns} \\ - f_{\text{CLK}} = 1 \text{ GHz} &\Leftrightarrow T_{\text{CLK}} = \frac{1}{1 \times 10^9} = \frac{1}{1'000'000'000} = 1 \text{ ns} \end{aligned}$$

⌚ Execution Time (CPU Time)

The **execution time** (also called **CPU time**) is the **total time** needed by a processor to complete a given program. Two main ways to compute CPU time:

- **Basic form:**

$$\text{CPU Time} = \text{Clock Cycles} \times T_{\text{CLK}} = \frac{\text{Clock Cycles}}{f_{\text{CLK}}} \quad (14)$$

- **Expanded form based on instructions:**

$$\begin{aligned} \text{CPU Time} &= \text{Instruction Count} \times \text{Cycles Per Instruction} \times T_{\text{CLK}} \\ &= \frac{\text{Instruction Count} \times \text{Cycles Per Instruction}}{f_{\text{CLK}}} \end{aligned} \quad (15)$$

- **Weighted CPU Time.** If the program is composed of different instruction types, with different CPI values, the **total CPU time must be weighted** by how much each instruction type contributes:

$$\text{CPU Time} = \left(\sum_{i=1}^n (\text{CPI}_i \times I_i) \right) \times T_{\text{CLK}} \quad (16)$$

Where:

- I_i is the number of instructions of type i .
- T_{CLK} is the clock cycle time.

The total CPU time is the sum of the CPU times **for each class of instruction**, weighted by how many times each instruction appears.

Note that weighted CPU time can be calculated using the classic formula, but **using a weighted CPI** (shown on page 249):

$$\text{CPU Time} = \frac{\text{IC} \times \text{CPI Weighted}}{f_{\text{CLK}}}$$

⌚ Cycles Per Instruction (CPI) and Instructions Per Cycle (IPC)

- **Clocks Per Instruction (CPI)** (clock cycles per instruction or clocks per instruction) is the **average number of clock cycles** each instruction requires.

$$\text{CPI} = \frac{\text{Total Clock Cycles}}{\text{Instruction Count (IC)}} \quad (17)$$

Where the Total Clock Cycles is:

$$\text{Total Clock Cycles} = \text{IC} + 4 + \text{Stall Cycles}$$

And the Instruction Count is the total number of instructions executed.

- **Instructions Per Cycle (IPC)** is the reciprocal of CPI:

$$\text{IPC} = \frac{1}{\text{CPI}} \quad (18)$$

- **Mixed Programs.** In a real program, we have a mix of operations. Each instruction type can have a different CPI, thus the **total average performance depends on:**

1. **How often** each instruction type appears (its **frequency**).
2. **How many cycles** each type needs (its **CPI**).

Because **not all instructions appear equally**, we must give **more weight** to instructions that appear more frequently. The **Weighted Average CPI** formula is:

$$\text{CPI} = \sum_{i=1}^n (\text{CPI}_i \times F_i) \quad (19)$$

Where:

- CPI_i is the CPI of instruction type i .
- F_i is the **frequency of instruction** type i (percentage). It is calculated as:

$$F_i = \frac{I_i}{\text{IC}} \quad (20)$$

Where:

- * I_i is the number of instructions of type i .
- * IC is the total **Instruction Count** (total number of instructions in the program).

Example 3: CPI and CPU Time Calculation

Given a program with **100 instructions** and the following instruction mix (500 MHz clock frequency):

Instruction Type	Frequency	CPI
ALU	43%	1
Load	21%	4
Store	12%	4
Branch	12%	2
Jump	12%	2

First, convert frequencies to decimals:

- 43% → 0.43
- 21% → 0.21
- 12% → 0.12
- 12% → 0.12
- 12% → 0.12

Now apply the CPI weighted formula 19:

$$\begin{aligned}
 \text{CPI} &= \sum_{i=1}^n (\text{CPI}_i \times F_i) \\
 &= (0.43 \cdot 1) + (0.21 \cdot 4) + (0.12 \cdot 4) + (0.12 \cdot 2) + (0.12 \cdot 2) \\
 &= 0.43 + 0.84 + 0.48 + 0.24 + 0.24 \\
 &= 2.23
 \end{aligned}$$

Thus, the **average CPI** is 2.23. For CPU time, we can use the weighted CPI, but we must calculate the Clock Cycle Time T_{CLK} :

$$\begin{aligned}
 \text{CPU Time} &= \text{IC} \times \text{CPI Weighted} \times T_{\text{CLK}} \\
 &= 100 \times 2.23 \times \frac{1}{(500 \times 10^6)} \\
 &= 223 \times \frac{1}{(500 \times 10^6)} \\
 &= 223 \times 0.000'000'002 \text{ sec} \\
 &= 223 \times 2 \text{ ns} \\
 &= 446 \text{ ns}
 \end{aligned}$$

⚡ MIPS (Millions of Instructions Per Second)

MIPS measures how many millions of instructions a processor can execute per second. Two equivalent formulas:

- Using frequency and CPI:

$$\text{MIPS} = \frac{f_{\text{CLK}}}{\text{CPI} \times 10^6} \quad (21)$$

- Using instruction count and execution time:

$$\text{MIPS} = \frac{\text{Instruction Count}}{\text{Execution Time} \times 10^6} \quad (22)$$

Where the Execution Time is the CPU Time.

Higher MIPS does not necessarily mean better performance, because MIPS does not account for instruction complexity or actual work done.

4.2 Amdahl's Law

Amdahl's Law describes the **limits** of performance improvement when only **part of a system** is enhanced. Key principles are:

- **Speedup** of a program is **limited** by the part of the program that **cannot be improved**.
- **Making the common case fast** gives the best returns.

Thus, **no matter how much we optimize a part of a system**, the total speedup is **bounded**.

Formal Definition

Let:

- F : fraction of execution time **affected** by the enhancement.
- S : speedup **factor** for the enhanced part.

Then:

- **New Execution Time**:

$$\text{ExTime(w/ E)} = \left((1 - F) + \frac{F}{S} \right) \times \text{ExTime(w/o E)} \quad (23)$$

Where:

- “w/ E” = “with Enhancement”
- “w/o E” = “without Enhancement”

- **Overall Speedup**:

$$\text{Speedup(E)} = \frac{1}{(1 - F) + \frac{F}{S}} \quad (24)$$

Where:

- $(1 - F)$ is a portion of execution time **not improved**.
- $\frac{F}{S}$ is the improved portion, now **accelerated**.

In many books, we also see:

- **FractionE**: fraction of computation time accelerated.
- **SpeedupE**: speedup factor for the enhanced portion.

Thus:

$$\text{Speedup}_{\text{overall}} = \frac{1}{(1 - \text{FractionE}) + \frac{\text{FractionE}}{\text{SpeedupE}}} \quad (25)$$

Example 4: Amdahl's Law

Let us consider an enhancement for a CPU resulting 10 time faster on computation than the original one but the original CPU is busy with computation only 40% of the time. What is the overall speedup gained by introducing the enhancement?

Given:

- FractionE: 0.4 (40% of the time can be accelerated).
- SpeedupE: 10 (enhancement is 10x faster)

Apply Amdahl's Law:

$$\begin{aligned}\text{Speedup}_{\text{overall}} &= \frac{1}{(1 - 0.4) + \frac{0.4}{10}} \\ &= \frac{1}{0.6 + 0.04} \\ &= \frac{1}{0.64} \\ &= 1.56\end{aligned}$$

Thus, even if the improvement is 10x, the overall speedup is only about 1.56x.

If we **only speed up a small part** (small F), the **impact is small**. To achieve **large speedup**, we must:

- Speed up a **large fraction** of the execution time.
- Or make the enhancement extremely fast ($S \rightarrow \infty$).

To reach the maximum theoretical speedup, if $S \rightarrow \infty$ (perfect acceleration), then:

$$\text{Speedup}_{\max} = \frac{1}{1 - F} \quad (26)$$

For example if only 40% can be improved, even with infinite speedup, the maximum is:

$$\text{Speedup}_{\max} = \frac{1}{1 - 0.4} = \frac{1}{0.6} \approx 1.67$$

This shows **the real bottleneck**: the non-improved part.

4.3 Pipelined Processors

Pipelining affects CPU performance **differently** than simple frequency or instruction count.

- ✓ **Pipelining increases instruction Throughput** (number of instructions completed per unit time).
- ✗ **Pipelining does not reduce Execution Time** (latency) of a single instruction.

In fact slight latency **increase** can happen due to:

- **Imbalance** among **pipeline stages** (some stages are slower, limiting the clock frequency).
- **Overhead** from **pipeline control** (registers, clock skew, pipeline latch delays).

Also, we must remember that the **clock cycle time must be long enough for the slowest stage to finish**. Otherwise, the slowest stage will not finish its work before the clock ticks again, and the pipeline will become inconsistent or crash. So:

$$T_{CLK} > \text{time of the slowest stage}$$

So, if we **want to make the CPU faster**, we must **balance the stages** (make all stages take about the same time). Otherwise, one slow stage will limit the entire clock speed, even if other stages are very fast. This is why sometimes CPU designers split a slow stage into two smaller stages to balance the pipeline better!

≡ Performance Metrics for Pipelining

Define:

- IC as Instruction Count
- CPI as Cycles Per Instruction
- IPC as Instructions Per Clock $\frac{1}{CPI}$

We have:

- **Number of clock cycles in pipelined execution:**

$$\text{Clock Cycles} = \text{IC} + \text{Stall Cycles} + \text{Pipeline fill penalty (4 cycles)} \quad (27)$$

For a 5-stage pipeline like MIPS or RISC-V. And the CPI is:

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{IC}} = \frac{(\text{IC} + \text{Stall Cycles} + 4)}{\text{IC}} \quad (28)$$

- **MIPS calculation:**

$$\text{MIPS} = \frac{f_{clock}}{\text{CPI} \times 10^6} \quad (29)$$

Example 5

Given:

- IC = 5 instructions
- Stall cycles = 2
- Clock = 500 MHz

Step-by-step:

- Clock cycles:

$$5 + 2 + 4 = 11$$

- CPI:

$$\frac{11}{5} = 2.2$$

- MIPS:

$$\frac{500 \times 10^6}{2.2 \times 10^6} \approx 227$$

Thus, the effective throughput is about **227 MIPS**.

⌚ Performance in Loops

Suppose we have:

- ***n* iterations**
- Loop of ***m* instructions** per iteration
- ***k* stalls** per iteration

Formulas per iteration:

- **IC per iteration:**

$$IC_{\text{per iter}} = m \quad (30)$$

- **Clock Cycles per iteration:**

$$\# \text{ Clock Cycles}_{\text{per iter}} = m + k + 4 \quad (31)$$

- **Cycles Per Instruction per iteration:**

$$CPI_{\text{per iter}} = \frac{(m + k + 4)}{m} \quad (32)$$

- **MIPS per iteration:**

$$MIPS = \frac{f_{\text{clock}}}{CPI_{\text{per iter}} \times 10^6} \quad (33)$$

∞ Asymptotic Performance (Many Iterations)

If the loop runs for $n \rightarrow \infty$ iterations:

- **Total Instruction Count per iteration:**

$$IC_{AS} = \text{Instruction Count}_{AS} = m \times n \quad (34)$$

- **Total Clock Cycles per iteration:**

$$\begin{aligned} \# \text{ Clock Cycles}_{AS} &= IC_{AS} + \# \text{ Stall Cycles}_{AS} + 4 \\ &= (m \times n) + (k \times n) + 4 \end{aligned} \quad (35)$$

- **Cycles Per Instruction per iteration:**

$$CPI_{AS} = \lim_{n \rightarrow \infty} \frac{\# \text{ Clock Cycles}_{AS}}{IC_{AS}} = \lim_{n \rightarrow \infty} \frac{m \times n + k \times n + 4}{m \times n} = \frac{m + k}{m} \quad (36)$$

The fixed 4 cycles become irrelevant when n becomes very large.

- **MIPS per iteration:**

$$MIPS_{AS} = \frac{f_{clock}}{CPI_{AS} \times 10^6} \quad (37)$$

¶ Ideal vs Realistic Pipelining

The ideal CPI on a pipelined processor would be 1, but stalls cause the pipeline performance to degrade from the ideal performance, so we have:

- **Ideal pipelining:**

$$\text{Ideal CPI} = 1 \quad (38)$$

1 instruction per cycle after filling the pipeline.

- **Realistic pipelining:**

$$CPI = 1 + (\text{stall cycles per instruction}) \quad (39)$$

Stalls are caused by:

- Structural Hazards (e.g., hardware resources conflict)
- Data Hazards (e.g., instruction dependencies)
- Control Hazards (e.g., branches)
- Memory Stalls (e.g., cache misses)

► Pipeline Speedup

Speedup measures how much **pipelining** performance compared to **unpipelined execution**. General formula:

$$\text{Pipeline Speedup} = \frac{\text{Avg Exec Time Unpipelined}}{\text{Avg Exec Time Pipelined}} = \frac{\text{CPI}_{\text{unp}} \times T_{\text{clk, unp}}}{\text{CPI}_{\text{pipe}} \times T_{\text{clk, pipe}}} \quad (40)$$

If we **ignore the cycle time overhead** of pipelining and we assume the **stages are perfectly balanced**, the clock cycle time of unpipelined/pipelined processors can be equal, so:

$$\text{Pipeline Speedup} = \frac{\text{CPI}_{\text{unp}}}{1 + \text{CPI}_{\text{pipe}}} = \frac{\text{CPI}_{\text{unp}}}{1 + \text{Stall Cycles per Instruction}} \quad (41)$$

If we assume that **each instruction takes the same number of cycles**, which must be **equal to the number of pipeline stages** (called pipeline depth):

$$\text{Pipeline Speedup} = \frac{\text{Pipeline Depth}}{1 + \text{Stall Cycles per Instruction}} \quad (42)$$

If there were **no pipeline stalls (ideal case)**, this leads to the intuitive result that pipelining improves performance by the depth of the pipeline:

$$\text{Speedup} = \text{Pipeline Depth} \quad (43)$$

Thus, a 5-stage pipeline **could at best achieve 5x speedup only if there are no stalls**.

↳ Impact of Branches

Branches cause **pipeline stalls**, and their effect is quantified as:

$$\text{Pipeline Speedup} = \frac{\text{Pipeline Depth}}{1 + \text{Branch Frequency} \times \text{Branch Penalty}} \quad (44)$$

- **Branch Frequency:** fraction of **instructions that are branches**.
- **Branch Penalty:** number of **stall cycles** per branch.

Where the Pipeline Stall Cycles per Instruction due to Branches (PSCIB) is:

$$\text{PSCIB} = \text{Branch Frequency} \times \text{Branch Penalty} \quad (45)$$

4.4 Memory Hierarchy

In modern processors we organize memories as a **hierarchy**.

- At the **top**, **smallest** but **fastest** memories (e.g., L1 cache, CPU registers).
- At the **bottom**, **largest** but **slowest** memories (e.g., DRAM, disk storage).

The idea is to keep most data in fast memories for efficiency.

- **Hit**. The data we want is **found** in the upper (faster) memory level.
- **Hit Rate**. Probability (fraction) of memory accesses that result in a hit.

$$\text{Hit Rate} = \frac{\# \text{ hits}}{\# \text{ memory accesses}} \quad (46)$$

- **Hit Time**. Time to **access** the upper memory level and check whether it's a hit.
- **Miss**. The data is **not found** in the upper memory, we must access the slower lower level.
- **Miss Rate**. Probability (fraction) of accesses that result in a miss.

$$\text{Miss Rate} = \frac{\# \text{ misses}}{\# \text{ memory accesses}} \quad (47)$$

Important property:

$$\text{Hit Rate} + \text{Miss Rate} = 1 \quad (48)$$

- **Miss Time**. Total time when there is a miss:

$$\text{Miss Time} = \text{Hit Time} + \text{Miss Penalty} \quad (49)$$

Where **Miss Penalty** is extra time needed to fetch data from the lower level and update upper cache.

AMAT (Average Memory Access Time)

AMAT (Average Memory Access Time) measures the **average time the processor needs to access memory** (whether there's a hit or a miss).

$$\text{AMAT} = \text{Hit Rate} \times \text{Hit Time} + \text{Miss Rate} \times \text{Miss Time}$$

It combines both: the fast accesses (hits) and the slow accesses (misses). Since equation 49, we can substitute into the AMAT formula:

$$\text{AMAT} = \text{Hit Rate} \times \text{Hit Time} + \text{Miss Rate} \times (\text{Hit Time} + \text{Miss Penalty})$$

But by definition:

$$\text{Hit Rate} + \text{Miss Rate} = 1$$

Thus:

$$\text{AMAT} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty} \quad (50)$$

This is the most important and most used AMAT formula!

How to improve cache performance?

To reduce AMAT (and improve system speed), we can:

1. **Reduce Hit Time:** Make the cache smaller, simpler, or closer to the CPU.
2. **Reduce Miss Rate:** Improve cache organization (better replacement, associativity, prefetching).
3. **Reduce Miss Penalty:** Use faster lower levels (e.g., add an L2 or L3 cache).

Unified Cache vs Separate Instruction

There are two major architectures for **L1 Cache**:

- **Unified Cache:** Same L1 cache for both instructions and data.
- **Separate I\$ and D\$ (Harvard Architecture):** Different L1 caches for instructions (I\$) and data (D\$).

For **separate caches** (Harvard Architecture):

- Instruction cache and Data cache have **different miss rates**.
- **Accesses are divided** between instruction and data operations.

This is the AMAT (Average Memory Access Time) for Harvard architectures:

$$\text{AMAT}_{\text{Harvard}} = (\% \text{Instr}) \times (\text{Hit Time} + \text{Miss Rate}_{\text{I\$}} \times \text{Miss Penalty}) + (\% \text{Data}) \times (\text{Hit Time} + \text{Miss Rate}_{\text{D\$}} \times \text{Miss Penalty}) \quad (51)$$

Example 6: Harvard Architecture

Assumptions:

- 16 KB I\$ with Miss Rate = 0.64%
- 16 KB D\$ with Miss Rate = 6.47%
- 32 KB Unified Cache with Aggregate Miss Rate = 1.99%
- 33% of accesses are loads/stores, so:
 - 75% accesses are instructions
 - 25% accesses are data
- Hit Time = 1 cycle, Miss Penalty = 50 cycles
- **Data hit has 1 more stall for Unified cache (only one port)**

Which cache is better?

Calculate Harvard AMAT (equation 51):

$$\begin{aligned}
 \text{AMAT}_{\text{Harvard}} &= 0.75 \times (1 + 0.0064 \times 50) + 0.25 \times (1 + 0.0647 \times 50) \\
 &= 0.75 \times (1 + 0.32) + 0.25 \times (1 + 3.235) \\
 &= 0.75 \times 1.32 + 0.25 \times 4.235 \\
 &= 0.99 + 1.05875 \\
 &= 2.04875 \approx 2.05 \text{ cycles}
 \end{aligned}$$

Calculate AMAT (Unified cache):

$$\begin{aligned}
 \text{AMAT} &= 0.75 \times (1 + 0.0199 \times 50) + 0.25 \times (1 + 1 + 0.0199 \times 50) \\
 &= 0.75 \times (1 + 0.995) + 0.25 \times (2 + 0.995) \\
 &= 0.75 \times 1.995 + 0.25 \times 2.995 \\
 &= 1.49625 + 0.74875 \\
 &= 2.245 \approx 2.24 \text{ cycles}
 \end{aligned}$$

So **Harvard (separate I\$ and D\$)** gives **better performance** than Unified Cache, because AMAT (Average Memory Access Time) is less.

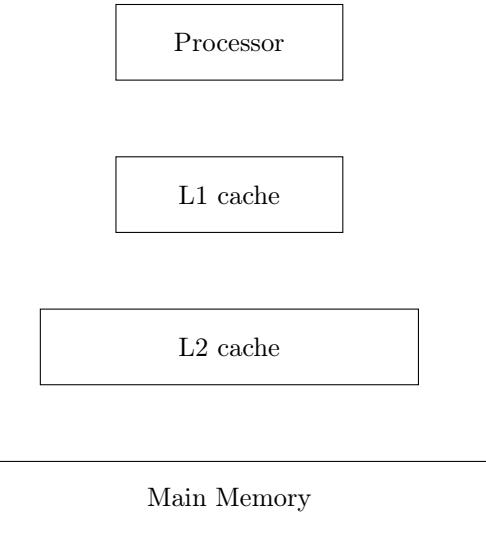
⌚ Miss Penalty Reduction: Second Level Cache (L2)

L1 cache must be **very small and fast** (to match the fast CPU cycle time). But small cache, relatively **high miss rate**. Every miss in L1 would normally mean **access to slow main memory** (very expensive!).

The solution is to **insert a larger, slower L2 cache** between L1 and main memory. Thus:

- If L1 misses, we first **try to find the data in L2**.
- Only if L2 also misses, we go to **main memory**.

This **reduces the effective miss penalty** seen by the CPU.



When using both L1 and L2 caches, first:

$$\text{AMAT} = \text{Hit Time}_{L1} + \text{Miss Rate}_{L1} \cdot \text{Miss Penalty}_{L1}$$

where:

$$\text{Miss Penalty}_{L1} = \text{Hit Time}_{L2} + \text{Miss Rate}_{L2} \cdot \text{Miss Penalty}_{L2}$$

Thus, expanding:

$$\begin{aligned} \text{AMAT} = & \text{Hit Time}_{L1} + \text{Miss Rate}_{L1} \cdot \\ & (\text{Hit Time}_{L2} + \text{Miss Rate}_{L2} \cdot \text{Miss Penalty}_{L2}) \end{aligned}$$

or simplified the **AMAT** is:

$$\begin{aligned} \text{AMAT} = & \text{Hit Time}_{L1} + \\ & \text{Miss Rate}_{L1} \cdot \text{Hit Time}_{L2} + \\ & \text{Miss Rate}_{L1} \cdot \text{Miss Rate}_{L2} \cdot \text{Miss Penalty}_{L2} \end{aligned} \tag{52}$$

Each term captures: immediate L1 access, possible L2 access, and final access to main memory if L2 also misses.

⌚ Local vs Global Miss Rates

- **Local Miss Rate:** misses divided accesses at that specific cache level.
For example:

- L1 Local Miss Rate = Misses at L1 ÷ Accesses to L1
- L2 Local Miss Rate = Misses at L2 ÷ Accesses to L2

- **Global Miss Rate:** misses at a level relative to **total CPU accesses**.

- For L1:

$$\text{Global Miss Rate}_{L1} = \text{Miss Rate}_{L1} \quad (53)$$

- For L2:

$$\text{Global Miss Rate}_{L2} = \text{Miss Rate}_{L1} \cdot \text{Miss Rate}_{L2} \quad (54)$$

Global miss rate is what really matters, because it tells us **what percentage of memory accesses are so unlucky** that they: miss in L1 (fast small cache), miss in L2 (bigger slower cache), go down to main memory (very slow).

Example 7: Local vs Global Miss Rates

Let us consider a computer with a L1 cache and L2 cache memory hierarchy. Suppose that in 1000 memory references there are 40 misses in L1 and 20 misses in L2. *What are the various miss rates?*

- L1 Miss Rate:

$$\frac{40}{1000} = 4\%$$

- L2 Local Miss Rate:

$$\frac{20}{40} = 50\%$$

- L2 Global Miss Rate:

$$\text{Miss Rate}_{L1} \cdot \text{Miss Rate}_{L2} = 4\% \cdot 50\% = 2\%$$

Thus, only 2% of all CPU memory accesses end up reaching main memory.

Adding an L2 cache:

- ✓ Increases hit rate (fewer accesses to slow main memory),
- ✓ Reduces effective miss penalty,
- ✓ Improves CPU performance by reducing memory stall cycles.

Impact of Memory Hierarchy on CPU Execution Time

Without cache every memory access would go directly to main memory (very slow). With cache:

- Most accesses are **fast** (cache hits),
- Only a few accesses are **slow** (cache misses).

Thus, **memory stalls** (waiting for memory) become part of the CPU execution time. Therefore, the CPU total execution time depends on:

- **CPU execution cycles** (doing normal instruction work),
- **Memory stall cycles** (waiting for slow memory).

The formal formula for CPU time is:

$$\text{CPU Time} = (\text{CPUexec Cycles} + \text{Memory Stall Cycles}) \times T_{CLK}$$

Where:

- T_{CLK} : clock cycle time.
- **CPUexec Cycles**: cycles needed for normal instruction execution (ALU operations, Load/Store assuming all hits).

$$\text{CPUexec Cycles} = \text{IC} \times \text{CPI}_{\text{exec}} \quad (55)$$

Where:

- IC: Instruction Count.
- CPI_{exec}: ideal CPI **without considering cache misses**.

- **Memory Stall Cycles**: extra cycles lost because of memory misses. Memory stall cycles are caused by misses times penalty per miss. Specifically:

$$\text{Memory Stall Cycles} = \text{IC} \times \text{Misses per Instruction} \times \text{Miss Penalty} \quad (56)$$

And:

$$\begin{aligned} \text{Misses per Instruction} &= \frac{\# \text{ misses}}{\text{IC}} \times \frac{\# \text{ mem. accesses}}{\text{IC}} \times \text{Miss Rate} \\ &= \text{MAPI} \times \text{Miss Rate} \end{aligned} \quad (57)$$

(where MAPI is Memory Accesses per Instruction); Therefore:

$$\text{Memory Stall Cycles} = \text{IC} \times \text{MAPI} \times \text{Miss Rate} \times \text{Miss Penalty} \quad (58)$$

Putting it all together, we get the **Full CPU Time formula**:

$$\text{CPU Time} = \text{IC} \cdot (\text{CPI}_{\text{exec}} + \text{MAPI} \cdot \text{Miss Rate} \cdot \text{Miss Penalty}) \cdot T_{CLK} \quad (59)$$

This formula shows **both**: the “normal” CPU work, and the “extra time” lost due to cache misses. These are also (**ideal**) **special cases**:

- ✓ **Ideal cache (100% hits):**

$$\text{CPU Time} = \text{IC} \times \text{CPI}_{\text{exec}} \times T_{\text{CLK}}$$

- ✗ **No cache (100% misses):**

$$\text{CPU Time} = \text{IC} \times (\text{CPI}_{\text{exec}} + \text{MAPI} \times \text{Miss Penalty}) \times T_{\text{CLK}}$$

A Including Pipeline Stalls. If we also want to consider **pipeline stalls** (caused by hazards like data dependencies), the full CPU Time becomes:

$$\begin{aligned} \text{CPU Time} = & \text{IC} \times (\text{CPI}_{\text{exec}} + \text{Stalls per Instruction} + \\ & \text{MAPI} \times \text{Miss Rate} \times \text{Miss Penalty}) \times \\ & T_{\text{CLK}} \end{aligned} \quad (60)$$

Where **Stalls per Instruction** is the **average extra cycles lost per instruction due to pipeline hazards**. So we can see that **pipeline hazards** and **memory “misses”** both **contribute to processor slowdowns**.

This is important because even if we build the fastest CPU core in the world, if our memory system is slow, our real performance will be bad! Memory system design (good caches, good pipelines) is **essential** for real CPU performance.

⚠️ Memory Stalls in L1 and L2 Caches

In a **hierarchical cache system**:

- A memory access **first** tries L1,
- If **L1 misses**, tries **L2**,
- If **L2 also misses**, goes to **main memory** (very slow).

Each **level of cache** can **introduce extra stall cycles**:

- Stall cycles after an L1 miss → time to access L2.
- Stall cycles after an L2 miss → time to access main memory.

Average **memory stall cycles per instruction** (considering both L1 and L2) is:

$$\text{Memory stall cycles per instruction} = (\text{Misses}_{L1} \cdot \text{Hit Time}_{L2}) + (\text{Misses}_{L2} \cdot \text{Miss Penalty}_{L2}) \quad (61)$$

Where:

- Misses_{L1} = Misses per instruction at L1.
- Misses_{L2} = Misses per instruction at L2.
- Hit Time_{L2} = Time to access L2 cache (after L1 miss).
- Miss Penalty_{L2} = Time to access main memory (after L2 miss).
- $(\text{Misses}_{L1} \cdot \text{Hit Time}_{L2})$ = The **first term** is for accesses that hit in L2 (still slower than L1 but faster than main memory).
- $(\text{Misses}_{L2} \cdot \text{Miss Penalty}_{L2})$ = The **second term** is for accesses that miss in both L1 and L2 (very slow).

If we want to **find misses per instruction**, they depend on MAPI (Memory Accesses Per Instruction) and miss rates:

- **Misses at L1 per instruction**

$$\text{Misses}_{L1} = \text{MAPI} \times \text{Miss Rate}_{L1}$$

- **Misses at L2 per instruction**

$$\text{Misses}_{L2} = \text{MAPI} \times \text{Global Miss Rate}_{L2}$$

Where:

$$\text{Global Miss Rate}_{L2} = \text{Miss Rate}_{L1} \times \text{Miss Rate}_{L2}$$

Considering everything, the **Final Full CPU Time Formula** is:

$$\text{CPU Time} = IC \times (\text{CPI}_{\text{exec}} + \text{Memory stall cycles per instruction}) \times T_{\text{CLK}} \quad (62)$$

Where:

$$\begin{aligned} \text{Memory stall cycles per instruction} = & (\text{MAPI} \times \\ & \text{Miss Rate}_{L1} \times \\ & \text{Hit Time}_{L2}) + \\ & (\text{MAPI} \times \quad (63) \\ & \text{Miss Rate}_{L1} \times \\ & \text{Miss Rate}_{L2} \times \\ & \text{Miss Penalty}_{L2}) \end{aligned}$$

Putting it all together, the CPU Time is:

$$IC \cdot (\text{CPI}_{\text{exec}} + \text{MAPI} \cdot (\text{MR}_{L1} \cdot \text{HT}_{L2} + \text{MR}_{L1} \cdot \text{MR}_{L2} \cdot \text{MP}_{L2})) \cdot T_{\text{CLK}} \quad (64)$$

This is the most complete formula that accounts for:

- L1 and L2 cache effects,
- Main memory access,
- How memory stalls degrade CPU performance.

5 VLIW (Very Long Instruction Word)

5.1 Introduction

Traditional compilers use **static code scheduling** to exploit Instruction-Level Parallelism (ILP). Key tasks of the compiler:

- Detect **parallelizable** instructions considering: Hardware resource constraints and Data dependencies.
- **Schedule** instructions to execute **in parallel** when possible.
- Otherwise, **insert NOPs** (No Operations) if no safe parallel execution is possible.

Statically scheduled processors trust the compiler to “fill the pipeline” and avoid hazards at compile time.

▀ VLIW Processors: Alternative Way to Extract ILP

VLIW (Very Long Instruction Word) processors are a class of architectures designed to **execute multiple operations in parallel during a single clock cycle**, but unlike superscalar processors, they **rely on the compiler** rather than on dynamic hardware mechanisms to **detect and schedule parallelism**.

The fundamental idea behind VLIW is to **group several independent operations together into a single long instruction word**, called a **bundle**. This bundle is **composed of several fixed slots**, each one **corresponding to a different functional unit in the processor**, such as an integer ALU, a floating-point unit, a load/store unit, or a branch unit.

For example, in a 5-issue VLIW processor, a single bundle would typically carry up to five operations: integer, floating-point, load/store, branch, etc.

The **key difference** compared to **traditional superscalar** processors lies in who decides what can run in parallel. In superscalar designs, the hardware dynamically analyzes dependencies between instructions at runtime, which requires complex circuitry for hazard detection and scheduling. In VLIW architectures, this **analysis is performed entirely at compile time**. The **compiler is responsible for**:

- **Statically identify** independent operations.
- **Solve structural hazards** (e.g., two operations trying to use the same hardware unit).
- **Solve data hazards** (dependencies between instructions).
- Insert **NOPs** when necessary.

When no useful instruction is available for a particular slot, the compiler inserts a NOP (no-operation). In cases where conflicts cannot be avoided, NOPs are inserted into the bundle to fill empty slots.

💡 Why move to Compiler?

The problem with Superscalar is that the hardware requires complex dynamic scheduling and dependency checking, which **costs area and power**.

VLIW Idea is:

- Push **complexity** to the **compiler**.
- **Compiler groups parallel operations** into a **single bundle**.
- No need for runtime dependency checking anymore.

The VLIW paradigm is characterized by the use of **very wide instruction words**, each containing multiple independent operations (“syllables”). A multiple-issue VLIW processor typically features specialized units for integer, floating-point, memory access, and branch instructions. For example, a 4-issue VLIW processor will have four operation slots per bundle, each connected to a different unit.

☰ Multiple-issue VLIW: Operation Latencies

An important aspect in VLIW scheduling is the **management of operation latencies**. Each operation type may require a **different number of clock cycles to complete**. Integer operations, memory accesses, and floating-point computations may all have varying latencies, and these differences must be carefully considered during scheduling. The **compiler must plan the execution** so that operations complete in the correct order and without causing unnecessary stalls in the pipeline.

In summary, **VLIW architectures represent a shift of complexity from the hardware to the compiler**, aiming for more efficient and simpler processor designs while demanding sophisticated compiler techniques to fully exploit available parallelism.

⌚ Single Program Counter and Branch Management

In a VLIW processor, even though multiple operations are issued simultaneously, the architecture still uses a **single program counter** to fetch instructions. Each instruction fetched corresponds to a **bundle**, which can contain multiple parallel operations.

Importantly, within each bundle, there can be **at most one branch instruction** that affects control flow. This constraint simplifies the handling of branches, making it easier for the processor to predict and manage the program’s execution path.

■ Shared Multi-Ported Register File

Since a VLIW processor issues multiple operations in parallel, the **register file must support multiple simultaneous reads and writes**. In a typical 4-issue VLIW machine, the register file needs to have enough ports to read **eight source operands** and write **four destination results** every clock cycle.

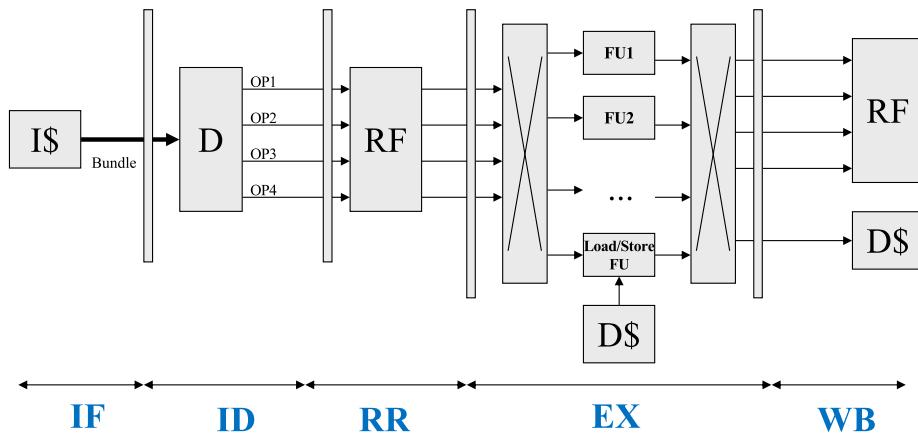
This requirement implies a **shared, multi-ported register file**, which is one of the non-trivial aspects of the hardware design. It ensures that all functional units can access their operands without creating bottlenecks.

■ Importance of Parallelism and the Use of NOPs

To achieve high performance, the **compiler must ensure that the source code has enough parallelism** to fill all the available operation slots in each bundle. When **sufficient independent instructions are not available**, **NOPs are inserted** into the empty slots. This insertion preserves the fixed structure of the bundle but can lead to inefficiencies if the application does not expose enough ILP.

■ Pipeline Organization

VLIW processors often use a **pipelined execution model** similar to traditional RISC architectures. A standard pipeline might include stages like **Instruction Fetch (IF)**, **Instruction Decode (ID)**, **Register Read (RR)**, **Execution (EX)**, and **Write Back (WB)**.



In the previous figure, a **5-stage pipeline** is used, with each bundle passing through these stages. Each operation within a bundle proceeds independently through the functional units connected to the pipeline.

■ Operation Dispatch and Decode

If the architecture dedicates **one functional unit per slot**, the decode stage becomes relatively simple. Each operation is directed straight to its corresponding functional unit for execution without the need for complex arbitration.

However, if there are **more functional units than slots**, meaning there is a surplus of parallel hardware resources, the architecture must include a **dispatch network**. This network is responsible for forwarding each operation, along with its operands, to the appropriate available functional unit. In this way, the design of the dispatch system depends heavily on the organization of functional units relative to the number of issue slots defined by the instruction bundle format.

5.2 Data dependencies

In VLIW architectures, **data dependencies**:

- True dependencies (RAW: Read After Write)
- Anti-dependencies (WAR: Write After Read)
- Output dependencies (WAW: Write After Write)

Are **handled entirely by the compiler**. During the compilation process, the **compiler analyzes the functional unit latencies and rearranges the instruction sequence to eliminate conflicts**. When dependencies cannot be resolved by reordering, the compiler inserts NOPs to delay operations and maintain correct execution order.

This static management of hazards is **crucial** because the hardware in VLIW processors does not dynamically detect or resolve data hazards during execution.

⌚ Operation Latency Management

In order to correctly schedule instructions, the compiler must also consider **operation latencies**.

If an **operation**, such as a multiplication, **requires multiple cycles** to complete, the **compiler must account for this delay explicitly**. It can either rearrange independent instructions to fill the delay or insert NOPs to stall dependent instructions until the result is ready.

For instance, if an instruction $C = A * B$ has a latency of two cycles, the compiler may schedule a NOP after it, before allowing an instruction that depends on C to execute. This **careful latency management ensures that data hazards are avoided** and that results are available when needed, even when using pipelined functional units.

⚠ WAR, WAW, and Structural Hazards

Besides true dependencies, **anti-dependencies (WAR)** and **output dependencies (WAW)** are also statically resolved at compile time. The compiler can manage these by:

- ✓ Appropriately choosing the timing of operations.
- ✓ Using **register renaming** techniques to avoid conflicts.

Structural hazards, such as two operations trying to use the same hardware resource simultaneously, are similarly **resolved** by the compiler, either **through scheduling decisions or resource management strategies**.

In addition to data and structural hazards, the compiler may provide **static branch prediction hints**. However, if the **prediction is wrong**, it is the **hardware's responsibility** to detect the misprediction and flush the pipeline, just as in *traditional architectures*.

☰ Maintaining In-Order Write-Back

An **important constraint** in VLIW processors is that **all operations in a bundle should complete their write-back stage at the same clock cycle**. This **synchronization prevents structural hazards in accessing the register file and avoids WAR and WAW hazards during register writes**.

If operations within a bundle have **different latencies**, they are **all forced to align to the longest latency operation** among them. Otherwise, if operations were allowed to complete independently, **out-of-order write-backs** would occur, and the **processor would require additional hardware to track and resolve register file conflicts**, precisely the kind of complexity VLIW designs try to avoid.

⚠ Register Pressure

One important issue in VLIW architectures is the problem of **register pressure**. Because VLIW bundles can issue multiple operations simultaneously, and because many **operations have multicycle latencies**, a **large number of registers may be occupied at the same time by intermediate results waiting to complete**.

For instance, if a bundle contains two floating-point operations, each requiring five clock cycles to complete, their destination registers remain occupied for the entire duration of those five cycles. Since a VLIW processor continues issuing new bundles each cycle, this quickly accumulates a large number of active registers. In the given example, five bundles are issued while the first two results are still pending, resulting in up to twenty operations competing for register resources.

The situation is further exacerbated by **register renaming**, a technique used by the compiler to solve WAR (Write After Read) and WAW (Write After Write) hazards statically. Register renaming:

- ✓ **Avoids conflicts**
- ✗ **Increases the number of physical register needed**, thereby amplifying register pressure.

Managing register pressure is therefore a **crucial task** for the compiler, often requiring careful allocation strategies or even **spilling values temporarily to memory** when the **register file becomes saturated**.

⚠ Dynamic Events in VLIW

Although VLIW processors rely heavily on static compilation, **some dynamic events cannot be fully predicted or controlled at compile time**.

- A common example is the **data cache miss**. While the latency of a cache miss is known in general terms (e.g., memory access delay), whether a particular memory access will cause a miss depends on runtime conditions. If a cache miss occurs, the processor experiences a stall, which was not anticipated during static scheduling.
- Another dynamic event is the **branch misprediction**. Although the compiler may provide static branch hints, actual program execution might differ. When a branch is mispredicted, the processor must dynamically flush the pipeline to discard speculative instructions that were fetched and partially executed under the wrong control flow assumption.

In both cases, while VLIW simplifies hardware by eliminating dynamic scheduling for regular instruction flow, it still requires mechanisms to handle these unpredictable runtime events.

Deepening: What is a Data Cache Miss?

A **Data Cache Miss** occurs when the processor tries to **read** or **write** data and the requested data is **not found in the data cache** (typically L1 data cache, page 258).

Formally, a **Data Cache Miss** happens when a load or store instruction accesses a memory address whose data is **not present in the cache**, forcing the processor to fetch it from a **lower level of the memory hierarchy** (e.g., L2 cache, main memory), incurring additional latency.

5.3 Statically Scheduled Processors

In statically scheduled architectures, the compiler is responsible for arranging instructions to exploit **Instruction-Level Parallelism (ILP)**. By carefully analyzing data dependencies and resource constraints, the compiler reorders operations to maximize parallel execution and minimize stalls.

A **Basic Block** is the main unit considered during static scheduling. It consists of a **straight-line code sequence without branches except at the entry and exit points**. However, in practice, basic blocks tend to be **small**: for typical programs, only around 4 to 7 instructions are found between branches. Moreover, even within a basic block, **true data dependencies (RAW, Read After Write)** further **limit the amount of exploitable parallelism**, forcing some instructions to be issued sequentially.

As a result, achieving substantial performance improvements requires techniques that go **beyond individual basic blocks**, aiming to exploit ILP across larger regions of code.

✓ VLIW Processors: Main Advantages

- VLIW architectures represent a **powerful solution for statically exploiting ILP**. The extensive use of compiler optimizations enables VLIW processors to achieve **high performance**, even **without the complex runtime scheduling hardware** used in superscalar designs.
- Since the compiler performs *dependency checking* and *scheduling at compile time*, it has the **advantage of analyzing the program with a much larger instruction window** than would be feasible in hardware. This allows **more opportunities to detect parallelism**, especially when aggressive code transformations, such as loop unrolling or software pipelining, are applied.
- By transferring complexity from hardware to software, VLIW processors achieve a **significant reduction in hardware complexity**. Smaller die areas lead to **cheaper production costs** and **lower power consumption**, making VLIWs attractive for **embedded applications**.
Moreover, the **fixed format of VLIW instructions simplifies the decode logic**, and scaling to a larger number of functional units **becomes easier**.

⚠ Open Challenges of VLIW Architectures

Despite these advantages, VLIW designs face several open challenges:

- First, they rely heavily on **strong compiler technology**. Sophisticated compilation techniques are necessary not only to detect parallelism within basic blocks but also to **manage parallelism across larger code regions**, which increases the complexity of compiler design.

- Another issue is the **increase in code size**. Because the instruction bundles have a fixed structure, and NOPs must be inserted to handle scheduling gaps, **increase in code size**. While code compression techniques can mitigate this effect, they **add extra decoding complexity to the processor**.
- **Register management** is also a concern. The extensive use of **register renaming** to avoid WAR and WAW hazards **increases the number of required registers** and complicates the organization of the register file. Clustered VLIW designs have been proposed to partially address this problem.
- Finally, **binary incompatibility** is a **major limitation**.

Even small changes in the number of slots, the types of functional units, or their latencies can render code compiled for one **VLIW processor incompatible with another**, even if they share the same instruction set architecture.

This **lack of portability contrasts with more dynamic architectures**, where binary compatibility is preserved across generations.

Although solutions such as **Just-In-Time compilation** have been explored, they **introduce additional complexity and are rarely adopted in practice**. For this reason, VLIW processors are primarily **employed in embedded systems**, where binaries can be compiled specifically for the target hardware and do not require high portability.

5.4 Code Scheduling

As we saw on page 191, Code Scheduling is the process of rearranging program instructions to maximize parallelism during execution. The **primary goal** of **Code Scheduling** is to **statically reorder instructions** within object code in such a way that:

- ✓ Execution time is minimized.
- ✓ Semantic correctness is preserved.

In simple terms, we want the program to run as **fast as possible without changing what it computes**. This is essential because in architectures like VLIW (Very Long Instruction Word), the hardware doesn't handle instruction reordering dynamically. Instead, the **compiler is fully responsible** for finding parallelism and deciding the scheduling.

② Motivation in VLIW Architectures

In VLIW architectures, the **responsibility of exploiting instruction-level parallelism (ILP)** is **shifted from the hardware to the compiler**. This shift implies several things:

- The **hardware is simpler**: no complex dynamic scheduling (like Tomasulo or Scoreboarding).
- **Parallel instructions are issued together** as part of a wide instruction word.
- **Each instruction** slot in a VLIW word maps to a specific **functional unit**.

② Why do we need scheduling in this context? Because a VLIW machine expects **parallel instructions to be packed statically** in the same bundle. If the compiler fails to schedule efficiently, the result is many **NOPs** and a waste of hardware resources.

⚠ Semantic Correctness and Performance

It's not enough to just reorder instructions for performance. We must also preserve:

- **True data dependencies (RAW)**, instructions must respect the order of reads after writes.
- Avoid introducing errors by violating:
 - **Anti-dependencies (WAR)**
 - **Output dependencies (WAW)**

The scheduler (compiler) must:

- ✓ Carefully **analyze the dependency graph** of the instructions.
- ✓ **Reorder only when it's safe**.
- ✓ **Possibly apply renaming** or other transformations to eliminate false dependencies.

5.4.1 Scheduling Basics

A **Basic Block** is the smallest unit of code for scheduling purposes. It's a straight-line code sequence that has:

- **One single entry point** (i.e., no jumps into the middle).
- **One single exit point** (i.e., only the last instruction may jump elsewhere).

In practice:

- No branches except at the **start (entry)** and **end (exit)**.
- All instructions inside a basic block **execute sequentially and always together**.
- Used by compilers to isolate code regions where **safe reordering** can be performed.

This **isolation** is essential for local scheduling techniques like **loop unrolling** and **software pipelining**.

Dependence Graph (DAG, Directed Acyclic Graph)

Once a basic block is identified, the compiler builds a dependence graph to model how instructions depend on each other.

- Each **node** in the graph represents an **instruction**.
- Each **edge** denotes a **data dependence** between two instructions.
- The **type of edge** can correspond to:
 - **RAW** (true dependence).
 - **WAR, WAW** (name dependencies, can often be eliminated via register renaming).

This **graph helps determine**:

- ✓ Which instructions are **ready to execute**.
- ✓ Which instructions must **wait for others to complete**.
- ✓ The **longest dependency chain**, which limits parallelism.

Why is this graph useful?

There are two main reasons:

1. It **exposes parallelism**: independent instructions can be scheduled in the same cycle or in parallel execution units.
2. It defines the **Critical Path** of execution: the **longest path from a source to a sink node** (i.e., sequential dependencies that can't be parallelized).

5.4.2 Dependence Graph and Critical Path

A **Dependence Graph** (also called a **Data Dependence Graph** or **Instruction Dependence Graph**) is a **directed acyclic graph (DAG)** used by the compiler to represent **data dependencies among instructions** in a **basic block**. It is fundamental in instruction scheduling, especially in VLIW and other statically scheduled architectures.

💡 How to Build the Dependence Graph

To build a **Dependence Graph** from a basic block:

1. **Create a node** for each individual instruction in the basic block.
2. **Add edges** between nodes based on dependencies: an edge from instruction $I_1 \rightarrow I_2$, means that I_2 depends on the results of I_1 , and I_2 cannot execute before I_1 .
3. **Label each edge** with the type of dependency (usually RAW is critical, WAR/WAW can often be resolved by renaming).
4. **Annotate each node** with the instruction **latency** (how long it takes to execute) and other scheduling metadata (like priority).

Each path in the graph shows **which instructions must wait for which**. The deeper the graph, the more serialization is required.

❖ Longest Path Computation

Each node is annotated with:

$$\text{Longest Path}(i) = \max [\text{Longest Path}(p)] + \text{Latency}(i) \quad (65)$$

Where p are all the **predecessors** of instruction i .

- This gives the **maximum cumulative latency** needed to reach instruction i .
- It reflects **how deep a node is in terms of data dependencies**.

Once all node values are computed, the **maximum among them gives the Critical Path**.

Example 1: Exam - 11 february 2025

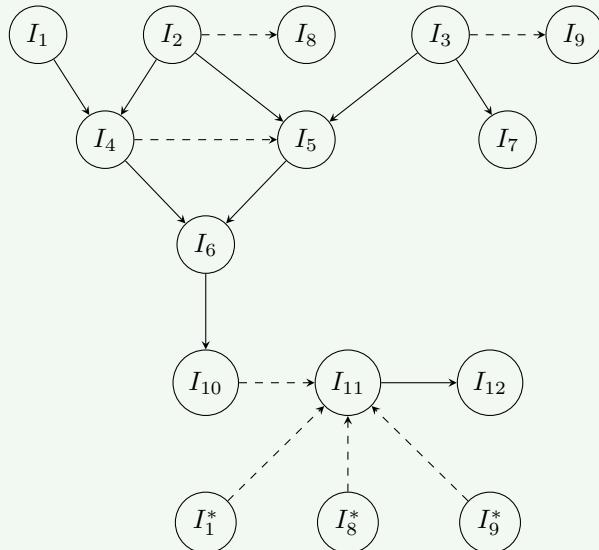
Let's consider the following LOOP code:

```

1 LOOP:   LD F1, A(R1)
2       LD F2, A(R2)
3       LD F3, A(R3)
4       FADD F1, F1, F2
5       FADD F2, F2, F3
6       FMUL F1, F1, F2
7       FADD F3, F3, F3
8       ADDUI R2, R1, 8
9       ADDUI R3, R1, 8
10      SD F1, B(R1)
11      ADDUI R1, R1, 4
12      BNE R1, R6, LOOP

```

The DAG is:



An oriented line indicates a RAW dependency. A dashed line indicates a WAR or WAW dependency. For this reason, labels are not inserted. Additionally, since the latency is calculated by hand, no text is inserted to avoid confusion in the graph. The asterisk indicates that these nodes were redrawn for readability.

❓ What is the Critical Path?

The **Critical Path** is the **longest path** through the dependence graph, starting from an entry instruction and ending at an instruction with no successors.

- It represents the **minimum number of clock cycles** needed to execute the basic block, even with unlimited resources.
- Any attempt to speed up execution (through scheduling or parallelism) cannot beat this lower bound.

If our goal is to schedule instructions as efficiently as possible, the critical path is our **primary constraint**.

5.4.3 ASAP Scheduling Algorithm (As Soon As Possible)

ASAP (As Soon As Possible) Scheduling is a greedy scheduling algorithm used in instruction scheduling. Its main idea is to **schedule each operation as early as possible** once its **dependencies are satisfied**.

- It **does not consider resource constraints**.
- It focuses only on **data readiness**.
- It is useful for computing **lower bounds** on execution time and for **priority estimation**.

❖ How ASAP Scheduling Works

1. Build the **dependence graph** (page 277).
2. Initialize a **ready list** with all **source nodes** (instructions with no predecessors). A **Ready List** is a list of all instructions **ready to be scheduled** at a given cycle. An **instruction is ready** if:
 - All its **predecessors** have been **scheduled**.
 - Its **operands are available** (i.e., computed and propagated).
3. For each cycle:
 - Schedule every instruction in the ready list.
 - Update successors (check if all their predecessors are now scheduled).
 - Move newly ready instructions into the next cycle's ready list.
4. Repeat until all instructions are scheduled.

ASAP produces **one of the shortest possible schedules**, assuming no resource constraints. The result gives the **earliest cycle** at which each instruction can execute.

❓ Why use ASAP Scheduling if it ignores resource constraints?

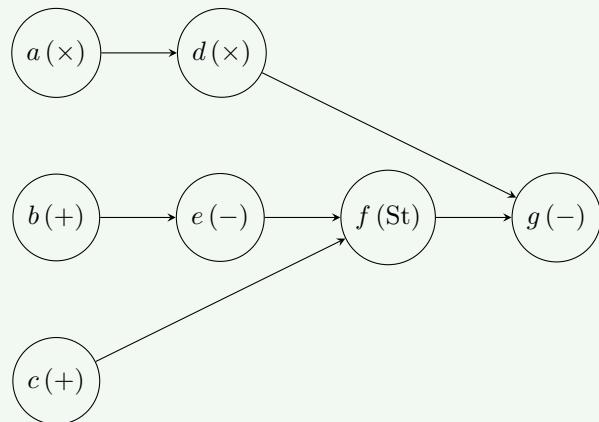
1. **It Computes a Theoretical Lower Bound.** ASAP gives us the earliest possible execution cycle for each instruction, assuming infinite resources. The critical path of the dependence graph emerges directly from ASAP: it is the shortest possible execution time, even for a perfectly parallel machine. We use **ASAP as a baseline** because **we know we can't do better than this in terms of schedule length**.
2. **It Helps Assign Priorities for Real Scheduling Algorithms.** In list-based scheduling, instructions are chosen from the ready list based on priority. One way to assign this priority is using the Longest Path to Sink (from ASAP): instructions deeper in the graph (closer to the end) are more urgent. So **even if ASAP isn't the final schedule, it informs how to choose wisely in realistic algorithms**.

3. **It's a Building Block for More Advanced Scheduling.** Many algorithms, like List Scheduling, start by computing the ASAP schedule. Then they incorporate resource constraints and adjust from there. Think of it as the “first approximation”, later refined under constraints.
4. **Compiler Simplicity in Early Stages.** In early compiler passes (before hardware-specific optimization), ASAP can: help restructure code, estimate ILP, and guide unrolling or pipelining transformations.

A great **analogy**: ASAP is like planning a trip **assuming we hit green lights at every intersection**. It's not realistic, but it **gives us a best-case travel time**, which is useful for planning and comparison.

Example 2: ASAP Scheduling Algorithm

Consider the following dependence graph:



Each node has the corresponding operation in brackets. For example, “St” is the store operation. Multiply operations take three latency cycles; ALU operations (addition and subtraction) take one; and the store operation takes two.

Ready List		Resource Reservation Table				
Cycle	Ready List	Cycle	ALU1	ALU2	L/S	MUL
1	a, b, c	1	b	c		a
2	e	2	e			a
3	f	3	f		f	a
4	d	4			f	d
5		5				d
6		6				
7	g	7	g			

5.4.4 List-Based Scheduling Algorithm

The **List-Based Scheduling Algorithm** is a realistic and practical scheduling technique that improves over ASAP by taking into account **resource constraints**.

Unlike ASAP (which assumes infinite resources), **list-based scheduling**:

- Works with a limited number of functional units (ALU, MUL, etc.).
- Ensures that at each cycle:
 - ✓ We do not oversubscribe any hardware unit.
 - ✓ We only schedule instructions whose operands are ready.

This reflects **real hardware limitations** in VLIW superscalar, or statically scheduled pipelined processors.

Algorithm Steps

At each cycle, the algorithm maintains a **Ready Set**, similar to ASAP. But now, when multiple instructions are ready, and **not all can be scheduled** in parallel, it must **choose** based on a **priority**.

1. Build the dependence graph.
2. Compute **priorities** for all nodes (typically with longest-path-to-sink). Nodes closer to the end of the graph are **more critical** (less slack). The compiler calculates for each instruction:

$$\text{priority}(i) = \text{Max path length from } i \text{ to a sink node}$$

Higher priority, more urgent to schedule.

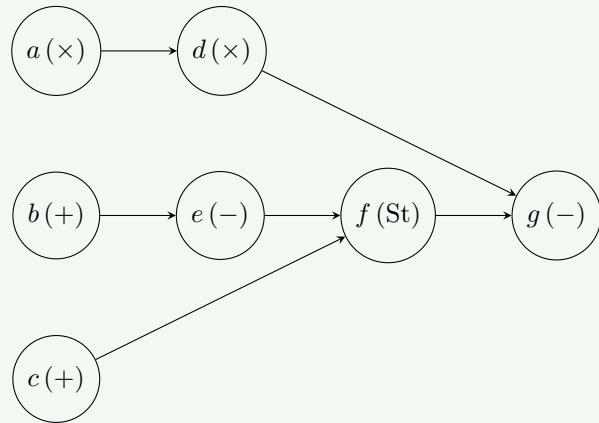
3. Initialize the **Ready Set** with nodes that have no predecessors.
4. For each cycle:

- While resources are available and Ready Set is not empty:
 - Pick the **highest-priority instruction** that fits into available resources.
 - Assign it to the current cycle.
 - Mark it as scheduled
- Update the Ready Set: add successors whose dependencies are now satisfied.

5. Repeat until all instructions are scheduled.

Example 3: List-Based Scheduling Algorithm

Consider the following dependence graph:



Each node has the corresponding operation in brackets. For example, “St” is the store operation. Multiply operations take three latency cycles; ALU operations (addition and subtraction) take one; and the store operation takes two.

Ready List		Resource Reservation Table			
Cycle	Ready List	Cycle	ALU1	L/S	MUL
1	a, b, c	1	b		a
2	c, e	2	c		a
3	e	3	e	f	a
4	d, f	4		f	d
5		5			d
6		6			d
7	g	7	g		

VLIW Code			
Cycle	ALU1	L/S	MUL
1	b		a
2	c		
3	e		
4		f	d
5			
6			
7	g		

5.4.5 Local vs Global Scheduling

In modern compilers, especially for statically scheduled architectures like VLIW, **instruction scheduling can be applied at two levels**:

- **Local Scheduling: Within Basic Blocks.** Local scheduling operates inside a single basic block, where:

- The control flow is **linear** (no internal branches).
- The compiler has **full knowledge** of all instructions and dependencies.

◎ The **objective** is to:

- Maximize instruction-level parallelism (ILP).
- Minimize execution time of each block.
- Efficiently fill VLIW slots with real instructions (reduce NOPs).

❖ Techniques

- ASAP scheduling (page 279).
- List-based scheduling (page 281).
- Loop unrolling (page 285).
- Software pipelining (page 289).

Local scheduling is **simpler** and **effective**, but limited:

- ✖ It cannot move **instructions** across block boundaries.
- ✖ ILP is bounded by the size and dependencies of the block.

- **Global Scheduling: Across Basic Blocks.** Global Scheduling considers a **wider scope**, across **multiple basic blocks**, allowing **reordering of instructions**:

- Across **branches**.
- Across **loop boundaries**.
- Even across **entire control paths**.

◎ The **goal** is to:

- Exploit **more ILP than local scheduling** allows.
- Hide or overlap long-latency operations.
- Restructure code for better pipeline utilization.

🚫 Challenges

- Must preserve **control and data dependences**.
- Must generate **compensation code** to fix incorrect speculative moves.
- More **complex** and **expensive** to implement in the compiler.

Technique	Type	Description
Loop Unrolling	Local	Replicates loop body to increase basic block size and expose ILP.
Software Pipelining	Local	Reorganizes instructions from different iterations to overlap them.
Trace Scheduling	Global	Predicts a frequent path (trace) and aggressively schedules it.
Superblock Scheduling	Global	Extension of trace scheduling; allows multiple exits, single entry.

Table 23: Techniques Overview.

5.4.6 Local Scheduling Techniques

5.4.6.1 Loop Unrolling

Loop Unrolling is a **compiler optimization** that duplicates the **body of a loop multiple times**, reducing the number of loop control instructions and exposing **more parallelism** within the loop body.

❷ What is the goal?

- To **increase the size of the basic block** inside a loop.
- To **expose more independent instructions** to the scheduler.
- To **reduce the number of branches and loop counters**, lowering control.

✓ Benefits

1. **More instruction-level parallelism (ILP)**, more instructions can be scheduled in parallel.
2. **Reduced loop overhead**, fewer BNE, SUBI, etc., executed per iteration.
3. **Better resource utilization**, more functional units are kept busy (fewer NOPs).

✗ Downsides

1. **Increased register pressure**, more temporary values, so more registers needed.
2. **Increased code size**, can lead to **instruction cache misses**.

Example 4: No Unrolling vs. Loop Unrolling

Consider the following code:

```

1 for (i=1000; i>0; i--)
2   x[i] = x[i] + s;
3

```

- No Unrolling, one iteration in Assembly:

```

1 Loop: LD    F0, 0(R1)
2      ADD   F4, F0, F2
3      SD    F4, 0(R1)
4      SUBI R1, R1, 8
5      BNE   R1, R2, Loop

```

- Execution time: **10 cycles per iteration**
- Efficiency: 5 instructions, 10 cycles, then **50%** (IPC = 0.5)

- 4× Loop Unrolling (not yet optimized):

```

1 Loop: LD F0, 0(R1)
2     ADD F4, F0, F2
3     SD F4, 0(R1)
4
5     LD F0, -8(R1)
6     ADD F4, F0, F2
7     SD F4, -8(R1)
8
9     LD F0, -16(R1)
10    ADD F4, F0, F2
11    SD F4, -16(R1)
12
13    LD F0, -24(R1)
14    ADD F4, F0, F2
15    SD F4, -24(R1)
16
17    SUBI R1, R1, 32
18    BNE R1, R2, Loop

```

Only **true dependences** remain between SUBI and BNE. But there's a problem: we are **reusing** F0, F4, and leads to **WAW** and **WAR** hazards.

✓ Register Renaming to Resolve Name Dependencies

As we saw in the previous example, there are WAW and WAR hazards due to the use of the same register during loop unrolling. To remove false dependencies (WAW and WAR), the compiler performs **register renaming**:

```

1 Loop: LD F0, 0(R1)
2     ADD F4, F0, F2
3     SD F4, 0(R1)
4
5     LD F6, -8(R1)
6     ADD F8, F6, F2
7     SD F8, -8(R1)
8
9     LD F10, -16(R1)
10    ADD F12, F10, F2
11    SD F12, -16(R1)
12
13    LD F14, -24(R1)
14    ADD F16, F14, F2
15    SD F16, -24(R1)
16
17    SUBI R1, R1, 32
18    BNE R1, R2, Loop

```

Now all instructions use **unique registers**, and only **true dependencies remain**. It allows **maximum scheduling freedom**.

- **Execution time:** 16 cycles per 4 iterations, then 4 cycles per iteration.
- **Loop overhead:** 4 cycles per 4 iterations, then 1 cycle per iteration.
- **Efficiency:** 14 instructions, 16 cycles, then $14 \div 16 = 87.5\%$

Effect on ILP After renaming, instructions are independent and can be reordered or scheduled in parallel:

- Original loop: 1 useful instruction per cycle.
- After unrolling and renaming: up to 4 useful instructions per cycle (in theory)

This enables better pipelining, more efficient use of wide VLIW issue slots, and less stalling.

Concept	Description
Loop unrolling	Duplicates loop body to expose more ILP.
Benefit	Reduces loop overhead, increases block size.
Drawback	Increases register pressure and code size.
Register renaming	Eliminates name dependencies (WAW/WAR) to enable parallel scheduling.
Result	More efficient and parallel scheduling possible.

Performance Metrics of Unrolling

When applying **loop unrolling**, the goal is not just to make the code longer, it's to execute faster. We can measure this improvement through several performance metrics:

- **Execution Time.** The total execution time per loop iteration reflects how efficiently the instructions are executed.
 - Execution time includes useful operations plus stalls/NOPs.
 - Lower execution time = better use of ILP and scheduling.

We typically express it as:

$$\text{Execution time} = \frac{\text{Cycles to execute unrolled loop}}{\text{Number of iterations covered}} \quad (66)$$

- **Loop Overhead.** Loop overhead includes instructions that are not part of the loop's core computation:

- SUBI (loop counter update).
- BNE (branch).
- Any NOP needed to avoid hazards.

Reducing overhead:

- Makes more room for useful instructions.
- Boosts efficiency.

In unrolling we do 4× the work but only 1 loop update, loop overhead is amortized over more iterations.

- **Code Efficiency.** Code efficiency **quantifies how well instruction slots are used**. Defined as:

$$\text{Efficiency} = \frac{\text{Number of useful operations}}{\text{Total number of instruction slots}} \quad (67)$$

This tells us **how many issue slots are actually contributing to the computation**.

5.4.6.2 Software Pipelining

Software Pipelining is a scheduling technique where instructions from **different iterations** of a loop are **reordered and overlapped** to keep functional units busy every cycle.

Instead of executing one full iteration at a time, the processor **starts a new iteration every cycle**, with different stages of previous iterations still in progress, similar to hardware pipelining.

Comparison with Loop Unrolling

Aspect	Loop Unrolling	Software Pipelining
Structure	Duplicates code	Keeps loop body same size
ILP exposure	Gained by replicating loop body	Gained by reordering across iterations
Code size	Increases (proportional to unrolling factor)	Remains compact
Scheduling scope	Within a block	Across multiple iterations
Register pressure	High (due to duplication)	Very high (many live temporaries)
Throughput	Good (bounded by unrolling factor)	Potentially optimal (1 iteration per cycle)

Table 24: Comparison with Loop Unrolling.

Phases of Software Pipelined Loop

A software-pipelined loop executes in **three distinct phases**:

1. **Startup (Prologue)**. Initial cycles where instructions from the **first few iterations** begin execution. Not all stages are filled yet.
2. **Steady State**. The **core phase**: every cycle completes one iteration, while other iterations are in progress. This is where **ILP is fully exploited**.
3. **Drain (Epilogue)**. Final cycles where instructions from the **last few iterations** finish. No new iterations start; pipeline “drains” out.

Concept	Description
Software pipelining	Reorders loop instructions from different iterations.
Objective	Maximize ILP by filling the execution pipeline every cycle.
Benefits	Compact code, steady throughput, ideal for VLIW.
Trade-offs	High register pressure, scheduling complexity.
Execution phases	Startup (filling), Steady (full overlap), Drain (finishing up).

Example 5: Software Pipelining

We start with the original scalar loop:

```

1 Loop: LD    F0, 0(R1)
2      ADD   F4, F0, F2
3      SD    F4, 0(R1)
4      SUBI  R1, R1, 8
5      BNE   R1, R2, Loop

```

1. Loads a value from memory.
2. Adds a scalar value F2.
3. Stores it back.
4. Updates the pointer.
5. Branches if not done.

Execution time per iteration: 5-6 cycles. So only **one useful result per loop iteration**, poor ILP.

To apply software pipelining, we **reorder instructions from different iterations to fill functional units**. Assume all dependencies are respected (RAW), and rename registers to avoid WAR/WAW hazards.

```

1 LD    F0, 0(R1)      # I0
2 ADD   F4, F0, F2      # I1 from previous iteration
3 SD    F4, -8(R1)      # I2 from 2 iterations ago

```

This version performs LD for current iteration, ADD for previous, SD for 2 iterations ago. This matches the **steady state** of a pipelined execution: **each cycle does useful work from 3 iterations simultaneously**. We now issue **one full result per cycle** after startup.

Now let's handle address computation correctly. Assume:

- R1 points to current element.
- We track offset correctly for each instruction.

```

1 # Prologue (Startup)
2 LD    F0, 0(R1)
3
4 # Loop body (Steady state)
5 Loop: LD    F6, -8(R1)      # LD for next iteration
6     ADD   F4, F0, F2       # ADD from previous iteration
7     SD    F4, 0(R1)        # SD from 2 iterations ago
8     SUBI  R1, R1, 8
9     BNE   R1, R2, Loop
10    MOV   F0, F6          # prepare F0 for next ADD

```

Operation	From Iteration
LD F6	i + 1
ADD F4	i
SD F4	i - 1

- F0 is shifted forward via MOV.
- Pipelined stages are handled via register flow.
- Control remains simple, and code is compact.

Version	Throughput	Notes
Classic Loop	1 result every 5 cycles	Poor resource utilization.
Pipelined (steady)	1 result per cycle	Maximized ILP.
Code size	Small (no unrolling)	More compact than unrolling.

5.4.7 Global Scheduling

5.4.7.1 Trace Scheduling

Trace Scheduling is a **global instruction scheduling technique** that aims to optimize the **most frequently executed paths** (called **traces**) through a program's control flow graph (CFG).

It allows the compiler to move instructions **across basic block boundaries**, including across branches, to expose more **instruction-level parallelism (ILP)** and fill wide VLIW instruction words more efficiently.

- **Step 1: Trace Selection.** A **trace** is a likely path through the code, a **linear sequence of basic blocks** that are expected to execute one after the other **most of the time**.

💡 How traces are selected

- Based on **profiling information**, static heuristics, or feedback from runs.
- Hot paths (e.g., common branches of `if-else` or main loop bodies) are prioritized.
- A trace may include: **forward branches**, **backward branches**, multiple blocks treated as a **single scheduling unit**.

We can think of a trace as a “**super-block**” of code that is optimized as a single unit, assuming that it is frequently executed in that order.

- **Step 2: Trace Compaction (Scheduling).** Once the trace is selected, the compiler:

- Builds a **dependence graph** for the entire trace.
- Applies **list-based scheduling** across all blocks in the trace.
- Moves instructions **upward** across branches if possible, to fill unused slots (VLIW-friendly).

This reordering allows:

- ✓ Filling of issue slots across multiple basic blocks.
- ✓ Earlier execution of long-latency operations (e.g., memory loads).

⚠ Speculation and Compensation Code

✖ Problem. When instructions are **moved across branches**, they might be executed **in paths where they shouldn't**. This is called **Speculative Execution**:

- The compiler **speculates** that a certain path is taken.
- It moves instructions **before a branch**, even if they are only **valid in one branch**.

✓ **Solution: Compensation Code.** To preserve correctness, the compiler adds **compensation code** in **off-trace blocks** to:

- Recompute or restore the correct values.
- Prevent side effects (e.g., stores) from speculates

Definition 1: Compensation Code

Compensation Code is extra code inserted by the compiler to preserve program correctness after performing speculative instruction movement during global scheduling (e.g., in trace scheduling).

Example 6: Compensation Code

Suppose this trace is selected:

```
1 if (x > 0)
2     A = B + C; // part of the trace
3 else
4     A = D + E
```

Compiler speculatively moves $A = B + C$ above the `if`. Then in the `else` block, it inserts:

```
1 A = D + E; // compensation code
```

This ensures correctness if the speculation fails.

Step	Action
Trace Selection	Choose most likely sequence of basic blocks (the “hot path”).
Trace Compaction	Reorder instructions across blocks to improve ILP and scheduling.
Speculation	Move instructions before branches, assuming likely path.
Compensation Code	Inserted to preserve semantics if speculation was wrong.

Table 25: Trace Scheduling summary.

5.4.7.2 Superblock Scheduling

Superblock Scheduling is an evolution of **Trace Scheduling**, designed to simplify speculation management while still providing **high ILP** for **VLIW** or statically scheduled superscalar processors.

It is a **global instruction scheduling technique** that extends Trace Scheduling, with two key design changes:

1. It restricts code regions to have **only one entry point**.
2. It allows **multiple exits**.

This enables the compiler to **aggressively schedule instructions across basic blocks**, including across branches, but with **simpler and safer control-flow management** than full trace scheduling.

⚠️ Superblocks vs Traces

Feature	Trace Scheduling	Superblock Scheduling
Entry points	Multiple (any predecessor block).	Single entry point (one predecessor).
Exit points	Multiple.	Multiple.
Control flow	General, complex.	Simplified, controlled .
Speculation	Really aggressive, harder to manage.	More manageable.
Compensation code	May be needed on both sides of branches.	Needed only on one side (outside the superblock).

❓ Why is this useful?

In trace scheduling, we may have **multiple entry points** into the trace, which means any speculative movement of instructions must be **repaired in many paths** (with compensation code). In **Superblock Scheduling**, by **enforcing a single-entry rule**, and using **tail duplication** to eliminate side entries, the compiler can:

- ✓ Speculate more **safely**.
- ✓ Insert **fewer compensation instructions**.
- ✓ Retain most of the **performance benefit** of traces.
- ✓ Keep control flow **more manageable**.

❖ Typical Structure of a Superblock

A **superblock** is a **linear sequence of basic blocks** that:

- Starts from a **unique entry point**.
- Ends at **multiple potential exits**.
- Is built around a **frequently executed path** (like a hot loop body or branch).

❓ How is it built?

1. **Profiling** (or prediction) selects a hot path.
2. **Tail Duplication** is applied to remove alternative entries. Compiler **duplicates the tails** (exit blocks) of a region, because this ensures all control flow into the superblock happens from **a single point**.
3. The resulting region becomes a **superblock** (safe to schedule aggressively).

We can think of it as a **safer, more controlled trace**, scheduled to maximize instruction-level parallelism (ILP) on VLIW or superscalar machines.

Example 7: Superblock

Original code:

```

1 if (x > 0)
2     A = B + C;
3 else
4     A = D + E;
```

We want to **schedule more instructions in parallel**. But the compiler sees a branch:

- In one path: $A = B + C$
- In the other: $A = D + E$

To do **Superblock Scheduling**, we assume that **the first path ($x > 0$) is more frequent** (we get this from profiling).

1. We take the **most frequent path ($x > 0$)**.
2. We include its block: $A = B + C$.
3. We move it **before the branch**, so it's scheduled early (speculative move).
4. We make the region a **superblock**: a linear code region with one entry.

```

1 A = B + C           // speculated! Scheduled early
2 if (x <= 0)
3     A = D + E      // compensation code
```

Now:

- If the branch is *not taken* ($x > 0$), then $A = B + C$ is correct.
- If the branch *is taken* ($x \leq 0$), then we **fix A** using compensation code.

It is called a “superblock” because we treat that piece of code as one large block when scheduling. The compiler: optimizes this region together, tries to issue as many instructions per cycle as possible, and assumes most of the time the “hot path” will be followed.

6 Advanced Memory

6.1 Introduction

Modern processors operate at very high speeds, and any delay in accessing data can drastically reduce performance. Yet, fast memory is expensive and limited in size, while large memory is slow. This chapter provides strategies to bridge the gap between processor speed and memory latency, introducing techniques to optimize access time, reduce misses, and manage complexity.

Understanding **advanced memory hierarchy** allows us to:

- Reduce the **Average Memory Access Time (AMAT)**.
- Optimize system performance while managing cost and power.
- Design efficient **cache architectures**, essential in high-performance systems (from embedded CPUs to data centers)

⌚ Motivation for Hierarchical Memory

Programmers naturally wish for: “*an unlimited amount of memory, as fast as the CPU*”. But this ideal is physically and economically unfeasible:

- **Fast memories** (SDRAM) are **very expensive** and **power-hungry**.
- **Slow memories** (DRAM) are **cheaper** and **larger**, but much **slower**.

🕒 **Solution: A memory hierarchy.** Build a layered system of memories, where each layer:

- Gets **faster** and **smaller** as we go **upward**.
- Gets **slower** and **larger** as we go **downward**.

Each upper level stores a **subset** of data from the lower level. This way, **frequently accessed data** stays in **fast memory (cache)** and **less frequent data** is stored deeper (L2, L3, DRAM, Disk).

⚡ Illusion of a Large, Fast Memory

Even though the real system consists of many memory levels, **the processor** “sees” a **single unified memory**. Thanks to the hierarchy and locality principles:

- Most accesses hit the **fast upper-level caches**.
- Rarely used data is retrieved from lower levels (with longer latency).

This *illusion* is what enables high-speed execution even with cost-effective memory.

❓ There is a sort of hierarchy with different levels of access to memory benefits. Do developers manually choose fast or slow memory?

No, typically developer do *not* manually choose which memory level (L1, L2, DRAM, etc.) stores the data. Instead, the **hardware and compiler manage this automatically**, using a memory hierarchy. So what's going on?

- The developer writes code as if there's **one big memory**.
- The system **automatically** stores:
 - Frequently used data in **fast memory (L1 cache)**.
 - Less-used data in **slower memory (L2/L3 cache or RAM)**.
- The **processor checks the fast levels first** and, only if needed, moves to slower levels.

This is what we call the **illusion of a large and fast memory**.

❓ How is locality linked to the hierarchy?

Locality is a behavioral pattern of how programs access memory. The memory hierarchy is *designed to exploit locality*. There are two types of locality in programs:

- **Temporal Locality**: “*if we used this data recently, we will likely use it again soon*”. For example:

```
1 for (int i = 0; i < 100; i++) {
2     sum += a[i];    // temporal locality for variable 'sum'
3 }
```

Where `sum` is read and written **many times in a short span**, so keep it in the **L1 cache** to avoid loading from RAM every time.

- **Spatial Locality**: “*if we access this memory address, we will probably access nearby ones soon*”. For example:

```
1 for (int i = 0; i < 1000; i++) {
2     // spatial locality: accessing sequential array elements
3     x = arr[i];
4 }
```

If we access `arr[0]`, we will soon access `arr[1]`, `arr[2]`, etc. The system loads **entire blocks** (not just one word) into cache, anticipating the next accesses.

The **hierarchy is effective only if programs exhibit locality**. Fast memory (caches) are **small**, so they **rely on locality** to keep relevant data. **Without locality**, caches would be **useless**: every access would go to slow memory.

Example 1: Memory Hierarchy Analogy

Think of a **bookshelf system**:

- We keep our most-used books on our desk (L1).
- Others in our room shelf (L2).
- The rest in the university library (main memory).
- The rarest in the city archive (disk).

If we're smart about what to keep close (using locality), we'll avoid most slow trips.

6.2 Principle of Locality

Cache memories are effective because programs tend to access data and instructions in predictable patterns. These patterns are captured by the **Principle of Locality**:

- **Temporal Locality** (“Time-based reuse”). If a memory location is accessed, it’s likely to be accessed again *soon*. So recently accessed data stays in cache, because it’s likely to be accessed again.

For example, repeated access to loop variables:

```
1 for (int i = 0; i < n; ++i) {
2     sum += a[i];
3 }
```

Or reuse of stack frames in recursive functions or instruction fetches inside tight loops.

✖ **Cache strategy.** Keep recently accessed **blocks** in cache. Don’t replace them unless absolutely necessary (see LRU in later sections).

- **Spatial Locality** (“Nearby reuse”). If a memory location is accessed, it’s likely that *nearby locations* will be accessed soon. In other words, if we need $a[i]$, we’ll probably also need $a[i+1]$, $a[i+2]$, etc.

For example, sequential instruction execution, or accessing elements of an array, or scanning a matrix row by row.

✖ **Cache strategy.** Fetch **entire blocks**, not single words (e.g., 64-byte lines). Use **block size > 1 word** to bring adjacent memory into cache.

⌚ How Caches Exploit Locality

With **Cache Block** we refer to the smallest unit of data moved from main memory to cache. It is typically 32, 64 or 128 bytes. It helps bring **spatial neighbors** into the cache.

Locality Type	Cache Mechanism
Temporal	Keep recently accessed data in cache
Spatial	Fetch blocks (containing multiple nearby elements)

Most modern cache architectures assume that the programmer writes code that follows locality principles, which is why **compiler optimizations and programming style** matter.

6.3 So, what is a Cache?

Definition 1: Cache

A **Cache** is a small, fast **memory** located **close to the processor** that temporarily **stores copies of data from the main memory**, aiming to reduce the average time needed to access data and instructions.

Modern memory systems are structured hierarchically:

- **Cache (upper level)**: small, fast, but expensive (SRAM). It is very close to the CPU and stores a subset of main memory to reduce access latency.
- **Main memory (lower level)**: larger, slower, but cheaper (DRAM). It stores the complete working set and is accessed only if data is not found in cache.

The cache acts as a **buffer** between fast CPU and slow main memory.

Terminology

- **Cache Block** or **Cache Line**. The smallest unit of data transferred between memory and cache. It is typically 32 to 128 bytes. Since a cache is divided into multiple blocks, we use:

$$\text{Number of blocks} = \frac{\text{Cache size}}{\text{Block size}} \quad (68)$$

To calculate the number of blocks in the cache.

Example 2

With a 64 KB cache, and 16-byte per block:

$$\frac{64 \times 1024}{16} = 4096 \text{ blocks}$$

✓ **Cache Hit**. Occurs when the requested memory address is **already present in the cache**. Fast access, minimal latency.

✗ **Cache Miss**. Occurs when the requested address is **not in the cache**. The block must be fetched from lower-level memory.

⚠ Consequences of a miss

1. **Stall** CPU.
2. Fetch block from memory.
3. **Write** block to cache.
4. Retry access (which now becomes a hit).

- **Hit Time.** It is the time to:

1. Access the cache.
2. Determine if it's a hit.
3. Return data (if present).

It is fast, typically 1-3 CPU cycles.

- **Miss Penalty.** It is the time to:

1. Fetch the block from main memory (or next-level cache).
2. Update cache.
3. Resume the CPU.

Much **slower** than hit time, often 10s to 100s of cycles.

- **Average Memory Access Time (AMAT)**

$$\text{AMAT} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty} \quad (69)$$

This is the **expected cost of accessing memory**, accounting for both hits and misses.

Example 3

For example:

- Hit Time: 1 cycle.
- Miss Rate: 5%.
- Miss Penalty: 100 cycles.

$$\text{AMAT} = 1 + 0.05 \times 100 = 6 \text{ cycles}$$

6.4 Cache Performance Metrics

This section is about quantifying how effective a cache is, and how to improve it. Similar topics were covered on page 257.

The **Average Memory Access Time (AMAT)** is the average time the CPU needs to access memory, considering both hits and misses. So, when accessing memory, we either get a **hit** or a **miss**:

$$\text{AMAT} = (\text{Probability of Hit}) \cdot (\text{Time if Hit}) + \\ (\text{Probability of Miss}) \cdot (\text{Time if Miss})$$

Which is written as:

$$\text{AMAT} = \text{Hit Rate} \cdot \text{Hit Time} + \text{Miss Rate} \cdot \text{Miss Time} \quad (70)$$

This is the **most general formula**, where:

- **Hit Rate**: fraction of memory accesses that result in a hit.
- **Hit Time**: time to check cache and return data on a hit.
- **Miss Rate**: fraction of memory accesses that result in a miss.
- **Miss Time**: time to
 1. **Check the cache** (same time as a *Hit Time*).
 2. **Fetch the data from main memory** (same time as a *Miss Penalty*)

So it is reasonable to write the Miss Time as follows:

$$\text{Miss Time} = \text{Hit Time} + \text{Miss Penalty} \quad (71)$$

However, if the Miss Time is equal to the sum of the Hit Time and the Miss Penalty, then the most general formula can be **simplified**:

$$\begin{aligned} \text{AMAT} &= \text{Hit Rate} \cdot \text{Hit Time} + \text{Miss Rate} \cdot \text{Miss Time} \\ &= \text{Hit Rate} \cdot \text{Hit Time} + \text{Miss Rate} \cdot \\ &\quad (\text{Hit Time} + \text{Miss Penalty}) \\ &= \text{Hit Rate} \cdot \text{Hit Time} + \text{Miss Rate} \cdot \text{Hit Time} + \\ &\quad \text{Miss Rate} \cdot \text{Miss Penalty} \\ &= \text{Hit Time} \cdot \underbrace{(\text{Hit Rate} + \text{Miss Rate})}_{\% \text{ Miss} + \% \text{ Hit} = 100\% = 1} + \text{Miss Rate} \cdot \text{Miss Penalty} \end{aligned}$$

So the **simplified and most common** version is:

$$\text{AMAT} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty} \quad (72)$$

Even if a cache **misses**, we still pay the **Hit Time**, because we have to *check* the cache to know that it's a miss. So every access:

 Pays **Hit Time**.

 Plus, if it's a miss, it pays **Miss Penalty**.

6.5 Cache Architecture

Each **Cache Block** (also called **Cache Line**) is a row in the cache that stores part of memory.

A typical cache line has **three fields**:

- **Valid Bit**: Indicates if the **block contains valid** (i.e., initialized) **data**.
- **Tag**: The high-order bits (most significant bits, leftmost side) of the memory address, used to **identify the block**.
- **Data**: The actual words/bytes of memory copied into the cache line.

At **startup**, all **valid bits = 0** (cache is empty).

Example 4: Visual Representation

V	TAG	DATA
1	0x001AE	[Word0, Word1, Word2...]

- V: valid bit.
- TAG: identifies *which memory block* this is.
- DATA: contains the content of that memory block (usually several words).

⌚ Why the TAG?

Let's say the CPU requests a memory address like 0x1AE023. The cache checks:

1. **Index**: determines (extract) which row to look in.
2. **Tag**: is this row the one holding 0x1AE?
3. **Valid**: is the content even initialized?

If both Tag and Valid match, it's a **hit**.

⌚ Four key cache design questions

- ⌚ **Block placement** - *Where can a block be placed?* Page 305.
- ⌚ **Block identification** - *How is a block found?* Page 308.
- ⌚ **Replacement strategy** - *Which block should be replaced?* Page 309.
- ⌚ **Write strategy** - *What happens on a write?* Page 311.

6.5.1 Block Placement: *Where can a block be placed?*

This question is about how memory addresses map to positions in the cache. It defines the **mapping policy** between **main memory blocks** and **cache lines**.

There are **three classic cache organizations**:

1. **Direct-Mapped Cache.** It is the simplest and most intuitive. Each block from memory **can go in only one location** in the cache.

$$\text{Cache Index} = (\text{Block Address}) \mod (\text{Number of Cache Blocks}) \quad (73)$$

This means:

- ✓ Easy and **fast lookup** (just compute index).
- ✗ But high risk of **conflict misses** (collision): two blocks that map to the same index will evict each other repeatedly.

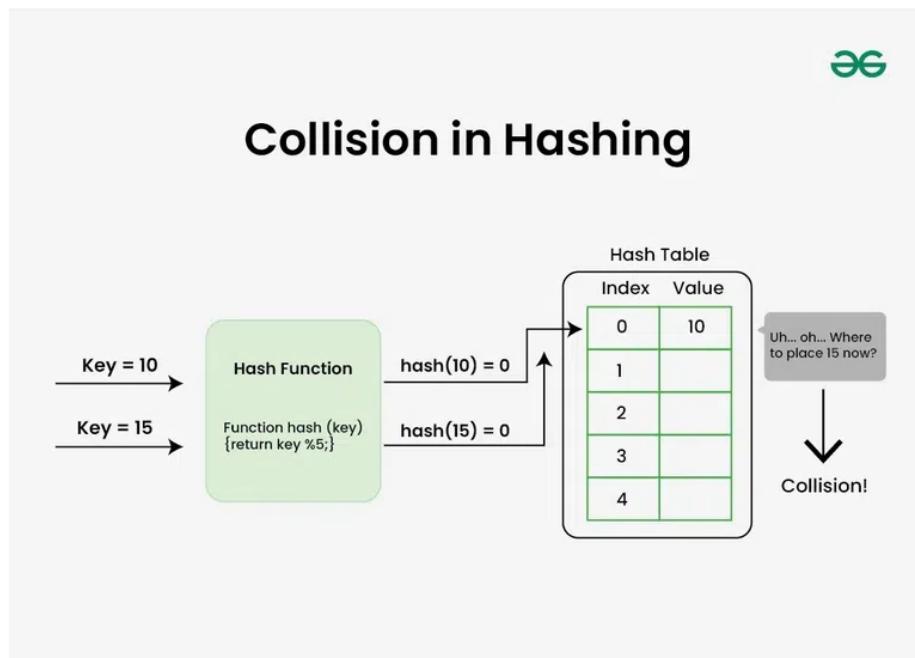


Figure 28: A great analogy is in a hashing environment, where a collision occurs if a key is the same. In our case, it's the same: same index, collision! [1]

Example 5: Direct-Mapped Cache

Cache has 8 blocks: indexes from 0 to 7. The memory block of 12 is calculated as follows:

$$12 \mod 8 = 4$$

And it can only go in **cache block 4**.

Memory Block	Cache Block
4	4
12	4
20	4

So if we alternate access to blocks 12 and 20, we get repeated evictions, then poor performance.

2. **Fully Associative Cache.** A block from memory **can be placed in any cache line**. There is **no index in the address** (no needed), and the mapping rule is a **full TAG comparison across all cache lines**. This means:

- ✓ **No conflict misses** (since any block can go anywhere).
- ✓ Ideal flexibility.
- ✗ **Expensive hardware:** must compare TAGs for *all* cache lines. Because to do this, the hardware includes:
 - **N comparators**, each comparing the input tag with the tag stored in one of the N cache lines.
 - A **priority encoder** or logic to select the matching line (if any).
- ✗ Slower access.

3. **N-Way Set-Associative Cache.** A compromise between direct-mapped and fully associative.

- Cache is divided into **sets**.
- Each set contains n **lines** (ways).
- A memory block maps to **exactly one set**, but can be placed in **any of the n lines** within that set.

The mapping rule is:

$$\text{Set Index} = (\text{Block Address}) \mod (\text{Number of Sets}) \quad (74)$$

This reduces conflict misses **while keeping lookup cost low** (only compare tags inside the set).

Example 6: N-Way Set-Associative Cache

Imagine a 2-way set-associative cache with 4 sets (each set has 2 blocks). For example, the **block 12** goes into **Set 0**:

$$12 \mod 4 = 0$$

In Set 0 it can go into **either Way 0 or Way 1**.

Memory Block	Set Index
0	0
4	0
8	0
12	0

Each of these blocks competes for 2 slots in Set 0, so **fewer conflicts** than direct-mapped.

Summary of Block Placement Policies:

- **Direct-Mapped**

Mapping Rule	Flexibility	Hardware Cost	Risk of Conflict Miss
Block Address mod #Blocks	● Rigid (1 line)	✓ Simple	● High

- **Fully Associative**

Mapping Rule	Flexibility	Hardware Cost	Risk of Conflict Miss
Any block to any line (compare all)	● Full Freedom	● Expensive	● None

- **N-Way Set-Associative**

Mapping Rule	Flexibility	Hardware Cost	Risk of Conflict Miss
Block Address mod #Sets	● Medium	● Moderate	● Controlled

6.5.2 Block Identification: *How is a block found?*

When the CPU requests a memory address, the cache must **quickly determine** whether that **address is stored inside it**, and if so, in **which line**. This process depends on the **cache mapping type** that we studied in the previous section.

■ General Method

Regardless of mapping type, the process involves:

1. **Index**. Determines *where to look* (which set or which specific line in direct-mapped caches).
2. **Tag**. Stored in the cache line; must match the tag bits extracted from the CPU's requested address.
3. **Valid Bit**. Ensures the entry contains real, initialized data.

If **index matches**, **tag matches**, and **valid bit = 1**, we have a **hit**. Otherwise a **miss**.

☒ How it works in each mapping type

- **Direct-Mapped Cache**

1. Use **index bits** from the address to locate exactly **one cache line**.
2. Compare **tag bits** in that line to the requested address's tag.
3. Check **valid bit**.
4. **Hit** if both match; otherwise, **miss**.

- **Set-Associative Cache**

1. Use **index bits** to select the **set**. Inside that set, there are n **lines** (ways).
2. Compare n **lines** of each way in that set in parallel.
3. If one matches and **valid = 1** \rightarrow **hit**; otherwise, **miss**.

- **Fully Associative Cache**

1. **No index bits**; the whole cache is **one big set**.
2. Compare the requested tag to **every tag** in the cache in parallel.
3. **Hit** if any match with **valid = 1**; otherwise **miss**.

Mapping Type	How Block is Found
Direct-Mapped	Use index to find 1 line, compare its tag.
Set-Associative	Use index to find 1 set, compare tags of n ways in that set.
Fully Associative	Compare tag with every line in the cache.

6.5.3 Replacement Strategy: *Which block should be replaced?*

When a **cache miss occurs**, the cache must bring a new block from the lower memory level. If the cache (or the relevant set) is **full**, one of the existing blocks must be **evicted** to make space. But, *which block should be replaced?*

⌚ Depends on Mapping Type

This decision depends on the type of mapping we have in the cache. With a direct-mapped cache, the answer is clear: always replace. However, with a set-associative or fully associative cache, it depends on the replacement policy.

1. **Direct-Mapped Cache.** No choice: the mapping rule already determines **exactly one line** for the block. The existing block in that line is **always replaced**. So more conflict misses.
2. **Set-Associative or Fully Associative Cache.** There are multiple possible slots (n lines in the set, or all lines for fully associative). A **replacement policy** decides which one to evict. A good replacement policy (like LRU, see below) can significantly reduce misses, especially in workloads with strong temporal locality.

⚡ Common Replacement Policies

A **Set-Associative** or **Fully Associative Cache** can adopt one of the following common replacement policies:

1. **Random Replacement.** Pick any block in the set at random.
 - ✓ Simple hardware and fast.
 - ✓ Avoids always evicting the same block in some pathological patterns.
 - ✗ It doesn't exploit locality because it is random and doesn't follow a heuristic. It **may evict a frequently used block**.
2. **FIFO (First-In First-Out).** Evict the block that has been in the cache the longest (oldest arrival).
 - ✓ Easy to implement with a queue per set.
 - ✗ Ignores recent usage, so might **evict a frequently used block** if it's old.
3. **LRU (Least Recently Used).** Evict the block that has **not been used for the longest time**.
 - ✓ Matches the idea of **temporal locality**: recently used data is likely to be used again.
 - ✗ More **complex hardware**: must track usage order for each block in the set.

Example 7: 2-Way Set-Associative Cache (LRU Policy)

Set 0 has:

- Way 0 → Block 4 (last used 5 cycles ago).
- Way 1 → Block 8 (last used 20 cycles ago).

CPU requests Block 12 (maps to Set 0), miss occurs. LRU chooses Way 1 (Block 8) to replace, since it was used least recently.

6.5.4 Write Strategy: *What happens on a write?*

When the CPU writes to a memory address and the block is cached, we must decide:

1. **Where does the write go? (Write Policies)**
 - Only in the cache?
 - In both cache and main memory?
2. **What if the address is not in the cache? (Write Miss Policies)**
 - Do we bring it into the cache before writing?
 - Or write directly to memory without caching it?

Step 1: Write Policies (*Where does the write go?*)

- **Write-Through.** The value is written **to both**:
 - The cache block (if present).
 - The corresponding location in main memory.

This ensures that the **memory** is always **up to date**. But it needs a **write buffer** to avoid stalling the CPU while memory is updated.

- **Write-Back.** The value is written **only to cache**. The modified cache block is marked **dirty** (with a **Dirty Bit**). And the main memory is updated **only when the dirty block is evicted**.
 - ✓ Less **memory traffic** (multiple writes to same block cost only one memory update).
 - ✓ **Faster in write-intensive workloads.**
 - ✗ More **complex** (need dirty bits and eviction logic).
 - ✗ Main memory may be **out of date** until eviction.

■ Step 2: Write Miss Policies (*What if the block is NOT in the cache?*)

- **Write Allocate** (aka **Fetch on Write**). On a write miss:

1. Load the block into the cache (same as a read miss).
2. Then perform the write in the cache.

Good for temporal locality: if we write once, we might write again soon.

- **No Write Allocate** (aka **Write Around**). On a write miss, **do not** load block into cache, but write directly to main memory. So, skip the cache and go directly to the main memory. **Good when writes are rare or sequential** (no need to keep data around).

Cache Policy	Write Miss Policy	Why
Write-Back	Write Allocate	Writes are kept in cache, so allocate makes sense.
Write-Through	No Write Allocate	Avoids unnecessary block fetch before write.

Table 26: Typical combinations in real systems.

Aspect	Write-Through	Write-Back
Write Location	Cache + Memory	Cache only
Memory Coherence	Always up to date	Updated only on eviction
Memory Traffic	High	Low
Complexity	Low	High (dirty bit)

Miss Policy	Write Allocate	No Write Allocate
Action on Miss	Bring block into cache	Write directly to memory
Best With	Write-Back	Write-Through

⌚ Write Operation Timeline

We'll use:

- **CPU**: issuing a write to memory address X.
- **Cache**: may or may not contain block X.
- **Memory**: slower DRAM.

1. Write Hit (data is already in cache).

- **Write-Through**

- (a) CPU writes new data into cache line for block X.
- (b) Cache immediately sends the same write to main memory.
- (c) Memory is always **coherent** (same as cache).

⌚ **Latency**: cache hit time (plus background memory update if write buffer exists).

⌚ **Memory traffic**: 1 write to memory.

- **Write-Back**

- (a) CPU writes new data into cache line for block X.
- (b) Set **dirty bit** = 1 for that cache line.
- (c) Main memory is **not updated now**, will be updated when block X is evicted.

⌚ **Latency**: just the cache hit time.

⌚ **Memory traffic**: 0 writes now (delayed until eviction).

2. Write Miss (data is not in cache).

- **Case A: Write Allocate**

- **Write-Back + Write Allocate (most common)**:

- (a) Cache fetches block X from memory into a cache line.
- (b) CPU writes new data into cache line.
- (c) Dirty bit is set.
- (d) Memory is updated **later** when block is evicted.

⌚ **Memory traffic**: 1 block read (fetch) + 1 block write later at eviction.

- **Write-Through + Write Allocate**:

- (a) Cache fetches block X from memory into cache.
- (b) CPU writes into cache line.
- (c) Cache immediately writes to main memory as well.

⌚ **Memory traffic**: 1 block read + 1 word write immediately.

- **Case B: No Write Allocate**

- **Write-Through + No Write Allocate** (*common combination*):

- (a) Cache does not load block X.
 - (b) CPU writes directly to memory.
 - (c) Cache content is unchanged.

- Memory traffic:** 1 word write only.

- **Write-Back + No Write Allocate:**

- (a) Cache does not load block X.
 - (b) CPU writes directly to memory.
 - (c) Cache content is unchanged.

- Memory traffic:** 1 word write only, no dirty bit set.

6.6 Miss Penalty Reduction

When a cache miss happens, the CPU must fetch the required block from a lower memory level, which takes much longer than a hit. This extra time is called the **Miss Penalty**. Because miss penalties can be **tens or hundreds of cycles**, reducing them has a huge impact on overall performance.

⌚ Role of the L2 Cache

The most common way to reduce miss penalty is to introduce a **Second-Level Cache (L2)** between the L1 cache and main memory.

Its main job is to **catch L1 cache misses** before they reach main memory. Since main memory is much slower than L1 or L2, resolving miss in L2 instead of DRAM dramatically reduces the miss penalty.

 **Larger than L1:** typically hundreds of KB to several MB.

 **Slower than L1, but still much faster than DRAM.**

- Usually **unified**: stores both instructions and data.
- Implemented in **SRAM** like L1, but with slightly longer access time.
- Can be **on-chip** (modern CPUs) or **off-chip** (older designs).

When an L1 miss occurs, instead of going directly to DRAM:

1. **CPU request** → goes to L1 cache.
2. If **L1 hit** → data returned immediately.
3. If **L1 miss** → request sent to L2 cache.
4. If **L2 hit** → data returned quickly to L1 (and CPU).
5. If **L2 miss** → data fetched from main memory (very slow).

It works because L1 caches are small (to keep them fast), so they miss more often. L2 caches are bigger and can hold more blocks, so they can keep data that L1 had to evict. This way, a **large fraction of L1 misses are “saved” by L2 before going to DRAM**.

⌚ AMAT in Presence of L2

Let's define:

- $HT_1 = L1$ hit time.
- $MR_1 = L1$ miss rate.
- $HT_2 = L2$ hit time.

- $MR_2 = \text{L2 Local Miss Rate}$ (fraction of L1 misses that also miss in L2):

$$MR_2 = \frac{\text{L2 misses}}{\text{L1 misses}} \quad (75)$$

It measures how good L2 is at catching L1 misses.

- $MP_2 = \text{Miss Penalty from L2 to main memory}.$
- $MR_{1,2} = \text{Global Miss Rate}$ (to main memory):

$$MR_{1,2} = MR_1 \times MR_2 = \frac{(\text{L2 misses})}{(\text{Total CPU Accesses})} \quad (76)$$

This tells us how often we actually have to access main memory.

The formula becomes:

$$AMAT = HT_1 + MR_1 \times (HT_2 + MR_2 \times MP_2) \quad (77)$$

Every L1 misses costs at least an **L2 access time** (HT_2). If L2 also misses (MR_2), we pay the extra penalty of going to main memory (MP_2).

Example 8: AMAT in Presence of L2

Imagine:

- $HT_1 = 1$ cycle
- $MR_1 = 5\%$
- $HT_2 = 10$ cycles
- $MR_2 = 20\%$
- $MP_2 = 100$ cycles

Step-by-step:

1. L1 hit $\rightarrow 1$ cycle.
2. L1 miss (5% of cases) \rightarrow go to L2 (10 cycles).
3. In 20% of those L2 accesses \rightarrow miss again and pay 100 cycles to access DRAM.

$$\begin{aligned} AMAT &= 1 + 0.05 \times (10 + 0.20 \times 100) \\ &= 1 + 0.05 \times (10 + 20) \\ &= 1 + 0.05 \times 30 \\ &= 1 + 1.5 = 2.5 \text{ cycles} \end{aligned}$$

Without L2:

$$AMAT = 1 + 0.05 \times 100 = 6 \text{ cycles}$$

L2 reduced the AMAT from **6 cycles** to **2.5 cycles**.

6.7 Design Space of Cache

Designing a cache involves tuning several parameters that all interact with each other. Changing one parameter usually impacts hit time, miss rate, and miss penalty, so it's always about trade-offs.

1. Cache Size

- **Bigger cache**

- ✓ Lower miss rate (can store more of the working set).
- ✗ Higher hit time (more bits to check, longer wires, more complex lookup).
- ✗ More area and power consumption.

- **Smaller cache**

- ✓ Faster hit time, lower power.
- ✗ Higher miss rate.

We can't make L1 huge without slowing it down, that's why big caches are placed at L2 or L3.

2. Block Size (Cache Line Size)

- **Larger block**

- ✓ Exploits spatial locality better (fetches more neighboring data on each miss).
- ✗ Higher miss penalty (more bytes to transfer on a miss).
- ✗ Possible increase in **conflict misses** if large blocks reduce the number of total blocks.

- **Smaller block**

- ✓ Lower miss penalty (less data to fetch per miss).
- ✓ More blocks → potentially lower conflict misses.
- ✗ Less spatial locality exploitation.

3. Associativity

- **Direct-Mapped**

- ✓ Simple, fast hit time.
- ✗ High conflict miss rate.

- **Fully Associative**

- ✓ No conflict misses.
- ✗ Slow and expensive to implement for large caches.

- **N-Way Set Associative**

- ✓ Balance between conflict misses and complexity.
- ✗ Higher hit time than direct-mapped.

General trend: Increasing associativity reduces **conflict misses** but increases **hit time** and hardware complexity.

4. **Replacement Policy.** Determines which block is evicted on a miss (in associative caches):

- **LRU (Least Recently Used):** good with temporal locality but expensive for high associativity.
- **Random:** simple, avoids pathological patterns, but ignores usage.
- **FIFO:** simple, but can evict still-hot blocks.

Trend: For low associativity (2-4 ways), LRU is common; for high associativity, pseudo-LRU or random is used.

5. **Write Policy.** Two decisions here:

- **Write-Through vs Write-Back**
 - **Write-through:** simpler, always update memory; higher memory traffic.
 - **Write-back:** update cache only, mark dirty; write to memory only on eviction.
- **Write Allocate vs No Write Allocate**
 - **Write Allocate:** bring the block into cache on a write miss; better with write-back.
 - **No Write Allocate:** write directly to memory on a miss; better with write-through.

Parameter	Bigger/More	Smaller/Less
Cache size	↓ miss rate, ↑ hit time	↑ miss rate, ↓ hit time
Block size	↑ spatial locality, ↑ penalty	↓ penalty, ↓ spatial gain
Associativity	↓ conflict misses, ↑ hit time	↑ conflict misses, ↓ hit time
Replacement	↑ LRU accuracy, ↑ hardware cost	↓ complexity, ↑ random evictions
Write policy	↓ traffic (write-back), ↑ complexity	↑ traffic (write-through), ↓ complexity

Table 27: Design Trade-offs.

6.8 Cache Miss Classification

Whenever the CPU requests data, the **cache may or may not contain it**:

- ✓ If it's there ⇒ **Cache Hit**
- ✗ If it's not ⇒ **Cache Miss**

There are **different causes** for cache misses. Understanding their classification helps in designing **better memory hierarchies**.

■ The 3 Caches Model of Cache Misses

This classic classification includes:

- **Compulsory Misses** (aka **Cold Start Misses**) occurs the **first time** a block is accessed, so it's **not in the cache yet**. It occurs because the cache starts empty and we must load the block from main memory. For example:

```
1 // first time a[0] is accessed → compulsory miss
2 int x = a[0];
```

These **always occur**, even with infinite cache. Can be **reduced with prefetching or larger blocks** (use spatial locality).

- **Capacity Misses**. The cache is **too small** to hold all the data the program is working on. So, **useful blocks are evicted** and later needed again, causing misses. For example:

```
1 for (i = 0; i < 100000; i++) {
2     sum += arr[i];
3 }
```

If the array **arr** is larger than the cache, older blocks will be replaced, causing misses later. Reducing these misses requires a **larger cache**. Still subject to **cost, power, and hit time trade-offs**.

- **Conflict Misses** (aka **Collision** or **Interference Misses**). Happens when **two or more blocks map to the same cache set**, and they overwrite each other, **even if the cache is not full**. Occurs in **direct-mapped caches** and **set-associative caches** with limited associativity. For example:

```
1 // Suppose these addresses map to the same set:
2 int a = mat[0][0];           // → cache set X
3 int b = mat[512][0];         // → same cache set X
```

They keep evicting each other: **ping-pong effect**. Can be reduced by: **higher associativity, better mapping, using victim caches**.

- (in multiprocessors) **Coherence Misses**. Introduced later with **shared-memory multiprocessors**. It occurs when **another processor** modifies a **cache block** that a core was using. The block must be **invalidated**, leading to a **miss on next access**. For example:

- Processor A and B both cache variable x .
- A writes to x , B's copy is invalidated.
- B tries to read x , then coherence miss.

These are handled by **cache coherence protocols** (like MESI), and it is only relevant in **multi-core systems**.

Type	Cause	How to reduce
Compulsory	First-time access	Larger blocks, prefetching
Capacity	Working set > cache size	Larger cache
Conflict	Multiple blocks map to the same location	Higher associativity, victim cache
Coherence	Invalidation from another processor	Coherence protocol (MESI, MOESI, etc.)

Table 28: Cache Miss Classification.

6.9 Improving Cache Performance

One of the main challenges in memory system design is achieving **low-latency memory access** while maintaining **high throughput and efficiency**. Since memory is a major performance bottleneck, architects aim to **minimize delays caused by cache misses**.

▣ Key Metric: Average Memory Access Time (AMAT)

The **Average Memory Access Time (AMAT)** equation defines how effective a memory hierarchy is:

$$\text{AMAT} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty} \quad (78)$$

Each term in this formula captures a different aspect of performance:

- **Hit Time:** Time to access data in the cache (when it's a hit).
The **unit is cycles**.
- **Miss Rate:** Fraction of memory accesses that result in a miss.
The **unit is %**.
- **Miss Penalty:** Time to fetch data from the next memory level.
The **unit is cycles**.

◎ Goal: Minimize AMAT

We can improve cache performance through **three orthogonal strategies**, each targeting one component of AMAT:

1. **Reduce Miss Rate.** This means **making the cache more effective** in storing the right data. **Techniques include:**

- Larger caches.
- Larger blocks (but with care).
- Higher associativity.
- Software and hardware prefetching.
- Compiler optimizations.

✓ Result: fewer accesses go to slower memory levels.

2. **Reduce Miss Penalty.** Once a miss happens, the delay can be high. We aim to **reduce the time spent waiting** by:

- Prioritizing reads.
- Using victim caches.
- Sub-blocking.
- Early restart or critical word first.
- Second-level caches.

- Non-blocking caches.

✓ **Result:** the *impact* of a miss is reduced.

3. **Reduce Hit Time.** Even on a hit, we want the **fastest possible access**. This can be achieved by:

- Making the L1 cache small and simple.
- Avoiding address translation delays.
- Using pipelined writes.
- Virtually-indexed, physically-tagged caches.

✓ **Result:** faster access when hit, then lower AMAT baseline.

Strategy	Targeted Term	Goal
Reduce Miss Rate	Miss Rate	Prevent misses from occurring
Reduce Miss Penalty	Miss Penalty	Make misses less costly
Reduce Hit Time	Hit Time	Make hits as fast as possible

6.9.1 Reducing Miss Rate Techniques

6.9.1.1 Increasing Cache Size

One intuitive way to reduce the number of cache misses is to simply make the cache **larger**. A bigger cache can store **more blocks**, which increases the chance that **future accesses will hit**, especially in programs with large working sets.

⌚ Effect on Miss Rate

Increasing cache size reduces capacity misses. If a miss happens because a previously used block was evicted to make room, then increasing the cache may **prevent such eviction**. Miss rate **typically decreases** with increasing cache size, but with **diminishing returns**. Empirical studies show that doubling the cache size doesn't halve the miss rate.

⚠ Trade-offs and Drawbacks

While it's effective, this approach isn't free. Larger caches come with:

Factor	Effect
Hit Time	✖ Usually increases (larger cache, more complex logic and longer access path).
Area	✖ Grows significantly (more memory cells, more tag storage).
Power Consumption	✖ Higher due to increased access energy.
Cost	✖ More silicon area, higher production cost.

So, a larger cache **may reduce miss rate**, but could also **increase hit time** (page 257), hurting the *overall AMAT*.

In practice:

- **L1 caches** are kept **small** and fast (to preserve CPU lock frequency).
- **L2/L3 caches** are larger and slower, used to **absorb the remaining misses**.

Increasing cache size is an **effective way to reduce capacity misses**, but it **comes with trade-offs in area, latency, power, and cost**. It **must be balanced** with the needs of the processor pipeline and application behavior.

6.9.1.2 Increasing Block Size

When a program accesses a memory location, it is **very likely** to soon access **neighboring addresses**: this is **Spatial Locality** (page 298).

By **increasing the cache block size** (i.e., the amount of data brought in with each cache miss), we **preload nearby data into the cache before it's even requested**.

⌚ Effect on Miss Rate

Compulsory misses⁸ are reduced: since a single access brings in **more than one word**, **fewer unique block fetches** are needed. The **effectiveness depends on how well the program exhibits spatial locality**⁹.

Example 9

```
1 for (int i = 0; i < 100; i++)
2     sum += a[i];
```

If $a[i]$ and $a[i + 1]$ are in the same larger block, we benefit from fetching both in one miss.

⚠ Drawbacks and Trade-offs

While increasing block size helps with **spatial locality**, there are serious **down-sides**:

Trade-Off	Impact
Increased Miss Penalty	Larger blocks take more time to transfer from lower levels → slower cache refill.
Fewer Blocks in Cache	For a fixed cache size, larger blocks = fewer blocks stored → more conflict/capacity misses .
Wasted Bandwidth	If spatial locality is weak , much of the block may be unused .
Power and Energy	More data fetched and written increases energy usage.

⚠ At some point, **increasing block size hurts more than it helps**.

⁸Compulsory Misses, or Cold Start Misses, occurs the first time a block is accessed, so it's not in the cache yet (see page 319 for more).

⁹Spatial Locality says: “if we access this memory address, we will probably access nearby ones soon”.

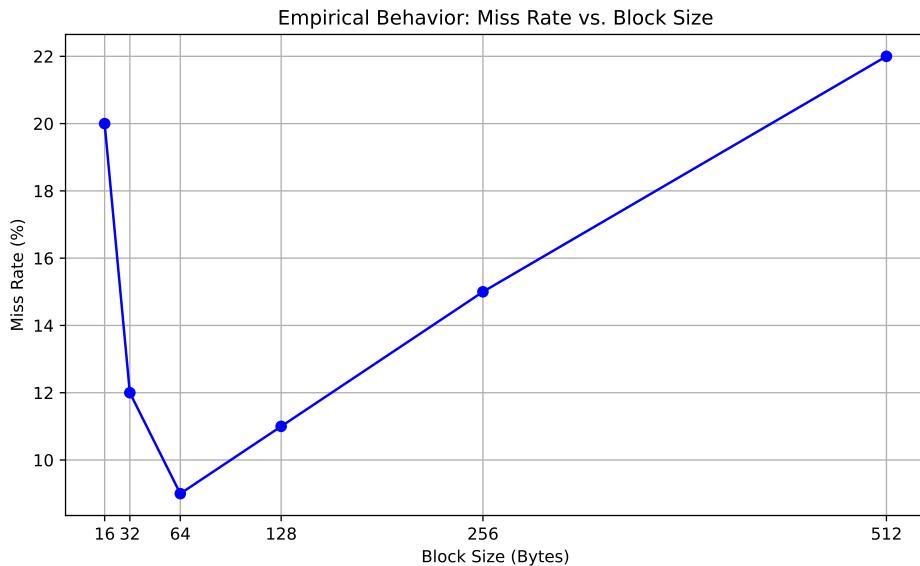
Empirical Behavior

Empirical behavior refers to how something actually behaves in practice, based on measured data, experiments, or real-world observations, not just theoretical analysis.

In our context (Block Size vs. Miss Rate), theory says: “*increasing block size should reduce compulsory misses thanks to spatial locality*”. But when researchers actually test it on real systems and real programs, they observe that:

- Miss rate **initially decreases** as block size increases.
- But **beyond an optimal size**, it begins to **increase again** due to:
 - **Fewer total blocks** (higher conflict/capacity pressure).
 - **Higher miss penalty**.

This forms a **U-shaped curve** when plotting miss rate vs. block size.



- Common block sizes: **16, 32, 64 bytes** (chosen based on workload).
- Larger blocks may be beneficial for **instruction caches** (sequential fetch).
- Data caches are more sensitive, **balance carefully**.

Increasing block size helps **✓ reduce compulsory misses** by exploiting **spatial locality**, but only up to a point. It can **✗ increase conflict and capacity misses** and make **✗ misses more expensive**.

6.9.1.3 Increasing Associativity

💡 What is Associativity?

In the cache memory context, **Cache Associativity** describes how many places in the cache a given block of main memory can be stored. It answers the question: “if we need to store this memory block in the cache, how many cache locations are possible candidates?”. There are three main types:

- **Direct-Mapped** (page 305): Associativity = 1.
- **Fully Associative** (page 306): Associativity = Number of lines in the cache.
- **N-Way Set-Associative** (page 306): Associativity = n -way, with $n > 1$.

So, associativity is essentially the “number of candidate cache lines per set” that a block can occupy. Higher associativity reduces conflict misses but increases hardware cost and lookup time.

In **direct-mapped caches** (page 305), each memory block maps to exactly **one cache location**. This simplicity results in **fast access**, but it can also cause many **conflict misses**, especially when multiple frequently used addresses map to the same index.

💡 Solution: Use **set-associative caches** (page 306), where each set contains **multiple slots (ways)** for blocks to be stored. This **reduces conflicts** by offering **more choices** for where a block can go.

📘 Terminology Refresher

Cache Type	Behavior
Direct-Mapped	1 block per set (1-way associative).
2-Way Set-Associative	2 blocks per set.
Fully Associative	Any block can go anywhere in the cache (all in one set).

💡 Effect on Miss Rate

- As associativity increases, **conflict misses decreases**.
- **Capacity and compulsory misses stay the same**.
- **Fully associative caches** (page 306) **eliminate all conflict misses**, but they are expensive and slower.

Example 10

If the cache is direct-mapped, these blocks **overwrite each other**, causing many **conflict misses**:

```
1 a[i] → cache set 5
2 b[i] → cache set 5
3 c[i] → cache set 5
```

But with a **4-way set-associative cache**, all 3 blocks could coexist in **set 5**, avoiding those misses.

A Trade-offs and Drawbacks

Trade-Off	Effect
Hit Time Increases	More ways = more tag comparisons = slower access.
Area and Complexity Increase	More comparators, multiplexers, and logic.
Power Usage Increases	More hardware is active on each access.

Design trade-off: There's a **sweet spot**; usually **2-way or 4-way** associativity gives good results without hurting hit time too much.

E Rule of Thumb: 2:1 Cache Rule

A direct-mapped cache of size N has **about the same miss rate** as a 2-way set-associative cache of size $\frac{N}{2}$. This tells us that we can **trade associativity for size**: higher associativity sometimes compensates for a smaller cache.

In general, increasing associativity **reduces conflict misses** by allowing more blocks to coexist in the same set, but it comes at the cost of **higher hit time, area, and power**. Designers must **balance associativity with hit time and hardware complexity**.

6.9.1.4 Victim Cache

Even with set-associative caches, **conflict misses**¹⁰ can still occur when frequently accessed blocks map to the same set and keep evicting each other.

Instead of increasing associativity (which slows down hit time and increases complexity), we can add a small buffer to catch those recently evicted blocks.

💡 Basic Idea

A **Victim Cache** is:

- **Small** (e.g., 4-16 entries).
- **Fully Associative** cache.
- Placed **between the main cache and the next lower level** in the hierarchy.

Process:

1. If a block is evicted from the main cache, it is moved into the Victim Cache.
2. On a cache miss in the main cache, **check the victim cache**:
 - If found there → **swap** it with the block in the main cache (**Victim Block Swap**).
 - If not found → fetch from lower memory level.

⌚ Effect on Miss Rate

Victim caches **reduce conflict misses** significantly, especially for **direct-mapped caches**. For example, if two memory blocks keep evicting each other from the same line in a direct-mapped cache, the victim cache allows them to **coexist** without forcing an immediate main-memory access.

⚠ Trade-offs and Costs

Factor	Impact
✓ Miss Rate Reduction	Very effective for small conflict sets.
✓ No Change to Main Cache Hit Time	The main cache remains fast.
✗ Extra Hardware	Additional storage, tag comparators, swap logic.
✗ Area and Power	Still small, but non-zero overhead.

¹⁰Conflict Miss happens when two or more blocks map to the same cache set, and they overwrite each other, even if the cache is not full.

✖ Simplified Flow

1. CPU Request.
2. Main Cache Hit?
3. Yes → Done.
4. No, check Victim Cache.
5. Found?
6. Yes → Swap into Main Cache (Fast Recovery).
7. No → Fetch from Lower Level.

A **victim cache** is a low-cost, high-impact technique for **reducing conflict misses** without hurting hit time. It acts as a **safety net** for blocks that would otherwise be lost to the next memory level.

6.9.1.5 Pseudo-Associativity & Way Prediction

Increasing associativity reduces **conflict misses**, but it also **slows down the hit time** because the cache must compare **more tags in parallel**.

These techniques aim to **get some of the benefits of associativity without paying the full hit-time penalty**.

- **Way Prediction** (for Set-Associative Caches).

A The Problem with Set-Associative Caches. In an **n-way set-associative cache**:

- Each **set** has n **ways** (cache lines).
- When we access memory, we:
 1. Find the correct **set** using the index bits.
 2. **Compare the tag** against all n ways in parallel to see if one matches (a *hit*).

The main issue, however, is that comparing all n ways in parallel requires n **comparators** and n **tag lookups** per cycle. This **increases both power consumption and access time**.

✓ The Idea of Way Prediction. Instead of checking *all* ways at once, **predict** which may be likely to have the block:

- Hardware keeps a “**way predictor**” (like a small table) that records the *last used way* for each set. The **Way Predictor Table (WPT)**:
 - * **Indexed by:** the **set index bits** from the address.
 - * **Entry size:** enough bits to encode one of the n ways. For example, 4-way cache, need 2 bits per entry.
 - * **Storage size:**

$$\text{Number of sets} \times \text{Bits per entry}$$

- * **Contents:** “*Last hit way*” for that set.

Visually:

Set Index	Predict Way (2 bits)	Meaning
0	00	Way 0
1	10	Way 2
2	01	Way 1
3	11	Way 3
4	10	Way 2
5	00	Way 0

Each row corresponds to a set in the cache. The Predicted Way column stores a small code (2 bits here) telling which way is most likely to have the next hit for that set. The Meaning column shows the human-readable interpretation of those bits (Way 0, Way 1, etc.).

This tiny table sits alongside the cache, indexed by the same set index bits as the cache itself.

For example, 4-way cache with 1024 sets:

- * Index bits: 10 bits (for 1024 sets).
 - * Way encoding: 2 bits.
 - * Way Predictor Table size: $1024 \times 2 = 2048$ bits ≈ 256 bytes.
 - When accessing memory:
 1. **Address split:**
 - * **Tag bits.**
 - * **Index bits**, select set in both the cache and WPT.
 2. **WPT lookup:** read the predicted way (e.g., “*Way 2*”).
 3. **Access predicted way:** fetch **tag** and **data** from only that way.
 4. **Tag check:**
 - ✓ If tag matches \rightarrow ✓ **Correct Prediction Hit**.
 - ✗ If tag mismatches \rightarrow ✗ **Prediction Failure**:
 - (a) The **remaining $n - 1$ ways** in the same set are read sequentially or in parallel (depends on design) to locate the block.
 - (b) **Two possibilities:**
 - **Block is found in one of these ways** \rightarrow **Misprediction-Hit**. For example, predicted *Way 2*, but the block is in *Way 0* (same set!).
 - **Block not found in any way** \rightarrow *Miss* \rightarrow fetch from lower level (L2, DRAM).
- The WPT is **updated only on hits**.

Example 11

Let's assume a **4-way set-associative cache**, with:

$$\text{WPT}(\text{Set } 12) = \text{Way } 3$$

* Access 1 - Correct Prediction Hit

- Address \rightarrow Set 12, Tag matches *Way 3* \rightarrow Data returned in 1 cycle.
- WPT(12) stays *Way 3*.

* Access 2 - Misprediction-Hit

- Address \rightarrow Set 12, WPT says *Way 3*.
- Tag mismatch \rightarrow Search remaining ways: Check *Way 0* \rightarrow match found.
- Data returned with +1 cycle penalty.
- Update WPT(12) = *Way 0*.

* Access 3 - Miss

- Address → Set 12, WPT says *Way 0*.
- Tag mismatch → Search remaining ways: No match → Miss.
- Load block into *Way 2* (chosen by LRU).
- Update WPT(12) = *Way 2*.

This predict → check predicted way → search specific alternative way(s) → update table loop is the real hardware logic.

? Why this Works. Many programs have **temporal locality**, so if a block is in the cache now, it will likely still be in the cache when accessed again. This means the **last used way** for a set is often the **next** way that will be hit. Accuracy of way prediction can be 80-95%, so most accesses are fast.

- ✓ Reduces **power** (only one way read at a time for most hits).
- ✓ Potentially reduces **access time** (one comparator instead of n).
- ✓ Keeps the **flexibility** of set-associative caches.
- ✗ **Penalty** for misprediction: 1-2 extra cycles.
- ✗ Needs **way predictor table** (small overhead).
- ✗ Works best when there's strong locality, worse with random access.

Example 12: The Hotel Room Analogy

Imagine a big hotel with **many floors (sets)** and **several rooms per floor (ways)**.

- Each **guest** is like a **data block**.
- We are the **CPU**, looking for a specific guest.

✗ Without way prediction

Every time we visit a floor, we knock on the door of **every single room** on that floor in parallel to see if our guest is inside.

- Fast in a weird “everyone answers at once” way,
- But it’s noisy, energy-draining, and requires a lot of “staff” to knock on all doors simultaneously.

✓ With way prediction

The receptionist keeps a **small notebook** (the *way predictor*) that remembers **the last room each guest stayed in** on each floor.

- When we arrive at a floor, the receptionist says: “Last time this guest was on this floor, they were in **Room 3**”.
- We knock **only on Room 3** first.
 - * If they’re still there → we’re done in 1 knock (fast, quiet).
 - * If they’ve moved → we knock on the other doors (a little slower).

❷ Why it works

Guests usually stay in the **same room** until checkout (*temporal locality*). Even if they move, most of the time we guess right, so we save knocking on every door.

Way prediction is **not** about magically knowing where *new guests* will go, it’s about remembering where **current guests** are likely to be, and checking there first.

- **Pseudo-Associativity** (or **Column-Associativity**, for Direct-Mapped caches) is a clever trick to make a direct-mapped cache behave a bit like a 2-way set-associative cache without doubling the parallel hardware.

⚠ The Problem it solves. A **direct-mapped cache** has **exactly one possible location** for each memory block. If two blocks map to the same line, they constantly evict each other (**conflict misses**). A **2-way set-associative cache** would fix that, but it requires:

- Two tag comparisons in parallel.
- Two data arrays to read in parallel.
- More area, power, and complexity.

✓ The Idea of Pseudo-Associativity. Instead of having two ways checked in parallel, a pseudo-associative cache checks them **sequentially**:

- For each index, store **two possible locations** (primary and secondary).
- On an access:
 1. **Check the primary location** (like a normal direct-mapped cache). If tag matches → **hit** in 1 cycle.
 2. If not a match → check the **secondary location** (the “pseudo-way”).
 - * If tag matches → **hit** but with an extra cycle of latency.
 - * If still no match → miss → fetch from lower memory.

✗ How the secondary location is found. Often, the secondary location is **another index computed from the address** (e.g., by flipping the MSB of the index or XORing certain bits). This ensures that if two

blocks map to the same primary index, they *probably* don't also share the same secondary index.

⌚ **Why it works.** Most accesses still hit in the primary location, so it is fast. If there's a conflict, there's a good chance the block is in the secondary location, then conflict miss avoided, but with slightly higher hit time. Greatly reduces **conflict misses** with little extra hardware compared to a full set-associative design.

⚠️ Key Difference

- **Way Prediction:** Works for set-associative caches → reduces hit time.
- **Pseudo-Associativity:** Works for direct-mapped → reduces conflict misses.

Both techniques are **hybrid approaches**. They try to **combine the speed of direct-mapped** with the **miss rate benefits of associativity**. Best suited for designs where **hit time is critical** and full associativity is too costly.

6.9.1.6 Hardware Prefetching (Instructions & Data)

Even with optimized cache size, associativity, and victim caches, **compulsory misses** will still occur: the first time a block is accessed, it has to be fetched.

Idea: If we can *predict* what data or instructions will be needed soon, we can bring them into the cache early. When the CPU asks for them, they're already there, so **no miss penalty**.

- **Instruction Prefetching.** Instruction streams often have **strong spatial locality** (instructions are usually executed sequentially). **Hardware prefetch logic** detects sequential access patterns and preloads **the next block of instructions** while the current one is being executed.

- ✓ Reduces **compulsory misses** for sequential code.
- ✓ Keeps the instruction pipeline fed.

Example 13: Alpha AXP 21064

The **Alpha AXP 21064** was a **high-performance microprocessor** from Digital Equipment Corporation (DEC) released in the early 1990s, one of the first in the Alpha architecture family.

It was one of the earliest commercially available CPUs to implement **stream buffer-based hardware prefetching** for caches.

The 21064's data cache was direct-mapped, which is fast but prone to **conflict misses**. To help hide memory latency, DEC implemented **hardware prefetching using stream buffers**:

- When the CPU missed in the cache for an address, the hardware **predicted the next sequential address** that would be needed.
- It loaded them into **stream buffers** in advance, so when the CPU accessed them, they were ready.

This prefetch mechanism was **triggered automatically** by the detection of sequential access patterns (common in loops and array processing).

Before this, hardware prefetching wasn't common in general-purpose CPUs, prefetching was mostly a software/compiler job. The Alpha 21064's implementation showed that **simple hardware mechanisms could significantly reduce memory stall cycles**. Later CPUs (from MIPS, PowerPC, Intel, AMD, etc.) adopted similar or more advanced prefetching techniques.

In summary, on a miss:

1. Fetch the **requested block** into the instruction cache.
2. Also fetch the **next sequential block** into a separate **instruction stream buffer**.

If the next fetch hits in the stream buffer, the block is moved to the cache, and the next block after that is prefetched.

- **Data Prefetching.** Data Prefetching means **bringing data into the cache before the CPU actually requests it**, so that when it needs it, it's already there. This **hides memory latency**. Prefetching can be done by:

- **software** (compiler or programmer inserts prefetch instructions).
- **Hardware** (CPU detects patterns and fetches ahead automatically).

Data prefetching is more challenging because data access patterns can be irregular. However, in many workloads (e.g., scientific computing, multimedia), data is accessed in predictable sequences.

1. CPU detects a **pattern of accesses**, usually sequential or strided addresses (e.g., $A[0]$, $A[1]$, $A[2]$, ...).
2. It predicts the **next address** we'll want.
3. It sends a request to memory/cache hierarchy **in advance**.
4. When we actually access that address later, it's already in the cache, then **hit** instead of miss.

More specifically:

1. **Cache Miss.** CPU requests address X . It's not in the cache, then **miss**. Cache controller fetches block X from memory **into the cache**.
2. **Prefetch the next block(s).** The hardware prefetcher guesses the **next block(s)** we'll need (usually sequential, if we missed on block X , next will be $X+1$, $X+2$, etc). These predicted blocks are **not** put into the main cache immediately. Instead, they are stored in the **D-stream buffer**.

❓ What is a “hardware data prefetcher” with D-stream buffers?

It's a **small dedicated hardware unit** sitting next to the cache. A **D-Stream Buffer** is a temporary holding area for **prefetched data** that isn't in the cache yet, but might be needed soon.

It's a sort of “waiting room” for data that the CPU hasn't requested yet, but that the prefetcher predicts it will.

3. **CPU requests prefetched data.** If later the CPU requests block $X+1$:
 - (a) Prefetcher sees it's already in the D-stream buffer.
 - (b) Moves it **quickly** from the D-stream buffer into the cache → **fast hit**.

❓ Why not put prefetched data directly into the cache?

Prefetched data might never be used. If we store it directly in cache, we risk evicting useful data (**cache pollution**). Stream buffers keep it separate until it's actually needed, so only *useful* prefetched data enters the cache.

Example 14: Data Prefetching Analogy

Think of the **cache** as the table we're eating at, and the **D-stream buffer** as a **side tray** where the waiter puts the next dish they think we'll want.

- If they guessed right, they slide it onto our plate instantly.
- If they guessed wrong, they take it away without messing with our current plate.

Hardware prefetching **guesses future accesses** and proactively brings data or instructions closer to the CPU, reducing **compulsory** and some **capacity misses**, but must be carefully tuned to avoid **wasting bandwidth**.

6.9.1.7 Software Prefetching (Compiler-Controlled)

Hardware prefetchers work automatically but **cannot always detect complex or irregular access patterns**. A **compiler** (or programmer) often has more information about upcoming memory accesses and can explicitly request that certain data be loaded in advance.

This **targets compulsory and some capacity misses** by overlapping memory fetches with useful computation.

Idea

The **compiler analyzes the program** (often using profiling data) to identify memory accesses that could cause misses. It inserts **special prefetch instructions** into the code, placed *before* the data is actually needed. These instructions **fetch the block into cache or into a register**, giving enough time for the load to complete before use.

Types of Software Prefetching

1. **Register Prefetching.** Prefetch into **CPU registers** for immediate use. For example, HP PA-RISC “load” instruction. Useful for small, predictable data needs.
2. **Cache Prefetching.** Prefetch into **cache** not into registers. For example, MIPS IV, PowerPC, SPARC v9 have dedicated cache prefetch instructions. Data stays in cache until needed by normal load.

Example 15

Here, the compiler issues a prefetch for `A[i+8]` while the CPU is still working on `A[i]`.

- Without Prefetching

```
1 for (i = 0; i < N; i++) {
2     sum += A[i]; // May cause many misses
3 }
```

- With Compiler-Inserted Prefetch

```
1 for (i = 0; i < N; i++) {
2     __prefetch(&A[i+8]); // Load ahead of time
3     sum += A[i];
4 }
```

Performance Considerations

Prefetch instructions **occupy instruction slots**, so the CPU must be able to issue them without stalling important work. In **superscalar processors**, the difficulty is reduced because multiple instructions can be issued in parallel. The **distance between prefetch and actual use** is crucial:

- Too close → miss penalty still occurs.
- Too far → data might be evicted before use.

Trade-offs

Pros	Cons
<ul style="list-style-type: none"> ✓ Reduces compulsory/capacity misses for predictable accesses. ✓ Works for irregular patterns hardware may miss. ✓ Can be tuned by the compiler based on profiling. 	<ul style="list-style-type: none"> ✗ Wastes bandwidth if data not used. ✗ Adds extra instructions (execution overhead). ✗ Needs accurate prediction of future accesses.

Software prefetching leverages **compiler insight** to prepare data before it's needed. It is especially useful for **predictable loops** and **regular patterns**, but must be carefully tuned to avoid wasted bandwidth and unnecessary instruction overhead.

6.9.1.8 Compiler Optimizations (for Cache Performance)

Even with prefetching, a program's **structure** might still cause many cache misses. The compiler can transform the code and data layout to **increase spatial and temporal locality**, thus reducing misses **before the hardware even runs the code**.

▣ Basic Idea

1. **Profile the program** to understand typical memory access patterns.
2. **Reorganize code and restructure data** so that:
 - Related data is stored **together** (improves spatial locality).
 - Frequently reused data stays **in cache longer** (improves temporal locality).
3. Generate machine code that exploits these new layouts.

☒ Techniques

- **Merging Arrays.** Combine multiple arrays into one array of structures. Reduces conflict misses and improves spatial locality.

— Before

```
1 int val[SIZE];
2 int key[SIZE];
```

— After

```
1 struct {
2     int val;
3     int key;
4 } merged_array[SIZE];
```

Both **val** and **key** for the same index end up in the **same cache block**.

- **Loop Interchange.** Swap the order of nested loops to match memory layout (row-major or column-major). Improves spatial locality by accessing consecutive memory locations.

— Before

```
1 for (j = 0; j < 100; j++)
2     for (i = 0; i < 5000; i++)
3         x[i][j] = 2 * x[i][j];
```

— After

```
1 for (i = 0; i < 5000; i++)
2     for (j = 0; j < 100; j++)
3         x[i][j] = 2 * x[i][j];
```

Now accesses move **sequentially through memory**.

- **Loop Fusion.** Combine two loops that iterate over the same data into one loop. Reduces repeated loading of the same data.

– Before

```

1 for (i = 0; i < N; i++)
2   A[i] = B[i] + 1;
3 for (i = 0; i < N; i++)
4   C[i] = A[i] * 2;
```

– After

```

1 for (i = 0; i < N; i++) {
2   A[i] = B[i] + 1;
3   C[i] = A[i] * 2;
4 }
```

Improves **temporal locality**, A[i] stays in cache between operations.

- **Loop Blocking (Tiling).** Process data in **small chunks (blocks)** that fit in the cache. Common in matrix operations.

– Before

```

1 for (i = 0; i < N; i++)
2   for (j = 0; j < N; j++)
3     process(A[i][j]);
```

– After

```

1 for (ii = 0; ii < N; ii += B)
2   for (jj = 0; jj < N; jj += B)
3     for (i = ii; i < ii + B; i++)
4       for (j = jj; j < jj + B; j++)
5         process(A[i][j]);
```

Improves **temporal locality**, elements in the block are reused before being evicted.

⚠ Trade-offs

Pros	Cons
✓ Works for predictable patterns.	✗ Requires compiler complexity.
✓ Reduces both compulsory & capacity misses.	✗ May increase code complexity & size.
✓ No hardware cost.	✗ Limited effect on unpredictable access patterns.

Compiler optimizations restructure both **data** and **loops** to **align program access patterns with cache organization**. When combined with hardware techniques, they significantly reduce miss rate **without changing the hardware**.

6.9.2 Reducing Miss Penalty Techniques

6.9.2.1 Read Priority over Write on Miss

A cache miss can happen in two main ways:

- **Read miss:** The CPU needs data from memory (load instruction).
- **Write miss:** The CPU wants to store data to memory (store instruction) and the block isn't in the cache.

When a miss happens, the cache must send a **request to the next memory level** (L2, DRAM). But memory is slow and usually **only one outstanding miss can be serviced at a time per port**.

Idea: when both reads and writes compete for memory access, **give priority to reads** because:

- Reads **stall the CPU** immediately (we can't execute without the data).
- Writes often go to a **write buffer** and can be delayed without stopping execution.

② **So are we assuming that a read and a write miss happen at the same time?** Yes, we're talking about a situation where two independent memory operations are outstanding in the cache/memory system at the same time. That can happen in a modern CPU because pipelines are out-of-order and non-blocking. For example, imagine:

1. Cycle 10 → CPU issues load from address A → **read miss**.
2. Cycle 12 → CPU issues store to address B → **write miss**.
3. Both go into the miss queue.
4. Memory controller sees both pending, **picks read first** to reduce CPU stall.
5. Write sits in the write buffer until read completes, then gets sent to memory.

The “*simultaneous*” read miss and write miss doesn't mean they happen in the *exact same cycle*, it means they are **both pending before memory has finished handling the first one**, so the controller has to choose an order.

■ Basic Concept

- Case 1: **Write-Through with Write Buffer**. Writes go immediately to memory but are buffered to avoid stalling.
 1. CPU issues a **write** → cache updated (write-through, value is written to cache block and main memory) → **write buffer** stores the memory update so the CPU doesn't wait.
 2. Later, CPU issues a **read miss**.

3. Before fetching from memory, the cache controller **checks the write buffer**:
 - If the read address matches an address in the write buffer → **read directly from the buffer** so we don't get stale data.
 - If no match → bypass the write requests in the buffer and **immediately send the read miss to memory** (read priority).
4. Writes in the buffer are sent to memory later, in background, when the bus is free.

⌚ **Why it works:** ensures read misses aren't delayed by pending writes and still guarantees correctness if the needed data is in the write buffer.

- Case 2: **Write-Back with Dirty Block Replacement.** Writes are deferred until a dirty block is evicted.

1. Cache uses **write-back** policy → writes only update cache and mark block as dirty.
2. On a **read miss**, we need to bring a new block into the cache.
3. If the victim block is **clean** (Dirty Bit is 0) → just replace it.
4. If the victim block is **dirty** (Dirty Bit is 1):

✗ **Without read priority:** we'd write the dirty block back to memory *before* fetching the new block → long delay before CPU gets the new data.

✓ **With read priority:**

- (a) Copy the dirty block into the **write buffer** (so it can be written to memory later).
- (b) **Immediately start fetching the new block** for the read miss.
- (c) Once memory is free, the dirty block in the write buffer is written back in the background.

⌚ **Why it works:** the CPU gets the missing data sooner and the dirty block write-back doesn't block the read fetch.

These “*read priority*” methods aren't about *reordering a generic queue arbitrarily*, they're about **changing the handling sequence inside the cache controller's normal process** so that when a **read miss** and a **write-related action** would normally compete for memory access, the **read gets serviced first** within that specific policy's procedure.

In other words:

- In **write-through with write buffer**: instead of draining the write buffer first, we pause it (unless the read needs its data) and go straight to the read miss fetch.
- In **write-back with dirty block replacement**: instead of writing the dirty victim to memory *before* fetching the new block, we park it in the write buffer and fetch the read's block immediately.

✓ Benefits

- ✓ Reduces **CPU stall time** by overlapping write operations with read miss handling.
- ✓ Keeps the processor busy instead of waiting for slow write completions.

⚠ Challenges

- ✗ **RAW hazard through memory**. If the read request matches an address in the write buffer, we must read from there instead of memory.
- ✗ **Write buffer overflow**. If too many writes accumulate, read requests can still be delayed.
- ✗ **Control complexity**. Cache controller must track dependencies between writes and reads.

6.9.2.2 Sub-block Placement

Normally, when a cache miss occurs, the cache **fetches the entire block** (e.g., 64 bytes) from the next memory level. This **increases miss penalty** because:

- The transfer time is proportional to block size.
- We might not even need all the words in that block.

Idea: Break each cache block into **smaller pieces (sub-blocks)** with **individual valid bits**. This way, on a miss, we can **load only the needed sub-block** instead of the entire block.

❖ How it works

- Each block is divided into **sub-blocks** (e.g., 4 words per block, becomes 4 sub-blocks).
- Each sub-block has its own **Valid Bit**.
- On a miss:
 - Only the requested sub-block is fetched.
 - Other sub-blocks in the same block are marked invalid until fetched.

✓ Benefits

- ✓ **Lower Miss Penalty:** Only part of the block is transferred, then less latency.
- ✓ **Better Bandwidth Usage:** Avoids fetching unused data.
- ✓ Can be useful for **irregular access patterns** where spatial locality is weak.

It is useful when **memory bandwidth is limited** or when **access patterns are sparse** (e.g., random access over large arrays). It is also useful in systems where the **miss penalty is high** and spatial locality is low.

⚠ Drawbacks

- ✗ **Lower Spatial Locality Exploitation:** Since we don't fetch the whole block, nearby addresses may miss later.
- ✗ **Extra Tag/Valid Bit Storage:** Need valid bits for each sub-block.
- ✗ **More Complex Control Logic:** Cache controller must track which sub-blocks are valid.

Sub-block placement trades off **spatial locality** for **lower miss penalty** and better bandwidth usage. It's especially useful when **large block sizes** would otherwise cause excessive transfer delays.

6.9.2.3 Early Restart & Critical Word First

Early Restart & Critical Word First are two related techniques that aim to let the CPU resume execution faster when a cache miss occurs.

In a normal cache miss:

1. The cache requests the **entire block** from the next memory level.
2. The CPU waits until **all words** in the block have arrived.
3. Only then does the CPU resume execution.

⚠ Problem: This wastes cycles if the needed word arrives before the rest of the block.

✓ Solution: Send the **critical word** (the one the CPU actually requested) first, so execution can resume earlier.

☰ Techniques

- **Early Restart**

- Fetch the block **in normal order** from memory.
- As soon as the **requested word** arrives, **send it to the CPU** and let it resume.
- Continue transferring the rest of the block into the cache in the background.

For example:

- Block size: 4 words (W0, W1, W2, W3).
- CPU requests W2.
- Memory returns W0 → W1 → **W2** → W3.
- On W2's arrival → **CPU resumes**, W3 still coming.

✓ Benefit

- ✓ Works well when **memory returns words sequentially** in block order.
- ✓ Reduces **CPU stall time** for some misses.

- **Critical Word First** (aka **Wrapped Fetch**).

- When a miss occurs, **start fetching at the requested word**, then wrap around to the start of the block.
- The **requested word is fetched first**, not just delivered first.

For example:

- Block size: 4 words (W0, W1, W2, W3).

- CPU requests W2.
- Memory returns **W2 → W3 → W0 → W1**.
- W2 is delivered to CPU immediately → stall time minimized.

✔ Benefit

- ✓ More effective than Early Restart when **requested word is late in block order**.
- ✓ Particularly useful when **miss penalty is dominated by first-word latency**.

⚠ Trade-offs

- ✗ **Controller Complexity:** Critical Word First requires non-sequential wrap-around fetch logic.
- ✗ **Lower Spatial Locality Use:** If other words are needed soon, re-ordering fetch might not help.
- ✗ **Memory Bus Efficiency:** Some memory systems prefer sequential transfers for efficiency.

Both Early Restart and Critical Word First **reduce miss penalty by delivering the requested word sooner**, but Critical Word First goes further by **fetching it first** from memory.

6.9.2.4 Non-Blocking Caches

⚠ The problem with classic (blocking) caches

In a **blocking cache**, when we get a **miss**, the cache controller stalls the CPU until the miss is resolved and the data is in the cache. Even if the CPU could access **other data already in the cache**, it can't. It must wait for the miss to finish.

```
1 sum += A[i];    // miss → wait...
2 sum += B[0];    // could be a hit, but must wait anyway
```

✓ The idea of Non-Blocking Caches

A **Non-Blocking Cache** lets the CPU **continue to service hits even when a miss is being processed**.

- This means the cache can have **multiple memory transactions in-flight**.
- The hardware must track each outstanding miss (via **Miss Status Holding Registers (MSHRs)**).
- The CPU can keep executing instructions that do not depend on the missing data.

_HIT UNDER MISS

Hit Under Miss is the **simplest level of non-blocking cache capability**. The cache can handle hits to other addresses while one miss is outstanding.

1. Miss on address X → request sent to memory.
2. While waiting, CPU requests address Y.
3. If Y is in the cache → serve it immediately (**hit under miss**).
4. If Y is also miss → depends on capability (next topic, miss under miss).

✓ **Benefit:** hides some of the miss by overlapping **independent hits** with miss handling.

❓ Going beyond: “Miss Under Miss”

A more advanced feature is **Miss Under Miss** (aka **Multiple Outstanding Misses**):

- Allows **starting a second miss while the first one is still being serviced**.
- Requires **multiple Miss Status Holding Registers (MSHRs)** to track all pending misses, and **non-blocking memory controllers** and **out-of-order handling** of cache fills.

✓ **Benefit:** maximizes **memory-level parallelism (MLP)** and is particularly useful for high-latency main memory or multi-core systems.

⚠ Trade-offs

- ✗ **Hardware Complexity:** Tracking multiple misses is non-trivial..
- ✗ **Extra Area & Power:** Miss Status Holding Registers (MSHRs) consume silicon and energy.
- ✗ **Consistency Handling:** Must ensure memory ordering rules are respected.

Non-blocking caches **reduce effective miss penalty** by overlapping miss servicing with useful work:

- **Blocking cache:** One miss stops *all* accesses.
- **Hit Under Miss:** One miss stops *new misses*, but hits to other addresses still work.
- **Miss Under Miss:** Multiple misses can be serviced in parallel.

6.9.2.5 Second-Level and Multi-Level Caches

Even with all optimizations, an **L1 cache miss penalty** can be very large, typically **tens to hundreds of CPU cycles** if it goes directly to main memory (DRAM).

Idea: Insert an **intermediate cache level** (L2, or even L3) between L1 and main memory:

- L2 is **larger** and **slower** than L1, but much **faster** than DRAM.
- L3 is even larger and slower, often shared between cores in multicore systems.

This hierarchy **catches many L1 misses** before they reach main memory, greatly reducing **average miss penalty**.

Access path:

CPU → L1 Cache → L2 Cache → L3 Cache → Main Memory

On an L1 miss:

1. Check L2 cache.
2. If hit → fetch from L2 (much faster than DRAM).
3. If miss → proceed to L3 or main memory.

Effect on AMAT

Let's extend the AMAT formula for two levels:

$$\text{AMAT} = \text{Hit Time}_{L1} + \text{Miss Rate}_{L1} \times (\text{Hit Time}_{L2} + \text{Miss Rate}_{L2} \times \text{Main Memory Penalty}) \quad (79)$$

If **L2 has a low miss rate**, the main memory penalty is rarely paid. Even if L2 is slower than L1, it's **orders of magnitude faster** than DRAM.

Design Considerations

1. Size

- L2/L3 caches are **much larger** (hundreds of KB to MBs) to reduce **capacity misses**.
- Larger size = slower hit time → acceptable because they're not on the CPU's critical cycle path.

2. Inclusion Policy

- **Inclusive:** L2 contains all of L1's contents (**simplifies coherence in multiprocessors**).
- **Exclusive:** L2 and L1 store different data (**maximizes total cache capacity**).
- **Non-inclusive, non-exclusive (NINE):** No strict relationship.

3. **Associativity.** L2/L3 are often higher associativity (e.g., 8-way, 16-way) to minimize conflict misses.

4. Shared vs. Private

- **Private L2:** One per core (reduces access contention).
- **Shared L3:** Serves multiple cores (improves sharing, helps coherence).

Example 16

If:

- L1 miss rate = 5%.
- L2 hit rate = 90%.
- L1 miss penalty (to L2) = 10 cycles.
- L2 miss penalty (to DRAM) = 100 cycles.

Then:

$$\begin{aligned} \text{AMAT} &= \text{L1 Hit Time} + 0.05 \times (10 + 0.10 \times 100) \\ &= \text{L1 Hit Time} + 0.05 \times (10 + 10) \\ &= \text{L1 Hit Time} + 1 \end{aligned}$$

Main memory accesses are now rare, reducing overall penalty.

⚠ Trade-offs

Pros	Cons
✓ Large reduction in average miss penalty.	✗ Extra latency for L1 misses (must check L2).
✓ Reduces main memory bandwidth demand.	✗ Increased area and power.
✓ Helps multi-core coherence (with inclusion).	✗ More complexity in cache hierarchy management.

Multi-level caches act as **latency buffers** between the fast CPU and slow DRAM. By catching most L1 misses in L2 (and L2 misses in L3), they **drastically reduce the effective miss penalty** and improve system throughput.

6.9.2.6 Merging Write Buffers

When a **write-through cache** is used (see page 0), every store operation that misses in the cache writes data to the **write buffer** before going to memory. The write buffer:

- Temporarily holds these writes so the CPU doesn't stall waiting for memory.
- Sends them to memory **in the background**.

⚠ Problem: If multiple writes target **adjacent addresses** in the same memory block:

- Without merging → each write generates a **separate memory transaction**.
- This **wastes bandwidth** and increases **miss penalty** for other operations waiting for the bus.

✓ Solution: Equip the write buffer with the ability to **merge** multiple writes to the **same block** into a **single memory access**.

▣ Structure of a merging write buffer entry

Each entry in a merging write buffer stores:

1. **Block address**. Aligned to the cache block size (e.g., 32B or 64B), it is used to **detect if a new write is to the same block**.
2. **Data field**. Space for an entire cache block. Initially empty except for the bytes written so far.
3. **Byte-valid mask**. One bit per byte (or per word) in the block. It indicates which parts of the block have valid, updated data from the CPU.

❖ How it works

Let's say cache block size is 32 bytes, and CPU writes 8 bytes at a time.

1. **CPU write hits the cache**
 - If **write-through cache**: update the cache line and send the same write to the write buffer.
 - If **write-back cache with eviction**: when a dirty line is evicted, it goes into the write buffer.
2. **Write buffer search**. When sending the write to the write buffer:
 - (a) **Align the address** to the block boundary (e.g., address 0x1048, block base 0x1040).
 - (b) Search the write buffer for an **entry with the same block base address**.

2a. **Merge into existing entry (same block)**. If found:

- Place the new 8 bytes into the **correct offset** inside the existing data field.
- Set the corresponding bits in the **byte-valid mask** to 1.
- No new memory transaction is started yet, the entry remains until it's full or needs to be flushed.

2b. **Create new entry (different block)**. If not found:

- Allocate a new entry in the write buffer.
- Store:
 - Block base address.
 - Data in the correct offset.
 - Byte-valid mask with those bytes set.

3. **Flushing to memory**. When the write buffer decides to send data to memory (due to bus availability or eviction):

- (a) If the byte-valid mask = “all 1s” → send the full block.
- (b) If mask is partial → send a *partial write* (if memory bus supports it) or first read the block from memory, merge, then write.

After sending, the entry is removed from the buffer.

Example 17

Cache block size is 32 bytes, addresses 0x1000-0x101F:

1. Write 8B @ 0x1000. New entry:

- Base = 0x1000
- Data[0:7] = *value*
- Mask = 00000011 (binary per word).

2. Write 8B @ 0x1008. Same base, then merge into existing entry:

- Data[8:15] = *value*
- Mask updated

3. Write 8B @ 0x1010. Same base, merge again.

4. When buffer flushes: send one **32B write** to memory.

➊ Benefits

- ✓ Fewer **memory transactions** → less contention on the memory bus.
- ✓ Better **bandwidth efficiency**.

- ✓ **Lower miss penalty** for reads waiting on bus availability.

Particularly effective for write-through caches. Also helps in write-back caches when dirty blocks are evicted and written in pieces.

⚠ Trade-offs

- ✗ **More Complex Hardware:** Need matching logic and byte masking.
- ✗ **Extra Area:** Each entry must store masks and partial data.
- ✗ **Delay in Writes:** If the buffer is full, CPU may still stall.

Merging write buffers **reduces the number of memory transactions** for adjacent writes, saving **bandwidth** and **reducing contention**, thus lowering effective **miss penalty** for both reads and writes.

6.9.3 Reducing Hit Time Techniques

6.9.3.1 Small and Simple L1 Caches

Even if a cache hit is **much faster than a miss**, it still lies on the **critical path** of every memory access. The **L1 cache** is accessed **every CPU cycle**, so even a **1-cycle delay** in hit time can slow down the processor's maximum clock frequency.

✓ **Idea:** Keep the L1 cache **small and simple** so that:

- Tag comparison is quick.
- Access time is minimal.
- It can run at the CPU's full clock speed.

💡 Why Small Caches Are Faster

- **Shorter wires** inside the cache array.
- Fewer sets → simpler indexing logic.
- Smaller tag array → faster tag comparison.
- Lower associativity → fewer parallel tag comparators.

Example 18

A 32KB, 8-way L1 cache might take **2-3 cycles** to access. A 16KB, direct-mapped L1 cache could be accessed in **1 cycle**. The smaller L1 will have a **higher miss rate**, but the much lower hit time may still **improve overall AMAT**.

⚠ Trade-off: Miss Rate vs. Hit Time

Cache Size	Hit Time	Miss Rate
Small	✓ Low	✗ High
Large	✗ High	✓ Low

This is why L1 caches are **small** (16-64 KB) and **fast**, while **L2/L3 caches** are much larger but slower.

❖ Design Practice

- Keep **L1 cache** small enough for **1-cycle accesses** at CPU clock speed.
- Use **L2 and L3 caches** to absorb most misses from L1.
- Choose associativity carefully, often **2-4 way** in L1 to keep hit time low.

A **small and simple L1 cache** ensures **minimum hit time**, allowing the CPU to run at full speed. Although the miss rate is slightly higher, the **overall performance** often improves because the **baseline access time** is much lower.

6.9.3.2 Avoid address translation (VIPT caches)

Remark 1: Virtual Page and Physical Page

Before starting this chapter, we need to review some concepts from the Operating Systems topic. The VIPT and VIVT caches are based on virtual and physical pages. But what are these pages, exactly?

❷ Why paging exists

Programs think they have a **large, continuous address space** (from 0 up to 2^{32} or 2^{64}). In reality, RAM is **finite** and fragmented. Paging is the Operating System and Hardware mechanism that:

- Divides memory into **fixed-size blocks**.
- Maps **virtual addresses** (from the program's point of view) onto **physical addresses** (real RAM).

▀ Virtual Page (VP)

A **Virtual Page (VP)** is a block of the **virtual address space**. All virtual pages have the same fixed size (commonly 4 KB, but can be larger like 2 MB, 1 GB). Each process has its **own set of virtual pages** that ensures **isolation** (process A can't see B's memory). Think of a virtual page as **a numbered box in the program's illusion of memory**.

▀ Physical Page (Frame)

A **Physical Page** (often called a **Frame**) is a block of **physical RAM** of the same size as a virtual page. The OS maps virtual pages to physical frames. Multiple virtual pages (from different processes, or even the same process) may map to the **same physical frame** (shared memory). So a physical page is **a real storage slot in RAM**.

⌚ Virtual $\xrightarrow{\text{to}}$ Physical mapping

Translation is stored in the **page table** (per process). Each entry contains:

Virtual Page Number (VPN) = Physical Frame Number (PFN) + flags

The **page offset** (low-order bits) is unchanged. For example, (4 KB pages, then 12-bit offset):

- **VA** = [VPN | Offset] = 0x7FFE1 | 0x234
- Page table: VPN 0x7FFE1 \rightarrow PFN 0x01AB0
- **PA** = [PFN | Offset] = 0x01AB0 | 0x234

❓ Why the offset doesn't change

Memory is divided into **pages** of fixed size (e.g. 4 KB). A page is like a **block of consecutive bytes**. The **page offset** tells us *where inside that page* the byte/word lives. When the OS translates a **virtual page** to a **physical frame**, it only changes the **page number part** (VPN to PFN). The **offset inside the page stays identical**, because the location inside the block doesn't change. So the mapping looks like this:

$$\begin{array}{lcl} \text{Virtual Address} & = & [\underbrace{\text{VPN}}_{\text{translated}} \quad | \quad \underbrace{\text{Offset}}_{\text{same}}] \\ \\ \text{Physical Address} & = & [\underbrace{\text{PFN}}_{\text{translated}} \quad | \quad \underbrace{\text{Offset}}_{\text{same}}] \end{array}$$

❓ What the offset is for

The offset is the **position inside the page**. For example, if page size is 4 KB, that is 2^{12} bytes, then the offset is 12 bits. It tells us: “*go to byte #X inside this 4 KB block*”. So the **VPN** chooses **which block of memory** and the **offset** chooses **which byte within that block**.

Think of a city with identical hotel buildings, each having rooms numbered 0-4095. The **hotel building number** is like the page number (VPN $\xrightarrow{\text{to}}$ PFN translation may change this). The **room number inside the hotel** is like the page offset, it never changes. If we were in room 0x234 before translation, we'll still be in room 0x234 after translation, just in a different hotel building.

Example 19: Offset

Imagine:

- Page size = 4 KB
- Virtual Address = 0x7FFE_1234

Breakdown:

- VPN = 0x7FFE1
- Offset = 0x234

Page table says: VPN 0x7FFE1 $\xrightarrow{\text{to}}$ PFN 0x01AB0. So physical address is 0x01AB0_234. Notice the **low 12 bits** (0x234) didn't change.

Quick Summary

- **Virtual Page (VP)**: block in program's virtual memory.
- **Physical Page (Frame)**: block in actual RAM.
- **Page offset**: same in both, ensures byte position inside page.
- **Page table**: maps VP → Physical Frame.

Analogy

- Virtual Pages = hotel room numbers on a booking site (illusion of continuous availability).
- Physical Pages = real hotel rooms across different hotels (real memory).
- Page table = reception desk list mapping booking numbers to actual room keys.
- TLB (see the explanation below) = post-it note with our current room key to avoid asking the receptionist each time.

A The problem. Modern CPUs use **virtual addresses** (what our program sees) instead of physical addresses (actual DRAM locations). The mapping from virtual to physical is done by the **page table** in memory. But doing a page table lookup in memory **every time** would be **too slow** (extra memory access per instruction!).

Q The solution. The **Translation Lookaside Buffer (TLB)** is a **small, fast cache** that stores the **most recent virtual-to-physical address translations**. On every memory access:

1. CPU generates a **virtual address**.
2. The **TLB is checked** to see if the translation is cached.
 - ✓ **TLB hit** → use physical address immediately.
 - ✗ **TLB miss** → walk the page table in memory (slow), then insert result into TLB.

⚠ Another Problem...

If we always wait for the TLB before touching the cache, every cache hit would be delayed by the **TLB translation time**. For example:

- Cache hit time alone: 1 cycle.
- TLB lookup: 1 cycle.
- Sequential: 2 cycles hit time, this doubles the critical latency.

Since cache hits happen the vast majority of the time ($> 90\%$), even one extra cycle is a **huge slowdown**. The motivation is simple: we want to **avoid adding TLB latency to every cache hit**.

✓ The Solution: Virtually Indexed, Physically Tagged (VIPT) Caches

A **Virtually Indexed, Physically Tagged (VIPT) Cache** is a cache design where:

- The **index** (which set to look at) is taken from the **virtual address**.
- The **tag** (to confirm if it's the right block) comes from the **physical address** after TLB translation.

This allows:

- ✓ **Cache access to start immediately** using the virtual address.
- ✓ **Correctness guaranteed** by checking the physical tag once translation is ready.

❖ How VIPT Works

1. CPU issues a **virtual address**.
2. The cache controller uses the **index bits from the virtual address** to pick a cache set right away (no need to wait for TLB).
3. Meanwhile, the **TLB** translates the virtual page number into the physical page number.
4. The **physical tag** is compared against the tags stored in the chosen set.
5. If there's a match → **cache hit**.

❓ But wait. Why does VIPT use an index and a tag?

- If we wait for the TLB before cache access → **slow** (extra cycle or two).
- If we skip the TLB and just use virtual addresses in the cache (Virtually Indexed, Virtually Tagged design, VIVT) → **synonyms problem**.

⚠ Synonyms Problem

Virtual memory allows **different Virtual Addresses (VAs)** to map to the **same Physical Address (PA)**. This can happen with: shared memory between processes; memory-mapped files; multiple mappings (aliases) of the same physical page. So **two VAs can point to the same PA**.

❓ Why is this a problem for caches? Suppose the cache is indexed/tagged by **virtual address (VIVT design)**:

- $VA1 \rightarrow PA = 0x1000$
- $VA2 \rightarrow PA = 0x1000$ (alias, same physical page)

But since the cache only sees the *virtual tags*:

- $VA1$ might go into Set 3 with Tag $VA1$.
- $VA2$ might go into Set 7 with Tag $VA2$.

Now the *same physical memory block* is **duplicated** in two cache lines.

⚠ Where is the danger? Imagine:

- CPU writes to $VA1$ (updates cache line in Set 3).
- CPU later reads $VA2$ (different cache line in Set 7).

As a result, the CPU gets **obsolete data**, because the write through $VA1$ wasn't visible via $VA2$. That breaks memory coherence, which is unacceptable.

✅ How VIPT avoids this. In a **VIPT cache**:

- **Index** comes from **page offset bits** (same in VA and PA; see the remark box on page 357).
- **Tag** checked against physical address (after TLB).

So even if $VA1$ and $VA2$ differ in their higher bits, as long as they map to the same PA:

- They land in the **same set** (because offset bits are identical, same in Virtual Address and Physical Address).
- And only one of them matches in the tag comparison (because tag uses PA, not VA).

Result: no duplication, no stale data.

❓ Why is “the offset doesn't change” essential for VIPT? Since offset bits are guaranteed to be identical between VA and PA: we can safely use them to index into the cache before translation finishes. No risk of mismatch. That's the **page offset trick**.

Example 20: Why is VIPT better than VIVT?

Imagine:

- Page size = 4 KB, then 12-bit page offset.
- Cache block size = 64 B, then 6-bit block offset.
- Cache = 4 KB, direct-mapped → needs 64 sets → 6 index bits.

So address breakdown:

$$\begin{aligned}\text{Virtual Address (VA)} &= [\text{Virtual Page Number} \mid \text{Page Offset}] \\ \text{Page Offset (12 bits)} &= [\text{Index}(6) \mid \text{Block offset}(6)]\end{aligned}$$

Translation:

- VA upper bits (Virtual Page Number) → via TLB → Physical Page Number (PPN).
- Page Offset stays the same in VA and PA.

⚠ Case: Two aliases

Suppose the OS maps two different virtual pages to the **same physical page** (shared memory):

- VA1 = 0x4000_1234 → PA = 0x1000_1234
- VA2 = 0x8000_1234 → PA = 0x1000_1234

Both point to the same physical memory at 0x1000_1234.

✗ If cache is Virtually Indexed, Virtually Tagged (VIVT)

- VA1 index = bits [11 : 6] of 0x4000_1234 → suppose = Set 10.
- VA2 index = bits [11 : 6] of 0x8000_1234 → suppose = Set 42.

So:

- VA1 loads physical block → stored in Set 10 with Tag = VA1.
- VA2 loads the same physical block → stored in Set 42 with Tag = VA2.

Now the same physical data lives **twice** in the cache. If CPU writes through VA1, the copy in Set 42 (VA2) is stale, then **synonym problem!**

✓ If cache is Virtually Indexed, Physically Tagged (VIPT)

Index still comes from Page Offset bits [11 : 6]. Since VA1 and VA2 share the **same Page Offset** (0x234), they both map to Set 10. Tag check uses **Physical Page Number** (0x1000) from TLB. So:

- VA1 loads block → goes into Set 10, Tag = 0x1000.

- VA2 accesses → goes to Set 10, Tag = 0x1000 → **hits the same line.**

Only one copy exists, consistent for both aliases, then **synonym problem solved!**

✓ Key takeaway

- **Virtually Indexed, Virtually Tagged (VIVT):** Index and tag from Virtual Address, then duplicates possible.
- **Virtually Indexed, Physically Tagged (VIPT):** Index from Virtual Address's page offset (safe), tag from Physical Address, then no duplicates.

⚠ Ok, now we can sleep peacefully knowing that the Synonym Problem has been solved.

Not at all! There is still a **second problem**, the **cache size restriction**, which comes directly from the bit-level math of VIPT.

First, let's review how indexing works. Cache indexing uses some bits of the **address**:

[Tag | Index | Block Offset]

- **Block Offset:** selects the byte/word inside the cache line (e.g. 6 bits for 64B line).
- **Index:** selects the cache set (which row of the cache to look at).
- **Tag:** checked against stored tag to confirm it's the right line.

We said:

- The **index bits must come only from the Page Offset** (because those bits are identical in VA and PA).
- BUT, if the cache needs *more index bits* than what fits inside the Page Offset, we'll accidentally use **VPN bits**, which may change after translation, then **synonym problem reappears**.

Therefore, we **must restrict the cache size** to ensure that the index bits are a subset or a set of the page offset bits.

✓ The Cache Size Restriction formula

$$\begin{aligned} \text{Cache size per way} &\leq \text{Page size} \\ \text{Total cache size} \div \text{Associativity} &\leq \text{Page size} \end{aligned} \quad (80)$$

Why? Because the size of the cache and the associativity is equal to indexing per way. If that exceeds the page size, then the index needs bits outside the page offset.

❓ How CPUs deal with this

Designers choose **L1 sizes and associativity** carefully so the restriction is met. For example, many modern CPUs use 32 KB, 8-way L1 caches with 4 KB pages, that is safe, because $4 \text{ KB} \times 8 = 32 \text{ KB}$. Higher-level caches (L2, L3) are usually **physically indexed, physically tagged (PIPT)**, so they don't suffer this limitation.

So the second problem is VIPT is only safe if the cache isn't “*too big per way*”. Otherwise, synonym issues come back.

Example 21: Second VIPT problem

Imagine:

- Page size = 4 KB = 2^{12} = 12 offset bits.
- Block size = 64 B 2^6 = 6 block offset bits.

That leaves 6 bits from the offset for **indexing**. So **64 sets per way** max. Now, if cache is 32 KB, 4-way associative, we have 8 KB per way:

- Needs 7 index bits.
 - But offset gave us only $12 - 6 = 6$ safe index bits.
- ⚠️** 1 index bit would come from VPN, so potential synonym!

So max L1 size in this case:

$$\text{Page size (4 KB)} \times \text{Associativity (4)} = 16 \text{ KB}$$

If we want a 32 KB VIPT L1 cache, we must increase associativity (e.g., 8-way).

Key Takeaways: VIPT Caches

- **Motivation:** Avoid adding TLB latency to every cache hit.
- **VIPT Idea:**
 - Use **virtual address offset bits** to index the cache (safe, identical in VA and PA).
 - Use **physical address bits** (after TLB) for tag comparison.
- **Synonym Problem:** Two different VAs can map to the same PA.
 - In VIVT: block duplicated in different sets → stale data risk.
 - In VIPT: same offset ensures same set; physical tag ensures one unique copy.
- **Page Offset Trick:** Offset doesn't change in VA→PA mapping, so it can safely be used for indexing.

- **Cache Size Restriction:**

$$\frac{\text{Cache size}}{\text{Associativity}} \leq \text{Page size}$$

Otherwise index bits may come from VPN, reintroducing synonym issues.

- **Practical Note:** CPUs carefully size L1 caches (e.g., 32KB, 8-way with 4KB pages is safe). Higher-level caches (L2, L3) are usually PIPT.

6.9.3.3 Pipelined Writes

Modern CPUs are deeply **pipelined**, they execute multiple instructions per cycle, often out-of-order. If the **cache write path** is slow (due to checking tags, updating data arrays, handling dirty bits, etc.), it could:

- Increase the **hit time** (because reads and writes compete for resources).
- Stall the **pipeline**, lowering CPU throughput.

Idea: Pipeline the cache write operations, so multiple writes can be **in flight simultaneously**, just like instructions in a CPU pipeline.

❖ How it works

Instead of treating a **write** as one monolithic operation, **split it into stages** (decode, tag check, data array update, buffer, etc.). New writes can be **accepted every cycle** even if the previous one hasn't finished yet. Reads can often proceed independently while writes are being finalized. This allows:

- ✓ Increases **throughput** of cache writes (important for store-intensive workloads).
- ✓ Prevents **write bottlenecks** in superscalar and out-of-order pipelines.
- ✓ Allows **parallelism**, reads can often proceed even while earlier writes are completing.
- ✓ **Reduced average hit time**, since the CPU doesn't have to wait for each write to complete before issuing the next one.

⚠ Challenges

- ✗ **Hazard detection:** Must ensure reads don't bypass pending writes to the same address.
- ✗ **Extra control logic:** More complex write pipeline control.
- ✗ **Area/Power:** Additional buffering and pipeline registers.

However, **pipelined writes** improve cache performance by allowing multiple writes to overlap, reducing their effect on **hit time** and preventing them from slowing down the CPU pipeline.

Example 22

Imagine a write normally takes 3 cycles: Address decode, Tag check, Data write. **Without pipelining:** The CPU must wait three cycles; then, the throughput is one write every three cycles. But **with pipelining**, a new write can enter each cycle, then throughput becomes 1 write per cycle (latency still 3, but overlapped).

6.9.3.4 Small Sub-Blocks for Write-Through Caches

In a **write-through cache** (page 0):

- Every store operation updates **both the cache and main memory**.
- This can generate a **lot of traffic** to memory (or to the write buffer if one is present).
- If the cache block size is large (e.g., 64B), updating the entire block for a single small write is **wasteful** and increases **hit time**.

Idea: Divide each cache block into **smaller sub-blocks**, each with its own **valid/dirty bit**, so that **only the sub-block being written** needs to be updated.

❖ How it works

A cache block (say 64B) is divided into multiple **sub-blocks** (e.g., 4 sub-blocks of 16B). Each sub-block has:

- Its own **Valid Bit** (to track if it contains valid data).
- Its own **Dirty/Modified Indicator** (if using write-back).

On a write:

- Only the affected sub-block is updated in cache.
- Only that sub-block's portion is sent to memory (write-through).

✓ Benefits

- ✓ Reduces **memory traffic** for small writes.
- ✓ Improves **hit time** in write-through caches since less data moves.
- ✓ Works well with workloads that update small pieces of data frequently.

It is particularly useful in **write-through caches** because every write goes to memory, and smaller granularity avoids wasting bandwidth. Less common in **write-back caches** since writes accumulate in cache and only dirty blocks are written back, so the benefit is smaller.

⚠ Trade-offs

- ✗ **Extra control logic:** Need valid/dirty bits for each sub-block.
- ✗ **More tag complexity:** Cache must track multiple sub-block states.
- ✗ **Spatial locality underused:** If later sub-blocks are needed, multiple transfers occur.

However, using **small sub-blocks** in write-through caches **reduces write bandwidth** and improves **hit time**, at the cost of slightly more control logic and reduced spatial locality exploitation.

6.9.4 Summary

The memory hierarchy is designed to give the *illusion* of a large, fast memory, by combining small/fast and large/slow storage. Programs naturally exhibit **temporal locality** (recently used data is reused soon) and **spatial locality** (nearby data is likely used together). Cache optimizations exploit these principles by reducing miss rate, miss penalty, or hit time.

✖ Reducing Miss Rate

- **Increase Cache Size** (page 323): Larger caches hold more blocks, reducing capacity misses, but slower and more complex.
- **Increase Block Size** (page 324): Larger blocks capture spatial locality, but too large blocks increase miss penalty and internal waste.
- **Increase Associativity** (page 326): Higher associativity reduces conflict misses, though hit time and hardware cost grow.
- **Victim Cache** (page 328): A small fully-associative buffer keeps recently evicted blocks, catching conflict misses.
- **Pseudo-Associativity / Way Prediction** (page 330): Simulates higher associativity with lower hit time by predicting the correct way.
- **Hardware Prefetching** (page 335): Dedicated hardware anticipates sequential or streaming accesses, hiding compulsory misses.
- **Software Prefetching** (page 338): Compiler inserts prefetch instructions for predictable patterns.
- **Compiler Optimizations** (page 340): Loop blocking, reordering, and array alignment improve temporal and spatial locality.

☒ Reducing Miss Penalty

- **Read Priority over Write** (page 342): On a miss, reads are prioritized so computation continues sooner.
- **Sub-block Placement** (page 345): Blocks divided into smaller sub-blocks; only the needed sub-block is fetched.
- **Early Restart and Critical Word First** (page 346): The CPU resumes as soon as the requested word arrives, without waiting for the full block.
- **Non-blocking Caches** (page 348): Allow hits during a miss (*hit under miss*) or multiple misses in parallel (*miss under miss*) using MSHRs.
- **Second-Level and Multi-Level Caches** (page 350): Extra cache levels catch L1 misses before they reach main memory, dramatically lowering effective penalty.
- **Merging Write Buffers** (page 353): Consecutive writes to the same block are merged, reducing memory traffic and contention.

⌚ Reducing Hit Time

- **Small and Simple L1** (page 356): Keeping L1 small (16-64 KB) and low-associativity enables single-cycle access, despite a slightly higher miss rate.
- **Virtually Indexed, Physically Tagged (VIPT)** (page 357): Starts cache lookup using virtual address bits while TLB translation proceeds in parallel; avoids TLB latency in the critical path.
- **Pipelined Writes** (page 366): Splits writes into pipeline stages, allowing multiple writes in flight and improving throughput.
- **Small Sub-blocks for Write-Through Caches** (page 367): Dividing cache lines into sub-blocks avoids writing entire blocks for small updates, saving time and bandwidth.

Cache design always balances three factors:

- **Miss Rate** (avoid misses with larger/associative caches and prefetching).
- **Miss Penalty** (recover faster with multi-level caches, early restart, non-blocking designs).
- **Hit Time** (keep L1 as fast as possible with simplicity and VIPT).

No single technique dominates; architects combine them carefully to match the workload and technology constraints.

Technique	Miss Rate	Miss Penalty	Hit Time	Complexity
Increase Cache Size	↓ (fewer capacity)	–	↑ (slower)	↑ (larger arrays)
Increase Block Size	↓ (spatial locality)	↑ (bigger transfer)	↑	Moderate
Increase Associativity	↓ (conflict misses)	–	↑ (tag checks)	↑
Victim Cache	↓ (conflict misses)	↓ (slightly)	=	Low–Moderate
Pseudo-Associativity / Way Prediction	↓ (like 2-way)	–	= (fast if predicted)	Moderate
HW Prefetching	↓ (compulsory/capacity)	–	=	↑ (stream buffers)
SW Prefetching	↓ (predictable patterns)	–	↑ (extra instr.)	Moderate (compiler)
Compiler Optimizations	↓ (better locality)	–	=	↑ (compiler effort)
Read Priority over Write	–	↓ (fewer stalls)	=	Low
Sub-block Placement	–	↓ (partial fetch)	=	↑ (valid bits)
Early Restart	–	↓ (resume earlier)	=	Low
Critical Word First	–	↓ (critical word first)	=	Moderate
Non-blocking Caches	–	↓↓ (overlap misses)	=	↑↑ (MSHRs)
Second-Level / Multi-Level	↓ (fewer L1 misses to DRAM)	↓↓	↑ (extra level)	↑↑
Merging Write Buffers	–	↓ (less traffic)	=	Moderate
Small & Simple L1	↑ (slightly more misses)	–	↓↓ (fast hits)	Low
VIPT Caches	–	–	↓ (parallel TLB/cache)	↑ (alias handling)
Pipelined Writes	–	–	↓ (higher throughput)	↑
Small Sub-blocks (Write-through)	–	–	↓ (less write work)	↑

Table 29: Cache optimization techniques and their impact on Miss Rate, Miss Penalty, Hit Time, and Complexity.

7 Multithreading (TLP)

7.1 Introduction

So far, we have seen how processors exploit Instruction-Level Parallelism (**ILP**). The central question was: *how many instructions from a single thread can be executed at the same time?*

There are two fundamental strategies we already saw:

1. **Superscalar (dynamic scheduling)**, page 95. The processor fetches multiple instructions per cycle. Hardware (using mechanisms like Tomasulo's algorithm, page 140, reorder buffers, page 222, register renaming, page 186, speculation) decides in real-time which instructions can issue and execute without violating dependencies.
 - ✓ **Advantage:** the compiler doesn't need to know the microarchitecture in detail; the processor dynamically extracts parallelism.
 - ✗ **Limitation:** this requires very complex hardware, because every cycle the CPU must check dependencies across many instructions, select ready ones, handle multiple wakeups, arbitrate functional units.
2. **VLIW (Very Long Instruction Word, static scheduling)**. The compiler does the hard work: it schedules instructions into “bundles” that can execute in parallel. Each bundle is one wide instruction word containing multiple operations.
 - ✓ **Advantage:** simpler hardware, just execute the bundles as they are.
 - ✗ **Limitation:** the compiler must know exact latencies, and it has to speculate on what parallelism is available. If the program has little ILP, many slots inside the wide word stay empty (wasted).

⚠ Why both superscalar and VLIW hit a wall

In both cases, the goal is to exploit parallelism inside one thread. But in practice, this parallelism is **limited**:

- **True dependencies (RAW hazards):** some instructions simply must wait for the result of previous ones.
- **Control dependencies (branches):** the processor doesn't know which path will be taken, so it may waste cycles.
- **Memory stalls:** loads that miss in the cache introduce huge delays.

As a result, even if the processor can *issue 8 instructions per cycle*, a single thread rarely provides that many independent instructions at the right time. This is why **real processors rarely go beyond 4-6 issue slots**. Further expansion would require more complex scheduling logic and consume more energy, yet most of those slots would remain unused.

In the issue diagrams we've already seen, this appears as **bubbles** (stalls) or **empty slots** (NO-OP): spaces where no instruction issues because none are ready, even though the hardware could handle more.

⌚ The motivation for Multithreading

This is exactly where multithreading enters the story. The reasoning is:

- If one thread alone cannot keep the machine busy, why not let **another thread** provide instructions in the meantime?
- Instead of waiting for a stalled instruction (e.g., a cache miss that takes 200 cycles), the processor can **switch to a different thread** whose instructions are ready.
- This way, the **unused slots of ILP** can be filled by exploiting **Thread-Level Parallelism (TLP)**.

So multithreading does not replace ILP, it complements it.

- ILP: look for parallelism inside a single instruction stream.
- TLP: look for parallelism across multiple instruction streams (threads).

⌚ First idea of hardware support

To do this, the processor needs:

- One **register state and program counter per thread** (so it can keep track of each thread's progress).
- A fast way to **select which thread to run** each cycle.
- Shared functional units and memory system (since we don't want to duplicate the entire core).

This is much cheaper than adding ever more ILP hardware (which grows super-linearly in complexity).

7.2 What is a Thread?

At the highest level, we write a **program**. When we run it, the operating system creates a **process**: an instance of that program in execution. A process has:

- Its own **address space** (the memory image of the program: code, heap, stack, data).
- **Operating system resources** (file descriptors, sockets, etc.).

Inside a process, we may have one or multiple **threads of execution**.

A **Thread** is the **smallest unit of execution that can be scheduled by the operating system**. Each threads has:

- Its own **program counter** (so it knows which instruction it is about to execute).
- Its own **register state** (so it can track local variables, intermediate values, return addresses).

But threads **share the same address space** of the process. This means:

- They see the **same code and global data**.
- They can **communicate cheaply via shared memory** (no need for explicit message passing as with processes).
- On the flip side, **incorrect synchronization can cause data races**.

In short:

- **Process**: **container** (resources + address space).
- **Thread**: **active execution flow** inside that container.

Thread Grain Size

A thread is not a fixed concept, it can be very **large** (millions of instructions, like an OS thread) or very **small** (a few instructions, as in hardware fine-grain MT). The **grain size** sets the stage for the type of multithreading we are interested in at the *architectural* level.

When we talk about “grain size” of a thread, we are asking: “*how much work does a thread represent before we would switch to another one?*”.

Hardware Multithreading Typology

- **Fine-Grain Thread**. Represents a **very small amount of work**, maybe only a **few instructions or a small function call**.
 - ✓ Useful when we want to expose **as much parallelism as possible** (so the hardware can interleave many threads).

✖ But the **overhead of switching and synchronization** increases.

- **Coarse-Grain Thread**. Represents a larger unit of work, like a loop iteration or an entire function execution.
 - ✓ Easier to manage, lower synchronization overhead.
 - ✖ But harder to keep hardware busy if that single thread stalls (e.g., on a cache miss).

Hardware multithreading usually works with **fine- to medium-grain threads**:

- ✖ It doesn't want to wait for an operating system context switch (thousands of cycles).
- ✓ Instead, it keeps multiple architectural states locally and can switch in a few cycles or even every cycle (depending on the MT strategy).

So both coarse-grain and fine-grain are hardware multithreading schemes and the difference is how frequently the hardware switches between threads (coarse, only on long stalls, fine, every cycle).

❓ But if HW Multithreading works with fine- and medium-grain threads, why do coarse-grain threads exist?

When hardware designers first explored multithreading, they wanted something **simple** and **cheap**. Duplicating register files for each thread is straightforward. Switching threads only when one blocks on a *very long latency event* (like a DRAM miss) avoids the cost of wasting cycles in a stall.

If a thread misses in memory, we flush or “drain” the pipeline, load in another thread's register, and continue. It's not cycle-by-cycle interleaving, it's “*run until we hit a brick wall, then switch*”. So Coarse-Grain makes sense in machines where:

- Stalls are very long (hundreds of cycles).
- Switching overhead (drain/refill of pipeline) is small compared to that stall.
- The hardware issue width is not very wide (so we don't need multiple threads every cycle to fill slots).

Finally, Coarse-Grain MT exists in theory and historically (some early machines, some GPUs in specific modes), but **most commercial CPUs with multithreading today use fine or medium approaches**, often in the form of **SMT** (simultaneous multithreading, which is essentially “very fine” grain because multiple threads issue in the same cycle).

- **Coarse-grain MT**: simple, only switches on long stalls.
- **Fine-grain MT**: switch every cycle, great for latency hiding.
- **Medium-grain MT**: umbrella term for designs between the two extremes (switch every few cycles or on moderate stalls).

⌚ Why can't we simply use the Operating System Context Switch?

When the **operating system** switches between two threads (or processes), it has to:

- Save the entire register file, program counter, and status of the old thread.
- Restore the state of the new thread.
- Update memory mappings, caches, etc.

This is called a **context switch**, and it usually takes **thousands of CPU cycles**. That's fine for multitasking at the OS level, but it's far too slow if our goal is to hide a cache miss (hundreds of cycles).

⌚ **Hardware Multithreading does it differently.** In hardware multithreading, the **processor itself** holds the state of multiple threads *at the same time*.

- Each thread has its own copy of architectural state: its own **program counter** and **register file**.
- These copies are kept in hardware, ready to use.

So when the processor wants to change which thread is executing:

- ✗ It doesn't have to call the OS or save/restore memory.
- ✓ It just selects another PC + register file (already in hardware).

This takes a **few cycles or even just one cycle**, depending on the type of multithreading (coarse-grain vs fine-grain vs SMT).

⌚ **Why “fine- to medium-grain” threads?** Because the hardware switch is so fast, the unit of work a thread can represent is very small:

- We might switch threads every **few instructions** (medium-grain).
- Or even **every cycle** (fine-grain).

This is in contrast to OS threads, where switching is so slow that we want to run a thread for millions of instructions before switching.

⌚ Why Threads Matter for Hardware Multithreading

A superscalar or VLIW core is designed to extract ILP **within one thread**. But if that thread has stalls, or if it simply doesn't provide enough independent instructions, the processor wastes resources. If we have **multiple threads ready**, then the processor can:

- Switch to another thread to hide latency (coarse- or fine-grain multithreading).
- Or issue instructions from multiple threads in parallel (simultaneous multithreading, SMT).

That's why the thread abstraction is the **basic currency of multithreading at the hardware level**.

7.3 Multithreading Basics

⚠ The Problem to Solve

Modern superscalar/VLIW processors are designed with **multiple functional units (FUs)**: integer ALUs, FP units, load/store units, branch units, etc.

⌚ **Goal:** execute several instructions in parallel each cycle.

⚠ **Reality:** a single thread often **cannot provide enough ready instructions** because of:

- Cache misses → hundreds of cycles stall.
- Branch mispredictions → pipeline flush.
- Data dependencies → must wait for results.

👎 **Result:** many FUs stay **idle**, then wasted silicon, wasted power.

✅ The Goal of Hardware Multithreading

Multithreading's central idea is simple: **when one thread is stalled, another can run, keeping the hardware busy**. Instead of *widening issue width* (which hits ILP limits), we **interleave multiple threads** on the same datapath. Benefits:

- ✓ **Hide stalls** (memory latency, branch penalties).
- ✓ **Improve utilization** of functional units.
- ✓ **Increase throughput** (more instructions completed per cycle across threads).

The improvements are more visible when compared to an architecture without multithreading:

✗ Without Multithreading

- One thread flows through the pipeline.
- If it stalls (say, waiting for L2 miss), the entire pipeline bubbles.

✓ With Multithreading

- Several threads have state “loaded” in the CPU.
- Each cycle, the front-end can pick an instruction from whichever thread is ready.
- If Thread A stalls, Thread B (or C, D, etc.) keeps the pipeline fed.

⌚ Throughput vs Latency

With Multithreading, **Throughput** (instructions per cycle overall) increases because the pipeline rarely idles. Instead, **Latency** (time for a single thread's instruction to finish) may increase, since our thread sometimes has to wait its turn behind others. So:

- Multithreading is **throughput-oriented** (maximize utilization).
- Not necessarily **latency-oriented** (our thread alone may not get faster).

This is why multithreading is great for **server workloads** (many independent threads), but less useful if we only run a **single-threaded application**. In other words, multithreading make the **core as busy as possible overall**, even if individual threads slow down.

7.4 Hardware Support for Multithreading

⌚ What must be duplicated?

A thread, as we defined earlier, is basically:

- A **Program Counter (PC)** (to know where it is in the instruction stream).
- A **Register File (RF)** with all architectural registers (general-purpose, FP, control/status).
- Status bits like condition codes.

To support multiple threads simultaneously, the CPU must keep **a separate copy of this state for each hardware thread**. For example: if one core supports 4 hardware threads, it will contain 4 PCs and 4 register files. These register files don't all sit in a single bank, but usually they are **physically replicated** so the **CPU can “switch” instantly** by pointing to a different set.

⌚ What can be shared?

The **expensive datapath and memory hierarchy** are **not duplicated**:

- ⌚ **Function Units (FUs)**: ALUs, FPUs, Load/Store units.
- ⌚ **Caches (L1, L2, L3)**: shared among threads of the core.
- ⌚ **TLB & Memory System**: usually shared, sometimes partitioned logically.

⌚ **Why?** Because duplicating those would make multithreading as expensive as adding more cores. The point is to keep utilization high *without* replicating all the silicon.

☑ Why a Thread is a smart idea

Adding a whole **new core costs lots of area and power** (extra datapath, extra L1 caches, extra pipeline logic). Adding a **hardware thread** only costs extra register files + PC logic. This is much cheaper in silicon area, but still allows the processor to stay busy by **switching between threads when one is waiting**.

⌚ **What does it achieve?** It does **not double performance**, since threads still share execution resources. But it can improve throughput significantly, especially when:

- One thread stalls on memory → the other fills in.
- Threads use different resources (e.g., one integer-heavy, one floating-point-heavy).

It's a **clever way to parallelize** at low cost: instead of building another core, let one core host multiple *active threads*.

➡ Thread Switching

Thread Switching happens when the CPU stops executing one thread and starts executing another. The main **goal** is to **switch to another thread quickly enough to hide stalls**.

❓ How?

- Hardware keeps all thread contexts “live”.
- Switching is just select another Program Counter and Register File in the next cycle.
- No need to save/restore to memory like the OS does

⌚ Latency

- **Fine-Grained MT** (page 373): can switch **every cycle**.
- **Coarse-Grained MT** (page 374): switches on stalls → costs a few cycles for pipeline drain/refill, but much cheaper than OS context switches.
- **Simultaneous MT** (page 381): no “switch”, but multiple threads issue in the same cycle, so contexts are interleaved continuously.

⚠ Resource Sharing Challenges

Since threads share FUs and caches, conflicts can occur:

- Two threads want the ALU in the same cycle → one must wait.
- Threads can evict each other’s cache lines (cache thrashing).

Simultaneous MT especially must use **smart scheduling policies** to prevent one thread from starving others. We will review each of these scenarios.

7.5 Typology of Hardware Multithreading

Hardware Multithreading is a technique used in modern processors to improve resource utilization and overall throughput by **allowing multiple threads to be executed concurrently within a single core**.

This section introduces the main types of hardware multithreading:

- **Coarse-Grained Multithreading (CGMT)**. The processor executes one thread until it hits a **long-latency stall** (e.g., cache miss to DRAM). Then it **switches to another thread**. The switch requires **draining/refilling the pipeline**, which costs a few cycles, but that's negligible compared to hundreds of cycles of stall.

✓ **Advantages**

- ✓ Very simple to implement (just duplicate PC + registers).
- ✓ Effective at hiding *long* stalls (memory, I/O).

✗ **Disadvantages**

- ✗ Useless against *short* stalls (like data hazards or branch mispredicts).
- ✗ Pipeline bubbles still occur frequently.

Example 1: CGMT Analogy

Imagine one cashier per supermarket checkout. When a cashier has to wait for a price check (long stall), they temporarily leave and another cashier steps in. But for small pauses (handing change), the switch isn't worth it.

- **Fine-Grained Multithreading (FGMT)**. The processor **switches thread every cycle** (usually round-robin). If one thread is stalled, it's skipped. Each thread makes slower progress (because it only issues one instruction every few cycles), but the **pipeline never sits idle**.

✓ **Advantages**

- ✓ Hides all kinds of stalls, even short ones.
- ✓ Keeps functional units busy almost every cycle.

✗ **Disadvantages**

- ✗ Increases **per-thread latency** (our program feels slower if it's the only one running).
- ✗ Requires very fast thread-selection logic.

Example 2: CGMT Analogy

Imagine a classroom where students (threads) answer questions in strict rotation. Even if one student hesitates, the next gets the turn immediately, so the teacher (pipeline) never waits.

- **Simultaneous Multithreading (SMT).** The processor is **wide superscalar** (can issue 4-8 instructions per cycle). But one thread alone rarely fills all slots (ILP limit). So the CPU can issue instructions from **different threads in the same cycle**. Combines ILP (multiple instructions per thread) and TLP (multiple threads).

 **Advantages**

- ✓ Maximizes **throughput per cycle**.
- ✓ Utilizes execution units that would otherwise be idle.

 **Disadvantages**

- ✗ Threads compete for shared resources (caches, ALUs, branch predictors).
- ✗ Per-thread performance can vary unpredictably (depends on what the sibling thread does).
- ✗ Hardware complexity is high (must decide *which thread's instructions* win each slot).

Example 3: CGMT Analogy

Think of a wide highway with 8 lanes (issue slots). A single car (thread) can't use them all. With multiple cars (threads), the lanes are filled more efficiently, but they can also block each other if not well managed.

The figure 29, page 382, makes clear the **design philosophy shift** we introduced in Section 7.3, page 376.

- In the **superscalar single-thread case**, the CPU tries to minimize *latency* for one thread, but many execution slots remain empty due to ILP limits.
- **Coarse-grain multithreading** improves throughput only when a thread is completely blocked by a long stall, but short stalls still waste slots.
- **Fine-grain multithreading** fills the pipeline continuously, maximizing **throughput** but stretching *per-thread latency* (since each thread issues less often).
- **SMT** goes further: by issuing multiple threads simultaneously, it combines ILP and TLP to achieve the **highest throughput**, at the cost of unpredictable latency per thread (threads interfere for shared resources).

Thus, the figure visually captures the **central idea of multithreading**: stop focusing exclusively on speeding up a *single* thread (ILP); instead, keep the core busy across *multiple* threads, so the processor's valuable resources are rarely idle.

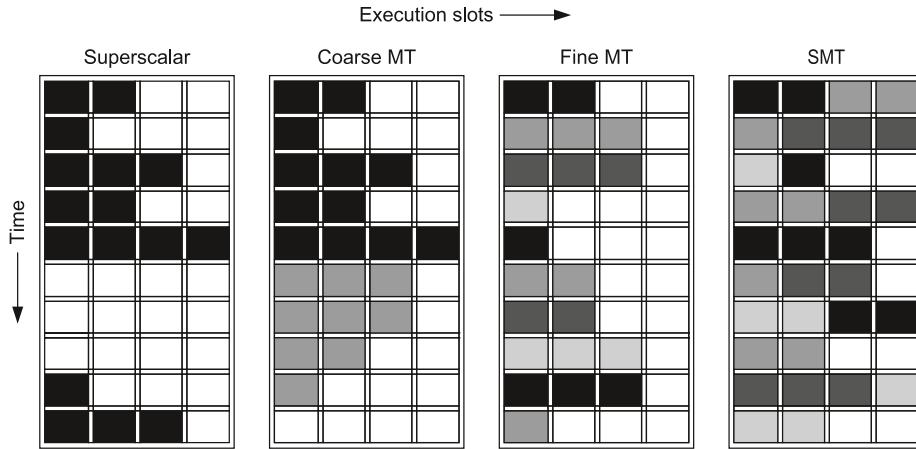


Figure 29: Execution Slot Utilization. [2] This figure shows how a superscalar pipeline’s issue slots are used under four scenarios:

1. **Superscalar without multithreading.** Only one thread. Many slots are **empty** (white) because the single thread suffers stalls: dependencies, cache misses, branch mispredicts. Even if the core is 4-issue wide, often only 1-2 instructions can issue per cycle.
2. **Coarse-Grained Multithreading (CGMT).** The figure shows long stretches of slots in one color (Thread A), then a big stall (white), then long stretches in another color (Thread B). Here, each thread runs until it encounters a long stall (e.g., a memory miss). Then the hardware switches to another thread. Empty slots still exist between switches, because the pipeline must “drain/refill”.
3. **Fine-Grained Multithreading (FGMT).** Threads alternate every cycle (different colors interleaved). This keeps slots filled continuously, even if one thread stalls. However, each thread progresses more slowly, since it only issues every N cycles (with N threads).
4. **Simultaneous Multithreading (SMT).** Multiple threads share the pipeline **within the same cycle**. Different colors appear **side-by-side in the same cycle row**, showing two (or more) threads issuing simultaneously. This achieves the **highest utilization** of slots: both ILP (within one thread) and TLP (across threads) contribute to filling the issue width.

In the superscalar case, black represents one reference thread, gray represents instructions from different threads, and white represents empty slots, i.e., unused pipeline capacity.

⌚ Summary

- **Coarse-Grain MT:** simple, hides *long* stalls only.
- **Fine-Grain MT:** frequent switching, hides *all* stalls, but per-thread latency grows.
- **SMT:** most advanced, fills superscalar issue slots with multiple threads, best throughput, but high complexity and contention.

7.6 SMT on Wide Superscalars

A superscalar is a CPU that can issue multiple instructions per cycle from a single thread. A wide superscalar is a very wide design, e.g., 4-issue, 6-issue, even 8-issue (e.g., AMD Zen). However, there is a problem: a single thread's ILP is limited, so many of those issue slots remain **unused** (the “white bubbles” we saw in the diagrams).

⚠ The Problem of Wide Superscalars

Designers wanted to make processors that can issue **4, 6, or even 8 instructions per cycle**. But a **single thread rarely has enough ILP** to keep all slots filled. Out-of-order execution helps, but it cannot magically invent independent instructions beyond the program's inherent limits. Result: even wide OoO cores show **empty issue slots** (low utilization).

✓ **Solution:** SMT (Simultaneous Multithreading) is a hardware technique in which **multiple threads of execution share the same core at the same time**. Instructions from different threads can be issued in the same cycle, side by side. With a **wide superscalar core** (e.g., 4-issue OoO), instead of depending on one thread to fill three slots, **SMT allows multiple threads to supply instructions in the same cycle**. So the SMT logic is added to the existing wide superscalar design.

💡 Why SMT Helps

Simultaneous Multithreading (SMT) allows multiple threads to **share the same cycle**:

- One cycle: Thread A issues 3 instructions, Thread B issues 1.
- Next cycle: Thread A is stalled, Thread B fills 4 slots.

SMT **combines ILP and TLP**:

- If a thread has enough ILP → it uses many slots.
- If not → other threads “fill in the gaps”.

This is why SMT fits naturally with wide superscalars: the more slots we have, the more room there is for multiple threads to share.

⌚ Typical Speedups

✓ **Throughput benefit:** a 2-way SMT core (two hardware threads per core) often achieves **20-30% more throughput** on multithreaded workloads. It is **not a 2× speedup** because threads still complete for shared caches and functional units.

⚠ **Single-thread performance:** can degrade slightly when SMT is active (competition for caches, branch predictor pollution). But most Operating Systems allow us to disable SMT per workload if latency is critical.

⌚ Energy and 📈 Area Trends

Adding SMT costs relatively **little extra hardware**:

- ➕ Duplicate register files, rename tables, PCs.
- ⌚ Execution units, caches, memory system are shared.

Energy per instruction **improves**: more work is done with the same datapath.
But cache contention can increase **miss rates**, sometimes hurting performance if workloads are not well-balanced.

7.7 Multicore + SMT

When you combine **multicore** (multiple physical cores) with **SMT** (multiple hardware threads per core):

$$\text{Total hardware threads} = (\# \text{ cores}) \times (\text{SMT threads per core}) \quad (81)$$

For example, Intel Core i7 with 8 cores and 2-way SMT, has **16 hardware threads**, and IBM POWER9 with 24 cores and 4-way SMT has **96 hardware threads**. This is what Task Manager (Windows) or `lscpu` (Linux) reports as **logical CPUs**.

Example 4: `lscpu`

For example (on my Thinkpad T430), when we run the `lscpu` command on a random Ubuntu computer, we get the following output:

```

1 andre@andre:~$ lscpu
2 Architecture:           x86_64
3 CPU op-mode(s):         32-bit, 64-bit
4 Address sizes:          36 bits physical, 48 bits virtual
5 Byte Order:              Little Endian
6 CPU(s):                 8
7 On-line CPU(s) list:    0-7
8 Vendor ID:              GenuineIntel
9 Model name:             Intel(R) Core(TM) i7-3632QM CPU @
                           2.20GHz
10 CPU family:             6
11 Model:                  58
12 Thread(s) per core:    2
13 Core(s) per socket:    4
14 Socket(s):              1
15 Stepping:                9
16 CPU(s) scaling MHz:   51%
17 CPU max MHz:            3200.0000
18 CPU min MHz:            1200.0000
19 ...

```

A socket is the physical connector on the motherboard where a CPU chip (package) is installed. Each socket can hold one CPU chip (the “package” with its silicon die(s) inside). For example, servers can have multiple sockets (e.g., 2-socket or 4-socket motherboards).

In this example, there is 1 physical CPU package with 4 physical cores. Each core supports 2 hardware threads via Intel Hyper-Threading technology (SMT).

So Linux scheduler sees 8 logical CPUs. The OS can schedule up to 8 software threads to run truly *simultaneously*. But remember: the extra thread per core (Hyper-Threading) is used because we are working with Intel architecture) doesn’t double performance. It just helps keep the core busy when one thread is stalled.

⌚ Why this matters

More hardware threads increase the number of OS threads that can execute **truly simultaneously**. But scaling threads per core vs scaling cores affects performance differently, depending on the workload:

- **More cores** helps when threads are **truly independent** and need their own compute capacity.
- **More SMT per core** helps when a single core has idle slots to fill, but not enough ILP from one thread.

⌚ Cache Sharing Policies

When multiple threads share a core, they inevitably also share caches:

- **L1 cache**: almost always **shared between SMT threads of the same core**.
 - ✓ **Pro**: keeps design simple, avoids duplication.
 - ✗ **Con**: threads can evict each other's cache lines (cache thrashing).
- **L2 cache**: depends on the architecture. In some CPUs it is **private** to each core (like in most Intel/AMD), in others it is **shared among a group of cores** (some ARM clusters, some older Intel Atom).
- **L3 cache**: often **shared across all cores** (chip-level cache).

A Implications: with SMT, **cache pressure increases**. Two (or more) active threads double the working set of the L1/L2. Designers sometimes add policies:

- **Partitioning** cache space between threads.
- **QoS (quality of service)** mechanisms to prevent one thread from starving the other.
- Replacement policies aware of thread ID.

Multicore scaling increases raw compute capacity. **SMT scaling** improves utilization of each core. Both combined give high *thread counts* on modern CPUs, but cache sharing makes resource management critical.

8 SIMD & Vector Architectures

8.1 Why DLP now? Limits of ILP

Throughout the '90s and '00s, we pursued better performance by extracting **Instruction-Level Parallelism (ILP)** through wider pipelines, out-of-order (OoO) scheduling, branch prediction, and speculation. That path **hit three walls**:

1. **Control wall (branches).** When control flow is irregular, speculation wastes work; deeper pipelines make each misprediction more expensive. **Even with good predictors, we still pay frequent bubbles.**
2. **Memory wall (cache misses + irregularity).** ILP engines require a large number of independent operands that are readily available. However, real code often involves pointer chases or unpredictable misses. OoO can hide some latency, but only if the instruction window “sees” other independent work. This is an **increasingly power-hungry proposition**.
3. **Complexity/energy wall.** The machines that find ILP (renaming, large ROB/RS, etc.) require more energy and area as the issue is widened or the windows are enlarged. The result is that **a lot of energy is spent to squeeze out a few more IPCs**.

We **can't keep relying on ILP**; we need other parallelism models, such as TLP, DLP, and RLP.

✓ The alternative: exploit Data-Level Parallelism (DLP)

Many kernels are **same operation, many elements** (images, linear algebra, DSP, ML). With **SIMD/vector**, one instruction drives many element operations in lockstep, *a single instruction stream controls multiple Processing Elements (PEs)*, so we amortize the control overhead and keep the datapaths busy.

Data-Level Parallelism (DLP) exists when the **same operation** can be applied **independently** to many data items at once. In other words, we don't need to wait for the result of one element before processing the next because all elements are parallelizable.

The classic example: adding two arrays element by element. Each sum $x[i] + y[i]$ is independent, so all iterations could, in principle, run at the same time. This is why we often say DLP is “*same instruction, multiple data*”.

An in-order vector processor can **match or surpass the performance of a complex out-of-order (OoO) core** on such workloads by relying on data locality (DLP) and achieving better energy efficiency.

❓ **Why the energy win?** Because **SIMD only needs to fetch one instruction per data operation**. We fetch/decode once; the work spreads out over many lanes and elements. That's less front-end activity, less rename/issue overhead, and more of our power going into real arithmetic. In other words, **one instruction controls many ALUs, so fetch/decode overhead is amortized**.

8.2 Flynn's taxonomy

Michael Flynn (1966) classified computers based on how many **instruction streams** and **data streams** they handle simultaneously:

- **SISD (Single Instruction, Single Data)**. One instruction stream, one data stream. It's a classic uniprocessor model with a scalar RISC pipeline, an in-order MIPS, or even an out-of-order (OoO) superscalar, because despite multiple instructions being issued per cycle, there's still one program counter and one stream of instructions.
- **SIMD (Single Instruction, Multiple Data)**. One instruction stream, many data streams. A control unit broadcasts the same instruction to many Processing Elements (PEs). Each PE has its own registers (and sometimes local memory) but all execute the *same* op on different data. Includes:
 - Classic vector processors (Cray, VMIPS model).
 - Packed-SIMD extensions in CPUs (MMX, SSE, AVX).
 - GPUs in their SIMD “warp” execution style.
- **MISD (Multiple Instruction, Single Data)**. Rarely used in real systems. Hypothetical model where multiple different instructions operate on the same data stream. Mostly academic (pipeline reliability systems sometimes cited).
- **MIMD (Multiple Instruction, Multiple Data)**. Multiple independent processors, each with its own PC, instruction stream, and data. This is the model of **multicore CPUs, clusters, distributed-memory systems**. Threads can be tightly or loosely coupled.

⌚ Where vector machines sit

Vectors are **one form of SIMD**. We issue **one vector instruction**; each element in the vector register is processed in lockstep across functional units. The programmer still sees **one PC, one instruction stream**, but the datapath replicates across elements. This makes vector machines the **canonical, textbook example of SIMD**.

⚠ Practical clarification

There's often confusion:

- **Superscalar CPUs (ILP)**: not SIMD in Flynn's sense, because although they issue multiple instructions, they all belong to **the same stream** (still SISD).
- **⌚ Why are superscalar CPUs SISD and not SIMD?** Flynn's classification is about *streams*, not width. Flynn cares about how many **independent instruction streams** and **independent data streams** the machine has. A superscalar CPU still fetches from a **single program counter (PC)**, meaning **one instruction stream**. Even if it can

issue 4, 6, or 8 instructions per cycle, those instructions all come from the *same stream*.

In SIMD, we fetch **one instruction** (ADDV V1, V2, V3), and that instruction **applies to many data elements at once** ($V1[i] = V2[i] + V3[i]$ for all i). In superscalar, we fetch multiple **distinct instructions**, each working on their own operands, not one instruction replicated over multiple data. So:

- SIMD is **parallelism across data elements (DLP)**.
- Superscalar is **parallelism across independent instructions (ILP)**.
- **Multithreaded/multicore (TLP)**: that's MIMD, not SIMD. Each thread has its own instruction stream.
- **SIMD/vector/GPU**: add data-level parallelism, one PC controlling many elements.

The taxonomy is coarse, but it helps us distinguish between OoO ILP and DLP SIMD. For us: **in Flynn's taxonomy, vector processors are SIMD**.

8.3 SIMD Architecture

In a **SIMD machine**, there is a **single control unit** (the “brain”) with **one Program Counter (PC)**. Each instruction is **fetched once** from this PC and then **broadcast** to all the **Processing Elements (PEs)**. Each PE executes that instruction **in lockstep**, but on its own **local data**.

Example 1: Analogy

Think of a classroom. One teacher (controller) gives instructions (“add 5 to your number”), and 32 students (PEs) each apply it to their own notebooks (registers). Everyone does the same thing at the same time, but on different numbers.

💡 Muscle of a SIMD machine: Processing Elements (PEs)

A **Processing Element (PE)** is the **replicated compute unit** in a SIMD array.

- Each PE executes the **same instruction** broadcast by the controller, but on its **own local data**.
- Together, many PEs form the “data-parallel fabric” of the machine.

So if the controller says “*add 5*”, every PE performs an addition in lockstep, but on its own input value.

❓ **What a PE typically contains.** A PE is usually quite simple compared to a full CPU core:

- **Registers:** its own small set of registers, holding local operands and results.
- **ALU/FPU:** an arithmetic unit (sometimes both integer and floating point).
- **Optional Local Memory:** in some SIMD arrays (like early Connection Machines), each PE had a small local memory slice. In vector-register architectures, the “vector register file” plays this role.

Contrast with a general-purpose core: a core has its own PC, branch predictor, caches, etc. A PE does **not**, it **relies on the central controller**.

❓ **How PEs execute.** All PEs step through the **same instruction** simultaneously. For example, suppose we have 4 PEs, each storing a number in a register:

- Controller issues: ADD R1, R2, R3
- PE0 does $R1_0 = R2_0 + R3_0$
- PE1 does $R1_1 = R2_1 + R3_1$
- PE2 does $R1_2 = R2_2 + R3_2$

- PE3 does $R1_3 = R2_3 + R3_3$

The **operation is identical**, but the **data are distinct**.

❷ Why one PC is powerful

- **Energy-efficient**: only one instruction fetch/decode per operation, instead of per PE.
- **Code compactness**: the program is written once, not duplicated per PE.
- **Simplicity**: no need for synchronization between PEs, since they all step in lockstep.

This is why SIMD machines historically delivered **high performance per watt** compared to pushing ILP in superscalars.

☒ Mixing SISD + SIMD

Most real systems are **hybrids**:

- A **scalar unit** executes normal SISD instructions.
- SIMD/vector instructions are executed by **PE arrays**.
- For example, in MIPS + VMIPS, we have both normal scalar registers and vector registers; in modern Intel, we have scalar x86 instructions plus AVX-512 SIMD ops.

So the system still looks like “one CPU” to the programmer, but when SIMD instructions are hit, the PEs all fire in parallel.

⤳ Comparison to MIMD

- **MIMD (multicores)**: each core has its own PC, they can all run different instructions on different data. Synchronization becomes our problem (locks, barriers).
- **SIMD**: one PC only, *no divergence*. Either all PEs do the same instruction, or some are “masked off” (we’ll cover masking later).

This explains why **GPUs**, though they manage thousands of threads, are essentially **SIMD under the hood**: groups of threads (warps) share a PC and execute in lockstep.

❷ Why vector processors are a clean SIMD implementation

Vector machines (like Cray, or our teaching model VMIPS) embody this model:

- The “one controller” issues **vector instructions**.

- The “many PEs” are abstracted as **lanes** inside vector functional units (VFUs).
- The “one PC” is the scalar control flow, but a vector instruction keeps PEs busy for many cycles.

This is why we say: **vector processors = classical SIMD**, just presented in a more programmer-friendly way (vector registers, vector instructions).

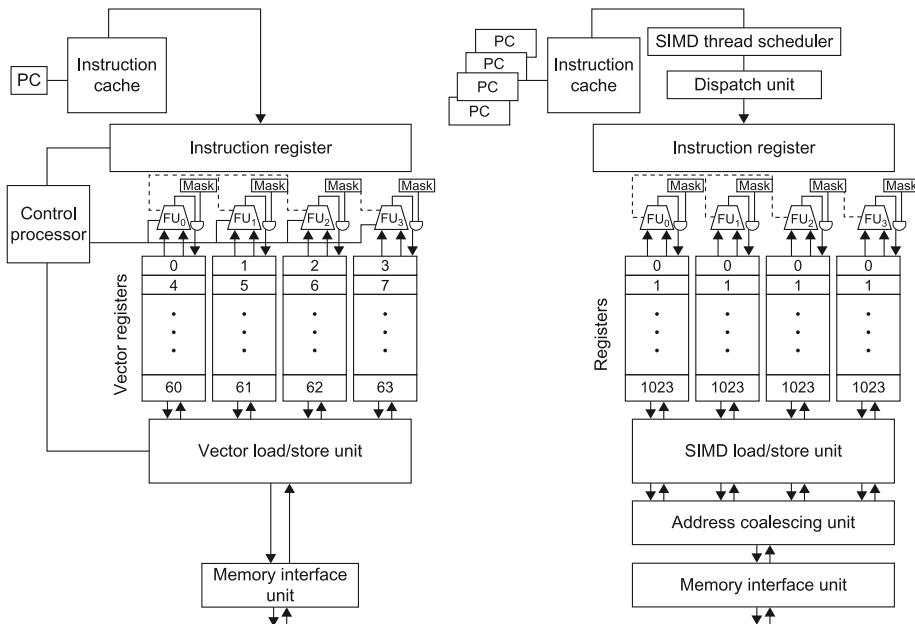


Figure 30: Vector Processor vs GPU SIMD Processor (side by side). [2]

This figure (30) is a comparison diagram and it places two architectures next to each other to show that **both are SIMD machines**, but they realize the model differently.

- **Vector Processor** (left side). This is the “one controller, many PEs, one PC” architecture that we have emphasized (this is what we are studying for the exam).
 - **Vector Register File (VRF)**: Large registers hold many elements (e.g., 64 floats).
 - **Vector Functional Units (VFUs)**: Operate in *lanes* (each lane = a PE). When we issue one vector instruction, the controller streams elements from the VRF through the VFUs.
 - **SIMD property**: One instruction (ADDV) controls many operations in parallel across lanes.
 - **Control**: Still a single PC driving both scalar and vector instructions.

- **GPU SIMD Processor** (right side). The GPU realization is **multi-threaded SIMD**: more flexible, hides memory latency, but still “one controller + many lanes per warp”.
 - **Registers per thread**: Each SIMD “lane” is associated with its own registers. Instead of a shared vector register file, each thread has its private state.
 - **Warp Scheduler**: The GPU groups threads into *warps* (e.g., 32 threads). Each warp has a single PC. The scheduler chooses which warp runs, hiding latency by switching between them.
 - **SIMD property**: Within a warp, one instruction is broadcast to all lanes (threads), still SIMD execution.
 - **Control**: Many warps = many PCs, but within each warp it’s pure SIMD.

Vectors and GPUs are two faces of SIMD. Vector machines present SIMD as “vector registers + vector instructions”, and GPUs present SIMD as “threads grouped into warps, each warp executing SIMD-style”. Both fit Flynn’s SIMD category because: **one instruction stream (per vector op or per warp)**; **multiple data elements processed in parallel by replicated datapaths (lanes/PEs)**.

8.4 Vector Architectures vs Generic SIMD

Not all SIMD is created equal. Vector processors are one way to implement SIMD and solve many problems that generic SIMD arrays face.

- **Generic SIMD (Flynn-style)**

≡ **Programming model:** The programmer sees many PEs arranged in an array. One instruction is broadcast to all. Data may live in each PE's local memory.

≣ **Pipeline behavior:** Each PE executes the operation independently, but control is centralized. Latency is exposed: if an operation takes multiple cycles, PEs may stall until it finishes.

⚠ **Limitations**

- ✗ Not very programmer-friendly (we must think in terms of PEs, local memory, communication).
- ✗ Latency of pipelines shows up as wasted cycles.
- ✗ Branching and irregular control flow are both awkward.

This is why such pure SIMD arrays (Illiad IV, Connection Machine) are more of **historical interest**.

- **Vector architectures (Cray-style, VMIPS model)**

≡ **Programming model:**

- * We write in terms of **vector registers** (V1, V2, V3), not individual PEs.
- * A vector instruction (`ADDV V1, V2, V3`) specifies an operation on *whole arrays of elements*.
- * The hardware hides the fact that many simple datapaths (lanes) do the work in parallel.

≣ **Pipeline behavior:**

- * Functional units are deeply pipelined. Once started, they produce one result per cycle.
- * **Element independence** (no cross-element dependency), easy to pipeline at high frequency.
- * Latency of the first result is amortized over the whole vector.

So instead of stalling for pipeline latency, the programmer only cares about **vector length** and throughput.

Difference between Generic SIMD and Vector architectures

The difference is in **how the machine executes each instruction internally**, and that makes generic SIMD slower in practice.

- **Generic SIMD (“array of PEs” model).** Imagine we have N PEs, each with its own ALU. The controller broadcasts: ADD R1, R2, R3. Each PE executes this instruction **exactly like a tiny CPU would**.

Now suppose the **ADD instruction has a latency of 6 cycles** (because of pipelining). In a scalar CPU, we can overlap new instructions while waiting. But in a simple SIMD array, all PEs must **wait 6 cycles** before they can do the next instruction, because each PE is “frozen” until the operation finishes.

Result: latency is exposed. Throughput is lower because every vector operation pays the full latency cost.

- **Vector architectures (Cray-style).** We don’t address individual PEs. We issue a **vector instruction**: ADDV V1, V2, V3. Inside the vector unit, the operation is **pipelined**:

- Cycle 1: first pair of elements enters the adder.
- Cycle 2: second pair enters, first is still computing.
- Cycle 3: third pair enters, first is nearly done.
- And so on.

After the latency of 6 cycles, we start getting **1 result per cycle**. If our vector has 64 elements:

- Startup cost: 6 cycles.
- Then 64 results flow out, one per cycle.
- Total = 70 cycles.
- Throughput ≈ 1 per cycle (after startup).

Result: latency is hidden; throughput is maximized.

The core difference:

- **Generic SIMD:** think of it as N *little scalar processors* all stalling on latency. Each vector op takes latency \times vector length.
- **Vector processor:** think of it as a *conveyor belt (pipeline)*. Latency is only paid once; then results stream out every cycle.

That’s why vector processors were a revolution compared to naïve SIMD arrays.

8.5 Anatomy of a Classic Vector Machine

When people say “vector processor”, they usually mean **Cray-style architecture**. The model we use in class (VMIPS) is a simplified version of this design. The following are the choices that designers, especially Seymour Cray and his descendants, made, and the **reasons why they work so well for DLP**.

Deepening: Honorable mention goes to Seymour Cray

Seymour Cray (1925 - 1996) was an American computer architect, often called the “**father of supercomputing**”.

- Worked at **Control Data Corporation (CDC)**, where he designed the **CDC 6600** (1964), considered the first successful supercomputer.
- Founded **Cray Research** in the 1970s, producing the **Cray-1** (1976), the first commercially successful **vector supercomputer**.

Cray-1 and vector processing. Introduced the idea of **vector instructions** with **vector registers** and **vector pipelines**, exactly the model we call “Cray-style SIMD”. His machine could perform arithmetic on entire arrays of data with unprecedented speed. Famous for the **C-shaped chassis** of the Cray-1 (actually a practical cooling solution).

Philosophy. Believed in **simplicity and regularity** over complexity. While mainstream microprocessors pursued **ILP and caches**, Cray bet on **bandwidth, pipelines, and vector registers** for scientific workloads. Preferred **low-latency, high-throughput memory** (banked DRAM) instead of large caches.

Legacy. The **Cray-1** defined the blueprint for vector machines (and later the VMIPS teaching model). Nearly every **vector ISA today** (RISC-V V, ARM SVE, Intel AVX-512) can trace its lineage back to Cray’s principles. He is remembered not just as an engineer, but as an icon of **supercomputing culture**.

In summary, Seymour Cray was the architect who turned the theoretical idea of SIMD into a **practical, commercial success**. His Cray-1 supercomputer (1976) pioneered **vector registers, pipelined vector units, and interleaved memory**; exactly the anatomy we’re studying. That’s why when we say “Cray-style vector machine”, we’re acknowledging his design as the reference model.

█ Choices designers made

1. **Register-to-Register Operations (Vector Registers).** Vector instructions read source operands from **Vector Register Files (VRFs)** and write results back into VRFs.

❓ Why (vs memory-memory designs)

- ✓ **Latency hiding:** VRFs buffer data so the pipeline can stream without stalling on memory.
- ✓ **Bandwidth efficiency:** Each element is loaded once into VRF, then reused for many ops (saves repeated memory fetches).
- ✓ **Pipeline friendliness:** Functional Units (VFUs) consume registers element by element in order, sustaining throughput.

2. **Deeply Pipelined Vector Functional Units (VFUs).** VFUs (add, multiply, divide, etc.) are **long pipelines** that accept a new element each cycle. Latency is non-zero, but throughput is what matters: one result per cycle once the pipeline is full. **Thanks to element independence**, we can **pipeline very deeply without hazards**. This is what makes vector processors extremely fast on long loops: startup cost amortized across many elements.

3. Interleaved Main Memory (Banked Memory)

⚠ **Problem:** we need to fetch/store many consecutive elements per cycle to keep pipelines busy. A single memory bank cannot deliver one word every cycle.

✓ **Solution:** Split memory into **multiple banks**, each can serve one request per cycle. Addresses are spread (“interleaved”) across banks. For example, with 16 banks, consecutive words map to different banks, then we can sustain 1 word per cycle. If stride matches bank count (e.g., every access hits bank 0), conflicts happen; this is why “prime number of banks” or “address hashing” is used.

Finally, our memory provides bandwidth rather than low latency (i.e., cache-like hiding).

4. **Little or No Data Cache**, because vector memory access patterns are highly predictable (sequential, strided). Caches add overhead but bring little benefit for streaming. Interleaved DRAM banks provide enough throughput directly. However, there may be a scalar cache for scalar instructions, but the vector side bypasses it. So our memory system is simple and provides high-bandwidth thanks to vector access.

5. Other Key Components

- **Vector Mask Registers:** allow predication (some elements disabled, used in conditionals).
- **Multiple lanes:** replicate VFUs to scale throughput (e.g., 4 lanes = 4 elements per cycle).
- **Scalar Unit:** executes scalar instructions, controls flow, handles scalar data.

6. **Why these choices make sense (historical + performance)**

- **Cray's philosophy:** don't waste transistors on speculation or caches; spend them on **wide datapaths and memory bandwidth**.
- **Target workloads:** scientific codes with large arrays (matrix ops, PDE solvers, simulations) is the perfect match for DLP.
- **Energy/performance:** much higher throughput per watt than OoO ILP, because no giant instruction windows, no speculative logic.

8.6 VMIPS as Didactic Reference Machine

Vector MIPS (VMIPS) is not a commercial computer, but a teaching model of a vector processor created by Hennessy & Patterson [2] to illustrate how vector architectures work.

It starts from **MIPS**, a clean RISC scalar ISA we already know, then it adds a **vector unit**: a Vector Register File (VRF), a Vector Functional Units (VFUs), and a Load/Store Unit (LSU) that fetches/stores vectors from memory. The **scalar MIPS core** remains present, handling sequential code, address calculations, and control flow. So VMIPS is like saying: “*what if our simple MIPS machine also had a Cray-style vector coprocessor?*”.

If we think of MIPS as the “teaching model” for scalar RISC pipelines, then **VMIPS is the teaching model for vector processors**.

⌚ Why VMIPS?

- **MIPS as a base:** they extended MIPS because it’s a **minimal, orthogonal RISC ISA**, already widely used in education. Students know MIPS from scalar pipelines, so adding vectors to it is natural.
- **Faithful to Cray design:** VMIPS adopts all the key Cray-style choices: vector register file, pipelined functional units, chaining, interleaved memory banks. So that reasoning about performance (like counting cycles for DAXPY) reflects reality.
- **Simplicity over full realism:** Real Cray machines had more complex instruction sets, memory systems, and register counts. VMIPS reduces this to the essentials so that students can do cycle-by-cycle performance calculations without drowning in detail.
- **Bridge to modern SIMD/GPU:** Once we understand VMIPS, it’s easier to see the connection to modern extensions (AVX, SVE, RISC-V V) and GPUs. It provides the conceptual foundation: vector registers, pipelined FUs, interleaved memory.

⌚ Three pillars of any vector machine

- **Vector Register File (VRF).** The **Vector Register File (VRF)** is a special register bank designed to hold **entire vectors of data** rather than single scalars.
 - In a scalar machine (like MIPS), each register holds **one word** (e.g., 32 bits).
 - In a vector machine, each **vector register** holds **many elements** (e.g., 64 elements × 64-bit each in VMIPS).

So, instead of saying “*register R1 contains the value 5*”, in a vector machine we say “*vector register V1 contains the 64 values of array X*”.

A **Vector Register File** is large, think of it as **8-32 registers**, each containing **64-128 elements**. It has **multiple read and write ports** (e.g., 16 read, 8 write in VMIPS) because many functional units (add, multiply, load/store) may want to read/write concurrently. Ports are connected via a **crossbar** to the Vector Functional Units (VFUs). However, the VRF’s role in execution is to **act as a compiler-controlled buffer between the memory and compute units** (elements are streamed out to the functional units one per cycle).

- **Vector Functional Units (VFUs).** A **Vector Functional Unit (V FU)** is a **pipelined arithmetic unit** (adder, multiplier, divider, logical unit, etc.) that operates on elements of a vector register **one after another, in a streaming fashion**. Think of them as the “datapath lanes” of the vector processor: they take operands from the VRF, process them in a pipeline, and return results into another VRF.

⌚ **Chaining.** One of the most important features is the **Chaining**: The output of one VFU can feed directly into another **without waiting for the full vector to finish**.

- **Vector Load/Store Unit (LSU).** The **Vector Load/Store Unit (LSU)** is the hardware that transfers whole vectors of data between **main memory** and the **Vector Register File (VRF)**.
 - Scalar load/store moves one word.
 - Vector load/store moves an **entire vector** (e.g., 64 elements in VMIPS) efficiently, in a pipelined, streaming fashion.

❖ How it works

- **Vector Load (LV):** fetches a sequence of words from memory into a vector register (LV V1, addr). Memory latency is paid **once**, then elements stream into the VRF one per cycle.
- **Vector Store (SV):** writes all elements of a vector register to memory (SV addr, V1). Again, throughput is 1 word per cycle.

The LSU sustains the **high bandwidth** required to keep vector pipelines busy.

The memory system is composed of **interleaved memory banks** because a single memory bank cannot deliver one word per cycle. Therefore, VMIPS assumes **multiple banks**, with each bank servicing one access per cycle. Addresses are interleaved across banks (e.g., word 0 goes to bank 0, word 1 goes to bank 1, and so on). This allows sequential vector loads and stores to map nicely across banks, achieving continuous bandwidth.

The **access modes** are smarter. There are two main modes of access:

- **Strided access** (for structured data): we can load every k -th element into a vector register (important for traversing rows/columns in matrices). For example: `LWVS V1, addr, stride=4` loads every 4th word.
- **Gather/Scatter** (for irregular patterns): With an **index vector**, we can load or store non-contiguous elements. Useful for sparse matrices or irregular data structures.

Everything is important because, without a high-bandwidth LSU, the vector pipelines (VFUs) would starve. Caches are also less useful here since vector codes have predictable, streaming patterns. This is why vector machines rely on banked memory and LSU instead of data caches.

- **Scalar Side.** The scalar side makes VMIPS (and real vector machines) **general-purpose**:

- If a program is **not** vectorizable, it still runs fine (just slower, scalar-only).
- If **part** of a program is vectorizable (like inner loops), the scalar processor sets it up and hands the heavy work to the vector unit.

This hybrid model is exactly what we see today: CPUs with SIMD units, GPUs with scalar “Schedulers” and SIMD “warps”.

We introduce VRF, VFUs, LSU, and Scalar Side in VMIPS because they are the three fundamental building blocks of a vector machine. Understanding them gives us a simple but accurate model of how vectors are executed, how latency is hidden, and how memory bandwidth is sustained. Exactly what we need to analyze performance and compare to other architectures.

8.7 Case Study: DAXPY

DAXPY (Double-precision A·X Plus Y) computes the operation:

$$Y[i] = a \cdot X[i] + Y[i] \quad (82)$$

For each element of vectors X and Y , with a a scalar constant. This is a **very common kernel in linear algebra**, part of the **BLAS (Basic Linear Algebra Subprograms)** library, widely used in scientific computing.

❷ Why is DAXPY a good case study?

1. **Simplicity.** It has a very regular structure: one multiply, one add, per element. It avoids complex dependencies, so it is perfect to isolate the effects of vectorization.
2. **Representative of real workloads.** DAXPY is a **building block in linear algebra and machine learning**. Appears in the **Linpack benchmark**, which is used to rank supercomputers.
3. **Highlights the difference between scalar and vector execution**
 - **Scalar MIPS version:** ≈ 578 instructions for 64 elements (plus stalls).
 - **Vector VMIPS version:** just 6 instructions for the whole loop!

Shows how **instruction count, loop overhead, and pipeline stalls** disappear with vectorization.
4. **Perfect to introduce performance metrics.** Lets us compare execution time in **clock cycles**:
 - Scalar: ≈ 600 cycles.
 - Vector (without chaining): ≈ 320 cycles.
 - Vector (with chaining): ≈ 192 cycles.

Later, we can extend to **multiple lanes, convoys, chimes**, etc.

The DAXPY case study is **didactic** because it connects the abstract concepts (VRF, VFUs, chaining, convoys) with a concrete **numerical example**. It shows **why SIMD/Vector machines are powerful** (fewer instructions, fewer stalls, higher throughput) and provides a **benchmark loop** to quantify speedups when new architectural features (chaining, multiple lanes, strip mining, predication) are introduced. Finally, it's the **entry point** to understanding how to reason about **vector performance models**.

8.8 Chaining, Convoys, and Chimes

⚠ The Problem

In scalar pipelines we saw that **data dependencies** (RAW hazards) cause **stalls** unless we can forward results. In vector processors, something similar happens: many vector instructions in a loop depend on one another. For example, in DAXPY, the ADDV depends on the result of MULV. So we must understand:

- *How do dependent vector instructions overlap?*
- *How do we count the execution time (cycles) in the presence of vector dependencies?*

▣ Convoys

A **Convoy** is the set of vector instructions that can **start in the same cycle**, because they use different functional units (FUs).

- Vector machines usually have multiple **pipelines** (for load/store, multiply, add, etc.).
- Instructions in the same convoy can all **start together**, each in its own pipeline.

However only **one vector instruction per FU** per convoy. Order must respect dependencies. Convoys is like **issue packets** in VLIW, but dynamic.

▣ Chimes

A **Chime** is one **round of execution of a convoy**.

- Think of it as: one “tick” in which all FUs execute one element of their vector instruction.
- The **number of chimes** is the **length of the vector instruction sequence** (in convoys), not the number of elements.
- For example, if we need 3 convoys to cover all instructions in a loop, then execution takes **3 chimes × vector length**.

Chimes is “time steps” of a vector loop execution.

▣ Chaining

Without **Chaining**, dependent vector instructions **must wait until the whole vector result is written back** before the consumer starts. With chaining: as soon as the **first element** of the producer is available, it can be forwarded to the consumer’s FU. This allows **overlap of dependent vector ops**, element by element. It’s the **vector analogue of scalar forwarding**. The result is fewer convoys and chimes and much lower latency.

Example 2: DAXPY revisited

Let's model $Y[i] = a \cdot X[i] + Y[i]$.

1. Instructions

- Load vector X .
- Load vector Y .
- Multiply vector ($a \cdot X$).
- Add vector (result + Y).
- Store result.

2. Convoys without chaining

- Convoy 1: Load X , Load Y .
- Convoy 2: Multiply.
- Convoy 3: Add.
- Convoy 4: Store.
- 4 convoys \rightarrow 4 chimes \times vector length.

3. Convoys with chaining

- Convoy 1: Load X , Load Y .
- Convoy 2: Multiply + Add (because chaining allows overlap).
- Convoy 3: Store.
- 3 convoys \rightarrow 3 chimes \times vector length.

8.9 Multiple Lanes

So far, in our vector model (VMIPS) we assumed:

- Each **vector functional unit (VFU)** handles **one element per cycle**.
- A vector of length n therefore takes n cycles to process (plus pipeline latency).

This is simple, but it **limits throughput**. If $n = 64$, a single vector add still requires ≈ 64 cycles, even if FUs are fully pipelined.

Q The idea of multiple lanes

To scale performance, we **replicate the hardware pipelines into lanes**.

- Each lane is an **independent datapath** with its own **slice of the vector register file** and functional unit(s).
- If we have L lanes, then L **elements can be processed per cycle**.
- Vector length n then takes $\lceil \frac{n}{L} \rceil$ cycles.

This is analogous to **SIMD width** in modern processors.

Example 3

Suppose:

- Vector length (VL) = 64
- We want to execute $Y[i] = a \cdot X[i] + Y[i]$

1. With 1 lane:

- Each element takes 1 cycle.
- Multiply: 64 cycles.
- Add: 64 cycles.
- Total dominated by vector length.

2. With 4 lane:

- Each cycle, 4 elements are processed in parallel.
- Multiply: 16 cycles.
- Add: 16 cycles.
- Execution shrinks by a factor of 4.

■ Architectural consequences

- **Vector Register File (VRF).** Must provide more ports: each lane needs **read and write access**. **Bandwidth grows with number of lanes.**
- **Load/Store Unit.** Needs to deliver multiple words per cycle. Often paired with **interleaved/banked memory** to sustain bandwidth.
- **Instruction semantics.** ISA does **not** change: the programmer still writes ADDV, MULV etc. The hardware executes them faster thanks to wider lanes.

This is why vector machines are said to scale “transparently” with hardware width, unlike scalar unrolling.

8.10 Handling Vector Lengths

Vector processors execute operations on **vectors of fixed hardware length**: the **Maximum Vector Length (MVL)**. For VMIPS, MVL is 64 elements. But in **real programs, the loop trip count n may not equal MVL**:

- Sometimes **shorter** (e.g., process only 31 elements).
- Sometimes **longer** (e.g., process 1000 elements).

We need mechanisms to **handle this mismatch** between **program data size** and **hardware register size**.

Maximum Vector Length (MVL)

Maximum Vector Length (MVL) is a hardware constant: the **maximum number of elements that a vector register can hold in a given vector processor**. It is the **physical width** of the vector register file. For example, in VMIPS, the MVL is 64 elements, so **each vector register can store 64 elements**, each 64 bits long.

There are several reasons why MVL is necessary. One of those reasons is the hardware boundary, because the vector register file (VRF) is finite in size. Thus, MVL specifies the limits of the hardware in order to balance cost, area, and memory bandwidth (**hardware upper bound**).

Vector Length Register (VLR)

The **Vector Length Register (VLR)** is a **special control register** in a vector processor. It specifies the **active number of elements** that vector instructions will operate on, **for the current loop segment**. From this definition, we can derive a logical constraint:

$$0 \leq \text{VLR} \leq \text{MVL} \quad (83)$$

In other words, the number of active elements must be less than or equal to the maximum number of elements that a vector register can hold, known as the MVL or hardware upper bound.

Classic vector machines handle “*less than MVL*” cases elegantly with VLR, thus avoiding a special scalar loop. This is why Cray-style vectors are considered more flexible than fixed-width SIMD. For instance, if an MVL is 64 and a loop has 100 elements, execution is divided into two steps. First, VLR is equal to 64, and the first 64 elements are run. Then, VLR is equal to 36, and the last 36 elements are run. In this example, the hardware still works with 64-entry registers, but only the first VLR entries are “active”.

Why is the Vector Length Register (VLR) really needed?

Suppose our machines has MVL equals to 64. Every vector instruction would always operate on all 64 elements of a vector register. Now imagine our program loop has 100 elements. Execution attempt:

1. First iteration: process elements 0-63. Ok.
2. Second iteration: Process elements 64-127. However, our loop only has 100 elements. Reading/writing beyond the array bounds would result in incorrect program behavior.

With VLR: In the second iteration, the VLR is set to 36 and only elements from 64 to 99 are processed. Unused entries are ignored. This eliminates the need for separate scalar loops, and the vector ISA works cleanly for any loop size.

Strip Mining

What if the **loop trip count n is larger than MVL**? We use **Strip Mining**, a **compiler technique** used to transform a loop with an arbitrary number of iterations n into **chunks of size MVL** (the hardware maximum vector length):

1. **Partition the loop** into chunks of size MVL.
2. Execute the body for each full chunk.
3. For the last remainder ($n \bmod (\text{MVL})$), set the VLR to that remainder size.

In other words:

$$\text{Loop of size } n \Rightarrow \left\lfloor \frac{n}{\text{MVL}} \right\rfloor \text{ full strips} + (n \bmod \text{MVL}) \text{ remainder strip} \quad (84)$$

For example $n = 1000$, $\text{MVL} = 64$:

- Do 15 iterations of 64 elements ($15 \times 64 = 960$).
- Last strip is 40 elements ($1000 \bmod 64 = 40$).

It is essential because **real program loops can be much longer than MVL**.

Together, they make vector machines **flexible** for any loop length, while maintaining **regular, efficient execution**.

8.11 Control Divergence Inside Loops

⚠ The problem: control dependence

Consider a loop with a conditional inside:

```

1 for (i = 0; i < 64; i++) {
2     if (X[i] != 0)
3         X[i] = X[i] - Y[i];
4 }
```

This loop is **data-parallel** (each iteration is independent), but the **if** makes it **non-uniform**:

- Some elements of X satisfy the condition, others don't.
- A naïve vector processor cannot execute this with a single vector instruction, because it would apply the operation to **all** elements indiscriminately.

✓ The solution: Vector Mask Registers (VMRs)

A **Vector Mask Register (VMR)** is a **Boolean vector** with one bit per element (length = MVL). Each bit says whether the corresponding element is “active” (1) or “disabled” (0). When enabled, **all vector instructions are executed under the mask**. Instead, inactive elements are skipped, leaving data unchanged. This is **predication at the vector level**.

For example:

1. Load X and Y into vector registers.
2. Compare X with zero and then produce a mask vector:
 - $\text{Mask}[i] = 1$ if $X[i] \neq 0$
 - $\text{Mask}[i] = 0$ otherwise
3. Apply subtraction under the mask:
 - If $\text{Mask}[i] = 1$, do $X[i] = X[i] - Y[i]$
 - If $\text{Mask}[i] = 0$, leave $X[i]$ unchanged

✓ Pros and ✗ Cons

✓ Pros

- ✓ **Keeps loops vectorizable**, even with conditionals.
- ✓ Avoids scalar fallback code.
- ✓ Masking is orthogonal: **any** vector instruction can be predicated.

✗ Cons

- ✗ Cycles are still spent for **all** elements, even masked ones (though inactive results are discarded).

- ✖ If many elements are masked off, the efficiency drops.
- ✖ Still, better than serializing into scalar execution.

In conclusion, **vector mark registers solve the problem of control divergence in vector loops**. These registers enable predicated vector execution, maintaining the effectiveness of vectorization even when loops contain `if` conditions. Conceptually, this solution can be written as follows: “perform this operation only for elements with a 1 in the mask bit”.

8.12 Memory System for Vectors

Vector machines can **consume and produce data at very high rates** (e.g., one element per cycle per lane). This puts enormous pressure on memory. To keep up, their memory subsystems were designed differently from scalar machines.

Memory Banks

Instead of a single wide memory, **vector processors use multiple memory banks** (interleaving). Each bank can provide one word per cycle. By spreading addresses across banks, the system can serve many concurrent loads/stores. Addresses are distributed across banks in a round-robin way:

$$\text{Bank number} = \frac{\text{Address}}{\text{WordSize}} \mod \# \text{ banks} \quad (85)$$

So ideally, when we load a vector, all elements land in different banks (full parallelism).

Q Why Memory Banks? Because if we need to fetch 64 elements of a vector, we don't want to wait 64 cycles serially. Instead, with 64 banks, each element can come from a different bank, so potentially all 64 delivered in parallel (after latency).

Stride (LVWS)

Access pattern in vector loads is often not contiguous. **Stride** is the spacing between consecutive vector elements in memory.

- Stride = 1, contiguous elements (row-major array).
- Stride = k , every k -th element (e.g., column of a matrix stored row-major).

The instruction used is: **LVWS V1, (R1, R2)**. This means to load the vector into V1, starting at the base address R1 and with a stride of R2. This allows vector processors to efficiently access **matrix columns** or other non-unit stride data.

Bank Conflicts

A **Bank Conflict** happens when **two or more consecutive elements of a vector map to the same bank** in the same cycle.

- ✖ Those banks cannot serve multiple requests simultaneously.
- ✖ Accesses to that bank must be **serialized**, introducing **stalls**.
- ✖ Effective throughput drops dramatically.

The number of distinct banks used per stride access is given by:

$$\frac{\# \text{ banks}}{\gcd(\text{stride}, \# \text{ banks})} \quad (86)$$

If stride and # banks are relatively prime (greatest common denominator = 1), we get full parallelism. If gcd is greater than 1, then the effective parallelism is reduced.

Example 4: Analogy

Imagine we have 8 supermarket checkout lanes (banks).

- If 8 shoppers each choose a different lane, then everyone checks out in parallel.
- If all 8 shoppers choose lane 0 (stride 8 case), then they must queue sequentially, wasting the other 7 lanes.

That queuing is bank conflict.

Bank conflicts occur when stride addressing causes multiple vector elements to hit the same memory bank in the same cycle. They **serialize memory access** and can reduce effective bandwidth from “one element per cycle per bank” to just “one element per cycle total”.

Example 5: Access Patterns and Bank Conflicts

Bank conflicts explain why accessing arrays along the dimension that is contiguous in memory (row-major rows, column-major columns) is better for performance, and why such loops are easier to parallelize.

Contiguous (stride = 1). If we access memory in **row-major order** (like C arrays), and we walk **across the row** (columns), each successive element is one word away in memory. This is stride = 1. With interleaved banks, each element lands in a different bank, so **no conflict, maximum throughput**. This is why compilers and programmers try to vectorize **inner loops over contiguous elements**.

Jumping by row (stride = row length). Now imagine we want to traverse **a column** of the matrix in row-major memory. Each new element is RowSize words apart. If RowSize is, say, 128 and we have 8 banks: $128 \bmod 8 = 0$. Then, every column element maps to the **same bank**. Result: massive conflict and serialized access. This is why **column accesses are “bad” in row-major layout**.

Fixes

1. **Change traversal order:** Access the array in row order (contiguous) instead of column order, if possible.

2. **Padding:** Add 1 dummy element at the end of each row (`RowSize` = 129 instead of 128). Now stride = 129. $\text{gcd}(129, 8 \text{ banks}) = 1$, then data spreads evenly and no conflict.
3. **Different layout:** Use **column-major** (like Fortran/Matlab) if our algorithm accesses columns more often. Or even more advanced formats (tiling, SoA vs AoS).

So, to avoid stalls, we usually want to access **consecutive addresses in memory** (stride = 1). But when our algorithm forces non-unit strides (e.g., columns in row-major storage), we may need **padding or different layout** to spread accesses across banks.

💡 Sparse Access: Gather and Scatter

Dense vectors are nice, but real applications often need **sparse matrices**. Elements are not stored contiguously; instead, we have an **index vector** that tells us where to fetch/store each element.

- **Gather:** Load a vector from non-contiguous memory locations given by an index vector.
- **Scatter:** Store a vector into non-contiguous locations given by an index vector.

This makes vector processors **flexible**: they can handle irregular data layouts too, though with more overhead.

8.13 SIMD ISA Extensions vs True Vectors

Starting from the mid-1990s, general-purpose CPUs (x86, ARM, PowerPC) added **SIMD instruction sets**: Intel MMX (1996), then SSE, AVX, AVX-512; ARM NEON. They let a scalar CPU operate on **short fixed-width vectors** packed into a register (128-, 256-, 512-bit). Each instruction applies the same op to several elements in parallel (e.g., 4 doubles in AVX2, 8 doubles in AVX-512).

Aspect	True Vector Machines (Cray, VMIPS)	SIMD ISA Extensions (SSE/AVX)
Vector length	Arbitrary, controlled by VLR , up to MVL.	Fixed by hardware register width (e.g., 4, 8, 16 elements).
Scalability	Same code runs efficiently if hardware MVL increases (e.g., 64 → 128 elements).	Code must be recompiled or rewritten when register width changes (e.g., SSE 128-bit vs AVX 256-bit).
Memory access	Rich modes: unit stride , non-unit stride (LVWS), gather/scatter .	Limited: mostly contiguous aligned loads/stores; gather/scatter only in AVX-512 (slower).
Masking / Predication	Supported via vector mask registers .	Only added in AVX-512 (k -mask registers); older SIMD had no clean predication.
Instruction count	One vector instruction covers VL operations, regardless of MVL.	One instruction covers only as many elements as fit in the register; multiple instructions needed for longer vectors.
Loop handling	Elegant with strip mining + VLR ⇒ same loop handles any size.	Leftovers require scalar “remainder loop” or separate vector cleanup.
Philosophy	Architecturally first-class vector model.	Scalar ISA “extended” with packed data ops

Table 30: Comparison between True Vector Machines and SIMD ISA Extensions.

✓ What we *gain* with SIMD extensions

- ✓ **Compatibility with scalar CPUs:** no need for separate vector hardware, just extra datapaths in the same ALU pipeline.
- ✓ **Widespread adoption:** present in all modern CPUs → compilers and libraries can target them.
- ✓ **Good for multimedia/data-parallel kernels:** audio, video, ML inner loops.

✗ What we *lose* compared to true vectors

- ✗ **Lack of flexibility:** fixed width means code isn't performance-portable across future wider registers.
- ✗ **More compiler complexity:** must generate special loops for remainders.
- ✗ **Limited memory addressing:** strided or sparse accesses are clumsy (or impossible pre-AVX-512).
- ✗ **Fewer instructions:** vector ISAs had elegant rich semantics; SIMD extensions are “bolt-ons” with awkward encodings.

9 GPGPU Computing

Since the GPU aspects have been thoroughly covered in the Parallel Computing course, we provide a link to the PDF and invite anyone to refer to the notes. The slides used in this course are the same as those used in the Parallel Computing course, so the notes are reliable.

Notes on Parallel Computing



10 Multiprocessors

10.1 What is a multiprocessor?

A **Multiprocessor** is a computer system with **tightly coupled processors**, coordinated by a **single operating system**, and typically offering a **shared memory abstraction** (shared address space). The idea is that **processors work in parallel** but **remain part of one coherent machine**, not a loose cluster. The execution paradigm is **MIMD (Multiple Instruction, Multiple Data)**: each processor/core runs its own instruction stream on its own data (page 388).

◎ Goals of Multiprocessors

Multiprocessors emerged to address the limits of uniprocessor performance:

1. **Performance**: Achieve *very high throughput* by exploiting thread-level parallelism.
2. **Scalability**: System performance should grow (ideally linearly) as more processors are added.
3. **Reliability**: Faults can sometimes be tolerated thanks to redundancy.
4. **Versatility**: Found in domains ranging from **embedded systems** to **high-end servers**.

≡ Taxonomy

We distinguish multiprocessors based on integration and scale:

- **Multicores (on-chip multiprocessors)**: Several cores on the same chip. Called **manycores** when the number of cores grows beyond 32. For example, Intel i7, AMD Ryzen, ARM bit.LITTLE.
- **SMP - Symmetric Multiprocessors (centralized shared memory)**: Typically ≤ 32 processors. All processors share a **single centralized memory** with **Uniform Memory Access (UMA)** time. Most existing multicore chips today are SMPs. Cache coherence is handled with **snooping protocols**.
- **Large-Scale Multiprocessors (Distributed Shared Memory)**: Needed when scaling beyond 32 processors. Memory is **physically distributed**, leading to **NUMA (Non-Uniform Memory Access)** latencies. Typically managed with **directory-based cache coherence**.

Multiprocessors are not just “more cores”. They are about how we **organize and coordinate processors, memory, and communication**. The shift from single-core \rightarrow multicore \rightarrow manycore reflects the industry response to ILP and frequency walls. The choice between **SMP vs. NUMA** reflects scalability limits: centralized memory works up to a point, beyond which distribution is necessary.

10.2 Key Design Questions

Designing a multiprocessor is a **multi-dimensional problem**: Hardware architects must balance **scalability** (network, memory placement, coherence protocols) with **programmability** (shared vs. message passing); Software (algorithms, compilers) and hardware (networks, caches) must co-design to extract parallelism and hide communication latency.

- ***How many processors?***

- Small-scale: 2-8 cores (single bus, simple cache snooping, UMA).
- Medium-scale: 8-32 cores (banked shared caches, interconnects like crossbars, NUMA).
- Large-scale: hundreds/thousands (NUMA, DSM, or message passing).
- Choice depends on target workload: servers, embedded, HPC, etc.;

- ***How powerful are processors?***

- Use few **complex/out-of-order cores** vs. many **simple/in-order cores** (as in IBM Blue Gene).
- Trade-off: single-thread performance vs. parallel throughput.
- Efficiency (performance per watt) is critical at large scales.

- ***How to connect processors?***

- **Bus**: cheap, simple, but scales poorly (saturation at ≈ 36 processors).
- **Interconnection networks**: rings, meshes, crossbars, hypercubes; each with cost/performance trade-offs.
- Network choice determines bandwidth, latency, and scalability.

- ***How do processors share data?***

- **Shared memory model**: implicit via loads/stores \rightarrow natural but coherence needed.
- **Message passing model**: explicit send/receive \rightarrow simpler hardware, but harder programming.
- Both are “Turing complete”¹¹ and can simulate each other. Even if they feel very different (implicit vs. explicit communication), each can be used to implement the other. So from a *theoretical* standpoint, neither is more “expressive”. The real difference is **practicality**: performance, programmability, and ease of reasoning. If a

¹¹A system (programming language, machine model, computational paradigm) is called **Turing complete** if it can compute **anything that is computable** in principle, i.e. anything that a *Turing machine* can compute. **Turing Machine** is an abstract mathematical model of computation, invented by Alan Turing, consisting of an infinite tape (memory), a head that reads/writes symbols, and a set of rules (program). It captures the essence of algorithmic computation. An example of turing complete are C, Python, Java. If a model can simulate a Turing machine, then it is “as powerful” as any other general-purpose computer; it can compute every computable function (given enough time and memory, **Turing Completeness**).

model is Turing complete, we know we can write any algorithm in it; So, choosing shared memory vs. message passing does not limit *what* can be computed, only *how*.

- ***Where to place physical memory?***

- **Centralized memory (UMA)**: uniform access, easy for programmers, but bottlenecked.
- **Distributed memory (NUMA)**: scalable, but non-uniform latency.
- Choice depends on processor count and workload locality.

- ***How do processors cooperate and coordinate?***

- Synchronization primitives (locks, barriers, atomic operations).
- Key bottleneck: synchronization costs can limit scalability.

- ***How to program processors?***

- Multiprogramming (many independent jobs).
- Shared memory (threads).
- Message passing (MPI processes).
- Data-parallel/SPMD (bulk synchronous phases).
- Programming model influences ease of use vs. scalability.

- ***How to maintain cache coherence?***

- Caches reduce contention and latency but introduce *coherence problems*.
- **Snooping protocols**: work on small-scale SMPs.
- **Directory protocols**: scalable solution for large NUMA systems.

- ***How to maintain memory consistency?***

- Defines the *order* in which writes become visible across processors.
- Sequential consistency (simple but slow) vs. relaxed models (complex but efficient).

- ***How to evaluate system performance?***

- Scalability: speedup vs. # processors (Amdahl's Law, Gustafson's Law).
- Latency of communication (local vs remote).
- Bandwidth (network, memory).
- Performance/watt, cost/performance ratio.

10.3 From single bus to interconnection networks

The **bus** was sufficient for early multiprocessors (workstation, small servers). For scaling (tens/hundreds of processors), the **interconnection network** is mandatory. The “sweet spot” explains why small SMPs (e.g., 4-16 cores) dominated desktops/servers, while large-scale HPC systems adopt meshes, tori, and fat trees.

SINGLE BUS Multiprocessors

In **Single Bus Multiprocessors**, all processors and memory modules are connected via a **single bus**. Therefore, every memory access uses the bus, and **all processors compete for the same shared medium**. The **main limitation is saturation**. Commercial systems maxed out at ≈ 36 processors before the bus saturated. Beyond that point, performance *flattens* because adding CPUs doesn’t help; the bus is the bottleneck.

NETWORK-CONNECTED Multiprocessors

In **Network-Connected Multiprocessors**, each processor has **its own local memory**. The **interconnection network** is used *only* for **inter-processor communication**, not for every single load/store. The main **advantage** is that it removes the shared bus from the critical path and can **scale to many more processors**. We will examine some of these topologies, such as rings, meshes, hypercubes, and crossbars.

COST/PERFORMANCE Trade-offs

- **Bus-based MPs**

- Cheap for small N (2-8 CPUs).
- Linear performance growth until the bus saturates.
- **Sweet spot:** 8-16 processors, after that, performance plateaus regardless of added CPUs.

- **Network-based MPs**

- Initially costlier (switches, links).
- Scale better: performance per processor stays roughly constant as N grows.
- Always win once the bus reaches its plateau.

10.4 Network performance metrics

When comparing interconnection networks (bus, ring, mesh, hypercube, etc.), we need **quantitative metrics** to evaluate scalability. The two canonical ones are:

1. **Total Network Bandwidth** (Best Case): the **sum of the bandwidths of all links** in the network.

$$\text{Total BW} = M \times b \quad (87)$$

Where:

- M is the number of links.
- b is the bandwidth of one link.

For example, the total bandwidth of a single bus is equal to the bandwidth of one link. Total Network Bandwidth represents the **maximum aggregate throughput that the network could deliver** if all links were used simultaneously and were perfectly balanced.

2. **Bisection Bandwidth** (Worst Case): the **minimum total bandwidth of links that must be cut** to split the network into two equal halves (each with half the nodes).

$$\text{Bisection BW} = \min \sum b_{\text{crossing links}} \quad (88)$$

It is the **worst-case metric** because, for **asymmetric** topologies, we have to **try all possible cuts and select the smallest one**. For example:

- **Single bus**: splitting cuts just the bus $\rightarrow b$.
- **Ring**: any cut crosses 2 links $\rightarrow 2b$.
- **Crossbar** (P nodes): $\left(\frac{P}{2}\right)^2 \times b$.
- **2D Mesh** ($P = N^2$ nodes): $N \times b$.
- **Hypercube** ($P = 2^N$): $2(N - 1) \times b$.

It measures *worst-case throughput* when half the processors communicate with the other half.

⚠ Best vs. Worst Case

- ✓ **Total bandwidth (best case)**: optimistic, assumes all links are equally used.
- ✗ **Bisection bandwidth (worst case)**: pessimistic, captures global contention.

In practice, a bus topology is sufficient for **small SMPs** (it is low cost, but has low bisection bandwidth). For **large-scale systems**, however, one must choose a topology with a higher bisection bandwidth, such as meshes, hypercubes, or fat trees, to avoid bottlenecks.

10.5 Classic Multiprocessor Topologies

A **topology** is the way processors and memories are **interconnected** in a multiprocessor. It defines:

- **How processors communicate** (direct neighbors vs. global links).
- **How scalable the system is** (bandwidth, latency).
- **How much it costs** (wiring, switches).

“Classic” refers to the **fundamental interconnection structures** studied since the 70s-90s and still forming the basis of today’s multicore **Networks-on-Chip (NoCs)** and HPC interconnects.

The five canonical interconnection structures used to classify and reason about multiprocessor networks are as follows:

1. **Single Bus:** All processors connected to one shared bus (one link).

- **Total Bandwidth**

$$BW_{\text{total}} = b \quad (89)$$

- **Bisection Bandwidth**

$$BW_{\text{bisect}} = b \quad (90)$$

For example, with 4 processors and $b = 1$ GB/s, both total and bisection bandwidth is 1 GB/s.

- ✖ Very low scalability.
- ✖ Saturates quickly.

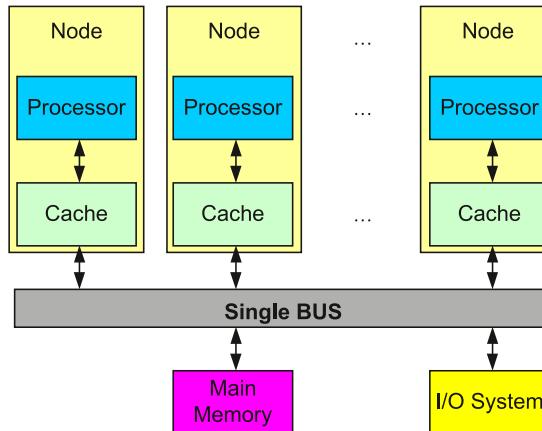


Figure 31: Single Bus Architecture. [4]

2. **Ring:** Each processor connected to two neighbors in a loop.

- **Total Bandwidth**

$$\text{BW}_{\text{total}} = P \cdot b \quad (91)$$

- **Bisection Bandwidth**

$$\text{BW}_{\text{bisect}} = 2b \quad (92)$$

For example 8-node ring, $b = 1$ GB/s, total BW = 8 GB/s, bisection BW = 2 GB/s.

- ✓ Good best-case parallelism.
- ✗ Worst-case is only $2 \times$ bus.

3. **Crossbar (Fully Connected Switch Network):** Each processor has a dedicated bidirectional link to every other.

- **Total Bandwidth**

$$\text{BW}_{\text{total}} = \frac{P(P - 1)}{2} \cdot b \quad (93)$$

- **Bisection Bandwidth**

$$\text{BW}_{\text{bisect}} = \left(\frac{P}{2}\right)^2 \cdot b \quad (94)$$

For example, $P = 4$, $b = 1$, total BW is 6 and bisection BW is 4.

- ✓ Non-blocking.
- ✓ High performance.
- ✗ Very high cost (scales P^2).

4. **2-D Mesh:** Processors arranged in an $\sqrt{P} \times \sqrt{P}$ grid.

- **Total Bandwidth**

$$\text{BW}_{\text{total}} = 2N(N - 1) \cdot b \quad \text{with } N = \sqrt{P} \quad (95)$$

- **Bisection Bandwidth**

$$\text{BW}_{\text{bisect}} = N \cdot b \quad (96)$$

For example, $P = 16$ ($N = 4$), $b = 1$, total BW is 24 and bisection BW is 4.

- ✓ Good scalability.
- ✓ Cost-effective.
- ✗ Multiple hops increase latency.

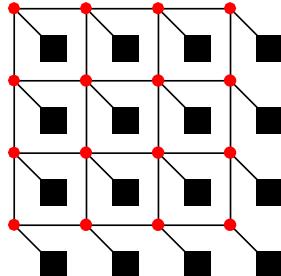


Figure 32: 2-D Mesh Architecture. [4]

5. **Hypercube (N-cube)**: Boolean N-cube, with $P = 2^N$ nodes. Each node connects to N neighbors.

- **Total Bandwidth**

$$\text{BW}_{\text{total}} = \frac{N \cdot P}{2} \cdot b \quad (97)$$

- **Bisection Bandwidth**

$$\text{BW}_{\text{bisect}} = 2(N - 1) \cdot b \quad (98)$$

For example:

- $P = 4$ ($N = 2$): total = $4b$, bisection = $2b$.
 - $P = 8$ ($N = 3$): total = $12b$, bisection = $4b$.
- ✓ Very good connectivity.
 - ✓ Logarithmic diameter.
 - ✗ Wiring complexity grows.

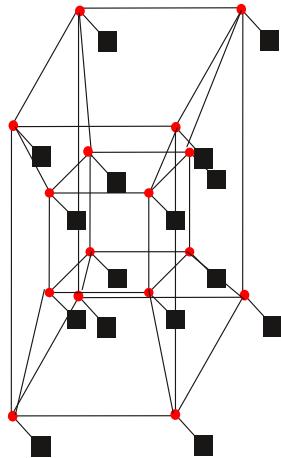


Figure 33: Hypercube Architecture. [4]

10.6 Address-Space Models

An **Address Space** is the set of memory addresses that a processor can use in its instructions. For example, when a CPU executes `LOAD R1, [0x1000]`, the value `0x1000` refers to an address in *its* address space. Different processors may or may not share the same address space.

Address-Space Models in Multiprocessors

In multiprocessors, the key question is: “do all processors share one global address space, or does each processor have its own private one?”

- **Shared Address Space (Shared Memory Model).** All processors see **one single logical memory**. Any processor can load/store from any address. A variable `X` at address `0x2000` is the same for all processors. So communication is **implicit**. If CPU A writes `X = 5`, CPU B can later just read `X`, no explicit “send” is required. Programs look similar to uniprocessor programs, just with multiple threads. Example of programming models: Pthreads, OpenMP.

Advantages

- ✓ **Ease of programming:** programmers just use variables and loads/stores.
- ✓ **Dynamic & irregular patterns:** works well when communication structure is unpredictable.
- ✓ **Lower latency for small items:** implicit communication avoids message overhead.
- ✓ **Compiler friendliness:** compilers can optimize with familiar memory model.
- ✓ **Hardware caching:** caches can automatically reduce remote traffic.

Disadvantages

- ✗ **Hardware complexity:** building scalable, coherent shared memory is *hard*.
- ✗ **Synchronization overhead:** locks/barriers needed.
- ✗ **Data placement control:** programmer has little control over where shared data is cached, which may hurt NUMA performance.

- **Private Address Spaces (Message Passing Model).** Each processor has its **own separate memory space**. Processor A’s `0x2000` is not the same as Processor B’s `0x2000`. No direct loads/stores across processors. So communication must be **explicit**. To share data, CPU A must **send a message** to CPU B. For example: `MPI_Send()` and `MPI_Recv()` in MPI. Example of programming models: MPI, Erlang actors.

Advantages

- ✓ **Simple hardware:** no need for cache coherence protocols.

- ✓ **Explicit communication:** programmer knows *when* and *how much* data is transferred.
- ✓ **Natural synchronization:** every send/receive acts as synchronization.
- ✓ **Data placement control:** programmer decides what data goes where.
- ✓ **Fault isolation:** easier in clusters (looser coupling).

✖ **Disadvantages**

- ✖ **Harder programming model:** must structure program around explicit messages.
- ✖ **Higher latency for fine-grain sharing:** sending small items has overhead.
- ✖ **More effort:** programmer bears responsibility for decomposing and distributing data.

The chosen model matters because it affects how **easy it is to program** and how **difficult it is to build the hardware**:

- Shared memory: easier for programmers, harder for hardware (needs cache coherence).
- Message passing: easier hardware, harder for programmers (must structure program explicitly).

Both models can express any algorithm (they're Turing complete), but the **trade-offs in latency, control, and scalability** are different.

10.7 Physical Memory Organization

This is about **where memory is physically placed** in a multiprocessor system, and how access latency behaves.

- **Centralized Memory: UMA (Uniform Memory Access).** **UMA (Uniform Memory Access)** refers to multiprocessor systems where **all** processors share a single centralized physical memory with the **same access latency and bandwidth**, regardless of which processor initiates the access. These systems are also called **Symmetric Multiprocessors (SMPs)** or **Centralized Shared-Memory Multiprocessors**.

❖ Architectural Properties

- **Centralized memory:** Memory is physically located in one place (often a set of banks connected by a shared bus or crossbar). Each processor connects to this memory through an interconnect.
- **Uniform latency:** Any processor accessing any memory location experiences the same access time (hence **uniform**). This makes programming easier: developers don't need to care whether data is "local" or "remote".
- **Cache hierarchy:** Usually, each core has private L1/L2 caches. Sometimes, a shared L3 cache sits before the main memory (especially in multicore chips). Cache coherence (future topics) mechanisms (snooping protocols) are required, since multiple processors can access the same memory.

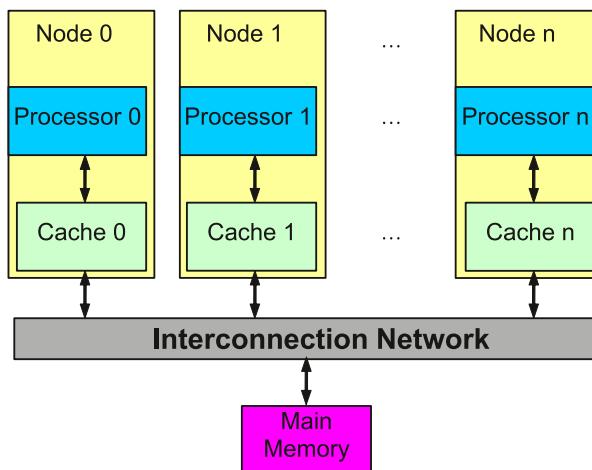


Figure 34: The **UMA system** was presented during the course. This should be familiar because it could happen on the exam. [4]

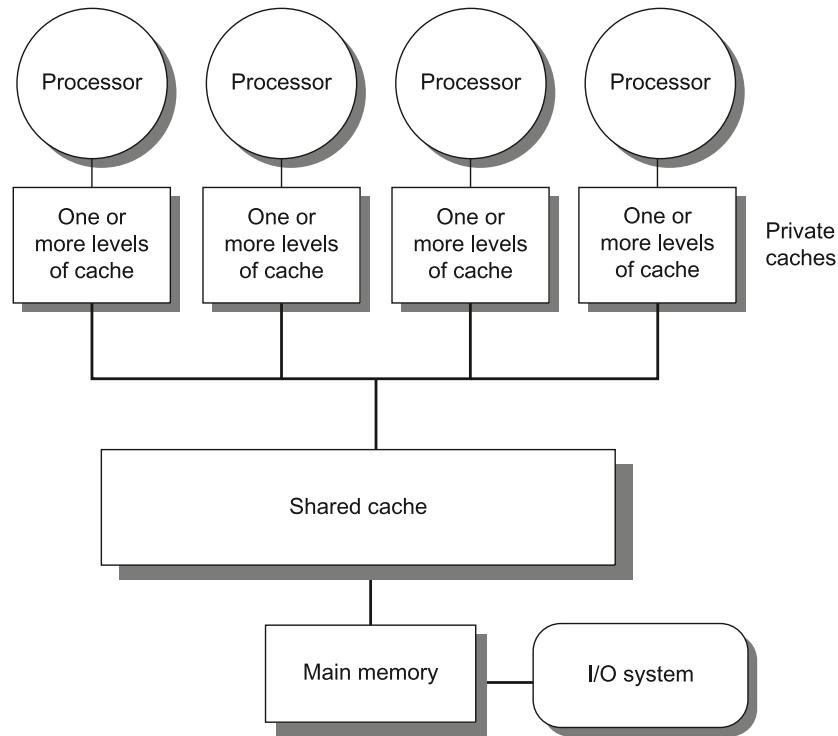


Figure 35: This is the **UMA system** picture taken from the Hennessy & Patterson book. [2] In our opinion, it is more realistic because it shows processors and private caches, sometimes with a shared last-level cache (L3) before memory. It is slightly busier, but it captures the modern SMP organization.

✓ UMA Advantages

- ✓ **Simplicity in programming:** no need to optimize data placement across nodes.
- ✓ **Shared-memory model:** natural for many applications, easy to port sequential code.
- ✓ **Lower software complexity:** the OS and programmer can assume a uniform view of memory.

✗ UMA Limitations

- ✗ **Scalability bottleneck:** as the number of processors increases, the single shared memory (and interconnect bus) can't provide enough bandwidth.
- ✗ **Contention:** multiple processors trying to access memory simultaneously can saturate the interconnect.
- ✗ Typically limited to ≈ 32 processors in classic SMP designs.

- **Distributed Memory: NUMA (Non-Uniform Memory Access).**

NUMA (Non-Uniform Memory Access) is a multiprocessor organization where memory is physically distributed across nodes, but logically forms a single shared address space (DSM, Distributed Shared Memory). Access time to memory depends on whether the data resides in the **local memory** of the processor's node or in a **remote node**. In other words:

- **Local Memory Access:** fast (low latency).
- **Remote Memory Access:** slower (high latency).

❖ Architectural Properties

- **Distributed physical memory:** Each node:

node = Processor(s) + private caches + local memory module

All nodes are connected through an **interconnection network** (mesh, torus, crossbar, etc.). The union of all local memories is one global shared address space.

- **Non-uniform latency:** If P0 accesses data in MM0 (its own memory), latency is low. If P0 accesses data in MM3 (remote memory), it must cross the interconnect, so latency is higher. Hence, *non-uniform* memory access times.
- **Cache coherence:** Since it is still a shared-memory system, caches must remain coherent. Centralized snooping does not scale, so NUMA systems typically use **directory-based protocols** to track which caches hold copies of which memory blocks.

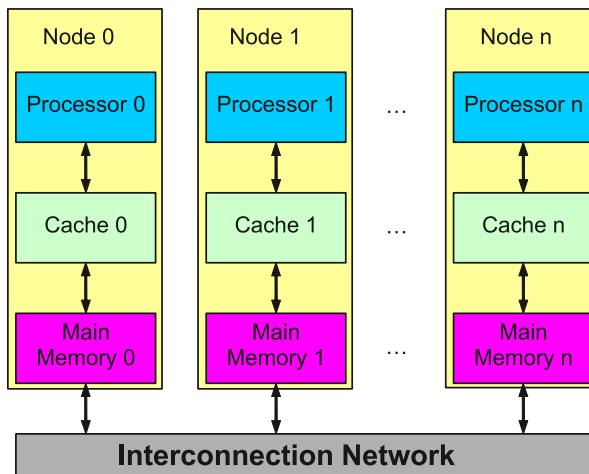


Figure 36: The **NUMA system** was presented during the course. This should be familiar because it could happen on the exam. [4]

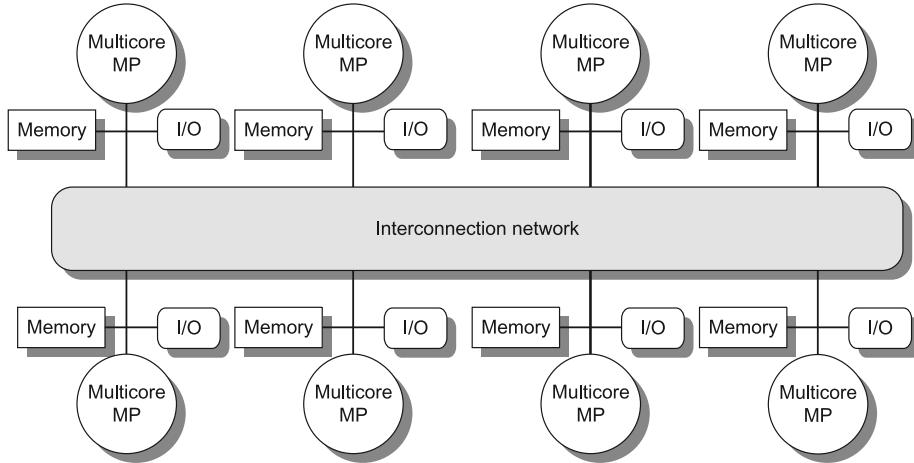


Figure 37: This is the **NUMA system** picture taken from the Hennessy & Patterson book. [2]

✓ NUMA Advantages

- ✓ **Scalability:** avoids the single bottleneck of centralized UMA memory. Systems can scale to dozens or hundreds of cores/nodes.
- ✓ **High bandwidth:** each memory module provides additional aggregate bandwidth.

✗ NUMA Limitations

- ✗ **Complex programming:** performance depends on data placement. If data is not local, performance may degrade.
- ✗ **More complex hardware:** requires scalable interconnection and directory-based cache coherence.

❑ Orthogonality with Address Space

Address-space model (shared vs. private) and **physical memory organization** (centralized vs. distributed) are **orthogonal dimensions**. We can have:

- Shared address space + Centralized physical memory (UMA). Classic SMP.
- Shared address space + Distributed physical memory (NUMA). Distributed Shared Memory (DSM, cc-NUMA).
- Private address spaces + Centralized memory. Rare, but possible in small message-passing systems.
- Private address spaces + Distributed memory. Clusters, large-scale MPI machines.

Orthogonal means that the two concepts are **independent dimensions**. We can **combine them freely**: choosing one does not force the other. The two “axes” are:

1. Address-space model

- *Shared address space*: all processors see the same addresses (e.g., UMA, DSM).
- *Private address spaces*: each processor has its own local space (message passing).

2. Physical memory organization

- *Centralized memory*: one big memory block (UMA).
- *Distributed memory*: divided among nodes (NUMA, clusters).

10.8 Small-scale multicores

Small-scale multicore processors are multicore CPUs with a **relatively small number of cores** (typically up to 8, sometimes 16). They implement a **Symmetric Multiprocessor (SMP) model on a single chip**:

- All cores are identical (symmetric).
- All cores have **uniform access** to a centralized main memory (UMA).
- The OS and software see them as a shared-memory machine.

Most commercial multicore processors today, including those from Intel, AMD, and ARM, as well as System on a Chip (SoC) processors, are Small-scale Symmetric Multiprocessing (SMP) systems and have these features:

- Less than 8 cores (typical laptops, desktops).
- Centralized shared memory (UMA model).
- One shared L3 cache and private L1 and L2 per core.

Each processor core is identical and symmetric: any core can execute any thread, and all have equal access to memory.

Cache hierarchy organization

- **Private L1 and L2 caches.** Very small, per-core and ultra-low latency. Store local working sets and reduce contention¹².
- **Shared L3 cache.** A larger, slower cache accessible by all cores. Reduces memory traffic to DRAM and acts as a buffer for inter-core communication: if Core A writes to shared data, Core B can see it via L3 (with coherence enforced).

Cache coherence: Snooping. Since multiple caches may hold copies of the same memory block, coherence is needed. One common protocol, which will be discussed later, is called the snooping protocol. **Snooping protocols** are typically used for small-scale SMP:

- Each cache controller “snoops” (monitors) the **shared bus or interconnect**.
- When one core issues a read/write, others observe it and update/invalidate their cache lines accordingly.
- Works well for **small numbers of cores (≤ 8)**, where bus-based broadcasting is affordable.

¹²**Contention** means that **two or more processors (or components) try to use the same resource at the same time**. Since the resource is finite (bus, memory bank, cache port, interconnect link), they cannot all proceed simultaneously, but some must **wait**. This waiting increases **latency** and reduces **throughput**.

⌚ Scaling beyond small-scale: banked caches & NoCs

As the number of cores grows (> 8 , toward manycores), a **single shared L3** or memory controller becomes a bottleneck. To scale, designers use **banked caches**:

- L3 is partitioned into several **banks**, each connected to different parts of the chip.
- Improves **aggregate bandwidth**, since multiple banks can serve requests in parallel.

The interconnect evolves into a **Network-on-Chip (NoC)**. Instead of a single bus, cores and cache banks are linked by mesh, ring, or crossbar topologies. NoC allows scalable communication without saturating a single bus. These types of protocols will be covered in future sections.

10.9 Distributed Shared Memory

Distributed Shared Memory (DSM) is a multiprocessor architecture where **physical memory** is **distributed** among nodes (each node has its own local memory), but the system provides a **global shared address space**, so any processor can address any memory location transparently. This is also called **Cache-Coherent NUMA (CC-NUMA)**:

- “Distributed”: physical location of memory.
- “Shared”: logical single address space for programming.

⚠ UMA Problem

In **UMA SMP**, all processors are connected to a **single centralized memory**. One memory controller, one set of DRAMs, and access latency is the same for everyone. The main problem is that it doesn’t scale. Having too many processors causes too much contention on one memory system.

✅ DSM / NUMA idea

To scale, designers spread memory **physically**. Each node has:

- **Processor(s)** (with caches).
- **Local memory module** (part of the global memory).

All nodes are linked by an **interconnection network**. So instead of a single huge memory, the total memory is **partitioned** and attached to nodes.

Now, if memory is physically distributed, how do we still let programmers think it is *one shared memory*? That’s where DSM gives the **illusion**. Using a **global shared address space**, any core can issue a load/store to any address. The system (hardware + OS) automatically decides:

- If the address belongs to its **local memory** (fast).
- If the address belongs to a **remote node**, so the request travels over the interconnect to that node’s memory (slow).

This is why it’s called **Distributed Shared Memory** (*distributed* physically, *shared* logically).

Example 1: Distributed Shared Memory Analogy

Imagine a **university library system**:

- Every faculty (node) has its **own local library** (memory).
- All libraries together form the **global catalog** (shared address space).

- A student (processor) can request *any* book (address).
 - If the book is in their faculty's library, **local access** (fast).
 - If the book is in another faculty, request goes over the campus shuttle (interconnect) to fetch it, **remote access** (slower).

To ensure everyone has the latest version of a book (cache coherence), the **central catalog** (directory) tracks who borrowed which copy.

Properties

1. Non-uniform memory access (NUMA):

- Access to **local memory** (same node) is fast.
- Access to **remote memory** (different node) is slower.
- Latency depends on the distance over the interconnect.

2. Single global address space:

- Programmers can use **load/store instructions** without worrying about where memory physically resides.
- The hardware/OS ensures that references are routed correctly.

3. Scalability:

- Because memory is distributed, the system avoids the bottleneck of a single centralized memory controller (as in UMA SMPs).
- Larger multiprocessors (8-32 cores or 100s of cores in modern servers) rely on DSM.

⇒ Why Directory-based protocols?

In small SMPs (≤ 8 cores), coherence is handled by **snooping**: every cache "listens" to a shared bus. But in DSM/NUMA there is **no single shared bus**, instead, there is a scalable interconnection network. Also, **broadcast snooping would not scale** (too much traffic, too many cores). The solution is **directory-based protocols**.

- A **directory** keeps track of which caches have a copy of each memory block.
- Stored alongside each memory block (in the local memory module).
- When a block is read or written, the directory is updated, and only the relevant caches are notified.

This avoids global broadcasts.

⌚ Role of the interconnect

The **interconnection network** ties everything together. Responsibilities:

- **Data movement:** transfer cache lines between nodes.
- **Directory lookups:** route coherence requests to the memory module that “owns” a block.
- **Scalability:** provide enough bandwidth/low latency for dozens of nodes.

Topologies used: Crossbar (small systems), 2D mesh, torus, fat-tree, and often implemented as a Network-on-Chip (NoC) for chip multiprocessors.

Key Takeaways: Important Concepts of Distributed Shared Memory Chapter

In this box, we will address some common questions about this topic. In particular, what does it mean that the memory is shared physically but not logically? Why and how can programmers not see this difference?

❓ What does *global shared address space* mean?

Imagine in a single-processor system: memory is a **linear array of addresses** (0x0000 ... up to max). In a multiprocessor with **DSM**, we want to **preserve this illusion**:

- Every processor sees **the same set of addresses**.
- Address 0x1000 always refers to the *same logical location* for all processors, no matter which node owns it physically.

In DSM, the hardware and OS cooperate so that any load/store instruction can be applied to *any address* in that global space.

❓ How is this built if memory is distributed?

Physical memory is **partitioned across nodes**:

- Node 0 “owns” addresses 0x0000 → 0x0FFF.
- Node 1 “owns” addresses 0x1000 → 0x1FFF.
- Node 2 “owns” addresses 0x2000 → 0x2FFF.
- ... and so on.

Each node’s memory controller is responsible for its **portion** of the global space. If a processor in Node 0 accesses address 0x2200:

1. Its cache misses → request goes to the interconnect.
2. The interconnect routes the request to **Node 2**, because Node 2 owns the range 0x2000-0x2FFF.
3. Node 2’s memory module returns the data.

So it's **transparent to the programmer** because they don't care where data is physically located.

❷ Why should a processor ever access *remote* memory?

Because **parallel programs share data**. Some examples:

- **Shared data structures:** Suppose Node 0 initializes an array. Later, Node 1 must process part of it. Even if that part resides in Node 0's memory, Node 1 can access it via DSM.
- **Synchronization variables:** Locks, semaphores, and barriers must be visible to all processors. A lock variable in Node 3's memory must be accessible by Node 0 when it tries to acquire the lock.
- **Work distribution:** In scientific simulations, the global domain (e.g., a 3D grid of cells) is spread across nodes. But some computations require reading/writing data from neighbors in another node.

If no processor could access remote memory, it wouldn't be "shared memory" anymore; it would be pure **message passing**.

❸ Why not just use private memory (clusters, next section)?

In message-passing clusters, the programmer must **manually orchestrate communication** (e.g., MPI):

- Node 0 does `send(array, Node 1)`.
- Node 1 does `receive(array, Node 0)`.

In DSM/CC-NUMA, the **hardware handles this automatically**. A simple load/store suffices. This makes DSM **much easier to program**, at the cost of more complex hardware.

In conclusion, DSM provides a global shared address space by mapping each physical memory partition into one big address map. Any processor can issue a load/store to any address and the system routes the request to the right node. This allows transparent sharing of data structures and synchronization variables, without explicit message passing.

10.10 Distributed-Memory Clusters

A **Cluster** is a set of independent computers (called nodes), each with:

- Its **own processors** (often multicore CPUs).
- Its **own private memory** (not shared with others).
- Local disk and OS instance.

Nodes are connected via a **scalable interconnection network** (Ethernet, InfiniBand, high-speed fabrics). Unlike DSM/CC-NUMA, here there is **no shared address space**. Each node has its **own private address space**, and memory of one node cannot be directly accessed by another node via loads/stores.

❓ How can two nodes communicate?

Communication between nodes happens via **explicit message passing**:

- A process in Node A calls `send(message, Node B)`.
- A process in Node B calls `receive(message, Node A)`.

The software support is MPI (Message Passing Interface), the de facto standard in HPC.

Deepening: MPI

MPI (Message Passing Interface) is a **standardized library interface** for communication in distributed-memory systems. It defines how **processes exchange messages** (data) explicitly, so they can work together on a parallel program. It is not a language, but a **specification + libraries** available in C, C++, Fortran, Python bindings, etc.

❓ Why MPI exists

In **distributed-memory clusters**, processes do **not share memory**. Each process only sees its **own private address space**. To cooperate, processes must **send and receive messages** across the interconnect. MPI standardizes this, so parallel programs are **portable** across machines and vendors.

❓ When to use distributed-memory clusters

- **Scalability:** Clusters can grow from a few nodes to tens of thousands of nodes. Adding memory is easy, each node brings its own DRAM.
- **Commodity hardware:** Built from off-the-shelf PCs/servers connected with networking gear. Cheaper than custom supercomputers.
- **HPC workloads:** Large-scale scientific computing (climate simulation, astrophysics, molecular dynamics). Data analytics, AI training at scale.

- **Cloud datacenters:** The dominant architecture, each server is one node, applications distribute data across them.

✓ Pros and ✗ Cons

✓ Advantages

- ✓ Easy to scale: just add more nodes.
- ✓ Memory grows linearly with the number of nodes.
- ✓ Fault isolation: one node can fail without crashing the whole system.
- ✓ Cheaper and easier to build with commodity hardware.

✗ Disadvantages

- **Programming is harder:** No shared memory; programmer must explicitly partition data and orchestrate communication. Think in terms of “who owns the data” and “when to send it”.
- **Communication costs are high:** Network latency is much larger than local memory access. Bandwidth between nodes is lower than bandwidth to local DRAM.
- **Synchronization overhead:** Every send/receive pair involves coordination.

10.11 Communication/Programming Models

When we program for parallel machines, we need a **model** that defines:

- ***How processors communicate:***
 - Do they exchange data via a common shared memory?
 - Or do they send explicit messages?
 - Or do they all run the same program on different pieces of data?
- ***How programmers express parallelism:***
 - What primitives and abstractions they use (threads, send/recv, barriers, etc.).

In other words, a **Communication/Programming Model** is the **software-level abstraction** that maps onto the underlying **hardware organization**.

➊ The simplest model: Shared-Memory

The **Shared-Memory Model** is the simplest one, processors communicate via **shared variables** in a common address space, using loads/stores + synchronization.

➌ Hardware

- Multiple processors/cores all access the **same global address space** (UMA or NUMA).
- Any processor can read/write any memory location.

➍ Programming model

- Processes/threads communicate **implicitly** by accessing shared variables.
- Example: if Thread 0 writes $x = 42$, Thread 1 can later read x from memory.
- Synchronization is done via **locks, semaphores, barriers, atomic operations** to avoid race conditions.

➎ Frameworks

- **OpenMP** (compiler pragmas for parallel loops).
- POSIX threads (**pthreads**).
- Java/C# threads.

➏ Advantages

- ✓ Natural and easy for programmers used to sequential programming.
- ✓ No need to explicitly send messages, just share variables.
- ✓ Good for small- to medium-scale multiprocessors (SMPs, multicores).

✖ Disadvantages

- ✖ **Scalability bottleneck:** cache coherence traffic grows with core count.
- ✖ **Data placement issues:** in NUMA systems, memory might be local or remote → programmer must sometimes optimize placement for performance.
- ✖ Synchronization overhead can reduce parallel efficiency.

▀ The Message-Passing Model

In the **Message-Passing Model**, processes communicate via **explicit messages** rather than shared memory. It scales extremely well, but shifts responsibility to the programmer. Used in **distributed-memory clusters** with MPI as the standard tool.

▀ Hardware

- Each processor (or node) has its **own private memory**.
- There is **no global address space**, so one process cannot directly load/store data from another process's memory.
- Nodes are connected via a **network** (Ethernet, InfiniBand, custom HPC interconnect).

This is the model used in **distributed-memory clusters**.

leftrightarrow Programming side

- Communication is **explicit**: processes must use primitives like `send` and `receive`.
- A process cannot just “look” at another’s variables, it must request them.
- Typical flow:
 - * Process A: `MPI_Send(data, dest=B)`
 - * Process B: `MPI_Recv(data, source=A)`
- Larger groups use **collective operations**: broadcast, scatter/gather, reduce, barriers.

The most widely used library is **MPI (Message Passing Interface)**.

✓ Advantages

- ✓ **Scalability:** works for thousands of nodes, and no coherence problem since memories are private.
- ✓ **Data locality control:** programmer decides explicitly what data to send, when, and where. It avoids hidden memory-traffic surprises.
- ✓ **Fault isolation:** each node is independent; one node crash doesn’t kill the whole memory system.

✖ Disadvantages

- ✖ **More complex programming:** must carefully partition data among processes, and must explicitly coordinate all communication.
- ✖ **Communication overhead:** network latency and bandwidth are much worse than local memory.
- ✖ **Synchronization cost:** processes may need to wait for each other at send/receive points.

⌚ The Data-Parallel Model

The core idea of the **Data-Parallel Model** is **data-parallelism**: perform the **same operation** simultaneously on **different elements of a large dataset**. All processors execute the **same program**, but on **different portions of the data**. This is why it's also called **SPMD (Single Program, Multiple Data)**.

Example 2: Data-Parallel Model

If we want to add two vectors A and B of length 1000, each processor works on a different slice:

- Proc: indices 0-249
- Proc: indices 250-499
- Proc: indices 500-749
- Proc: indices 750-999

Each runs the **same addition loop**, just on its own chunk.

Data-parallel programs often follow **bul-synchronous phases**:

1. **Local computation:** each processor works on its local data independently.
2. **Global communication / synchronization:** processors exchange results or synchronize via a barrier.
3. Next phase begins.

This model was formalized by Leslie Valiant as the **BSP Model (Bulk-Synchronous Parallel)**.

Example 3: Hardware and Programming frameworks

Some hardware examples:

- **SIMD/vector machines:** each instruction applies to multiple data elements at once.

- **GPUs today:** CUDA / OpenCL launch thousands of threads that all run the same kernel on different data elements.
- **Multicore CPUs** with OpenMP pragmas like `#pragma omp parallel for`: each thread runs the same loop on different iterations.
- **Distributed systems:** MapReduce or Spark frameworks apply the same function across partitions of big datasets.

Some programming frameworks: CUDA, OpenCL, OpenMP, and MPI.

✓ Pros and ✗ Cons

✓ Advantages

- ✓ **Simplicity:** one program, replicated everywhere.
- ✓ **High scalability:** works well if the dataset is large and can be evenly split.
- ✓ **Hardware efficiency:** GPUs and vector units thrive on data-parallel workloads.

✗ Disadvantages

- ✗ **Load balancing:** if one processor gets more work than others, all must wait at the synchronization point.
- ✗ **Synchronization bottlenecks:** bulk barriers can slow down if stragglers exist.
- ✗ **Limited applicability:** best suited for regular, array-like computations (linear algebra, ML, scientific simulations). Harder for irregular data structures (e.g., trees, graphs).

10.12 Why caches matter

This is an introductory section, the details come later.

In modern CPUs, main memory is **too slow** relative to processor speed. Caches reduce **average memory access latency** and increase **effective bandwidth**. In multiprocessors, each core has **private caches** (L1, L2) and possibly a shared L3. Caches allow parallel performance, but also introduce **new problems** when data is shared across cores.

⚠ The cache problem in multiprocessors

When multiple processors **share memory**:

- Each may **cache copies** of the same memory block.
- If one processor updates its cached copy, others may see an **old value**.
- Without extra mechanisms, this breaks the **shared-memory programming model**.

This problem is called **Cache Coherence problem**.

⚠ Coherence vs. Consistency

In a multiprocessor, there are **two different correctness problems** when using caches:

1. *What value do we get when we read a single variable that multiple processors are updating?* This is the cache coherence problem (see above).
2. *In what order do we see writes to different variables from other processors?* This is the **Memory Consistency problem**.

They sound similar, but they control **different dimensions of correctness**.

- **Cache Coherence**: the **scope** is a single memory location. It ensures:
 1. **Write propagation**: all processors eventually see the most recent value of a variable.
 2. **Write serialization**: all processors see writes to that variable in the same order.

For example, P0 writes $X = 1$. Later, P1 reads X . **Coherence guarantees** that P1 eventually sees 1, not some stale old value.

- **Memory Consistency**: the **scope** is the ordering of operations across multiple memory locations. Defines *when* the result of a write by one processor becomes visible to others. Determines the rules for **ordering of loads and stores** across addresses. For example, P0 executes:

```

1 A = 1;
2 B = 2;
```

P1 executes:

```
1 print(B);
2 print(A);
```

But if P1 sees $B = 2$, must it also see $A = 1$? Some consistency models (like sequential consistency) say yes, instead weaker models (like relaxed consistency) may allow no.

So even if the system is **coherent**, without consistency rules we don't know *when* other processors see writes, or in what order.

Note that memory consistency is important; otherwise, programs could behave unpredictably due to instruction reordering, buffering, or weak memory ordering.

Protocol Families for Cache Coherence

When multiple processors have caches, we need a **protocol** to keep all copies of memory blocks consistent. Two big families exist, depending on the hardware organization:

- **Snooping Protocols.** Used in **small-scale SMP/UMA** with a **shared bus**. Each cache controller “snoops” (listens) to the bus for transactions (read, write). If another core writes to a block, snooping ensures invalidation or update of local copies. It scales poorly because broadcasting traffic grows with the number of cores.
- **Directory Protocols.** Used in **large-scale DMS/NUMA** systems with interconnect networks. A **directory** (kept alongside each memory block) tracks which caches hold copies. Only the relevant caches are notified on writes. Scales better than snooping (no broadcast).

Cache Coherence protocols: Invalidate vs. Update

The Invalidate and Update protocols are two **cache coherence protocols**. They ensure that **all processors see a consistent view of memory when caches hold copies of the same block**. These protocols **define what happens when one processor writes to a block that may exist in multiple caches**.

Despite being cache coherence protocols, the invalidate and update protocols define the **coherence policy of the caches**. The snooping and directory protocols, on the other hand, define the **hardware organization** (bus versus scalable network).

- **Invalidate Protocols.** When one processor writes to a block, all **other cached copies** are **invalidated**. Other processors no longer trust their copy; they must fetch a fresh copy from memory (or from the writer) if they need it later. For example, P0 and P1 both cache $X = 0$. P0 writes $X = 1$. Protocol sends “invalidate”, then P1's copy of X becomes *invalid*. If P1 later reads X , it must fetch the new value.

- **Update Protocols.** When one processor writes to a block, the **new value is immediately sent** to all caches that hold a copy. Everyone's copy is updated in place, so no invalidations. For example, P0 and P1 both cache $X = 0$. P0 writes $X = 1$. Protocol sends “update”, then P1’s cached copy is changed to 1. If P1 later reads X later, it already has the correct value.

10.13 Why Cache Coherence?

In a **shared-memory multiprocessor**, each processor has **private caches** to reduce latency and offload the memory bus. This works perfectly for **private data** (no other core cares). But when data are **shared** among processors, the same memory block may exist in **multiple caches**. That raises the **Cache Coherence problem**:

- Processors might see **different values** for the same memory location, depending on which cache they read from.
- For example, Core A writes $X=0 \rightarrow X=1$ in its cache, but Core B still has $X=0$ in its own cache. If B now reads, it gets the **stale value**.

So the central question: **How do we ensure that all processors have a coherent view of shared memory?**

⌚ What “coherence” means

A memory system is called **Coherent** if the following hold:

1. **Program order preserved locally**: if a processor writes X then reads X (without another processor intervening), it must read back its **own last write**.
2. **Read sees the latest write (globally)**: a read by processor P2 after processor P1 wrote to X should return P1's value; provided enough time has passed and no other writes to X happened in between.
3. **Write serialization**: all processors must see writes to the **same location** in the **same order**. If P1 writes X=1 and then P2 writes X=2, no core should ever observe the sequence 2 then later 1.

These three properties **define coherence**. The precise timing of when values become visible is part of **consistency models** (later topic). In summary:

- **Reads** must return the most recent value written to that address.
- **Writes** must be **serialized**, only one writer can “own” the data at a time; all other cached copies must be invalidated or updated.

✅ Two protocol families to fix it

To enforce coherence, we need a **protocol** that ensures caches are synchronized:

1. **Write-Invalidation protocol**. Before writing, a processor sends an **invalidate** to all other caches holding that block. From now on, the writer has **exclusive access**.
 - ✓ Only the **first write** to a block generates traffic; subsequent writes are silent.
 - ✓ Good for write-heavy workloads (exploits temporal locality).

2. **Write-Update protocol.** On each write, the new value is **pushed to all other caches** (and often memory).

- ✓ Readers see the new value immediately (lower latency).
- ✗ Every write consumes **bandwidth**.

Almost all modern multiprocessors use **write-back caches + write-invalidate protocols** (not update) to reduce bus traffic.

T	Processor A	Processor B	Memory (X)
0			1
1	Read X (1) → cache A=1		1
2	Cache: A=1	Read X (1) → cache B=1	1
3	Write X = 0 → cache A=0	Cache: B=1	0

Table 31: The **cache coherence problem** for a single memory location (X), read and written by two processors (A and B). [2]

T	Processor A	Processor B	Mem (X)
0			0
1	Read X → cache miss → A=0		0
2	Cache: A=0	Read X → cache miss → B=0	0
3	Write X = 1 → cache A=1	Cache invalidated A	0
4	Cache: A=1	Read X → cache miss → B=1	1

Table 32: An example of an **invalidation protocol** (write-invalidate) working on a snooping bus for a single cache block (X) with **write-back caches**. [2]

10.14 Snooping (bus-based) Protocols

10.14.1 Snoopy bus idea, broadcast, and scalability limits

In early multiprocessors, all CPUs and memory are connected via a **single bus**. Every **cache controller** “snoops” (listens) to the bus: if another processor issues a **read or write request** for a block that the cache holds, the snooper reacts appropriately (invalidate, supply data, etc.). In other words: caches **cooperate** by monitoring the bus traffic, so that all copies of a memory block can be kept coherent. This is why it’s called a **Snooping Protocol**, because each cache controller acts like a “spy” on the bus.

Broadcast mechanism

The bus is a **broadcast medium** because any request placed on it is seen by **all processors** simultaneously. For example, when the processor P1 writes to address X, the bus carries the request, and all other caches check if they have X, they **invalidate** or **update** depending on the protocol. This **broadcast simplicity** make snooping straightforward:

- **Only one bus transaction** is needed for all caches to know what’s happening.
- No separate “directory” or centralized record is required.

Scalability limits

Snooping is elegant but doesn’t scale well:

1. **Bus bandwidth bottleneck.** Every coherence transaction is broadcast to **all processors**. With more processors, bus traffic grows quickly, saturating available bandwidth.
2. **Electrical limits.** A single shared bus has **limited speed** and **fan-out**¹³ (signal must reach all caches). Adding more processors increases capacitance, then slower bus cycles.
3. **Serialization.** Bus provides a single global order (all requests are serialized). This is good for correctness but bad for performance since no parallel coherence actions are possible.

In practice, snooping works well for **small-scale simultaneously multiprocessors (SMPs)**, like 2-8 CPUs. For larger systems (tens or hundreds of cores), it becomes inefficient; that’s why **directory protocols** (covered later) are needed.

In summary, snooping leverages the **shared bus as a natural broadcast medium**. Each cache controller **snoops** to maintain coherence. It’s **simple and effective**, but **not scalable**: bus traffic and electrical limits make it suitable only for small multiprocessors.

¹³**Fan-out** is the **number of loads on the bus**. It is the **electrical loading** on the bus. In cache context, each cache controller (and memory module) connected to the shared bus is a **receiver**. The more devices we attach, the **greater the capacitance** and resistance that the driving signal must overcome. This increases the **propagation delay** of the bus: the time it takes for a transition at point A to be seen cleanly at the farthest point B.

10.14.2 Practicalities: duplicate tag to avoid stalling the core

In a **snooping system**, every bus transaction (read, write, invalidate, etc.) must be checked against the cache tags to see if the block is present. That means, whenever another processor issues a request, *our* cache must search its tag store.

⚠ Here comes the problem. The **cache tag array** is already heavily used by the **local processor** (to check hits/misses on every load/store). If we force the processor to stop whenever the bus needs a tag lookup, we create **conflicts**:

1. Either the processor stalls waiting for the snoop check;
2. Or the snoop is delayed (risking coherence violations).

✓ The solution: Duplicate tag (“snoop tag”)

Designers add a **duplicate copy** of the cache’s tag array, called **Snoop Tag Store**. It is a much smaller, parallel structure. However, this duplicate is used **only for coherence snooping**:

- Whenever a bus transaction occurs, the snoop tag is checked to see if this cache has the block.
- Meanwhile, the main tag array is still free to serve the local processor.

This **prevents** the processor pipeline from **stalling** due to bus snoop checks and **minimizes snooping overhead**.

❖ Implementation details

The **snoop tag array** doesn’t need data, only **address + state bits** (valid/dirty/ownership). For example, in an **MSI/MESI protocol** (covered later), each line of the snoop tag contains:

- The block’s tag;
- Its coherence state (M, E, S, I.)

On a bus transaction, if the tag matches and the line is not **Invalid**, the cache controller reacts (invalidate, flush data, etc.). This check can happen **every cycle** without disturbing the processor’s own cache access.

⚠ Trade-off

- ✗ Cost: extra silicon area for a second tag store (smaller than the full cache, since only tags are duplicated).
- ✓ Benefit: avoids **structural hazard** between processor accesses and snooping.

Without snoop tags, high-performance multiprocessors would suffer **pipeline stalls** just from coherence traffic.

In summary, snoopy coherence requires frequent tag lookups. To avoid interfering with the CPU's own cache access, systems use **duplicate “snoop tags”**. This is a small hardware trick, but **essential** for performance in bus-based coherence systems.

10.14.3 Write-invalidate vs write-update (bandwidth vs latency)

When a processor writes a new value into its cache, **other caches may also hold that block**. **How do we keep them consistent?** Two families of protocol exist:

- **Write-Invalidate protocol**

- Before performing the write, the processor sends an **invalidate message** on the bus.
- All other caches that hold the block mark it as **Invalid**.
- Now the writer has **exclusive ownership** of the block.
- Subsequent writes to the block are **local only** (no more bus traffic) until another processor requests the block.

Bus traffic is generated *only* on the **first write** to a block shared by multiple caches. Multiple consecutive writes are *silent* (great if a processor writes a lot to the same location).

⚠ Latency trade-off. A reader that later wants the block must fetch the updated version from the writer or from memory, which introduces *latency*.

✖ Usage. Dominant in practice (e.g., MSI/MESI) and good for **write-intensive workloads** (spatial/temporal locality exploited).

- **Write-Update protocol** (or **Write-Broadcast protocol**)

- On every write, the processor broadcasts the **new value** to all caches that hold the block.
- Each cache updates its local copy immediately.
- Memory is usually updated too.

Every single write produces **bus traffic**. Bad if one processor is writing many times to the same block.

⚠ Latency trade-off. readers never see stale data: if they access later, they already have the fresh value. Good for **read-mostly workloads** where a variable is written by one processor but read by many others.

✖ Usage. Rare in modern multiprocessors (too much bandwidth cost), but more common in small-scale systems or in some **vector machines / GPUs** where data are broadcasted naturally

In summary, **invalidate** protocol saves bandwidth, dominant in modern designs; instead, **update** protocol lowers read latency but consumes lots of bandwidth. The **invalidate vs update** decision is one of the core trade-offs in coherence protocol design.

10.14.4 MSI Protocol: Finite-State Machine (FSM)

MSI is a cache coherence protocol designed for snooping systems (bus-based multiprocessors). It specifies how each cache block behaves, depending on:

1. Its **current state** (Modified, Shared, Invalid);
2. A **local CPU request**;
3. Or a **bus request** from another processor.

Remark 1: Cache Block

Our cache (L1, L2, etc.) is divided into **blocks** (a.k.aa cache lines). The typical size is 32-128 bytes and each block holds the data from a contiguous memory region. For every block, the cache also keeps **metadata**:

- **Tag** (to know *which* address is cached).
- **Valid bit** (does this block hold useful data?).
- **Dirty bit** (has this block been modified and not written back?).
- And in multiprocessors: **Coherence state bits** (M, S, I).

Each cache block in our private cache is controlled by a **small Finite-State Machine (FSM)**¹⁴. Stimuli come from **two sources**:

- **The CPU attached to this cache** (local read, local write, called “*CPU events*”).
- **The snooping bus** (read miss, write miss, invalidate placed on the bus, called “*bus events*”).

The FSM tracks the coherence state of the block:

- **M - Modified.** The cache has the **only copy**, and it's **dirty** (different from memory). Must write back to memory before eviction.
- **S - Shared.** The cache has a **clean copy**, **identical to memory**. Other caches may also have it.
- **I - Invalid.** The cache does not have a valid copy of the block.

Each arrow in the diagram corresponds to:

stimulus → (possible bus action) → new state

¹⁴A **Finite-State Machine (FSM)** is a model of computation made of: (1) a **finite set of states** (e.g., M, S, I); (2) a **current state** (at any given moment, the system is in exactly one state); (3) a set of **event/inputs** (e.g., CPU read/write, bus read/write, etc.); (4) a set of **rules (transaction)** that say how to move from one state to another based on the current state and the input event.

Remark 2: Read/Write Hit vs Miss

- **Read Hit.** The processor requests a block that is **already in the cache**. Data returned immediately, no bus/memory access.
- **Read Miss.** The processor requests a block that is **not in the cache**. Cache issues a **bus read** to fetch the block from memory or another cache.
- **Write Hit.** The processor writes to a block **present in the cache**.
 - With **write-back**, only cache updated, state may change.
 - With **write-through**, cache and memory updated simultaneously.
- **Write Miss.** The processor writes to a block **not in the cache**. Cache issues a **bus write** (write miss) to fetch the block, possibly invalidating others, then performs the write.

💡 What triggers transitions?

- CPU requests
 - **Read hit:** no state change (stay in M or S).
 - **Write hit in M:** just update locally. Because we are in M, we are the only owner, so no bus action needed.
 - **Write hit in S:** must broadcast an **invalidate** so others drop their copies, then move to M.
 - **Read/Write miss:** place a request (read miss or write miss) on the bus. If we were in I, we move to S (read miss) or M (write miss). If we were in S, we remain in S.
- Snooped bus requests
 - If another core issues a **read miss** to a block we have in M → we must supply data and downgrade to S. Because we are supplying data, we must do a **write-back** so memory or requester has updated data.
 - If another core issues a **write miss/invalidate** to a block we have in S or M → we must invalidate (go to I).

Example 4: Transitions

Let's walk through a few “*how do we change state*” cases:

- **Invalid → Shared:** We get a CPU *read miss*. We fetch data from memory (or another cache) via the bus. End state = S.
- **Shared → Modified:** CPU does a *write hit*. Must place *invalidate* on bus so all other sharers drop their copies. Then upgrade to M.

At this point, we are the only owner of a dirty copy, all others are I.

- **Modified → Shared:** Another processor issues a *read miss* for that block. We supply the data (write-back), and we downgrade from M to S.

⌚ So, do we write back to memory?

It depends on the implementation of the snooping system, but the MSI definition requires that once a line is in **S (shared)** it is **clean** (identical to memory). So when we go from M (modified) → S (shared), the dirty copy in M must be “cleaned” somehow:

1. **Write-back to memory (simple but slower)^a.** On the *read miss* from another processor, we flush the block to memory. Memory is now up-to-date. Both us and the requester hold the line in S (shared).
2. **Cache-to-cache transfer (common in practice, less in theory).** Instead of updating memory, we *directly forward* the data to the requester. Both us and the requester enter S (shared).

⌚ **Buuu wait, memory is now stale. How can S mean “identical to memory”?** In real protocols, the bus guarantees that **the data supplied to the requester is the correct one**. After the transfer, the line is marked as “clean” in both caches, and from then on it behaves as if memory had the latest version. Some systems even update memory in the background, but it’s not strictly necessary as long as one clean copy exists.

⌚ **However, if we don’t write to the memory, but instead just forward it over the bus, the memory will still be stale!** That’s fine because the guarantee of coherence (page 447) means that *at least one cache has the correct clean copy*. If later all shared copies are evicted, **one of them will write back to memory** before discarding, so memory is eventually updated.

Coherence guarantees inter-cache correctness, not: “*memory is always the truth*”. In fact, in multiprocessors with write-back caches, **memory is often stale**. As long as the protocol ensures that no processor uses stale data, correctness is maintained. Since memory doesn’t always have the updated value, the MSI and MESI protocols (covered in the next chapter) ensure that caches communicate and that no processor ever sees stale memory.

- **Modified → Invalid:** Another processor issues a *write miss*. We must supply our dirty data (since memory is stale): we put the

block on the bus, and this can go to **memory** (write-back, so it becomes updated). Finally, drop our copy.

- **Shared → Invalid:** Another processor issues a *write miss* or *invalidate*. We simply drop our copy, no write-back needed since we are clean (identical to memory).

^aNote: this is the most common type of MSI on exams.

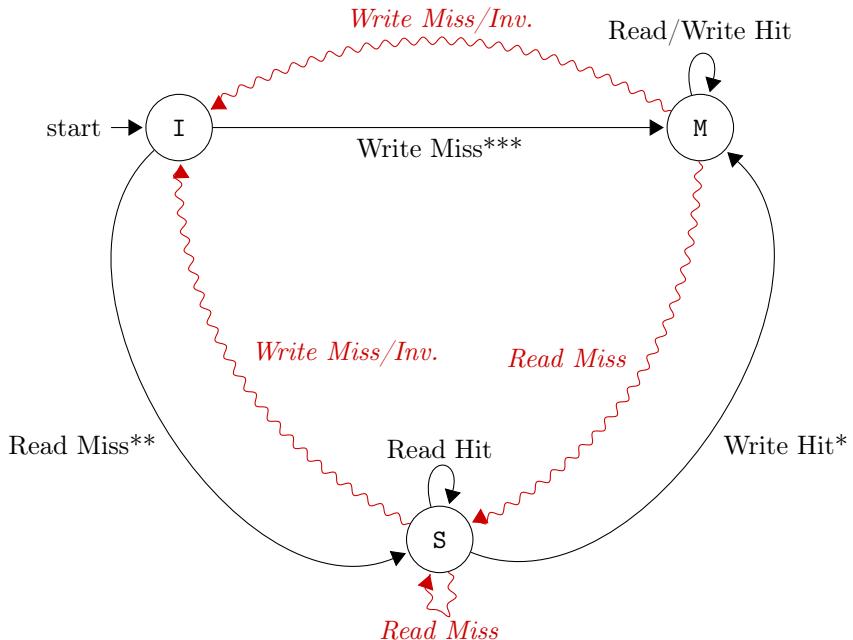


Figure 38: **MSI Protocol: Finite-State Machine (FSM)**. The asterisk (*) indicates that on a *write hit* in *S*, we must broadcast an *invalidate* on the bus so all other caches drop their copies. The red wavy arrows indicate bus events (snooped requests from other processors). Instead, the straight arrows indicate local CPU events. The double asterisk (**) indicates that on a *read miss* from *I*, we fetch data from memory via the bus. The triple asterisk (***)) indicates that on a *write miss* from *I*, we fetch data from memory via the bus, and we must also invalidate other caches that may have a copy.

Some notes about the diagram (Figure 38):

- There are no accepting/final states because the system runs continuously as the program executes, responding to events forever.
- Not all transitions are drawn:
 - Read/Write Hit in *I* is not possible because we don't have the block locally.
 - Read/Write Miss in *M* and *S* are not possible because we already have the block.

- There is bus traffic on some transitions, but not all. The **local CPU events** that cause bus traffic are:
 - (*) $S \rightarrow M$ (broadcast *Invalidate*).
 - (**) $I \rightarrow S$ (bus *Read Miss*, because data comes from memory).
 - (***) $I \rightarrow M$ (bus *Write Miss*, because data comes from memory and we must invalidate others).
- Obviously, the **snooped bus events** always cause bus traffic because they originate from the bus itself. These are:
 - $M \rightarrow S$ on a *Read Miss* from another core. We must supply the up-to-date data (write-back). Our line becomes clean and shared.
 - $M \rightarrow I$ on a *Write Miss* or explicit *Invalidate* from another core. We must supply the up-to-date data (since memory is stale), then drop our copy.
 - $S \rightarrow I$ on a *Write Miss* or *Invalidate* from another core. We are clean, so no data is supplied, just invalidate.
 - $S \rightarrow S$ on a *Read Miss* from another core. No state change, we remain shared. If memory or another cache supplies the data, we just stay in S .
- The **write-back** to memory occurs only when we go from $M \rightarrow S$ or $M \rightarrow I$, because in both cases we are supplying our dirty data to either memory or another cache. In all other transitions, no write-back is needed.

10.14.5 MESI Protocol

The **MESI Protocol** appeared in the early 1980s, notably with the Intel 486 (1989) being one of the first commercial CPUs to implement MESI. It was a natural refinement of MSI, created by **computer architects studying snooping coherence** for scalable multiprocessors. The design goal: reduce **bus contention**, improve **latency of private writes**, and keep protocols **simple enough for hardware implementation**.

⚠ Problem with MSI

In MSI, whenever a processor brings a block into the cache after a **Read Miss**, it always goes to **Shared (S)**, even if no one else has that block (see FSM diagram, page 456). But in that case, the cache *knows* it is the only copy (because the bus transaction didn't see any sharer response). Yet, MSI still marks it as "Shared", which forces:

- An unnecessary **Invalidate** on the next write (wasting bandwidth). Because the block is in state S (shared), the **next write must issue a trigger to invalidate other caches**, even though there are no other caches with that block.
- Extra latency before the processor can write. Because the block is in state S (shared), the **next write must wait for the Invalidate to complete before it can proceed**, even though there are no other caches with that block.

So MSI wasted both **bandwidth** (spurious invalidations) and **time** (delayed private writes).

💡 MESI Solution

To solve this problem, MESI adds a new state called **Exclusive (E)**. The key idea is that when a processor brings a block into the cache after a *Read Miss*, it can go to state **Exclusive (E)** if *no other cache has that block*. This captures the common case:

- **Most memory blocks are private** (i.e., not shared). For example, stack data and many heap objects are private to a single thread; or local variables in a function are private to that function; or unique data structures used by a single thread are private to that thread.
- **If a block is private and clean, the cache should not need to notify other caches on a write.** Instead, it can silently transition to **Modified (M)** without any bus transaction. This saves both bandwidth (no spurious invalidations) and time (no delayed writes).

Thus, *Read Misses* can go to state **Exclusive (E)** if no other cache has that block, and then the next write can transition to **Modified (M)** without any bus transaction (silently). **This optimization is the main benefit of MESI over MSI.**

≡ MESI States

MESI has four states:

- **Modified (M)**: The cache line is present only in the current cache and has been modified (is dirty). Memory is stale.
- **Exclusive (E)**: The cache line is present only in the current cache and is clean (matches memory). The cache can silently transition to *Modified (M)* on a write.
- **Shared (S)**: The cache line may be present in other caches and is clean (matches memory).
- **Invalid (I)**: The cache line is not valid (not present in the cache).

State	Description
Modified (M)	Valid, dirty, owned exclusively. Must write back to memory on eviction.
Exclusive (E)	Valid, clean, owned exclusively. Can be promoted to M without bus transaction.
Shared (S)	Valid, clean, may be in other caches.
Invalid (I)	Not valid or not present in cache.

Table 33: MESI Cache States.

≡ Key Transitions (compared to MSI)

The MESI protocol modifies some transitions from MSI and adds new ones:

- **On a Read Miss**
 - If no other cache responds, we need to fetch from memory → **E** (instead of S in MSI).
 - If some cache responds → **S**.
- **On a Write Hit**
 - **S** → **M** with bus invalidate (same as MSI).
 - **E** → **M** silently (no bus invalidate).
- **On snooped requests**
 - **E** → **S** if another processor reads.
 - **E** → **I** if another processor writes.

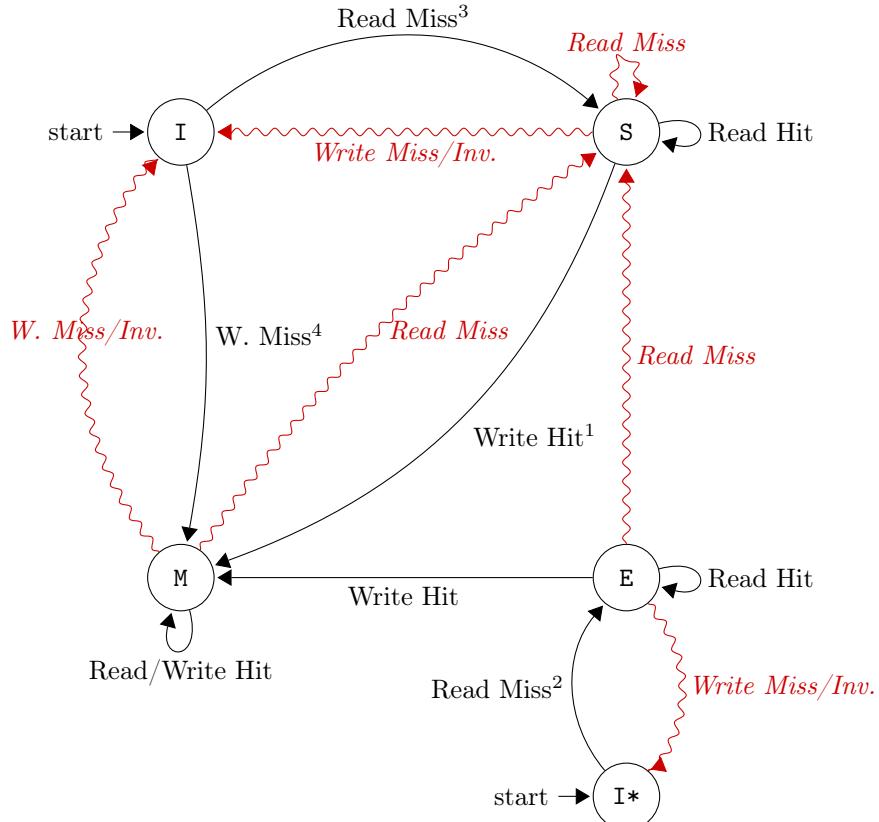


Figure 39: **MESI Protocol: Finite-State Machine (FSM)**. The states are: M (Modified), E (Exclusive), S (Shared), I (Invalid), and I* (Invalid, copy of state I for clarity; they are the same state). The transitions are triggered by local processor actions (solid arrows) or snooped bus transactions (wavy arrows). Notes:

- ¹ If the block is in state S, a write hit must issue an invalidate on the bus before transitioning to M;
- ² If no other cache has the block, it transitions to E;
- ³ If another cache has the block, it transitions to S;
- ⁴ Fetch data from memory via the bus (read-exclusive) and invalidate other caches.

❓ When to Write Back to Memory?

In MESI, a cache **writes back to memory only when evicting a dirty block**, i.e. when a line in **M (Modified)** is replaced or must be supplied to another core.

- **Eviction of an M line.** If the cache needs to evict a line in **Modified (M)**, it must **write it back to memory** (because memory is stale). Transition from **M → I**, with write-back.
- **Snooped Read Miss while in M.** Another processor requests the line. Our cache supplies the data (since memory is stale). Transition from **M → S**, with write-back.
- **Snooped Write Miss / Invalidate while in M.** Another processor wants ownership of the line. Our cache must **write back the block to memory** before invalidating its copy. Transition from **M → I**, with write-back.

There is absolutely no need to write back to memory when transitioning from **E (Exclusive)**, **S (Shared)**, or **I (Invalid)** states, since these states are clean (memory is up-to-date) or invalid (not present in cache). In other words, **the only time a write-back is needed is when evicting a dirty block in M state.**

10.15 Directory-Based Protocols

10.15.1 Basics

The **Directory-Based Protocol** is a cache coherence protocol that uses a centralized directory to keep track of the state of each memory block in a multiprocessor system. The directory maintains information about which caches have copies of each memory block and their states (e.g., modified, shared, invalid).

Unlike snooping protocols (which rely on broadcast), directory protocols store the sharing state of each memory block in a single, well-defined location, called **directory**.

☞ What is the Directory?

The **directory is hardware**, not an OS software structure. It is a **table** maintained in hardware inside the **memory controller of each node**. We can think of it as a **metadata** attached to each memory block.

In a **Distributed Shared-Memory (DSM)** system (more common in scalable multiprocessors), each memory module has its **own directory** stored next to it. So if a block resides in memory node N0, the directory for that block is in node N0's memory controller.

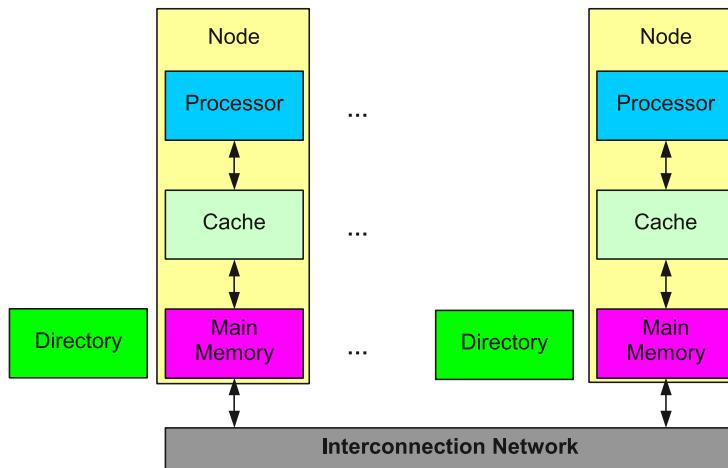


Figure 40: Directory-Based Protocol Overview from slides [4].

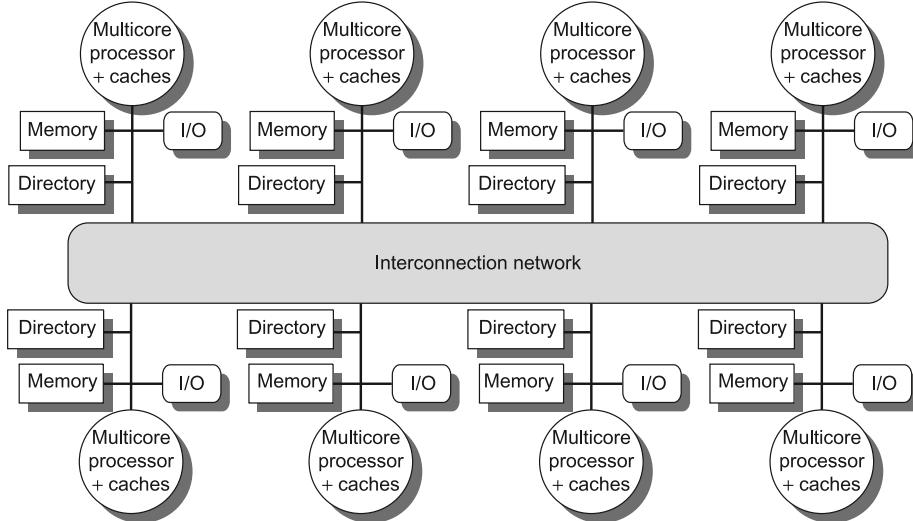


Figure 41: Directory-Based Protocol Overview from the book [2].

❷ What is stored in the Directory?

The directory keeps track of the **state of each memory block** and which nodes have cached copies of it. Directory keeps track of:

- The **coherence state** of the block (Uncached, Shared, Modified):
 - **U (Uncached)**: No cache has a copy of the block. Only the memory copy is valid (up-to-date).
 - **S (Shared)**: One or more caches may have it in read-only mode (because they issued a read request). Memory copy is still valid (**up-to-date**).
 - **M (Modified)**: One cache has a read-write copy of the block (the owner). Memory copy is **stale** (out-of-date).

These are conceptually analogous to the states in the MSI protocol (page 456), but tracked centrally¹⁵ in the directory rather than in each cache.

- A **sharer list**. It is usually implemented as a **bit vector** (one bit per node). If the bit node i is 1, it means node i 's cache currently has a valid **shared copy** of that block. If the bit is 0, node i does not cache that block.
 - **U state (Uncached)**: sharer list is all zeros.
 - **S state (Shared)**: sharer list may have several bits set to 1.
 - **M state (Modified)**: sharer list will have exactly one bit set to 1 (the owner).

¹⁵With “centrally” we mean that the directory has a global view of the state of each block, whereas in a distributed cache system, each cache would need to manage its own state independently.

- The **owner** of the block (if it is in the *Modified* state).

? Since the **owner field is filled only in the Modified state, and the share list contains only one bit set to 1 in the Modified state, why do we need both?** From a pure information theory point of view, the owner field is redundant. However, having a **dedicated owner field can simplify and speed up** certain operations, such as identifying the owner quickly without scanning the sharer list. This can be particularly useful in large systems where efficiency is critical.

So, in a DSM system, each node has its own local memory and a directory that tracks all the blocks physically mapped to that memory. For example, the node N0 has a directory that tracks all the blocks physically stored in its local memory and keeps track of which nodes have cached copies of those blocks.

This allows **point-to-point messaging** (no broadcast), which is essential for scalability in large multiprocessor systems (DSM, page 434).

👥 Roles of Nodes

In a directory-based protocol, nodes can take on different roles based on their interaction with memory blocks. For nodes that participate in a coherence transaction, we have **three roles**:

- 🏠 **Home Node (H)**: the node that “owns” the memory block (fixed by the physical address mapping). Stores both the **block in memory** and its **directory entry**. All requests for this block must eventually pass through its directory.
- 👉 **Local Node (L)**: the node where the **request originates**. For example, CPU at node N2 issues a read miss, then its cache controller is the local requester.
- 🌐 **Remote Node (R)**: a node that **already has a cached copy** of the block. (either *Shared* or *Modified*). If the copy is *Modified*, the most updated value resides there (not in the home memory), but the home directory knows which node is the owner and can request the updated value from it.

Obviously, for a given transaction, a node can take on multiple roles. For example, the local requester can also be a home node if it is requesting a block that it owns.

10.15.2 Message Types

In this section, we provide the “*vocabulary*” of directory protocols: the set of messages nodes exchange to enforce coherence.

Directory-based protocols are **message-oriented** (unlike snooping, which is broadcast-based). Each coherence event is handled by **explicit request and response messages** between Local, Home, and Remote nodes. The main message types are:

- **ReadMiss**: A Local node sends this message to the Home node to request read access to a cache block that is not present in its local cache.
 - ➔ **Source**: Local cache (requester).
 - ⬅ **Destination**: Home directory (where the block is located).
 - 📄 **Content**: Processor ID and address of the requested block.
 - ❓ **Purpose**: Add requester (processor ID) to the **sharer list** and supply data back to the requester.
- **DataValueReply**: The Home node sends this message to the Local node in response to a *ReadMiss* request, providing the requested data, or to a *WriteMiss* request if the block is in the Modified state.
 - ➔ **Source**: Home directory (or Remote, if block is Modified).
 - ⬅ **Destination**: Local cache (requester).
 - 📄 **Content**: Data of the requested block.
 - ❓ **Purpose**: Supply data to the requester.
- **Invalidate**: Two forms, depending on who sends it:
 - Sent by the Local node to the Home node to invalidate all shares of a block when it requests write access (via *WriteMiss*).
 - ➔ **Source**: Local cache (requester).
 - ⬅ **Destination**: Home directory.
 - 📄 **Content**: Address of the block to be invalidated.
 - Sent by the Home node to Remote nodes to **invalidate** their **copies** of a block when a Local node requests write access (via *WriteMiss*).
 - ➔ **Source**: Home directory.
 - ⬅ **Destination**: Remote caches (sharers).
 - 📄 **Content**: Address of the block to be invalidated.

Here, the **common purpose** is to remove stale shared copies so a requester can become exclusive owner. The Local node sends the initial invalidate request to the Home node, which then forwards it to all Remote nodes that have a copy of the block.

- **Fetch:** The Home node sends this message to a Remote node that has the block in the Modified state (the owner), requesting it to send the updated data back to the Home node.

→ **Source:** Home directory.
 ← **Destination:** Remote cache (owner).
 ⚑ **Content:** Address of the block to be fetched.
 ? **Purpose:** Satisfy a *ReadMiss* when memory is stale.

The effect of this message is that the Remote node passes from *Modified* to *Shared* state, as it must supply the data to the Home node but can keep a read-only copy.

- **Fetch/Invalidate (Fetch/Inv):** The Home node sends this message to a Remote node that has the block in the Modified state (the owner), requesting it to send the updated data back to the Home node and invalidate its copy. It is the combination of *Fetch* and *Invalidate*.

→ **Source:** Home directory.
 ← **Destination:** Remote cache (owner).
 ⚑ **Content:** Address of the block to be fetched and invalidated.
 ? **Purpose:** Satisfy a *WriteMiss* when memory is stale.

The effect of this message is that the Remote node passes from *Modified* to *Uncached* state, as it must supply the data to the Home node and invalidate its copy. Also, the requester will become the new owner (Modified state).

- **Data Write-Back:** A Remote node sends this message to the Home node to write back the updated data of a block it owns (in Modified state) when it is evicted from its cache.

→ **Source:** Remote cache (owner).
 ← **Destination:** Home directory.
 ⚑ **Content:** Address and updated data of the block being written back.
 ? **Purpose:** Update memory and directory state when a block is evicted.

This message ensures that the Home node has the most recent data and can update the directory accordingly.

The key insight is that **all coherence actions are explicitly managed through these messages**, allowing for precise control over cache states and memory consistency in a distributed environment. Each message type serves a specific role in maintaining coherence, ensuring that all nodes have a consistent view of memory. Despite the **snooping protocols being more intuitive** due to their broadcast nature, **directory-based protocols offer better scalability and efficiency** in large multiprocessor systems by avoiding unnecessary broadcasts and reducing contention on the memory bus.

Message	Source → Destination	Purpose
ReadMiss	Local → Home	Request block on <i>read miss</i>
DataValueReply	Home/Remote → Local	Return block data
Invalidate	Home → Remotes	Invalidate stale copies
Fetch	Home → Owner	Get block back for <i>read miss</i>
Fetch/Inv	Home → Owner	Get block back + invalidate (<i>write miss</i>)
Data Write-Back	Remote → Home	Return dirty block to memory

Table 34: Summary of Directory-Based Protocol Message Types.

10.15.3 Walkthroughs on a 4-node system

Let's consider a 4-node system with nodes N0, N1, N2, and N3. Each node has its own cache and local memory. We will walk through a series of operations to illustrate how a directory-based cache coherence protocol works in this system.

➊ Preliminary Setup

We assume the following **initial conditions**:

- We assume 4 nodes: N0, N1, N2, and N3.
- Each node has its own Processor (P), Cache (C), Main Memory (MM), and Directory (D).
- Physical address space is **statically partitioned** among the nodes, so each block has a **Home Node**.
- The directory entry at the Home Node contains:
 - The **state** of the block: U (uncached), S (shared), or M (modified).
 - **Sharer list / Owner bits**: vector marking which nodes hold copies.

So we only walk through the protocol for a single block, say block B0, which is mapped to N0 (i.e., N0 is the Home Node for block B0).

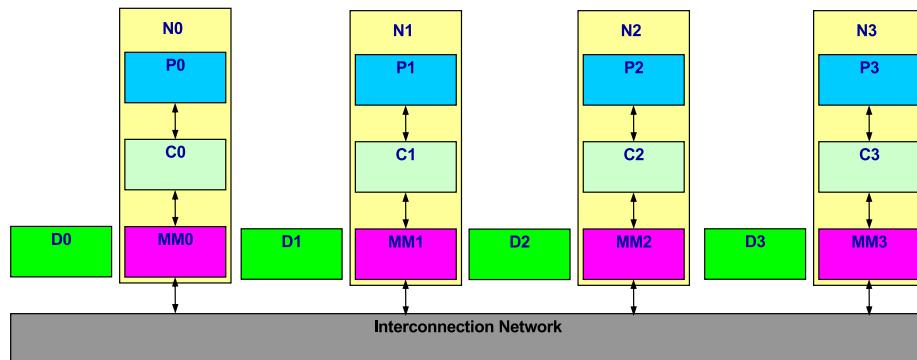


Figure 42: 4-node system with directory-based cache coherence protocol. Each node (N0, N1, N2, N3) has its own Processor (P0, P1, P2, P3), Cache (C0, C1, C2, C3), Main Memory (MM0, MM1, MM2, MM3), and Directory (D0, D1, D2, D3). The directory at each node keeps track of the state and sharers/owner of each block.

❶ Uncached State (U)

The block is not cached in any node (uncached state). The directory entry at N0 for block B0 is:

- **State:** U
- **Sharer list:** [0, 0, 0, 0] (no sharers)

$$B0 \mid \text{State: } U \mid \text{Sharers: } [0, 0, 0, 0]$$

From here, we can consider the following scenarios:

- **Read Miss from N1:**

1. N1 (Local) sends a request (**ReadMiss**) to N0 (Home) for block B0.

$$\text{Local (N1)} \xrightarrow{\text{ReadMiss}} \text{Home (N0)}$$

$$\text{ReadMiss(ProcID: 1, Addr: B0)}$$

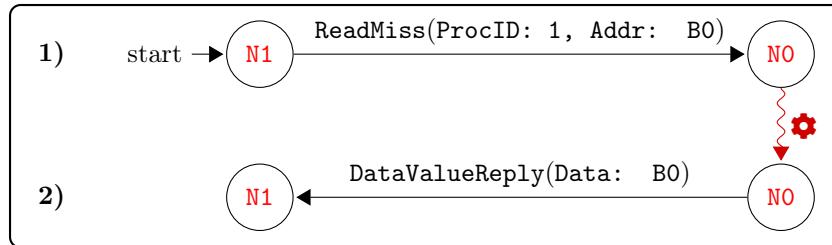
2. Home sends data from memory (because the data are not cached) to N1 (**DataValueReply**).

$$\text{Home (N0)} \xrightarrow{\text{DataValueReply}} \text{Local (N1)}$$

$$\text{DataValueReply(Data: B0)}$$

3. Directory (at N0) updates its state to S and adds N1 to the sharer list.

$$B0 \mid \text{State: } S \mid \text{Sharers: } [0, 1, 0, 0]$$



The effect is that N1 now has a shared copy of block B0. However, the memory copy is still valid.

- **Write Miss from N2:**

1. N2 (Local) sends a request (**WriteMiss**) to N0 (Home) for block B0.

$$\text{Local (N2)} \xrightarrow{\text{WriteMiss}} \text{Home (N0)}$$

$$\text{WriteMiss(ProcID: 2, Addr: B0)}$$

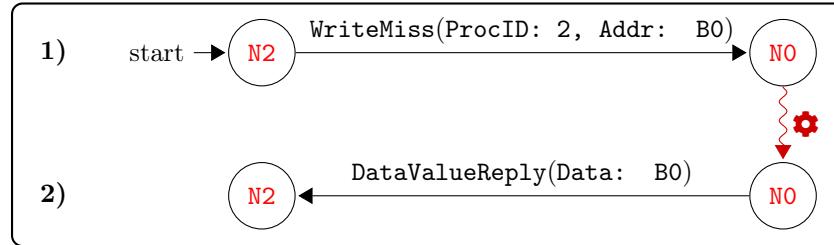
2. Home sends block from memory (because the data are not cached) to N2 (**DataValueReply**). It also sends **Invalidate** messages to all sharers (in this case, no one has it, so no invalidates are sent).

$$\text{Home (N0)} \xrightarrow{\text{DataValueReply}} \text{Local (N2)}$$

$$\text{DataValueReply(Data: B0)}$$

3. Directory (at N0) updates its state to M and sets N2 as the owner:

B0 | State: M | Sharers: [0, 0, 1, 0] | Owner: N2



The effect is that N2 now has exclusive ownership of block B0 in Modified state. The memory copy is now stale.

❶ Shared State (S)

The block is cached in one (or more) node(s) in Shared state. Continuing from the scenario above, the directory entry at N0 for block B0 is:

- State: S
- Sharer list: [0, 1, 0, 0] (only N1 has a copy)

B0 | State: S | Sharers: [0, 1, 0, 0]

From here, we can consider the following scenarios:

• Read Miss from N2:

1. N2 (Local) sends a request (`ReadMiss`) to N0 (Home) for block B0.

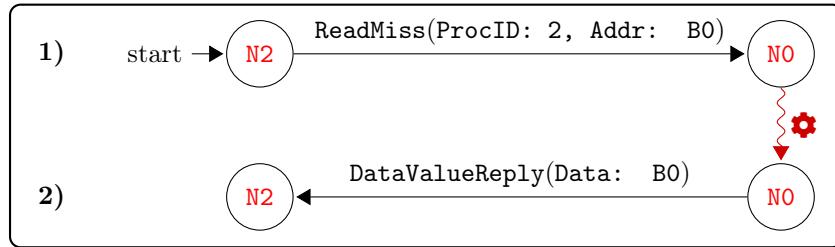
Local (N2) $\xrightarrow{\text{ReadMiss}}$ Home (N0)
ReadMiss(ProcID: 2, Addr: B0)

2. Home sends data from memory (because the data are not cached) to N2 (`DataValueReply`).

Home (N0) $\xrightarrow{\text{DataValueReply}}$ Local (N2)
DataValueReply(Data: B0)

3. Directory (at N0) updates its sharer list to include N2.

B0 | State: S | Sharers: [0, 1, 1, 0]



The effect is that both N1 and N2 now have shared copies of block B0. The memory copy is still valid.

- **Write Miss from N2:**

1. N2 (Local) sends a request (`WriteMiss`) to NO (Home) for block B0.

Local (N2) $\xrightarrow{\text{WriteMiss}}$ Home (NO)

`WriteMiss(ProcID: 2, Addr: B0)`

2. Home sends `Invalidate` messages to all sharers (in this case, N1) to invalidate their copies.

Home (NO) $\xrightarrow{\text{Invalidate}}$ Remote (N1)

`Invalidate(Addr: B0)`

3. N1 (Remote) invalidates its copy of block B0 and sends an acknowledgement back to Home.

Remote (N1) $\xrightarrow{\text{Ack}}$ Home (NO)

`Ack(Addr: B0)`

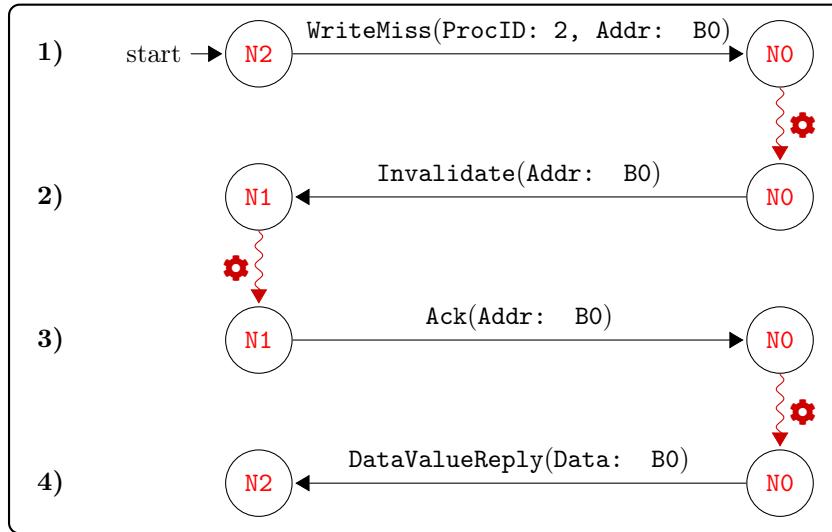
4. After receiving all acknowledgments, Home sends data from memory to N2 (`DataValueReply`).

Home (NO) $\xrightarrow{\text{DataValueReply}}$ Local (N2)

`DataValueReply(Data: B0)`

5. Directory (at NO) updates its state to M and sets N2 as the owner, removing all other sharers.

`B0 | State: M | Sharers: [0, 0, 1, 0] | Owner: N2`



The effect is that N2 now has a copy of block B0 in Modified state. The memory copy is now stale, and N1 no longer has a valid copy.

❶ Modified State (M)

The block is cached in one node in Modified state. Let's assume the directory entry at NO for block B0 is:

- State: M
- Sharer list: [0, 1, 0, 0] (only N1 has a copy)

B0 | State: M | Sharers: [0, 1, 0, 0] | Owner: N1

From here, we can consider the following scenarios:

- **Read Miss from N2:**

1. N2 (Local) sends a request (`ReadMiss`) to NO (Home) for block B0.

Local (N2) $\xrightarrow{\text{ReadMiss}}$ Home (NO)

`ReadMiss(ProcID: 2, Addr: B0)`

2. Home sends a `Fetch` request to the owner (N1) to get the data, because the owner (N1) has the most up-to-date copy and it is in Modified state (only one owner in M state).

Home (NO) $\xrightarrow{\text{Fetch}}$ Remote (N1)

`Fetch(Addr: B0)`

3. N1 (Remote) sends the data to Home (NO, `Data Write-Back`), and the owner (N1) transitions from Modified to Shared state because it must supply the data but can keep a read-only copy.

Remote (N1) $\xrightarrow{\text{Data Write-Back}}$ Home (NO)

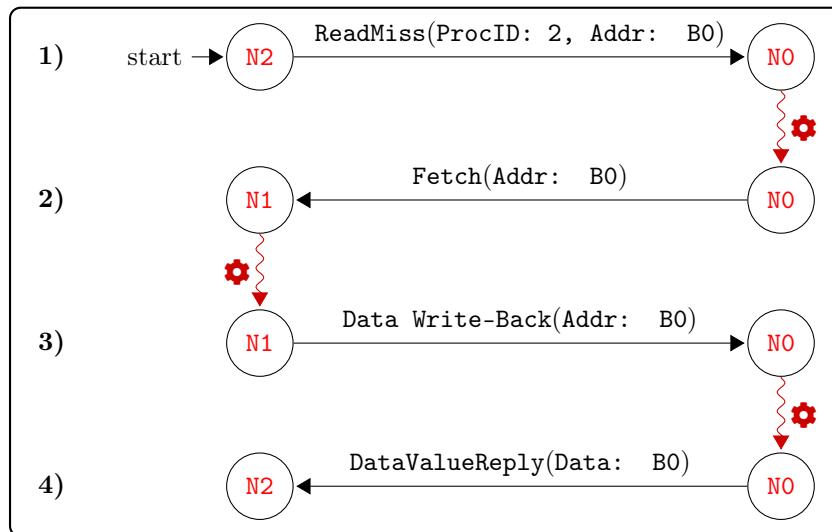
`Data Write-Back(Data: B0)`

4. Home sends data to N2 (**DataValueReply**). It is important to note that Home does not change the state to Modified because N2 is only requesting read access. The state would change to Modified only if N2 requested write access (via **WriteMiss**).

Home (N0) $\xrightarrow{\text{DataValueReply}}$ Local (N2)
DataValueReply(Data: B0)

5. Directory (at N0) updates its sharer list to include N2. The state becomes Shared (S) because now there are multiple sharers.

B0 | State: S | Sharers: [0, 1, 1, 0]



The effect is that both N1 and N2 now have shared copies of block B0. The memory copy is now stale, but N1 and N2 have valid copies; N1 has the most up-to-date data.

- **Write Miss from N2:**

1. N2 (Local) sends a request (**WriteMiss**) to NO (Home) for block B0.

Local (N2) $\xrightarrow{\text{WriteMiss}}$ Home (N0)
WriteMiss(ProcID: 2, Addr: B0)

2. Home sends a **Fetch/Inv** request to the owner (N1) to get the data and invalidate its copy, because the owner (N1) has the most up-to-date copy and it is in Modified state (only one owner in M state). The combination of **Fetch** and **Invalidate** is used because N2 is requesting write access.

Home (N0) $\xrightarrow{\text{Fetch/Inv}}$ Remote (N1)
Fetch/Inv(Addr: B0)

3. N1 (Remote) sends the data to Home (N0, Data Write-Back) and invalidates its copy, transitioning from Modified to Uncached state.

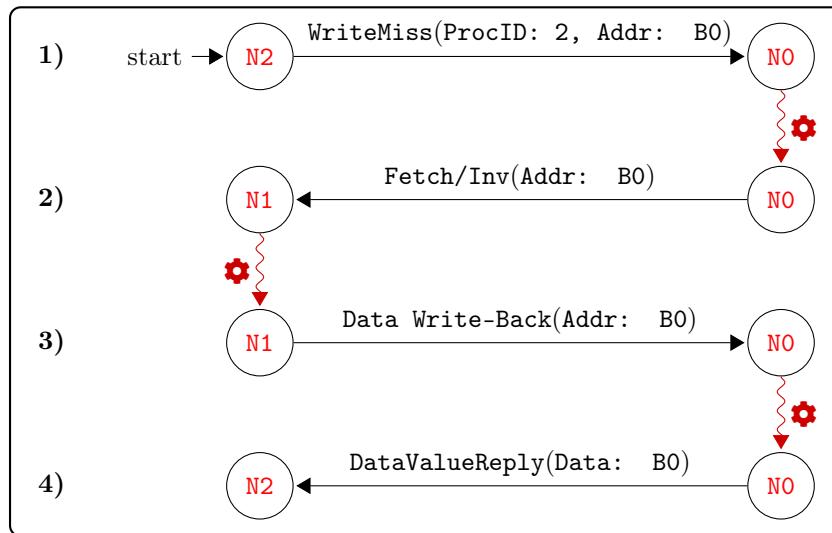
Remote (N1) $\xrightarrow{\text{Data Write-Back}}$ Home (N0)
Data Write-Back(Data: B0)

4. Home sends data to N2 (DataValueReply). Now, Home changes the state to Modified because N2 is requesting write access and will become the new owner.

Home (N0) $\xrightarrow{\text{DataValueReply}}$ Local (N2)
DataValueReply(Data: B0)

5. Directory (at N0) updates its state to M and sets N2 as the owner, removing all other sharers.

B0 | State: M | Sharers: [0, 0, 1, 0] | Owner: N2



The effect is that N2 now has a copy of block B0 in Modified state. The memory copy is now stale, and N1 no longer has a valid copy.

- **Eviction at the Owner (Replacement):**

1. N1 (Remote, Owner) decides to evict block B0 from its cache (e.g., due to cache replacement policy).
2. N1 sends a Data Write-Back message to Home (N0) to write back the updated data of block B0.

Remote (N1) $\xrightarrow{\text{Data Write-Back}}$ Home (N0)
Data Write-Back(Data: B0)

3. Home (at N0) receives the write-back data and updates memory with latest data.
4. Directory (at N0) updates its state to U and clears the sharer list and owner bits, as no node now has a valid copy. However, the memory copy is now up-to-date.

B0 | State: U | Sharers: [0, 0, 0]

The effect is that block B0 is now uncached in all nodes, and the memory copy is up-to-date. The directory reflects that no node has a valid copy.

10.15.4 Directory FSM and Invariants

A directory-based protocol can be described using a **finite state machine (FSM)** that captures the **states of cache lines and the transitions between these states based on various events** (e.g., read miss, write miss, invalidate, etc.). The FSM helps in understanding how the directory manages coherence among multiple caches in a multiprocessor system.

Each block of memory (at its Home node) has three possible states in the directory:

- **Uncached (U):** Block not present in any cache. Memory up-to-date.
- **Shared (S):** One or more caches have clean copies. Memory up-to-date.
- **Modified (M):** Exactly one cache has a dirty copy. Memory stale.

We describe this FSM **from the Home directory's perspective**, which tracks the state of each block and manages requests from other caches.

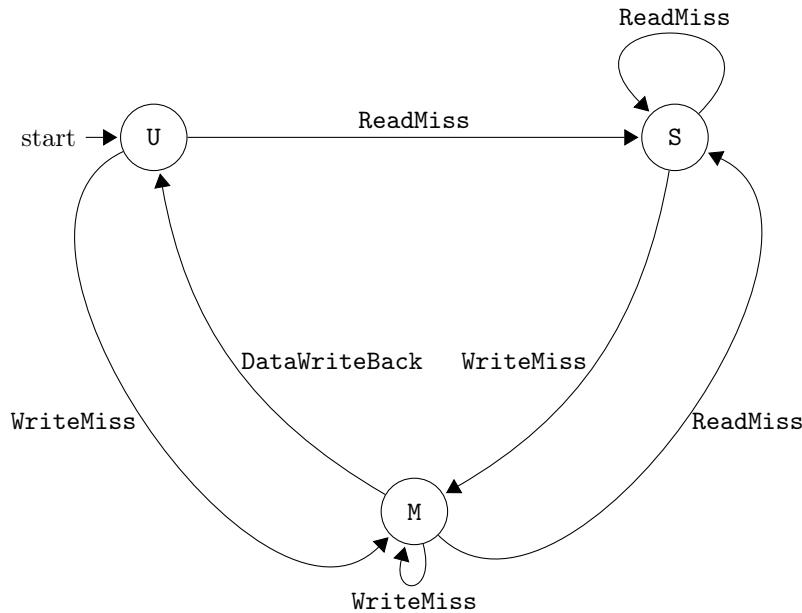


Figure 43: Directory FSM with states Uncached (U), Shared (S), and Modified (M). Transitions are triggered by events such as **ReadMiss**, **WriteMiss**, and **DataWriteBack**.

☰ State Transitions

- From **Uncached (U)**:

- $U \rightarrow S$. On **ReadMiss**: Transition to **Shared (S)**.
 - * The directory sends the **data to the requesting cache** and updates its state to **S**.
 - * The **cache is added** to the **sharer list**.
 - * The data is read from memory because no cache has it (see page 469 for more details).
- $U \rightarrow M$. On **WriteMiss**: Transition to **Modified (M)**.
 - * The directory sends an **invalidate** message to all eventual sharers, grants ownership to the requesting cache, and updates its state to **M**.
 - * The **cache is added** to the **sharer list**.
 - * The data is read from memory because no cache has it (see page 471 for more details).

- From **Shared (S)**:

- $S \rightarrow S$. On **ReadMiss**: Remain in **Shared (S)**.
 - * The directory sends the **data to the requesting cache** without changing the state.
 - * The **cache is added** to the **sharer list**.
 - * The data is read from one of the caches that already has it (see page 470 for more details).
- $S \rightarrow M$. On **WriteMiss**: Transition to **Modified (M)**.
 - * The directory sends **invalidate messages to all other caches holding the block**.
 - * **Grants ownership to the requesting cache**.
 - * Updates its state to **M**.
 - * Send the data to the requesting cache (see page 471 for more details).

- From **Modified (M)**:

- $M \rightarrow S$. On **ReadMiss**: Transition to **Shared (S)**.
 - * Send a **request to the current owner cache** for the data.
 - * The owner cache **sends the data** to the requesting cache and **dowgrades** its state to **S**.
 - * The directory **updates the sharer list** to include both caches and changes its state to **S**.
 - * The **memory copy remains stale**, but both caches have valid copies (see page 472 for more details).
- $M \rightarrow M$. On **WriteMiss**: Remain in **Modified (M)**.
 - * Send a **request to the current owner cache** for the data and to **invalidate** its copy.

- * The owner cache **sends the data** to the requesting cache and **invalidates** its copy.
 - * The directory **updates the sharer list** to include only the requesting cache (the new owner) and remains in state M.
 - * The **memory copy remains stale**, but the requesting cache has the most up-to-date data (see page 473 for more details).
- M → U. On DataWriteBack: Transition to Uncached (U).
- * The owner cache **sends the updated data back to memory**.
 - * The directory **clears the sharer list** and updates its state to U.
 - * This occurs when the owner cache evicts the block (see page 475 for more details).

⚠ Invariants

To ensure the correctness of the directory-based cache coherence protocol, several invariants must be maintained. These are **conditions that must always hold true** regardless of the sequence of operations performed on the caches and memory. The key invariants for a directory-based protocol are:

1. **Single owner invariant**: At most one cache can have a block in Modified (M).
2. **Memory freshness invariant**:
 - ✓ If state = S or U → memory is up-to-date.
 - ✗ If state = M → memory is stale.
3. **Sharer set invariant**: The sharer list must accurately reflect all nodes with cached copies.
4. **Owner membership invariant**: If a block is M → the owner must appear in the sharer/owner bits (only one).
5. **State-transition correctness**: All state changes occur only via defined message exchanges (ReadMiss, DataValueReply, Invalidate, Fetch, Fetch/Inv, Data Write-Back).
6. **Atomicity of updates**: The directory must process one request at a time for a given block (serialization).

11 Exams

11.1 2026

11.1.1 January 8

Exercise 1A - Scoreboard

PART A: Please consider the program in the table be executed on a CPU with dynamic scheduling based on **SCOREBOARD BASIC SCHEME** with the following assumptions and resources:

- Check for WAW hazards in the ISSUE stage;
- No Forwarding;
- 2 LOAD/STORE Units (LDU1, LDU2) with latency 3 cycles;
- 1 FP Unit (FPU1) with latency 3 cycles;
- 1 ALU/BR Unit (ALU1) with latency 1 cycle.

Instruction	I	RR	E	WB	Hazards	Type	Unit
I0:LD \$F2, VB(\$R6)	1	2	5	6			LDU1
I1:FADD \$F4, \$F2, \$F6							
I2:SD \$F4, VA(\$R7)							
I3:LD \$F4, VC(\$R6)							
I4:SD \$F4, VC(\$R7)							
I5:ADDUI \$R6, \$R6, 4							
I6:ADDUI \$R7, \$R7, 4							

Where I is the Issue stage, RR is the Read Operands stage, E is the Execute stage, WB is the Write Back stage.

Calculate the **CPI** of the program and draw the timing diagram showing the status of each instruction at each clock cycle.

Answer: The scoreboard (page 103) allows out-of-order execution while maintaining program correctness by tracking hazards and resource usage. It is similar to Tomasulo's algorithm but lacks register renaming and forwarding mechanisms. So, **in-order issue** is enforced, there is no forwarding, and **out-of-order execution and commit** are allowed (which can introduce WAW and WAR hazards, since there's no register renaming).

In a traditional pipeline, the Instruction Decode (ID) stage performs both decoding and operand fetching. But this assumes that operands are always available in the register file during the ID stage. However, with dynamic scheduling, operands might not be ready when an instruction is issued, leading to stalls if we wait for them to be fetched in the ID stage. So, the scoreboard splits the ID stage into two separate stages:

- **Issue Stage.** It decodes the instruction, checks for structural hazards, and enforces in-order issue. If there are structural hazards, the instruction is stalled here.
- **Read Operands Stage.** It waits for the operands to become available. It avoids RAW hazards by deferring operand reads until the register is no longer “reserved” by an active instruction writing to it. If there are RAW hazards, the instruction is stalled here. Since there is no forwarding, the instruction must wait until the writing instruction completes its Write Back stage.

So the **RAW** hazards are handled in the **Read Operands stage (RR)**, while **WAR** hazards are handled in the **Write Back stage (WB)**, **WAW** and **structural** hazards are handled in the **Issue stage (I)**. See subsubsection 3.6.4 (page 109) for more details on the scoreboard stages.

2. Cycle 2

- We do not consider the first cycle and the instruction I0 because they are already filled in the table.
- I1 can be issued in cycle 2 because there are no structural hazards (FPU1 is free) and no WAW hazards (no previous instruction writes to \$F4). So, I1 is in the Issue stage in cycle 2.
- Since I1 is issued, it reserves the unit FPU1 and the destination register \$F4.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6		2					FPU1
SD \$F4, VA(\$R7)							
LD \$F4, VC(\$R6)							
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

3. Cycle 3

- I1 would like to read its operands in cycle 3, but it cannot because of a RAW hazard on \$F2 (it is being written by I0 which has not yet completed its Write Back stage). So, I1 is stalled until cycle 6 (inclusive), when I0 completes its Write Back stage.
- I2 can be issued in cycle 3 because there are no structural hazards (LDU2 is free) and no WAW hazards (no previous instruction writes to \$F4). So, I2 is in the Issue stage in cycle 3.
- Since I2 is issued, it reserves the unit LDU2 and the destination register \$F4.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2				RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3						LDU2
LD \$F4, VC(\$R6)							
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

4. Cycle 4

- I1 is still stalled in the Read Operands stage due to the RAW hazard on \$F2.
- I2 would like to read its operands in cycle 4, but it cannot because of a RAW hazard on \$F4 (it is being written by I1 which has not yet completed its Write Back stage). So, I2 is stalled until cycle 9 (inclusive), when I1 completes its Write Back stage.
- *I3 would like to be issued in cycle 4, but it cannot because of a **structural hazard** on LDU1 (it is being used by I0) and a **WAW hazard** on \$F4 (it is being written by I1 which has not yet completed its Write Back stage). So, I3 is stalled.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2				RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3				RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)					*STR/WAW		
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

7. Cycle 7

- Now I1 can read its operands in cycle 7, since the RAW hazard on \$F2 has been resolved (it was written back by I0 in cycle 6). So, I1 is in the Read Operands stage in cycle 7.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7			RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3				RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)					STR/WAW		
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

12. Cycle 12

- We skip to cycle 12, but in cycle 10 the execution of I1 completes and in cycle 11 it writes back its result, resolving the RAW hazard

on \$F4 for I2. So, in cycle 12 I2 can read its operands. So, I2 is in the Read Operands stage in cycle 12.

- I3 can now be issued because the WAW hazard on \$F4 has been resolved (it was written back by I1 in cycle 11) and there are no structural hazards (LDU1 is free since I0 completed its Write Back stage in cycle 6). So, I3 is in the Issue stage in cycle 12.
- Since I3 is issued, it reserves the unit LDU1 and the destination register \$F4.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12			RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12				STR/WAW		LDU1
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

13. Cycle 13

- I3 reads its operands in cycle 13, since there is no RAW hazard on its \$R6 (it is not being written by any active instruction). So, I3 is in the Read Operands stage in cycle 13.
- I4 would like to be issued in cycle 13, but it cannot because of a **structural hazard** on LDU2 (it is being used by I2).

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12			RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13			STR/WAW		LDU1
SD \$F4, VC(\$R7)					STR \$F4 I3		
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

17. Cycle 17

- I3 execution completes in cycle 16 and it writes back its result in cycle 17.
- I4 can now be issued because the structural hazard on LDU2 has been resolved (it was freed by I2 which completed its Write Back stage in cycle 16). So, I4 is in the Issue stage in cycle 17.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12	15	16	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13	16	17	STR/WAW		LDU1
SD \$F4, VC(\$R7)	17				STR \$F4 I3		LDU2
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

18. Cycle 18

- I4 reads its operands in cycle 18, since there is no RAW hazard on its \$R7 (it is not being written by any active instruction). So, I4 is in the Read Operands stage in cycle 18.
- I5 can be issued in cycle 18 because there are no structural hazards (ALU1 is free) and no WAW hazards (no previous instruction writes to \$R6). So, I5 is in the Issue stage in cycle 18.
- Since I5 is issued, it reserves the unit ALU1 and the destination register \$R6.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12	15	16	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13	16	17	STR/WAW		LDU1
SD \$F4, VC(\$R7)	17	18			STR \$F4 I3		LDU2
ADDUI \$R6, \$R6, 4		18					ALU1
ADDUI \$R7, \$R7, 4							

19. Cycle 19

- I5 reads its operands in cycle 19, since there is no RAW hazard on its \$R6 (it is not being written by any active instruction). So, I5 is in the Read Operands stage in cycle 19.
- I6 cannot be issued in cycle 19 because of a **structural hazard** on ALU1 (it is being used by I5).

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12	15	16	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13	16	17	STR/WAW		LDU1
SD \$F4, VC(\$R7)	17	18			STR \$F4 I3		LDU2
ADDUI \$R6, \$R6, 4	18	19					ALU1
ADDUI \$R7, \$R7, 4					STR ALU1		

21. Cycle 21

- I4 completes its execution in cycle 20.
- I5 writes back its result in cycle 21, freeing up ALU1. There are no WAR hazards on \$R6 since no active instruction reads from it, so there are no stalls.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12	15	16	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13	16	17	STR/WAW		LDU1
SD \$F4, VC(\$R7)	17	18	21		STR \$F4 I3		LDU2
ADDUI \$R6, \$R6, 4	18	19	20	21			ALU1
ADDUI \$R7, \$R7, 4					STR ALU1		

22. Cycle 22

- I4 writes back its result in cycle 22, freeing up LDU2. There are no WAR hazards on \$F4 since no active instruction reads from it, so there are no stalls.
- I6 can now be issued because the structural hazard on ALU1 has been resolved (it was freed by I5 which completed its Write Back stage in cycle 21) and there are no WAW hazards (no previous instruction writes to \$R7). So, I6 is in the Issue stage in cycle 22.
- Since I6 is issued, it reserves the unit ALU1 and the destination register \$R7.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12	15	16	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13	16	17	STR/WAW		LDU1
SD \$F4, VC(\$R7)	17	18	21	22	STR \$F4 I3		LDU2
ADDUI \$R6, \$R6, 4	18	19	20	21			ALU1
ADDUI \$R7, \$R7, 4	22				STR ALU1		ALU1

25. Cycle 25

- The execution of I6 completes in cycle 24 and it writes back its result in cycle 25.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	7	10	11	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	12	15	16	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	12	13	16	17	STR/WAW		LDU1
SD \$F4, VC(\$R7)	17	18	21	22	STR \$F4 I3		LDU2
ADDUI \$R6, \$R6, 4	18	19	20	21			ALU1
ADDUI \$R7, \$R7, 4	22	23	24	25	STR ALU1		ALU1

The **CPI** (Cycles Per Instruction) can be calculated as follows:

$$\text{CPI} = \frac{\# \text{ clock cycles}}{\# \text{ instructions}} = \frac{25}{7} \approx 3.57$$

Exercise 1B - Scoreboard

PART B: Please consider the introduction of the following **Optimizations**:

- Check for WAW postponed to the WRITE BACK phase;
- Forwarding;

Instruction		I	RR	E	WB	Hazards	Type	Unit
I0:LD	\$F2, VB(\$R6)	1	2	5	6			LDU1
I1:FADD	\$F4, \$F2, \$F6							
I2:SD	\$F4, VA(\$R7)							
I3:LD	\$F4, VC(\$R6)							
I4:SD	\$F4, VC(\$R7)							
I5:ADDUI	\$R6, \$R6, 4							
I6:ADDUI	\$R7, \$R7, 4							

Where I is the Issue stage, RR is the Read Operands stage, E is the Execute stage, WB is the Write Back stage.

Fill the following table indicating, for each cycle, which instruction is in which stage, and which hazards (if any) are present. Highlight the stalls caused by hazards. Finally, compute the **CPI** (Cycles Per Instruction) achieved with these optimizations and the speedup with respect to the previous case (without optimizations).

Answer: With respect to Exercise 1A, here we have two optimizations: (1) WAW hazard checks are postponed to the Write Back stage, so instructions can be issued even if there is a WAW hazard; (2) Forwarding is implemented, so RAW hazards can be resolved earlier during the Execute stage instead of waiting for the Write Back stage. About forwarding, it simply means that when an instruction produces a result, that result can be forwarded directly to a subsequent instruction that needs it, without waiting for the first instruction to write it back to the register file.

In the following tables, we do not report all cycles, but only the ones where the optimizations have an effect.

6. Cycle 6

- In the cycle 6, I1 was stalled in the Read Operands stage due to the RAW hazard on \$F2. But now, with forwarding, I1 can read its operand \$F2 directly from the output of I0 during its Execute stage (cycle 5), instead of waiting for I0 to write it back in cycle 6. So, I1 can read its operands in cycle 6. It anticipates the Write Back of I0 by one cycle.

- I3 cannot be issued in cycle 6 because of a **structural hazard** on LDU1 (it is being used by I0). However, there is no WAW hazard on \$F4 because WAW checks are postponed to the Write Back stage. This will allow I3 to be issued earlier in a later cycle.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6			RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3				RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)					STR LDU1		
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

7. Cycle 7

- Thanks to WAW checks being postponed to the Write Back stage, I3 can now be issued in cycle 7, even if there is a WAW hazard on \$F4 with I1. So, I3 is in the Issue stage in cycle 7.
- Since I3 is issued, it reserves the unit LDU1 and the destination register \$F4.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6			RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3				RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7				STR LDU1	LDU1	
SD \$F4, VC(\$R7)							
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

8. Cycle 8

- I3 reads its operands in cycle 8, since there is no RAW hazard on its \$R6 (it is not being written by any active instruction). So, I3 is in the Read Operands stage in cycle 8.
- I4 cannot be issued in cycle 8 because of a **structural hazard** on LDU2 (it is being used by I2). I4 chooses the FU LDU1 because is the least recently used among the two load/store units.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6			RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3				RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8			STR LDU1	LDU1	
SD \$F4, VC(\$R7)					STR LDU2		
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

10. Cycle 10

- I1 completes its execution in cycle 9 and it writes back its result in cycle 10.
- I2 can now read its operands in cycle 10, since there is forwarding and it can get the value of \$F4 directly from the output of I1 during its Write Back stage (cycle 10), instead of waiting for I1 to write it back in cycle 11. So, I2 is in the Read Operands stage in cycle 10.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6	9	10	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	10			RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8			STR LDU1		LDU1
SD \$F4, VC(\$R7)					STR LDU2		
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

13. Cycle 13

- I3 completes its execution in cycle 11 and it writes back its result in cycle 12. There is a WAW hazard on \$F4 with I1, but since WAW checks are postponed to the Write Back stage, I3 can write back its result in cycle 12 without stalls (I1 already wrote back in cycle 10, so there is no conflict).
- I4 can now read its operands in cycle 13, since there is no more structural hazard on LDU2 (it was freed by I2 which completed its Execute stage in cycle 12). So, I4 is in the Read Operands stage in cycle 13.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6	9	10	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	10	13		RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8	11	12	STR LDU1		LDU1
SD \$F4, VC(\$R7)	13				STR LDU2	LDU1	
ADDUI \$R6, \$R6, 4							
ADDUI \$R7, \$R7, 4							

14. Cycle 14

- I2 completes its execution in cycle 13 and it writes back its result in cycle 14.
- I4 reads its operands in cycle 14, since there is no RAW hazard on its \$R6 (it is not being written by any active instruction). So, I4 is in the Read Operands stage in cycle 14.
- I5 starts its Issue stage in cycle 14. There are no structural hazards.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6	9	10	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	10	13	14	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8	11	12	STR LDU1		LDU1
SD \$F4, VC(\$R7)	13	14			STR LDU2		LDU1
ADDUI \$R6, \$R6, 4							ALU1
ADDUI \$R7, \$R7, 4							

15. Cycle 15

- I5 reads its operands in cycle 15, since there are no RAW hazards on its source register \$R6 (it is not being written by any active instruction). So, I5 is in the Read Operands stage in cycle 15.
- I6 cannot be issued in cycle 15 because of a **structural hazard** on ALU1 (it is being used by I5).

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6	9	10	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	10	13	14	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8	11	12	STR LDU1		LDU1
SD \$F4, VC(\$R7)	13	14			STR LDU2		LDU1
ADDUI \$R6, \$R6, 4							ALU1
ADDUI \$R7, \$R7, 4							STR ALU1

16. Cycle 16

- I4 completes its execution in cycle 17 and it writes back its result in cycle 18.
- I6 can now be issued because the structural hazard on ALU1 has been resolved (it was freed by I5 which completed its Execute stage in cycle 17). So, I6 is in the Issue stage in cycle 18.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6	9	10	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	10	13	14	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8	11	12	STR LDU1		LDU1
SD \$F4, VC(\$R7)	13	14	17	18	STR LDU2		LDU1
ADDUI \$R6, \$R6, 4	14	15	16	17			ALU1
ADDUI \$R7, \$R7, 4							STR ALU1

21. Cycle 21

- The execution of I6 completes in cycle 20 and it writes back its result in cycle 21.

Instruction	I	RR	E	WB	Hazards	Type	Unit
LD \$F2, VB(\$R6)	1	2	5	6			LDU1
FADD \$F4, \$F2, \$F6	2	6	9	10	RAW \$F2 I0		FPU1
SD \$F4, VA(\$R7)	3	10	13	14	RAW \$F4 I1		LDU2
LD \$F4, VC(\$R6)	7	8	11	12	STR LDU1		LDU1
SD \$F4, VC(\$R7)	13	14	17	18	STR LDU2		LDU1
ADDUI \$R6, \$R6, 4	14	15	16	17			ALU1
ADDUI \$R7, \$R7, 4	18	19	20	21	STR ALU1		ALU1

The **CPI** (Cycles Per Instruction) can be calculated as follows:

$$\text{CPI} = \frac{\# \text{ clock cycles}}{\# \text{ instructions}} = \frac{21}{7} = 3$$

The **speedup** with respect to the previous case (without optimizations) is:

$$\text{Speedup} = \frac{\text{CPI}_{\text{old}}}{\text{CPI}_{\text{new}}} = \frac{3.57}{3} \approx 1.19$$

So, with these optimizations, we achieve a speedup of approximately 1.19 times compared to the original implementation without optimizations, about a 19% performance improvement.

Exercise 2 - VLIW Scheduling

Let's consider the following LOOP code, where $\$R_i$ are integer registers and $\$F_i$ are floating-point registers.

```

1 FOR:   LD $F2, VB ($R6)    # I0
2        FADD $F4, $F2, $F6  # I1
3        SD  $F4, VA ($R7)  # I2
4        LD  $F4, VC ($R6)  # I3
5        SD  $F4, VC ($R7)  # I4
6        FADD $F8, $F8, $F8  # I5
7        ADDUI $R6, $R6, 4   # I6
8        ADDUI $R7, $R7, 4   # I7
9        BNE $R7, $R8, FOR   # I8

```

Given a 3-issue VLIW machine with fully pipelined functional units:

- 1 Memory Unit with 3 cycles latency
- 1 FP ALUs with 3 cycles latency
- 1 Integer ALU with 1 cycle latency to next Int/FP/L/S & 2 cycles latency to next Branch

The branch is completed with 1 cycle delay slot (branch solved in ID stage). **No branch prediction.** In the Register File, it is possible to read and write at the same address at the same clock cycle. Considering one iteration of the loop, complete the following table by using the **list-based scheduling** (do NOT introduce any software pipelining, loop unrolling and modifications to loop indexes) on the 3-issue VLIW machine including the BRANCH DELAY SLOT. Please do not write in NOPs.

Let's complete the following table indicating, for each cycle, which instructions are issued in the VLIW: And finally, answer the following questions:

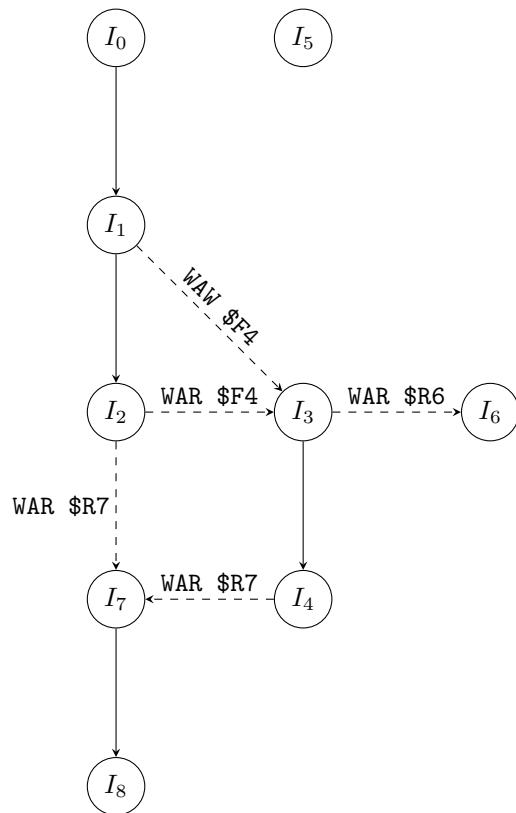
	Integer ALU	Memory Unit	FP Unit
C1		LD \$F2, VB(\$R6)	
C2			
C3			
C4			
C5			
C6			
C7			
C8			
C9			
C10			
C11			
C12			
C13			
C14			
C15			

- How long is the critical path for a single iteration?
- What performance did you achieve in FP ops per cycles?
- How much is the code efficiency for a single iteration?
- What CPI did you achieve?

Answer: To perform list-based scheduling for the given loop code on a 3-issue VLIW machine, we first need to draw the dependency graph for the instructions (see page 277). In the following list, we indicate the dependencies between instructions:

- RAW:
 - I1 depends on I0 (reads \$F2 written by I0)
 - I2 depends on I1 (reads \$F4 written by I1)
 - I4 depends on I3 (reads \$F4 written by I3)
 - I8 depends on I7 (reads \$R7 written by I7)
- WAR and WAW:
 - I3 has a WAW hazard with I1 (both write to \$F4)
 - I3 has a WAR hazard with I2 (writes to \$F4 read by I2)
 - I6 has a WAR hazard with I3 (writes to \$R6 read by I3)
 - I7 has a WAR hazard with I4 (writes to \$R7 read by I4) and I2 (writes to \$R7 read by I2)

Using this information, we draw the dependency graph below.



And the table is completed simply by scheduling the instructions while respecting the dependencies and the functional unit constraints (waiting for RAW hazards to be resolved and avoiding structural hazards):

	Integer ALU	Memory Unit	FP Unit
C1		LD \$F2, VB(\$R6)	FADD \$F8, \$F8, \$F8
C2			
C3			
C4			FADD \$F4, \$F2, \$F6
C5			
C6			
C7		SD \$F4, VA (\$R7)	
C8	ADDUI \$R6, \$R6, 4	LD \$F4, VC (\$R6)	
C9			
C10			
C11	ADDUI \$R7, \$R7, 4	SD \$F4, VC (\$R7)	
C12			
C13	BLT \$R7, \$R8, FOR (delay)		
C14			
C15			

The BLT instruction is placed in cycle 13 because the previous instruction (I_7) takes 2 cycles instead of 1 to prepare for the branch, as specified in the problem

statement. Additionally, a delay slot must be added in cycle 14 after the branch instruction, as required by the exercise.

Now, we can answer the questions:

- *How long is the critical path for a single iteration?* It corresponds to the last instruction to complete its Write Back stage: I8 in cycle 14. Therefore, the critical path is **14 cycles**.
- *What performance did you achieve in FP ops per cycles?* There are 2 floating-point operations (I1 and I5) completed in 14 cycles. Therefore, the performance is:

$$\text{FP ops/cycle} = \frac{2 \text{ FP ops}}{14 \text{ cycles}} \approx 0.14 \text{ FP ops/cycle}$$

- *How much is the code efficiency for a single iteration?* Code efficiency quantifies how well the issued instructions utilize the available issue slots. In our case, we have a 3-issue VLIW machine and we issued 9 instructions in 14 cycles (cycles C2, C3, C5, C6, C9, C10, and C12 are empty). Therefore, the code efficiency is:

$$\text{Code Efficiency} = \frac{\text{Instruction Count}}{\# \text{ cycles} \cdot \text{Issue Width}} = \frac{9}{14 \cdot 3} \approx 0.21$$

- *What CPI did you achieve?* The CPI (Cycles Per Instruction) can be calculated as follows:

$$\text{CPI} = \frac{\# \text{ clock cycles}}{\text{Instruction Count}} = \frac{14}{9} \approx 1.56$$

Exercise 3 - MESI Protocol

Let's consider the following access patterns on a dual processor system with a direct-mapped, write-back cache with one cache block per processor and a two cache block memory. Assume the **MESI protocol** is used, with **write-back** caches, **write-allocate**, and **write-invalidate** of other caches. Please complete the following table:

T	After Operation	P0 cache block state	P1 cache block state	Memory at block 0 up to date?	Memory at block 1 up to date?
0	P0: read block 1	Excl(1)	Invalid	Yes	Yes
1	P1: write block 0				
2	P0: write block 0				
3	P1: read block 0				
4	P0: read block 1				
5	P1: read block 1				
6	P0: read block 1				
7	P1: write block 1				
8	P0: read block 0				
9	P1: write block 1				
10	P1: write block 1				
11	P0: read block 1				
12	P1: write block 1				
13	P0: write block 1				
14	P1: read block 1				

Answer: The MESI protocol (page 458) is a cache coherence protocol used in multiprocessor systems to maintain consistency between caches. The four states in the MESI protocol are Modified (M, a cache line that has been modified and is different from main memory), Exclusive (E, a cache line that is present only in one cache and is the same as main memory), Shared (S, a cache line that may be present in multiple caches and is the same as main memory), and Invalid (I, a cache line that is not valid). Also, the cache uses:

- **Write-Back** policy: data is written to main memory only when it is evicted from the cache. In MESI, this means that a cache line in the Modified state must be written back to memory before being invalidated or replaced.
- **Write-Allocate** policy: on a write miss, the cache allocates a new cache line and loads the data from memory before performing the write.
- **Write-Invalidate** policy: when a processor writes to a cache line, it invalidates that line in other caches to maintain coherence.

T	After Operation	P0 cache block state	P1 cache block state	Memory at block 0 up to date?	Memory at block 1 up to date?
0	P0: read block 1	Excl(1)	Invalid	✓ Yes	✓ Yes
1	P1: write block 0	Excl(1)	Mod(0)	✗ No	✓ Yes
2	P0: write block 0	Mod(0)	Invalid	✗ No	✓ Yes
3	P1: read block 0	Shr(0)	Shr(0)	✓ Yes	✓ Yes
4	P0: read block 1	Excl(1)	Excl(0)	✓ Yes	✓ Yes
5	P1: read block 1	Shr(1)	Shr(1)	✓ Yes	✓ Yes
6	P0: read block 1	Shr(1)	Shr(1)	✓ Yes	✓ Yes
7	P1: write block 1	Invalid	Mod(1)	✓ Yes	✗ No
8	P0: read block 0	Excl(0)	Mod(1)	✓ Yes	✗ No
9	P1: write block 1	Excl(0)	Mod(1)	✓ Yes	✗ No
10	P1: write block 1	Excl(0)	Mod(1)	✓ Yes	✗ No
11	P0: read block 1	Shr(1)	Shr(1)	✓ Yes	✓ Yes
12	P1: write block 1	Invalid	Mod(1)	✓ Yes	✗ No
13	P0: write block 1	Mod(1)	Invalid	✓ Yes	✗ No
14	P1: read block 1	Shr(1)	Shr(1)	✓ Yes	✓ Yes

Where **Mod(x)** indicates that the cache block x is in the Modified state, **Excl(x)** indicates that the cache block x is in the Exclusive state, **Shr(x)** indicates that the cache block x is in the Shared state, and **Invalid** indicates that the cache block is in the Invalid state. For reference on state transitions, see the complete finite state machine (FSM) for the MESI protocol on page 460.

Question 1: Cache Coherence Protocols

Let's consider the SNOOPING and the DIRECTORY-BASED protocols used to maintain the cache coherence in modern multiprocessor. Answer the following questions:

- Explain the main characteristics of each protocol.

Answer: (see page 449 for snooping protocols and page 462 for directory-based protocols)

– **Snooping Protocols** are **cache coherence protocols based on broadcast communication**. Each cache controller monitors (snoops) all coherence transactions issued on a **shared communication medium** (typically a bus or a broadcast interconnect). The main characteristics include:

- * All coherence requests are broadcast to all caches.
- * Each cache independently checks whether it has a copy of the block.
- * Coherence actions (invalidate / update) are taken locally.
- * No centralized state is maintained.

– **Directory-Based Protocols** are **cache coherence protocols based on centralized or distributed directories that track the sharing state of each memory block**. For each memory block, a **directory entry** records:

- * Which processors currently cache the block.
- * The coherence state (e.g., shared, exclusive).

Main characteristics include:

- * Coherence requests are sent to the directory, not broadcast.
- * The directory selectively notifies only the caches involved.
- * Supports **point-to-point communication**.
- * Eliminates unnecessary broadcasts.

- For which type of multiprocessor architecture is each protocol suitable for?

Answer:

– **Snooping protocol** is suitable for **small-scale centralized shared-memory multiprocessors (SMPs)**. In these architectures, processors share a **single logical memory**, a **broadcast interconnect** (like a bus) is available, and the number of processors is relatively small (typically up to 16 or 32). The use of broadcast allows every cache to observe coherence requests, but the approach **does not scale** due to increasing traffic and contention on the bus as the number of processors increases.

– **Directory-based protocol** is suitable for **large-scale distributed shared-memory multiprocessors**, including **NUMA** and **distributed shared-memory systems**. In these architectures, the memory is **physically distributed**, a coherence is enforced using

point-to-point messages, and the directory tracks which processors cache each memory block. This approach avoids broadcast and therefore **scales well with the number of processors**. It is suitable for systems with a large number of processors (hundreds or thousands) where snooping would be impractical due to high communication overhead.

- What are the **main benefits** of each protocol?

Answer:

– **Benefits of Snooping Protocols:**

- * **Simplicity**, since no centralized structure is required.
- * **Low coherence latency**, thanks to immediate broadcast visibility.
- * **Fast detection of coherence events**, as every cache observes all requests.
- * **Low protocol overhead** for small-scale systems.

In short, snooping is **simple, low latency, broadcast visibility**, and **good for small systems**.

– **Benefits of Directory-Based Protocols:**

- * **High scalability**, since broadcast is avoided.
- * **Reduced coherence traffic**, due to targeted point-to-point messages.
- * **Efficient bandwidth usage**, especially in large systems.
- * **Support** for physically distributed memory architectures.

In short, directory-based protocols are **scalable, no broadcast, targeted communication**, and **good for large/NUMA systems**.

- Explain the **two types** of Snooping protocol depending on what happens on a write operation.

Answer: (see page 452)

– **Write-Invalidate Protocol**. When a processor writes to a shared cache block, it **invalidates all other cached copies** of that block. Mechanism:

1. The writing processor issues an **invalidate request** on the bus.
2. All other caches **snoop** the bus and **invalidate** their local copies of the block if they have it.
3. The writing processor then obtains **exclusive ownership** and can update the block locally.

This approach allows **multiple readers but only one writer**. Subsequent writes by the same processor do **not generate bus traffic** and reduces coherence traffic for repeated writes.

– **Write-Update (or Write-Broadcast) Protocol**. When a processor writes to a shared cache block, it **broadcast the new data value** to all other caches. Mechanism:

1. The writing processor places the **updated value** on the bus.

2. All caches **snoop** the write.
3. Caches that hold the block **update their local copy** with the new value.

This approach allows to keep all cached copies **consistent at all times**, but generates **bus traffic on every write**. It is suitable when shared data are frequently read by many processors after each write.

- *What are the possible coherence states of a block in the home directory and the meaning of the sharer bits?*

Answer: (see page 463) In a **directory-based cache coherence protocol**, each memory block has an associated **directory entry** stored at its **home node**. The directory entry records the **coherence state of the block** and a set of **sharer bits**.

Coherence states in the home directory. A directory typically maintains **three coherence states** for each block:

1. **Uncached (U)** (or **Invalid**). The block is **not cached in any processor** and the memory holds the only valid copy.
2. **Shared (S)**. The block is cached by **multiple processors** and all cached copies are **read-only**. The memory contains an up-to-date copy.
3. **Modified (M)**. The block is cached by **exactly one processor** and that processor has **write permission**. The memory may be **out-of-date** if the cached copy has been modified.

Sharer bits. The **sharer bits** are a bit-vector associated with each directory entry. Each bit corresponds to a processor (or cache) in the system. The meaning of the sharer bits is:

- A bit is set to **1** if the corresponding processor currently **caches a copy** of the block.
- A bit is set to **0** if the corresponding processor does **not cache** the block.

The sharer bits allow the directory to send **targeted invalidations** and **update or recall messages** only to the processors that currently cache the block, avoiding broadcast communication. So, these bits are essential for scalability (avoiding broadcast), point-to-point communication, and efficiency in large multiprocessor systems.

Question 2: Dynamic Branch Prediction

Let's consider the **Dynamic Branch Prediction** techniques used in modern microprocessors. Answer to the following questions:

- What are the **main benefits** to introduce the **dynamic branch prediction** in a modern microprocessor?

Answer: (see page 37) Dynamic branch prediction is introduced in modern microprocessors to **reduce the performance penalties caused by control hazards** (i.e., branch instructions that can change the flow of execution) and to **enable aggressive exploitation of Instruction-Level Parallelism (ILP)**. The main benefits are:

1. **Reduction of branch penalty.** Dynamic branch prediction allows the processor to **predict the outcome of a branch at runtime**, instead of waiting for the branch condition to be resolved. This reduces pipeline stalls, avoids frequent pipeline flushes, and improves overall instruction throughput.
2. **Adaptation to runtime behavior.** Unlike static prediction, dynamic branch prediction **learns from the actual execution behavior** of the program, and can adapt when branch behavior changes due to input data, loop iteration counts or program phase changes. This leads to **higher prediction accuracy**.
3. **Support for deep pipelines.** Modern processors often have deep pipelines to increase clock speeds. Dynamic branch prediction **helps to keep the pipeline filled** by predicting branches early, which is crucial for maintaining high performance in deep pipelines.
4. **Enabling speculative execution.** Dynamic branch prediction allows **instructions beyond unresolved branches** to be fetched and **executed speculatively**. This increases ILP, keeps function units busy, and improves overall performance. Speculative execution would not be effective without accurate dynamic prediction.
5. **Performance scalability with program complexity.** As programs become more branch-intensive and more irregular, dynamic branch prediction **provides robust performance**, independent of **compile-time heuristics**.

In short:

- Reduce branch penalty.
- Adapt to runtime behavior.
- Essential for deep pipelines.
- Enable speculation and ILP.
- Higher accuracy than static.

- Explain how works a **Branch History Table**.

Answer: (see page 56) A **Branch History Table (BHT)** is a hardware structure used in **dynamic branch prediction** to predict the outcome of conditional branches based on **past execution behavior**.

Structure. The BHT is a **small memory table** indexed using part of the **branch instruction address (PC)**. Each entry contains a **prediction state**, typically implemented using a **1-bit** or a **2-bit saturating counter** (i.e., a finite state machine) that indicates whether the branch is predicted to be taken or not taken.

Operation.

1. **During instruction fetch (IF stage),** the PC of the branch indexes the BHT. The corresponding entry provides a prediction: **taken or not taken**.
2. **Speculative execution.** The processor fetches and executes instructions according to the predicted outcome.
3. **When the branch outcome is resolved,** the actual outcome (taken or not taken) is compared with the prediction. The BHT entry is **updated** to reflect the new behavior.

In summary, the BHT is a table indexed by PC, storing past branch behavior to predict future branches using taken/not taken states. It is updated after each branch resolution and usually uses a 2-bit saturating counter (FSM) for better accuracy over simple 1-bit schemes.

- Explain the main purpose to introduce the **Branch Target Buffer**.

Answer: (see page 61) The **main purpose** of introducing a **Branch Target Buffer (BTB)** is to **eliminate the control hazard caused by the unknown target address of taken branches**, allowing the processor to continue instruction fetch **without stalling**.

In particular, the BTB provides the **predicted target address** of a branch **during the instruction fetch (IF) stage**, enables the processor to **fetch the next instruction from the predicted target in the same cycle**, and avoids waiting for branch target computation in later pipeline stages.

This is essential to support **efficient speculative execution**, especially in deeply pipelined processors.

- Explain the main concepts of the **correlating branch prediction technique**.

Answer: (see page 65) The **Correlating Predictor** (also known as **correlating branch prediction technique**) improves accuracy by **using the outcomes of recently executed branches to predict the current branch**.

How correlation is captured. To exploit correlation, the processor introduces a **history register**, called **Global History Register (GHR)**. It is a shift register of the last **N branch outcomes** (taken/not taken), where each bit represents whether a branch was taken (1) or not taken (0). This is **explicit history**, unlike the implicit FSM state of a basic BHT.

The correlating predictor combines the **branch address (PC)** with the **Global Branch History (GHR)** to select a prediction entry from a

Pattern History Table (PHT). The PHT contains multiple entries for each branch, each corresponding to a different history pattern. The selection of the PHT entry is based on both the PC and the GHR. It is done by **concatenating** or **XORing**, and the selected entry contains a **2-bit saturating counter (FSM)** for the prediction.

Possible exam answer structure. Correlating branch prediction improves dynamic branch prediction accuracy by exploiting the correlation between different branch outcomes. It uses a history register (called Global History Register, GHR) to record the outcomes of recently executed branches and combines this information with the branch address to select a prediction entry. This two-level mechanism allows the predictor to adapt predictions based on the execution context, enabling accurate prediction of branches whose behavior depends on previous branches.

- Explain what happens at runtime in a Speculative Tomasulo Architecture in the case of a **branch misprediction**.

Answer: At runtime, before the branch is resolved:

1. A branch is **predicted** (taken / not taken).
2. Instructions **after the branch** are fetched, issued, and executed **speculatively**.
3. Their results are written into **ROB entries** but **not** into the register file or memory.
4. These instructions are **marked as speculative** in the ROB.
5. Once the branch is resolved:
 - (a) If the prediction was **correct**, the speculative instructions are **committed** in order.
 - (b) If the prediction was **incorrect** (misprediction):
 - i. **Flush speculative instructions.** All instructions younger than the branch are invalidated and removed from the ROB, reservation stations, and functional units.
 - ii. **Restore architectural state.** Since the speculative instructions did not update the register file or memory, because they were only in the ROB, the architectural state remains consistent. The processor can discard the ROB entries of the mispredicted instructions without needing to roll back any changes, and restore the rename table to the state before the branch.
 - iii. **Redirect instruction fetch.** The PC is set to the correct target address of the branch (if taken) or the next sequential instruction (if not taken).
 - iv. **Resume correct execution.** New instructions are fetched and executed from the correct path.

In case of a branch misprediction, the only cost is **lost cycles** due to discarded speculative work and the time taken to fetch and execute instructions from the correct path.

Quizzes

1. **Question 1 (format Multiple Choice - Single answer).** Let's consider the following code executed by a Vector Processor with:

- Vector Register File composed of 32 vectors of 16 elements per 32 bits/element;
- Scalar FP Register File composed of 32 registers of 32 bits;
- One Load/Store Vector Unit with operation chaining and memory bandwidth 32 bits;
- One Add/Sub Vector Unit with operation chaining.
- One Mul/Div Vector Unit with operation chaining.

```

1 L.V V1, RX      # Load vector from memory address RX into
                   V1
2 ADDVS.D V1, V1, F0  # FP add vector V1 to scalar F0
3 MULV.D V2, V1, V1  # FP multiply vectors V1 and V1
4 DIVVS.D V2, V2, F0  # FP divide vector V2 to scalar F0
5 L.V V3, RY      # Load vector from memory address RY into
                   V3
6 ADDVV.D V3, V2, V3  # FP add vectors V2 and V3
7 S.V V2, RZ      # Store vector V2 into memory address RZ
8 S.V V3, RW      # Store vector V3 into memory address RW

```

How many convoys? How many clock cycles to execute the code?

- **Answer 1:** 3 convoys, 48 clock cycles
- **Answer 2:** 2 convoys, 32 clock cycles
- **Answer 3:** 4 convoys, 64 clock cycles
- **Answer 4:** 5 convoys, 80 clock cycles

Motivate your answer by completing the following table:

	Load/Store	Add/Sub	Mul/Div
1 convoy			
2 convoy			
3 convoy			
4 convoy			
5 convoy			

Answer: We start from the table and then we can answer the question. The code can be scheduled as follows:

	Load/Store	Add/Sub	Mul/Div
1 convoy	L.V V1, RX	ADDVS.D V1, V1, F0	MULV.D V2, V1, V1
2 convoy	L.V V3, RY	ADDVV.D V3, V2, V3	DIVVS.D V2, V2, F0
3 convoy	S.V V2, RZ		
4 convoy	S.V V3, RW		
5 convoy			

Each convoy takes 16 clock cycles to complete, as the vector length is 16 elements and the memory bandwidth is 32 bits (1 element per cycle). Therefore, the total number of clock cycles to execute the code is:

$$\text{Total clock cycles} = 4 \text{ convoys} \times 16 \text{ cycles/convoy} = 64 \text{ clock cycles}$$

Final answer: 4 convoys, 64 clock cycles (Answer 3).

2. **Question 2 (format Multiple Choice - Single answer).** *Assume to apply a processor optimization resulting ten time faster on computation than the original mode of execution. What is the fraction of computation needed to quadruplicate the overall speedup?*

- **Answer 1:** 55%
- **Answer 2:** 83%
- **Answer 3:** 45%
- **Answer 4:** 40%
- **Answer 5:** 25%

Motivate your answer.

Answer: We can use Amdahl's Law (page 251) to solve this problem. Amdahl's Law states that the overall speedup S of a system is given by:

$$S = \frac{1}{(1-f) + \frac{f}{s}}$$

Where:

- f is the fraction of the execution time that can be improved (in this case, the computation fraction).
- s is the speedup of the improved portion (in this case, “ten”).

We want to find the fraction f such that the overall speedup S is quadrupled (i.e., $S = 4$). Plugging in the values, we have:

$$\begin{aligned} 4 &= \frac{1}{(1-f) + \frac{f}{10}} \\ (1-f) + \frac{f}{10} &= \frac{1}{4} \\ 10 - 10f + f &= \frac{10}{4} \\ -9f &= \frac{10}{4} - 10 \\ f &= -\left(\frac{10}{4} - 10\right) \cdot \frac{1}{9} \\ f &= \frac{30}{4} \cdot \frac{1}{9} \\ f &= \frac{30}{36} = \frac{5}{6} \approx 0.8333 \end{aligned}$$

Therefore, the fraction of computation needed to quadruplicate the overall speedup is approximately 83% (**Answer 2**).

3. **Question 3 (format Multiple Choice - Single answer).** Let us consider a computer with L1 and L2 caches with the following parameters:

- Processor Clock Frequency = 1 GHz;
- Hit Time L1 = 1 clock cycle; Hit Rate L1 = 95%;
- Hit Time L2 = 5 clock cycles; Hit Rate L2 = 90%; Miss Penalty L2 = 15 clock cycles;

How much is the Global Miss Rate for Last Level Cache?

- **Answer 1:** 5%
- **Answer 2:** 0.5%
- **Answer 3:** 10%
- **Answer 4:** 7.5%

Motivate your answer.

Answer: Given:

- L1 Hit Rate = 95% = 0.95 \Rightarrow L1 Miss Rate = $1 - 0.95 = 0.05 = 5\%$
- L2 Hit Rate = 90% = 0.90 \Rightarrow L2 Miss Rate = $1 - 0.90 = 0.10 = 10\%$

The Global Miss Rate for the Last Level Cache (L2) can be calculated as:

$$\text{Global Miss Rate} = \text{L1 Miss Rate} \times \text{L2 Miss Rate}$$

Plugging in the values:

$$\text{Global Miss Rate} = 0.05 \times 0.10 = 0.005 = 0.5\%$$

Therefore, the Global Miss Rate for the Last Level Cache is 0.5% (**Answer 2**).

Deepening: What is the Last Level Cache (LLC)?

The **Last Level Cache (LLC)** is the **lowest cache level before main memory**. If the hierarchy is L1 \rightarrow L2 \rightarrow L3 \rightarrow Main Memory, then L3 is the LLC. So here, in this question, L2 is the LLC.

4. **Question 4 (format Multiple Choice - Multiple answers).** Which of the following cache improvements are effective in reducing the miss penalty?

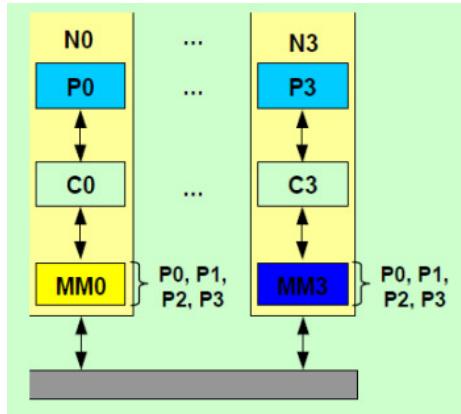
- **Answer 1:** Adopting higher associativity in designing the cache to reducing the impact of collisions when the correct block must be loaded in the cache.
- **Answer 2:** Avoiding address translation during indexing of the cache to minimize the management time of miss penalties.

- **Answer 3:** Combining a small and fast level-1 cache with a slower but larger level-2 cache to capture many accesses that would go to the main memory.
- **Answer 4:** Giving priority to read misses over writes, exploiting techniques to serve the reads before the writes have been completed.

Answer: The definition of **miss penalty** is the **additional time required to service a cache miss**, i.e. the time to fetch the block from the lower level (or memory) and make it usable by the processor. This should not be confused with the **miss rate**, which is the **fraction of memory accesses that result in a miss**. So we are looking for techniques that **shorten the time to recover from a miss**, or **avoid going all the way to main memory once a miss occurs**. Now, let's analyze each answer:

- ✗ **Answer 1:** *Adopting higher associativity...* - **Incorrect**. Higher associativity means **fewer conflict misses**, which **reduces the overall miss rate**, but it does **not directly reduce the miss penalty** itself. The time to service a miss remains the same regardless of associativity. So this affects **miss rate**, not **miss penalty**.
- ✓ **Answer 2:** *Avoiding address translation during indexing...* - **Correct**. During a cache miss, the processor must: (1) detect the miss, (2) translate the virtual address (TLB), (3) generate the physical address, and (4) access the lower-level cache or memory. If **address translation is on the critical path of the miss handling**, then the miss service starts later, the pipeline remains stalled longer, and the **effective miss penalty increases**. By **avoiding address translation during indexing**, part of the miss handling overlaps, translation latency is hidden and **total miss penalty is reduced**.
- ✓ **Answer 3:** *Combining a small and fast level-1 cache with a slower but larger level-2...* - **Correct**. When an L1 miss occurs, the access may be satisfied by **L2 instead of main memory**. Since L2 is **much faster than DRAM**, the effective miss penalty is **greatly reduced**. This is also the main motivation for **multi-level cache hierarchies** (L1, L2, L3). So this directly reduces **miss penalty**.
- ✗ **Answer 4:** *Giving priority to read misses over writes...* - **Incorrect**. While prioritizing read misses can improve overall performance by reducing stalls for read operations, it does **not reduce the miss penalty itself**. The time to service a miss remains unchanged; it only affects the order in which misses are handled. So this does not directly impact **miss penalty**.

5. Question 5 (format Multiple Choice - Single answer). What type of multiprocessor architecture does the following picture represent?



- **Answer 1:** Centralized Memory; Single Logically Shared Address Space
- **Answer 2:** Centralized Memory; Message Passing Architectures
- **Answer 3:** Distributed Memory; Single Logically Shared Address Space
- **Answer 4:** Distributed Memory; Message Passing Architectures

Answer: The correct answer is **Answer 3: Distributed Memory; Single Logically Shared Address Space**. Each vertical block contains processor (P_i), private cache (C_i), and local memory (MM_i). So physically there are **multiple memory modules**, each memory is **attached to a node** and there is **no single central memory**. So the first and second answers are incorrect because they refer to **centralized memory** architectures.

Now, each memory module is annotated with P_0, P_1, P_2, P_3 , indicating that **all processors can address all memory modules** and any processor can load/store from any MM_i . This indicates a **single logically shared address space**, where the entire memory is **logically shared among all processors**, even though it is physically distributed. Therefore, the third answer is correct.

The fourth answer is incorrect because in a **message passing architecture**, each processor can access **only its local memory**, and communication happens via send/receive operations and explicit messages. But in the picture, all processors access memory using normal memory operations, no explicit message queues or memory looks shared to software.

This picture represents a **Distributed Shared-Memory (DSM)** architecture, where memory is physically distributed but logically shared. Also known as **NUMA (Non-Uniform Memory Access)** architecture (page 429).

11.2 2025

11.2.1 Midterm - May 5

Exercise 1.A - Software Pipelining

Consider the following software pipelined loop *SP_LOOP* and the corresponding start-up and finish-up code:

```

1 START_UP:    LD F0, 0 (R1)
2             LD F2, 0 (R2)
3             FADD F4, F0, F0
4             FADD F6, F2, F2
5             LD F0, 8 (R1)
6             LD F2, 8 (R2)
7
8 SP_LOOP:     SD F4, 0 (R1)
9             SD F6, 0 (R2)
10            FADD F4, F0, F0
11            FADD F6, F2, F2
12            LD F0, 16 (R1)
13            LD F2, 16 (R2)
14            ADDUI R1, R1, 8
15            ADDUI R2, R2, 8
16            ADD R3, R2, R1
17            BNE R3, R4, SP_LOOP
18
19 FINISH_UP:   SD F4, 8 (R1)
20             SD F6, 8 (R2)
21             FADD F4, F0, F0
22             FADD F6, F2, F2
23             SD F4, 16 (R1)
24             SD F6, 16 (R2)

```

Reconstruct the original (non-pipelined) version of the code.

Answer: **Software Pipelining** is a loop optimization technique. Instead of processing iterations strictly sequentially, it overlaps instructions from multiple iterations. So in the pipelined version:

- The loop body does *parts of multiple iterations at the same time*.
- Because of this, we need extra **start-up** code (to prime the pipeline) and **finish-up** code (to flush the remaining operations).

The original loop must do **one iteration at a time**. In the original **non-pipelined loop**, we:

- Load the data for the current iteration.
- Do the computation.
- Store the results.
- Update pointers.
- Check whether you are done.
- Repeat.

In software pipelining we always have:

1. Start-up code
2. Steady-state loop
3. Finish-up code

Because the pipeline overlaps *parts of different iterations*. So we need:

- Extra code at the **start** to *prime* the pipeline with partial work.
- The **main loop** which does the steady overlapping work for all *middle* iterations.
- Extra code at the **end** to *drain* the pipeline, to finish the partial work left from the last overlapping iterations.

In other words, we need to reverse engineer the original code for this exercise.

The software pipelined loop usually overlaps:

- *Load* for **next** iteration.
- *Compute* for **current** iteration.
- *Store* for **previous** iteration.

The goal is trace the dependency chain:

- Which LD feeds which FADD.
- Which FADD feeds which SD.

So the key is to identify the loop-carried data flow and align addresses to match.

General 5-step method:

1. **Mark what the pipelined loop does.** In our case:

```

1 SP_LOOP:
2   SD F4, 0(R1)          # Store result of previous iteration?
3   SD F6, 0(R2)
4
5   FADD F4, F0, F0      # Compute for current iteration?
6   FADD F6, F2, F2
7
8   LD F0, 16(R1)        # Load for next iteration
9   LD F2, 16(R2)
10
11  ADDUI R1, R1, 8      # Increment base pointers
12  ADDUI R2, R2, 8

```

2. **Figure out the purpose**

- LD: prefetch next input.
- FADD: do the math for the loaded input.
- SD: store the result **from previous iteration**.

3. Sketch the timeline. Think conceptually:

Pipeline Stage	Operation
1. Load	Loads new data for $i + 1$
2. Compute	Uses old load for i
3. Store	Saves result for $i - 1$

4. Match start-up and finish-up

- **Start-up:** primes the pipeline, so gets first F0, F2 ready.
- **Loop:** repeats the middle chunk.
- **Finish-up:** drains the leftover results.

5. Write the original loop. For each iteration:

- Load F0, F2 \leftarrow data for i
- Compute $F4 = F0 + F0, F6 = F2 + F2 \leftarrow$ process for i
- Store F4, F6 \leftarrow save result for i
- Increment pointers
- Check loop condition

So, here's the final solution:

```

1 LOOP:   LD F0, 0 (R1)
2           LD F2, 0 (R2)
3           FADD F4, F0, F0
4           FADD F6, F2, F2
5           SD F4, 0 (R1)
6           SD F6, 0 (R2)
7           ADDUI R1, R1, 8
8           ADDUI R2, R2, 8
9           ADD R3, R2, R1
10          BNE R3, R4, LOOP

```

Exercise 1.B - Software Pipelining

Give the same software pipelined loop of **Exercise 1.A**:

```

1 SP_LOOP:    SD F4, 0 (R1)
2                 SD F6, 0 (R2)
3                 FADD F4, F0, F0
4                 FADD F6, F2, F2
5                 LD F0, 16 (R1)
6                 LD F2, 16 (R2)
7                 ADDUI R1, R1, 8
8                 ADDUI R2, R2, 8
9                 ADD R3, R2, R1
10                BNE R3, R4, SP_LOOP

```

Consider a **3-issue VLIW** machine with **fully pipelined functional units**:

- 1 Memory Units with 3 cycles latency
- 1 FP ALUs with 3 cycles latency
- 1 Integer ALU with 1 cycle latency to next Int/FP & 2 cycle latency to next Branch

The branch is completed with 1 cycle delay slot (branch solved in ID stage). **No branch prediction**. In the Register File, it is possible to read and write at the same address at the same clock cycle.

1. Considering one iteration of the **SP_LOOP**, complete the following table by using the **list-based scheduling** (do NOT introduce any loop unrolling and modifications to loop indexes) on the 3-issue VLIW machine including the **BRANCH DELAY SLOT**. Please do not write in NOPs.

Answer: In a **VLIW architecture** (Very Long Instruction Word), a single instruction word encodes **multiple operations** that execute *in parallel*. The compiler does the heavy lifting: it schedules independent instructions side by side in the same very long instruction word. The hardware just *fires* all operations in parallel, it does **not** dynamically check for dependencies at runtime (unlike out-of-order superscalar).

❓ **What does 3-issue mean?** It is the issue width and is the number of operations the processor *fetches, decodes, and executes* in parallel **per cycle**. So, with 3-issue, the processor can execute up to **3 instructions per clock cycles**. They must be **independent** enough for the compiler to pack them together.

❓ **What does fully pipelined functional units mean?** Each type of operation is done in its own **functional unit** (e.g., integer ALU, floating-point adder, load/store unit, etc.). **Fully pipelined** means multiple operations can be in *different pipeline stages* of the same unit simultaneously. So, every clock cycle a new operation can start, even if a previous one isn't finished yet. For example, if an FADD unit is fully pipelined and the FADD takes 4 cycles to complete, it can accept **one new FADD per cycle**, the result pops our 4 cycles later.

Q What is *list-based scheduling*? **List-based scheduling** is a classic compiler scheduling algorithm used to assign operations to *time slots* (cycles) and *resources* (e.g., functional units) **while respecting dependencies**.

- It tries to pack **as many independent instructions as possible** into each cycle.
- It tries to **fill all slots** of a wide-issue machine (VLIW, superscalar).
- It tries to respect resource limits (e.g., 1 integer ALU, 1 FP adder).

It is called list-based because the compiler first builds a **dependency graph** (where the nodes are the operations and the edges are the dependencies) and then creates a **list of ready operations**. These are instructions **whose inputs are ready and whose predecessors have been completed**. Each cycle:

- It picks operations from the list.
- It assigns them to available slots/resources.
- It marks successors ready when dependencies clear.

A What does *In the Register File, it is possible to read and write at the same address in the same clock cycle* mean? The **register file** supports *simultaneous read and write* to the **same register**. So, in one cycle:

- We might **read register F0** to feed an FADD.
- **AND write back a result** to F0 in the same cycle.

The hardware guarantees the read gets the **old value**, the write updates it for **later cycles**.

If the register file did not allow this, we'd have to schedule dependent instructions in separate cycles to avoid conflict (e.g., we couldn't read F0 to feed FADD while overwriting F0 with a new LD). Otherwise, if it is allowed, in same cycle we can:

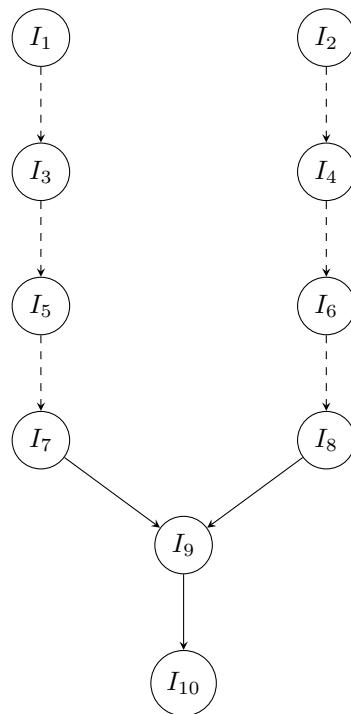
- LD F0, ... (writes F0)
- FADD F4, F0, F0 (reads old F0)

Here are some notes about the solution:

- We are using software pipelining, so the SD and FADD operators can start together because the SD operator stores the results of the previous iteration (or the initial loop, if it is the first iteration).
- The ADDUI operator starts when the first SD store finishes, thanks to the register file. In the same cycle, the load operation begins, as it loads the operands for the next iteration.
- Cycle 6 stalls because the BNE operator must wait for the result of the previous ADD and cannot start in cycle 6 because it reads old values.

❓ What should we do about our loop?

- (a) Build the **DAG** for one iteration. A Data Dependency Graph (DAG) for scheduling has nodes, which are operations or instructions, and edges, which are directed and show true dependencies (RAW). For example, an edge from A to B means that B must wait for A to finish.
- (b) Draw a dashed line if there are WAR or WAW hazards. This means we can perform these two instructions in the same cycle if a FU free is available.
- (c) In the table, write the instructions from start to finish.



	Memory Unit	Floating Point Unit	Integer Unit
C1	SD F4, 0 (R1)	FADD F4, F0, F0	
C2	SD F6, 0 (R2)	FADD F6, F2, F2	
C3	LD F0, 16 (R1)		ADDUI R1, R1, 8
C4	LD F2, 16 (R2)		ADDUI R2, R2, 8
C5			ADD R3, R2, R1
C6			
C7			BNE R3, R4, SP_LOOP
C8			(branch delay)

2. How long is the critical path for a single iteration?

Answer: in **instruction scheduling**, the **critical path** is the longest chain of dependent operations that must be executed *in sequence* because of true data dependencies (RAW). It determines the **minimum possible latency** to complete *one full logical iteration*. The critical path length depends on the data dependencies and the assumed functional unit latencies, **not** the scheduling algorithm.

In our case, there are 8 cycles for each iteration, so the **critical path is 8**.

3. What performance did you achieve in CPI?

Answer: The Clock Per Instruction (CPI) can be calculated using the following formula:

$$\text{CPI} = \frac{\# \text{ Clock Cycles}}{\text{IC}} = \frac{\# \text{ Clock Cycles}}{\# \text{ Instructions}} = \frac{8}{10} = 0.8$$

4. What performance did you achieve in FP ops per cycles?

Answer: FP ops per cycle measures **how well we are utilizing the FP functional units**. Here, there are 2 floating-point additions (FADD F4, FADD F6):

$$\text{FP ops / cycle} = \frac{\# \text{ Floating-Point ops}}{\# \text{ Clock Cycles}} = \frac{2}{8} = 0.25$$

5. How much is the code efficiency?

Answer: When asked for **code efficiency** in this VLIW/ILP context, they usually mean:

$$\text{Efficiency} = \frac{\text{Achieved IPC}}{\text{Maximum Theoretical IPC}}$$

So we measure **how much of the machine's peak parallelism we actually use**. The achieved Instruction Per Cycle (IPC) is:

$$\text{Achieved IPC} = \frac{\# \text{ Instructions}}{\# \text{ Clock Cycles}} = \frac{10}{8} = 1.25$$

Here, VLIW is 3-issue, so it can execute up to **3 ops per cycle** if there's no data or resource limit (Peak IPC = 3). So, code efficiency is:

$$\text{Efficiency} = \frac{1.25}{3} = 0.4167 \approx 41.67\% \approx 42\%$$

We're achieving about 41.67% ($\approx 42\%$) of the maximum possible execution parallelism.

Exercise 2.A - Dependency Analysis

Consider the same software pipelined loop of **Exercise 1.A** containing multiple types of intra-loop dependences. Complete the following table by inserting all types of true-data-dependences, anti-dependences and output dependences for each instruction:

Answer:

I#	Type of instruction	Analysis of dependences:
		1. True data dependence with I# for \$Fx 2. Anti-dependence with I# for \$Fy 3. Output-dependence with I# for \$Fz
I1	SD F4, 0 (R1)	None
I2	SD F6, 0 (R2)	None
I3	FADD F4, F0, F0	Anti-dependence with I1 for F4
I4	FADD F6, F2, F2	Anti-dependence with I2 for F6
I5	LD F0, 16 (R1)	Anti-dependence with I3 for F0
I6	LD F2, 16 (R2)	Anti-dependence with I4 for F2
I7	ADDUI R1, R1, 8	Anti-dependence with I1 for R1 Anti-dependence with I5 for R1
I8	ADDUI R2, R2, 8	Anti-dependence with I2 for R2 Anti-dependence with I6 for R2
I9	ADD R3, R2, R1	True data dependence with I7 for R1 True data dependence with I8 for R2
I10	BNE R3, R4, SP_LOOP	True data dependence with I9 for R3

Exercise 2.B - Tomasulo

Consider the same assembly code to be executed on a CPU with dynamic scheduling based on **Tomasulo algorithm** with all cache HITS, a single Common Data Bus and:

- 2 RESERVATION STATIONS (**RS1, RS2**) with 2 LOAD/STORE units (**LDU1, LDU2**) with latency 4
- 2 RESERVATION STATIONS (**RS3, RS4**) with 1 FP unit (**FPU1**) with latency 4
- 2 RESERVATION STATIONS (**RS5, RS6**) with 2 INT_ALU/BR units (**ALU1, ALU2**) with latency 2

Please complete the following table:

Answer: Before solving the Tomasulo, we need to remember how it is composed (see page 149).

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0						
I4: FADD F6, F2, F2						
I5: LD F0, 16 (R1)						
I6: LD F2, 16 (R2)						
I7: ADDUI R1, R1, 8						
I8: ADDUI R2, R2, 8						
I9: ADD R3, R2, R1						
I10: BNE R3, R4, LOOP						

1. Cycle 1:

	Vj	Qj	Vk	Qk
RS1	F4	-		
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

2. Cycle 2:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0						
I4: FADD F6, F2, F2						
I5: LD F0, 16 (R1)						
I6: LD F2, 16 (R2)						
I7: ADDUI R1, R1, 8						
I8: ADDUI R2, R2, 8						
I9: ADD R3, R2, R1						
I10: BNE R3, R4, LOOP						

	Vj	Qj	Vk	Qk
RS1	F4	-		
RS2	F6	-		
LDU1	F4	-		
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	4
LDU2	
FPU1	
ALU1	
ALU2	

3. **Cycle 3:** Here, we have an anti-dependence hazard (WAR) because the third instruction writes the sum to register F4 and the first instruction reads F4. In Tomasulo, though, we don't care about this because the issue stage stalls if and only if there are no reservation stations, the Load/Store queue is full (for memory ops), or for speculative instructions (for prediction, which is not our case). Therefore, we issue and start execution. At the end of execution, we will determine whether to write the result or wait.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3			WAR I1, F4	RS3	
I4: FADD F6, F2, F2						
I5: LD F0, 16 (R1)						
I6: LD F2, 16 (R2)						
I7: ADDUI R1, R1, 8						
I8: ADDUI R2, R2, 8						
I9: ADD R3, R2, R1						
I10: BNE R3, R4, LOOP						

	Vj	Qj	Vk	Qk
RS1	F4	-		
RS2	F6	-		
LDU1	F4	-		
LDU2	F6	-		

	Vj	Qj	Vk	Qk
RS3	F0		F0	
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	3
LDU2	4
FPU1	
ALU1	
ALU2	

4. **Cycle 4:** It's the same situation as before with the hazard (WAR), but now it's I4 with I2.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4		WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4			WAR I2, F6	RS4	
I5: LD F0, 16 (R1)						
I6: LD F2, 16 (R2)						
I7: ADDUI R1, R1, 8						
I8: ADDUI R2, R2, 8						
I9: ADD R3, R2, R1						
I10: BNE R3, R4, LOOP						

	Vj	Qj	Vk	Qk
RS1	F4	-		
RS2	F6	-		
LDU1	F4	-		
LDU2	F6	-		

	Vj	Qj	Vk	Qk
RS3	F0		F0	
RS4	F2		F2	
FPU1	F0		F0	

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	2
LDU2	3
FPU1	4
ALU1	
ALU2	

5. **Cycle 5:** Instruction 4 cannot begin because there are no available functional units. Therefore, it remains idle at reservation station 4. Instructions 5 and 6 cannot start either because there are no load/store reservation units available. ADDUI operations cannot start because instructions 5 and 6 need to read the R1 and R2 registers (WAR hazard).

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4		WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4			WAR I2, F6	RS4	
I5: LD F0, 16 (R1)				WAR I3, F0		
I6: LD F2, 16 (R2)				WAR I4, F2		
I7: ADDUI R1, R1, 8				WAR I5, R1		
I8: ADDUI R2, R2, 8				WAR I6, R2		
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	F4	-		
RS2	F6	-		
LDU1	F4	-		
LDU2	F6	-		

	Vj	Qj	Vk	Qk
RS3	F0		F0	
RS4	F2		F2	
FPU1	F0		F0	

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	1
LDU2	2
FPU1	3
ALU1	
ALU2	

6. Cycle 6:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4		WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4			WAR I2, F6	RS4	
I5: LD F0, 16 (R1)				WAR I3, F0		
I6: LD F2, 16 (R2)				WAR I4, F2		
I7: ADDUI R1, R1, 8				WAR I5, R1		
I8: ADDUI R2, R2, 8				WAR I6, R2		
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	F4	-		
RS2	F6	-		
LDU1	F4	-		
LDU2	F6	-		

	Vj	Qj	Vk	Qk
RS3	F0		F0	
RS4	F2		F2	
FPU1	F0		F0	

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	0
LDU2	1
FPU1	2
ALU1	
ALU2	

7. Cycle 7:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4		WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4			WAR I2, F6	RS4	
I5: LD F0, 16 (R1)	7			WAR I3, F0	RS1	
I6: LD F2, 16 (R2)				WAR I4, F2		
I7: ADDUI R1, R1, 8				WAR I5, R1		
I8: ADDUI R2, R2, 8				WAR I6, R2		
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	16		R1	
RS2	F6		-	
LDU1				
LDU2	F6		-	

	Vj	Qj	Vk	Qk
RS3	F0		F0	
RS4	F2		F2	
FPU1	F0		F0	

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	
LDU2	0
FPU1	1
ALU1	
ALU2	

8. Cycle 8:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4			WAR I2, F6	RS4	
I5: LD F0, 16 (R1)	7	8		WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8			WAR I4, F2	RS2	
I7: ADDUI R1, R1, 8				WAR I5, R1		
I8: ADDUI R2, R2, 8				WAR I6, R2		
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	16		R1	
RS2	16		R2	
LDU1	16		R1	
LDU2				

	Vj	Qj	Vk	Qk
RS3	F0		F0	
RS4	F2		F2	
FPU1	F0		F0	

	Vj	Qj	Vk	Qk
RS5				
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	4
LDU2	
FPU1	0
ALU1	
ALU2	

9. Cycle 9:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9		WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8		WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9		WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9			WAR I5, R1	RS5	
I8: ADDUI R2, R2, 8				WAR I6, R2		
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	16		R1	
RS2	16		R2	
LDU1	16		R1	
LDU2	16		R2	

	Vj	Qj	Vk	Qk
RS3				
RS4			F2	F2
FPU1		F2		F2

	Vj	Qj	Vk	Qk
RS5		R1		8
RS6				
ALU1				
ALU2				

Unit	Remaining cycles
LDU1	3
LDU2	4
FPU1	4
ALU1	
ALU2	

10. Cycle 10:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9		WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8		WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9		WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10		WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8				WAR I6, R2	RS6	
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	16		R1	
RS2	16		R2	
LDU1	16		R1	
LDU2	16		R2	

	Vj	Qj	Vk	Qk
RS3				
RS4			F2	F2
FPU1			F2	F2

	Vj	Qj	Vk	Qk
RS5	R1		8	
RS6	R2		8	
ALU1	R1		8	
ALU2				

Unit	Remaining cycles
LDU1	2
LDU2	3
FPU1	3
ALU1	2
ALU2	

11. Cycle 11:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9		WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8		WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9		WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10		WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11		WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	16		R1	
RS2	16		R2	
LDU1	16		R1	
LDU2	16		R2	

	Vj	Qj	Vk	Qk
RS3				
RS4			F2	F2
FPU1			F2	F2

	Vj	Qj	Vk	Qk
RS5	R1		8	
RS6	R2		8	
ALU1	R1		8	
ALU2	R2		8	

Unit	Remaining cycles
LDU1	1
LDU2	2
FPU1	2
ALU1	1
ALU2	2

12. **Cycle 12:** “*The instruction 7 has finished executing, and it wants to write the result back. Why can't it do that?*” Because instruction 5 has higher priority, being the first in chronological order, and the common data bus can host only one value at a time.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9		WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9		WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10		WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11		WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1	16		R1	
RS2	16		R2	
LDU1	16		R1	
LDU2	16		R2	

	Vj	Qj	Vk	Qk
RS3				
RS4			F2	F2
FPU1			F2	F2

	Vj	Qj	Vk	Qk
RS5		R1		8
RS6		R2		8
ALU1		R1		8
ALU2		R2		8

Unit	Remaining cycles
LDU1	0
LDU2	1
FPU1	1
ALU1	0
ALU2	1

13. **Cycle 13:** Each functional unit has finished executing and is ready to write back the result. Again, since there is only one common data bus, we issue each instruction in chronological order. In this case, it is instruction 4.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9		WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10		WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11		WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1				
RS2	16		R2	
LDU1				
LDU2	16		R2	

	Vj	Qj	Vk	Qk
RS3				
RS4	F2		F2	
FPU1	F2		F2	

	Vj	Qj	Vk	Qk
RS5	R1		8	
RS6	R2		8	
ALU1	R1		8	
ALU2	R2		8	

Unit	Remaining cycles
LDU1	
LDU2	0
FPU1	0
ALU1	0 (-1)
ALU2	0

14. Cycle 14:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10		WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11		WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1				
RS2	16		R2	
LDU1				
LDU2	16		R2	

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5	R1		8	
RS6	R2		8	
ALU1	R1		8	
ALU2	R2		8	

Unit	Remaining cycles
LDU1	
LDU2	0 (-1)
FPU1	
ALU1	0 (-2)
ALU2	0 (-1)

15. Cycle 15:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10	15	WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11		WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1				RAW R1, R2		
I10: BNE R3, R4, LOOP				RAW I9, R3		

	Vj	Qj	Vk	Qk
RS1				
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5	R1		8	
RS6	R2		8	
ALU1	R1		8	
ALU2	R2		8	

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
ALU1	0 (-3)
ALU2	0 (-2)

16. Cycle 16:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10	15	WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11	16	WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1	16			RAW R1, R2	RS5	
I10: BNE R3, R4, LOOP				RAW R19, R3		

	Vj	Qj	Vk	Qk
RS1				
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5	R2		R1	
RS6	R2		8	
ALU1				
ALU2	R2		8	

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
ALU1	
ALU2	0 (-3)

17. **Cycle 17:** This is the power of Tomasulo and its tag. This algorithm doesn't block the tenth instruction because it lacks the operand R3. Instead, it says, “*You cannot retrieve the value of register R3. No problem. Meanwhile, you can issue the instruction. When the R3 value is ready, it will be sent to the Common Data Bus.*” Thus, instruction 10 waits on the Common Data Bus for the value of register R3. “*But how does it know when the value is ready?*” Simple. Use the tag logic. The value will be sent on the Common Data Bus with the key ALU1. Instruction 10 writes the tag “ALU1” to the reservation station, and when the Common Data Bus sends that tag, it takes the value.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10	15	WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11	16	WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1	16	17		RAW R1, R2	RS5	ALU1
I10: BNE R3, R4, LOOP	17			RAW R1, R3	RS6	

	Vj	Qj	Vk	Qk
RS1				
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5	R2		R1	
RS6		ALU1	R4	
ALU1	R2		R1	
ALU2				

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
ALU1	2
ALU2	

19. **Cycle 19:** We skip cycle 18 because no operations have been performed. In instruction 10, the reservation station reads the value from the common data bus and stores it in memory. It is now ready to be executed!

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10	15	WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11	16	WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1	16	17	19	RAW R1, R2	RS5	ALU1
I10: BNE R3, R4, LOOP	17			RAW I9, R3	RS6	

	Vj	Qj	Vk	Qk
RS1				
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5	R2		R1	
RS6	R3		R4	
ALU1	R2		R1	
ALU2				

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
ALU1	0
ALU2	

20. Cycle 20:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10	15	WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11	16	WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1	16	17	19	RAW R1, R2	RS5	ALU1
I10: BNE R3, R4, LOOP	17	20		RAW I9, R3	RS6	ALU2

	Vj	Qj	Vk	Qk
RS1				
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5				
RS6	R3		R4	
ALU1				
ALU2	R3		R4	

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
ALU1	
ALU2	2

22. Cycle 22:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
I1: SD F4, 0 (R1)	1	2	6	None	RS1	LDU1
I2: SD F6, 0 (R2)	2	3	7	None	RS2	LDU2
I3: FADD F4, F0, F0	3	4	8	WAR I1, F4	RS3	FPU1
I4: FADD F6, F2, F2	4	9	13	WAR I2, F6	RS4	FPU1
I5: LD F0, 16 (R1)	7	8	12	WAR I3, F0	RS1	LDU1
I6: LD F2, 16 (R2)	8	9	14	WAR I4, F2	RS2	LDU2
I7: ADDUI R1, R1, 8	9	10	15	WAR I5, R1	RS5	ALU1
I8: ADDUI R2, R2, 8	10	11	16	WAR I6, R2	RS6	ALU2
I9: ADD R3, R2, R1	16	17	19	RAW R1, R2	RS5	ALU1
I10: BNE R3, R4, LOOP	17	20	22	RAW I9, R3	RS6	ALU2

	Vj	Qj	Vk	Qk
RS1				
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				

	Vj	Qj	Vk	Qk
RS5				
RS6	R3		R4	
ALU1				
ALU2	R3		R4	

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
ALU1	
ALU2	0

Calculate the CPI.

Answer:

$$\text{CPI} = \frac{\# \text{ Clock Cycles}}{\# \text{ Instructions}} = \frac{22}{10} = 2.2$$

Question 3: Loop Unrolling

Let's consider the following loop code where registers **R1** and **R2** are initialized to **0** and **R3** is initialized to **400**:

```

1 LOOP:   LD F2, 0 (R1)
2       LD F4, 0 (R2)
3       FADD F6, F2, F4
4       SD F6, 0 (R1)
5       ADDUI R1, R1, 4
6       ADDUI R2, R2, 4
7       BNE R1, R3, LOOP

```

Answer to the following questions:

- How many iterations of the loop?

Answer: 100 iterations. This is because register R1 is initialized to 0 and is incremented by 4 each iteration until it reaches 400:

$$\text{Iterations} = \frac{400}{4} = 100$$

So the loop executes 100 iterations before R1 equals R3 and the loop exits. Note that the SD F6, 0(R1) instruction stores F6 back into memory at the address in R1; it does not write to R1. Therefore, only the ADDUI operator modifies the R1 register.

- How many instructions per iteration?

Answer: The number of instructions per iteration equals the number of instructions in the loop block, which is 7.

- How many instructions due to the loop overhead per iteration?

Answer: When we write a loop in assembly, **some instructions do useful work, and some instructions are only there to control the loop.**

- **Useful work:** instructions that do the actual computation we want (load data, do the calculation, store result).
- **Overhead:** instructions that *only* exist to keep the loop going; updating counters, changing pointers, checking conditions, branching.

In our case, we have 4 useful operations (2 loads, 1 add, and 1 store), and 3 that *only* exist to make the loop repeat: incrementing pointers/counters (ADDUI R1, ADDUI R2) and testing the loop condition (BNE). Loop overhead is extra work we must do every iteration to **make the loop repeat**, but it does not contribute to the real useful computation.

- How many instructions are executed globally?

Answer: There are 7 instructions for each iteration. Since there are 100 iterations, **700 instructions** (7×100) are executed globally.

- How many branch instructions are executed globally?

Answer: There is 1 branch instruction for each iteration, so 100 branch instructions (1×100) are executed globally.

Let's consider a loop unrolling version of the code with unrolling factor 2.

- *How many iterations of the unrolled loop?*

Answer: Instead of repeating a small body many times, we manually replicate the body multiple times inside the loop; so each iteration does more useful work, but we run the loop fewer times (loop unrolling).

If we unroll by a factor of 2, the loop body is rewritten to do 2 operations per iteration:

```

1 LOOP_UNROLLING_2:    LD F2, 0(R1)
2                      LD F4, 0(R2)
3                      FADD F6, F2, F4
4                      SD F6, 0(R1)
5
6                      LD F8, 4(R1)
7                      LD F10, 4(R2)
8                      FADD F12, F8, F10
9                      SD F12, 4(R1)
10
11                     ADDUI R1, R1, 8
12                     ADDUI R2, R2, 8
13                     BNE R1, R3, LOOP_UNROLLING_2

```

So each iteration now does twice as much useful work and we only need half as many loop iterations:

$$\text{New Iteration Count} = \frac{\text{Original Iteration Count}}{\text{Unrolling Factor}} = \frac{100}{2} = 50$$

- *How many instructions per iteration?*

Answer: The instructions are now 11.

- *How many instructions due to the loop overhead per iteration?*

Answer: There are always 3 because we have unrolled useful work, not overhead work.

- *How many instructions are executed globally?*

Answer: There are 11 instructions for each iteration. We repeat the cycle 50 times, so a total of 550 instructions are executed.

- *How many branch instructions are executed globally?*

Answer: There is only one branch instruction per cycle and 50 iterations, so a total of 50 branch instructions are executed globally.

- *How many more registers are needed due to register renaming?*

Answer: Original loop has 3 FP registers (F2 holds first operand, F4 holds second operand, and F6 holds result). When we unroll the loop by factor 2, we do two sums at the same time. If we use the same registers, we overwrite the result of the first sum before we store it. So we need new registers for the second pair.

Therefore, we can conclude that:

- For the first sum: F2, F4, F6.
- For the second sum: F8, F10, F12.

So we need 3 more FP registers. In general, with an unrolling factor X , we need:

$$(X - 1) \times (\text{number of live registers per body}) \quad \text{more registers}$$

- *What are the registers that need to be renamed?*

Answer: As mentioned earlier, the registers that need to be renamed are F2, F4, and F6.

- *How much is the global instruction count decrease?*

Answer: Before unrolling, the global instructions to execute was 700, and now 550. So we have 150 instructions less than before, that correspond to:

$$700 : 100 = 150 : x \Rightarrow \frac{150 \cdot 100}{700} = 21.428571429 = 21.43\%$$

Let's consider a loop unrolling version of the code with unrolling factor 4.

- *How many iterations of the unrolled loop?*

Answer: If we unroll by a factor of 4, the loop body is rewritten to do 4 operations per iteration:

```

1 LOOP_UNROLLING_4:    LD F2, 0(R1)
2                      LD F4, 0(R2)
3                      FADD F6, F2, F4
4                      SD F6, 0(R1)
5
6                      LD F8, 4(R1)
7                      LD F10, 4(R2)
8                      FADD F12, F8, F10
9                      SD F12, 4(R1)
10
11                     LD F14, 8(R1)
12                     LD F16, 8(R2)
13                     FADD F18, F14, F16
14                     SD F18, 8(R1)
15
16                     LD F20, 12(R1)
17                     LD F22, 12(R2)
18                     FADD F24, F20, F22
19                     SD F24, 12(R1)
20
21                     ADDUI R1, R1, 16
22                     ADDUI R2, R2, 16
23                     BNE R1, R3, LOOP_UNROLLING_4

```

$$\text{New Iteration Count} = \frac{\text{Original Iteration Count}}{\text{Unrolling Factor}} = \frac{100}{4} = 25$$

- *How many instructions per iteration?*

Answer: The instructions are now 19.

- *How many instructions due to the loop overhead per iteration?*

Answer: There are always 3 instructions due to the loop overhead per iteration.

- *How many instructions are executed globally?*

Answer:

$$19 \text{ instructions} \cdot 25 \text{ iterations} = 475$$

- *How many branches are executed?*

Answer: There is always 1 branch for each iteration, so a total of 25 branches are executed.

- *How many more registers are needed due to register renaming?*

Answer:

$$(4 - 1) \cdot 3 = 9 \text{ more registers}$$

- *What are the registers that need to be renamed?*

Answer: It's always the three registers F2, F4, and F6, but three times.

- *How much is the global instruction count decrease?*

Answer:

$$700 : 100 = (700 - 475) : x \Rightarrow \frac{225 \cdot 100}{700} = 32.142857143 = 32.14\%$$

- *What is the best unrolling factor between 2 and 4? Explain why.*

Answer: Factor 4 is preferable to Factor 2 because it reduces loop control overhead per useful operation even further. The number of iterations is halved again, from 50 to 25, and the overhead instructions per sum drop from 1.5 to 0.75.

However, a higher unrolling factor requires more register names due to overlapping live values. In this case, factor 4 increases the demand for floating-point registers from 3 to 12. This increase must be handled by the hardware's register renaming; otherwise, spilling will occur. Therefore, factor 4 is preferable if the register file and rename logic are large enough to store all live values in registers without spilling to memory. Otherwise, factor 2 may be safer.

Quiz 4 - Branch Prediction

During the execution of a program, the number of mispredictions depends on the initialization and the size of the Branch History Table (BHT). True or False? Motivate.

Answer: A BHT stores the *history* (past behavior) of branches, usually with n -bit saturating counters (page 56). The size of the BHT determines **how many unique branches can be tracked independently**

- Small BHT → more aliasing (page 72) → different branches share the same entry → more chance of misprediction.
- Large BHT → fewer collisions → better prediction accuracy.

At program start, the BHT entries are **initialized** (often all taken or all not taken, or all weakly taken). If the actual branch pattern is opposite to the initialization, we'll have **mispredictions** until the counters adapt. So initialization affects startup behavior (how many initial mispredictions we get) and size of BHT affects aliasing (whether unrelated branches interfere with each other, more or fewer mispredictions overall).

The correct answer is: **True**, because the number of mispredictions depends on how the BHT entries are initialized and on the size of the BHT, which determines how many branches can be tracked without aliasing.

Quiz 5 - Branch Prediction

Consider the following assembly code executed by a processor with 1-entry 1-bit Branch History Table initialized as Taken:

```

1      ADDI R1, R0, 0
2      ADDI R2, R0, 100
3
4 LOOP: BEQ R1, R2, DONE
5      ADD R5, R5, R4
6      ADDI R1, R1, 1
7      J LOOP
8
9 DONE:

```

Assuming only the branch instruction accesses the BHT:

- How many accesses to the Branch History Table?
- How many mispredictions?

(Single Answer)

Answer 1: 101 accesses, 2 mispredictions

Answer 2: 100 accesses, 1 misprediction

Answer 3: 99 accesses, 1 misprediction

Answer 4: 201 accesses, 2 mispredictions

Answer: The correct answer is **Answer 1: 101 accesses, 2 mispredictions.** Let's explain why. See also: page 56.

- **How many BHT accesses?** Only the **conditional branch BEQ R1, R2, DONE** accesses the BHT (the J LOOP does **not**, by assumption).
 - R1 starts at 0, R2 = 100.
 - The branch is executed once per loop check for $R1 = 0, 1, 2, \dots, 100$.
 - * For $R1 = 0 \dots 99$: branch outcome is **Not Taken**.
 - * For $R1 = 100$: branch outcome is **Taken** and we exit the loop.

So the number of executions (thus BHT accesses) of the branch is:

$$100 \text{ (for } R1 = 0 \text{ to } 99\text{)} + 1 \text{ (for } R1 = 100\text{)} = 101 \text{ accesses}$$

- **How many mispredictions with a 1-entry, 1-bit BHT initialized as Taken?** A 1-bit BHT remembers the **last outcome** (Taken or Not Taken) of the branch. If correct it stays, if wrong it flips. The initial state is **Predict Taken**. Then, let's walk through the loop:

1. **First time** branch is encountered $R1 = 0$:
 - ✗ Prediction: **Taken** (initial state, written on the exercise text).
 - ✓ Actual outcome: **Not Taken** (since $R1 \neq R2$)
 - ✗ **Misprediction!** Update BHT to **Predict Not Taken** (flip the bit).
2. From $R1 = 1$ to $R1 = 99$ (99 iterations):
 - ✗ Prediction: **Not Taken** (from last update).
 - ✓ Actual outcome: **Not Taken** (since $R1 \neq R2$)
 - ✓ **Correct prediction!** BHT remains **Predict Not Taken**.
3. **Exit check** when $R1 = 100$:
 - ✗ Prediction: **Not Taken** (from last update).
 - ✓ Actual outcome: **Taken** (since $R1 == R2$)
 - ✗ **Misprediction!** Update BHT to **Predict Taken** (flip the bit).

So in total, we have **2 mispredictions** (first encounter and exit check).

The answer is as expected, as it aligns with the **traditional limitation of 1-bit predictors**. They often overlook a “mostly one direction” pattern when the outcome changes. In this case, the outcome is first *Not Taken* and then *Taken*. This results in two misses in this specific scenario: one at the beginning due to poor initialization and one at the final change.

Quiz 6 - Superscalar Processors

Considering superscalar processors, answer **TRUE** or **FALSE** to the following questions, **motivating your answers**.

- Doubling the number of issues reduces the CPI metric.

Answer: True. In superscalar processors, the CPI (Cycles Per Instruction) is ideally:

$$\text{CPI}_{\text{ideal}} = \frac{1}{\text{issue width}}$$

So if we double the issue width, theoretical CPI reduces. This is true only under the assumption of perfect conditions: infinite ILP, no hazards, no cache misses, no branch mispredictions, perfect decode, dispatch, execution bandwidth.

- Out-of-order execution can be generated by data cache misses.

Answer: True. Cache misses introduce stalls for dependent instructions; out-of-order execution allows independent instructions to bypass stalled loads, maximizing utilization of execution units. In other words, data cache misses do **not prevent** execution; in an out-of-order superscalar processor, they often **enable visible out-of-order behavior** by stalling one instruction and allowing others to bypass it.

Deepening: Motivation

In a **superscalar processor** (page 95), instructions are fetched and decoded in program order, but, if the microarchitecture supports it, they may **execute out of order** as long as **data and control dependencies are respected**.

A **data cache miss** (page 272, e.g., a load that misses in L1/L2) introduces a **long and unpredictable latency**. When this happens:

- The instruction that triggered the miss **cannot complete** (it is stalled waiting for data from lower levels of the memory hierarchy).
- Independent instructions that are **already in the instruction window / reservation stations** and **do not depend on the missing load** can still be: issued, executed, and possibly completed **before** the stalled load completes.

This behavior is precisely **out-of-order execution**: instructions that appear *later* in program order are executed *earlier* in time.

From a hardware point of view:

- The cache miss **creates a stall in one dependency chain**.
- The **dynamic scheduler** (e.g., Tomasulo logic with reservation stations and register renaming) exploits **instruction-**

level parallelism (ILP) by continuing execution on other ready instructions.

Therefore, although the cache miss itself is not a scheduling mechanism, it is a **triggering event** that **induces out-of-order execution** by freeing the scheduler to execute independent instructions while the miss is serviced.

- *The introduction of dynamic branch predictors improves the CPI metric.*

Answer: **True.** Dynamic branch predictors reduce the number of branch mispredictions by learning from past branch behavior. Fewer mispredictions mean fewer pipeline flushes and stalls, which directly lowers the average Cycles Per Instruction (CPI) and improves overall processor performance.

Deepening: Motivation

The **Cycles Per Instruction (CPI)** metric captures both the ideal execution rate and the penalties due to pipeline hazards, including **control hazards** caused by branches (page 32). In a pipelined or superscalar processor:

- A **branch misprediction** causes:
 - * Pipeline flushing (all instructions after the mispredicted branch are discarded).
 - * Squashing of wrong-path instructions (i.e., instructions that were fetched and possibly executed based on the incorrect prediction are invalidated).
 - * And a **stall penalty** proportional to pipeline depth and issue width.
- This penalty directly **increases CPI**, since extra cycles are spent without retiring useful instructions.

A **dynamic branch predictor** (e.g., 1-bit, 2-bit saturating counters, local/global history predictors):

- Learns branch behavior at runtime.
- Increases **prediction accuracy** compared to static schemes (e.g., always taken / always not taken).
- **Reduces the misprediction rate.**

Since CPI can be expressed as:

$$\text{CPI} = \text{CPI}_{\text{ideal}} + \text{Branch frequency} \times \text{Misp. rate} \times \text{Misp. penalty}$$

Improving prediction accuracy **reduces the branch penalty term**, thus **lowering CPI**. Therefore, introducing **dynamic branch predictors** improves (i.e., reduces) the CPI metric, assuming: branches are present (which is almost always true), and the predictor does not introduce prohibitive overheads.

11.2.2 February 11

Exercise 1 - Dependency Analysis + Tomasulo

Let's consider the following assembly code containing multiple types of intra-loop dependences. Complete the following table by inserting all types of true-data-dependences, anti-dependences and output dependences for each instruction:

Answer:

I#	Type of instruction	Analysis of dependences:
		1. True data dependence with I# for \$Fx 2. Anti-dependence with I# for \$Fy 3. Output-dependence with I# for \$Fz
I0	LD \$F0, A(\$R0)	None
I1	LD \$F2, B(\$R0)	None
I2	FADD \$F4, \$F0, \$F2	True data dependence with I0 for \$F0 True data dependence with I1 for \$F2
I3	FADD \$F6, \$F4, \$F4	True data dependence with I2 for \$F4
I4	SD \$F4, C(\$R0)	True data dependence with I2 for \$F4
I5	SD \$F6, D(\$R0)	True data dependence with I3 for \$F6
I6	ADDI \$R0, \$R0, 4	Anti-dependence with I0 for \$R0 Anti-dependence with I1 for \$R0 Anti-dependence with I4 for \$R0 Anti-dependence with I5 for \$R0
I7	BNE \$R0, \$R1, LOOP	True data dependence with I6 for \$R0

Let's consider the previous assembly code to be executed on a CPU with dynamic scheduling based on TOMASULO algorithm with all cache HITS, a single Common Data Bus and:

- 2 RESERVATION STATIONS (**RS1, RS2**) with 2 LOAD/STORE units (**LDU1, LDU2**) with latency **6**
- 2 RESERVATION STATION (**RS3, RS4**) with 2 FP unit12 (**FPU1, FPU2**) with latency **2**
- 1 RESERVATION STATION (**RS5**) with 1 INT_ALU/BR unit (**ALU1**) with latency **1**

Answer:

1. Cycle 1:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2						
FADD \$F6, \$F4, \$F4						
SD \$F4, C(\$R0)						
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1	A		\$R0	
RS2				
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				
FPU2				

	Vj	Qj	Vk	Qk
RS5				
ALU1				

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	
FPU2	
ALU1	

2. Cycle 2:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2						
FADD \$F6, \$F4, \$F4						
SD \$F4, C(\$R0)						
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1	A		\$R0	
RS2	B		\$R0	
LDU1	A		\$R0	
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				
FPU2				

	Vj	Qj	Vk	Qk
RS5				
ALU1				

Unit	Remaining cycles
LDU1	6
LDU2	
FPU1	
FPU2	
ALU1	

3. **Cycle 3:** Here, instruction 2 has a RAW hazard with instructions 0 and 1 for registers F0 and F2. Therefore, it uses the tag mechanism to wait for their values. In other words, the instruction is issued, but will not start until the values of F0 and F2 are forwarded to the Common Data Bus.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3			RAW IO, I1	RS3	
FADD \$F6, \$F4, \$F4						
SD \$F4, C(\$R0)						
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1	A		\$R0	
RS2	B		\$R0	
LDU1	A		\$R0	
LDU2	B		\$R0	

	Vj	Qj	Vk	Qk
RS3			LDU1	LDU2
RS4				
FPU1				
FPU2				

Unit	Remaining cycles
LDU1	5
LDU2	6
FPU1	
FPU2	
ALU1	

4. **Cycle 4:** It's a similar situation to before. The second addition cannot be completed because it requires the result of the previous instruction. Therefore, it waits for the result on the common data bus.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3			RAW I0, I1	RS3	
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)						
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1	A		\$R0	
RS2	B		\$R0	
LDU1	A		\$R0	
LDU2	B		\$R0	

	Vj	Qj	Vk	Qk
RS3			LDU1	LDU2
RS4			FPU1	FPU1
FPU1				
FPU2				

	Vj	Qj	Vk	Qk
RS5				
ALU1				

Unit	Remaining cycles
LDU1	4
LDU2	5
FPU1	
FPU2	
ALU1	

5. **Cycle 5:** The fourth instruction cannot be issued here because there are no reservation stations available, which constitutes a structural hazard. This means Tomasulo stops execution because it guarantees in-order issuance, meaning the instructions are issued in order. Some updates will be seen at the end of the first instruction (cycle 8).

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3			RAW I0, I1	RS3	
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)				STRUCTURAL		
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1	A		\$R0	
RS2	B		\$R0	
LDU1	A		\$R0	
LDU2	B		\$R0	

	Vj	Qj	Vk	Qk
RS3			LDU1	LDU2
RS4			FPU1	FPU1
FPU1				
FPU2				

	Vj	Qj	Vk	Qk
RS5				
ALU1				

Unit	Remaining cycles
LDU1	3
LDU2	4
FPU1	
FPU2	
ALU1	

8. Cycle 8:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3			RAW I0, I1	RS3	
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)				STRUCTURAL		
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1	A		\$R0	
RS2	B		\$R0	
LDU1	A		\$R0	
LDU2	B		\$R0	

	Vj	Qj	Vk	Qk
RS3			\$F0	LDU2
RS4				FPU1
FPU1				FPU1
FPU2				

	Vj	Qj	Vk	Qk
RS5				
ALU1				

Unit	Remaining cycles
LDU1	0
LDU2	1
FPU1	
FPU2	
ALU1	

9. Cycle 9:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3			RAW I0, I1	RS3	
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)	9			STR., RAW I2, \$F4	RS1	
SD \$F6, D(\$R0)						
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1		RS3	\$R0	
RS2	B		\$R0	
LDU1				
LDU2	B		\$R0	

	Vj	Qj	Vk	Qk
RS3	\$F0		\$F2	
RS4		FPU1		FPU1
FPU1				
FPU2				

Unit	Remaining cycles
LDU1	
LDU2	0
FPU1	
FPU2	
ALU1	

10. Cycle 10:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10		RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)	9			STR., RAW I2, \$F4	RS1	
SD \$F6, D(\$R0)	10			RAW I3, \$F6	RS2	
ADDI \$R0, \$R0, 4						
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1		RS3	\$R0	
RS2		RS4	\$R0	
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3	\$F0		\$F2	
RS4		FPU1		FPU1
FPU1	\$F0		\$F2	
FPU2				

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	2
FPU2	
ALU1	

11. Cycle 11:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10		RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)	9			STR., RAW I2, \$F4	RS1	
SD \$F6, D(\$R0)	10			RAW I3, \$F6	RS2	
ADDI \$R0, \$R0, 4	11				RS5	
BNE \$R0, \$R1, LOOP						

	Vj	Qj	Vk	Qk
RS1		RS3	\$R0	
RS2		RS4	\$R0	
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3	\$F0		\$F2	
RS4		FPU1		FPU1
FPU1	\$F0		\$F2	
FPU2				

	Vj	Qj	Vk	Qk
RS5	\$R0		4	
ALU1				

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	1
FPU2	
ALU1	

12. **Cycle 12:** The last operation cannot be issued due to a structural hazard in the ALU1 functional unit. Note: The professor's solution has an error. The structural hazard must be on the RS5 because it is the only one capable of performing the operation.

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4			RAW I2, F6	RS4	
SD \$F4, C(\$R0)	9			STR., RAW I2, \$F4	RS1	
SD \$F6, D(\$R0)	10			RAW I3, \$F6	RS2	
ADDI \$R0, \$R0, 4	11	12			RS5	ALU1
BNE \$R0, \$R1, LOOP				STRUCTURAL		

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2		RS4	\$R0	
LDU1				
LDU2				

	Vj	Qj	Vk	Qk
RS3	\$F0		\$F2	
RS4	\$F4		\$F4	
FPU1	\$F0		\$F2	
FPU2				

	Vj	Qj	Vk	Qk
RS5	\$R0		4	
ALU1	\$R0		4	

Unit	Remaining cycles
LDU1	
LDU2	
FPU1	0
FPU2	
ALU1	1

13. Cycle 13:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13		RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13		STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10			RAW I3, \$F6	RS2	
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP				STRUCTURAL		

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2		RS4	\$R0	
LDU1	\$F4		\$R0	
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4	\$F4		\$F4	
FPU1				
FPU2	\$F4		\$F4	

	Vj	Qj	Vk	Qk
RS5	\$R0		4	
ALU1	\$R0		4	

Unit	Remaining cycles
LDU1	6
LDU2	
FPU1	
FPU2	2
ALU1	0

14. Cycle 14:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13		RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13		STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10			RAW I3, \$F6	RS2	
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14			STRUCTURAL	RS5	

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2		RS4	\$R0	
LDU1	\$F4		\$R0	
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4	\$F4		\$F4	
FPU1				
FPU2	\$F4		\$F4	

Unit	Remaining cycles
LDU1	5
LDU2	
FPU1	
FPU2	1
ALU1	

15. Cycle 15:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13	15	RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13		STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10			RAW I3, \$F6	RS2	
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14	15		STRUCTURAL	RS5	ALU1

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2	\$F4		\$R0	
LDU1	\$F4		\$R0	
LDU2				

	Vj	Qj	Vk	Qk
RS3				
RS4	\$F4		\$F4	
FPU1				
FPU2	\$F4		\$F4	

	Vj	Qj	Vk	Qk
RS5	\$R0		\$R1	
ALU1	\$R0		\$R1	

Unit	Remaining cycles
LDU1	4
LDU2	
FPU1	
FPU2	0
ALU1	1

16. Cycle 16:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13	15	RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13		STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10	16		RAW I3, \$F6	RS2	LDU2
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14	15	16	STRUCTURAL	RS5	ALU1

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2	\$F4		\$R0	
LDU1	\$F4		\$R0	
LDU2	\$F4		\$R0	

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				
FPU2				

	Vj	Qj	Vk	Qk
RS5		\$R0		\$R1
ALU1		\$R0		\$R1

Unit	Remaining cycles
LDU1	3
LDU2	6
FPU1	
FPU2	
ALU1	0

17. Cycle 17:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13	15	RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13		STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10	16		RAW I3, \$F6	RS2	LDU2
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14	15	16	STRUCTURAL	RS5	ALU1

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2	\$F4		\$R0	
LDU1	\$F4		\$R0	
LDU2	\$F4		\$R0	

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				
FPU2				

	Unit	Remaining cycles
RS5	LDU1	2
ALU1	LDU2	5
	FPU1	
	FPU2	
	ALU1	

19. Cycle 19:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13	15	RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13	19	STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10	16		RAW I3, \$F6	RS2	LDU2
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14	15	16	STRUCTURAL	RS5	ALU1

	Vj	Qj	Vk	Qk
RS1	\$F4		\$R0	
RS2	\$F4		\$R0	
LDU1	\$F4		\$R0	
LDU2	\$F4		\$R0	

	Vj	Qj	Vk	Qk
RS3				
RS4				
FPU1				
FPU2				

	Unit	Remaining cycles
RS5	LDU1	0
ALU1	LDU2	3
	FPU1	
	FPU2	
	ALU1	

22. Cycle 22:

Instruction	Issue	Start Exec	Write Res	Hazard	RSi	Unit
LD \$F0, A(\$R0)	1	2	8	None	RS1	LDU1
LD \$F2, B(\$R0)	2	3	9	None	RS2	LDU2
FADD \$F4, \$F0, \$F2	3	10	12	RAW I0, I1	RS3	FPU1
FADD \$F6, \$F4, \$F4	4	13	15	RAW I2, F6	RS4	FPU2
SD \$F4, C(\$R0)	9	13	19	STR., RAW I2, \$F4	RS1	LDU1
SD \$F6, D(\$R0)	10	16	22	RAW I3, \$F6	RS2	LDU2
ADDI \$R0, \$R0, 4	11	12	13		RS5	ALU1
BNE \$R0, \$R1, LOOP	14	15	16	STRUCTURAL	RS5	ALU1

	Vj	Qj	Vk	Qk		Vj	Qj	Vk	Qk
RS1					RS3				
RS2	\$F4		\$R0		RS4				
LDU1					FPU1				
LDU2	\$F4		\$R0		FPU2				
					Unit	Remaining cycles			
	Vj	Qj	Vk	Qk	LDU1				
RS5					LDU2	0			
ALU1					FPU1				
					FPU2				
					ALU1				

Calculate the **CPI**.

Answer:

$$\text{CPI} = \frac{\# \text{ Clock Cycles}}{\# \text{ Instructions}} = \frac{22}{8} = 2.75$$

Exercise 2 - VLIW Scheduling

Let's consider the following LOOP code:

```

1 LOOP:   LD F1, A(R1)
2           LD F2, A(R2)
3           LD F3, A(R3)
4           FADD F1, F1, F2
5           FADD F2, F2, F3
6           FMUL F1, F1, F2
7           FADD F3, F3, F3
8           ADDUI R2, R1, 8
9           ADDUI R3, R1, 8
10          SD F1, B(R1)
11          ADDUI R1, R1, 4
12          BNE R1, R6, LOOP

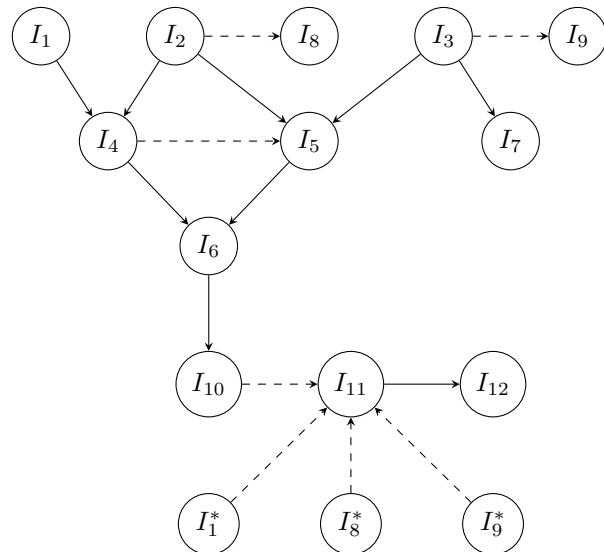
```

Given a **3-issue VLIW machine with fully pipelined functional units**:

- **1 Memory Unit with 3 cycles latency**
- **1 FP ALU with 2 cycles latency**
- **1 Integer ALU with 1 cycle latency to next Int/FP & 2 cycle latency to next Branch**

The branch is completed with 1 cycle delay slot (branch solved in ID stage).
No branch prediction. In the Register File, it is possible to read and write at the same address at the same clock cycle. *Considering one iteration of the loop, complete the following table by using the list-based scheduling* (do NOT introduce any software pipelining, loop unrolling and modifications to loop indexes) on the *4-issue VLIW machine including the BRANCH DELAY SLOT*. Please do not write in NOPs.

Answer:



	Memory Unit	Floating Point Unit	Integer Unit
C1	LD F1, A(R1)		
C2	LD F2, A(R2)		ADDUI R2, R1, 8
C3	LD F3, A(R3)		ADDUI R3, R1, 8
C4			
C5		FADD F1, F1, F2	
C6		FADD F2, F2, F3	
C7		FADD F3, F3, F3	
C8		FMUL F1, F1, F2	
C9			
C10	SD F1, B(R1)		ADDUI R1, R1, 4
C11			
C12			BNE R1, R2, LOOP
C13			(delay slot)
C14			
C15			

1. *How long is the critical path?*

Answer: The critical path is determined by the latency of the last instruction, BNE, so it is **13**.

2. *How much is the code efficiency?*

Answer: Efficiency is measured as follows:

$$\text{Efficiency} = \frac{\text{Instruction Count}}{\# \text{ cycles} \times \# \text{ issues}} = \frac{12}{13 \cdot 3} = 0.307692308 = 30.77\%$$

Exercise 3 - MESI Protocol

Let's consider the following access patterns on a **4-processor** system with a direct-mapped, write-back cache with one cache block per processor and a two-block memory.

Assume the **MESI** protocol is used with **write-back** caches, **write-allocate**, and **write-invalidate** of other caches. Please COMPLETE the following table:

Cycle	After Op.	P0	P1	P2	P3	Mem.	Mem.
		cache	cache	cache	cache	at bl. 0 up	at bl. 1 up
		block	block	block	block	to date?	to date?
state	state	state	state	state	state		
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes
3	P3: Read Bl. 0						
4	P1: Write Bl. 0						
5	P0: Write Bl. 1						
6	P3: Read Bl. 1						

Answer: **MESI** is a **cache coherence protocol**. It guarantees that multiple caches in a **shared-memory multiprocessor** see a **coherent view** of main memory.

❷ **Why is it needed?** When multiple processors have private caches and share a memory space, they may each hold *local copies* of the same memory block. If one processor modifies its copy, the other copies may become **stale** (incoherent). The MESI protocol ensures that all copies are kept consistent, meaning we never use an old value when we shouldn't.

■ **MESI States** Each block in a cache can be in one of **4 states**:

- **M (Modified)**: Block is valid, *dirty* (modified). It *differs* from main memory. This cache *owns* the latest value.
- **E (Exclusive)**: Block is valid, *clean*. It matches memory. This cache is the *only* owner.
- **S (Shared)**: Block is valid, clean, and may be in *multiple* caches. All have the same value as memory.
- **I (Invalid)**: Block is *not* valid in this cache. Must be fetched on access.

Coherence is enforced by invalidating or updating other caches:

- If a cache writes to a block, it must invalidate all other caches that have that block (**Write-Invalidate**).
- If a cache needs a block held dirty by another, that cache must supply the up-to-date data (and possibly write back to memory).

② **What is Direct-Mapped Cache?** **Direct-Mapped** is the simplest cache mapping:

- Each memory block maps to **exactly one** cache block.
- There's only one place in the cache where a given memory block can go.

In this exercise, each processor has a direct-mapped cache with **only one block**, so each processor can store **only one block at a time**.

③ **What does Write-Back mean?** Caches can be:

- **Write-Through:** Every write to the cache is also immediately written to main memory. Simpler, but more traffic.
- **Write-Back:** Writes go only to the cache at first. The block is marked *dirty*. The updated block is written back to main memory only when it is *evicted* or needed by another processor.

In MESI: The **Modified (M)** state marks blocks that are *dirty*. These blocks must be written back before eviction or when another processor requests them.

④ **What does Write-Allocate mean?** **Write-Allocate** means that when a processor writes to a block that is not present in its cache (*write miss*), it:

- **Allocates:** It *loads* the block into its cache.
- Then performs the write in the cache.

This is typical with write-back caches: the block must be in the cache so that it can later be written back to memory if needed.

⑤ **What does “One cache block per processor” mean?** In this exercise:

- Each processor has **one slot in its cache**.
- That slot can hold *one memory block* at a time.
- So if a processor needs to read/write a new block, it may need to *evict* the old one (and write it back if it's dirty).

⑥ **What does “Two-cache-block memory” mean?** This means that **main memory** has **only 2 data blocks** available: **block 0** and **block 1**. So:

- The whole example focuses only on those two blocks.
- Any cache action involves either block 0 or block 1.
- The goal is to track how these two blocks move between the processors' caches and the main memory.

In summary, the entire scenario is as follows:

- A *tiny system* with 4 processors.
 - Each has a *tiny cache* (1 block).
 - The system has only 2 memory blocks.
 - Using MESI, caches cooperate to keep shared blocks consistent.
 - *Write-back*: modified blocks may not immediately update memory.
 - *Write-allocate*: writes on misses bring the block in.
0. In the **initial state**, everything is up to date, but each block is invalid.

1. **Operation 1**: “P0: Read block 1”. Since all the blocks are invalid, this is a *cache miss* operation. Therefore, processor 0 (P0) loads block 1 from memory. Since no one else has it, the state is **Exclusive**. The memory is still up to date because it was “*read-by*”, not “*written-to*”.

Cycle	After Op.	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
0		Invalid	Invalid	Invalid	Invalid	Yes	Yes
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes

2. **Operation 2**: “P2: Read Bl. 0”. Processor 2 has an invalid cache block, resulting in a *cache miss*. Since nobody else has block 0, processor 2 loads it. It gets it in **Exclusive** mode. Again, the memory is up to date because it was read.

Cycle	After Op.	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
0		Invalid	Invalid	Invalid	Invalid	Yes	Yes
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes

3. **Operation 3:** “P3: Read Bl. 0”. Processor 3 has an invalid cache block, resulting in a *cache miss*. This time, however, a copy of block 0 is in processor 2’s cache, so processor 3 can request a copy of it. Now, since the value of block 0 is shared by different caches, we transition from **Exclusive** to **Shared**. Finally, the memory remains updated (no writes).

Cycle	After Op.	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
0		Invalid	Invalid	Invalid	Invalid	Yes	Yes
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes
3	P3: Read Bl. 0	Excl (1)	Invalid	Shr (0)	Shr (0)	Yes	Yes

4. **Operation 4:** “P1: Write Bl. 0”. Since Processor 1 does not have Block 0, there is a cache miss. Block 0 is in processors 2 and 3’s caches in Shared mode. So, MESI:

- (a) P1 needs exclusive access, so it invalidates all other copies.
- (b) P2 and P3 are invalidated by P1 and become invalid.
- (c) P1 loads block 0, writes it, and transitions from **Exclusive** to **Modified**.

Since it’s write-back, the memory for block 0 becomes stale (marked as out-of-date).

Cycle	After Op.	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
0		Invalid	Invalid	Invalid	Invalid	Yes	Yes
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes
3	P3: Read Bl. 0	Excl (1)	Invalid	Shr (0)	Shr (0)	Yes	Yes
4	P1: Write Bl. 0	Excl (1)	Mod (0)	Invalid	Invalid	No	Yes

5. **Operation 5:** “P0: Write Bl. 1”. Processor 0 already has Block 1 in **Exclusive** mode. When it writes, it switches from Exclusive to **Modified** mode. Because a write was made, the memory is no longer up to date. Since no one has block 1, it doesn’t invalidate the other caches.

Cycle	After Op.	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
0		Invalid	Invalid	Invalid	Invalid	Yes	Yes
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes
3	P3: Read Bl. 0	Excl (1)	Invalid	Shr (0)	Shr (0)	Yes	Yes
4	P1: Write Bl. 0	Excl (1)	Mod (0)	Invalid	Invalid	No	Yes
5	P0: Write Bl. 1	Mod (1)	Mod (0)	Invalid	Invalid	No	No

6. **Operation 6:** “P3: Read Bl. 1”. Since processor 3 does not have block 1, there is a *cache miss*. Processor 0 has block 1 in modified mode.

- (a) Another processor needs the data.
- (b) P0 must supply the *latest value*, so P0 does *write-back* or sends directly.
- (c) Result: P3 loads up-to-date block, and both now have it in **Shared**.

Now the memory for block 1 is clean again (assuming P0 did a write-back).

Cycle	After Op.	P0 cache block state	P1 cache block state	P2 cache block state	P3 cache block state	Mem. at bl. 0 up to date?	Mem. at bl. 1 up to date?
0		Invalid	Invalid	Invalid	Invalid	Yes	Yes
1	P0: Read Bl. 1	Excl (1)	Invalid	Invalid	Invalid	Yes	Yes
2	P2: Read Bl. 0	Excl (1)	Invalid	Excl (0)	Invalid	Yes	Yes
3	P3: Read Bl. 0	Excl (1)	Invalid	Shr (0)	Shr (0)	Yes	Yes
4	P1: Write Bl. 0	Excl (1)	Mod (0)	Invalid	Invalid	No	Yes
5	P0: Write Bl. 1	Mod (1)	Mod (0)	Invalid	Invalid	No	No
6	P3: Read Bl. 1	Shr (1)	Mod (0)	Invalid	Shr (1)	No	Yes

Note: Read from **Modified** forces a write-back (or direct transfer) to maintain consistency.

Exercise 4 - Reordered Buffer

Let's consider the following assembly loop where registers \$R1 and \$R2 are initialized at 0 and 40 respectively:

```

1 L0:      LD $F2, 0 ($R1)          # datapath: load data
2 L1:      ADDD $F4, $F2, $F2       # datapath: compute
3 L2:      SD $F4, 0 ($R1)          # datapath: store data
4 L3:      ADDI $R1, $R1, 4         # control: index update
5 L4:      BNE $R1, $R2, L0        # control: conditional branch

```

- How many loop iterations?

Answer: The register R1 is incremented by 4 at each iteration:

$$\text{Iterations} = \frac{40}{4} = 10$$

- How many instructions per iteration?

Answer: 5 instructions.

- How many control vs datapath instructions?

Answer: To answer the question, we count how many instructions in the loop control the flow vs how many do actual data processing. In our loop:

```

1 L0:      LD $F2, 0 ($R1)          # datapath: load data
2 L1:      ADDD $F4, $F2, $F2       # datapath: compute
3 L2:      SD $F4, 0 ($R1)          # datapath: store data
4 L3:      ADDI $R1, $R1, 4         # control: index update
5 L4:      BNE $R1, $R2, L0        # control: conditional branch

```

- Data path instructions: LD, ADDD, SD, 3 instructions.
- Control instructions: ADDI (updates loop index), BNE (branch), 2 instructions.

So, the answer is 2 control vs 3 datapath instructions. In general:

- Control instructions: branches, jumps, index updates.
- Datapath instructions: load, store, ALU/FPU operations that compute or move data.

Write the unrolled version of the loop with unrolling factor 2 by using Register Renaming.

Answer:

```

1 L0:    LD $F2, 0($R1)      # datapath: load data
2 L1:    ADDD $F4, $F2, $F2    # datapath: compute
3 L2:    SD $F4, 0($R1)      # datapath: store data
4 L3:    LD $F6, 4($R1)      # datapath: load data
5 L4:    ADDD $F8, $F6, $F6    # datapath: compute
6 L5:    SD $F8, 0($R1)      # datapath: store data
7 L6:    ADDI $R1, $R1, 8      # control: index update
8 L7:    BNE $R1, $R2, L0      # control: conditional branch

```

- How many loop iterations?

Answer: The register R1 is incremented by 8 at each iteration:

$$\text{Iterations} = \frac{40}{8} = 5$$

- How many instructions per iteration?

Answer: 8 instructions.

- How many control vs datapath instructions?

Answer: 2 control vs 6 datapath instructions.

Execute the unrolled version of the loop by the **Speculative Tomasulo** architecture with a 10-entry ROB and:

- 4 Load Buffers (Load1, Load2, Load3, Load4);
- 4 FP Reservation Stations (FP1, FP2, FP3, FP4)
- 2 Integer Reservation Stations (Int1, Int2)

Complete the ROB and the Rename Table until the ROB becomes **full** while the first instruction is still in execution due to a cache miss (*):

ROB#	Instruction	Dest.	Res.	Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1		No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1		No, issued	No
ROB2						
ROB3						
ROB4						
ROB5						
ROB6						
ROB7						
ROB8						
ROB9						

ROB Table

Reg #	ROB #
\$F0	-
\$F2	
\$F4	
\$F6	
\$F8	

*Rename Table***Answer:**

1. First, we allocate the entries given by the exercise in the Rename Table:

Reg #	ROB #
\$F0	-
\$F2	ROB0
\$F4	ROB1
\$F6	
\$F8	

Rename Table

2. We issue the L2 instruction. The operation's destination is memory. The status has not been issued yet because we are filling the ROB due to a cache miss, and the processor is waiting for the result of the operation.

ROB#	Instruction	Dest.	Res.	Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1		No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1		No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-		No, issued	No
ROB3						
ROB4						
ROB5						
ROB6						
ROB7						
ROB8						
ROB9						

ROB Table

3. We issue the L3 instruction. The destination is the left operand, register F6. We use Load 2 as the functional unit.

ROB#	Instruction	Dest.	Res.	Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1		No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1		No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-		No, issued	No
ROB3	L3: LD \$F6, 4(\$R1)	\$F6	Load2		No, issued	No
ROB4						
ROB5						
ROB6						
ROB7						
ROB8						
ROB9						

ROB Table

Reg #	ROB #
\$F0	-
\$F2	ROB0
\$F4	ROB1
\$F6	ROB3
\$F8	

Rename Table

4. We issue the L4 instruction.

ROB#	Instruction	Dest.	Res. Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1	No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1	No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-	No, issued	No
ROB3	L3: LD \$F6, 4(\$R1)	\$F6	Load2	No, issued	No
ROB4	L4: ADDD \$F8, \$F6, \$F6	\$F8	FP2	No, issued	No
ROB5					
ROB6					
ROB7					
ROB8					
ROB9					

ROB Table

Reg #	ROB #
\$F0	-
\$F2	ROB0
\$F4	ROB1
\$F6	ROB3
\$F8	ROB4

Rename Table

5. We issue the L5 instruction.

ROB#	Instruction	Dest.	Res. Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1	No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1	No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-	No, issued	No
ROB3	L3: LD \$F6, 4(\$R1)	\$F6	Load2	No, issued	No
ROB4	L4: ADDD \$F8, \$F6, \$F6	\$F8	FP2	No, issued	No
ROB5	L5: SD \$F8, 0(\$R1)	Mem	-	No, issued	No
ROB6					
ROB7					
ROB8					
ROB9					

ROB Table

6. We issue the L6 instruction.

ROB#	Instruction	Dest.	Res.	Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1		No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1		No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-		No, issued	No
ROB3	L3: LD \$F6, 4(\$R1)	\$F6	Load2		No, issued	No
ROB4	L4: ADDD \$F8, \$F6, \$F6	\$F8	FP2		No, issued	No
ROB5	L5: SD \$F8, 0(\$R1)	Mem	-		No, issued	No
ROB6	L6: ADDI \$R1, \$R1, 8	\$R1	Int1		No, issued	No
ROB7						
ROB8						
ROB9						

ROB Table

7. We issue the L7 instruction.

ROB#	Instruction	Dest.	Res.	Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1		No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1		No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-		No, issued	No
ROB3	L3: LD \$F6, 4(\$R1)	\$F6	Load2		No, issued	No
ROB4	L4: ADDD \$F8, \$F6, \$F6	\$F8	FP2		No, issued	No
ROB5	L5: SD \$F8, 0(\$R1)	Mem	-		No, issued	No
ROB6	L6: ADDI \$R1, \$R1, 8	\$R1	Int1		No, issued	No
ROB7	L7: BNE \$R1, \$R2, L0	-	Int2		No, issued	No
ROB8						
ROB9						

ROB Table

8. We issue L0 and L1. Since we need to fill out the entire ROB table and we don't know whether the branch will be true or false, we made a prediction. Therefore, each speculative entry will be set to true.

ROB#	Instruction	Dest.	Res.	Alloc.	Ready/Status	Spec.
ROB0	L0: LD \$F2, 0 (\$R1)	\$F2	Load1		No, exec.(*)	No
ROB1	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP1		No, issued	No
ROB2	L2: SD \$F4, 0(\$R1)	Mem	-		No, issued	No
ROB3	L3: LD \$F6, 4(\$R1)	\$F6	Load2		No, issued	No
ROB4	L4: ADDD \$F8, \$F6, \$F6	\$F8	FP2		No, issued	No
ROB5	L5: SD \$F8, 0(\$R1)	Mem	-		No, issued	No
ROB6	L6: ADDI \$R1, \$R1, 8	\$R1	Int1		No, issued	No
ROB7	L7: BNE \$R1, \$R2, L0	-	Int2		No, issued	No
ROB8	L0: LD \$F2, 0(\$R1)	\$F2	Load3		No, issued	Yes
ROB9	L1: ADDD \$F4, \$F2, \$F2	\$F4	FP3		No, issued	Yes

ROB Table

Reg #	ROB #
\$F0	-
\$F2	ROB0 ROB8
\$F4	ROB1 ROB9
\$F6	ROB3
\$F8	ROB4

Rename Table

Question 1: Instruction-Level and Thread-Level Parallelism

Modern processors exploit **Instruction Level Parallelism** and **Thread Level Parallelism**. *Answer to the following questions:*

- Explain the main concepts for each approach (Instruction Level Parallelism ILP, Thread Level Parallelism TLP).

Answer:

– **Instruction-Level Parallelism (ILP)**. ILP means **executing multiple independent instructions from a single instruction stream in parallel**. The goal is to exploit *fine-grain* parallelism *within* a single thread.

Many instructions in a program do not depend on each other, so they can be executed simultaneously if the hardware can exploit this.

- * **Pipelining**: Overlap the execution of multiple instructions by dividing execution into stages (fetch, decode, execute, memory, write-back). While one instruction is being decoded, the next can be fetched, and so on.
- * **Superscalar execution**: Issue multiple instructions per clock cycle. A superscalar processor has multiple pipelines.
- * **Out-of-order execution**: Hardware reorders instructions dynamically to maximize parallelism while preserving correctness.
- * **Speculation**: Predict branches and execute along the predicted path before knowing if it is correct.
- * **Register renaming**: Remove false dependencies (WAW, WAR hazards) by dynamically mapping logical registers to physical registers.

ILP aims to increase *single-thread performance*. Performance improvement depends on the amount of independent instructions and on how well the hardware can extract and schedule them.

– **Thread-Level Parallelism (TLP)**. TLP means **executing multiple independent instruction streams (threads) in parallel**. This exploits coarse-grain parallelism between threads.

Many programs or tasks can be decomposed into multiple independent threads. If there are multiple cores, these threads can run simultaneously.

- * **Multithreading**: Hardware executes instructions from multiple threads on a single core to better utilize pipeline resources (e.g., SMT, Simultaneous Multithreading).
- * **Multiprocessing (Multicore)**: Multiple cores, each with its own pipeline(s), execute different threads.
- * **Clusters or Shared Memory Multiprocessors**: Multiple processors/cores share memory; they communicate and synchronize to solve larger problems.

To coordinate threads, programmers use synchronization primitives (locks, barriers). Memory consistency models define how memory updates become visible.

TLP increases *throughput* and/or *responsiveness*. It is the main driver behind modern multicore processors. It shifts the burden to software: the program must be parallelized.

Aspect	ILP	TLP
Parallelism granularity	Fine-grain: within a single thread	Coarse-grain: multiple threads
Handled by	Mostly hardware	Both hardware & software
Example	Pipelining, superscalar, out-of-order	Multicore CPUs, multithreading
Goal	Speed up single program thread	Speed up multi-threaded workload

Table 35: Comparison ILP vs TLP.

- Which type of technique can be applied in superscalar processors to combine both ILP and TLP?

Answer: In a superscalar processor, the main technique to combine both ILP and TLP is **Simultaneous Multithreading (SMT)**.

- A superscalar processor can issue multiple instructions per cycle, this exploits Instruction-Level Parallelism (ILP).
- However, sometimes a single thread cannot supply enough independent instructions to keep all issue slots busy (due to data dependencies, cache misses, or control hazards).
- To keep the pipelines full, Simultaneous Multithreading allows **instructions from multiple hardware threads to be fetched and issued in the same cycle**.
- In other words, SMT **interleaves instructions from different threads into the same superscalar pipeline**, and this combines ILP (*issuing multiple instructions per cycle*) with TLP (*multiple threads*).

The most common SMT implementation is [Intel's Hyper-Threading technology](#).

In summary, **SMT** leverages the *wide superscalar pipeline* to execute instructions from multiple threads *at the same time*. This effectively *combines ILP and TLP* on a single core.

- Explain what type of **instruction scheduling** is used in superscalar processors to combine both ILP and TLP?

Answer: In superscalar processors, instruction scheduling is about **deciding when and in what order instructions issue to the multiple pipelines**. Scheduling ensures that dependencies are respected and hardware resources are efficiently used.

- **Static scheduling:** The **compiler** reorders instructions **before runtime**. Used in simple in-order superscalar pipelines (like early RISC designs).
- **Dynamic scheduling:** The **processor** **reorders instructions at runtime** to exploit more ILP. This is the standard for modern high-performance superscalar cores.

Almost all modern PC superscalar cores (Intel, AMD, ARM) do dynamic scheduling, with out-of-order execution. The most common algorithm for dynamic scheduling is the Tomasulo algorithm, which uses register renaming to eliminate false dependencies.

When the processor implements **Simultaneous Multithreading** (SMT):

- The **instruction window** holds instructions from *multiple hardware threads* simultaneously.
- The **dynamic scheduling logic** treats all the instructions *from all active threads* together.
- At each cycle, it issues the *ready instructions*, from whichever threads have instructions ready, to the available execution units.

This means the same dynamic out-of-order scheduler works across multiple threads.

- *What are the main hardware modifications required for a generic superscalar processor to support both ILP and TLP?*

Answer: To combine ILP and TLP **within each core**, we typically add Simultaneous Multithreading (SMT). So, what we need is:

- **Multiple architectural register sets.** Each hardware thread needs its own copy of program-visible registers:
 - * Every thread needs its own *program counter (PC)*.
 - * Every thread needs its own *architectural registers* (e.g. x86 has 16 GPRs, ARM has 31 GPRs, plus FP regs).

If we didn't keep these separate, two threads would overwrite each other's registers.

⌚ **How is this solved in hardware?** The physical register file is big enough to hold *many* physical registers. Each thread has a logical map that says: “*My ISA register RAX is currently mapped to physical register #45*”. This is the *register renaming table*. **SMT adds one register map per hardware thread.** This lets the hardware:

- * Track which physical registers belong to which thread's logical registers.
- * Keep architectural state for each thread separate, but reuse the same physical register file.

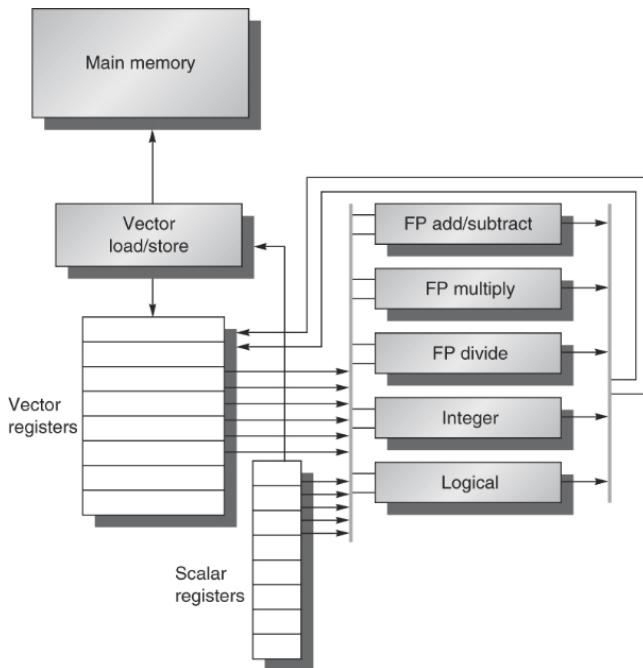
This is needed so the processor can fetch, decode, and track multiple *independent* contexts. For example, Intel's SMT (Hyper-Threading) duplicates only the architectural state per thread, but the physical register file is shared and dynamically renamed.

- **Thread ID tagging in pipeline.** Each instruction must carry a **Thread ID (TID)** through the pipeline. Reservation stations, reorder buffer (ROB), load-store queues, all must track which thread each instruction belongs to. The commit logic uses the TID to retire instructions into the correct thread's architectural state.

- **Thread-aware instruction fetch.** The front-end fetch unit must support fetching from multiple program counters (PCs) in the same cycle or round-robin across active threads. For example, Intel Hyper-Threading allows two threads, front-end selects which one to fetch from each cycle (based on fairness, resource usage, branch miss, stalls).
- **Larger shared resources.** *Instruction window, reservation stations, ROB, rename table,* all must be **larger** to hold more total in-flight instructions. E.g., if a single thread can fill 200 μ ops in the ROB, with SMT we might need to handle 2 \times or 4 \times as many μ ops in-flight to maintain high throughput.
- **Thread scheduling logic.** The issue logic must pick instructions not just by readiness and dependencies, but also across *multiple threads*. This keeps execution units busy when one thread stalls (e.g., on a cache miss).
- **Shared cache + coherent pipeline queues.** L1/L2 caches must be shared among the SMT threads to allow fine-grained switching. Coherence is easier than in multicore because threads share the same core and same cache hierarchy.
- **Interrupt/exception support per thread.** Exceptions and interrupts must be tracked *per thread*. The pipeline must be able to squash/rollback only the instructions belonging to the faulting thread.

Question 2: Vector Processors

Consider a vector processor architecture, as VMIPS shown in the figure.



- Present the main concepts;
- Present the main advantages of vector execution with respect to scalar execution, detailing the features of the vector architecture that provide the advantage.

Answer:

Deepening: What is VMIPS?

Vector MIPS is an educational extension of the classic MIPS RISC architecture, designed to illustrate how a vector processor works. The “V” means it extends the scalar MIPS with vector instructions, vector registers, and vector functional units. The core idea is one vector instruction applies an operation to all elements of a vector register (element-wise) in hardware. For example, let's say we have a vector register V1 holding 64 floating point numbers, a scalar register F2 holding the scalar value 10, then we could write: VADDVS V2, V1, F2; that it corresponds to:

```

1 for i = 0 to V1_length-1:
2     V2[i] = V1[i] + F2

```

The hardware streams the elements of V1 through a pipelined adder. Each cycle, one pair: (V1[i], scalar) is read, added, and the result is written to V2[i]. After an initial pipeline latency, we get one result per cycle, so throughput is very high.

The figure shows a *vector processor datapath*, this type of architecture is typical of Cray-style supercomputers and is the basis for modern SIMD units.

- **Vector Registers:** Large registers that hold whole *vectors*, each vector register stores multiple elements (e.g. 64 or 128 floats). This enables operations on entire arrays in a single instruction.
- **Scalar Registers:** Hold normal single-word scalar values (e.g., loop bounds, offsets).
- **Vector Load/Store Unit:** Moves entire vectors between *main memory* and the vector registers. Memory is accessed in *strided* or *sequential* patterns.
- **Multiple Vector Functional Units:** Specialized vector ALUs: FP add/subtract unit, FP multiply unit, FP divide unit, Integer unit, Logical unit. These execute vector operations in parallel, multiple elements flow through the pipeline concurrently.
- **Pipelined Execution:** Each functional unit is deeply pipelined, one element per cycle throughput after an initial latency.

A single **vector instruction** specifies an operation on an entire vector register, the hardware pipelines iterate over the elements automatically.

Advantages of vector execution vs. scalar execution:

- ✓ **High throughput via data-level parallelism**
 - Scalar: One instruction operates on one element at a time → many instructions → high instruction bandwidth needed.
 - Vector: One vector instruction does N operations → **reduces instruction fetch, decode, and issue overhead**.
- ✓ **Better memory bandwidth efficiency.** Vector loads/stores transfer *entire contiguous blocks* → hardware can stream efficiently. Strided accesses help with matrices. Memory latency is amortized over large transfers.
- ✓ **Deep pipelining with high utilization.** Pipelines stay full because the same operation repeats over many elements. High utilization of functional units → no pipeline bubbles for loop overhead.
- ✓ **Compact code → fewer branch hazards.** Fewer loop control instructions → fewer branches. Branch misprediction cost is lower. Control flow overhead (loop index increments, tests) is eliminated for the vector part.
- ✓ **Easier to overlap computation and memory.** Because operations are predictable and regular, hardware can prefetch next vectors while current computation runs. Software and hardware together can schedule operations to hide memory latency.

A vector processor like VMIPS is specialized for **data-level parallelism**: a single instruction describes operations on many data elements. This reduces instruction overhead, improves pipeline utilization, maximizes memory throughput, and delivers higher performance than scalar execution for the same silicon area when data-parallel problems are available.

Quizzes

1. **Question 1 (format Multiple Choice - Single answer).** Let's consider the following code executed by a Vector Processor with:

- Vector Register File composed of 32 vectors of 8 elements per 64 bits/element;
- Scalar FP Register File composed of 32 registers of 64 bits;
- One Load/Store Vector Unit with operation chaining and memory bandwidth 64 bits;
- One ADD/SUB Vector Unit with operation chaining;
- One MUL/DIV Vector Unit with operation chaining.

```

1 L.V V1, RA      # Load vector from memory address RA into V1
2 L.V V3, RB      # Load vector from memory address RB into V3
3 MULVS.D V1, V1, F0 # FP multiply vector V1 to scalar F0
4 ADDVV.D V2, V1, V1 # FP add vectors V1 and V1
5 MULVS.D V2, V2, F0 # FP multiply vector V2 to scalar F0
6 ADDVV.D V3, V2, V3 # FP add vectors V2 and V3
7 S.V V1, RX      # Store vector V3 into memory address RX
8 S.V V2, RY      # Store vector V3 into memory address RY
9 S.V V3, RZ      # Store vector V3 into memory address RZ

```

How many convoys? How many clock cycles to execute the code? (SINGLE ANSWER)

- ✗ **Answer 1:** 2 convoys; 16 clock cycles
- ✗ **Answer 2:** 3 convoys; 24 clock cycles
- ✗ **Answer 3:** 4 convoys; 32 clock cycles
- ✓ **Answer 4:** 5 convoys; 40 clock cycles
- ✗ **Answer 5:** 6 convoys; 48 clock cycles

A **Convoy** is a group of vector instructions that can be issued together in the same cycle, respecting structural and data dependency constraints, so that the vector processor's functional units are fully utilized without conflicts.

- A convoy is like a bundle of operations that run in lock-step on different functional units.
- The vector processor may have separate pipelines for Load/Store, Vector Add/Subtract, Vector Multiply/Divide, etc.
- If there is no resource conflict (e.g., we don't need two loads at the same time if we only have one load unit) and no data hazard (e.g., we don't need the result of an instruction that hasn't finished yet), we can pack them into the same convoy.
- Convoys are executed sequentially, but the instructions within a convoy execute in parallel, overlapping their pipelines.

Chaining is a hardware feature in vector processors that allows different vector functional units to overlap execution on dependent vector operations, by directly forwarding partial results element by element, without waiting for the entire vector result to be written back to the vector register file first.

Vector pipelines are deep, and each vector operation processes one element per cycle once the pipeline is full. Normally, if Instruction B depends on Instruction A:

- **Without chaining:** B must wait until A computes the entire vector result and writes it back to a vector register.
- **With chaining:** B can start as soon as the first element from A is produced, passing the result directly between functional units for each element.

So we get element-wise overlap → more parallelism, shorter total execution time, fewer convoys.

Motivate your answer by completing the following table:

	Load/Store Vector Unit	Add/Sub Vector Unit	Mul/Div Vector Unit
1^ convoy	L.V V1, RA		
2^ convoy	L.V V3, RB		MULVS.D V1, V1, F0
3^ convoy	S.V V1, RX	ADDVV.D V2, V1, V1	
4^ convoy			MULVS.D V2, V2, F0
5^ convoy	S.V V2, RY	ADDVV.D V3, V2, V3	
6^ convoy	S.V V2, RY		

Without chaining.

	Load/Store Vector Unit	Add/Sub Vector Unit	Mul/Div Vector Unit
1^ convoy	L.V V1, RA	ADDVV.D V2, V1, V1	MULVS.D V1, V1, F0
2^ convoy	S.V V1, RX		
3^ convoy	L.V V3, RB	ADDVV.D V3, V2, V3	MULVS.D V2, V2, F0
4^ convoy	S.V V2, RY		
[.3em] 6^ convoy			
5^ convoy	S.V V3, RZ		

With chaining (solution).

2. **Question 2 (format True/False).** In VLIW architectures, the compiler can detect parallelism only in basic blocks of the code.

Answer: False. A basic block is a straight-line piece of code with no branches in except the entry, and no branches out except the exit. In a basic block, the compiler easily sees all data dependencies. But real programs have branches, loops, and procedure calls, so a compiler that only detects ILP inside each basic block finds very limited parallelism.

VLIW compilers do:

- **Basic block scheduling:** VLIW machines rely on **static compile-time scheduling**. They issue wide instructions with many slots (e.g., 4-8 operations in one VLIW bundle). Most basic blocks in normal code are small, often only 5-15 instructions.
- **Global code scheduling:**
 - **Trace scheduling (Fisher):** Combines multiple basic blocks along the most likely execution paths to form a longer region → reorders instructions globally.
 - **Software pipelining (modulo scheduling):** Especially in loops, overlaps iterations → exposes more ILP than is visible in a single basic block.
 - **Loop unrolling and reordering:** same idea, bigger region, more ILP to fill the VLIW slots.

3. **Question 3 (format True/False).** A 4-issue VLIW processor requires 4 Program Counters to load the necessary instructions in the 4 parallel lanes

Answer: False. Because in VLIW, we don't have multiple independent program counters. We have one program counter that points to a very wide instruction word which contains all the parallel operations to be issued together.

4. **Question 4 (format Multiple Choice - Single answer).** In the *Speculative Tomasulo Architecture*, what type of hardware block is used to undo speculative instructions in case of a mispredicted branch?

- ✓ **Answer 1:** Reorder Buffer;
- ✗ **Answer 2:** Instruction Dispatcher;
- ✗ **Answer 3:** Store Buffer;
- ✗ **Answer 4:** Reservation Stations;
- ✗ **Answer 5:** Load Buffers;

Motivate your answer by explaining what happens in the case of a mispredicted branch.

Answer: The hardware block is the ROB (Reorder Buffer). During a misprediction, the hardware can simply invalidate all ROB entries younger than the branch because all results are only in the ROB. The renaming table shows instructions made thanks to a prediction because they have a speculative flag set to true. Finally, the PC is reset to the correct target.

References

- [1] GeeksforGeeks. Collision resolution techniques [image]. <https://www.geeksforgeeks.org/dsa/collision-resolution-techniques/>, July 23 2025. Image illustrating collision in hashing.
- [2] J.L. Hennessy and D.A. Patterson. *Computer Architecture: A Quantitative Approach*. ISSN. Elsevier Science, 2017.
- [3] Cristina Silvano. Lesson 1, pipelining. Slides from the HPC-E master's degree course on Politecnico di Milano, 2024.
- [4] Cristina Silvano. Advanced computer architecture. Slides from the HPC-E master's degree course on Politecnico di Milano, 2024-2025.

Index

Symbols

1-bit Branch History Table (1-bit BHT)	56
2-D Mesh	423
2-bit Branch History Table (BHT)	59
2-level Predictors	65

A

Address Aliasing	73
Address Space	425
Address Unit	229
Aliasing	72
AMAT (Average Memory Access Time)	257
AMAT in Presence of L2	315
Amdahl's Law	101, 251
Anti-Dependence (Write After Read - WAR)	76
ASAP (As Soon As Possible) Scheduling	279
Average Memory Access Time (AMAT)	302, 303, 321

B

Backward Taken Forward Not Taken (BTFTNT)	44
Bank Conflict	411
Basic Block	273, 276
Bisection Bandwidth	421
Branch Always Not Taken	39
Branch Delay Slot	46
Branch History Register (BHR)	71
Branch History Table (BHT)	56
Branch Outcome	30
Branch Outcome Predictor (BOP)	54
Branch Prediction	37
Branch Prediction Buffer	56
Branch Predictor Aliasing	73
Branch Target Address (BTA)	30
Branch Target Buffer (BTB)	41, 54, 61, 63
BSP Model (Bulk-Synchronous Parallel)	442

C

Cache	301
Cache - Block identification	304
Cache - Block placement	304
Cache - Replacement strategy	304
Cache - Write strategy	304
Cache Associativity	326
Cache Block	300, 301, 304
Cache Coherence problem	444, 447
Cache Hit	301, 319
Cache Line	301, 304
Cache Miss	301, 319

Cache Prefetching	338
Cache Replacement Policy: FIFO (First-In First-Out)	309
Cache Replacement Policy: LRU (Least Recently Used)	309
Cache Replacement Policy: Random Replacement	309
Cache Write Miss Policies: Fetch on Write	312
Cache Write Miss Policies: No Write Allocate	312
Cache Write Miss Policies: Write Allocate	312
Cache Write Miss Policies: Write Around	312
Cache Write Policies: Write-Back	311
Cache Write Policies: Write-Through	311
Cache Write-Back	563
Cache Write-Invalidate	562
Cache Write-Through	563
Cache-Coherent NUMA (CC-NUMA)	434
Capacity Misses	319
Centralized Shared-Memory Multiprocessors	427
Chaining	400, 403, 581
Chime	403
Cluster	438
Coarse-Grain Thread	374
Coarse-Grained Multithreading (CGMT)	380
Code Efficiency	288
Code Scheduling	191, 275
Coherence Misses	320
Coherent	447
Cold Start Misses	319
Collision	319
Column-Associativity	333
Common Data Bus (CDB)	145, 150
Communication/Programming Model	440
Compensation Code	293
Compiler Scheduling	25
Compulsory Misses	319
Conflict Misses	319
Contention	432
Control Dependencies	79
Control Hazards	20, 21, 32
Convoy	403, 580
Correlating Predictor	499
Correlating Predictors	65
Critical Path	276, 278
Critical Word First	346
Crossbar (Fully Connected Switch Network)	423
D	
D-Stream Buffer	336
Data	304
Data Cache Miss	272
Data Dependence Graph	277
Data Dependencies	75

Data Hazards	20, 21
Data Prefetching	336
Data-Level Parallelism (DLP)	387
Data-Parallel Model	442
DAXPY (Double-precision A·X Plus Y)	402
Delayed Branch	46
Delayed Branch Scheduling: From After	53
Delayed Branch Scheduling: From Before	48
Delayed Branch Scheduling: From Fall-Through	51
Delayed Branch Scheduling: From Target	49
Dependence Graph	277
Direct-Mapped	563
Direct-Mapped Cache	305
Directory-Based Protocol	462
Dirty Bit	311
Distributed Shared Memory (DSM)	434
Dual-Issue Pipeline	90
Dynamic Branch Prediction Techniques	37
Dynamic Loop Unrolling	193
Dynamic Out-of-Order Parallelism	193
Dynamic Renaming - Hardware-side	78
Dynamic Scheduling	87
Dynamic Speculative Execution	193
E	
Early Branch Evaluation	35
Early Restart	346
EX (Execution)	9
EX/EX Path	23
Execution Stage	110
Execution Time	287
Execution Time (Latency, Response Time)	246
F	
False Dependencies	76, 142
Fan-out	449
Fine-Grain Thread	373
Fine-Grained Multithreading (FGMT)	380
Finite-State Machine (FSM)	453
Floating Point (FP)	150
Floating-Point Unit (FPU)	149
Formula Memory Hierarchy: AMAT	257
Formula Memory Hierarchy: AMAT (L1, L2)	260
Formula Memory Hierarchy: AMAT Harvard	258
Formula Memory Hierarchy: Final Full CPU Time Formula	265
Formula Memory Hierarchy: Full CPU Time	262
Formula Memory Hierarchy: Hit Rate	257
Formula Memory Hierarchy: Memory stall cycles per instruction	264
Formula Memory Hierarchy: Miss Rate	257
Formula Memory Hierarchy: Miss Time	257

Formula Pipeline Loops (AS): IC _{AS} per iter	255
Formula Pipeline Loops (AS): MIPS per iteration	255
Formula Pipeline Loops (AS): Total Clock Cycles per iter	255
Formula Pipeline Loops AS: Cycles Per Instruction per iter	255
Formula Pipeline Loops: Clock Cycles per iter	254
Formula Pipeline Loops: Cycles Per Instruction per iter	254
Formula Pipeline Loops: IC per iter	254
Formula Pipeline Loops: MIPS per iter	254
Formula Pipeline: MIPS	253
Formula Pipeline: Number of clock cycles	253
Formula Speedup: Pipeline Speedup	256
Formula: Clock Cycle Time f_{CLK}	247
Formula: Clock Cycle Time T_{CLK}	247
Formula: Clocks Per Instruction (CPI)	27, 248
Formula: CPU Time	248
Formula: Execution Time	248
Formula: Instruction Count (IC)	27
Formula: Instruction Frequency	249
Formula: Instructions Per Cycle (IPC)	248
Formula: MIPS	250
Formula: Performance	247
Formula: Relative Performance	247
Formula: Total Clock Cycles	27
Formula: Weighted Average CPI	249
Formula: Weighted CPU Time	248
Forwarding	22
Frame	357
Fully Associative Cache	306
Function Units (FUs)	150
Functional Unit Status Table	113

G

Gather	413
Global Adaptive (GA)	72
Global History Register (GHR)	65, 499
Global Miss Rate	261, 316
Graduation	241
GShare	72

H

Hardware Multithreading	380
Hazard	20
Hit	257
Hit Time	257, 302, 321
Hit Under Miss	348
Hypercube (N-cube)	424

I

ID (Instruction Decode)	9
IF (Instruction Fetch)	9

Imprecise Exceptions	85
In-Order Issue	80
Instruction Dependence Graph	277
Instruction Per Clock (IPC)	27
Instruction Prefetching	335
Instruction Queue	149
Instruction Status Table	113
instruction window	101
Instruction-Level Parallelism (ILP)	37, 74, 574
Interference Misses	319
Issue Stage	105
Issue Stage (In-Order)	109
L	
L1 Cache	258
L2 cache	260
L2 Local Miss Rate	316
Large-Scale Multiprocessors (Distributed Shared Memory)	417
Last Level Cache (LLC)	503
List-based scheduling	510
List-Based Scheduling Algorithm	281
Load Buffers	147, 229
Load/Store Queue	150
Local Miss Rate	261
Locality	298
Loop Blocking (Tiling)	341
Loop Fusion	341
Loop Interchange	340
Loop Overhead	287
Loop Unrolling	189, 285
M	
Maximum Vector Length (MVL)	407
ME (Memory Access)	9
MEM/EX Path	24
MEM/MEM Path	24
Memory Accesses per Instruction (MAPI)	262
Memory Aliasing	73
Memory Consistency problem	444
Memory Disambiguation	202
Merging Arrays	340
MESI	562
MESI Protocol	458
MESI Protocol: Finite-State Machine (FSM)	460
Message Passing Model	425
Message-Passing Model	441
Millions of Instructions Per Second (MIPS)	28
MIMD (Multiple Instruction, Multiple Data)	388
MISD (Multiple Instruction, Single Data)	388
Misprediction-Hit	331

Misprediction	37
Miss	257
Miss Penalty	257, 302, 315, 321
Miss Rate	321
Miss Status Holding Registers (MSHRs)	348
Miss Time	303
Miss Under Miss	348
MPI (Message Passing Interface)	438
MSI Protocol	453
MSI Protocol: Finite-State Machine (FSM)	456
Multi-Cycle In-Order Pipeline	81
Multi-Cycle Out-of-Order (OoO) Pipeline	84
Multicores (on-chip multiprocessors)	417
Multiple Outstanding Misses	348
Multiple-Issue Processors	89
Multiple-Issue Width	95
Multiprocessor	417
N	
N-Way Set-Associative Cache	306
Name Dependencies	76
Network-Connected Multiprocessors	420
Non-Blocking Cache	348
NUMA (Non-Uniform Memory Access)	429
O	
Operating System Context Switch	375
Output Dependence (Write After Write - WAW)	76
P	
Pattern History Table (PHT)	71
Physical Page	357
Pipeline Registers	19
Pipelining	8
Predicted Target Address (PTA)	54
Principle of Locality	300
Private Address Spaces	425
Processing Element (PE)	390
Profile-Driven Prediction	45
Pseudo-Associativity	333
Pseudo-Dependencies	76
R	
RAW (Read After Write)	21
Read After Write (RAW)	75
Read Operands Stage (Out-of-Order)	110
Read Operands Stage (RR)	105
Ready Bit	230
Ready List	279
Ready Set	281

Register Aliasing	73
Register File (RF)	150
Register Prefetching	338
Register Renaming	77, 186
Register Result Status Table	113
Rename Table	232
Reorder Buffer (ROB)	222, 223
Reservation Station (RS)	146
Reservation Stations	144
Reservation Stations (RS)	143
Reservation Stations (RSS)	149
Ring	423
RISC-V Data Path	10
ROB Head pointer	231
ROB number	226
ROB Tail pointer	231
S	
Scatter	413
Scoreboard	103
Second-Level Cache (L2)	315
Shared Address Space	425
Shared Memory Model	425
Shared-Memory Model	440
SIMD (Single Instruction, Multiple Data)	388
Simultaneous Multithreading (SMT)	381, 575
Single Bus	422
Single Bus Multiprocessors	420
SISD (Single Instruction, Single Data)	388
Small-scale multicore	432
SMP - Symmetric Multiprocessors (centralized shared memory)	417
Snoop Tag Store	450
Snooping Protocol	449
Software Pipelining	289, 506
Spatial Locality	298, 300
Speculative Execution	292
Speculative Flag	230
SPMD (Single Program, Multiple Data)	442
Static Branch Prediction Techniques	37
Static Renaming	142
Static Renaming - Software-side	78
Store Buffers	147
Stride	411
Strip Mining	408
Strongly Not Taken (SNT)	59
Strongly Taken (ST)	59
Structural Hazards	20, 21
Superblock Scheduling	294
Superscalar Processor	95
Symmetric Multiprocessor (SMP)	432

Symmetric Multiprocessors (SMPs)	427
T	
Tag	304
Tail Duplication	295
Temporal Locality	298, 300
Temporary Names (tags)	143
Thread	373
Thread Switching	379
Thread-Level Parallelism (TLP)	574
Throughput (Bandwidth)	246
Tomasulo's Algorithm	140
Total Network Bandwidth	421
Trace Compaction (Scheduling)	292
Trace Scheduling	292
Trace Selection	292
Translation Lookaside Buffer (TLB)	359
True Data Dependencies	75
Turing complete	418
Turing Completeness	418
Turing Machine	418
Two-Level Adaptive Branch Predictors	71
U	
UMA (Uniform Memory Access)	427
V	
Valid Bit	304
Vector Functional Unit (VFU)	400
Vector Functional Units (VFUs)	392
Vector Length Register (VLR)	407
Vector Load (LV)	400
Vector Load/Store Unit (LSU)	400
Vector Mask Register (VMR)	409
Vector MIPS (VMIPS)	399
Vector Register File (VRF)	392, 399
Vector Store (SV)	400
Victim Block Swap	328
Victim Cache	328
Virtual Page (VP)	357
Virtually Indexed, Physically Tagged (VIPT) Cache	360
VLIW (Very Long Instruction Word)	98, 266
VLIW (Very Long Instruction Word) Processor	95
W	
WAR (Write After Read)	21
WAW (Write After Write)	21
Way Prediction	330
Way Predictor Table (WPT)	330
WB (Write Back)	9

Weakly Not Taken (WNT)	59
Weakly Taken (WT)	59
Wrapped Fetch	346
Write Result Stage (Out-of-Order)	110
Write-Allocate	563
Write-Broadcast protocol	452
Write-Invalidate protocol	447, 452
Write-Update protocol	448, 452