

# Contents

<b>1</b>	<b>PRAM</b>	<b>6</b>
1.1	Prerequisites . . . . .	6
1.2	Definition . . . . .	6
1.3	How it works . . . . .	7
1.3.1	Computation . . . . .	7
1.3.2	PRAM Classification . . . . .	7
1.3.3	Strengths of PRAM . . . . .	8
1.3.4	How to compare PRAM models . . . . .	8
1.4	MVM algorithm . . . . .	10
1.5	SPMD sum . . . . .	12
1.6	MM algorithm . . . . .	16
1.7	PRAM variants and Lemmas . . . . .	17
1.8	PRAM implementation . . . . .	18
1.9	Amdahl's and Gustafson's Laws . . . . .	20
<b>2</b>	<b>Fundamentals of architecture</b>	<b>23</b>
2.1	Introduction . . . . .	23
2.1.1	Simplest processor . . . . .	23
2.1.2	Superscalar processor . . . . .	24
2.1.3	Single Instruction, Multiple Data (SIMD) processor . . . . .	25
2.1.4	Multi-Core Processor . . . . .	25
2.2	Accessing Memory . . . . .	26
2.2.1	What is a memory? . . . . .	26
2.2.2	How to reduce processor stalls . . . . .	28
2.2.2.1	Cache . . . . .	28
2.2.2.2	Multi-threading . . . . .	28
<b>3</b>	<b>Programming models</b>	<b>31</b>
3.1	Implicit SPMD Program Compiler (ISPC) . . . . .	31
3.2	Shared Address Space Model . . . . .	35
3.3	Message Passing model of communication . . . . .	36
3.4	Data-Parallel model . . . . .	37
<b>4</b>	<b>Parallel Programming Models and pthreads</b>	<b>39</b>
4.1	How to create parallel algorithms and programs . . . . .	39
4.2	Analyze parallel algorithms . . . . .	41
4.3	Technologies . . . . .	44
4.4	Threads . . . . .	47
4.4.1	Flynn's taxonomy . . . . .	47
4.4.2	Definition . . . . .	47
4.4.3	pthreads API . . . . .	49
4.4.3.1	Creation . . . . .	49
4.4.3.2	Termination . . . . .	50
4.4.3.3	Joining . . . . .	51
4.4.3.4	Detaching . . . . .	52
4.4.3.5	Joining through Barriers . . . . .	53
4.4.3.6	Mutexes . . . . .	54
4.4.3.7	Condition variables . . . . .	54

<b>5</b>	<b>OpenMP v5.2</b>	<b>55</b>
5.1	Introduction . . . . .	55
5.2	Basic syntax . . . . .	57
5.3	Work sharing . . . . .	60
5.3.1	For . . . . .	60
5.3.1.1	Reduction . . . . .	65
5.3.2	Sections . . . . .	67
5.3.3	Single/Master . . . . .	68
5.3.4	Tasks . . . . .	69
5.3.4.1	Task dependences . . . . .	72
5.4	Synchronization . . . . .	76
5.5	Data environment . . . . .	79
5.6	Memory model . . . . .	87
5.7	Nested Parallelism . . . . .	90
5.8	Cancellation . . . . .	94
5.9	SIMD Vectorization . . . . .	97
<b>6</b>	<b>GPU Architecture</b>	<b>100</b>
6.1	Introduction . . . . .	100
6.2	GPU compute mode . . . . .	101
6.3	CUDA . . . . .	103
6.3.1	Basics of CUDA . . . . .	103
6.3.2	Memory model . . . . .	107
6.3.3	NVIDIA V100 Streaming Multiprocessor (SM) . . . . .	109
6.3.4	Running a CUDA program on a GPU . . . . .	112
6.3.5	Implementation of CUDA abstractions . . . . .	118
6.3.6	Advanced thread scheduling . . . . .	121
6.3.7	Memory and Data Locality in Depth . . . . .	126
6.3.8	Tiling Technique . . . . .	135
6.3.8.1	Tiled Matrix Multiplication . . . . .	138
6.3.8.2	Implementation Tiled Matrix Multiplication . . . . .	143
6.3.8.3	Any size matrix handling . . . . .	148
6.3.9	Optimizing Memory Coalescing . . . . .	153
<b>7</b>	<b>CUDA</b>	<b>162</b>
7.1	Introduction . . . . .	162
7.2	CUDA Basics . . . . .	166
7.2.1	GPGPU Best Practices . . . . .	168
7.2.2	Compilation . . . . .	170
7.2.3	Debugging . . . . .	172
7.2.4	CUDA Kernel . . . . .	175
7.3	Execution Model . . . . .	178
7.4	Querying Device Properties . . . . .	180
7.5	Thread hierarchy . . . . .	182
7.6	Memory hierarchy . . . . .	185
7.7	Streams . . . . .	193
7.8	CUDA and OpenMP or MPI . . . . .	197
7.8.1	Motivations . . . . .	197
7.8.2	CUDA API for Multi-GPUs . . . . .	202
7.8.3	Memory Management with Multiple GPUs . . . . .	205

7.8.4	Batch Processing and Cooperative Patterns with OpenMP	211
7.8.5	OpenMP for heterogeneous architectures . . . . .	213
7.8.6	MPI-CUDA applications . . . . .	216
<b>8</b>	<b>Memory Consistency</b>	<b>220</b>
8.1	Coherence vs Consistency . . . . .	220
8.2	Definition . . . . .	223
8.3	Sequential Consistency Model . . . . .	225
8.4	Memory Models with Relaxed Ordering . . . . .	229
8.4.1	Allowing Reads to Move Ahead of Writes . . . . .	230
8.4.2	Allowing writes to be reordered . . . . .	232
8.4.3	Allowing all reorderings . . . . .	234
8.5	Languages Need Memory Models Too . . . . .	236
8.6	Implementing Locks . . . . .	238
8.6.1	Introduction . . . . .	238
8.6.2	Test-and-Set based lock . . . . .	240
8.6.3	Test-and-Test-and-Set lock . . . . .	244
<b>9</b>	<b>Heterogeneous Processing</b>	<b>248</b>
9.1	Energy Constrained Computing . . . . .	250
9.2	Compute Specialization . . . . .	251
9.3	Challenges of heterogeneous designs . . . . .	265
9.4	Reducing energy consumption . . . . .	268
<b>10</b>	<b>Patterns</b>	<b>271</b>
10.1	Dependencies . . . . .	271
10.2	Parallel Patterns . . . . .	280
10.2.1	Nesting Pattern . . . . .	281
10.2.2	Serial Control Patterns . . . . .	282
10.2.3	Parallel Control Patterns . . . . .	284
10.2.4	Serial Data Management Patterns . . . . .	289
10.2.5	Parallel Data Management Patterns . . . . .	292
10.2.6	Other Parallel Patterns . . . . .	295
10.3	Map Pattern . . . . .	297
10.3.1	What is a Map? . . . . .	297
10.3.2	Optimizations . . . . .	299
10.3.2.1	Sequences of Maps . . . . .	299
10.3.2.2	Code Fusion . . . . .	300
10.3.2.3	Cache Fusion . . . . .	301
10.3.3	Related Patterns . . . . .	302
10.3.4	Scaled Vector Addition (SAXPY) . . . . .	304
10.4	Collectives operations . . . . .	307
10.4.1	Reduce (or Reduction) Pattern . . . . .	308
10.4.2	Scan Pattern . . . . .	314
10.5	Gather Pattern . . . . .	325
10.5.1	What is a Gather? . . . . .	325
10.5.2	Shift . . . . .	329
10.5.3	Zip . . . . .	331
	<b>Index</b>	<b>330</b>

---

### 10.5.3 Zip

The **Zip operation** is a special case of the gather pattern where **two (or more) arrays are combined by interleaving their elements**. It functions like a zipper, **taking one element from each array in sequence to form a new combined array**. It is important to note that it **works with different types**, so we can zip elements of different types, like integers and floats, or even complex objects.

#### ✂ How does it work?

The operation **takes an element from the first array**, then one from the **second array**, another from the third, and so on, and **repeats the process**. The **output is the combined sequence**.

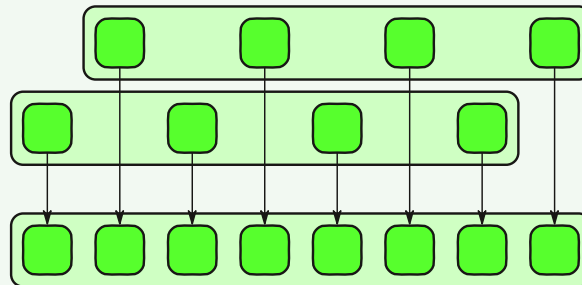
#### Example 46: zip operation

Consider two arrays:

1. Real Parts: contains real numbers.
2. Imaginary Parts: contains imaginary numbers.

The output is a combined sequence like:

[Real0, Imag0, Real1, Imag1, Real2, Imag2, ...]



#### 🧑‍🤖 Parallelism

Each pair of elements (one from each array) can be **combined independently**. This independence allows **parallel execution** since **there's no dependency between the operations for different pairs**.