

# Contents

|          |  |           |
|----------|--|-----------|
| <b>1</b> | <b>PRAM</b>  | <b>7</b>  |
| 1.1      | Prerequisites . . . . .                                      | 7         |
| 1.2      | Definition . . . . .   | 7         |
| 1.3      | How it works . . . . .                                       | 8         |
| 1.3.1    | Computation . . . . .  | 8         |
| 1.3.2    | PRAM Classification . . . . .                                | 8         |
| 1.3.3    | Strengths of PRAM . . . . .                                  | 9         |
| 1.3.4    | How to compare PRAM models . . . . .                         | 9         |
| 1.4      | MVM algorithm . . . . .                                      | 11        |
| 1.5      | SPMD sum . . . . .   | 13        |
| 1.6      | MM algorithm . . . . .                                       | 17        |
| 1.7      | PRAM variants and Lemmas . . . . .                           | 18        |
| 1.8      | PRAM implementation . . . . .                                | 19        |
| 1.9      | Amdahl's and Gustafson's Laws . . . . .                      | 21        |
| <b>2</b> | <b>Fundamentals of architecture</b>                          | <b>24</b> |
| 2.1      | Introduction . . . . .                                       | 24        |
| 2.1.1    | Simplest processor . . . . .                                 | 24        |
| 2.1.2    | Superscalar processor . . . . .                              | 25        |
| 2.1.3    | Single Instruction, Multiple Data (SIMD) processor . . . . . | 26        |
| 2.1.4    | Multi-Core Processor . . . . .                               | 26        |
| 2.2      | Accessing Memory . . . . .                                   | 27        |
| 2.2.1    | What is a memory? . . . . .                                  | 27        |
| 2.2.2    | How to reduce processor stalls . . . . .                     | 29        |
| 2.2.2.1  | Cache . . . . .  | 29        |
| 2.2.2.2  | Multi-threading . . . . .                                    | 29        |
| <b>3</b> | <b>Programming models</b>                                    | <b>32</b> |
| 3.1      | Implicit SPMD Program Compiler (ISPC) . . . . .              | 32        |
| 3.2      | Shared Address Space Model . . . . .                         | 36        |
| 3.3      | Message Passing model of communication . . . . .             | 37        |
| 3.4      | Data-Parallel model . . . . .                                | 38        |
| <b>4</b> | <b>Parallel Programming Models and pthreads</b>              | <b>40</b> |
| 4.1      | How to create parallel algorithms and programs . . . . .     | 40        |
| 4.2      | Analyze parallel algorithms . . . . .                        | 42        |
| 4.3      | Technologies . . . . .                                       | 45        |
| 4.4      | Threads . . . . .  | 48        |
| 4.4.1    | Flynn's taxonomy . . . . .                                   | 48        |
| 4.4.2    | Definition . . . . .   | 48        |
| 4.4.3    | pthreads API . . . . .                                       | 50        |
| 4.4.3.1  | Creation . . . . .   | 50        |
| 4.4.3.2  | Termination . . . . .  | 51        |
| 4.4.3.3  | Joining . . . . .  | 52        |
| 4.4.3.4  | Detaching . . . . .  | 53        |
| 4.4.3.5  | Joining through Barriers . . . . .                           | 54        |
| 4.4.3.6  | Mutexes . . . . .  | 55        |
| 4.4.3.7  | Condition variables . . . . .                                | 55        |

|          |  |            |
|----------|--|------------|
| <b>5</b> | <b>OpenMP v5.2</b>                                   | <b>56</b>  |
| 5.1      | Introduction . . . . .                               | 56         |
| 5.2      | Basic syntax . . . . .                               | 58         |
| 5.3      | Work sharing . . . . .                               | 61         |
| 5.3.1    | For . . . . .  | 61         |
| 5.3.1.1  | Reduction . . . . .                                  | 66         |
| 5.3.2    | Sections . . . . .                                   | 68         |
| 5.3.3    | Single/Master . . . . .                              | 69         |
| 5.3.4    | Tasks . . . . .                                      | 70         |
| 5.3.4.1  | Task dependences . . . . .                           | 73         |
| 5.4      | Synchronization . . . . .                            | 77         |
| 5.5      | Data environment . . . . .                           | 80         |
| 5.6      | Memory model . . . . .                               | 88         |
| 5.7      | Nested Parallelism . . . . .                         | 91         |
| 5.8      | Cancellation . . . . .                               | 95         |
| 5.9      | SIMD Vectorization . . . . .                         | 98         |
| <b>6</b> | <b>GPU Architecture</b>                              | <b>101</b> |
| 6.1      | Introduction . . . . .                               | 101        |
| 6.2      | GPU compute mode . . . . .                           | 102        |
| 6.3      | CUDA . . . . .                                       | 104        |
| 6.3.1    | Basics of CUDA . . . . .                             | 104        |
| 6.3.2    | Memory model . . . . .                               | 108        |
| 6.3.3    | NVIDIA V100 Streaming Multiprocessor (SM) . . . . .  | 110        |
| 6.3.4    | Running a CUDA program on a GPU . . . . .            | 113        |
| 6.3.5    | Implementation of CUDA abstractions . . . . .        | 119        |
| 6.3.6    | Advanced thread scheduling . . . . .                 | 122        |
| 6.3.7    | Memory and Data Locality in Depth . . . . .          | 127        |
| 6.3.8    | Tiling Technique . . . . .                           | 136        |
| 6.3.8.1  | Tiled Matrix Multiplication . . . . .                | 139        |
| 6.3.8.2  | Implementation Tiled Matrix Multiplication . . . . . | 144        |
| 6.3.8.3  | Any size matrix handling . . . . .                   | 149        |
| 6.3.9    | Optimizing Memory Coalescing . . . . .               | 154        |
| <b>7</b> | <b>CUDA</b>  | <b>163</b> |
| 7.1      | Introduction . . . . .                               | 163        |
| 7.2      | CUDA Basics . . . . .                                | 167        |
| 7.2.1    | GPGPU Best Practices . . . . .                       | 169        |
| 7.2.2    | Compilation . . . . .                                | 171        |
| 7.2.3    | Debugging . . . . .                                  | 173        |
| 7.2.4    | CUDA Kernel . . . . .                                | 176        |
| 7.3      | Execution Model . . . . .                            | 179        |
| 7.4      | Querying Device Properties . . . . .                 | 181        |
| 7.5      | Thread hierarchy . . . . .                           | 183        |
| 7.6      | Memory hierarchy . . . . .                           | 186        |
| 7.7      | Streams . . . . .                                    | 194        |
| 7.8      | CUDA and OpenMP or MPI . . . . .                     | 198        |
| 7.8.1    | Motivations . . . . .                                | 198        |
| 7.8.2    | CUDA API for Multi-GPUs . . . . .                    | 203        |
| 7.8.3    | Memory Management with Multiple GPUs . . . . .       | 206        |

|           |   |            |
|-----------|---|------------|
| 7.8.4     | Batch Processing and Cooperative Patterns with OpenMP | 212        |
| 7.8.5     | OpenMP for heterogeneous architectures . . . . .      | 214        |
| 7.8.6     | MPI-CUDA applications . . . . .                       | 217        |
| <b>8</b>  | <b>Memory Consistency</b>                             | <b>221</b> |
| 8.1       | Coherence vs Consistency . . . . .                    | 221        |
| 8.2       | Definition . . . . .                                  | 224        |
| 8.3       | Sequential Consistency Model . . . . .                | 226        |
| 8.4       | Memory Models with Relaxed Ordering . . . . .         | 230        |
| 8.4.1     | Allowing Reads to Move Ahead of Writes . . . . .      | 231        |
| 8.4.2     | Allowing writes to be reordered . . . . .             | 233        |
| 8.4.3     | Allowing all reorderings . . . . .                    | 235        |
| 8.5       | Languages Need Memory Models Too . . . . .            | 237        |
| 8.6       | Implementing Locks . . . . .                          | 239        |
| 8.6.1     | Introduction . . . . .                                | 239        |
| 8.6.2     | Test-and-Set based lock . . . . .                     | 241        |
| 8.6.3     | Test-and-Test-and-Set lock . . . . .                  | 245        |
| <b>9</b>  | <b>Heterogeneous Processing</b>                       | <b>249</b> |
| 9.1       | Energy Constrained Computing . . . . .                | 251        |
| 9.2       | Compute Specialization . . . . .                      | 252        |
| 9.3       | Challenges of heterogeneous designs . . . . .         | 266        |
| 9.4       | Reducing energy consumption . . . . .                 | 269        |
| <b>10</b> | <b>Patterns</b>                                       | <b>272</b> |
| 10.1      | Dependencies . . . . .                                | 272        |
| 10.2      | Parallel Patterns . . . . .                           | 281        |
| 10.2.1    | Nesting Pattern . . . . .                             | 282        |
| 10.2.2    | Serial Control Patterns . . . . .                     | 283        |
| 10.2.3    | Parallel Control Patterns . . . . .                   | 285        |
| 10.2.4    | Serial Data Management Patterns . . . . .             | 290        |
| 10.2.5    | Parallel Data Management Patterns . . . . .           | 293        |
| 10.2.6    | Other Parallel Patterns . . . . .                     | 296        |
| 10.3      | Map Pattern . . . . .                                 | 298        |
| 10.3.1    | What is a Map? . . . . .                              | 298        |
| 10.3.2    | Optimizations . . . . .                               | 300        |
| 10.3.2.1  | Sequences of Maps . . . . .                           | 300        |
| 10.3.2.2  | Code Fusion . . . . .                                 | 301        |
| 10.3.2.3  | Cache Fusion . . . . .                                | 302        |
| 10.3.3    | Related Patterns . . . . .                            | 303        |
| 10.3.4    | Scaled Vector Addition (SAXPY) . . . . .              | 305        |
| 10.4      | Collectives operations . . . . .                      | 308        |
| 10.4.1    | Reduce (or Reduction) Pattern . . . . .               | 309        |
| 10.4.2    | Scan Pattern . . . . .                                | 315        |
| 10.5      | Gather Pattern . . . . .                              | 326        |
| 10.5.1    | What is a Gather? . . . . .                           | 326        |
| 10.5.2    | Shift . . . . .                                       | 330        |
| 10.5.3    | Zip . . . . .   | 332        |
| 10.5.4    | Unzip . . . . .                                       | 333        |
| 10.6      | Scatter Pattern . . . . .                             | 334        |

|           |   |            |
|-----------|---|------------|
| 10.6.1    | What is a Scatter? . . . . .                              | 334        |
| 10.6.2    | Avoid race conditions . . . . .                           | 337        |
| 10.6.2.1  | Atomic Scatter . . . . .                                  | 337        |
| 10.6.2.2  | Permutation Scatter . . . . .                             | 339        |
| 10.6.2.3  | Merge Scatter . . . . .                                   | 341        |
| 10.6.2.4  | Priority Scatter . . . . .                                | 343        |
| 10.7      | Pack Pattern . . . . .                                    | 344        |
| 10.7.1    | What is a Pack? . . . . .                                 | 344        |
| 10.7.2    | Split . . . . .   | 347        |
| 10.7.3    | Unsplit . . . . .   | 348        |
| 10.7.4    | Bin . . . . .   | 349        |
| 10.7.5    | Expand . . . . .  | 350        |
| 10.8      | Partitioning Data . . . . .                               | 351        |
| 10.9      | AoS vs. SoA . . . . .                                     | 352        |
| 10.10     | Stencil Pattern . . . . .                                 | 357        |
| 10.10.1   | What is a Stencil? . . . . .                              | 357        |
| 10.10.2   | Implementing stencil with shift . . . . .                 | 358        |
| 10.10.3   | Cache optimizations . . . . .                             | 360        |
| 10.10.4   | Communication optimizations . . . . .                     | 362        |
| <b>11</b> | <b>Parallel Patterns in OpenMP and CUDA</b>               | <b>364</b> |
| 11.1      | OpenMP . . . . .  | 364        |
| 11.2      | Histogram Pattern . . . . .                               | 366        |
| 11.3      | Reduction Pattern . . . . .                               | 375        |
| 11.4      | Scan Pattern . . . . .                                    | 379        |
| <b>12</b> | <b>Heterogeneous Computing - DSLs and HLS</b>             | <b>383</b> |
| 12.1      | Introduction to Heterogeneous Computing . . . . .         | 383        |
| 12.2      | Heterogeneous parallel programming . . . . .              | 384        |
| 12.3      | DSLs and Halide . . . . .                                 | 385        |
| 12.4      | Scheduling & Performance Optimization in Halide . . . . . | 389        |
|           | <b>Index</b>  | <b>384</b> |

## 12 Heterogeneous Computing - DSLs and HLS

### 12.1 Introduction to Heterogeneous Computing

**Heterogeneous Computing** (or **Heterogeneous Processing**) refers to systems that use multiple types of processors or accelerators to handle different workloads more efficiently.

- In contrast to traditional homogeneous systems (which only use CPUs), heterogeneous systems combine different processing units such as CPUs, GPUs, DSPs, and FPGAs.
- The goal is to match the right processor to the right task, achieving higher performance and energy efficiency.

#### Example 1: Heterogeneous Processing

A self-driving car requires CPUs for decision-making, GPUs for image recognition, and FPGAs for real-time sensor fusion.

### 🔌 Energy-Efficient Computing Strategies

When designing a heterogeneous system, performance isn't the only goal; **energy efficiency is just as critical**. Given a fixed power budget, simply **increasing performance without considering power constraints is inefficient**. Specialized hardware (e.g., FPGAs, ASICs) achieves better performance per watt than general-purpose processors.

There are two main strategies for improving energy efficiency:

1. **Use Specialized Processors.** CPUs are not energy-efficient due to instruction decoding, branch handling, and pipeline management overhead. Specialized hardware (FPGAs, ASICs) reduces overhead, leading to more computations per joule.

$$\text{Power} = \frac{\text{Op}}{\text{second}} \times \frac{\text{Joules}}{\text{Op}}$$

2. **Minimize Data Movement.** Memory access consumes more energy than computation! Optimizing data locality reduces power consumption. For example, moving computation closer to memory (e.g., using tensor core inside GPUs) significantly reduces energy cost.

## 12.2 Heterogeneous parallel programming

### ⚠ Challenges of Writing Portable and Efficient Parallel Code

Writing parallel programs for heterogeneous systems is difficult due to the following reasons:

1. **Diverse Hardware Architectures.** A CPU, GPU, and FPGA all have different programming models. **Code written for one hardware type may not perform well on another.**
2. **Performance vs. Productivity Trade-offs.**
  - **Performance:** Low-level programming (e.g., **CUDA, OpenCL, Verilog**) allows fine-tuned optimizations but **is hard to program.**
  - **Productivity:** High-level abstractions (e.g., **OpenMP, DSLs**) improve productivity but **may introduce performance overhead.**
3. **Memory Management.** Different memory models (shared vs. distributed) require different optimizations. Data movement between CPU and GPU memory can be costly if not handled efficiently.
4. **Scalability Issues.** Some **programs scale well on GPUs but poorly on CPUs** due to synchronization and memory bandwidth limitations.

### ✔ The Ideal Parallel Programming Language

An ideal parallel programming model should provide a balance of:

- ✔ **Performance.** Optimized execution across different hardware.
- ✔ **Productivity.** Easy to use and develop.
- ✔ **Generality.** Works across different architectures.

However, **most existing languages optimize only one or two** of these factors, leading to trade-offs.

| Approach             | Performance | Productivity | Generality |
|----------------------|-------------|--------------|------------|
| CUDA/OpenCL          | ✔ High      | ✗ Low        | ✗ Low      |
| OpenMP (CPU)         | ✔ High      | ✔ Medium     | ✗ Low      |
| MPI (Distributed)    | ✔ High      | ✗ Low        | ✔ High     |
| FPGA/Verilog/VHDL    | ✔ Very High | ✗ Very Low   | ✗ Low      |
| High-Level Synthesis | ✔ High      | ✔ Medium     | ✗ Low      |

### ❓ Why is this important?

If we want **portable parallel programs**, we need **new high-level abstractions** like Domain-Specific Languages (DSLs), which will be covered in the next section.

## 12.3 DSLs and Halide

### 🔗 What are Domain-Specific Languages (DSLs)?

A **Domain-Specific Language (DSL)** is a **specialized programming language** designed for a **specific application domain**. The main characteristics of DSLs are:

- **Restricted expressiveness** (focused on a single domain)
- **High-level, declarative syntax** (easier than general purpose languages)
- **Optimized performance** for the target domain
- **May be standalone or embedded** in another language

| DSL Name     | Target Domain    | Key Benefits   |
|--------------|------------------|--|
| Halide       | Image Processing | Separates algorithm from scheduling for optimized execution. |
| TensorFlow   | Machine Learning | Optimized computation graphs for AI workloads.               |
| SQL          | Databases        | Declarative queries for efficient data retrieval.            |
| Verilog/VHDL | Hardware Design  | Describes digital circuits for synthesis.                    |

Table 11: Examples of DSLs.

### ⚖️ Embedded vs. External DSLs

DSLs can be classified as:

- **External** DSLs:
  - 📖 Have **their own syntax** and *compiler/interpreter*.
  - 🔗 **Example:** SQL, Halide, Verilog.
  - ✅ **Advantages:** can be **more optimized** but require custom compilers.
- **Embedded** DSLs:
  - 📖 **Built inside another general-purpose language.**
  - 🔗 **Example:** TensorFlow (embedded in Python).
  - ✅ **Advantages:** benefit from **integration** with the host language

### ⚡ DSL Use Case: Halide for Image Processing

**Halide** is a Domain-Specific Language (DSL) for high-performance image processing.

#### ❓ Why does image processing need DSLs?

- 🔗 Image processing is **data-intensive** and **requires high performance**.
- 🗨️ **Traditional solutions** (C++, CUDA, OpenCV) **require manual optimizations**.
- 😞 Optimizing code for **parallelism and memory efficiency** is difficult.

#### ❓ Why Halide?

- ✂ Separates “*what*” is computed from “*how*” it is executed.
- 🔗 Expresses **computations at a high level**, leaving optimizations to the compiler.
- 🌐 **Portable** across CPUs, GPUs, and FPGAs.

### ✂ How Halide Works: Separating Algorithm from Schedule

In Halide, a **key feature is the separation of what** a program computes (*computation/algorithm*) from **how** it executes (*schedule*). This means:

- The **algorithm** specifies **what operations should be performed**. In other words, specifies **what to compute** (like a mathematical formula).
- The **schedule** defines **how those operations should be executed efficiently** on the hardware. In other words, specifies **how to execute the computation** (parallelism, memory layout, vectorization).

Now we see the difference between the traditional approach we have always used and the halide approach:

- **Traditional Approach (C++, CUDA, OpenCV)**. In traditional programming languages (e.g., C++, OpenCV, CUDA), the **algorithm and execution strategy are mixed together**. This means that if we want to **change parallelization or memory access optimizations**, we must **rewrite parts of the algorithm itself**. This makes it *hard to experiment* with different optimizations.

#### ⚠ Problems with the Traditional Approach

1. **If we want to optimize** for vectorization, parallel execution, or memory layout, we **must modify the algorithm itself**.
2. The **same code cannot easily be reused** for different architectures (e.g., CPU, GPU, FPGA).



**Example 2: Problems with the Traditional Approach**

```

1 void box_blur(const Image &in, Image &out) {
2     for (int y = 1; y < in.height() - 1; y++) {
3         for (int x = 1; x < in.width() - 1; x++) {
4             out(x, y) = (
5                 in(x-1, y) + in(x, y) + in(x+1, y)
6             ) / 3;
7         }
8     }
9 }

```

The problem here is that we need to specify the “*what*”, i.e. what operations should be performed, but also “*how*” these operations should be performed.

- **Halide’s Approach: Separate Algorithm from Execution.** Halide splits the computation into two parts:

1. **Computation** (Algorithm) - *What to Compute*
  - Defines the mathematical computation.
  - Remains unchanged across different hardware targets.
2. **Schedule** - *How to Execute Efficiently*
  - Controls memory layout, parallelization, and optimization.
  - Can be changed without modifying the algorithm.

**Example 3: Box Blur in Halide**

The computation part is separated from the scheduling! So a change in the algorithm can be made without affecting the control of the execution. Therefore, we define simply:

1. Computation (Algorithm), stays the same:

```

1 Var x, y;
2 Func blurx, blurry;
3
4 // First pass: horizontal blur
5 blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))
6               / 3;
7
8 // Second pass: vertical blur
9 blurry(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(
10                x, y+1)) / 3;

```

This part **only describes the math**, not how it should run.

2. Schedule, controls execution (can be changed easily):

```

1 blurry.tile(x, y, xi, yi, 256, 32)
2     // Vectorized execution for SIMD
3     .vectorize(xi, 8)
4     // Parallel execution over y-dimension
5     .parallel(y);

```

```
6
7 blurx.compute_at(blury, x) // Compute blurx only
   when needed by blury
8   .vectorize(x, 8);
```

This part **controls execution strategy** but does **not modify the algorithm**. The same algorithm can now **run efficiently on different hardware architectures** just by changing the schedule.

### ✔ Why DSLs Matter for Performance and Productivity: Advantages

- ✔ **Performance Optimization**. A Halide program can be better than **hand-optimization C++ code**. Scheduling decisions affect **parallel execution, memory locality, and vectorization**.
- ✔ **Productivity**. Instead of manually optimizing, **Halide allows rapid exploration of different schedules**. Easier to **port to different architectures** (CPU, GPU, FPGA).

In conclusion, DSLs like Halide **automate low-level optimizations**, enabling *faster* and *more efficient* code for specialized domains.

## 12.4 Scheduling & Performance Optimization in Halide

This section focuses on **how different scheduling strategies in Halide affect performance** and how to **choose the best schedule for different hardware architectures**.

### ? Why Scheduling Matters

Scheduling is the **key to optimizing performance in Halide**.

- A poorly scheduled program can be 10× slower than an optimized one.
- Memory access, cache locality, parallelism, and vectorization all **depend on scheduling**.

As we have already seen, in traditional languages (C++, CUDA, OpenMP), scheduling decisions must be hard-coded into the algorithm. In Halide, the schedule is separate and can be changed without changing the algorithm.

### ? How Scheduling Affects Performance

Different **scheduling strategies** impact how the computation is executed on hardware:

| Scheduling Strategy  | Impact  |
|----------------------|---|
| Serial Execution     | Simple, but slow. No parallelism.                     |
| Parallel Execution   | Uses multiple CPU cores. Good for multi-core CPUs.    |
| Vectorization (SIMD) | Uses wide registers for efficiency (e.g., AVX).       |
| Tiling               | Improves cache locality by processing data in chunks. |
| Compute-at           | Controls when intermediate results are computed.      |
| Store-at             | Controls where intermediate results are stored.       |


### ≡ Scheduling Strategies in Halide

Let's take the Box Blur Algorithm as an example. The **Box Blur** is a simple **image processing technique used for smoothing or blurring an image**. It performs two main steps:

1. **Each pixel** in the output image is computed as the **average of its neighboring pixels**.
2. It applies a **moving average filter** over a **small window** (e.g.,  $3 \times 3$  or  $5 \times 5$  pixels).

It is used to reduce noise in images, to create a smooth, blurry effect and, above all, because it is fast and efficient, since it involves only two operations: adding and dividing.

However, we will now analyze the algorithm implemented in Halide using different strategies:

 **Default Serial Execution (Slow).** Without scheduling, Halide will execute in **serial order** (one pixel at a time).

---




```


1 Var x, y;
2 Func blurx, blurry;
3
4 blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y)) / 3;
5 blurry(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1)) /
    3;

```

---

### Problems

-  No parallelism or vectorization.
-  Poor memory access patterns.
-  Slow execution.

 **Parallel Execution.** We can **add parallelism** to use multiple CPU cores:

---



```


1 blurry.parallel(y);

```

---

### Advantages

-  Halide automatically splits the work across CPU cores.
-  Useful for multi-threaded execution on CPUs.

 **Vectorization (SIMD).** Modern CPUs support **SIMD instructions** (e.g., AVX, NEON) to process multiple pixels at once:

---



```


1 blurry.vectorize(x, 8);

```

---

### Advantages

-  Uses **SIMD registers** for faster execution.
-  Works best for **data-parallel workloads** like image processing.

 **Tiling for Better Cache Performance.** Instead of processing the whole image at once, we **divide it into smaller tiles**:

---


```

1 blurry.tile(x, y, xi, yi, 256, 32);

```

---

### Advantages

-  Each tile fits better in cache, **reducing memory latency**.

✓ Improves locality of reference (less cache thrashing).

🔧 **Optimized Schedule: Combining Techniques.** We can combine multiple scheduling strategies for maximum performance:

```

1 // Process in 256x32 tiles
2 blurry.tile(x, y, xi, yi, 256, 32)
3   // Vectorized execution
4   .vectorize(xi, 8)
5   // Parallel execution across CPU cores
6   .parallel(y);
7
8 // Compute blurx only when needed
9 blurx.compute_at(blurry, x)
10  .vectorize(x, 8);

```

#### ✓ Advantages

- ✓ Breaks image into tiles for cache efficiency.
- ✓ Uses SIMD vectorization for fast execution.
- ✓ Runs in parallel on multiple CPU cores.
- ✓ Intermediate results (blurx) are computed only when needed.

#### ⚖️ Trade-offs in Scheduling

But *how do we choose the right scheduling?* Well, we need to find a good trade-off. Different scheduling choices affect **performance trade-offs**:

| Scheduling Strategy  | Performance Impact                                     |
|----------------------|--|
| Parallel Execution   | Increases throughput, uses multiple cores.             |
| Vectorization (SIMD) | Improves performance on CPUs/GPUs.                     |
| Tiling               | Improves cache locality, reduces memory overhead.      |
| Compute-at           | Avoids redundant computations.                         |
| Store-at             | Reduces memory footprint but increases re-computation. |

Table 12: Performance trade-offs in Scheduling.

#### ✓ Conclusion

In conclusion, the key takeaways are:

- **Scheduling** is the key to performance in Halide.
- **Parallel execution, vectorization, and tiling** significantly improve performance.
- **Halide's flexibility** allows quick experimentation with different schedules.
- The **right schedule depends on hardware constraints** (CPU, GPU, FPGA).